



Pure Language and Library Documentation

Release 0.46

Albert Gräf (Editor)

November 14, 2010

Contents

The Pure Manual	3
1 Introduction	3
2 Invoking the Pure Interpreter	4
3 Pure Overview	11
4 Rule Syntax	23
5 Examples	29
6 Macros	39
7 Declarations	44
8 Exception Handling	57
9 C Interface	59
10 Standard Library	71
11 Interactive Usage	72
12 Batch Compilation	87
13 Caveats and Notes	94
14 Copying	110
15 Author	111
16 See Also	111
 Pure Library Manual	 113
1 Prelude	113
2 Mathematical Functions	146
3 Container Types	150
4 System Interface	161
 pure-doc	 179
1 Copying	179
2 Installation	179
3 Usage	180
4 Literate Programming	181
5 Hyperlink Targets and Index Generation	182
6 Generating and Installing Local Documentation	183
7 Formatting Tips	185
 pure-ffi	 187

1	Copying	187
2	Installation	187
3	Usage	188
4	TODO	189
pure-gen: Pure interface generator		191
1	Synopsis	191
2	Options	191
3	Description	193
4	Filtering	194
5	Name Mangling	195
6	Generating C Code	195
7	Dealing with C Structs	197
8	Notes	199
9	Example	200
10	License	200
11	Authors	200
12	See Also	200
pure-readline		201
pure-sockets: Pure Sockets Interface		203
1	Installation	203
2	Usage	203
3	Example	206
Gnumeric/Pure: A Pure Plugin for Gnumeric		209
1	Introduction	209
2	Installation	210
3	Setup	211
4	Basic Usage	212
5	Interactive Pure Shell	213
6	Defining Your Own Functions	214
7	Gnumeric/Pure Interface	215
8	Advanced Features	218
9	Copying	226
Pure-GLPK - GLPK interface for the Pure programming language		229
1	Installation	229
2	Error Handling	230
3	Further Information and Examples	231
4	Interface description	231
5	Descriptions of interface functions	231
Gnuplot bindings		321
1	Copying	321
2	Introduction	321
3	Function Reference	321

pure-gsl - GNU Scientific Library Interface for Pure	325
1 Polynomials	326
2 Special Functions	328
3 Matrices	343
4 Least-Squares Fitting	347
5 Statistics	349
6 Random Number Distributions	353
7 Sorting	362
pure-octave	365
1 Introduction	365
2 Copying	365
3 Installation	365
4 Basic Usage	366
5 Direct Function Calls	367
6 Data Conversions	368
7 Calling Back Into Pure	369
8 Caveats and Notes	370
Pure-Rational - Rational number library for the Pure programming language	371
1 Copying	372
2 Installation	372
3 Introduction	372
4 The Rational Type	373
5 Arithmetic	376
6 Mathematical Functions	378
7 Special Rational Functions	380
8 $\mathbf{Q} \rightarrow \mathbf{Z}$ — Rounding	384
9 Rounding to Multiples	385
10 $\mathbf{Q} \rightarrow \mathbf{R}$ — Conversion / Casting	387
11 $\mathbf{R} \rightarrow \mathbf{Q}$ — Approximation	387
12 Decomposition	394
13 Continued Fractions	394
14 Rational Complex Numbers	395
15 String Formatting and Evaluation	401
16 $\mathbf{Q} \leftrightarrow$ Fraction String (“i + n/d”)	404
17 $\mathbf{Q} \leftrightarrow$ Recurring Numeral Expansion String (“I.FR”)	406
18 $\mathbf{Q} \leftrightarrow$ Numeral Expansion String (“I.F $\times 10^E$ ”)	408
19 Numeral String $\rightarrow \mathbf{Q}$ — Approximation	412
Pure-CSV - Comma Separated Value Interface for the Pure Programming Language	415
1 Installation	415
2 Usage	415
pure-fastcgi: FastCGI module for Pure	423
1 Copying	423
2 Installation	423

3	Usage	424
Pure-ODBC - ODBC interface for the Pure programming language		427
1	Copying	427
2	Installation	428
3	Opening and Closing a Data Source	428
4	Getting Information about a Data Source	430
5	Executing SQL Queries	431
6	Low-Level Operations	433
7	Lazy Processing	434
8	Error Handling	435
9	Caveats and Bugs	435
10	Further Information and Examples	436
Pure-Sql3		437
1	Introduction	437
2	Copying	439
3	Installation	439
4	Data Structure	439
5	Core Database Operations	440
6	Advanced Features	448
7	Threading Modes	452
Pure-XML - XML/XSLT interface		455
1	Copying	455
2	Installation	455
3	Usage	456
4	Data Structure	456
5	Operations	459
pure-g2		467
Pure OpenGL Bindings		469
1	Copying	469
2	Installation	470
3	Using the GL Bindings	470
4	Regenerating the Bindings	471
Pure GTK+ Bindings		473
1	Copying	473
2	Installation	473
3	Usage	474
pure-tk		475
1	Introduction	475
2	Copying	476
3	Installation	476
4	Basic Usage	476

5	Callbacks	477
6	The Main Loop	477
7	Accessing Tcl Variables	477
8	Conversions Between Pure and Tcl Values	478
9	Tips and Tricks	479
faust2pd: Pd Patch Generator for Faust		483
1	Copying	483
2	Requirements	484
3	Installation	484
4	Quickstart	485
5	Control Interface	486
6	Examples	487
7	Wrapping DSPs with faust2pd	488
8	Conclusion	490
9	Appendix: faustxml	491
pd-pure: Pd loader for Pure scripts		493
1	Copying	493
2	Installation	493
3	Usage	494
4	Advanced Features	497
5	Programming Interface	499
6	Interactive Facilities	500
pure-audio		503
1	Installation	503
2	License	504
pure-faust: Run Faust programs in Pure scripts		505
1	Copying	505
2	Installation	505
3	Usage	506
4	Acknowledgements	508
pure-liblo		509
1	Copying	509
2	Description	509
pure-midi		511
1	Installation	511
2	License	512
Index		513

This manual collects all of Pure's online documentation: the Pure Manual which covers the Pure language and the operation of the Pure interpreter; the Pure Library Manual which describes the standard library modules included in the distribution of the Pure interpreter; and all available documentation for the various addon modules which can be downloaded as separate packages from the [Pure website](#).

Most of the Pure documentation is distributed under the [GNU Free Documentation License](#). The authors of the current edition are listed below. (This just lists the primary section authors in alphabetical order; please check the different parts of this manual for additional authorship and licensing information.)

- Albert Gräf (*The Pure Manual; Pure Library Manual; various addon manuals*)
- Rob Hubbard (*Pure-Rational - Rational number library for the Pure programming language*)
- Kay-Uwe Kirstein (*Gnuplot bindings*)
- Eddie Rucker (*Pure-CSV - Comma Separated Value Interface for the Pure Programming Language; pure-gsl - GNU Scientific Library Interface for Pure*)
- Jiri Spitz (*Pure-GLPK - GLPK interface for the Pure programming language*)
- Peter Summerland (*Pure-Sql3*)

The Pure programming system is free and open source software. The interpreter runtime, the standard library and most of the addon modules are distributed under the [GNU Lesser General Public License](#) or the 3-clause [BSD License](#) which allow for commercial applications. Some parts of the system also use the [GNU General Public License](#) (typically because they interface to other GPL'ed software such as Gnumeric, GSL and Octave). Details about authorship and license conditions can be found in the sources or in the various manual sections.

For more information, discussions, feedback, questions, suggestions etc. please see:

- Pure website: <http://pure-lang.googlecode.com>
- Pure mailing list: <http://groups.google.com/group/pure-lang>

The Pure Manual

Version 0.46, November 14, 2010

Albert Gräf <Dr.Graef@t-online.de>

Copyright (c) 2009-2010 by Albert Gräf. This document is available under the [GNU Free Documentation License](#). Also see the [Copying](#) section for licensing information of the software.

This manual describes the Pure programming language and how to invoke the Pure interpreter program. To read the manual inside the interpreter, just type `help` at the command prompt. See the [Online Help](#) section for details.

There is a companion to this manual, the *Pure Library Manual* which contains the description of the standard library operations. More information about Pure can be found under the following URLs:

- Pure website: <http://pure-lang.googlecode.com>
- Pure mailing list: <http://groups.google.com/group/pure-lang>

1 Introduction

Pure is a modern-style functional programming language based on term rewriting. Pure programs are basically collections of equational rules used to evaluate expressions in a symbolic fashion by reducing them to normal form. An overview of the language can be found in the [Pure Overview](#) section below, and subsequent sections discuss most language features in detail.

The Pure interpreter has an [LLVM](#) backend which JIT-compiles Pure programs to machine code, hence programs run blazingly fast and interfacing to C modules is easy, while the interpreter still provides a convenient, fully interactive environment for running Pure scripts and evaluating expressions. You can also compile your scripts to standalone executables if you prefer that.

Pure programs (a.k.a. scripts) are just ordinary text files containing Pure code. They must be encoded in UTF-8 (which subsumes 7 bit ASCII), other encodings such as Latin-1 are not supported. A bunch of syntax highlighting files and programming modes for various popular text editors are included in the Pure sources. There's no difference between the

Pure programming language and the input language accepted by the interpreter, except that the interpreter also understands some special commands when running in interactive mode; see the [Interactive Usage](#) section for details.

(In case you're wondering, the name "Pure" actually refers to the adjective. But you can also write it as "PURE" and take this as a recursive acronym for the "Pure Universal Rewriting Engine".)

2 Invoking the Pure Interpreter

The Pure interpreter is invoked as follows:

```
pure [options ...] [script ...] [-- args ...]
pure [options ...] -x script [args ...]
```

2.1 Options

The interpreter accepts various options which are described in more detail below:

- c** Batch compilation.
- ctags, --etags** Create a tags file in ctags (vi) or etags (emacs) format.
- eager-jit** Enable eager JIT compilation. This requires LLVM 2.7 or later, otherwise this flag will be ignored.
- fPIC, -fpic** Create position-independent code (batch compilation).
- g** Enable symbolic debugging.
- help, -h** Print help message and exit.
- i** Force interactive mode (read commands from stdin).
- I directory** Add a directory to be searched for included source scripts.
- L directory** Add a directory to be searched for dynamic libraries.
- l libname** Library to be linked in batch compilation.
- noediting** Disable command-line editing.
- noprelude, -n** Do not load the prelude.
- norc** Do not run the interactive startup files.
- o filename** Output filename for batch compilation.
- q** Quiet startup (suppresses sign-on message in interactive mode).
- T filename** Tags file to be written by --ctags or --etags.
- u** Do not strip unused functions in batch compilation.

- v[level]** Set verbosity level. See below for details.
- version** Print version information and exit.
- w** Enable compiler warnings.
- x** Execute script with given command line arguments.
- Stop option processing and pass the remaining command line arguments in the `argv` variable.

(Besides these, the interpreter also understands a number of other command line switches for setting various code generation options; please see [Code Generation Options](#) below for details.)

2.2 Overview of Operation

If any source scripts are specified on the command line, they are loaded and executed, after which the interpreter exits. Otherwise the interpreter enters the interactive read-eval-print loop, see [Running Interactively](#) below. You can also use the `-i` option to enter the interactive loop (continue reading from `stdin`) even after processing some source scripts.

Options and source files are processed in the order in which they are given on the command line. Processing of options and source files ends when either the `-` or the `-x` option is encountered. The `-x` option must be followed by the name of a script to be executed, which becomes the “main script” of the application. In either case, any remaining parameters are passed to the executing script by means of the global `argc` and `argv` variables, denoting the number of arguments and the list of the actual parameter strings, respectively. In the case of `-x` this also includes the script name as `argv!0`. The `-x` option is useful, in particular, to turn Pure scripts into executable programs by including a “shebang” like the following as the first line in your main script. (This trick only works with Unix shells, though.)

```
#!/usr/local/bin/pure -x
```

On startup, the interpreter also defines the `version` variable, which is set to the version string of the Pure interpreter, and the `sysinfo` variable, which provides a string identifying the host system. These are useful if parts of your script depend on the particular version of the interpreter and the system it runs on. (Moreover, Pure 0.21 and later also define the variable `compiling` which indicates whether the program is executed in a batch compilation, see [Compiling Scripts](#) below.)

If available, the prelude script `prelude.pure` is loaded by the interpreter prior to any other definitions, unless the `-n` or `-noprelude` option is specified. The prelude is searched for in the directory specified with the `PURELIB` environment variable. If the `PURELIB` variable is not set, a system-specific default is used. Relative pathnames of other source scripts specified on the command line are interpreted relative to the current working directory. In addition, the executed program may load other scripts and libraries via a `using` declaration in the source, which are searched for in a number of locations, including the directories named with the `-I` and `-L` options; see the [Declarations](#) and [C Interface](#) sections for details.

2.3 Compiling Scripts

The interpreter compiles scripts, as well as definitions that you enter interactively, automatically. This is done in an incremental fashion, as the code is needed, and is therefore known as JIT (**just in time**) compilation. Thus the interpreter never really “interprets” the source program or some intermediate representation, it just acts as a frontend to the compiler, taking care of compiling source code to native machine code before it gets executed.

Pure’s LLVM backend does “lazy JIT compilation” by default, meaning that each function (global or local) is compiled no sooner than it is run for the first time. With the `-eager-jit` option, however, it will also compile all other (global or local) functions that may be called by the compiled function. (The `PURE_EAGER_JIT` environment variable, when set to any value, has the same effect, so that you do not have to specify the `-eager-jit` option each time you run the interpreter.) Eager JIT compilation may be more efficient in some cases (since bigger chunks of compilation work can be done in one go) and less efficient in others (e.g., eager JITing may compile large chunks of code which aren’t actually called later, except in rare circumstances).

Note that the eager JIT mode is only available with LLVM 2.7 or later; otherwise this option will be ignored.

It is also possible to compile your scripts to native code beforehand, using the `-c` batch compilation option. This option forces the interpreter to non-interactive mode (unless `-i` is specified as well, which overrides `-c`). Any scripts specified on the command line are then executed as usual, but after execution the interpreter takes a snapshot of the program and compiles it to one of several supported output formats, LLVM assembler (.ll) or bitcode (.bc), native assembler (.s) or object (.o), or a native executable, depending on the output filename specified with `-o`. If the output filename ends in the .ll extension, an LLVM assembler file is created which can then be processed with the LLVM toolchain. If the output filename is just ‘-’, the assembler file is written to standard output, which is useful if you want to pass the generated code to the LLVM tools in a pipeline. If the output filename ends in the .bc extension, an LLVM bitcode file is created instead.

The .ll and .bc formats are supported natively by the Pure interpreter, no external tools are required to generate these. If the target is an .s, .o or executable file, the Pure interpreter creates a temporary bitcode file on which it invokes the LLVM tools `opt` and `llc` to create a native assembler file, and then uses `gcc` to assemble and link the resulting program (if requested). You can also specify additional libraries to be linked into the executable with the `-l` option. If the output filename is omitted, it defaults to `a.out` (`a.exe` on Windows).

The `-c` option provides a convenient way to quickly turn a Pure script into a standalone executable which can be invoked directly from the shell. One advantage of compiling your script is that this eliminates the JIT compilation time and thus considerably reduces the startup time of the program. Another reason to prefer a standalone executable is that it lets you deploy the program on systems without a full Pure installation (usually only the runtime library is required on the target system). On the other hand, compiled scripts also have some limitations, mostly concerning the use of the built-in `eval` function. Please see the [Batch Compilation](#) section for details.

The `-v64` (or `-v0100`) verbosity option can be used to have the interpreter print the commands it executes during compilation, see [Verbosity and Debugging Options](#) below. When creating an object file, this also prints the suggested linker command (including all the dynamic modules loaded by the script, which also have to be linked in to create a working executable), to which you only have to add the options describing the desired output file.

2.4 Tagging Scripts

Pure programs often have declarations and definitions of global symbols scattered out over many different source files. The `-ctags` and `-etags` options let you create a `tags` file which allows you to quickly locate these items in text editors such as `vi` and `emacs` which support this feature.

If `-ctags` or `-etags` is specified, the interpreter enters a special mode in which it only parses source files without executing them and collects information about the locations of global symbol declarations and definitions. The collected information is then written to a tags file in the `ctags` or `etags` format used by `vi` and `emacs`, respectively. The desired name of the tags file can be specified with the `-T` option; it defaults to `tags` for `-ctags` and `TAGS` for `-etags` (which matches the default tags file names used by `vi` and `emacs`, respectively).

The tags file contains information about the global constant, variable, macro, function and operator symbols of all scripts specified on the command line, as well as the prelude and other scripts included via a `using` clause. Tagged scripts which are located in the same directory as the tags file (or, recursively, in one of its subdirectories) are specified using relative pathnames, while scripts outside this hierarchy (such as included scripts from the standard library) are denoted with absolute pathnames. This scheme makes it possible to move an entire directory together with its tags file and have the tags information still work in the new location.

2.5 Running Interactively

If the interpreter runs in interactive mode, it repeatedly prompts you for input (which may be any legal Pure code or some special interpreter commands provided for interactive usage), and prints computed results. This is also known as the **read-eval-print** loop and is described in much more detail in the [Interactive Usage](#) section. To exit the interpreter, just type the quit command or the end-of-file character (`^D` on Unix) at the beginning of the command line.

The interpreter may also source a few additional interactive startup files immediately before entering the interactive loop, unless the `-norc` option is specified. First `.purerc` in the user's home directory is read, then `.purerc` in the current working directory. These are ordinary Pure scripts which can be used to provide additional definitions for interactive usage. Finally, a `.pure` file in the current directory (containing a dump from a previous interactive session) is loaded if it is present.

When the interpreter is in interactive mode and reads from a tty, unless the `-noediting` option is specified, commands are usually read using `readline(3)` or some compatible replacement,

providing completion for all commands listed under [Interactive Usage](#), as well as for symbols defined in the running program. When exiting the interpreter, the command history is stored in `~/.pure_history`, from where it is restored the next time you run the interpreter.

As of Pure 0.22, the interpreter also provides a simple source level debugger when run in interactive mode, see [Debugging](#) for details. To enable the interactive debugger, you need to specify the `-g` option when invoking the interpreter. This option causes your script to run *much* slower and also disables tail call optimization, so you should only use this option if you want to run the debugger.

2.6 Verbosity and Debugging Options

The `-v` option is useful for debugging the interpreter, or if you are interested in the code your program gets compiled to. The level argument is optional; it defaults to 1. Seven different levels are implemented at this time (one more bit is reserved for future extensions). For most purposes, only the first two levels will be useful for the average Pure programmer; the remaining levels are most likely to be used by the Pure interpreter developers.

1 (0x1, 001) denotes echoing of parsed definitions and expressions.

2 (0x2, 002) adds special annotations concerning local bindings (de Bruijn indices, subterm paths; this can be helpful to debug tricky variable binding issues).

4 (0x4, 004) adds descriptions of the matching automata for the left-hand sides of equations (you probably want to see this only when working on the guts of the interpreter).

8 (0x8, 010) dumps the “real” output code (LLVM assembler, which is as close to the native machine code for your program as it gets; you definitely don’t want to see this unless you have to inspect the generated code for bugs or performance issues).

16 (0x10, 020) adds debugging messages from the bison(1) parser; useful for debugging the parser.

32 (0x20, 040) adds debugging messages from the flex(1) lexer; useful for debugging the lexer.

64 (0x40, 0100) turns on verbose batch compilation; this is useful if you want to see exactly which commands get executed during batch compilation (`-c`).

These values can be or’ed together, and, for convenience, can be specified in either decimal, hexadecimal or octal. Thus `0xff` or `0777` always gives you full debugging output (which isn’t likely to be used by anyone but the Pure developers). Some useful flag combinations for experts are (in octal) `007` (echo definitions along with de Bruijn indices and matching automata), `011` (definitions and assembler code) and `021` (parser debugging output along with parsed definitions).

Note that the `-v` option is only applied after the prelude has been loaded. If you want to debug the prelude, use the `-n` option and specify the `prelude.pure` file explicitly on the command line. Verbose output is also suppressed for modules imported through a `using` clause. As a remedy, you can use the interactive `show` command (see the [Interactive Usage](#) section) to list definitions along with additional debugging information.

The `-w` option enables some additional warnings which are useful to check your scripts for possible errors. Right now it will report implicit declarations of function symbols which might indicate missing or mistyped symbols that need to be fixed, see [Symbol Lookup and Creation](#) for details.

2.7 Code Generation Options

Besides the options listed above, the interpreter also understands some additional command line switches and corresponding environment variables to control various code generation options. The options take the form `--opt` and `--noopt`, respectively, where `opt` denotes the option name (see below for a list of supported options). By default, these options are all enabled; `--noopt` disables the option, `--opt` reenables it. In addition, for each option `opt` there is also a corresponding environment variable `PURE_NOOPT` (with the option name in uppercase) which, when set, disables the option by default. (Setting this variable to any value will do, the interpreter only checks whether the variable exists in the environment.)

For instance, the `checks` option controls stack and signal checks. Thus `--nochecks` on the command line disables the option, and setting the `PURE_NOCHECKS` environment variable makes this the default, in which case you can use `--checks` on the command line to reenable the option.

Each code generation option can also be used as a *pragma* (compiler directive) in source code so that you can control it on a per-rule basis. The pragma must be on a line by itself, starting in column 1, and takes the following form (using `--nochecks` as an example):

```
#! --nochecks // line-oriented comment may go here
```

Currently, the following code generation options are recognized:

- checks, --nochecks** Enable or disable various extra stack and signal checks. By default, the interpreter checks for stack overflows (if the `PURE_STACK` environment variable is set) and pending signals on entry to every function, see [Stack Size and Tail Recursion](#) and [Handling of Asynchronous Signals](#) for details. This is needed to catch these conditions in a reliable way, so we recommend to leave this enabled. However, these checks also make programs run a little slower (typically some 5%, YMMV). If performance is critical then you can disable the checks with the `--nochecks` option. (Even then, a minimal amount of checking will be done, usually on entry to every global function.)
- const, --noconst** Enable or disable the precomputing of constant values in batch compilation (cf. [Compiling Scripts](#)). If enabled (which is the default), the values of constants in `const` definitions are precomputed at compile time (if possible) and then stored in the generated executable. This usually yields faster startup times but bigger executables. You can disable this option with `--noconst` to get smaller executables at the expense of slower startup times. Please see the [Batch Compilation](#) section for an example.
- fold, --nofold** Enable or disable constant folding in the compiler frontend. This means that constant expressions involving `int` and `double` values and the usual arithmetic and logical operations on these are precomputed at compile time. (This is mostly for

cosmetic purposes; the LLVM backend will perform this optimization anyway when generating machine code.) For instance:

```
> foo x = 2*3*x;
> show foo
foo x = 6*x;
```

Disabling constant folding in the frontend causes constant expressions to be shown as you entered them:

```
> #! --nofold
> bar x = 2*3*x;
> show bar
bar x = 2*3*x;
```

--tc, --notc Enable or disable tail call optimization (TCO). TCO is needed to make tail-recursive functions execute in constant stack space, so we recommend to leave this enabled. However, at the time of this writing LLVM's TCO support is still bug-ridden on some platforms, so the `--notc` option allows you to disable it. (Note that TCO can also be disabled when compiling the Pure interpreter, in which case these options have no effect; see the installation documentation for details.)

Besides these, there are the following special pragmas affecting the code generation of some given function, which is specified in the pragma. These pragmas can only be used in source code, there are no command line options for them.

--eager fun Instruct the interpreter to JIT-compile the given function eagerly. This means that native code will be created for the function, as well as all other (global or local) functions that may be called by the compiled function, as soon as the function gets recompiled. This avoids the hiccups you get when a function is compiled on the fly if it is run for the first time, which is particularly useful for functions which are to be run in realtime, such as in audio and MIDI applications. Please note that, in difference to the `-eager-jit` option, this feature is available for all LLVM versions (it doesn't require LLVM 2.7 or later).

--required fun Inform the batch compiler (cf. [Compiling Scripts](#)) that the given function symbol `fun` should never be stripped from the program. This is useful, e.g., if a function is never called explicitly but only through `eval`. Adding a `--required` pragma for the function then makes sure that the function is always linked into the program. Please see the [Batch Compilation](#) section for an example.

2.8 Startup Files

The interpreter may source various files during its startup. These are:

`~/.pure_history` Interactive command history.

`~/.purerc`, `.purerc`, `.pure` Interactive startup files. The latter is usually a dump from a previous interactive session.

prelude.pure Standard prelude. If available, this script is loaded before any other definitions, unless `-n` was specified.

2.9 Environment

Various aspects of the interpreter can be configured through the following shell environment variables:

BROWSER If the `PURE_HELP` variable is not set (see below), this specifies a colon-separated list of browsers to try for reading the online documentation. See <http://catb.org/~esr/BROWSER/>.

PURELIB Directory to search for library scripts, including the prelude. If `PURELIB` is not set, it defaults to some location specified at installation time.

PURE_EAGER_JIT Enable eager JIT compilation (same as `--eager-jit`), see [Compiling Scripts](#) for details.

PURE_HELP Command used to browse the Pure manual. This must be a browser capable of displaying html files. Default is `w3m(1)`.

PURE_INCLUDE Additional directories (in colon-separated format) to be searched for included scripts.

PURE_LIBRARY Additional directories (in colon-separated format) to be searched for dynamic libraries.

PURE_MORE Shell command to be used for paging through output of the `show` command, when the interpreter runs in interactive mode. `PURE_LESS` is the same for evaluation results printed by the interpreter.

PURE_PS Command prompt used in the interactive command loop ("`>`" by default).

PURE_STACK Maximum stack size in kilobytes (default: 0 = unlimited).

Besides these, the interpreter also understands a number of other environment variables for setting various code generation options (see [Code Generation Options](#) above) and commands to invoke different LLVM compilers on inline code (see [Inline Code](#)).

3 Pure Overview

Pure is a fairly simple yet powerful language. Programs are basically collections of rewriting rules and expressions to be evaluated. For convenience, it is also possible to define global variables and constants, and for advanced uses Pure offers macro functions as a kind of preprocessing facility. These are all described below and in the following sections.

Here's a first example which demonstrates how to define a simple recursive function in Pure, entered interactively in the interpreter (note that the `'>'` symbol at the beginning of each input line is the interpreter's default command prompt):

```
> // my first Pure example
> fact 0 = 1;
> fact n::int = n*fact (n-1) if n>0;
> let x = fact 10; x;
3628800
```

3.1 Lexical Matters

Pure is a free-format language; whitespace is insignificant, except if it serves to delimit other symbols. Hence, as shown above, definitions and expressions at the toplevel have to be terminated with a semicolon, even in interactive mode.

Comments have the same syntax as in C++ (using `//` for line-oriented and `/* ... */` for multiline comments; the latter may not be nested). Lines beginning with `#!` are treated as comments, too; as already discussed above, on Unix-like systems this allows you to add a “shebang” to your main script in order to turn it into an executable program.

There are a few reserved keywords which cannot be used as identifiers:

```
case const def else end extern if infix infixl infixr let namespace
nonfix of otherwise outfix postfix prefix private public then using
when with
```

The customary notations for identifiers, numbers and strings are all provided. In addition, Pure also allows you to define your own operator symbols. Pure fully supports Unicode, so that you can write your programs in almost any language and make good use of the special symbols in the Unicode character set, provided that you encode your scripts in UTF-8. To keep this simple, besides the ASCII punctuation characters, Pure also considers the following code points in the Unicode repertoire as punctuation: U+00A1 through U+00BF, U+00D7, U+00F7, and U+20D0 through U+2BFF. This comprises the special symbols in the Latin-1 repertoire, as well as the Combining Diacritical Marks for Symbols, Letterlike Symbols, Number Forms, Arrows, Mathematical Symbols, Miscellaneous Technical Symbols, Control Pictures, OCR, Enclosed Alphanumerics, Box Drawing, Blocks, Geometric Shapes, Miscellaneous Symbols, Dingbats, Miscellaneous Mathematical Symbols A, Supplemental Arrows A, Supplemental Arrows B, Miscellaneous Mathematical Symbols B, Supplemental Mathematical Operators, and Miscellaneous Symbols and Arrows. This should cover almost everything you’d ever want to use in an operator symbol. All other extended Unicode characters are effectively treated as “letters” which can be used as identifier constituents.

3.2 Definitions and Expression Evaluation

On the surface, Pure is quite similar to other modern functional languages like [Haskell](#) and [ML](#). But under the hood it is a much more dynamic language, more akin to Lisp. In particular, Pure is dynamically typed, so functions can be fully polymorphic and you can add to the definition of an existing function at any time. For instance, we can extend the example above to make the `fact` function work with floating point numbers, too:

```
> fact 0.0 = 1.0;
> fact n::double = n*fact (n-1) if n>0;
> fact 10.0;
3628800.0
> fact 10;
3628800
```

Note the `n::double` construct on the left-hand side of the second equation, which means that the equation is only to be applied for (double precision) floating point values `n`. This construct is also called a “type tag” in Pure parlance, which is actually a simple form of pattern matching (see below). Similarly, our previous definition at the beginning of this section employed the `int` tag to indicate that the `n` parameter is an integer value. The `int` and `double` types are built into the Pure language, but it is also possible to introduce your own type tags for user-defined data structures. This will be explained in more detail under [Type Tags](#) in the [Rule Syntax](#) section below.

Expressions are generally evaluated from left to right, innermost expressions first, i.e., using call by value semantics. Pure also has a few built-in special forms (most notably, conditional expressions, the short-circuit logical connectives `&&` and `||`, the sequencing operator `$$`, the lazy evaluation operator `&`, and the quote) which take some or all of their arguments unevaluated, using call by name.

Like in Haskell and ML, functions are often defined by pattern matching, i.e., the left-hand side of a definition is compared to the target expression, binding the variables in the pattern to their actual values accordingly:

```
> foo (bar x) = x-1;
> foo (bar 99);
98
```

Due to its term rewriting semantics, Pure goes beyond most other functional languages in that it can do symbolic evaluations just as well as “normal” computations:

```
> square x = x*x;
> square 4;
16
> square (a+b);
(a+b)*(a+b)
```

Leaving aside the built-in support for some common data structures such as numbers and strings, all the Pure interpreter really does is evaluate expressions in a symbolic fashion, rewriting expressions using the equations supplied by the programmer, until no more equations are applicable. The result of this process is called a **normal form** which represents the “value” of the original expression. Keeping with the tradition of term rewriting, there’s no distinction between “defined” and “constructor” function symbols in Pure. Consequently, any function symbol or operator can be used *anywhere* on the left-hand side of an equation, and may act as a constructor symbol if it happens to occur in a normal form term. This enables you to work with algebraic rules like associativity and distributivity in a direct fashion:

```
> (x+y)*z = x*z+y*z; x*(y+z) = x*y+x*z;  
> x*(y*z) = (x*y)*z; x+(y+z) = (x+y)+z;  
> square (a+b);  
a*a+a*b+b*a+b*b
```

Note that, in contrast, languages like Haskell and ML always enforce the so-called “constructor discipline”, which stipulates that only pure constructor symbols (without any defining equations) may occur as a subterm on the left-hand side of a definition. Thus equational definitions like the above are forbidden in these languages. In Pure such definitions are just normal business, which makes the language directly usable as a tool for symbolic algebra.

3.3 Parameters in Equations

Taking a look at the above examples, you might have been wondering how the Pure interpreter figures out what the parameters (a.k.a. “variables”) in an equation are. This is quite obvious in rules involving just variables and special operator symbols, such as $(x+y)*z = x*z+y*z$. However, what about an equation like `foo (foo bar) = bar`? Since most of the time we don’t declare any symbols in Pure, how does the interpreter know that `foo` is a literal function symbol here, while `bar` is a variable?

The answer is that the interpreter considers the different positions in the left-hand side expression of an equation. Basically, a Pure expression is just a tree formed by applying expressions to other expressions, with the atomic subexpressions like numbers and symbols at the leaves of the tree. (This is even true for infix expressions like `x+y`, since in Pure these are always equivalent to a function application of the form `(+) x y` which has the atomic subterms `(+)`, `x` and `y` at its leaves.)

Now the interpreter divides the leaves of the expression tree into “head” (or “function”) and “parameter” (or “variable”) positions based on which leaves are leftmost in a function application or not. Thus, in an expression like `f x y z`, `f` is in the head or function position, while `x`, `y` and `z` are in parameter or variable positions. (Note that in an infix expression like `x+y`, `(+)` is the head symbol, not `x`, as the expression is really parsed as `(+) x y`, see above.)

Identifiers in head positions are taken as literal function symbols by the interpreter, while identifiers in variable positions denote, well, variables. We also refer to this convention as the **head = function rule**. It is quite intuitive and lets us get away without declaring the variables in equations. (There are some corner cases not covered here, however. In particular, Pure allows you to declare special constant symbols, if you need a symbol to be recognized as a literal even if it occurs in a variable position. This is done by means of a `nonfix` declaration, see [Symbol Declarations](#) for details.)

3.4 Expression Syntax

The Pure language provides built-in support for machine integers (32 bit), bigints (implemented using [GMP](#)), floating point values (double precision IEEE 754), character strings (UTF-8 encoded) and generic C pointers. The latter don’t have a syntactic representation in Pure, except that the predefined constant `NULL` may be used to denote a generic null pointer;

other pointer values need to be created with external C functions. Truth values are encoded as machine integers (as you might expect, zero denotes *false* and any non-zero value *true*). Pure also provides some built-in support for lists and matrices, although most of the corresponding operations are actually defined in the prelude.

Expressions consist of the following elements:

Numbers: `4711`, `4711L`, `1.2e-3` The usual C notations for integers (decimal: `1000`, hexadecimal: `0x3e8`, octal: `01750`) and floating point values are all provided. Integers can also be denoted in base 2 by using the `0b` or `0B` prefix: `0b1111101000`. Integer constants that are too large to fit into machine integers are promoted to bigints automatically. Moreover, integer literals immediately followed by the uppercase letter `L` are always interpreted as bigint constants, even if they fit into machine integers. This notation is also used when printing bigint constants, to distinguish them from machine integers.

Strings: `"Hello, world!\n"` String constants are double-quoted and terminated with a null character, like in C. In difference to C, strings are always encoded in UTF-8, and character escapes in Pure strings have a more flexible syntax (borrowed from the author's Q language) which provides notations to specify any Unicode character. In particular, the notation `\n`, where `n` is an integer literal written in decimal (no prefix), hexadecimal (`0x` prefix), octal (`0` prefix) or binary (`0b` prefix) notation, denotes the Unicode character (code point) `#n`. Since these escapes may consist of a varying number of digits, parentheses may be used for disambiguation purposes; thus, e.g. `"\ (123)4"` denotes character `#123` followed by the ASCII character `4`. The usual C-like escapes for special non-printable characters such as `\n` are also supported. Moreover, you can use symbolic character escapes of the form `\&name;`, where `name` is any of the XML single character entity names specified in the [XML Entity definitions for Characters](#). Thus, e.g., `"\©"` denotes the copyright character (code point `0x000A9`).

Function and variable symbols: `foo`, `foo_bar`, `BAR`, `foo:bar` These consist of the usual sequence of letters (including the underscore) and digits, starting with a letter. Case is significant, thus `foo`, `Foo` and `F00` are distinct identifiers. The `'_'` symbol, when occurring on the left-hand side of an equation, is special; it denotes the **anonymous variable** which matches any value without actually binding a variable. Identifiers can also be prefixed with a namespace identifier, like in `foo:bar`. (This requires that the given namespace has already been created, as explained under [Namespaces](#) in the [Declarations](#) section.)

Operator and constant symbols: `x+y`, `x==y`, `not x` For convenience, Pure also provides you with a limited means to extend the syntax of the language with special operator and constant symbols by means of a corresponding **fixity** declaration, as discussed in section [Symbol Declarations](#). Besides the usual infix, prefix and postfix operators, Pure also provides outfix (bracket) and nonfix (constant) symbols. (Nonfix symbols actually work more or less like ordinary identifiers, but the `nonfix` attribute tells the compiler that when such a symbol occurs on the left-hand side of an equation, it is always to be interpreted as a literal constant, cf. [Parameters in Equations](#).)

Operator and constant symbols may take the form of an identifier or a sequence of punctuation characters. They must always be declared before use. Once declared, they are always special, and can't be used as ordinary identifiers any more. However, like

in Haskell, by enclosing an operator in parentheses, such as `(+)` or `(not)`, you can turn it into an ordinary function symbol. Also, operators and constant symbols can be qualified with a namespace just like normal identifiers.

Lists: `[x,y,z]`, `x:xs`, `x..y`, `x:y..z` Pure's basic list syntax is the same as in Haskell, thus `[]` is the empty list and `x:xs` denotes a list with head element `x` and tail list `xs`. (The infix constructor symbol `'.'` is declared in the prelude.) The usual syntactic sugar for list values in brackets is provided, thus `[x,y,z]` is exactly the same as `x:y:z:[]`.

There's also a way to denote arithmetic sequences such as `1..5`, which denotes the list `[1,2,3,4,5]`. (Haskell users should note the missing brackets. In difference to Haskell, Pure doesn't use any special syntax for arithmetic sequences, the `'.'` symbol is just an ordinary infix operator declared and defined in the prelude.) Sequences with arbitrary stepsizes can be written by denoting the first two sequence elements using the `':'` operator, as in `1.0:1.2..3.0`. (To prevent unwanted artifacts due to rounding errors, the upper bound in a floating point sequence is always rounded to the nearest grid point. Thus, e.g., `0.0:0.1..0.29` actually yields `[0.0,0.1,0.2,0.3]`, as does `0.0:0.1..0.31`.)

Tuples: `x,y,z` Pure's tuples are a bit unusual: They are constructed by just "pairing" things using the `'.'` operator, for which the empty tuple `()` acts as a neutral element (i.e., `() , x` is just `x`, as is `x , ()`). Pairs always associate to the right, meaning that `x,y,z == x,(y,z) == (x,y),z`, where `x,(y,z)` is the normalized representation. This implies that tuples are always flat, i.e., there are no nested tuples (tuples of tuples); if you need such constructs then you should use lists instead. Also note that parentheses are generally only used to group expressions and are *not* part of the tuple syntax in Pure, *except* if you need to include a tuple in a list or matrix. E.g., `[(1,2),3,(4,5)]` is a three element list consisting of the tuple `1,2`, the integer `3`, and another tuple `4,5`. Likewise, `[(1,2,3)]` is a list with a single element, the tuple `1,2,3`.

Matrices: `{1.0,2.0,3.0}`, `{1,2;3,4}`, `{1L,y+1;foo,bar}` Pure also offers matrices, a kind of two-dimensional arrays, as a built-in data structure which provides efficient storage and element access. These work more or less like their Octave/MATLAB equivalents, but using curly braces instead of brackets. As indicated, commas are used to separate the columns of a matrix, semicolons for its rows. In fact, the `{...}` construct is rather general and allows you to construct new matrices from any collection of individual elements ("scalars") and submatrices, provided that all dimensions match up. Here, any expression which doesn't yield a matrix denotes a scalar, which is considered to be a `1x1` matrix for the purpose of matrix construction. The comma arranges submatrices in columns, while the semicolon arranges them in rows. So, if both `x` and `y` are `nxm` matrices, then `{x,y}` becomes an `n x 2*m` matrix consisting of all the columns of `x` followed by all the columns of `y`. Likewise, `{x;y}` becomes a `2*n x m` matrix (all the rows of `x` above of all rows of `y`). In addition, `{...}` constructs can be nested to an arbitrary depth. Thus `{{1;3},{2;4}}` is another way to write the `2x2` matrix `{1,2;3,4}` in a kind of "column-major" format (however, internally all matrices are stored in C's row-major format).

It is important to note here that, while the `[...]` and `{...}` constructs look superficially similar, they work in very different ways. The former is just syntactic sugar for a

corresponding constructor term and can thus be used as a pattern on the left-hand side of an equation. In contrast, the latter is special syntax for a built-in operation which creates objects of a special matrix type. Thus matrix expressions can *not* be used as patterns (instead, matrix values can be matched as a whole using the special matrix type tag, see the [Rule Syntax](#) section for details).

Pure supports both numeric and symbolic matrices. The former are homogeneous arrays of double, complex double or (machine) int matrices, while the latter can contain any mixture of Pure expressions. Pure will pick the appropriate type for the data at hand. If a matrix contains values of different types, or Pure values which cannot be stored in a numeric matrix, then a symbolic matrix is created instead (this also includes the case of bigints, which are considered as symbolic values as far as matrix construction is concerned). Numeric matrices use an internal data layout that is fully compatible with the [GNU Scientific Library](#) (GSL), and can readily be passed to GSL routines via the C interface. (The Pure interpreter does not require GSL, however, so numeric matrices will work even if GSL is not installed.)

More information about matrices and corresponding examples can be found in the [Examples](#) section below.

Comprehensions: `[x,y | x=1..n; y=1..m; x<y], {f x | x=1..n}` Pure provides both list and matrix comprehensions as a convenient means to construct list and matrix values from a “template” expression and one or more “generator” and “filter” clauses. The former bind a pattern to values drawn from a list or matrix, the latter are just predicates determining which generated elements should actually be added to the result. Both list and matrix comprehensions are in fact syntactic sugar for a combination of nested lambdas, conditional expressions and “catmaps” (a collection of operations which combine list or matrix construction and mapping a function over a list or matrix, defined in the prelude), but they are often much easier to write.

Matrix comprehensions work pretty much like list comprehensions, but produce matrices instead of lists. List generators in matrix comprehensions alternate between row and column generation so that most common mathematical abbreviations carry over quite easily. Examples of both kinds of comprehensions can be found in the [Examples](#) section below.

Function and operator applications: `foo x y z, -x, x+y, (+x)` As in other modern FPLs, function applications are written simply as juxtaposition (i.e., in “curried” form) and associate to the left. This means that in fact all functions only take a single argument. Multi-argument functions are represented as chains of single-argument functions. For instance, in `f x y = (f x) y` first the function `f` is applied to the first argument `x`, yielding the function `f x` which in turn gets applied to the second argument `y`. This makes it possible to derive new functions from existing ones using **partial applications** which only specify some but not all arguments of a function. For instance, taking the `max` function from the prelude as an example, `max 0` is the function which, for a given `x`, returns `x` itself if it is nonnegative and zero otherwise. This works because `(max 0) x = max 0 x` is the maximum of 0 and `x`.

Operator applications are written using prefix, postfix, outfix or infix notation, as the declaration of the operator demands, but are just ordinary function applications in dis-

guise. As already mentioned, enclosing an operator in parentheses turns it into an ordinary function symbol, thus $x+y$ is exactly the same as $(+) \ x \ y$. For convenience, partial applications of infix operators can also be written using so-called **operator sections**. A *left section* takes the form $(x+)$ which is equivalent to the partial application $(+) \ x$. A *right section* takes the form $(+x)$ and is equivalent to the term `flip (+) x`. (This uses the `flip` combinator from the prelude which is defined as `flip f x y = f y x`.) Thus $(x+)$ y is equivalent to $x+y$, while $(+x)$ y reduces to $y+x$. For instance, $(1/)$ denotes the reciprocal and $(+1)$ the successor function. (Note that, in contrast, $(-x)$ always denotes an application of unary minus; the section $(+-x)$ can be used to indicate a function which subtracts x from its argument.)

Conditional expressions: `if x then y else z` Evaluates to y or z depending on whether x is “true” (i.e., a nonzero integer). An exception is raised if the condition is not an integer.

Lambdas: `\x -> y` These denote anonymous functions and work pretty much like in Haskell. Pure supports multiple-argument lambdas (e.g., `\x y -> x*y`), as well as pattern-matching lambda abstractions which match one or more patterns against the lambda arguments, such as `\(x,y) -> x*y`. An exception is raised if the actual lambda arguments do not match the given patterns.

Case expressions: `case x of rule; ... end` Matches an expression, discriminating over a number of different cases, similar to the Haskell case construct. The expression x is matched in turn against each left-hand side pattern in the rule list, and the first pattern which matches x gives the value of the entire expression, by evaluating the corresponding right-hand side with the variables in the pattern bound to their corresponding values. An exception is raised if the target expression doesn’t match any of the patterns.

When expressions: `x when rule; ... end` An alternative way to bind local variables by matching a collection of subject terms against corresponding patterns, similar to Aardappel’s `when` construct. A single binding such as `x when u = v end` is equivalent to `case v of u = x end`, but the former is often more convenient to write. A `when` clause may contain multiple definitions, which are processed from left to right, so that later definitions may refer to the variables in earlier ones. This is exactly the same as several nested single definitions, with the first binding being the “outermost” one.

With expressions: `x with rule; ... end` Defines local functions. Like Haskell’s `where` construct, but it can be used anywhere inside an expression (just like Aardappel’s `where`, but Pure uses the keyword `with` which better lines up with `case` and `when`). Several functions can be defined in a single `with` clause, and the definitions can be mutually recursive and consist of as many equations as you want.

3.5 Operators and Precedence

Expressions are parsed according to the following precedence rules: Lambda binds most weakly, followed by `when`, `with` and `case`, followed by conditional expressions (`if-then-else`), followed by the **simple expressions**, i.e., all other kinds of expressions involving oper-

ators, function applications, constants, symbols and other primary expressions. Precedence and associativity of operator symbols are given by their declarations (cf. [Symbol Declarations](#)), and function application binds stronger than all operators. Parentheses can be used to override default precedences and associativities as usual.

The common operator symbols like `+`, `-`, `*`, `/` etc. are all declared at the beginning of the prelude, see the [Pure Library Manual](#) for a list of these. Arithmetic and relational operators mostly follow C conventions. However, out of necessity (`!`, `&` and `|` are used for other purposes in Pure) the logical and bitwise operations, as well as the negated equality predicates are named a bit differently: `~`, `&&` and `||` denote logical negation, conjunction and disjunction, while the corresponding bitwise operations are named `not`, `and` and `or`. Moreover, following these conventions, inequality is denoted `~=`. Also note that `&&` and `||` are special forms which are evaluated in short-circuit mode (see below), whereas the bitwise connectives receive their arguments using call-by-value, just like the other arithmetic operations.

3.6 Special Forms

As already mentioned, some operations are actually implemented as special forms. In particular, the conditional expression `if x then y else z` is a special form with call-by-name arguments `y` and `z`; only one of the branches is actually evaluated, depending on the value of `x`. Similarly, the logical connectives `&&` and `||` evaluate their operands in short-circuit mode. Thus, e.g., `x&&y` immediately becomes false if `x` evaluates to false, without ever evaluating `y`. The built-in definitions of these operations work as if they were defined by the following equations (but note that the second operand is indeed passed “by name”, so that it isn’t evaluated unless its value is actually needed):

```
x::int && y = if x then y else x;
x::int || y = if x then x else y;
```

(Note that this isn’t quite the same as in C, as the results of these operations are *not* normalized, i.e., they may return nonzero values other than 1 to denote “true”. This has the advantage that `&&` and `||` can be implemented tail-recursively, see [Stack Size and Tail Recursion](#). Thus, if you need a normalized truth value then you’ll have to make sure that either both operands are already normalized, or you’ll have to normalize the result yourself.)

If the built-in definitions of `&&` and `||` fail then both operands are evaluated and the resulting symbolic term is reduced using the equations specified by the programmer, which allows you to extend the definitions of `&&` and `||` just as with other builtins. Also note that if the result term is in normal form then it is returned as is, thus `&&` and `||` can be used in symbolic evaluations as usual. For instance, the following equations let you compute the disjunctive normal form of logical expressions:

```
// eliminate double negations:
~~a      = a;

// de Morgan's laws:
~(a || b) = ~a && ~b;
~(a && b)  = ~a || ~b;
```

```
// distributivity:
a && (b || c) = a && b || a && c;
(a || b) && c = a && c || b && c;

// associativity:
(a && b) && c = a && (b && c);
(a || b) || c = a || (b || c);
```

Example:

```
> a || ~(b || (c && ~d));
a || ~b&&~c || ~b&&d
```

The sequencing operator `$$` evaluates its left operand, immediately throws the result away and then goes on to evaluate the right operand which gives the result of the entire expression. This operator is useful to write imperative-style code such as the following prompt-input interaction:

```
> using system;
> puts "Enter a number:" $$ scanf "%g";
Enter a number:
21
21.0
```

We mention in passing here that the same effect can be achieved with a `when` clause, which also allows you to execute a function solely for its side-effects and just ignore the return value:

```
> scanf "%g" when puts "Enter a number:" end;
Enter a number:
21
21.0
```

The `&` operator does lazy evaluation. This is the only postfix operator defined in the standard prelude, written as `x&`, where `x` is an arbitrary Pure expression. It turns its operand into a kind of parameterless anonymous closure, deferring its evaluation. These kinds of objects are also commonly known as **thunks** or **futures**. When the value of a future is actually needed (during pattern-matching, or when the value becomes an argument of a C call), it is evaluated automatically and gets memoized, i.e., the computed result replaces the thunk so that it only has to be computed once. Futures are useful to implement all kinds of lazy data structures in Pure, in particular: lazy lists a.k.a. streams. A **stream** is simply a list with a thunked tail, which allows it to be infinite. The Pure prelude defines many functions for creating and manipulating these kinds of objects; further details and examples can be found in the [Examples](#) section below.

Last but not least, the special form `quote` quotes an expression, i.e., `quote x` (or, equivalently, `'x`) returns just `x` itself without evaluating it. The prelude also provides a function `eval` which can be used to evaluate a quoted expression at a later time. For instance:

```
> let x = '(2*42+2^12); x;  
2*42+2^12  
> eval x;  
4180.0
```

The quote also inhibits evaluation inside matrix values, including the “splicing” of embedded submatrices:

```
> '{1,2+3,2*3};  
{1,2+3,2*3}  
> '{1,{2,3},4};  
{1,{2,3},4}
```

The quote should be well familiar to Lisp programmers. However, there are some notable differences, please see [The Quote](#) in the [Caveats and Notes](#) section for details and more examples.

3.7 Toplevel

At the toplevel, a Pure program basically consists of rewriting rules (which are used to define functions and macros), constant and variable definitions, and expressions to be evaluated:

Rules: **lhs = rhs;** Rewriting rules always combine a left-hand side pattern (which must be a simple expression) and a right-hand side (which can be any kind of Pure expression described above). The same format is also used in `with`, `when` and `case` expressions. In toplevel rules, `with` and `case` expressions, this basic form can also be augmented with a condition `if guard` tacked on to the end of the rule, where `guard` is an integer expression which determines whether the rule is applicable. Moreover, the keyword `otherwise` may be used to denote an empty guard which is always true (this is syntactic sugar to point out the “default” case of a definition; the interpreter just treats this as a comment). Pure also provides some abbreviations for factoring out common left-hand or right-hand sides in collections of rules; see the [Rule Syntax](#) section for details.

Macro rules: **def lhs = rhs;** A rule starting with the keyword `def` defines a macro function. No guards or multiple left-hand and right-hand sides are permitted here. Macro rules are used to preprocess expressions on the right-hand side of other definitions at compile time, and are typically employed to implement user-defined special forms and simple kinds of optimization rules. See the [Macros](#) section below for details and examples.

Global variable bindings: **let lhs = rhs;** Binds every variable in the left-hand side pattern to the corresponding subterm of the right-hand side (after evaluating it). This works like a `when` clause, but serves to bind global variables occurring free on the right-hand side of other function and variable definitions.

Constant bindings: **const lhs = rhs;** An alternative form of `let` which defines constants rather than variables. (These are not to be confused with nonfix symbols which simply

stand for themselves!) Like `let`, this construct binds the variable symbols on the left-hand side to the corresponding values on the right-hand side (after evaluation). The difference is that `const` symbols can only be defined once, and thus their values do not change during program execution. This also allows the compiler to apply some special optimizations such as constant folding.

Toplevel expressions: `expr`; A singleton expression at the toplevel, terminated with a semicolon, simply causes the given value to be evaluated (and the result to be printed, when running in interactive mode).

3.8 Scoping Rules

A few remarks about the scope of identifiers and other symbols are in order here. Like most modern functional languages, Pure uses **lexical** or **static** binding for local functions and variables. What this means is that the binding of a local name is completely determined at compile time by the surrounding program text, and does not change as the program is being executed. In particular, if a function returns another (anonymous or local) function, the returned function captures the environment it was created in, i.e., it becomes a (lexical) **closure**. For instance, the following function, when invoked with a single argument `x`, returns another function which adds `x` to its argument:

```
> foo x = bar with bar y = x+y end;
> let f = foo 99; f;
bar
> f 10, f 20;
109,119
```

This works the same no matter what other bindings of `x` may be in effect when the closure is invoked:

```
> let x = 77; f 10, (f 20 when x = 88 end);
109,119
```

Global bindings of variable and function symbols work a bit differently, though. Like many languages which are to be used interactively, Pure binds global symbols **dynamically**, so that the bindings can be changed easily at any time during an interactive session. This is mainly a convenience for interactive usage, but works the same no matter whether the source code is entered interactively or being read from a script, in order to ensure consistent behaviour between interactive and batch mode operation.

So, for instance, you can easily bind a global variable to a new value by just entering a corresponding `let` command:

```
> foo x = c*x;
> foo 99;
c*99
> let c = 2; foo 99;
198
> let c = 3; foo 99;
297
```

This works pretty much like global variables in imperative languages, but note that in Pure the value of a global variable can only be changed with a `let` command at the toplevel. Thus referential transparency is unimpaired; while the value of a global variable may change between different toplevel expressions, it will always take the same value in a single evaluation.

Similarly, you can also add new equations to an existing function at any time:

```
> fact 0 = 1;
> fact n::int = n*fact (n-1) if n>0;
> fact 10;
3628800
> fact 10.0;
fact 10.0
> fact 1.0 = 1.0;
> fact n::double = n*fact (n-1) if n>1;
> fact 10.0;
3628800.0
> fact 10;
3628800
```

(In interactive mode, it is even possible to completely erase a definition, see section [Interactive Usage](#) for details.)

So, while the meaning of a local symbol never changes once its definition has been processed, toplevel definitions may well evolve while the program is being processed, and the interpreter will always use the latest definitions at a given point in the source when an expression is evaluated. This means that, even in a script file, you have to define all symbols needed in an evaluation before entering the expression to be evaluated.

4 Rule Syntax

Basically, the same rule syntax is used in all kinds of global and local definitions. However, some constructs (specifically, `when`, `let`, `const` and `def`) use a restricted rule syntax where no guards or multiple left-hand and right-hand sides are permitted. When matching against a function or macro call, or the subject term in a case expression, the rules are always considered in the order in which they are written, and the first matching rule (whose guard evaluates to a nonzero value, if applicable) is picked. (Again, the `when` construct is treated differently, because each rule is actually a separate definition.)

4.1 Patterns

The left-hand side of a rule is a special kind of simple expression, called a **pattern**. Patterns consist of function and operator applications as well as any of the “atomic” expression types (symbols, numbers, strings and list values). *Not* permitted are any of the special expression types (lambda, case, `when`, `with`, conditional expressions, as well as list and matrix comprehensions). For technical reasons, the current implementation also forbids matrix values in

patterns, but it is possible to match a matrix value as a whole using the `matrix` type tag, see below.

As already mentioned, the `'_'` symbol is special in patterns; it denotes the **anonymous variable** which matches an arbitrary value (independently for all occurrences) without actually binding a variable. For instance:

```
foo _ _ = 0;
```

This will match the application of `foo` to any combination of two arguments (and just ignore the values of these arguments).

Constants in patterns must be matched literally. For instance:

```
foo 0 = 1;
```

This will only match an application of `foo` to the machine integer `0`, not `0.0` or `0L` (even though these compare equal to `0` using the `'=='` operator).

In difference to Haskell, patterns may contain repeated variables (other than the anonymous variable), i.e., they may be **non-linear**. Thus rules like the following are legal in Pure, and will only be matched if all occurrences of the same variable in the left-hand side pattern are matched to the same value:

```
> foo x x = x;  
> foo 1 1;  
1  
> foo 1 2;  
foo 1 2
```

Non-linear patterns are particularly useful for computer algebra where you will frequently encounter rules such as the following:

```
> x*y+x*z = x*(y+z);  
> a*(3*4)+a*5;  
a*17
```

The notion of “sameness” employed here is that of syntactical identity, which means that the matched subterms must be identical in structure and content. The prelude provides syntactic equality as a function *same* and a comparison predicate `'=='`. Thus the above definition of `foo` is roughly equivalent to the following:

```
foo x y = x if same x y;
```

It is important to note the differences between syntactic equality embodied by *same* and `'=='`, and the “semantic” equality operator `'=='`. The former are always defined on all terms, whereas `'=='` is only available on data where it has been defined explicitly, either in the prelude or by the programmer. Also note that `'=='` may assert that two terms are equal even if they are syntactically different. Consider, e.g.:

```
> 0==0.0;  
1
```



```
> 0===0.0;  
0
```

This distinction is actually quite useful. It gives the programmer the flexibility to define ‘==’ in any way that he sees fit, which is consistent with the way the other comparison operators like ‘<’ and ‘>’ are handled in Pure.

Patterns may also contain the following special elements which are not permitted in right-hand side expressions:

- A Haskell-style “**as**” **pattern** of the form *variable* @ *pattern* binds the given variable to the expression matched by the subpattern *pattern* (in addition to the variables bound by *pattern* itself). This is convenient if the value matched by the subpattern is to be used on the right-hand side of an equation.
- A left-hand side variable (including the anonymous variable) may be followed by a **type tag** of the form `:: name`, where *name* is either one of the built-in type symbols `int`, `bigint`, `double`, `string`, `matrix`, `pointer`, or an existing identifier denoting a custom constructor symbol for a user-defined data type. The variable can then match only values of the designated type. Thus, for instance, ‘`x :: int`’ only matches machine integers. See the [Type Tags](#) section below for details.

Syntactically, both “as” patterns and type tags are primary expressions. If the subpattern of an “as” pattern is a compound expression, it *must* be parenthesized. For instance, the following function duplicates the head element of a list:

```
foo xs@(x:_) = x:xs;
```

Note that if you accidentally forget the parentheses around the subpattern `x:_,` you still get a syntactically correct definition:

```
foo xs@x:_ = x:xs;
```

But this gets parsed as `(foo xs@x):_ = x:xs`, which is most certainly *not* what you want. It is thus a good idea to just always enclose the subpattern with parentheses in order to prevent such glitches.

Another potential pitfall is that the notation `foo :: bar` is also used to denote “qualified symbols” in Pure, cf. [Namespaces](#). Usually this will be resolved correctly, but if `foo` happens to also be a valid namespace then most likely you’ll get an error message about an undeclared symbol. You can always work around this by adding spaces around the ‘`::`’ symbol, as in `foo :: bar`. Spaces are never permitted in qualified symbols, so this makes it clear that the construct denotes a type tag.

4.2 Type Tags

Type tags are really nothing but a special form of “as” patterns which restrict variables to given data “domains”. Like Lisp, Pure is essentially a typeless language and doesn’t really have a notion of “data types”; all data belongs to the same universe of terms. Thus the only way to restrict the type of values which can be matched by a variable is to provide an “as”

pattern which matches objects of the desired domain and nothing else. However, in the special case of the built-in types (machine and big integers, double values, strings, matrices and pointers) there is no way to spell out all “constructors”, as there are infinitely many (or none, as in the case of `matrix` and `pointer` which are constructed and inspected using special primitives, but are otherwise “opaque” at the Pure level). As a remedy, an appropriate type tag makes it possible to match these values.

In order to generalize this to user-defined domains of data, Pure adopts the convention that any other tagged variable `x :: bar` is just a shorthand for the “as” pattern `x@(bar _)`. Thus a custom data type can be represented by designating a special constructor symbol (`bar` in the above example) which takes the actual data as its single argument. This works the same no matter what the internal structure of the data is (which in many cases you wouldn’t want to depend on anyway, for the sake of data abstraction).

Note that this is merely a convention, but it works reasonably well and makes up for the fact that Pure doesn’t support data types at the language level. For instance, we might represent points in the plane using a constructor symbol `Point` which gets applied to pairs of coordinates. We equip this data type with an operation `point` to construct a point from its coordinates, and two operations `xcoord` and `ycoord` to retrieve the coordinates:

```
point x y = Point (x,y);
xcoord (Point (x,y)) = x;
ycoord (Point (x,y)) = y;
```

Now we might define a function `translate` which shifts the coordinates of a point by a given amount in the `x` and `y` directions as follows:

```
translate (x,y) p::Point = point (xcoord p+x) (ycoord p+y);
```

Note the use of `Point` as a type tag on the `p` variable. By these means, we can ensure that the argument is actually an instance of the point data type, without knowing anything about the internal representation. We can use these definitions as follows:

```
> let p::Point = point 3 3;
> p; translate (1,2) p;
Point (3,3)
Point (4,5)
```

Some data types in Pure’s standard library (specifically, the container data types) are actually represented in this fashion, see the [Pure Library Manual](#) for details.

4.3 General Rules

The most general type of rule, used in function definitions and case expressions, consists of a left-hand side pattern, a right-hand side expression and an optional guard. The left-hand side of a rule can be omitted if it is the same as for the previous rule. This provides a convenient means to write out a collection of equations for the same left-hand side which discriminates over different conditions:

```
lhs      = rhs if guard;
          = rhs if guard;
          ...
          = rhs otherwise;
```

For instance:

```
fact n   = n*fact (n-1) if n>0;
          = 1 otherwise;
```

Pure also allows a collection of rules with different left-hand sides but the same right-hand side(s) to be abbreviated as follows:

```
lhs      |
          ...
lhs      = rhs;
```

This is useful if you need different specializations of the same rule which use different type tags on the left-hand side variables. For instance:

```
fact n::int    |
fact n::double |
fact n         = n*fact(n-1) if n>0;
               = 1 otherwise;
```

In fact, the left-hand sides don't have to be related at all, so that you can also write something like:

```
foo x | bar y = x*y;
```

However, this construct is most useful when using an “as” pattern to bind a common variable to a parameter value after checking that it matches one of several possible argument patterns (which is slightly more efficient than using an equivalent type-checking guard). E.g., the following definition binds the `xs` variable to the parameter of `foo`, if it is either the empty list or a list starting with an integer:

```
foo xs@[ ] | foo xs@(_::int:_) = ... xs ...;
```

The same construct also works in case expressions, which is convenient if different cases should be mapped to the same value, e.g.:

```
case ans of "y" | "Y" = 1; _ = 0; end;
```

Sometimes it is useful if local definitions (when and with) can be shared by the right-hand side and the guard of a rule. This can be done by placing the local definitions behind the guard, as follows (we only show the case of a single when clause here, but of course there may be any number of when and with clauses behind the guard):

```
lhs = rhs if guard when defns end;
```

Note that this is different from the following, which indicates that the definitions only apply to the guard but not the right-hand side of the rule:

```
lhs = rhs if (guard when defns end);
```

Conversely, definitions placed *before* the guard only apply to the right-hand side but not the guard (no parentheses are required in this case):

```
lhs = rhs when defns end if guard;
```

An example showing the use of a local variable binding spanning both the right-hand side and the guard of a rule is the following quadratic equation solver, which returns the (real) solutions of the equation $x^2+px+q = 0$ if the discriminant $d = p^2/4-q$ is nonnegative:

```
> using math;
> solve p q = -p/2+sqrt d, -p/2-sqrt d if d>=0 when d = p^2/4-q end;
> solve 4 2; solve 2 4;
-0.585786437626905, -3.41421356237309
solve 2 4
```

Note that the above definition leaves the case of a negative discriminant undefined.

4.4 Simple Rules

As already mentioned, `when`, `let` and `const` use a simplified kind of rule syntax which just consists of a left-hand and a right-hand side separated by the equals sign. In this case the meaning of the rule is to bind the variables in the left-hand side of the rule to the corresponding subterms of the value of the right-hand side. This is also called a **pattern binding**.

Guards or multiple left-hand or right-hand sides are not permitted in these rules. However, it is possible to omit the left-hand side if it is just the anonymous variable `'_'` by itself, indicating that you don't care about the result. The right-hand side is still evaluated, if only for its side-effects, which is handy, e.g., for adding debugging statements to your code. For instance, here is a variation of the quadratic equation solver which also prints the discriminant after it has been computed:

```
> using math, system;
> solve p q = -p/2+sqrt d, -p/2-sqrt d if d>=0
> when d = p^2/4-q; printf "The discriminant is: %g\n" d; end;
> solve 4 2;
The discriminant is: 2
-0.585786437626905, -3.41421356237309
> solve 2 4;
The discriminant is: -3
solve 2 4
```

Note that simple rules of the same form `lhs = rhs` are also used in macro definitions (`def`), to be discussed in the [Macros](#) section. In this case, however, the rule denotes a real rewriting rule, not a pattern binding, hence the left-hand side is mandatory in these rules.

5 Examples

Here are a few examples of simple Pure programs.

The factorial:

```
fact n = n*fact (n-1) if n>0;
        = 1 otherwise;
let facts = map fact (1..10); facts;
```

The Fibonacci numbers:

```
fib n = a when a,b = fibs n end
        with fibs n = 0,1 if n<=0;
                = case fibs (n-1) of
                    a,b = b,a+b;
                end;
        end;
let fibs = map fib (1..30); fibs;
```

It is worth noting here that Pure performs tail call optimization so that tail-recursive definitions like the following will be executed in constant stack space (see [Stack Size and Tail Recursion](#) in the [Caveats and Notes](#) section for more details on this):

```
// tail-recursive factorial using an "accumulating parameter"
fact n = loop 1 n with
    loop p n = if n>0 then loop (p*n) (n-1) else p;
end;
```

Here is an example showing how constants are defined and used. Constant definitions take pretty much the same form as variable definitions with `let` (see above), but work more like the definition of a parameterless function whose value is precomputed at compile time:

```
> extern double atan(double);
> const pi = 4*atan 1.0;
> pi;
3.14159265358979
> foo x = 2*pi*x;
> show foo
foo x = 6.28318530717959*x;
```

Note that the compiler normally computes constant subexpression at compile time, such as `2*pi` in the `foo` function. This works with all simple scalars (machine ints and doubles), see [Constant Definitions](#) for details.

5.1 List Comprehensions

List comprehensions are Pure's main workhorse for generating and processing all kinds of list values. Here's a well-known example, a variation of Erathosthenes' classical prime sieve:

```
primes n      = sieve (2..n) with
  sieve []    = [];
  sieve (p:qs) = p : sieve [q | q = qs; q mod p];
end;
```

(This definition is actually rather inefficient, there are much better albeit more complicated implementations of this sieve.)

For instance:

```
> primes 100;
[2,3,5,7,11,13,17,19,23,29,31,37,41,43,47,53,59,61,67,71,73,79,83,89,97]
```

If you dare, you can actually have a look at the `catmap-lambda-if-then-else` expression the comprehension expanded to:

```
> show primes
primes n = sieve (2..n) with sieve [] = []; sieve (p:qs) = p:sieve
(catmap (\q -> if q mod p then [q] else []) qs) end;
```

List comprehensions are also a useful device to organize backtracking searches. For instance, here's an algorithm for the *n* queens problem, which returns the list of all placements of *n* queens on an *n* x *n* board (encoded as lists of *n* pairs (*i*,*j*) with *i* = 1..*n*), so that no two queens hold each other in check:

```
queens n      = search n 1 [] with
  search n i p = [reverse p] if i>n;
                = cat [search n (i+1) ((i,j):p) | j = 1..n; safe (i,j) p];
  safe (i,j) p = ~any (check (i,j)) p;
  check (i1,j1) (i2,j2)
    = i1==i2 || j1==j2 || i1+j1==i2+j2 || i1-j1==i2-j2;
end;
```

(Again, this algorithm is rather inefficient, see the examples included in the Pure distribution for a much better algorithm by Libor Spacek.)

5.2 Lazy Evaluation and Streams

As already mentioned, lists can also be evaluated in a “lazy” fashion, by just turning the tail of a list into a future. This special kind of list is also called a **stream**. Streams enable you to work with infinite lists (or finite lists which are so huge that you would never want to keep them in memory in their entirety). E.g., here's one way to define the infinite stream of all Fibonacci numbers:

```
> let fibs = fibs 0L 1L with fibs a b = a : fibs b (a+b) & end;
> fibs;
0L:#<thunk 0xb5d54320>
```

Note the `&` on the tail of the list in the definition of the local `fibs` function. This turns the result of `fibs` into a stream, which is required to prevent the function from recursing into

samadhi. Also note that we work with bigints in this example because the Fibonacci numbers grow quite rapidly, so with machine integers the values would soon start wrapping around to negative integers.

Streams like these can be worked with in pretty much the same way as with lists. Of course, care must be taken not to invoke “eager” operations such as `#` (which computes the size of a list) on infinite streams, to prevent infinite recursion. However, many list operations work with infinite streams just fine, and return the appropriate stream results. E.g., the `take` function (which retrieves a given number of elements from the front of a list) works with streams just as well as with “eager” lists:

```
> take 10 fibs;
0L: #<thunk 0xb5d54350>
```

Hmm, not much progress there, but that’s just how streams work (or rather they don’t, they’re lazy bums indeed!). Nevertheless, the stream computed with `take` is in fact finite and we can readily convert it to an ordinary list, forcing its evaluation:

```
> list (take 10 fibs);
[0L, 1L, 1L, 2L, 3L, 5L, 8L, 13L, 21L, 34L]
```

An easier way to achieve this is to cut a “slice” from the stream:

```
> fibs!!(0..10);
[0L, 1L, 1L, 2L, 3L, 5L, 8L, 13L, 21L, 34L, 55L]
```

Also note that since we bound the stream to a variable, the already computed prefix of the stream has been memoized, so that this portion of the stream is now readily available in case we need to have another look at it later. By these means, possibly costly reevaluations are avoided, trading memory for execution speed:

```
> fibs;
0L: 1L: 1L: 2L: 3L: 5L: 8L: 13L: 21L: 34L: 55L: #<thunk 0xb5d54590>
```

Let’s take a look at some of the other convenience operations for generating stream values. The prelude defines infinite arithmetic sequences, using `inf` or `-inf` to denote an upper (or lower) infinite bound for the sequence, e.g.:

```
> let u = 1..inf; let v = -1.0:-1.2..-inf;
> u!!(0..10); v!!(0..10);
[1,2,3,4,5,6,7,8,9,10,11]
[-1.0,-1.2,-1.4,-1.6,-1.8,-2.0,-2.2,-2.4,-2.6,-2.8,-3.0]
```

Other useful stream generator functions are `iterate`, which keeps applying the same function over and over again, `repeat`, which just repeats its argument forever, and `cycle`, which cycles through the elements of the given list:

```
> iterate (*2) 1!!(0..10);
[1,2,4,8,16,32,64,128,256,512,1024]
> repeat 1!!(0..10);
[1,1,1,1,1,1,1,1,1,1,1]
```

```
> cycle [0,1]!!(0..10);  
[0,1,0,1,0,1,0,1,0,1,0]
```

Moreover, list comprehensions can draw values from streams and return the appropriate stream result:

```
> let rats = [m,n-m | n=2..inf; m=1..n-1; gcd m (n-m) == 1]; rats;  
(1,1):#<thunk 0xb5d54950>  
> rats!!(0..10);  
[(1,1),(1,2),(2,1),(1,3),(3,1),(1,4),(2,3),(3,2),(4,1),(1,5),(5,1)]
```

Finally, let's rewrite our prime sieve so that it generates the infinite stream of *all* prime numbers:

```
all_primes      = sieve (2..inf) with  
  sieve (p:qs)  = p : sieve [q | q = qs; q mod p] &;  
end;
```

Note that we can omit the empty list case of sieve here, since the sieve now never becomes empty. Example:

```
> let P = all_primes;  
> P!!(0..20);  
[2,3,5,7,11,13,17,19,23,29,31,37,41,43,47,53,59,61,67,71,73]  
> P!299;  
1987
```

You can also just print the entire stream. This will run forever, so hit Ctrl-C when you get bored:

```
> using system;  
> do (printf "%d\n") all_primes;  
2  
3  
5  
...
```

(Make sure that you really use the `all_primes` function instead of the `P` variable to print the stream. Otherwise, because of memoization the stream stored in `P` will grow with the number of elements printed until memory is exhausted. Calling `do` on a fresh instance of the stream of primes allows `do` to get rid of each “cons” cell after having printed the corresponding stream element.)

5.3 Matrix Computations

Pure offers a number of basic matrix operations, such as matrix construction, indexing, slicing, as well as getting the size and dimensions of a matrix (these are briefly described in the [Standard Library](#) section below). However, it does *not* supply built-in support for matrix arithmetic and other linear algebra algorithms. The idea is that these can and should be

provided through separate libraries (please check the Pure website for the pure-gsl module which is an ongoing project to provide a full GSL interface for the Pure language).

But Pure's facilities for matrix and list processing also make it easy to roll your own, if desired. First, the prelude provides matrix versions of the common list operations like `map`, `fold`, `zip` etc., which provide a way to implement common matrix operations. E.g., multiplying a matrix `x` with a scalar `a` amounts to mapping the function `(a*)` to `x`, which can be done as follows:

```
> a * x::matrix = map (a*) x if ~matrixp a;
> 2*{1,2,3;4,5,6};
{2,4,6;8,10,12}
```

Likewise, matrix addition and other element-wise operations can be realized using `zipwith`, which combines corresponding elements of two matrices using a given binary function:

```
> x::matrix + y::matrix = zipwith (+) x y;
> {1,2,3;4,5,6}+{1,2,1;3,2,3};
{2,4,4;7,7,9}
```

Second, matrix comprehensions make it easy to express a variety of algorithms which would typically be implemented using for loops in conventional programming languages. To illustrate the use of matrix comprehensions, here is how we can define an operation to create a square identity matrix of a given dimension:

```
> eye n = {i==j | i = 1..n; j = 1..n};
> eye 3;
{1,0,0;0,1,0;0,0,1}
```

Note that the `i==j` term is just a Pure idiom for the Kronecker symbol. Another point worth mentioning here is that the generator clauses of matrix comprehensions alternate between row and column generation automatically, if values are drawn from lists as in the example above. (More precisely, the last generator, which varies most quickly, yields a row, the next-to-last one a column of these row vectors, and so on.) This makes matrix comprehensions resemble customary mathematical notation very closely.

Of course, matrix comprehensions can also draw values from other matrices instead of lists. In this case the block layout of the component matrices is preserved. For instance:

```
> {x,y|x={1,2};y={a,b;c,d}};
{(1,a),(1,b),(2,a),(2,b);(1,c),(1,d),(2,c),(2,d)}
```

Note that a matrix comprehension involving filters may fail because the filtered result isn't a rectangular matrix any more. E.g., `{2*x|x={1,2,3,-4};x>0}` works, as does `{2*x|x={-1,2,3,-4};x>0}`, but `{2*x|x={1,2,3,-4};x>0}` doesn't because the rows of the result matrix have different lengths.

As a slightly more comprehensive example (no pun intended!), here is a definition of matrix multiplication in Pure. The building block here is the "dot" product of two vectors which can be defined as follows:

```
> sum = foldl (+) 0;
> dot x::matrix y::matrix = sum $ zipwith (*) (rowvector x) (rowvector y);
> dot {1,2,3} {1,0,1};
4
```

The general matrix product now boils down to a simple matrix comprehension which just computes the dot product of all rows of `x` with all columns of `y` (the `rows` and `cols` functions are prelude operations found in `matrices.pure`):

```
> x::matrix * y::matrix = {dot u v | u = rows x; v = cols y};
> {0,1;1,0;1,1}*{1,2,3;4,5,6};
{4,5,6;1,2,3;5,7,9}
```

(For the sake of simplicity, this doesn't do much error checking. In production code you'd check at least the conformance of matrix dimensions, of course.)

Well, that was easy. So let's take a look at a more challenging example, Gaussian elimination, which can be used to solve systems of linear equations. The algorithm brings a matrix into "row echelon" form, a generalization of triangular matrices. The resulting system can then be solved quite easily using back substitution.

Here is a Pure implementation of the algorithm. Note that the real meat is in the pivoting and elimination step (step function) which is iterated over all columns of the input matrix. In each step, `x` is the current matrix, `i` the current row index, `j` the current column index, and `p` keeps track of the current permutation of the row indices performed during pivoting. The algorithm returns the updated matrix `x`, row index `i` and row permutation `p`.

```
gauss_elimination x::matrix = p,x
when n,m = dim x; p,_,x = foldl step (0..n-1,0,x) (0..m-1) end;

// One pivoting and elimination step in column j of the matrix:
step (p,i,x) j
= if max_x==0 then p,i,x
  else
    // updated row permutation and index:
    transp i max_i p, i+1,
    // the top rows of the matrix remain unchanged:
    x!!(0..i-1,0..m-1);
    // the pivot row, divided by the pivot element:
    {x!(i,l)/x!(i,j) | l=0..m-1};
    // subtract suitable multiples of the pivot row:
    {x!(k,l)-x!(k,j)*x!(i,l)/x!(i,j) | k=i+1..n-1; l=0..m-1}}
when
  n,m = dim x; max_i, max_x = pivot i (col x j);
  x = if max_x>0 then swap x i max_i else x;
end with
  pivot i x = foldl max (0,0) [j,abs (x!j)|j=i..#x-1];
  max (i,x) (j,y) = if x<y then j,y else i,x;
end;
```

Please refer to any good textbook on numerical mathematics for a closer description of the algorithm. But here is a brief rundown of what happens in each elimination step: First we

find the pivot element in column j of the matrix. (We're doing partial pivoting here, i.e., we only look for the element with the largest absolute value in column j , starting at row i . That's usually good enough to achieve numerical stability.) If the pivot is zero then we're done (the rest of the pivot column is already zeroed out). Otherwise, we bring it into the pivot position (swapping row i and the pivot row), divide the pivot row by the pivot, and subtract suitable multiples of the pivot row to eliminate the elements of the pivot column in all subsequent rows. Finally we update i and p accordingly and return the result.

In order to complete the implementation, we still need the following little helper functions to swap two rows of a matrix (this is used in the pivoting step) and to apply a transposition to a permutation (represented as a list):

```
swap x i j = x!!(transp i j (0..n-1),0..m-1) when n,m = dim x end;
transp i j p = [p!tr k | k=0..#p-1]
with tr k = if k==i then j else if k==j then i else k end;
```

Finally, let us define a convenient print representation of double matrices a la [Octave](#) (the meaning of the `__show__` function is explained in the [Caveats and Notes](#) section):

```
using system;
__show__ x::matrix
= strcat [printf (x!(i,j))|i=0..n-1; j=0..m-1] + "\n"
with printf 0 = sprintf "%10.5f"; printf _ = sprintf "%10.5f" end
when n,m = dim x end if dmatrixp x;
```

Example:

```
> let x = dmatrix {2,1,-1,8; -3,-1,2,-11; -2,1,2,-3};
> x; gauss_elimination x;
  2.00000  1.00000 -1.00000  8.00000
 -3.00000 -1.00000  2.00000 -11.00000
 -2.00000  1.00000  2.00000 -3.00000
[1,2,0],
  1.00000  0.33333 -0.66667  3.66667
  0.00000  1.00000  0.40000  2.60000
  0.00000  0.00000  1.00000 -1.00000
```

5.4 Symbolic Matrices

As already mentioned, matrices may contain not just numbers but any kind of Pure value, in which case they become *symbolic* matrices. Symbolic matrices are a convenient data structure for storing arbitrary collections of values which provides fast random access to its members. In particular, symbolic matrices can also be nested, and thus arrays of arbitrary dimension can be realized as nested symbolic vectors. However, you have to be careful when constructing such values, as the `{...}` construct normally combines submatrices to larger matrices. For instance:

```
> {{1,2},{3,4}};
{1,2,3,4}
```

As already mentioned, one way to inhibit this “splicing” of the submatrices in a larger matrix is to use the “quote” operator (cf. [Special Forms](#)):

```
> ' {{1,2},{3,4}};  
{{1,2},{3,4}}
```

(Note that this result is really different from `{1,2;3,4}`. The latter is a 2x2 integer matrix, while the former is a symbolic vector a.k.a. 1x2 matrix whose elements happen to be two integer vectors.)

Unfortunately, the quote operator in fact inhibits evaluation of *all* embedded subterms which may be undesirable if the matrix expression contains arithmetic (as in `' {{1+1,2*3}}`), so this method works best for constant matrices. A more general way to create a symbolic vector of matrices is provided by the `vector` function from the prelude, which is applied to a list of the vector elements as follows:

```
> vector [{1,2},{3,4}];  
{{1,2},{3,4}}
```

Calls to the `vector` function can be nested to an arbitrary depth to obtain higher-dimensional “arrays”:

```
> vector [vector [{1,2}],vector [{3,4}]];  
{{{1,2}},{3,4}}}
```

This obviously becomes a bit unwieldy for higher dimensions, but in Pure you can easily define yourself some more convenient notation if you like. For instance, the following macro may be used to define a pair of “non-splicing” vector brackets:

```
> outfix {: :};  
> def {: xs@(_,-) :} = vector (list xs);  
> def {: x :} = vector [x];  
> {:{:1,2}:},{3,4}:};  
{{{1,2}},{3,4}}}
```

(Both macros and `outfix` symbol declarations are described later in the appropriate sections, see [Macros](#) and [Symbol Declarations](#).)

5.5 Record Data

Symbolic matrices also provide a means to represent simple record-like data, by encoding records as symbolic vectors consisting of “hash pairs” of the form `key => value`. This kind of data structure is very convenient to represent aggregates with lots of different components. Since the components of records can be accessed by indexing with key values, you don’t have to remember which components are stored in which order, just knowing the keys of the required members is enough. In contrast, tuples, lists and other kinds of constructor terms quickly become unwieldy for such purposes.

The keys used for indexing the record data must be either symbols or strings, while the corresponding values may be arbitrary Pure values. The prelude provides some operations

on these special kinds of matrices, which let you retrieve vector elements by indexing and perform non-destructive updates, see the *Record Functions* section in the *Pure Library Manual* for details. Here are a few examples which illustrate how to create records and work with them:

```
> let r = {x=>5, y=>12};
> recordp r, member r x;
1,1
> r!y; r!![y,x];
12
{12,5}
> insert r (x=>99);
{x=>99,y=>12}
> insert ans (z=>77);
{x=>99,y=>12,z=>77}
> delete ans z;
{x=>99,y=>12}
```

Note the use of the “hash rocket” `=>` which denotes the key=>value associations in a record. The hash rocket is a constructor declared as an infix operator in the prelude, see the *Prelude* section in the *Pure Library Manual*. There’s one caveat here, however. Since neither `=>` nor `!` treat their key operand in a special way, symbols used as keys must not be bound to a value, or you’ll have to protect them from being evaluated by quoting them. Thus, to be on the safe side, you should actually write:

```
> let r = {'x'=>5, 'y'=>12};
> r!'y; r!!['y','x'];
12
{12,5}
```

It’s also possible to use strings as keys instead, which may actually be more convenient in some cases:

```
> let r = {"x"=>5, "y"=>12};
> keys r; vals r;
{"x","y"}
{5,12}
> update r "y" (r!"y"+1);
{"x"=>5,"y"=>13}
```

You can also mix strings and symbols as keys in the same record (but note that strings and symbols are always distinct, so `y` and `"y"` are really two different keys here):

```
> insert r (y=>99);
{"x"=>5,"y"=>12,y=>99}
```

As records are in fact just special kinds of matrices, the standard matrix operations can be used on record values as well. For instance, the matrix constructor provides an alternative way to quickly augment a record with a collection of new key=>value associations:

```
> let r = {x=>5, y=>12};
> let r = {r, x=>7, z=>3}; r;
```

```
{x=>5,y=>12,x=>7,z=>3}
> r!x, r!z;
7,3
> delete r x;
{x=>5,y=>12,z=>3}
> ans!x;
5
```

As the example shows, this may produce duplicate keys, but these are handled gracefully; indexing and updates will always work with the *last* association for a given key in the record. If necessary, you can remove duplicate entries from a record as follows; this will only keep the last association for each key:

```
> record r;
{x=>7,y=>12,z=>3}
```

In fact, the record operation not only removes duplicates, but also orders the record entries by keys. This produces a kind of normalized representation which is useful if you want to compare or combine two record values irrespective of the ordering of the fields. For instance:

```
> record {x=>5, y=>12} === record {y=>12, x=>5};
1
```

The record function can also be used to construct a normalized record directly from a list or tuple of hash pairs:

```
> record [x=>5, x=>7, y=>12];
{x=>7,y=>12}
```

Other matrix operations such as `map`, `foldl`, etc., and matrix comprehensions can be applied to records just as easily. This enables you to perform bulk updates of record data in a straightforward way. For instance, here's how you can define a function `maprec` which applies a function to all values stored in a record:

```
> maprec f = map (\(u=>v) -> u=>f v);
> maprec (*2) {x=>5,y=>12};
{x=>10,y=>24}
```

Another example: The following `ziprec` function collects pairs of values stored under common keys in two records (we also normalize the result here so that duplicate keys are always removed):

```
> ziprec x y = record {u=>(x!u,y!u) | u = keys x; member y u};
> ziprec {a=>3,x=>5,y=>12} {x=>10,y=>24,z=>7};
{x=>(5,10),y=>(12,24)}
```

Thus the full power of generic matrix operations is available for records, which turns them into a very versatile data structure, much more powerful than records in conventional programming languages which are usually limited to constructing records and accessing or modifying their components. Note that since the values stored in records can be arbitrary Pure values, you can also have mutable records by making use of Pure's expression refer-

ences (see *Expression References* in the library manual). And of course records can be nested, too:

```
> let r = {a => {b=>1,c=>2}, b => 2};  
> r!a, r!b, r!a!b;  
{b=>1,c=>2}, 2, 1
```

6 Macros

Macros are a special type of functions to be executed as a kind of “preprocessing stage” at compile time. In Pure these are typically used to define custom special forms and to perform inlining of function calls and other simple kinds of source-level optimizations.

Whereas the macro facilities of most programming languages simply provide a kind of textual substitution mechanism, Pure macros operate on symbolic expressions and are implemented by the same kind of rewriting rules that are also used to define ordinary functions in Pure. In difference to these, macro rules start out with the keyword `def`, and only simple kinds of rules without any guards or multiple left-hand and right-hand sides are permitted.

Syntactically, a macro definition looks just like a variable or constant definition, using `def` in lieu of `let` or `const`, but they are processed in a different way. Macros are substituted into the right-hand sides of function, constant and variable definitions. All macro substitution happens before constant substitutions and the actual compilation step. Macros can be defined in terms of other macros (also recursively), and are evaluated using call by value (i.e., macro calls in macro arguments are expanded before the macro gets applied to its parameters).

6.1 Optimization Rules

Here is a simple example, showing a rule which expands saturated calls of the `succ` function (defined in the prelude) at compile time:

```
> def succ x = x+1;  
> foo x::int = succ (succ x);  
> show foo  
foo x::int = x+1+1;
```

Rules like these can be useful to help the compiler generate better code. Note that a macro may have the same name as an ordinary Pure function, which is essential if you want to optimize calls to an existing function, as in the previous example. (Just like ordinary functions, the number of parameters in each rule for a given macro must be the same, but a macro may have a different number of arguments than the corresponding function.)

A somewhat more practical example is the following rule from the prelude, which eliminates saturated instances of the right-associative function application operator:

```
def f $ x = f x;
```

Like in Haskell, this low-priority operator is handy to write cascading function calls. With the above macro rule, these will be “inlined” as ordinary function applications automatically. Example:

```
> foo x = bar $ bar $ 2*x;
> show foo
foo x = bar (bar (2*x));
```

Here are two slightly more tricky rules from the prelude, which optimize the case of “throw-away” list comprehensions. This is useful if a list comprehension is evaluated solely for its side effects:

```
def void (catmap f x) = do f x;
def void (listmap f x) = do f x;
```

Note that the `void` function simply throws away its argument and returns `()` instead. The `do` function applies a function to every member of a list (like `map`), but throws away all intermediate results and just returns `()`, which is much more efficient if you don’t need those results anyway. These are both defined in the prelude.

Before we delve into this example, a few remarks are in order about the way list comprehensions are implemented in Pure. As already mentioned, list comprehensions are just syntactic sugar; the compiler immediately transforms them to an equivalent expression involving only lambdas and a few other list operations. Note that list comprehensions are essentially equivalent to piles of nested lambdas, filters and maps, but for various reasons they are actually implemented using two special helper operations, `catmap` and `listmap`.

The `catmap` operation combines `map` and `cat`; this is needed, in particular, to accumulate the results of nested generators, such as `[i, j | i = 1..n; j = 1..m]`. The same operation is also used to implement filter clauses, you can see this below in the examples. However, for efficiency simple generators like `[2*i | i = 1..n]` are translated to a `listmap` instead (which is basically just `map`, but works with different aggregate types, so that list comprehensions can draw values from aggregates other than lists, such as matrices).

Now let’s see how the rules above transform a list comprehension if we “voidify” it:

```
> using system;
> f = [printf "%g\n" (2^x+1) | x=1..5; x mod 2];
> g = void [printf "%g\n" (2^x+1) | x=1..5; x mod 2];
> show f g
f = catmap (\x -> if x mod 2 then [printf "%g\n" (2^x+1)] else []) (1..5);
g = do (\x -> if x mod 2 then [printf "%g\n" (2^x+1)] else []) (1..5);
```

Ok, so the `catmap` got replaced with a `do` which is just what we need to make this code go essentially as fast as a `for` loop in conventional programming languages (up to constant factors, of course). Here’s how it looks like when we run the `g` function:

```
> g;
3
9
33
()
```


It's not all roses, however, since the above macro rules will only get rid of the outermost `catmap` if the list comprehension binds multiple variables:

```
> u = void [puts $ str (x,y) | x=1..2; y=1..3];
> show u
u = do (\x -> listmap (\y -> puts (str (x,y))) (1..3)) (1..2);
```

If you're bothered by this, you'll have to apply `void` recursively, creating a nested list comprehension which expands to a nested `do`:

```
> v = void [void [puts $ str (x,y) | y=1..3] | x=1..2];
> show v
v = do (\x -> do (\y -> puts (str (x,y))) (1..3)) (1..2);
```

(It would be nice to have this handled automatically, but the left-hand side of a macro definition must be a simple expression, and thus it's not possible to write a macro which descends recursively into the lambda argument of `catmap`.)

6.2 Recursive Macros

Macros can also be recursive, in which case they usually consist of multiple rules and make use of pattern-matching like ordinary function definitions. As a simple example, let's implement a Pure version of Lisp's `quasiquote` which allows you to create a quoted expression from a "template" while substituting variable parts of the template. (For the sake of brevity, our definition is somewhat simplified and does not cover some corner cases. See the Pure distribution for a full version of this example.)

```
def quasiquote (unquote x)      = x;
def quasiquote (f@_ (splice x)) = foldl ($) (quasiquote f) x;
def quasiquote (f@_ x)          = quasiquote f (quasiquote x);
def quasiquote x                 = quote x;
```

(Note the `f@_`, which is an anonymous "as" pattern forcing the compiler to recognize `f` as a function variable, rather than a literal function symbol. See [Head = Function](#) in the [Caveats and Notes](#) section for an explanation of this trick.)

The first rule above takes care of "unquoting" embedded subterms. The second rule "splices" an argument list into an enclosing function application. The third rule recurses into subterms of a function application, and the fourth and last rule takes care of quoting the "atomic" subterms. Note that `unquote` and `splice` themselves are just passive constructor symbols, the real work is done by `quasiquote`, using `foldl` at runtime to actually perform the splicing. (Putting off the splicing until runtime makes it possible to splice argument lists computed at runtime.)

If we want, we can also add some syntactic sugar for Lisp weenies. (Note that we cannot have `'`, `'` for unquoting, so we use `'`, `'` instead.)

```
prefix 9 ' , $ , @ ;
def 'x = quasiquote x; def , $x = unquote x; def , @x = splice x;
```

Examples:

```
> '(2*42+2^12);
2*42+2^12
> '(2*42+,$(2^12));
2*42+4096.0
> 'foo 1 2 (,@'[2/3,3/4]) (5/6);
foo 1 2 (2/3) (3/4) (5/6)
> 'foo 1 2 (,@'args) (5/6) when args = '[2/3,3/4] end;
foo 1 2 (2/3) (3/4) (5/6)
```

We mention in passing here that, technically, Pure macros are just as powerful as (unconditional) term rewriting systems and thus they are Turing-complete. This implies that a badly written macro may well send the Pure compiler into an infinite recursion, which results in a stack overflow at compile time. See the [Caveats and Notes](#) section for information on how to deal with these by setting the PURE_STACK environment variable.

6.3 User-Defined Special Forms

The quasiquote macro in the preceding subsection also provides an example of how you can use macros to define your own special forms. This works because the actual evaluation of macro arguments is put off until runtime, and thus we can safely pass them to built-in special forms and other constructs which defer their evaluation at *runtime*. In fact, the right-hand side of a macro rule may be an arbitrary Pure expression involving conditional expressions, lambdas, binding clauses, etc. These are never evaluated during macro substitution, they just become part of the macro expansion (after substituting the macro parameters).

Here is another useful example of a user-defined special form, the macro `timex` which employs the system function `clock` to report the cpu time in seconds needed to evaluate a given expression, along with the computed result:

```
> using system;
> def timex x = (clock-t0)/CLOCKS_PER_SEC,y when t0 = clock; y = x end;
> sum = foldl (+) 0L;
> timex $ sum (1L..100000L);
0.43,5000050000L
```

Note that the above definition of `timex` wouldn't work as an ordinary function definition, since by virtue of Pure's basic eager evaluation strategy the `x` parameter would have been evaluated already before it is passed to `timex`, making `timex` always return a zero time value. Try it!

Here's yet another example, which is handy if you need to trace function calls. (As of Pure 0.22, the interpreter now has its own built-in debugging facility, see [Debugging](#). However, the following macro allows you to trace functions using your own custom output format, and may thus be useful in situations where the built-in debugger is not appropriate.)

```
using system;
def trace f x y = printf "** exit %s: %s -> %s\n" (str f,str x,str y) $$ y
when y = printf "** call %s: %s\n: " (str f,str x) $$ gets $$ y end;
```

This macro is invoked with the function to be traced, the arguments (or whatever you want to be printed as additional debugging information) and the actual function call as parameters. (This is a rather simplistic version, which just prints a prompt on function entry and the final reduction after the call. You can easily make this as elaborate as you like. E.g., you might want to keep track of recursive levels and profiling information, add various interactive commands to selectively enable and disable tracing during the evaluation, etc.)

We can still make this a bit more convenient by introducing the following ordinary function definition:

```
trace f x = trace f x (f x);
```

This lets us patch up a call to trace a given function, as shown below, without having to change the definition of the function at all. This trick only works with global functions; for local functions you'll have to add an explicit call of the `trace` macro to the local definition yourself. Also note that the definition above only works with functions taking a single parameter; see the `trace.pure` example in the distribution for the full version which can deal with any number of arguments.

```
// Uncomment this line to trace calls to the 'fact' function.
def fact n = trace fact n;
// Sample function to be traced.
fact n = if n>0 then n*fact(n-1) else 1;
```

Here's a trace of the `fact` function obtained in this fashion (hit carriage return after each ':' prompt to proceed with the computation):

```
> fact 2;
** call fact: 2
:
** call fact: 1
:
** call fact: 0
:
** exit fact: 0 -> 1
** exit fact: 1 -> 1
** exit fact: 2 -> 2
2
```

Note that by just removing the macro definition for `fact` above, you can make the function run untraced as usual again. This scheme is quite flexible, the only real drawback is that you have to explicitly add some code for each function you want to trace.

6.4 Macro Hygiene

Pure macros are lexically scoped, i.e., the binding of symbols in the right-hand-side of a macro definition is determined statically by the text of the definition, and macro parameter substitution also takes into account binding constructs, such as `with` and `when` clauses, in the right-hand side of the definition. Macro facilities with these pleasant properties are also known as **hygienic macros**. They are not susceptible to so-called "name capture," which

makes macros in less sophisticated languages bug-ridden and hard to use. (This is explained in more detail in the [Hygienic Macros](#) section.)

Pure macros also have their limitations. Specifically, the left-hand side of a macro rule must be a simple expression, just like in ordinary function definitions. This restricts the kinds of expressions which can be rewritten by a macro. But Pure macros are certainly powerful enough for most common preprocessing purposes, while still being robust and easy to use.

7 Declarations

Pure is a very terse language by design. Usually you don't declare much stuff, you just define it and be done with it. However, there are a few toplevel constructs which let you declare symbols with special attributes and manage programs consisting of several source modules:

- symbol declarations determine “scope” and “fixity” of a symbol;
- extern declarations specify external C functions (described in the [C Interface](#) section);
- using clauses let you include other scripts in a Pure script;
- namespace declarations let you avoid name clashes and thereby make it easier to manage large programs consisting of many separate modules.

7.1 Symbol Declarations

Scope declarations take the following form:

```
public symbol ...;  
private symbol ...;
```

This declares the listed symbols as public or private, respectively. Each symbol must either be an identifier or a sequence of punctuation characters. The latter kind of symbols *must* always be declared before use, whereas ordinary identifiers can be used without a prior declaration in which case they are declared implicitly and default to public scope, meaning that they are visible everywhere in a program. An explicit public declaration of ordinary identifiers is thus rarely needed (unless you want to declare symbols as members of a specific namespace, see [Namespaces](#) below). Symbols can also be declared private, meaning that the symbol is visible only in the namespace it belongs to. This is explained in more detail under [Private Symbols](#) in the [Namespaces](#) section below.

Note that to declare several symbols in a single declaration, you can list them all with white-space in between. The same syntax applies to the other types of symbol declarations discussed below. (Commas are *not* allowed as delimiters here, as they may occur as legal symbol constituents in the list of symbols.) The `public` and `private` keywords can also be used as a prefix in any of the special symbol declarations discussed below, to specify the scope of the declared symbols (if the scope prefix is omitted, it defaults to `public`).

The following “fixity” declarations are available for introducing special operator and constant symbols. This changes the way that these symbols are parsed and thus provides you with a limited means to extend the Pure language at the lexical and syntactical level.

Operator declarations: `infix level symbol ...`; Pure provides you with a theoretically unlimited number of different precedence levels for user-defined infix, prefix and postfix operators. Precedence levels are numbered starting at 0; larger numbers indicate higher precedence. (For practical reasons, the current implementation does require that precedence numbers can be encoded as 24 bit unsigned machine integers, giving you a range from 0 to 16777215, but this should be large enough to incur no real limitations on applications. Also, the operator declarations in the prelude have been set up to leave enough “space” between the “standard” levels so that you can easily sneak in new operator symbols at low, high or intermediate precedences.)

On each precedence level, you can declare (in order of increasing precedence) `infix` (binary non-associative), `infixl` (binary left-associative), `infixr` (binary right-associative), `prefix` (unary prefix) and `postfix` (unary postfix) operators. For instance, here is a typical excerpt from the prelude (the full table can be found in the [Prelude](#) section of the *Pure Library Manual*):

```
infix 1800 < > <= >= == ~= ;
infixl 2200 + - ;
infixl 2300 * / div mod ;
infixr 2500 ^ ;
prefix 2600 # ;
```

Instead of denoting the precedence by an explicit integer value, you can also specify an existing operator symbol enclosed in parentheses. Thus the following declaration gives the `++` operator the same precedence as `+`:

```
infixl (+) ++ ;
```

The given symbol may be of a different fixity than the declaration, but it must have a proper precedence level (i.e., it must be an infix, prefix or postfix symbol). E.g., the following declaration gives `^^` the same precedence level as the `infix ^` symbol, but turns it into a postfix operator:

```
postfix (^) ^^ ;
```

Outfix symbol declarations: `outfix left right ...`; Pure also provides unary outfix operators, which work like in Wm Leler’s constraint programming language [Bertrand](#). Outfix operators let you define your own bracket structures. The operators must be given as pairs of matching left and right symbols (which must be distinct). For instance:

```
outfix |: :| BEGIN END;
```

After this declaration you can write bracketed expressions like `|:x:|` or `BEGIN foo, bar END`. These are always at the highest precedence level (i.e., syntactically they work like parenthesized expressions). Just like other operators, you can turn outfix symbols

into ordinary functions by enclosing them in parentheses, but you have to specify the symbols in matching pairs, such as (BEGIN END).

Constant symbol declarations: `nonfix symbol ...`; Pure also has a notation for “nullary” operators, i.e., “operators without operands”, which are used to denote special constants. These are introduced using a `nonfix` declaration, e.g.:

```
nonfix red green blue;
```

Syntactically, these work just like ordinary identifiers, so they may stand wherever an identifier is allowed (no parentheses are required to “escape” them). The difference to ordinary identifiers is that `nonfix` symbols are always interpreted as literals, even if they occur in a variable position on the left-hand side of a rule. So, with the above declaration, you can write something like:

```
> foo x = case x of red = green; green = blue; blue = red end;  
> map foo [red,green,blue];  
[green,blue,red]
```

Thus `nonfix` symbols are pretty much like nullary constructor symbols in languages like Haskell. Non-fixity is just a syntactic attribute, however. Pure doesn’t enforce that such values are really “constant”, so you can still write a “constructor equation” like the following:

```
> red = blue;  
> map foo [red,green,blue];  
[blue,blue,blue]
```

Examples for all types of symbol declarations can be found in the prelude which declares a bunch of standard (arithmetic, relational, logical) operator symbols as well as the list and pair constructors ‘:’ and ‘,’ and a few `nonfix` symbols (mostly for denoting different kinds of exceptions).

One final thing worth noting here is that unary minus plays a special role in the syntax. Like in Haskell and following mathematical tradition, unary minus is the only prefix operator symbol which is also used as an infix operator, and is always on the same precedence level as binary minus, whose precedence may be chosen freely in the prelude. (The minus operator is the only symbol which gets that special treatment; all other operators must have distinct lexical representations.) Thus, with the standard prelude, `-x+y` will be parsed as `(-x)+y`, whereas `-x*y` is the same as `-(x*y)`. Also note that the notation `(-)` always denotes the binary minus operator; the unary minus operation can be denoted using the built-in `neg` function.

7.2 Modules and Imports

While Pure doesn’t offer separate compilation, the `using` declaration provides a simple but effective way to assemble a Pure program from several source modules. It takes the following form (note that in contrast to symbol declarations, the comma is used as a delimiter symbol here):

```
using name, ...;
```

This causes each given script to be included in the Pure program at the given point (if it wasn't already included before), which makes available all the definitions of the included script in your program. Note that each included script is loaded only *once*, when the first using clause for the script is encountered. Nested imports are allowed, i.e., an imported module may itself import other modules, etc. A Pure program then basically is the concatenation of all the source modules given as command line arguments, with other modules listed in using clauses inserted at the corresponding source locations.

(The using clause also has an alternative form which allows dynamic libraries and LLVM bitcode modules to be loaded, this will be discussed in the [C Interface](#) section.)

For instance, the following declaration causes the math.pure script from the standard library to be included in your program:

```
using math;
```

You can also import multiple scripts in one go:

```
using array, dict, set;
```

Moreover, Pure provides a notation for qualified module names which can be used to denote scripts located in specific package directories, e.g.:

```
using examples::libor::bits;
```

In fact this is equivalent to the following using clause which spells out the real filename of the script between double quotes (the .pure suffix can also be omitted in which case it is added automatically):

```
using "examples/libor/bits.pure";
```

Both notations can be used interchangeably; the former is usually more convenient, but the latter allows you to denote scripts whose names aren't valid Pure identifiers.

Script identifiers are translated to the corresponding filenames by replacing the '::' symbol with the pathname separator '/' and tacking on the '.pure' suffix. The following table illustrates this with a few examples.

Script identifier	Filename
math	"math.pure"
examples::libor::bits	"examples/libor/bits.pure"
::pure::examples::hello	"/pure/examples/hello.pure"

Note the last example, which shows how an absolute pathname can be denoted using a qualifier starting with '::'.

Unless an absolute pathname is given, the interpreter performs a search to locate the script. The search algorithm considers the following directories in the given order:

- the directory of the current script, which is the directory of the script containing the `using` clause, or the current working directory if the clause was read from standard input (as is the case, e.g., in an interactive session);
- the directories named in `-I` options on the command line (in the given order);
- the colon-separated list of directories in the `PURE_INCLUDE` environment variable (in the given order);
- finally the directory named by the `PURELIB` environment variable.

Note that the current working directory is not searched by default (unless the `using` clause is read from standard input), but of course you can force this by adding the option `-I.` to the command line, or by including `'.'` in the `PURE_INCLUDE` variable.

The directory of the current script (the first item above) can be skipped by specifying the script to be loaded as a filename in double quotes, prefixed with the special `sys:` tag. The search then starts with the “system” directories (`-I`, `PURE_INCLUDE` and `PURELIB`) instead. This is useful, e.g., if you want to provide your own custom version of a standard library script which in turn imports that library script. For instance, a custom version of `math.pure` might employ the following `using` clause to load the `math.pure` script from the Pure library:

```
using "sys:math";  
// custom definitions go here  
log2 x = ln x/ln 2;
```

The interpreter compares script names (to determine whether two scripts are actually the same) by using the *canonicalized* full pathname of the script, following symbolic links to the destination file (albeit only one level). Thus different scripts with the same basename, such as `foo/Utils.pure` and `bar/Utils.pure` can both be included in the same program (unless they link to the same file).

More precisely, canonicalizing a pathname involves the following steps:

- relative pathnames are expanded to absolute ones, using the search rules discussed above;
- the directory part of the pathname is normalized to the form returned by the `getcwd` system call;
- the `“.pure”` suffix is added if needed;
- if the resulting script name is actually a symbolic link, the interpreter follows that link to its destination, albeit only one level. (This is only done on Unix-like systems.)

The directory of the canonicalized pathname is also used when searching other scripts included in a script. This makes it possible to have an executable script with a shebang line in its own directory, which is then executed via a symbolic link placed on the system `PATH`. In this case the script search performed in `using` clauses will use the real script directory and thus other required scripts can be located there. This is the recommended practice for installing standalone Pure applications in source form which are to be run directly from the shell.

7.3 Namespaces

To facilitate modular development, Pure also provides namespaces as a means to avoid name clashes between symbols, and to keep the global namespace tidy and clean. Namespaces serve as containers holding groups of related identifiers and other symbols. Inside each namespace, symbols must be unique, but the same symbol may be used to denote different objects (variables, functions, etc.) in different namespaces. (Pure’s namespace system was heavily inspired by C++ and works in a very similar fashion. So if you know C++ you should feel right at home and skimming this section to pick up Pure’s syntax of the namespace constructs should be enough to start using it.)

The global namespace is always available. By default, new symbols are created in this namespace, which is also called the **default namespace**. Additional namespaces can be created with the namespace declaration, which also switches to the given namespace (makes it the *current* namespace), so that new symbols are then created in that namespace rather than the default one. The current namespace also applies to all kinds of symbol declarations, including operator and constant symbol declarations, as well as extern declarations (the latter are described in the [C Interface](#) section).

The basic form of the namespace declaration has the following syntax (there’s also a “scoped” form of the namespace declaration which will be discussed in [Scoped Namespaces](#) at the end of this section):

```
namespace name;  
// declarations and definitions in namespace 'name'  
namespace;
```

The second form switches back to the default namespace. For instance, in order to define two symbols with the same print name `foo` in two different namespaces `foo` and `bar`, you can write:

```
namespace foo;  
foo x = x+1;  
namespace bar;  
foo x = x-1;  
namespace;
```

We can now refer to the symbols we just defined using **qualified symbols** of the form `namespace::symbol`:

```
> foo::foo 99;  
100  
> bar::foo 99;  
98
```

This avoids any potential name clashes, since the qualified identifier notation always makes it clear which namespace the given identifier belongs to.

A namespace can be “reopened” at any time to add new symbols and definitions to it. This allows namespaces to be created that span several source modules. You can also create several different namespaces in the same module.

Similar to the `using` declaration, a namespace declaration accepts either identifiers or double-quoted strings as namespace names. E.g., the following two declarations are equivalent:

```
namespace foo;
namespace "foo";
```

The latter form also allows more descriptive labels which aren't identifiers, e.g.:

```
namespace "Private stuff, keep out!";
```

Note that the namespace prefix in a qualified identifier must be a legal identifier, so it isn't possible to access symbols in namespaces with such descriptive labels in a direct fashion. The only way to get at the symbols in this case is to use a namespace or `using namespace` declaration (for the latter see [Using Namespaces](#) below).

Using Namespaces

Since it is rather inconvenient if you always have to write identifiers in their qualified form outside of their "home" namespace, Pure allows you to specify a list of *search* namespaces which are used to look up symbols not in the default or the current namespace. This is done with the `using namespace` declaration, which takes the following form:

```
using namespace name1, name2, ...;
// ...
using namespace;
```

(As with namespace declarations, the second form without any namespace arguments gets you back to the default empty list of search namespaces.)

For instance, consider this example:

```
namespace foo;
foo x = x+1;
namespace bar;
foo x = x-1;
bar x = x+1;
namespace;
```

The symbols in these namespaces can be accessed unqualified as follows:

```
> using namespace foo;
> foo 99;
100
> using namespace bar;
> foo 99;
98
> bar 99;
100
```

This method is often to be preferred over opening a namespace with the namespace declaration, since `using namespace` only gives you "read access" to the imported symbols, so you

can't accidentally mess up the definitions of the namespace you're using. Another advantage is that the `using namespace` declaration also lets you search multiple namespaces at once:

```
using namespace foo, bar;
```

Be warned, however, that this brings up the very same issue of name clashes again:

```
> using namespace foo, bar;
> foo 99;
<stdin>, line 15: symbol 'foo' is ambiguous here
```

In such a case you'll have to resort to using namespace qualifiers again, in order to resolve the name clash:

```
> foo::foo 99;
100
```

To avoid this kind of mishap, you can also selectively import just a few symbols from a namespace instead. This can be done with a declaration of the following form:

```
using namespace name1 ( sym1 sym2 ... ), name2 ... ;
```

As indicated, the symbols to be imported can optionally be placed as a whitespace-delimited list inside parentheses, following the corresponding namespace name. For instance:

```
> using namespace foo, bar (bar);
> foo 99;
100
> bar 99;
100
> bar::foo 99;
98
```

Note that now we have no clash on the `foo` symbol any more, because we restricted the import from the `bar` namespace to the `bar` symbol, so that `bar::foo` has to be denoted with a qualified symbol now.

Symbol Lookup and Creation

Pure's rules for looking up and creating symbols are fairly straightforward and akin to those in other languages featuring namespaces. However, there are some intricacies involved, because the rewriting rule format of definitions allows "referential" use of symbols not only in the "body" (right-hand side) of a definition, but also in the left-hand side patterns. We discuss this in detail below.

The compiler searches for symbols first in the current namespace (if any), then in the currently active search namespaces (if any), and finally in the default (i.e., the global) namespace, in that order. This automatic lookup can be bypassed by using an *absolute* namespace qualifier of the form `::foo::bar`. In particular, `::bar` always denotes the symbol `bar` in the

default namespace, while `::foo::bar` denotes the symbol `bar` in the `foo` namespace. (Normally, the latter kind of notation is only needed if you have to deal with nested namespaces, see [Hierarchical Namespaces](#) below.)

If no existing symbol is found, a new symbol is created automatically, by implicitly declaring a public symbol with default attributes. New *unqualified* symbols are always created in the current namespace, while new *qualified* symbols are created in the namespace given by the namespace prefix of the symbol. However, note that in the latter case the compiler always checks that the given namespace prefix matches the current namespace:

```
> namespace foo;
> namespace;
> foo::bar x = 1/x;
<stdin>, line 3: undeclared symbol 'foo::bar'
```

Thus it's only possible to introduce a new symbol in a given namespace if that namespace is the current one. These error messages are somewhat annoying, but they provide at least some protection against typos and other silly mistakes and prevent you from accidentally clobbering the contents of other namespaces. To make these errors go away it's enough to just declare the symbols in their proper namespaces.

New symbols are also created if a global unqualified (and yet undeclared) symbol is being “defined” in a rewriting rule or `let/const` definition, even if a symbol with the same print name from another namespace is already visible in the current scope. To distinguish “defining” from “referring” uses of a global symbol, Pure uses the following (purely syntactic) notions:

- A **defining occurrence** of a global *function* or *macro symbol* is any occurrence of the symbol as the *head symbol* on the left-hand side of a rewriting rule.
- A **defining occurrence** of a global *variable* or *constant symbol* is any occurrence of the symbol in a *variable position* (as given by the “head = function” rule, cf. [Parameters in Equations](#)) on the left-hand side of a `let` or `const` definition.
- All other occurrences of global symbols on the left-hand side, as well as *all* symbol occurrences on the right-hand side of a definition are **referring occurrences**.

The following example illustrates these notions:

```
namespace foo;
bar (bar x) = bar x;
let x,y = 1,2;
namespace;
```

Here, the first occurrence of `bar` on the left-hand side `bar (bar x)` of the first rule is a *defining* occurrence, as are the occurrences of `x` and `y` on the left-hand side of the `let` definition. Hence these symbols are created as new symbols in the namespace `foo`. On the other hand, the other occurrences of `bar` in the first rule, as well as the `,` symbol on the left-hand side of the `let` definition are *referring* occurrences. In the former case, `bar` refers to the `bar` symbol defined by the rule, while in the latter case the `,` operator is actually declared in the prelude and thus imported from the global namespace.

As an additional safety measure against missing or mistyped symbols, the interpreter provides the option `-w` (see [Options](#)) to check your scripts for non-defining uses of undeclared unqualified function symbols. For instance:

```
$ pure -w
> puts "bla"; // missing import of system module
<stdin>, line 1: warning: implicit declaration of 'puts'
puts "bla"
```

For legitimate uses (such as forward uses of a symbol which is defined later), you can make these warnings go away by declaring the symbol before using it.

Note that special operator (and nonfix) symbols *always* require an explicit declaration. This works as already discussed in the [Symbol Declarations](#) section, except that you first switch to the appropriate namespace before declaring the symbols. For instance, here is how you can create a new `+` operation which multiplies its operands rather than adding them:

```
> namespace my;
> infixl 2200 +;
> x+y = x*y;
> 5+7;
35
```

Note that the new `+` operation really belongs to the namespace we created. The `+` operation in the default namespace works as before, and in fact you can use qualified symbols to pick the version that you need:

```
> namespace;
> 5+7;
12
> 5 ::+ 7;
12
> 5 my::+ 7;
35
```

Here's what you get if you happen to forget the declaration of the `+` operator:

```
> namespace my;
> x+y = x*y;
<stdin>, line 2: infixl symbol '+' was not declared in this namespace
```

Thus the compiler will never create a new instance of an operator symbol on the fly, an explicit declaration is always needed in such cases.

Note that if you *really* wanted to redefine the global `+` operator, you can do this even while the `my` namespace is current. You just have to use a qualified identifier in this case, as follows:

```
> namespace my;
> x ::+ y = x*y;
> a+b;
a*b
```

This should rarely be necessary (in the above example you might just as well enter this rule while in the global namespace), but it can be useful in some circumstances. Specifically, you

might want to “overload” a global function or operator with a definition that makes use of private symbols of a namespace (which are only visible inside that namespace; see [Private Symbols](#) below). For instance:

```
> namespace my;
> private bar;
> bar x y = x*y;
> x ::+ y = bar x y;
> a+b;
a*b
```

(The above is a rather contrived example, since the very same functionality can be accomplished much easier, but there are some situations where this approach is necessary.)

Private Symbols

Pure also allows you to have private symbols, as a means to hide away internal operations which shouldn’t be accessed directly outside the namespace in which they are declared. The scope of a private symbol is confined to its namespace, i.e., the symbol is only visible when its “home” namespace is current. Symbols are declared private by using the `private` keyword in the symbol declaration:

```
> namespace secret;
> private baz;
> // 'baz' is a private symbol in namespace 'secret' here
> baz x = 2*x;
> // you can use 'baz' just like any other symbol here
> baz 99;
198
> namespace;
```

Note that, at this point, `secret::baz` is now invisible, even if you have `secret` in the search namespace list:

```
> using namespace secret;
> // this actually creates a 'baz' symbol in the default namespace:
> baz 99;
baz 99
> secret::baz 99;
<stdin>, line 27: symbol 'secret::baz' is private here
```

The only way to bring the symbol back into scope is to make the `secret` namespace current again:

```
> namespace secret;
> baz 99;
198
> secret::baz 99;
198
```

Hierarchical Namespaces

Namespace identifiers can themselves be qualified identifiers in Pure, which enables you to introduce a hierarchy of namespaces. This is useful, e.g., to group related namespaces together under a common “umbrella” namespace:

```
namespace my;
namespace my::old;
foo x = x+1;
namespace my::new;
foo x = x-1;
```

Note that the namespace `my`, which serves as the parent namespace, must be created before the `my::old` and `my::new` namespaces, even if it does not contain any symbols of its own. After these declarations, the `my::old` and `my::new` namespaces are part of the `my` namespace and will be considered in name lookup accordingly, so that you can write:

```
> using namespace my;
> old::foo 99;
100
> new::foo 99;
98
```

This works pretty much like a hierarchy of directories and files, where the namespaces play the role of the directories (with the default namespace as the root directory), the symbols in each namespace correspond to the files in a directory, and the `using namespace` declaration functions similar to the shell’s `PATH` variable.

Sometimes it is necessary to tell the compiler to use a symbol in a specific namespace, bypassing the usual symbol lookup mechanism. For instance, suppose that we introduce another *global* `old` namespace and define yet another version of `foo` in that namespace:

```
namespace old;
foo x = 2*x;
namespace;
```

Now, if we want to access that function, with `my` still active as the search namespace, we cannot simply refer to the new function as `old::foo`, since this name will resolve to `my::old::foo` instead. As a remedy, the compiler accepts an **absolute** qualified identifier of the form `::old::foo`. This bypasses name lookup and thus always yields exactly the symbol in the given namespace (if it exists; as mentioned previously, the compiler will complain about an undeclared symbol otherwise):

```
> old::foo 99;
100
> ::old::foo 99;
198
```

Also note that, as a special case of the absolute qualifier notation, `::foo` always denotes the symbol `foo` in the default namespace.

Scoped Namespaces

Pure also provides an alternative scoped namespace construct which makes nested namespace definitions more convenient. This construct takes the following form:

```
namespace name with ... end;
```

The part between `with` and `end` may contain arbitrary declarations and definitions, using the same syntax as the toplevel. These are processed in the context of the given namespace, as if you had written:

```
namespace name;  
...  
namespace;
```

However, the scoped namespace construct always returns you to the namespace which was active before, and thus these declarations may be nested:

```
namespace foo with  
  // declarations and definitions in namespace foo  
  namespace bar with  
    // declarations and definitions in namespace bar  
  end;  
  // more declarations and definitions in namespace foo  
end;
```

Note that this kind of nesting does not necessarily imply a namespace hierarchy as discussed in [Hierarchical Namespaces](#). However, you can achieve this by using the appropriate qualified namespace names:

```
namespace foo with  
  // ...  
  namespace foo::bar with  
    // ...  
  end;  
  // ...  
end;
```

Another special feature of the scoped namespace construct is that using namespace declarations are always local to the current namespace scope (and other nested namespace scopes inside it). Thus the previous setting is restored at the end of each scope:

```
using namespace foo;  
namespace foo with  
  // still using namespace foo here  
  using namespace bar;  
  // now using namespace bar  
  namespace bar with  
    // still using namespace bar here  
    using namespace foo;  
    // now using namespace foo  
  end;
```



```
// back to using namespace bar  
end;  
// back to using namespace foo at toplevel
```

Finally, here's a more concrete example which shows how scoped namespaces might be used to declare two namespaces and populate them with various functions and operators:

```
namespace foo with  
  infixr (::^) ^;  
  foo x = x+1;  
  bar x = x-1;  
  x^y = 2*x+y;  
end;  
  
namespace bar with  
  outfix <: :>;  
  foo x = x+2;  
  bar x = x-2;  
end;  
  
using namespace foo(^ foo), bar(bar <: :>);  
  
// namespace foo  
foo x;  
x^y;  
  
// namespace bar  
bar x;  
<: x,y :>;
```

Pure's namespaces can thus be used pretty much like "modules" or "packages" in languages like Ada or Modula-2. They provide a structured way to describe program components offering collections of related data and operations, which can be brought into scope in a controlled way by making judicious use of `using namespace` declarations. They also provide an abstraction barrier, since internal operations and data structures can be hidden away employing private symbols.

Please note that these facilities are not Pure's main focus and thus they are somewhat limited compared to programming languages specifically designed for big projects and large teams of developers. Nevertheless they should be useful if your programs grow beyond a small collection of simple source modules, and enable you to manage most Pure projects with ease.

8 Exception Handling

Pure also offers a useful exception handling facility. To raise an exception, you just invoke the built-in function `throw` with the value to be thrown as the argument. To catch an exception, you use the built-in special form `catch` with the exception handler (a function to be applied to the exception value) as the first and the expression to be evaluated as the second (call-by-name) argument. For instance:

```
> catch error (throw hello_world);  
error hello_world
```

Exceptions are also generated by the runtime system if the program runs out of stack space, when a guard does not evaluate to a truth value, and when the subject term fails to match the pattern in a pattern-matching lambda abstraction, or a `let`, `case` or `when` construct. These types of exceptions are reported using the symbols `stack_fault`, `failed_cond` and `failed_match`, respectively, which are declared as constant symbols in the standard prelude. You can use `catch` to handle these kinds of exceptions just like any other. For instance:

```
> fact n = if n>0 then n*fact(n-1) else 1;  
> catch error (fact foo);  
error failed_cond  
> catch error (fact 100000);  
error stack_fault
```

(You'll only get the latter kind of exception if the interpreter does stack checks, see the discussion of the `PURE_STACK` environment variable in the [Caveats and Notes](#) section.)

Note that unhandled exceptions are reported by the interpreter with a corresponding error message:

```
> fact foo;  
<stdin>, line 2: unhandled exception 'failed_cond' while evaluating 'fact foo'
```

Exceptions also provide a way to handle asynchronous signals. Pure's system module provides symbolic constants for common POSIX signals and also defines the operation `trap` which lets you rebind any signal to a signal exception. For instance, the following lets you handle the `SIGQUIT` signal:

```
> using system;  
> trap SIG_TRAP SIGQUIT;
```

You can also use `trap` to just ignore a signal or revert to the system's default handler (which might take different actions depending on the type of signal, see `signal(7)` for details):

```
> trap SIG_IGN SIGQUIT; // signal is ignored  
> trap SIG_DFL SIGQUIT; // reinstalls the default signal handler
```

Note that when the interpreter runs interactively, for convenience most standard termination signals (`SIGINT`, `SIGTERM`, etc.) are already set up to produce corresponding Pure exceptions of the form `signal SIG` where `SIG` is the signal number. If a script is to be run non-interactively then you'll have to do this yourself (otherwise most signals will terminate the program).

Last but not least, exceptions can also be used to implement non-local value returns. For instance, here's a variation of our `n` queens algorithm which only returns the first solution. Note the use of `throw` in the recursive search routine to bail out with a solution as soon as we found one. The value thrown there is caught in the main routine. Also note the use of `void` in the second equation of `search`. This effectively turns the list comprehension into a simple

loop which suppresses the normal list result and just returns `()` instead. Thus, if no value gets thrown then the function regularly returns with `()` to indicate that there is no solution.

```
queens n      = catch reverse (search n 1 []) with
  search n i p = throw p if i>n;
              = void [search n (i+1) ((i,j):p) | j = 1..n; safe (i,j) p];
  safe (i,j) p = ~any (check (i,j)) p;
  check (i1,j1) (i2,j2)
    = i1==i2 || j1==j2 || i1+j1==i2+j2 || i1-j1==i2-j2;
end;
```

E.g., let's compute a solution for a standard 8x8 board:

```
> queens 8;
[(1,1),(2,5),(3,8),(4,6),(5,3),(6,7),(7,2),(8,4)]
```

9 C Interface

Pure makes it very easy to call C functions (as well as functions in a number of other languages supported by the GNU compiler collection). To call an existing C function, you just need an extern declaration of the function, as described below. By these means, all functions in the standard C library and the Pure runtime are readily available to Pure scripts. Functions can also be loaded from dynamic libraries and LLVM bitcode files at runtime. In the latter case, you don't even need to write any extern declarations, the interpreter will do that for you. As of Pure 0.45, you can also add inline C/C++ and Fortran code to your Pure scripts and have the Pure interpreter compile them on the fly, provided that you have the corresponding compilers from the LLVM project installed.

In some cases you will still have to rely on big and complicated third-party and system libraries which aren't readily available in bitcode form. It goes without saying that writing all the extern declarations for such libraries can be a daunting task. Fortunately, there is a utility to help with this, by extracting the extern declarations automatically from C headers. Please see [External C Functions](#) in the [Caveats and Notes](#) section for details.

9.1 Extern Declarations

To access an existing C function in Pure, you need an extern declaration of the function, which is a simplified kind of C prototype. The function can then be called in Pure just like any other. For instance, the following commands, entered interactively in the interpreter, let you use the `sin` function from the C library (of course you could just as well put the extern declaration into a script):

```
> extern double sin(double);
> sin 0.3;
0.29552020666134
```

An extern declaration can also be prefixed with a `public/private` scope specifier:

```
private extern double sin(double);
```

Multiple prototypes can be given in one extern declaration, separating them with commas:

```
extern double sin(double), double cos(double), double tan(double);
```

For clarity, the parameter types can also be annotated with parameter names (these only serve informational purposes and are for the human reader; they are effectively treated as comments by the compiler):

```
extern double sin(double x);
```

Pointer types are indicated by following the name of the element type with one or more asterisks, as in C. For instance:

```
> extern char* strchr(char *s, int c);  
> strchr "foo bar" (ord "b");  
"bar"
```

As you can see in the previous example, some pointer types get special treatment, allowing you to pass certain kinds of Pure data (such as Pure strings as `char*` in this example). This is discussed in more detail in [C Types](#) below.

The interpreter makes sure that the parameters in a call match; if not, then by default the call is treated as a normal form expression:

```
> extern double sin(double);  
> sin 0.3;  
0.29552020666134  
> sin 0;  
sin 0
```

This gives you the opportunity to augment the external function with your own Pure equations. To make this work, you have to make sure that the extern declaration of the function comes first. For instance, we might want to extend the `sin` function with a rule to handle integers:

```
> sin x::int = sin (double x);  
> sin 0;  
0.0
```

Sometimes it is preferable to replace a C function with a wrapper function written in Pure. In such a case you can specify an **alias** under which the original C function is known to the Pure program, so that you can still call the C function from the wrapper. An alias is introduced by terminating the extern declaration with a clause of the form `= alias`. For instance:

```
> extern double sin(double) = c_sin;  
> sin x::double = c_sin x;  
> sin x::int = c_sin (double x);  
> sin 0.3; sin 0;  
0.29552020666134  
0.0
```

As an alternative, you can also declare the C function in a special namespace (cf. [Namespaces](#) in the [Declarations](#) section):

```
> namespace c;
> extern double sin(double);
> c::sin 0.3;
0.29552020666134
```

Note that the namespace qualification only affects the Pure side; the underlying C function is still called under the unqualified name as usual. The way in which such qualified externs are accessed is the same as for ordinary qualified symbols. In particular, the using namespace declaration applies as usual, and you can declare such symbols as private if needed. It is also possible to combine a namespace qualifier with an alias:

```
> namespace c;
> extern double sin(double) = mysin;
> c::mysin 0.3;
0.29552020666134
```

Different aliases of the same external function (or, equivalently, instances of the function declared in different namespaces) can be declared in slightly different ways, which makes it possible to adjust the interpretation of pointer values on the Pure side. This is particularly useful for string arguments which, as described below, may be passed both as `char*` (which implies copying and conversion to or from the system encoding) and as `void*` (which simply passes through the character pointers). For instance:

```
> extern char *strchr(char *s, int c) = foo;
> extern void *strchr(void *s, int c) = bar;
> foo "foo bar" 98; bar "foo bar" 98;
"bar"
#<pointer 0x12c2f24>
```

Also note that, as far as Pure is concerned, different aliases of an external function are really different functions. In particular, they can each have their own set of augmenting Pure equations. For instance:

```
> extern double sin(double);
> extern double sin(double) = mysin;
> sin == sin;
1
> sin == mysin;
0
> sin 1.0; mysin 1.0;
0.841470984807897
0.841470984807897
> sin x::int = sin (double x);
> sin 1; mysin 1;
0.841470984807897
mysin 1
```

9.2 C Types

As indicated in the previous section, the data types in extern declarations are either C type names or pointer types derived from these. The special `expr*` pointer type is simply passed through; this provides a means to deal with Pure data in C functions in a direct fashion. For all other C types, Pure values are “marshalled” (converted) from Pure to C when passed as arguments to C functions, and the result returned by the C function is then converted back from C to Pure. All of this is handled by the runtime system in a transparent way, of course.

Note that, to keep things simple, Pure does *not* provide any notations for C structs or function types, although it is possible to represent pointers to such objects using `void*` or some other appropriate pointer types. In practice, this simplified system should cover most kinds of calls that need to be done when interfacing to C libraries, but there are ways to work around these limitations if you need to access C structs or call back from C to Pure, see [External C Functions](#) in the [Caveats and Notes](#) section for details.

Basic C Types

Pure supports the usual range of basic C types: `void`, `bool`, `char`, `short`, `int`, `long`, `float`, `double`, and converts between these and the corresponding Pure data types (machine ints, bigints and double values) in a straightforward way.

The `void` type is only allowed in function results. It is converted to the empty tuple `()`.

Both `float` and `double` are supported as floating point types. Single precision `float` arguments and return values are converted from/to Pure’s double precision floating point numbers.

A variety of C integer types (`bool`, `char`, `short`, `int`, `long`) are provided which are converted from/to the available Pure integer types in a straightforward way. In addition, the synonyms `int8`, `int16` and `int32` are provided for `char`, `short` and `int`, respectively, and `int64` denotes 64 bit integers (a.k.a. ISO C99 `long long`). Note that `long` is equivalent to `int32` on 32 bit systems, whereas it is the same as `int64` on most 64 bit systems. To make it easier to interface to various system routines, there’s also a special `size_t` integer type which usually is 4 bytes on 32 bit and 8 bytes on 64 bit systems.

All integer parameters take both Pure ints and bigints as actual arguments; truncation or sign extension is performed as needed, so that the C interface behaves as if the argument was “cast” to the C target type. Returned integers use the smallest Pure type capable of holding the result, i.e., `int` for the C `char`, `short` and `int` types, `bigint` for `int64`.

Pure considers all integers as signed quantities, but it is possible to pass unsigned integers as well (if necessary, you can use a `bigint` to pass positive values which are too big to fit into a machine int). Also note that when an unsigned integer is returned by a C routine, which is too big to fit into the corresponding signed integer type, it will “wrap around” and become negative. In this case, depending on the target type, you can use the `ubyte`, `ushort`, `uint`, `ulong` and `uint64` functions provided by the prelude to convert the result back to an unsigned quantity.

Pointer Types

The use of pointer types is also fairly straightforward, but Pure has some special rules for the conversion of certain pointer types which make it easy to pass aggregate Pure data to and from C routines, while also following the most common idioms for pointer usage in C. The following types of pointers are recognized both as arguments and return values of C functions.

Bidirectional pointer conversions:

- `char*` is used for string arguments and return values which are converted from Pure's internal utf-8 based string representation to the system encoding and vice versa. (Thus a C routine can never modify the raw Pure string data in-place; if this is required then you'll have to pass the string argument as a `void*`, see below.)
- `void*` is for any generic pointer value, which is simply passed through unchanged. When used as an argument, you can also pass Pure strings, matrices and bigints. In this case the raw underlying data pointer (`char*` in the case of strings, `int*`, `double*` or `expr*` in the case of numeric and symbolic matrices, and the GMP type `mpz_t` in the case of bigints) is passed, which allows the data to be modified in place (with care). In particular, passing bigints as `void*` makes it possible to call most GMP integer routines directly from Pure.
- `dmatrix*`, `cmatrix*` and `imatrix*` allow you to pass numeric Pure matrices of the appropriate types (double, complex, int). Here a pointer to the underlying GSL matrix structure is passed (not just the data itself). This makes it possible to transfer GSL matrices between Pure and GSL routines in a direct fashion without any overhead. (For convenience, there are also some other pointer conversions for marshalling matrix arguments to numeric C vectors, which are described in [Pointers and Matrices](#) below.)
- `expr*` is for any kind of Pure value. A pointer to the expression node is passed to or from the C function. This type is to be used for C routines which are prepared to deal with pristine Pure data, using the corresponding functions provided by the runtime. You can find many examples of this in the standard library.

All other pointer types are simply taken at face value, allowing you to pass Pure pointer values as is, without any conversions. This also includes pointers to arbitrary named types which don't have a predefined meaning in Pure, such as `FILE*`. As of Pure 0.45, the interpreter keeps track of the actual names of all pointer types and checks (at runtime) that the types match in an external call, so that you can't accidentally get a core dump by passing, say, a `FILE*` for a `char*`. (The call will then simply fail and yield a normal form, which gives you the opportunity to hook into the function with your own Pure definitions which may supply any desired data conversions.) Typing information about pointer values is also available to Pure scripts by means of corresponding library functions, please see the [Tagged Pointers](#) section in the *Pure Library Manual* for details.

Pointers and Matrices

The following additional pointer conversions are provided to deal with Pure matrix values in arguments of C functions, i.e., on the input side. These enable you to pass Pure matrices for certain kinds of C vectors. Note that in any case, you can also simply pass a suitable plain pointer value instead. Also, these types aren't special in return values, where they will simply yield a pointer value (with the exception of `char*` which gets special treatment as explained in the previous subsection). Thus you will have to decode such results manually if needed. The standard library provides various routines to do this, please see the [String Functions](#) and [Matrix Functions](#) sections in the [Pure Library Manual](#) for details.

Numeric pointer conversions (input only):

- `char*`, `short*`, `int*`, `int64*`, `float*`, `double*` can be used to pass numeric matrices as C vectors. This kind of conversion passes just the matrix data (not the GSL matrix structure, as the `dmatrix*` et al conversions do) and does conversions between integer or floating point data of different sizes on the fly. You can either pass an int matrix as a `char*`, `short*` `int*` or `int64*` argument, or a double or complex matrix as a `float*` or `double*` argument (complex values are then represented as two separate double numbers, first the real, then the imaginary part, for each matrix element).
- `char**`, `short**`, `int**`, `int64**`, `float**`, `double**` provide yet another way to pass numeric matrix arguments. This works analogously to the numeric vector conversions above, but here a temporary C vector of pointers is passed to the C function, whose elements point to the rows of the matrix.

Argv-style conversions (input only):

- `char**` and `void**` can be used to pass argv-style vectors as arguments to C functions. In this case, the Pure argument must be a symbolic vector of strings or generic pointer values. `char**` converts the string elements to the system encoding, whereas `void**` passes through character string data and other pointers unchanged (and allows in-place modification of the data). A temporary C vector of these elements is passed to the C function, which is always NULL-terminated and can thus be used for almost any purpose which requires such argv-style vectors.

Note that in the numeric pointer conversions, the matrix data is passed “per reference” to C routines, i.e., the C function may modify the data “in place”. This is true even for target data types such as `short*` or `float**` which involve automatic conversions and hence need temporary storage. In this case the data from the temporary storage is written back to the original matrix when the function returns, to maintain the illusion of in-place modification. Temporary storage is also needed when the GSL matrix has the data in non-contiguous storage. You may want to avoid this if performance is critical, by always using “packed” matrices (see `pack` in [Matrix Functions](#)) of the appropriate types.

Pointer Examples

Let's finally have a look at some instructive examples to explain some of the trickier pointer types.

First, the matrix pointer types `dmatrix*`, `cmatrix*` and `imatrix*` can be used to pass double, complex double and int matrices to GSL functions taking pointers to the corresponding GSL types (`gsl_matrix`, `gsl_matrix_complex` and `gsl_matrix_int`) as arguments or returning them as results. (Note that there is no special marshalling of Pure's symbolic matrix type, as these aren't supported by GSL anyway.) Also note that matrices are always passed by reference. Thus, if you need to pass a matrix as an output parameter of a GSL matrix routine, you should either create a zero matrix or a copy of an existing matrix to hold the result. The prelude provides various operations for that purpose (in particular, see the `dmatrix`, `cmatrix`, `imatrix` and `pack` functions in `matrices.pure`). For instance, here is how you can quickly wrap up GSL's double matrix addition function in a way that preserves value semantics:

```
> using "lib:gsl";
> extern int gsl_matrix_add(dmatrix*, dmatrix*);
> x::matrix + y::matrix = gsl_matrix_add x y $$ x when x = pack x end;
> let x = dmatrix {1,2,3}; let y = dmatrix {2,3,2}; x; y; x+y;
{1.0,2.0,3.0}
{2.0,3.0,2.0}
{3.0,5.0,5.0}
```

Most GSL matrix routines can be wrapped in this fashion quite easily. A ready-made GSL interface providing access to all of GSL's numeric functions is in the works; please check the Pure website for details.

For convenience, it is also possible to pass any kind of numeric matrix for a `char*`, `short*`, `int*`, `int64*`, `float*` or `double*` parameter. This requires that the pointer and the matrix type match up; conversions between `char`, `short`, `int64` and `int` data and, likewise, between `float` and `double` are handled automatically, however. For instance, here is how you can call the `puts` routine from the C library with an `int` matrix encoding the string "Hello, world!" as byte values (ASCII codes):

```
> extern int puts(char*);
> puts {72,101,108,108,111,44,32,119,111,114,108,100,33,0};
Hello, world!
14
```

Pure 0.45 and later also support `char**`, `short**`, `int**`, `int64**`, `float**` and `double**` parameters which encode a matrix as a vector of row pointers instead. This kind of matrix representation is often found in audio and video processing software (where the rows of the matrix might denote different audio channels, display lines or video frames), but it's also fairly convenient to do any kind of matrix processing in C. For instance, here's how to do matrix multiplication (the naive algorithm):

```
void matmult(int n, int l, int m, double **x, double **y, double **z)
{
    int i, j, k;
    for (i = 0; i < n; i++)
        for (j = 0; j < m; j++) {
            z[i][j] = 0.0;
            for (k = 0; k < l; k++)
                z[i][j] += x[i][k]*y[k][j];
        }
}
```

```
}
```

As you can see, this multiplies a n times l matrix x with a l times m matrix y and puts the result into the n times m matrix z :

```
> extern void matmult(int, int, int, double**, double**, double**);
> let x = {0.11,0.12,0.13;0.21,0.22,0.23};
> let y = {1011.0,1012.0;1021.0,1022.0;1031.0,1032.0};
> let z = dmatrix (2,2);
> matmult 2 3 2 x y z $$ z;
{367.76,368.12;674.06,674.72}
```

Also new in Pure 0.45 is the support for passing argv-style vectors as arguments. For instance, here is how you can use fork and execvp to implement a poor man's version of the C system function. (This is UNIX-specific and doesn't do much error-checking, but you get the idea.)

```
extern int fork();
extern int execvp(char *path, char **argv);
extern int waitpid(int pid, int *status, int options);

system cmd::string = case fork of
  // child: execute the program, bail out if error
  0 = execvp "/bin/sh" {" /bin/sh", "-c", cmd} $$ exit 1;
  // parent: wait for the child and return its exit code
  pid = waitpid pid status 0 $$ status!0 >> 8
    when status = {0} end if pid>=0;
end;

system "echo Hello, world!";
system "ls -l *.pure";
system "exit 1";
```

9.3 Importing Dynamic Libraries

By default, external C functions are resolved by the LLVM runtime, which first looks for the symbol in the C library and Pure's runtime library (or the interpreter executable, if the interpreter was linked statically). Thus all C library and Pure runtime functions are readily available in Pure programs. Other functions can be provided by adding them to the runtime, or by linking them into the runtime or the interpreter executable. Better yet, you can just "dlopen" shared libraries at runtime with a special form of the using clause:

```
using "lib:libname[.ext]";
```

For instance, if you want to call the functions from library libxyz directly from Pure:

```
using "lib:libxyz";
```

After this declaration the functions from the given library will be ready to be imported into your Pure program by means of corresponding extern declarations.

Shared libraries opened with using clauses are searched for in the same way as source scripts (see section [Modules and Imports](#) above), using the -L option and the PURE_LIBRARY environment variable in place of -I and PURE_INCLUDE. If the library isn't found by these means, the interpreter will also consider other platform-specific locations searched by the dynamic linker, such as the system library directories and LD_LIBRARY_PATH on Linux. The necessary filename suffix (e.g., .so on Linux or .dll on Windows) will be supplied automatically when needed. Of course you can also specify a full pathname for the library if you prefer that. If a library file cannot be found, or if an extern declaration names a function symbol which cannot be resolved, an appropriate error message is printed.

9.4 Importing LLVM Bitcode

As of Pure 0.44, the interpreter also provides a direct way to import LLVM bitcode modules in Pure scripts. The main advantage of this method over the “plain” C interface explained above is that the bitcode loader knows all the call interfaces and generates the necessary extern declarations automatically. This is more than just a convenience, as it also eliminates at least some of the mistakes in extern declarations that may arise when importing functions manually from dynamic libraries.

LLVM bitcode is loaded in a Pure script using the following special format of the using clause:

```
using "bc:modname[.bc]";
```

(Here the bc tag indicates a bitcode file, and the default .bc bitcode filename extension is supplied automatically. Also, the bitcode file is searched for on the usual library search path.)

That's it, no explicit extern declarations are required on the Pure side. The Pure interpreter automatically creates extern declarations (in the current namespace) for all the external functions defined in the LLVM bitcode module, and generates the corresponding wrappers to make the functions callable from Pure. (This also works when batch-compiling a Pure script. In this case, the bitcode file actually gets linked into the output code, so the loaded bitcode module only needs to be present at compile time.)

By default the imported symbols will be public. You can also specify the desired scope of the symbols explicitly, by placing the public or private keyword before the module name. For instance:

```
using private "bc:modname";
```

You can also import the same bitcode module several times, possibly in different namespaces. This will not actually reload the module, but it will create aliases for the external functions in different namespaces:

```
namespace foo;  
using "bc:modname";  
namespace bar;  
using private "bc:modname";
```

You can load any number of bitcode modules along with shared libraries in a Pure script, in any order. The JIT will try to satisfy external references in modules and libraries from other loaded libraries and bitcode modules. This is deferred until the code is actually JIT-compiled, so that you can make sure beforehand that all required libraries and bitcode modules have been loaded. If the JIT fails to resolve a function, the interpreter will print its name and also raise an exception at runtime when the function is being called from other C code. (You can then run your script in the debugger to locate the external visible in Pure from which the unresolved function is called.)

Let's take a look at a concrete example to see how this actually works. Consider the following C code which defines a little function to compute the greatest common divisor of two (machine) integers:

```
int mygcd(int x, int y)
{
    if (y == 0)
        return x;
    else
        return mygcd(y, x%y);
}
```

Let's say that this code is in the file `mygcd.c`, then you'd compile it to a bitcode module using `llvm-gcc` as follows:

```
llvm-gcc -emit-llvm -c mygcd.c -o mygcd.bc
```

Or, if you prefer to use [clang](#), the new LLVM-based C/C++ compiler:

```
clang -emit-llvm -c mygcd.c -o mygcd.bc
```

Note that the `-emit-llvm -c` options instruct `llvm-gcc` or `clang` to build an LLVM bitcode module. Of course, you can also add optimizations and other options to the compile command as desired.

You can now load the resulting bitcode module and run the `mygcd` function in the Pure interpreter simply as follows:

```
> using "bc:mygcd";
> mygcd 75 105;
15
```

To actually see the generated extern declaration of the imported function, you can use the interactive `show` command:

```
> show mygcd
extern int mygcd(int, int);
```

Some more examples showing how to use the bitcode interface can be found in the Pure sources. In particular, the interface also works with Fortran (using `llvm-gfortran`), and there is special support for interfacing to Grame's functional DSP programming language [Faust](#) (the latter uses a special variant of the bitcode loader, which is selected with the `dsp` tag in

the using clause). Please refer to the corresponding examples in the distribution for further details.

Please note that at this time the LLVM bitcode interface is still somewhat experimental, and there are some known limitations:

- LLVM doesn't distinguish between `char*` and `void*` in bitcode, so all `void*` parameters and return values in C code will be promoted to `char*` on the Pure side. Also, pointers to types which neither have a predefined meaning in Pure nor a proper type name in the bitcode file, will become a generic pointer type (`void*`, `void**`, etc.) in Pure. If this is a problem then you can just redeclare the corresponding functions under an alias *after* loading the bitcode module, giving the proper argument and result types (see [Extern Declarations](#) above).
- The bitcode interface is limited to the same range of C types as Pure's plain C interface. In practice, this should cover most C code, but it's certainly possible that you run into unsupported types for arguments and return values. The compiler will then print a warning; the affected functions will still be linked in, but they will not be callable from Pure. Also note that calling conventions for passing C structs *by value* depend on the host ABI, so you should have a look at the resulting extern declaration (using `show`) to determine how the function is actually to be called from Pure.

9.5 Inline Code

Instead of manually compiling source files to bitcode modules, you can also just place the source code into a Pure script, enclosing it in `%{ ... %}`. (Optionally, the opening brace may also be preceded with a `public` or `private` scope specifier, which is used in the same way as the scope specifier following the `using` keyword when importing bitcode files.)

For instance, here is a little script showing inline code for the `mygcd` function from the previous subsection:

```
%{
int mygcd(int x, int y)
{
    if (y == 0)
        return x;
    else
        return mygcd(y, x%y);
}
%}
```

```
mygcd 75 105;
```

The interpreter automatically compiles the inlined code to LLVM bitcode which is then loaded as usual. (Of course, this will only work if you have the corresponding LLVM compilers installed.) This method has the advantage that you don't have to write a Makefile and you can create self-contained Pure scripts which include all required external functions. The downside is that the inline code sections will have to be recompiled every time you run the script with the interpreter which may considerably increase startup times. If this is a

problem then it's usually better to import a separate bitcode module instead (see [Importing LLVM Bitcode](#)), or at least batch-compile your script to an executable (see [Batch Compilation](#)).

Currently, C, C++, Fortran and Faust are supported as foreign source languages, with `llvm-gcc`, `llvm-g++`, `llvm-gfortran` and `faust2` as the corresponding compilers. Alternatively, the LLVM `clang` and `clang++` compilers can be used for C/C++ compilation (this will actually be default if the Pure interpreter itself was compiled with `clang`). Examples for all of these can be found in the Pure sources.

C is the default language. The desired source language can be selected by placing an appropriate tag into the inline code section, immediately after the opening brace. (The tag is removed before the code is submitted to compilation.) For instance:

```
%{ -*- Fortran90 -*-  
function fact(n) result(p)  
  integer n, p  
  p = 1  
  do i = 1, n  
    p = p*i  
  end do  
end function fact  
%}
```

```
fact n::int = fact_ {n};  
map fact (1..10);
```

As indicated, the language tag takes the form `-*- lang -*-` where `lang` can currently be any of `c`, `c++`, `fortran` and `dsp` (the latter indicates the Faust language). Case is insignificant here, so you can also write `C`, `C++`, `Fortran`, `DSP` etc. For the `fortran` tag, you may also have to specify the appropriate language standard, such as `fortran90` which is used in the example above. The language tag can also be followed by a module name, using the format `-*- lang : name -*-`. This is optional for all languages except Faust (where the module name specifies the namespace for the interface routines of the Faust module). So, e.g., a Faust DSP named `test` would be specified with a `dsp:test` tag. Case is *significant* in the module name.

The Pure interpreter has some built-in knowledge on how to invoke the LLVM compilers to produce a working bitcode file ready to be loaded by the interpreter, so the examples above should work out of the box if you have the required compilers installed on your `PATH`. However, there are also some environment variables you can set for customization purposes. Specifically, `PURE_CC` is the command to invoke the C compiler. This variable lets you specify the exact name of the executable along with any debugging and optimization options that you may want to add. Likewise, `PURE_CXX`, `PURE_FC` and `PURE_FAUST` are used for the C++, Fortran and Faust compilers, respectively.

For instance, if you prefer to use `clang` as your C compiler, and you'd like to invoke it with the `-O3` optimization option, you would set `PURE_CC` to `"clang -O3"`. (To verify the settings you made, you can have the interpreter echo the compilation commands which are actually executed, by running Pure with the `-v0100` option, see [Verbosity and Debugging Options](#).)

10 Standard Library

Pure comes with a collection of Pure library modules, which includes the standard prelude (loaded automatically at startup time) and some other modules which can be loaded explicitly with a `using` clause. The prelude offers the necessary functions to work with the built-in types (including arithmetic and logical operations) and to do most kind of list processing you can find in ML- and Haskell-like languages. It also provides a collection of basic string and matrix operations. Please refer to the [Pure Library Manual](#) for details on the provided operations. Here is a very brief summary of some of the prelude operations which, besides the usual arithmetic and logical operators, are probably used most frequently:

- x+y** The arithmetic `+` operation is also used to denote list and string concatenation in Pure.
- x:y** This is the list-consing operation. `x` becomes the head of the list, `y` its tail. As `'::'` is a constructor symbol, you can use it in patterns on the left hand side of rewriting rules.
- x..y** Constructs arithmetic sequences. `x:y..z` can be used to denote sequences with arbitrary stepsize `y-x`. Infinite sequences can be constructed using an infinite bound (i.e., `inf` or `-inf`). E.g., `1:3..inf` denotes the stream of all odd integers starting at 1.
- x,y** This is the pair constructor, used to create tuples of arbitrary sizes. Tuples provide an alternative way to represent aggregate values in Pure. In difference to lists, tuples are always “flat”, so that `(x,y),z` and `x,(y,z)` denote the same triple `x,y,z`. (This is explained in more detail in the [Pure Overview](#) section.)
- #x** The size (number of elements) of the list, tuple, matrix or string `x`. In addition, `dim x` yields the dimensions (number of rows and columns) of a matrix.
- x!y** This is Pure’s indexing operation, which applies to lists, tuples, matrices and strings. Note that all indices in Pure are zero-based, thus `x!0` and `x!(#x-1)` are the first and last element of `x`. In the case of matrices, the subscript may also be a pair of row and column indices, such as `x!(1,2)`.
- x!!ys** This is the “slicing” operation, which returns the list, tuple, matrix or string of all `x!y` while `y` runs through the (list or matrix) `ys`. Thus, e.g., `x!!(i..j)` returns all the elements between `i` and `j` (inclusive). Indices which fall outside the valid index range are quietly discarded. The index range `ys` may contain any number of indices (also duplicates), in any order. Thus `x!![0|i=1..n]` returns the first element of `x` `n` times, and, if `ys` is a permutation of the range `0..#x-1`, then `x!!ys` yields the corresponding permutation of the elements of `x`. In the case of matrices the index range may also contain two-dimensional subscripts, or the index range itself may be specified as a pair of row/column index lists such as `x!!(i..j,k..l)`.

The prelude also offers support operations for the implementation of list and matrix comprehensions, as well as the customary list operations like `head`, `tail`, `drop`, `take`, `filter`, `map`, `foldl`, `foldr`, `scanl`, `scanr`, `zip`, `unzip`, etc., which make list programming so much fun in modern FPLs. In Pure, these also work on strings as well as matrices, although, for reasons of efficiency, these data structures are internally represented as arrays.

Besides the prelude, Pure’s standard library also comprises a growing number of additional library modules which we can only mention in passing here. In particular, the `math.pure`

module provides additional mathematical functions as well as Pure's complex and rational number data types. Common container data structures like sets and dictionaries are implemented in the `set.pure` and `dict.pure` modules, among others. Moreover, the (beginnings of a) system interface can be found in the `system.pure` module. In particular, this module also provides operations to do basic C-style I/O, including `printf` and `scanf`. These are all described in much more detail in the *Pure Library Manual*.

11 Interactive Usage

In interactive mode, the interpreter reads definitions and expressions and processes them as usual. You can use the `-i` option to force interactive mode when invoking the interpreter with some script files. Additional scripts can be loaded interactively using either a `using` declaration or the `interactive run` command (see the description of the `run` command below for the differences between these). Or you can just start typing away, entering your own definitions and expressions to be evaluated.

The input language is just the same as for source scripts, and hence individual definitions and expressions must be terminated with a semicolon before they are processed. For instance, here is a simple interaction which defines the factorial and then uses that definition in some evaluations. Input lines begin with `'>'`, which is the interpreter's default command prompt:

```
> fact 1 = 1;
> fact n = n*fact (n-1) if n>1;
> let x = fact 10; x;
3628800
> map fact (1..10);
[1,2,6,24,120,720,5040,40320,362880,3628800]
```

As indicated, in interactive mode the normal forms of toplevel expressions are printed after each expression is entered. We also call this the **read-eval-print loop**. Normal form expressions are usually printed in the same form as you'd enter them. However, there are a few special kinds of objects like anonymous closures, thunks ("lazy" values to be evaluated when needed) and pointers which don't have a textual representation in the Pure syntax and will be printed in the format `#<object description>` by default. It is also possible to override the print representation of any kind of expression by means of the `__show__` function, see the [Caveats and Notes](#) section for details.

11.1 Online Help

Online help is available in the interpreter with the `interactive help` command, see [Interactive Commands](#) below. You need to have a html browser installed for that. By default, the `help` command uses `w3m(1)`, but you can change this by setting either the `PURE_HELP` or the `BROWSER` environment variable accordingly.

As of Pure 0.46, the interpreter provides a much improved `help` command which gives you access to all the available documentation in html format, which includes this manual, the

Pure Library Manual, as well as all manuals of the addon modules available from the Pure website.

When invoked without arguments, the `help` command displays an overview of the available documentation, from which you can follow the links to the provided manuals:

```
> help
```

(If your browser gives you an error message when you do this then you haven't installed the documentation yet. The complete set of manuals is provided as a separate package at the Pure website, please see the Pure installation instructions for details.)

The `help` command also accepts a parameter which lets you specify a search term which is looked up in the global index, e.g.:

```
> help foldl
```

If the search term doesn't appear in the index, it is assumed to be a topic (a link target) in the Pure manual. Note that the `docutils` tools used to generate the html source of the Pure documentation mangle the section titles so that they are in lowercase and blanks are replaced with hyphens. So to look up the present section in this manual you'd have to type:

```
> help online-help
```

The help files are in html format and located in the docs subdirectory of the Pure library directory (i.e., `/usr/local/lib/pure/docs` by default). You can look up topics in any of the help files with a command like the following:

```
> help pure-gsl#matrices
```

Here `pure-gsl` is the basename of the help file (library path and `.html` suffix are supplied automatically), and `matrices` is a link target in that document. To just read the `pure-gsl.html` file without specifying a target, type the following:

```
> help pure-gsl#
```

(Note that just `help pure-gsl` won't work, since it would look for a search term in the index or a topic in the Pure manual.)

Last but not least, you can also point the help browser to any html document (either a local file or some website) denoted by a proper URL, provided that your browser program can handle these. For instance:

```
> help file:mydoc.html#foo
> help http://pure-lang.googlecode.com
```

11.2 Interactive Commands

When running interactively, the interpreter accepts a number of special commands useful for interactive purposes. Here is a quick rundown of the currently supported operations:

! **command** Shell escape.

break [**symbol** ...] Sets breakpoints on the given function or operator symbols. All symbols must be specified in fully qualified form, see the remarks below. If invoked without arguments, prints all currently defined breakpoints. This requires that the interpreter was invoked with the `-g` option to enable debugging support. See [Debugging](#) below for details.

bt Prints a full backtrace of the call sequence of the most recent evaluation, if that evaluation ended with an unhandled exception. This requires that the interpreter was invoked with the `-g` option to enable debugging support. See [Debugging](#) below for details.

cd **dir** Change the current working dir.

clear [**option** ...] [**symbol** ...] Purge the definitions of the given symbols (functions, macros, constants or global variables). All symbols must be specified in fully qualified form, see the remarks below. If invoked as `clear ans`, clears the `ans` value (see [Last Result](#) below). When invoked without any arguments, `clear` purges all definitions at the current interactive “level” (after confirmation) and returns you to the previous level, if any. (It might be a good idea to first check your current definitions with `show` or back them up with `dump` before you do that.) The desired level can be specified with the `-t` option. See the description of the `save` command and [Definition Levels](#) below for further details. A description of the common options accepted by the `clear`, `dump` and `show` commands can be found in [Specifying Symbol Selections](#) below.

del [**symbol** ...] Deletes breakpoints and tracepoints on the given function or operator symbols. All symbols must be specified in fully qualified form, see the remarks below. If invoked without arguments, clears all currently defined breakpoints and tracepoints (after confirmation). See [Debugging](#) below for details.

dump [**-n filename**] [**option** ...] [**symbol** ...] Dump a snapshot of the current function, macro, constant and variable definitions in Pure syntax to a text file. All symbols must be specified in fully qualified form, see the remarks below. This works similar to the `show` command (see below), but writes the definitions to a file. The default output file is `.pure` in the current directory, which is then reloaded automatically the next time the interpreter starts up in interactive mode in the same directory. This provides a quick-and-dirty way to save an interactive session and have it restored later, but note that this isn’t perfect. In particular, declarations of extern symbols won’t be saved unless they’re specified explicitly, and some objects like closures, thunks and pointers don’t have a textual representation from which they could be reconstructed. To handle these, you’ll probably have to prepare a corresponding `.purerc` file yourself, see [Interactive Startup](#) below.

A different filename can be specified with the `-n` option, which expects the name of the script to be written in the next argument, e.g: `dump -n myscript.pure`. You can then edit that file and use it as a starting point for an ordinary script or a `.purerc` file, or you can just run the file with the `run` command (see below) to restore the definitions in a subsequent interpreter session.

help [**topic**] Display online documentation. If a topic is given, it is looked up in the index. Alternatively, you can also specify a link target in any of the installed help files, or

any other html document denoted by a proper URL. Please see [Online Help](#) above for details.

ls [args] List files (shell `ls(1)` command).

mem Print current memory usage. This reports the number of expression cells currently in use by the program, along with the size of the freelist (the number of allocated but currently unused expression cells). Note that the actual size of the expression storage may be somewhat larger than this, since the runtime always allocates expression memory in bigger chunks. Also, this figure does not reflect other heap-allocated memory in use by the program, such as strings or malloc'ed pointers.

override Enter “override” mode. This allows you to add equations “above” existing definitions in the source script, possibly overriding existing equations. See [Definition Levels](#) below for details.

pwd Print the current working dir (shell `pwd(1)` command).

quit Exits the interpreter.

run [-g|script] When invoked without arguments or with the `-g` option, `run` does a “cold” restart of the interpreter, with the scripts and options given on the interpreter’s original command line. If just `-g` is specified as the argument, the interpreter is run with debugging enabled. Otherwise the interpreter is invoked without debugging support. (This overrides the corresponding option from the interpreter’s command line.) This command provides a quick way to rerun the interpreter after changes in some of the loaded script files, or if you want to enable or disable debugging on the fly (which requires a restart of the interpreter). You’ll also loose any definitions that you entered interactively in the interpreter, so you may want to back them up with `dump` beforehand.

When invoked with a script name as argument, `run` loads the given script file and adds its definitions to the current environment. This works more or less like a `using` clause, but only searches for the script in the current directory and places the definitions in the script at the current temporary level, so that `clear` can be used to remove them again. Also note that namespace and pragma settings of scripts loaded with `run` stick around after loading the script. This allows you to quickly set up your environment by just running a script containing the necessary namespace declarations and compiler directives. (Alternatively, you can also use the interpreter’s startup files for that purpose, see [Interactive Startup](#) below.)

save Begin a new level of temporary definitions. A subsequent `clear` command (see above) will purge the definitions made since the most recent `save` command. See [Definition Levels](#) below for details.

show [option ...] [symbol ...] Show the definitions of symbols in various formats. See [The show Command](#) below for details. All symbols must be specified in fully qualified form, see the remarks below. A description of the common options accepted by the `clear`, `dump` and `show` commands can be found in [Specifying Symbol Selections](#) below.

stats [-m] [on|off] Enables (default) or disables “stats” mode, in which some statistics are printed after an expression has been evaluated. Invoking just `stats` or `stats on` only prints the cpu time in seconds for each evaluation. If the `-m` option is specified, memory usage is printed along with the cpu time, which indicates the maximum amount of expression memory (in terms of expression cells) used during the computation. Invoking `stats off` disables stats mode, while `stats -m off` just disables the printing of the memory usage statistics.

trace [symbol ...] Sets tracepoints on the given function or operator symbols. This works like the `break` command (see above) but only prints rule invocations and reductions without actually interrupting the evaluation. See [Debugging](#) below for details.

underride Exits “override” mode. This returns you to the normal mode of operation, where new equations are added “below” previous rules of an existing function. See [Definition Levels](#) below for details.

Note that these special commands are only recognized at the beginning of the interactive command line (they are not reserved keywords of the Pure language). Thus it’s possible to “escape” identifiers looking like commands by entering a space at the beginning of the line.

Also note that symbols (identifiers, operators etc.) must always be specified in fully qualified form. No form of namespace lookup is performed by these commands, so they always work the same no matter what namespace and using namespace declarations are currently in effect.

11.3 Last Result

Another convenience for interactive usage is the `ans` function, which retrieves the most recent result printed in interactive mode. For instance:

```
> fact n = if n<=1 then 1 else n*fact (n-1);
> map fact (1..10);
[1,2,6,24,120,720,5040,40320,362880,3628800]
> scanl (+) 0 ans;
[0,1,3,9,33,153,873,5913,46233,409113,4037913]
```

Note that `ans` is just an ordinary function, defined in the prelude, not a special command. However, there is a special `clear ans` command which purges the `ans` value. This is useful, e.g., if you got a huge result which you want to erase from memory before starting the next computation.

```
> clear ans
> ans;
ans
```

11.4 Specifying Symbol Selections

The `clear`, `dump` and `show` commands all accept the following options for specifying a subset of symbols and definitions on which to operate. All symbols must be specified in fully qualified form. Options may be combined, thus, e.g., `show -mft` is the same as `show -m -f -t`. Some options specify optional numeric parameters; these must follow immediately behind the option character if present, as in `-t0`.

- c** Selects defined constants.
- f** Selects defined functions.
- g** Indicates that the following symbols are actually shell glob patterns and that all matching symbols should be selected.
- m** Select defined macros.
- pflag** Select only private symbols if *flag* is nonzero (the default), otherwise (*flag* is zero) select only public symbols. If this option is omitted then both private and public symbols are selected.
- tlevel** Select symbols and definitions at the given “level” of definitions and above. This is described in more detail below. Briefly, the executing program and all imported modules (including the prelude) are at level 0, while “temporary” definitions made interactively in the interpreter are at level 1 and above. Thus a level of 1 restricts the selection to all temporary definitions, whereas 0 indicates all definitions (i.e., everything, including the prelude). If *level* is omitted, it defaults to the current definitions level.
- v** Select defined variables.

In addition, the `-h` option prints a short help message describing all available options of the command at hand.

If none of the `-c`, `-f`, `-m` and `-v` options are specified, then all kinds of symbols (constants, functions, macros and variables) are selected, otherwise only the specified categories will be considered.

A reasonable default is used if the `-t` option is omitted. By default, if no symbols are specified, only temporary definitions are considered, which corresponds to `-t1`. Otherwise the command applies to all corresponding definitions, no matter whether they belong to the executing program, the prelude, or some temporary level, which has the same effect as `-t0`. This default choice can be overridden by specifying the desired level explicitly.

As a special case, just `clear` (without any other options or symbol arguments) always backs out to the previous definitions level (instead of level #1). This is inconsistent with the rules set out above, but is implemented this way for convenience and backward compatibility. Thus, if you really want to delete all your temporary definitions, use `clear -t1` instead. When used in this way, the `clear` command will only remove temporary definitions; if you need to remove definitions at level #0, you must specify those symbols explicitly.

Note that `clear -g *` will have pretty much the same disastrous consequences as the Unix command `rm -rf *`, so don't do that. Also note that a macro or function symbol may well have defining equations at different levels, in which case a command like `clear -tn foo` might only affect some part of `foo`'s definition. The `dump` and `show` commands work analogously (albeit less destructively). See [Definition Levels](#) below for some examples.

11.5 The show Command

The `show` command can be used to obtain information about defined symbols in various formats. Besides the common selection options discussed above, this command recognizes the following additional options for specifying the content to be listed and the format to use.

- a** Disassembles pattern matching automata. Works like the `-v4` option of the interpreter.
- d** Disassembles LLVM IR, showing the generated LLVM assembler code of a function. Works like the `-v8` option of the interpreter.
- e** Annotate printed definitions with lexical environment information (de Bruijn indices, subterm paths). Works like the `-v2` option of the interpreter.
- l** Long format, prints definitions along with the summary symbol information. This implies `-s`.
- s** Summary format, print just summary information about listed symbols.

Symbols are always listed in lexicographic order. Note that some of the options (in particular, `-a` and `-d`) may produce excessive amounts of information. By setting the `PURE_MORE` environment variable, you can specify a shell command to be used for paging, usually `more(1)` or `less(1)`.

For instance, to list all temporary definitions made in an interactive session, simply say:

```
> show
```

You can also list a specific symbol, no matter whether it comes from the interactive command line, the executing script or the prelude:

```
> show foldl
foldl f a x::matrix = foldl f a (list x);
foldl f a s::string = foldl f a (chars s);
foldl f a [] = a;
foldl f a (x:xs) = foldl f (f a x) xs;
```

Wildcards can be used with the `-g` option, which is useful if you want to print an entire family of related functions, e.g.:

```
> show -g foldl*
foldl f a x::matrix = foldl f a (list x);
foldl f a s::string = foldl f a (chars s);
foldl f a [] = a;
foldl f a (x:xs) = foldl f (f a x) xs;
foldl1 f x::matrix = foldl1 f (list x);
```

```
foldl1 f s::string = foldl1 f (chars s);
foldl1 f (x:xs) = foldl f x xs;
```

Or you can just specify multiple symbols as follows (this also works with multiple glob patterns when you add the `-g` option):

```
> show min max
max x y = if x>=y then x else y;
min x y = if x<=y then x else y;
```

You can also select symbols by category. E.g., the following command shows summary information about all the variable symbols along with their current values (using the “long” format):

```
> show -lv *
argc      var  argc = 0;
argv      var  argv = [];
compiling var  compiling = 0;
sysinfo   var  sysinfo = "x86_64-unknown-linux-gnu";
version   var  version = "0.37";
5 variables
```

Or you can list just private symbols of the namespace `foo`, as follows:

```
> show -pg foo::*
```

The following command will list each and every symbol that’s currently defined (instead of `-g *` you can also use the `-t0` option):

```
> show -g *
```

This usually produces a lot of output and is rarely needed, unless you’d like to browse through an entire program including all library imports. (In that case you might consider to use the `dump` command instead, which writes the definitions to a file which can then be loaded into a text editor for easier viewing. This may occasionally be useful for debugging purposes.)

Finally, there are two alternate forms of the `show` command: `show namespace` which lists the current and search namespaces, and `show namespaces` which lists all declared namespaces. These come in handy if you have forgotten what namespaces are currently active and which other namespaces are available in your program. For instance:

```
> show namespace
> show namespaces
namespace C;
namespace matrix;
> using namespace C;
> namespace my;
> show namespace
namespace my;
using namespace C;
```

11.6 Definition Levels

To help with incremental development, the interpreter offers some commands to manipulate the current set of definitions interactively. To these ends, definitions are organized into different subsets called **levels**. As already mentioned, the prelude, as well as other source programs specified when invoking the interpreter, are always at level 0, while the interactive environment starts at level 1. Each save command introduces a new temporary level, and each subsequent clear command (without any arguments) “pops” the definitions on the current level and returns you to the previous one (if any). This gives you a “stack” of temporary environments which enables you to “plug and play” in a (more or less) safe fashion, without affecting the rest of your program.

For all practical purposes, this stack is unlimited, so that you can create as many levels as you like. However, this facility does have its limitations. The interpreter doesn’t really keep a full history of everything you entered interactively, it only records the level a variable, constant, and function or macro rule belongs to so that the corresponding definitions can be removed again when the level is popped. On the other hand, intermediate changes in variable values are not recorded anywhere and cannot be undone. Moreover, global declarations (which encompasses using clauses, extern declarations and special symbol declarations) always apply to all levels, so they can’t be undone either.

That said, the temporary levels can still be pretty useful when you’re playing around with the interpreter. Here’s a little example which shows how to use clear to quickly get rid of a definition that you entered interactively:

```
> foo (x:xs) = x+foo xs;
> foo [] = 0;
> show
foo (x:xs) = x+foo xs;
foo [] = 0;
> foo (1..10);
55
> clear
This will clear all temporary definitions at level #1.
Continue (y/n)? y
> show
> foo (1..10);
foo [1,2,3,4,5,6,7,8,9,10]
```

We’ve seen already that normally, if you enter a sequence of equations, they will be recorded in the order in which they were written. However, it is also possible to override definitions in lower levels with the override command:

```
> foo (x:xs) = x+foo xs;
> foo [] = 0;
> show
foo (x:xs) = x+foo xs;
foo [] = 0;
> foo (1..10);
55
> save
```



```
save: now at temporary definitions level #2
> override
> foo (x:xs) = x*foo xs;
> show
foo (x:xs) = x*foo xs;
foo (x:xs) = x+foo xs;
foo [] = 0;
> foo (1..10);
warning: rule never reduced: foo (x:xs) = x+foo xs;
0
```

Note that the equation `foo (x:xs) = x*foo xs` was inserted before the previous rule `foo (x:xs) = x+foo xs`, which is at level #1. (The latter equation is now “shadowed” by the rule we just entered, hence the compiler warns us that this rule can’t be reduced any more.)

Even in override mode, new definitions will be added after other definitions at the *current* level. This allows us to just continue adding more high-priority definitions overriding lower-priority ones:

```
> foo [] = 1;
> show
foo (x:xs) = x*foo xs;
foo [] = 1;
foo (x:xs) = x+foo xs;
foo [] = 0;
> foo (1..10);
warning: rule never reduced: foo (x:xs) = x+foo xs;
warning: rule never reduced: foo [] = 0;
3628800
```

Again, the new equation was inserted above the existing lower-priority rules, but below our previous equation `foo (x:xs) = x*foo xs` entered at the same level. As you can see, we have now effectively replaced our original definition of `foo` with a version that calculates list products instead of sums, but of course we can easily go back one level to restore the previous definition:

```
> clear
This will clear all temporary definitions at level #2.
Continue (y/n)? y
clear: now at temporary definitions level #1
clear: override mode is on
> show
foo (x:xs) = x+foo xs;
foo [] = 0;
> foo (1..10);
55
```

Note that `clear` reminded us that override mode is still enabled (`save` will do the same if override mode is on while pushing a new definitions level). To turn it off again, use the `underride` command. This will revert to the normal behaviour of adding new equations below existing ones:

```
> underride
```

It's also possible to use `clear` to back out multiple levels at once, if you specify the target level to be cleared with the `-t` option. For instance:

```
> save
save: now at temporary definitions level #2
> let bar = 99;
> show
let bar = 99;
foo (x:xs) = x+foo xs;
foo [] = 0;
> // this scraps all our scribblings!
> clear -t1
This will clear all temporary definitions at level #1 and above.
Continue (y/n)? y
clear: now at temporary definitions level #1
> show
>
```

11.7 Debugging

The interpreter provides a simple but reasonably convenient symbolic debugging facility when running interactively. To make this work, you have to specify the `-g` option when invoking the interpreter (`pure -g`). If you're already at the interpreter's command line, you can also use the `run -g` command to enable the debugger. The `-g` option disables tail call optimization (see [Stack Size and Tail Recursion](#)) to make it easier to debug programs. It also causes special debugging code to be generated which will make your program run *much* slower. Therefore the `-g` option should only be used if you actually need the debugger.

One common use of the debugger is "post mortem" debugging after an evaluation ended with an unhandled exception. In such a case, the `bt` command of the interpreter prints a backtrace of the call sequence which caused the exception. Note that this only works if debugging mode was enabled. For instance:

```
> [1,2]!3;
<stdin>, line 2: unhandled exception 'out_of_bounds' while evaluating '[1,2]!3'
> bt
[1] (!): (x:xs)!n::int = xs!(n-1) if n>0;
      n = 3; x = 1; xs = [2]
[2] (!): (x:xs)!n::int = xs!(n-1) if n>0;
      n = 2; x = 2; xs = []
[3] (!): []!n::int = throw out_of_bounds;
      n = 1
>> [4] throw: extern void pure_throw(expr*) = throw;
      x1 = out_of_bounds
```

The last call, which is also marked with the `>>` symbol, is the call that raised the exception. The format is similar to the `p` command of the debugger, see below, but `bt` always prints a full backtrace. (As with the `show` command of the interpreter, you can set the `PURE_MORE`

environment variable to pipe the output through the corresponding command, or use *eval-cmd* to capture the output of *bt* in a string, cf. [Reflection](#).)

The debugger can also be used interactively. To these ends, you can set breakpoints on functions with the *break* command. The debugger then gets invoked as soon as a rule for one of the given functions is executed. Example:

```
> fact n::int = if n>0 then n*fact (n-1) else 1;
> break fact
> fact 1;
** [1] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
    n = 1
(Type 'h' for help.)
:
** [2] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
    n = 0
:
++ [2] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
    n = 0
    --> 1
** [2] (*): x::int*y::int = x*y;
    x = 1; y = 1
:
++ [2] (*): x::int*y::int = x*y;
    x = 1; y = 1
    --> 1
++ [1] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
    n = 1
    --> 1
1
```

Lines beginning with **** indicate that the evaluation was interrupted to show the rule (or external) which is currently being considered, along with the current depth of the call stack, the invoked function and the values of parameters and other local variables in the current lexical environment. In contrast, the prefix *++* denotes reductions which were actually performed during the evaluation and the results that were returned by the function call (printed as *--> return value*).

Sometimes you might also see funny symbols like *#<case>*, *#<when>* or *#<closure>* instead of the function name. These indicate lambdas and the special variable-binding environments, which are all implemented as anonymous closures in Pure. Also note that the debugger doesn't know about the argument names of external functions (which are optional in Pure and not recorded anywhere), so it will display the generic names *x1*, *x2* etc. instead.

At the debugger prompt *'*:*'* you can enter various special debugger commands, or just keep on hitting the carriage return key to walk through an evaluation step by step, as we did in the example above. (Command line editing works as usual at the debugger prompt, if it is enabled.) The usual commands are provided to walk through an evaluation, print and navigate the call stack, step over the current call, or continue the evaluation unattended until you hit another breakpoint. If you know other source level debuggers like *gdb* then you should feel right at home. You can type *h* at the debugger prompt to print the following

list:

```
: h
Debugger commands:
a      auto: step through the entire program, run unattended
c [f]  continue until next breakpoint, or given function f
h      help: print this list
n      next step: step over reduction
p [n]  print rule stack (n = number of frames)
r      run: finish evaluation without debugger
s      single step: step into reduction
t, b   move to the top or bottom of the rule stack
u, d   move up or down one level in the rule stack
x      exit the interpreter (after confirmation)
.      reprint current rule
! cmd  shell escape
? expr evaluate expression
<cr>   single step (same as 's')
<eof>  step through program, run unattended (same as 'a')
```

The command syntax is very simple. Besides the commands listed above you can also enter comment lines (`// comment text`) which will just be ignored. Extra arguments on commands which don't expect any will generally be ignored as well. The single letter commands all have to be separated from any additional parameters with whitespace, whereas the `!`, `?` and `.` commands count as word delimiters and can thus be followed immediately by an argument. For convenience, the `?` command can also be omitted if the expression to be evaluated doesn't start with a single letter or one of the special punctuation commands.

The debugger can be exited or suspended in the following ways:

- You can type `c` to continue the evaluation until the next breakpoint, or `c foo` in order to proceed until the debugger hits an invocation of the function `foo`.
- You can type `r` to run the rest of the evaluation without the debugger.
- The `a` ("auto") command single-steps through the rest of the evaluation, running unattended. This command can also be entered by just hitting the end-of-file key (`Ctrl-D` on Unix systems) at the debugger prompt.
- You can also type `x` to exit from the debugger *and* the interpreter immediately (after confirmation).

At the debugger prompt, you can use the `u` ("up"), `d` ("down"), `t` ("top") and `b` ("bottom") commands to move around on the current call stack. The `p` command prints a range of the call stack centered around the currently selected stack frame, which is indicated with the `>>` tag, whereas `**` denotes the current bottom of the stack (which is the rule to be executed with the single step command). The `p` command can also be followed by a numeric argument which indicates the number of stack frames to be printed (this will then become the default for subsequent invocations of `p`). The `n` command steps over the call selected with the stack navigation commands. For instance:

```
> fact 3;
** [1] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 3
: c *
** [4] (*): x::int*y::int = x*y;
      x = 1; y = 1
: p
  [1] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 3
  [2] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 2
  [3] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 1
** [4] (*): x::int*y::int = x*y;
      x = 1; y = 1
: u
>> [3] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 1
: u
>> [2] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 2
: p
  [1] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 3
>> [2] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 2
  [3] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 1
** [4] (*): x::int*y::int = x*y;
      x = 1; y = 1
: n
++ [2] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 2
      --> 2
** [2] (*): x::int*y::int = x*y;
      x = 3; y = 2
:
```

If you ever get lost, you can reprint the current rule with the `'.'` command:

```
: .
** [2] (*): x::int*y::int = x*y;
      x = 3; y = 2
```

Another useful feature is the `?` command which lets you evaluate any Pure expression, with the local variables of the current rule bound to their corresponding values. Like the `n` command, `?` applies to the current stack frame as selected with the stack navigation commands. The expression must be entered on a single line, and the trailing semicolon is optional. For instance:

```
> fact 3;
** [1] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
```

```
      n = 3
: c *
** [4] (*): x::int*y::int = x*y;
      x = 1; y = 1
: ?x+y
2
: u
>> [3] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 1
: n>0, fact n
1,1
```

A third use of the debugger is to trace function calls. For that the interpreter provides the trace command which works similarly to break, but sets so-called “tracepoints” which only print rule invocations and reductions instead of actually interrupting the evaluation. For instance, assuming the same example as above, let’s first remove the breakpoint on fact (using the del command) and then set it as a tracepoint instead:

```
> del fact
> trace fact
> fact 1;
** [1] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 1
** [2] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 0
++ [2] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 0
      --> 1
++ [1] fact: fact n::int = if n>0 then n*fact (n-1) else 1;
      n = 1
      --> 1
1
```

The break and trace commands can also be used in concert if you want to debug some functions while only tracing others.

The current sets of breakpoints and tracepoints can be changed with the break, trace and del commands, as shown above, and just break or trace without any arguments lists the currently defined breakpoints or tracepoints, respectively. Please see [Interactive Commands](#) above for details. Also note that these are really interpreter commands, so to enter them you first have to exit the debugger (using the a or r command) if an evaluation is currently in progress. (However, it’s possible to set a temporary breakpoint in the debugger with the c command, see above.)

11.8 Interactive Startup

In interactive mode, the interpreter also runs some additional scripts at startup, after loading the prelude and the scripts specified on the command line. This lets you tailor the interactive environment to your liking.

The interpreter first looks for a `.purerc` file in the user's home directory (as given by the `HOME` environment variable) and then for a `.purerc` file in the current working directory. These are just ordinary Pure scripts which may contain any additional definitions that you need. The `.purerc` file in the home directory is for global definitions which should always be available when running interactively, while the `.purerc` file in the current directory can be used for project-specific definitions.

Finally, you can also have a `.pure` initialization file in the current directory, which is usually created with the `dump` command (see above). This file is loaded after the `.purerc` files if it is present.

The interpreter processes all these files in the same way as with the `run` command (see above). When invoking the interpreter, you can specify the `-norc` option on the command line if you wish to skip these initializations.

12 Batch Compilation

As of Pure 0.21, the interpreter has a new `-c` option which provides a convenient means to turn Pure scripts into standalone executables. This feature is still a bit experimental right now. In particular, note that the compiled executable is essentially a *static snapshot* of your program which is executed on the "bare metal", without a hosting interpreter. Only a minimal runtime system is provided. This considerably reduces startup times, but also implies the following quirks and limitations:

- All toplevel expressions and `let` bindings are evaluated *after* all functions have been defined. This might cause inconsistent behaviour with an interpreted run of the same program, which executes expressions and variable definitions immediately, as the program is being processed. To avoid these semantic differences, you'll have to make sure that expressions are evaluated *after* all functions used in the evaluation have been defined completely.
- Toplevel expressions won't be of much use in a batch-compiled program, unless, of course, they are evaluated for their side-effects. Usually your program will have to include at least one of these to play the role of the "main program" in your script. In most cases these expressions are best placed after all the function and variable definitions, at the end of your program.
- The `eval` function can only be used to evaluate plain toplevel expressions. You can define local functions and variables in `with` and `when` clauses inside an expression, but you can't use `eval` to define new global variables and functions. In other words, anything which changes the executing program is "verboten". Moreover, the introspective capabilities provided by `evalcmd` (discussed under [Reflection](#) in the [Caveats and Notes](#) section) won't work either because the interactive commands are all disabled. If you need any of these capabilities, you have to run your program with the interpreter.
- Constant and macro definitions, being compile time features, aren't available in the compiled program. If you need to use these with `eval` at run time, you have to provide them through variable and function definitions instead. Also, the compiler usually

strips unused functions from the output code, so that only functions which are actually called somewhere in the static program text are available to `eval`. (The `-u` option and the `-required` pragma can be used to avoid this, see [Code Size and Unstripped Executables](#) below.)

- Code which gets executed to compute constant values at compile time will generally *not* be executed in the compiled executable, so your program shouldn't rely on side-effects of such computations (this would be bad practice anyway). There is an exception to this rule, however, namely if a constant value contains run time data such as pointers and local functions which requires an initialization at run time, then the batch compiler will generate code for that. (The same happens if the `-noconst` option is used to force computation of constant values at run time, see [Code Size and Unstripped Executables](#).)

What all this boils down to is that anything which requires the compile time or interactive facilities of the interpreter, is unavailable. These restrictions only apply at run time, of course. At compile time the program *is* being executed by the interpreter so you can use [eval](#) and [evalcmd](#) in any desired way. See the description of the `compiling` variable below for how to distinguish these cases in your script.

For most kinds of scripts, the above restrictions aren't really that much of an obstacle, or can easily be worked around. For the few scripts which actually need the full dynamic capabilities of Pure you'll just have to run the script with the interpreter. This isn't a big deal either, only the startup will be somewhat slower because the script is compiled on the fly. Once the JIT has done its thing the "interpreted" script will run every bit as fast as the "compiled" one, since in fact *both* are compiled (only at different times) to exactly the same code!

Also note that during a batch compilation, the compiled program is actually executed as usual, i.e., the script is also run *at compile time*. This might first seem to be a big annoyance, but it actually opens the door for some powerful programming techniques like [partial evaluation](#). It is also a necessity because of Pure's highly dynamic nature. For instance, Pure allows you to define constants by evaluating an arbitrary expression (see [Constant Definitions](#) below), and using `eval` a program can easily modify itself in even more unforeseeable ways. Therefore pretty much anything in your program can actually depend on previous computations performed while the program is being executed.

12.1 Example

For the sake of a concrete example, consider the following little script:

```
using system;

fact n = if n>0 then n*fact (n-1) else 1;

main n = do puts ["Hello, world!", str (map fact (1..n))];

if argc<=1 then () else main (sscanf (argv!1) "%d");
```


When invoked from the command line, with the number *n* as the first parameter, this program will print the string "Hello, world!" and the list of the first *n* factorials:

```
$ pure -x hello.pure 10
Hello, world!
[1,2,6,24,120,720,5040,40320,362880,3628800]
```

Note the condition on *argc* in the last line of the script. This prevents the program from producing an exception if no command line parameters are specified, so that the program can also be run interactively:

```
$ pure -i -q hello.pure
> main 10;
Hello, world!
[1,2,6,24,120,720,5040,40320,362880,3628800]
()
> quit
```

To turn the script into an executable, we just invoke the Pure interpreter with the *-c* option, using the *-o* option to specify the desired output file name:

```
$ pure -c hello.pure -o hello
$ ./hello 10
Hello, world!
[1,2,6,24,120,720,5040,40320,362880,3628800]
```

Next suppose that we'd like to supply the value *n* at *compile* rather than run time. To these ends we want to turn the value passed to the *main* function into a compile time constant, which can be done as follows:

```
const n = if argc>1 then sscanf (argv!1) "%d" else 10;
```

(Note that we provide 10 as a default if *n* isn't specified on the command line.)

Moreover, in such a case we usually want to skip the execution of the *main* function at compile time. The Pure runtime provides a special system variable *compiling* which holds a truth value indicating whether the program is actually running under the auspices of the batch compiler, so that it can adjust accordingly. In our example, the evaluation of *main* becomes:

```
if compiling then () else main n;
```

Our program now looks as follows:

```
using system;

fact n = if n>0 then n*fact (n-1) else 1;

main n = do puts ["Hello, world!", str (map fact (1..n))];

const n = if argc>1 then sscanf (argv!1) "%d" else 10;
if compiling then () else main n;
```

This script “specializes” `n` to the first (compile time) parameter when being batch-compiled, and it still works as before when we run it through the interpreter in both batch and interactive mode, too:

```
$ pure -i -q hello.pure
Hello, world!
[1,2,6,24,120,720,5040,40320,362880,3628800]
> main 5;
Hello, world!
[1,2,6,24,120]
()
> quit

$ pure -x hello.pure 7
Hello, world!
[1,2,6,24,120,720,5040]

$ pure -o hello -c -x hello.pure 7
$ ./hello
Hello, world!
[1,2,6,24,120,720,5040]
```

You’ll rarely need an elaborate setup like this, most of the time something like our simple first example will do the trick. But, as you’ve seen, Pure can easily do it.

12.2 Code Size and Unstripped Executables

By default, the batch compiler strips unused functions from the output code, to keep the code size small. You can disable this with the `-u` option, in which case the output code includes *all* functions defined in the compiled program or imported through a `using` clause, even if they don’t seem to be used anywhere. This considerably increases compilation times and makes the compiled executable much larger. For instance, on a 64 bit Linux systems with ELF binaries the executable of our `hello.pure` example is about thrice as large:

```
$ pure -o hello -c -x hello.pure 7 && ls -l hello
-rwxr-xr-x 1 ag users 178484 2010-01-12 06:21 hello
$ pure -o hello -c -u -x hello.pure 7 && ls -l hello
-rwxr-xr-x 1 ag users 541941 2010-01-12 06:21 hello
```

(Note that even the stripped executable is fairly large when compared to compiled C code, as it still contains the symbol table of the entire program, which is needed by the runtime environment.)

Stripped executables should be fine for most purposes, but you have to be careful when using `eval` in your compiled program. The compiler only does a *static* analysis of which functions might be reached from the initialization code (i.e., toplevel expressions and `let` bindings). It does *not* take into account code run via the `eval` routine. Thus, functions used only in `eval`ed code will be stripped from the executable, as if they were never defined at all. If such a function is then being called using `eval` at runtime, it will evaluate to a plain constructor symbol.

If this is a problem then you can either use the `-u` option to produce an unstripped executable, or you can force functions to be included in the stripped executable with the `--required` pragma (cf. [Code Generation Options](#)). For instance:

```
#!/ --required foo
foo x = bar (x-1);
eval "foo 99";
```

There is another code generation option which may have a substantial effect on code size, namely the `--noconst` option. Normally, constant values defined in a `const` definition are precomputed at compile time and then stored in the generated executable; this reduces startup times but may increase the code size considerably if your program contains big constant values such as lists. If you prefer smaller executables then you can use the `--noconst` option to force the value of the constant to be recomputed at run time (which effectively turns the constant into a kind of read-only variable). For instance:

```
#!/ --noconst
const xs = 1L..100000L;
sum = foldl (+) 0;

using system;
puts $ str $ sum xs;
```

On my 64 bit Linux system this produces a 187115 bytes executable. Without `--noconst` the code becomes almost an order of magnitude larger in this case (1788699 bytes). On the other hand, the smaller executable also takes a little longer to run since it must first recompute the value of the list constant at startup. So you have to consider the tradeoffs in a given situation. Usually big executables aren't much of a problem on modern operating systems, but if your program contains a lot of big constants then this may become an important consideration. However, if a constant value takes a long time to compute then you'll be better off with the default behaviour of precomputing the value at compile time.

12.3 Other Output Code Formats

Note that while the batch compiler generates native executables by default, it can just as well create object files which can be linked into other C/C++ programs and libraries:

```
$ pure -o hello.o -c -x hello.pure 7
```

The `.o` extension tells the compiler that you want an object file. When linking the object module, you also need to supply an initialization routine which calls the `__pure_main__` function in `hello.o` to initialize the compiled module. This routine is declared in C/C++ code as follows:

```
extern "C" void __pure_main__(int argc, char** argv);
```

As indicated, `__pure_main__` is to be invoked with two parameters, the argument count and NULL-terminated argument vector which become the `argc` and the `argv` of the Pure program, respectively. (You can also just pass 0 for both arguments if you don't need to

supply command line parameters.) The purpose of `__pure_main__` is to initialize a shell instance of the Pure interpreter which provides the minimal runtime support necessary to execute the Pure program, and to invoke all “initialization code” (variable definitions and toplevel expressions) of the program itself.

A minimal C main function which does the job of initializing the Pure module looks as follows:

```
extern void __pure_main__(int argc, char** argv);

int main(int argc, char** argv)
{
    __pure_main__(argc, argv);
    return 0;
}
```

If you link the main routine with the Pure module, don’t forget to also pull in the Pure runtime library. Assuming that the above C code is in `pure_main.c`:

```
$ gcc -c pure_main.c -o pure_main.o
$ g++ -o hello hello.o pure_main.o -lpure
$ ./hello
Hello, world!
[1,2,6,24,120,720,5040]
```

(The C++ compiler is used as the linker here so that the standard C++ library gets linked in, too. This is necessary because Pure’s runtime library is actually written in C++.)

In fact, this is pretty much what `pure -c` actually does for you when creating an executable.

If your script loads dynamic libraries (using “`lib:...`”;) then you’ll also have to link with those; *all* external references have to be resolved at compile time. This is taken care of automatically when creating executables. Otherwise it is a good idea to run `pure -c` with the `-v0100` verbosity option so that it prints the libraries to be linked (in addition to the commands which are invoked in the compilation process):

```
$ pure -v0100 -c hello.pure -o hello.o
opt -f -std-compile-opts hello.o.bc | llc -f -o hello.o.s
gcc -c hello.o.s -o hello.o
Link with: g++ hello.o -lpure
```

Well, we already knew that, so let’s consider a slightly more interesting example from Pure’s ODBC module:

```
$ pure -v0100 -c pure-odbc/examples/menagerie.pure -o menagerie.o
opt -f -std-compile-opts menagerie.o.bc | llc -f -o menagerie.o.s
gcc -c menagerie.o.s -o menagerie.o
Link with: g++ menagerie.o /usr/local/lib/pure/odbc.so -lpure
$ g++ -shared -o menagerie.so menagerie.o /usr/local/lib/pure/odbc.so -lpure
```

Note that the listed link options are necessary but might not be sufficient; `pure -c` just makes a best guess based on the Pure source. On most systems this will be good enough, but if it

isn't, you can just add options to the linker command as needed to pull in additional required libraries.

As this last example shows, you can also create shared libraries from Pure modules. However, on some systems (most notably x86_64), this requires that you pass the `-fPIC` option when batch-compiling the module, so that position-independent code is generated:

```
$ pure -c -fPIC pure-odbc/examples/menagerie.pure -o menagerie.o
```

Also note that even when building a shared module, you'll have to supply an initialization routine which calls `__pure_main__` somewhere.

Last but not least, `pure -c` can also generate just plain LLVM assembler code:

```
pure -c hello.pure -o hello.ll
```

Note the `.ll` extension; this tells the compiler that you want an LLVM assembler file. An LLVM bitcode file can be created just as easily:

```
pure -c hello.pure -o hello.bc
```

In these cases you'll have to have to handle the rest of the compilation yourself. This gives you the opportunity, e.g., to play with special optimization and code generation options provided by the LLVM toolchain. Please refer to the [LLVM documentation](#) (in particular, the description of the `opt` and `llc` programs) for details.

12.4 Calling Pure Functions From C

Another point worth mentioning here is that you can't just call Pure functions in a batch-compiled module directly. That's because in order to call a Pure function, at least in the current implementation, you have to set up a Pure stack frame for the function. However, there's a convenience function called `pure_funcall` in the runtime API to handle this. This function takes a pointer to the Pure function, the argument count and the arguments themselves (as `pure_expr*` objects) as parameters. For instance, here is a `pure_main.c` module which can be linked against the `hello.pure` program from above, which calls the `fact` function from the Pure program:

```
#include <stdio.h>
#include <pure/runtime.h>

extern void __pure_main__(int argc, char** argv);
extern pure_expr *fact(pure_expr *x);

int main()
{
    int n = 10, m;
    __pure_main__(0, NULL);
    if (pure_is_int(pure_funcall(fact, 1, pure_int(n)), &m))
        printf("fact %d = %d\n", n, m);
    return 0;
}
```

And here's how you can compile, link and run this program:

```
$ pure -o hello.o -c -x hello.pure 7
$ gcc -o pure_main.o -c pure_main.c
$ g++ -o myhello hello.o pure_main.o -lpure
$ ./myhello
Hello, world!
[1,2,6,24,120,720,5040]
fact 10 = 3628800
```

Note that the first two lines are output from the Pure program; the last line is what gets printed by the main routine in pure_main.c.

13 Caveats and Notes

This section is a grab bag of casual remarks, useful tips and tricks, and information on common pitfalls, quirks and limitations of the current implementation and how to deal with them.

13.1 Purity

People keep asking me what's so "pure" about Pure. The long and apologetic answer is that Pure tries to stay as close as possible to the spirit of term rewriting without sacrificing practicality. It's possible and in fact quite easy to write purely functional programs in Pure, and you're encouraged to use it that way when possible. On the other hand, Pure doesn't get in your way if you want to call external operations with side effects; it does allow you to call any C function after all.

The short answer is that I simply liked the name, and there wasn't any programming language named "Pure" yet (quite a feat nowadays), so there's one now. :)

13.2 Backward Compatibility

Pure 0.7 introduced built-in matrix structures, which called for some minor changes in the syntax of comprehensions and arithmetic sequences. Specifically, the template expression and generator/filter clauses of a comprehension are now separated with `|` instead of `;`. Moreover, arithmetic sequences with arbitrary stepsize are now written `x:y..z` instead of `x,y..z`, and the `'.'` operator now has a higher precedence than the `','` operator. This makes writing matrix slices like `x!!(i..j,k..l)` much more convenient.

In Pure 0.13 the naming of the logical and bitwise operations was changed, so that these are now called `~`, `&&`, `||` and `not/and/or`, respectively. (Previously, `~` was used for bitwise, not for logical negation, which was rather inconsistent, albeit compatible with the naming of the not operation in Haskell and ML.) Also, to stay in line with this naming scheme, inequality was renamed to `~=` (previously `!=`).

Pure 0.14 introduced the namespaces feature. Consequently, the scope of private symbols is now confined to a namespace rather than a source module; scripts making use of private symbols need to be adapted accordingly. Also note that syntax like `foo::int` may now also denote a qualified symbol rather than a tagged variable, if `foo` has been declared as a namespace. You can work around such ambiguities by renaming the variable, or by placing spaces around the `::` delimiter (these aren't permitted in a qualified symbol, so the construct `foo :: int` is always interpreted as a tagged variable, no matter whether `foo` is also a valid namespace).

Pure 0.26 extended the namespaces feature to add support for hierarchical namespaces. This means that name lookup works in a slightly different fashion now (see [Hierarchical Namespaces](#) for details), but old code which doesn't use the new feature should continue to work unchanged.

Pure 0.26 also changed the nullary keyword to `nonfix`, which is more consistent with the other kinds of fixity declarations. Moreover, the parser was enhanced so that it can cope with a theoretically unbounded number of precedence levels, and the system of standard operators in the prelude was modified so that it becomes possible to sneak in new operator symbols with ease; details can be found in the [Symbol Declarations](#) section.

Pure 0.41 added support for optimization of indirect tail calls, so that any previous restrictions on the use of tail recursion in indirect function calls and mutually recursive globals have been removed. Moreover, the logical operators `&&` and `||` are now tail-recursive in their second operand and can also be extended with user-defined equations, just like the other builtins. Note that this implies that the values returned by `&&` and `||` aren't normalized to the values 0 and 1 any more (this isn't possible with tail call semantics). If you need this then you'll have to make sure that either the operands are already normalized, or you'll have to normalize the result yourself.

Also, as of Pure 0.41 the batch compiler produces stripped executables by default. To create unstripped executables you now have to use the `-u` option, see [Code Size and Unstripped Executables](#) for details. The `-s` option to produce stripped executables is still provided for backward compatibility, but it won't have any effect unless you use it to override a previous `-u` option.

Pure 0.43 changed the rules for looking up symbols in user-defined namespaces. Unqualified symbols are now created in the current (rather than the global) namespace by default, see [Symbol Lookup and Creation](#) for details. The `-w` option can be used to get warnings about unqualified symbols which are resolved to a different namespace than previously. It also provides a means to check your scripts for implicit declarations which might indicate missing or mistyped function symbols.

Pure 0.45 added support for checking arbitrary pointer types in the C interface, so that you don't have to worry about passing the wrong kinds of pointers to system and library routines any more. Moreover, the interpretation of numeric pointer arguments (`int*` etc.) was changed to bring them in line with the other new numeric matrix conversions (`int**` etc.). In particular, the matrix data can now be modified in-place and type checking is more strict (`int*` requires an `int` matrix, etc.). Also, there's now support for `argv`-style vector arguments (`char**` and `void**`). Please see the [C Types](#) section for details.

13.3 Error Recovery

The parser uses a fairly simplistic panic mode error recovery which tries to catch syntax errors at the toplevel only. This seems to work reasonably well, but might catch some errors much too late. Unfortunately, Pure's terseness makes it rather difficult to design a better scheme. As a remedy, the parser accepts an empty definition (just `; by itself`) at the toplevel only. Thus, in interactive usage, if the parser seems to eat away your input without doing anything, entering an extra semicolon or two should break the spell, putting you back at the toplevel where you can start typing the definition again.

13.4 The `__show__` Function

As of Pure 0.6, the interpreter provides a "hook" to override the print representations of expressions at runtime by means of the `__show__` function, which works in a fashion similar to Haskell's `show` function. This feature is still a bit experimental, but seems to work reasonably well for the purposes for which it is intended.

`__show__` is just an ordinary Pure function expected to return a string with the desired custom representation of a normal form value given as the function's single argument. This function is not defined by default, so you are free to add any rules that you want. The interpreter prints the strings returned by `__show__` just as they are. It will not check whether they conform to Pure syntax and/or semantics, or modify them in any way.

Custom print representations are most useful for interactive purposes, if you're not happy with the default print syntax of some kinds of objects. One particularly useful application of `__show__` is to change the format of numeric values. Here are some examples:

```
> using system;
> __show__ x::double = sprintf "%.6f" x;
> 1/7;
0.142857
> __show__ x::int = sprintf "0x%x" x;
> 1786;
0x6fa
> using math;
> __show__ (x::double +: y::double) = sprintf "%.6f+%.6fi" (x,y);
> cis (-pi/2);
0.000000+-1.000000i
```

The prelude function `str`, which returns the print representation of any Pure expression, calls `__show__` as well:

```
> str (1/7);
"0.142857"
```

Conversely, you can call the `str` function from `__show__`, but in this case it always returns the default representation of an expression. This prevents the expression printer from going recursive, and allows you to define your custom representation in terms of the default one. E.g., the following rule removes the `L` suffixes from bigint values:


```
> __show__ x::bigint = init (str x);
> fact n = foldl (*) 1L (1..n);
> fact 30;
265252859812191058636308480000000
```

Of course, your definition of `__show__` can also call `__show__` itself recursively to determine the custom representation of an object.

One case which needs special consideration are thunks (futures). The printer will *never* use `__show__` for those, to prevent them from being forced inadvertently. In fact, you *can* use `__show__` to define custom representations for thunks, but only in the context of a rule for other kinds of objects, such as lists. For instance:

```
> nonfix ...;
> __show__ (x:xs) = str (x:...) if thunkp xs;
> 1:2:(3..inf);
1:2:3:...
```

Another case which needs special consideration are numeric matrices. For efficiency, the expression printer will always use the default representation for these, unless you override the representation of the matrix as a whole. E.g., the following rule for double matrices mimics Octave's default output format (for the sake of simplicity, this isn't perfect, but you get the idea):

```
> __show__ x::matrix =
>   strcat [printf j (x!(i,j))|i=0..n-1; j=0..m-1] + "\n"
> with printf 0 = sprintf "%10.5f"; printf _ = sprintf "%10.5f" end
> when n,m = dim x end if dmatrixp x;
> {1.0,1/2;1/3,4.0};
  1.00000  0.50000
  0.33333  4.00000
```

Finally, by just purging the definition of the `__show__` function you can easily go back to the standard print syntax:

```
> clear __show__
> 1/7; 1786; cis (-pi/2);
0.142857142857143
1786
6.12303176911189e-17+-1.0
```

Note that if you have a set of definitions for the `__show__` function which should always be loaded at startup, you can put them into the interpreter's interactive startup files, see [Interactive Usage](#).

13.5 Non-Linear Patterns

As explained in section [Patterns](#), Pure allows multiple occurrences of the same variable in a pattern (so-called non-linearities):

```
foo x x = x;
```

This rule will only be matched if both occurrences of `x` are bound to the same value. More precisely, the two instances of `x` will be checked for syntactic equality during pattern matching, using the *same* primitive provided by the prelude. This may need time proportional to the sizes of both argument terms, and thus become quite costly for big terms. In fact, *same* might not even terminate at all if the compared terms are both infinite lazy data structures, such as in `foo (1..inf) (1..inf)`. So you have to be careful to avoid such uses.

When using non-linearities in conjunction with “as” patterns, you also have to make sure that the “as” variable does not occur inside the corresponding subpattern. Thus a definition like the following is illegal:

```
> foo xs@(x:xs) = x;
<stdin>, line 1: error in pattern (recursive variable 'xs')
```

The explanation is that such a pattern couldn’t possibly be matched by a finite list anyway. Indeed, the only match for `xs@(x:xs)` would be an infinite list of `x`’s, and there’s no way that this condition could be verified in a finite amount of time. Therefore the interpreter reports a “recursive variable” error in such situations.

13.6 “As” Patterns

In the current implementation, “as” patterns cannot be placed on the “spine” of a function definition. Thus rules like the following, which have the pattern somewhere in the head of the left-hand side, will all provoke an error message from the compiler:

```
a@foo x y = a,x,y;
a@(foo x) y = a,x,y;
a@(foo x y) = a,x,y;
```

This is because the spine of a function application is not available when the function is called at runtime. “As” patterns in pattern bindings (`let`, `const`, `case`, `when`) are not affected by this restriction since the entire value to be matched is available at runtime. For instance:

```
> case bar 99 of y@(bar x) = y,x+1; end;
bar 99,100
```

13.7 Head = Function

“As” patterns are also a useful device if you need to manipulate function applications in a generic way. Note that the “head = function” rule means that the head symbol `f` of an application `f x1 ... xn` occurring on (or inside) the left-hand side of an equation, variable binding, or pattern-matching lambda expression, is always interpreted as a literal function symbol (not a variable). This implies that you cannot match the “function” component of an application against a variable, at least not directly. An anonymous “as” pattern like `f@_` does the trick, however, since the anonymous variable is always recognized, even if it occurs as

the head symbol of a function application. Here's a little example which demonstrates how you can convert a function application to a list containing the function and all arguments:

```
> foo x = a [] x with a xs (x@_ y) = a (y:xs) x; a xs x = x:xs end;  
> foo (a b c d);  
[a,b,c,d]
```

This may seem a little awkward, but as a matter of fact the “head = function” rule is quite useful since it covers the common cases without forcing the programmer to declare “constructor” symbols (except nonfix symbols). On the other hand, generic rules operating on arbitrary function applications are not all that common, so having to “escape” a variable using the anonymous “as” pattern trick is a small price to pay for that convenience.

Sometimes you may also run into the complementary problem, i.e., to match a function argument against a given function. Consider this code fragment:

```
foo x = x+1;  
foop f = case f of foo = 1; _ = 0 end;
```

You might expect foop to return true for foo, and false on all other values. Better think again, because in reality foop will always return true! In fact, the Pure compiler will warn you about the second rule of the case expression not being used at all:

```
> foop 99;  
warning: rule never reduced: _ = 0;  
1
```

This happens because an identifier on the left-hand side of a rule, which is neither the head symbol of a function application nor a nonfix symbol, is always considered to be a variable (cf. [Parameters in Equations](#)), even if that symbol is defined as a global function elsewhere. So foo isn't a literal name in the above case expression, it's a variable! (As a matter of fact, this is rather useful, since otherwise a rule like `f g = g+1` would suddenly change meaning if you happen to add a definition like `g x = x-1` somewhere else in your program, which certainly isn't desirable.)

A possible workaround is to “escape” the function symbol using an empty namespace qualifier:

```
foop f = case f of ::foo = 1; _ = 0 end;
```

This trick works in case expressions and function definitions, but fails in circumstances in which qualified variable symbols are permitted (i.e., in variable and constant definitions). A better solution is to employ the syntactic equality operator `===` defined in the prelude to match the target value against the function symbol. This allows you to define the foop predicate as follows:

```
> foop f = f===foo;  
> foop foo, foop 99;  
1,0
```

Another way to deal with the situation would be to just declare foo as a nonfix symbol. However, this makes the foo symbol “precious”, i.e., after such a declaration it cannot be

used as a local variable anymore. It's usually a good idea to avoid that kind of thing, at least for generic symbols, so the above solution is preferred in this case.

13.8 With and when

A common source of confusion is that Pure provides two different constructs to bind local function and variable symbols, respectively. This distinction is necessary because Pure does not segregate defined functions and constructors, and thus there is no magic to figure out whether an equation like `foo x = y` by itself is meant as a definition of a function `foo` with formal parameter `x` and return value `y`, or a pattern binding defining the local variable `x` by matching the pattern `foo x` against the value of `y`. The `with` construct does the former, when the latter. (As a mnemonic, you may consider that `when` conveys a sense of time, as the individual variable definitions in a `when` clause are executed in order, while the function definitions in a `with` clause are all done simultaneously.)

Another speciality is that `with` and `when` clauses are tacked on to the end of the expression they belong to. This mimics mathematical language and makes it easy to read and understand a definition in a "top-down" fashion. This style differs considerably from other block-structured programming languages, however, which often place local definitions in front of the code they apply to. To grasp the operational meaning of such nested definitions, it can be helpful to read the nested scopes "in reverse" (from bottom to top). Some people also prefer to write their programs that way. In difference to Haskell and ML which have `let` expressions to support that kind of notation, Pure doesn't provide any special syntax for this. But note that you can always write `when` clauses in the following style which places the "body" at the bottom of the clause:

```
result when
  y = foo (x+1);
  z = bar y;
  result = baz z;
end;
```

This doesn't incur any overhead, since the compiler will always eliminate the trivial "tail binding" for the result value. E.g., the above will compile to exactly the same code as:

```
baz z when
  y = foo (x+1);
  z = bar y;
end;
```

13.9 Numeric Calculations

If possible, you should decorate numeric variables on the left-hand sides of function definitions with the appropriate type tags, like `int` or `double`. This often helps the compiler to generate better code and makes your programs run faster. The `|` syntax makes it easy to add the necessary specializations of existing rules to your program. E.g., taking the polymorphic implementation of the factorial as an example, you only have to add a left-hand side with

the appropriate type tag to make that definition go as fast as possible for the special case of machine integers:

```
fact n::int    |
fact n        = n*fact(n-1) if n>0;
              = 1 otherwise;
```

(This obviously becomes unwieldy if you have to deal with several numeric arguments of different types, however, so in this case it is usually better to just use a polymorphic rule.)

Also note that `int` (the machine integers), `bigint` (the GMP “big” integers) and `double` (floating point numbers) are all different kinds of objects. While they can be used in mixed operations (such as multiplying an `int` with a `bigint` which produces a `bigint`, or a `bigint` with a `double` which produces a `double`), the `int` tag will only ever match a machine `int`, *not* a `bigint` or a `double`. Likewise, `bigint` only matches `bigints` (never `int` or `double` values), and `double` only `doubles`. Thus, if you want to define a function operating on different kinds of numbers, you’ll also have to provide equations for all the types that you need (or a polymorphic rule which catches them all). This also applies to equations matching against constant values of these types. In particular, a small integer constant like `0` only matches machine integers, not `bigints`; for the latter you’ll have to use the “big L” notation `0L`. Similarly, the constant `0.0` only matches `doubles`, but not `ints` or `bigints`.

13.10 Constant Definitions

Constants differ from variables in that they cannot be redefined (that’s their main purpose after all) so that their values, once defined, can be substituted into other definitions which use them. For instance:

```
> const c = 2;
> foo x = c*x;
> show foo
foo x = 2*x;
> foo 99;
198
```

While a variable can be rebound to a new value at any time, you will get an error message if you try to do this with a constant:

```
> const c = 3;
<stdin>, line 5: symbol 'c' is already defined as a constant
```

Note that in interactive mode you can work around this by purging the old definition with the `clear` command. However, this won’t affect any earlier uses of the symbol:

```
> clear c
> const c = 3;
> bar x = c*x;
> show foo bar
bar x = 3*x;
foo x = 2*x;
```

(You'll also have to purge any existing definition of a variable if you want to redefine it as a constant, or vice versa, since Pure won't let you redefine an existing constant or variable as a different kind of symbol. The same also holds if a symbol is currently defined as a function or a macro.)

Constants can also be used in patterns (i.e., on the left-hand side of a rule in a definition or a case expression), but only if you also declare the corresponding symbol as `nonfix`. This is useful, e.g., if you'd like to use constants such as `true` and `false` on the left-hand side of a definition, just like other `nonfix` symbols:

```
> show false true
const false = 0;
const true = 1;
> nonfix false true;
> check false = "no"; check true = "yes";
> show check
check 0 = "no";
check 1 = "yes";
> check (5>0);
"yes"
```

Note that without the `nonfix` declaration, the above definition of `check` wouldn't work as intended, because the `true` and `false` symbols on the left-hand side of the two equations would be interpreted as local variables. Also note that the standard library never declares any constant symbols as `nonfix`, since once a symbol is `nonfix` there's no going back. Thus the library leaves this to the programmer to decide.

As the value of a constant is known at compile time, the compiler can apply various optimizations to uses of such values. In particular, the Pure compiler inlines constant scalars (numbers, strings and pointers) by literally substituting their values into the output code, and it also precomputes simple constant expressions involving only (machine) integer and double values. Example:

```
> extern double atan(double);
> const pi = 4*atan 1.0;
> show pi
const pi = 3.14159265358979;
> foo x = 2*pi*x;
> show foo
foo x = 6.28318530717959*x;
```

In addition, the LLVM backend eliminates dead code automatically, so you can employ a constant to configure your code for different environments, without any runtime penalties:

```
const win = index sysinfo "mingw32" >= 0;
check boy = bad boy if win;
           = good boy otherwise;
```

In this case the code for one of the branches of `check` will be completely eliminated, depending on the outcome of the configuration check.

For efficiency, constant aggregates (lists, tuples, matrices and other kinds of non-scalar

terms) receive special treatment. Here, the constant is computed once and stored in a read-only variable which then gets looked up at runtime, just like an ordinary global variable. However, there's an important difference: If a script is batch-compiled (cf. [Batch Compilation](#)), the constant value is normally computed *at compile time only*; when running the compiled executable, the constant value is simply reconstructed, which is often much more efficient than recomputing its value. For instance, you might use this to precompute a large table whose computation may be costly or involve functions with side effects:

```
const table = [foo x | x = 1..1000000];
process table;
```

Note that this only works with `const` values which are completely determined at compile time. If a constant contains run time objects such as pointers and (local) functions, this is impossible, and the batch compiler will instead create code to recompute the value of the constant at run time. For instance, consider:

```
using system;
const p = malloc 100;
foo p;
```

Here, the value of the pointer `p` of course critically depends on its computation (involving a side effect which sets aside a corresponding chunk of memory). It would become unusable without actually executing the initialization, so the compiler generates the appropriate run time initialization code in this case. For all practical purposes, this turns the constant into a read-only variable. (There's also a code generation option to force this behaviour even for "normal" constants for which it's not strictly necessary, in order to create smaller executables; see [Code Size and Unstripped Executables](#) for details.)

13.11 External C Functions

The interpreter always takes your extern declarations of C routines at face value. It will not go and read any C header files to determine whether you actually declared the function correctly! So you have to be careful to give the proper declarations, otherwise your program might well give a segfault when calling the function. (This problem can to some extent be alleviated by using the bitcode interface. See [Importing LLVM Bitcode](#) and [Inline Code](#) in the [C Interface](#) section.)

Another limitation of the C interface is that it does not offer any special support for C structs and C function parameters. However, an optional addon module is available which interfaces to the [libffi](#) library to provide that kind of functionality, please see the description of the [pure-ffi](#) module on the Pure website for details.

Last but not least, to make it easier to create Pure interfaces to large C libraries, there's a separate [pure-gen](#) program available at the Pure website. This program takes a C header (.h) file and creates a corresponding Pure module with definitions and extern declarations for the constants and functions declared in the header. Please refer to the `pure-gen(1)` manual page for details.

13.12 Calling Special Forms

Special forms are recognized at compile time only. Thus the `catch` function, as well as `quote` and the operators `&&`, `||`, `$$` and `&`, are only treated as special forms in direct (saturated) calls. They can still be used if you pass them around as function values or in partial applications, but in this case they lose all their special call-by-name argument processing.

13.13 The Quote

Like in Lisp, the `quote` special form can be used to construct literal expressions which can then be manipulated before actually evaluating them using the built-in `eval` function. However, there are some notable differences to the Lisp version of `quote`:

- Local variables (i.e., variables bound by the left-hand side of an equation, a lambda expression or a `case` or `when` construct) can never be quoted, so the corresponding value gets substituted even into quoted expressions. Thus, e.g., `'(2*42+2^n)` when `n = 12` end evaluates to `2*42+2^12` rather than `2*42+2^n`.
- Only simple expressions can be quoted in Pure. Special constructs embedded in a quoted expression, such as conditionals and local bindings, are evaluated as usual. For instance, `'(2*42+(2^n when n = 2*6 end))` evaluates to `2*42+4096.0` rather than quoting the embedded `when` expression.

While it's possible to define a Lisp-like `quasiquote` in Pure (see the [Recursive Macros](#) section for a simplified version, a full implementation can be found in the Pure library sources), it's usually not needed. As discussed above, substitution of local variables is performed even in quoted expressions, so that it is quite easy to fill in variable parts in a quoted "template" expression. For instance:

```
> (\x -> '(2*x+1)) 99;  
2*99+1  
> foo x = '(2*x+1);  
> foo 99; foo $ '(7/y);  
2*99+1  
2*(7/y)+1
```

In fact, it is possible to perform arbitrary computations right in the middle of a quoted expression, using one of the special `if-then-else`, `case`, `when` and `with` expressions, since these are never quoted either. Example:

```
> '(x+1 when x = '(2*3) end);  
2*3+1
```

Another useful feature of Lisp's `quasiquote` is the capability to splice arguments into a function application. It is possible to achieve pretty much the same in Pure with the following variation of the `$` operator which "curries" its second (tuple) operand:

```
infixr 0 $@ ;  
f $@ () = f;
```



```
f $@ (x,xs) = f x $@ xs;  
f $@ x      = f x;
```

Now you can write, e.g.:

```
> '(foo 1 2) $@ '(2/3,3/4);  
foo 1 2 (2/3) (3/4)
```

One shortcoming of Pure’s quote is that there is no way to quote special constructs such as lambdas. Macro expansion *is* inhibited in quoted expressions, however, so it is possible to work around this limitation by defining a custom special form to be used as a symbolic representation for, say, a lambda expression, which reduces to a real lambda when evaluated. To these ends, the eval function can be invoked with a string argument as follows:

```
def lambda x y = eval $ "\\ " +str ('x') + " -> " +str ('y');
```

Example:

```
> let l = 'lambda x (x+1); l;  
lambda x (x+1)  
> let f = eval l; f; f 9;  
#<closure 0x7fdc3ca45be8>  
10
```

Other special constructs, such as case, when and with can be handled in a similar fashion.

13.14 Laziness

Pure does lazy evaluation in the same way as [Alice ML](#), providing an explicit operation (&) to defer evaluation and create a “future” which is called by need. However, note that like any language with a basically eager evaluation strategy, Pure cannot really support lazy evaluation in a fully automatic way. That is, coding an operation so that it works with infinite data structures usually requires additional thought, and sometimes special code will be needed to recognize futures in the input and handle them accordingly. This can be hard, but of course in the case of the prelude operations this work has already been done for you, so as long as you stick to these, you’ll never have to think about these issues. (It should be noted here that lazy evaluation has its pitfalls even in fully lazy FPLs, such as hidden memory leaks and other kinds of subtle inefficiencies or non-termination issues resulting from definitions being too lazy or not lazy enough. You can read about that in any good textbook on Haskell.)

The prelude goes to great lengths to implement all standard list operations in a way that properly deals with streams (a.k.a. lazy lists). What this all boils down to is that all list operations which can reasonably be expected to operate in a lazy way on streams, will do so. (Exceptions are inherently eager operations such as #, reverse and foldl.) Only those portions of an input stream will be traversed which are strictly required to produce the result. For most purposes, this works just like in fully lazy FPLs such as Haskell. However, there are some notable differences:

- Since Pure uses dynamic typing, some of the list functions may have to peek ahead one element in input streams to check their arguments for validity, meaning that these functions will be slightly more eager than their Haskell counterparts.
- Pure’s list functions never produce truly cyclic list structures such as the ones you get, e.g., with Haskell’s `cycle` operation. (This is actually a good thing, because the current implementation of the interpreter cannot garbage-collect cyclic expression data.) Cyclic streams such as `cycle [1]` or `fix (1:)` will of course work as expected, but, depending on the algorithm, memory usage may increase linearly as they are traversed.
- Pattern matching is always refutable (and therefore eager) in Pure. If you need something like Haskell’s irrefutable matches, you’ll have to code them explicitly using futures. See the definition of the `unzip` function in the prelude for an example showing how to do this.

Here are some common pitfalls with lazy data structures in Pure that you should be aware of:

- Laziness and side effects don’t go well together, as most of the time you can’t be sure when a given thunk will be executed. So as a general guideline you should avoid side effects in thunked data structures. If you can’t avoid them, then at least make sure that all accesses to the affected resources are done through a single instance of the thunked data structure. E.g., the following definition lets you create a stream of random numbers:

```
> using math;  
> let xs = [random | _ = 1..inf];
```

This works as expected if only a single stream created with `random` exists in your program. However, as the `random` function in the `math` module modifies an internal data structure to produce a sequence of pseudorandom numbers, using two or more such streams in your program will in fact modify the same underlying data structure and thus produce two disjoint subsequences of the same underlying pseudorandom sequence which might not be distributed uniformly any more.

- You should avoid keeping references to potentially big (or even infinite) thunked data structures when traversing them (unless you specifically need to memoize the entire data structure). In particular, if you assign such a data structure to a local variable, the traversal of the data structure should then be invoked as a tail call. If you fail to do this, it forces the entire memoized part of the data structure to stay in main memory while it is being traversed, leading to rather nasty memory leaks. Please see the `all_primes` function in [Lazy Evaluation and Streams](#) for an example.

13.15 Reflection

Pure versions since 0.12 offer some basic reflection capabilities via the `evalcmd` primitive. This function provides access to interactive commands like `clear`, `save` and `show`, which enable you to inspect and modify the running program. The only “canonical” way to represent an entire Pure program in Pure itself is the program text, hence `evalcmd` only provides a

textual interface at this time. But of course custom higher-level representations can be built on top of that, similar to those discussed in section [The Quote](#).

Here's an example showing what can be done using the `show` command and a little bit of trivial text processing. The following `sym_info` function retrieves information about a given collection of global symbols in a way which can be processed in a Pure program. The `cat` argument can be any combination of the letters "c", "v", "f" and "m" denoting the categories of constants, variables, functions and macros, respectively. (You can also just leave this empty if you don't care about the type of symbol.) The `pat` argument is a shell-like glob pattern for the name of symbols which should be listed (just "*" matches all symbols). The result is a list of tuples (name, value, cat, descr) with the name of the symbol and its value, as well as the category and description of the symbol, as provided by `show -s`.

```
using system;
sym_info cat::string pat::string
= [name,eval ("("+name+")"),descr | name,descr = info]
when
  // Get the info about matching symbols from the 'show' command.
  info = evalcmd $ sprintf "show -sg%s %s" (cat,pat);
  // Split into lines.
  info = if null info then [""] else split "\n" $ init info;
  // Get rid of the last line with the summary information.
  info = init info;
  // Retrieve the information that we need.
  info = [x | x@(s,_) = map fields info;
  // Get rid of extra lines containing extern and fixity declarations.
    s ~= "extern" && s ~= "nonfix" && s ~= "outfix" &&
    s ~= "prefix" && s ~= "postfix" && ~fnmatch "infix*" s 0];
end with
  // Regex call to split the summary information about one symbol, as
  // returned by 'show -s', into the name and description parts.
  fields s::string = tuple $
    [info!2 | info = tail $ regs $ reg_info $
      regex "([^\ ]+)[\ ]+([a-z]*)[ \ ]*(.*)" REG_EXTENDED s 0];
end;
```

E.g., this call retrieves information about all defined macros:

```
> sym_info "m" "*";
[("$",("$"),"mac","2 args, 1 rules"),(".",("."),"mac","3 args, 1 rules"),
("void","void","mac","1 args, 6 rules")]
```

13.16 Hygienic Macros

As mentioned in the [Macro Hygiene](#) section, Pure macros are lexically scoped and thus “hygienic”. Macro hygiene is a somewhat esoteric topic for most programmers, so let us take a brief look at what it's all about. The problem avoided by hygienic macros is that of *name capture*. There are actually two kinds of name capture which may occur in unhygienic macro systems:

- A free symbol in the macro *body* inadvertently becomes bound to the value of a local symbol in the context in which the macro is called.
- A free symbol in the macro *call* inadvertently becomes bound to the value of a local symbol in the macro body.

Pure's hygienic macros avoid both pitfalls. Here is an example for the first form of name capture:

```
> def G x = x+y;  
> G 10 when y = 99 end;  
10+y
```

Note that the expansion of the `G` macro correctly uses the global instance of `y`, even though `y` is locally defined in the context of the macro call. (Sometimes this form of name capture is actually used deliberately in order to make the macro use the binding of the symbol which is active at the point of the macro call. This never works in Pure, hence in such cases you will have to explicitly pass such symbols to the macro.)

In contrast, the second form of name capture is usually not intended, and is therefore more dangerous. Consider the following example:

```
> def F x = x+y when y = x+1 end;  
> F y;  
y+(y+1)
```

Pure again gives the correct result here. You'd have to be worried if you got $(y+1)+(y+1)$ instead, which would result from the literal expansion `y+y when y = y+1 end`, where the (free) variable `y` passed to `F` gets captured by the local binding of `y`. In fact, that's exactly what you get with C macros:

```
#define F(x) { int y = x+1; return x+y; }
```

Here `F(y)` expands to `{ int y = y+1; return y+y; }` which is usually *not* what you want.

There is also one Pure-related caveat here. The expression printer currently doesn't check for different bindings of the same variable identifier when it prints a (compile time) expression. E.g.:

```
> foo y = F y;  
> show foo  
foo y = y+y when y = y+1 end;
```

This *looks* as if `y` got captured, but in fact it's not, it's just the `show` command which displays the definition in an incorrect way. You can add the `-e` option to `show` which prints the deBruijn indices of locally bound symbols, then you see that the actual bindings are all right anyway (note that the number before the colon is the actual deBruijn index, the sequence of bits behind it is the subterm path):

```
> show -e foo  
foo y/*0:1*/ = y/*1:1*/+y/*0:*/ when y/*0:*/ = y/*0:1*/+1 end;
```

Alas, this means that if you use `dump` to write such a definition to a text file and read it back with `run` later, then you'll get the wrong definition. Currently you will have to correct this manually.

13.17 Stack Size and Tail Recursion

Pure programs may need a considerable amount of stack space to handle recursive function and macro calls, and the interpreter itself also takes its toll. So you should configure your system accordingly (8 MB of stack space is recommended for 32 bit systems, systems with 64 bit pointers probably need more). If the `PURE_STACK` environment variable is defined, the interpreter performs advisory stack checks on function entry and raises a Pure exception if the current stack size exceeds the given limit. The value of `PURE_STACK` should be the maximum stack size in kilobytes. Please note that this is only an advisory limit which does not change the program's physical stack size. Your operating system should supply you with a command such as `ulimit(1)` to set the real process stack size. (The `PURE_STACK` limit should be a little less than that, to account for temporary stack usage by the interpreter itself.)

Like Scheme, Pure does proper tail calls (if LLVM provides that feature on the platform at hand), so tail-recursive definitions should work fine in limited stack space. For instance, the following little program will loop forever if your platform supports the required optimizations:

```
loop = loop;  
loop;
```

This also works if your definition involves function parameters, guards and multiple equations, of course. Moreover, conditional expressions (`if-then-else`) are tail-recursive in both branches, and the logical operators `&&` and `||`, as well as the sequence operator `$$`, are tail-recursive in their second operand.

Also note that tail call optimization is *always* disabled if the debugger is enabled (`-g`). This makes it much easier to debug programs, but means that you may run into stack overflows when debugging a program that does deep tail recursion.

13.18 Handling of Asynchronous Signals

As described in section [Exception Handling](#), signals delivered to the process can be caught and handled with Pure's exception handling facilities. This has its limitations, however. Since Pure code cannot be executed directly from a C signal handler, checks for pending signals are only done on function entry. This means that in certain situations (such as the execution of an external C routine), delivery of a signal may be delayed by an arbitrary amount of time. Moreover, if more than one signal arrives between two successive signal checks, only the last one will be reported in the current implementation.

When delivering a signal which has been remapped to a Pure exception, the corresponding exception handler (if any) will be invoked as usual. Further signals are blocked while the exception handler is being executed.

A fairly typical case is that you have to handle signals in a tail-recursive function. This can be done with code like the following:

```
using system;

// Remap some common POSIX signals.
do (trap SIG_TRAP) [SIGHUP, SIGINT, SIGTERM];

loop = catch handler process $$ loop
with handler (signal k) = printf "Hey, I got signal %d.\n" k end;
process = sleep 1; // do something
```

Running the above loop function enters an endless loop reporting all signals delivered to the process. Note that to make this work, the tail-recursive invocation of loop must immediately follow the signal-handling code, so that signals don't escape the exception handler.

Of course, in a real application you'd probably want the loop function to carry around some data to be processed by the process routine, which then returns an updated value for the next iteration. This can be implemented as follows:

```
loop x = loop (catch handler (process x))
with handler (signal k) = printf "Hey, I got signal %d.\n" k $$ 0 end;
process x = printf "counting: %d\n" x $$ sleep 1 $$ x+1;
```

14 Copying

Pure comes with a fairly liberal license which lets you distribute your own Pure programs and extensions under a license of your choice and permits linking of commercial applications against the Pure runtime and the Pure standard library without requiring special permission. Moreover, the Pure interpreter (the pure main program), the Pure runtime library (libpure) and the Pure standard library (the Pure scripts in the lib folder distributed with the software) are distributed as free software, and you are welcome to modify and redistribute them under the appropriate license terms, as detailed below.

(The above explanations are not legal advice. Please read the full text of the licenses and consult qualified professional counsel for an interpretation of the license terms as they apply to you.)

The *Pure interpreter* is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

The *Pure runtime library* and the *Pure standard library* are also free software: you can redistribute them and/or modify them under the terms of the GNU Lesser General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

Pure is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR

PURPOSE.

Please see the [GNU General Public License](#) and the [GNU Lesser General Public License](#) for the precise license terms. You can also find the license conditions in the COPYING and COPYING.LESSER files accompanying the software. Also, please see the source code for the copyright and license notes pertaining to individual source files which are part of this software.

Pure uses [LLVM](#) as its compiler backend. LLVM is under Copyright (c) 2003-2010 by the University of Illinois at Urbana-Champaign, and is licensed under a 3-clause BSD-style license, please read COPYING.LLVM included in the distribution for the exact licensing terms. You can also find the LLVM license at the [LLVM](#) website.

15 Author

Albert Gräf <Dr.Graef@t-online.de>, Dept. of Computer Music, Johannes Gutenberg University of Mainz, Germany.

The author gratefully acknowledges the contributions by Scott E. Dillard, Rooslan S. Khayrov, Eddie Rucker, Libor Spacek, Jiri Spitz and Sergei Winitzki, as well as Toni Graffy, Michel Salim and Ryan Schmidt who maintain the SUSE Linux, Fedora Core and OSX packages, respectively. Thanks are also due to Vili Aapro, Alvaro Castro Castilla, John Cowan, Chris Double, Tim Haynes, Roman Neuhauser, Wm Leler, John Lunney and Max Wolf.

16 See Also

Aardappel Wouter van Oortmerssen's functional programming language based on term rewriting, <http://wouter.fov120.com/aardappel>.

Alice ML A version of ML (see below) from which Pure borrows its model of lazy evaluation, <http://www.ps.uni-sb.de/alice>.

Bertrand Wm Leler's constraint programming language based on term rewriting, <http://groups.google.com/group/bertrand-constraint>. See Wm Leler: Constraint Programming Languages: Their Specification and Generation. Addison-Wesley, 1988.

Clang The new C/C++/Objective C compiler designed specifically to work on top of LLVM, <http://clang.llvm.org>. Clang provides a comparatively light-weight alternative to llvm-gcc which is faster and has better and more friendly diagnostics than gcc.

Faust Grame's functional DSP programming language, <http://faust.grame.fr>.

GNU Multiprecision Library Free library for arbitrary precision arithmetic, <http://gmp.org>.

GNU Octave A popular high-level language for numeric applications and free MATLAB replacement, <http://www.gnu.org/software/octave>.

GNU Scientific Library A free software library for numeric applications, can be used with Pure's numeric matrices, <http://www.gnu.org/software/gsl>.

Haskell A popular non-strict FPL, <http://www.haskell.org>.

LLVM The LLVM code generator framework, <http://llvm.org>.

ML A popular strict FPL. See Robin Milner, Mads Tofte, Robert Harper, D. MacQueen: The Definition of Standard ML (Revised). MIT Press, 1997.

Pure Find the latest releases and the mailing list at the Pure website, <http://pure-lang.googlecode.com>.

Q Another term rewriting language by yours truly, <http://q-lang.sf.net>.

Pure Library Manual

Version 0.46, November 14, 2010

Albert Gräf <Dr.Graef@t-online.de>

Copyright (c) 2009-2010 by Albert Gräf. This document is available under the [GNU Free Documentation License](#).

This manual describes the operations in the standard Pure library, including the prelude and the other library modules which come bundled with the interpreter.

There is a companion to this manual, *The Pure Manual* which describes the Pure language and the operation of the Pure interpreter.

1 Prelude

The prelude defines the basic operations of the Pure language. This includes the basic arithmetic and logical operations, string, list and matrix functions, as well as the support operations required to implement list and matrix comprehensions. The string, matrix and record operations are in separate modules `strings.pure`, `matrices.pure` and `records.pure`, the primitive arithmetic and logical operations can be found in `primitives.pure`. Note that since the prelude module gets imported automatically (unless the interpreter is invoked with the `--no-prelude` option), all operations discussed in this section are normally available in Pure programs without requiring any explicit import declarations. The prelude also declares a signature of commonly used constant and operator symbols. This includes the truth values `true` and `false`. These are actually just integers in Pure, but sometimes it's convenient to refer to them using these symbolic constants. Note that if you also want to use these on the left-hand side of equations, you still have to declare them as `nonfix` symbols yourself, using a declaration like:

```
nonfix false true;
```

In addition, the following special exception symbols are provided:

- Built-in exception values: `failed_cond` (failed conditional in guard or if-then-else), `failed_match` (failed pattern match in lambda, case expression, etc.), `stack_fault` (not enough stack space, `PURE_STACK` limit exceeded) and `malloc_error` (memory allocation error).

- Value mismatches a.k.a. dynamic typing errors: `bad_string_value x`, `bad_matrix_value x`, `bad_list_value x`, `bad_tuple_value x`. These are thrown by some operations when they fail to find an expected value of the corresponding type.
- Index range exceptions: `out_of_bounds` is thrown by the index operator `!` if a list, tuple or matrix index is out of bounds.

Here's the list of predefined operator symbols. Note that the parser will automatically give unary minus the same precedence level as the corresponding binary operator.

```
infixl 1000  $$ ;           // sequence operator
infixr 1100  $ ;           // right-associative application
infixr 1200  , ;           // pair (tuple)
infix  1300  => ;          // key=>value pairs ("hash rocket")
infix  1400  .. ;          // arithmetic sequences
infixr 1500  || ;          // logical or (short-circuit)
infixr 1600  && ;          // logical and (short-circuit)
prefix 1700  ~ ;           // logical negation
infix  1800  < > <= >= == ~= ; // relations
infix  1800  === ~= ;      // syntactic equality
infixr 1900  : ;           // list cons
infix  2000  +: <: ;       // complex numbers (cf. math.pure)
infixl 2100  << >> ;       // bit shifts
infixl 2200  + - or ;      // addition, bitwise or
infixl 2300  * / div mod and ; // multiplication, bitwise and
infixl 2300  % ;           // exact division (cf. math.pure)
prefix 2400  not ;         // bitwise not
infixr 2500  ^ ;           // exponentiation
prefix 2600  # ;           // size operator
infixl 2700  ! !! ;        // indexing, slicing
infixr 2800  . ;           // function composition
prefix 2900  ' ;           // quote
postfix 3000  & ;          // thunk
```

1.1 Basic Combinators

The most important function combinators are `$` (right-associative application) and `.` (function composition), which are also defined as macros so that saturated calls of these are eliminated automatically. Examples:

```
> foo $ bar 99;
foo (bar 99)
> (foo.bar) 99;
foo (bar 99)
```

The customary identity and constant combinators from the combinatorial calculus are also available, in Pure these are named `id` and `cst`, respectively:

```
> map id (1..5);
[1,2,3,4,5]
```

```
> map (cst 0) (1..5);  
[0,0,0,0,0]
```

There's also a combinator `void` which is basically equivalent to `cst ()`, but with the special twist that it is also defined as a macro optimizing the case of “throwaway” list and matrix comprehensions. This is useful if a comprehension is evaluated solely for its side effects. E.g.:

```
> using system;  
> extern int rand();  
> foo = void [printf "%d\n" rand | _ = 1..3];  
> show foo  
foo = do (\_ -> printf "%d\n" rand) (1..3);  
> foo;  
1714636915  
1957747793  
424238335  
()
```

Note that the above list comprehension is actually implemented using `do` (instead of `map`, which would normally be the case), so that the intermediate list value of the comprehension is never constructed. This is described in more detail in section [Optimization Rules](#) of the Pure Manual.

In addition, Pure also provides the following combinators adopted from Haskell:

- `flip f` swaps arguments of a binary function `f`, e.g.:

```
> map (flip (/) 2) (1..3);  
[0.5,1.0,1.5]
```

This combinator is also used by the compiler to implement right operator sections, which allows you to write the above simply as:

```
> map (/2) (1..3);  
[0.5,1.0,1.5]
```

- `curry f` turns a function `f` expecting a pair of values into a curried function of two arguments:

```
> using system;  
> dowith (curry (printf "%d: %g\n")) (0..2) [0.0,2.718,3.14];  
0: 0  
1: 2.718  
2: 3.14  
()
```

- Conversely, `uncurry f` turns a curried function `f` expecting two arguments into a function processing a single pair argument:

```
> map (uncurry (*)) [(2,3),(4,5),(6,7)];  
[6,20,42]
```

- `curry3` and `uncurry3` are defined analogously, but are used to work with ternary functions.
- The (normal order) fixed point combinator `fix` allows you to create recursive anonymous functions. It takes another function `f` as its argument and applies `f` to `fix f` itself:

```
> let fact = fix (\f n -> if n<=0 then 1 else n*f (n-1));
> map fact (1..5);
[1,2,6,24,120]
```

See [Fixed point combinator](#) at Wikipedia for an explanation of how this magic works. Just like in Haskell, `fix` can be used to produce least fixed points of arbitrary functions. For instance:

```
> fix (cst bar);
bar
> let xs = fix (1:);
> xs;
1:#<thunk 0x7fe537fe2f90>
> xs!!(0..10);
[1,1,1,1,1,1,1,1,1,1,1]
```

1.2 Lists and Tuples

The prelude defines the list and tuple constructors `(x:y, x,y)`, as well as equality (`==`) and inequality (`~=`) on these structures. (The latter compare two lists or tuples by recursively comparing their members, so `==` must be defined on the list or tuple members if you want to use these operations.) It also provides the predicate `null x` which tests whether `x` is the empty list or tuple, the function `reverse x` which reverses a list or tuple, and the operators `#x` (size of a list or tuple), `x!i` (indexing), `x!!is` (slicing) and `x+y` (list concatenation). Note that there isn't a separate operation for concatenating tuples, since the pairing operator already does this:

```
> (1,2,3),(10,9,8);
1,2,3,10,9,8
```

This works because the `(,)` constructor is associative in Pure and will always produce right-recursive pairs. This also implies that tuples are always flat in Pure and can't be nested; if you need this, you should use lists instead. Also note that the empty tuple `()` acts as a neutral element with respect to `(,)`:

```
> (), (a,b,c);
a,b,c
> (a,b,c), ();
a,b,c
```

The prelude also provides another special kind of pairs called “hash pairs”, which take the form `key=>value`. These are used in various contexts to denote key-value associations. The

only operations on hash pairs provided by the prelude are equality testing (which recursively compares the components) and the functions `key` and `val` which extract the components of a hash pair:

```
> ("foo"=>99) == ("bar"=>99);
0
> key ("foo"=>99), val ("foo"=>99);
"foo", 99
```

Note that in difference to the tuple operator `(,)`, the “hash rocket” `(=>)` is non-associative, so nested applications *must* be parenthesized, and `(x=>y)=>z` is generally *not* the same as `x=>(y=>z)`. Also note that `(,)` has lower precedence than `(=>)`, so to include a tuple as key or value in a hash pair, the tuple must be parenthesized, as in `"foo"=>(1,2)` (whereas `"foo"=>1,2` denotes a tuple whose first element happens to be a hash pair). Lists are the usual right-recursive aggregates, pretty much the same as in Lisp or Prolog except that they use a Haskell-like syntax. In difference to Haskell, list concatenation is denoted `+`, and lists may contain an arbitrary mixture of arguments, i.e., they are fully polymorphic:

```
> 1:2:3:[];
[1,2,3]
> [1,2,3]+[u,v,w]+[3.14];
[1,2,3,u,v,w,3.14]
```

Lists are **eager** in Pure by default, but they can also be made **lazy**, see section [Lazy Evaluation and Streams](#) in the Pure Manual. Arithmetic sequences can be constructed with the infix `..` operator:

```
> 1..5;
[1,2,3,4,5]
> 1:3..11;
[1,3,5,7,9,11]
```

Note that the Pure syntax differs from Haskell in that there are no brackets around the construct and a step width is indicated by specifying the first two elements as `x:y` instead of `x,y`. Also, to specify infinite sequences you have to use an infinite upper bound (`inf` or `-inf`):

```
> 1:3..inf;
1: #<thunk 0x7f696cd2dbd8>
> -1:-3..-inf;
-1: #<thunk 0x7f696cd2fde8>
```

The lower bounds of an arithmetic sequence must always be finite. Non-lazy lists can be sorted using the `sort` function, which provides an interface to the C `qsort()` routine, please see your local `qsort(3)` documentation for details. An order predicate `p` is specified as the first argument, which must return a nonzero integer to indicate that its first argument is “less than” its second argument. Example:

```
> sort (>) (1..10);
[10,9,8,7,6,5,4,3,2,1]
> sort (<) ans;
[1,2,3,4,5,6,7,8,9,10]
```

You can convert between (finite) lists and tuples using the `list` and `tuple` operations:

```
> tuple (1..5);
1,2,3,4,5
> list (a,b,c);
[a,b,c]
```

The `list` function can also be used to turn a finite lazy list into an eager one:

```
> list $ take 10 (-1:-3..-inf);
[-1,-3,-5,-7,-9,-11,-13,-15,-17,-19]
```

You can also achieve the same effect somewhat more conveniently by slicing a finite part from a stream (see below):

```
> (-1:-3..-inf)!!(0..9);
[-1,-3,-5,-7,-9,-11,-13,-15,-17,-19]
```

Conversely, it is also possible to convert a list to a stream:

```
> stream (1..10);
1:#<thunk 0x7fe537fe2b58>
```

This might appear a bit useless at first sight, since all elements of the stream are in fact already known. However, this operation then allows you to apply other functions to the list and have them evaluated in a lazy fashion. Indexing of lists and tuples is always zero-based (i.e., indices run from 0 to `#x-1`), and an exception will be raised if the index is out of bounds:

```
> [1,2,3]!2;
3
> [1,2,3]!4;
<stdin>, line 34: unhandled exception 'out_of_bounds' while evaluating
'[1,2,3]!4'
```

The slicing operator `!!` takes a list or tuple and a list of indices and returns the list or tuple of the corresponding elements, respectively. Indices which are out of the valid range are silently ignored:

```
> (1..5)!!(3..10);
[4,5]
> (1,2,3,4,5)!!(3..10);
4,5
```

Indices can actually be specified in any order, so that you can retrieve any permutation of the members, also with duplicates. E.g.:

```
> (1..5)!![2,4,4,1];
[3,5,5,2]
```

This is less efficient than the case of contiguous index ranges (which is optimized so that it always works in linear time), because it requires repeated traversals of the list for each index. For larger lists you should hence use vectors or matrices instead, to avoid the quadratic complexity.

(The prelude actually implements the slicing operation in a fairly generic way, so that it works with any kind of container data structure which defines `!` in such a manner that it throws an exception when the index is out of bounds. It also works with any kind of index container that implements the *catmap* operation.)

1.3 List Functions

This mostly comes straight from the Q prelude which in turn was based on the first edition of the Bird/Wadler book, and is very similar to what you can find in the Haskell prelude. Some functions have slightly different names, though, and of course everything is typed dynamically.

Common List Functions

any *p xs* tests whether the predicate *p* holds for any of the members of *xs*

all *p xs* tests whether the predicate *p* holds for all of the members of *xs*

cat *xs* concatenate a list of lists

catmap *f xs* convenience function which combines *cat* and *map*; this is also used to implement list comprehensions

do *f xs* apply *f* to all members of *xs*, like *map*, but throw away all intermediate results and return `()`

drop *n xs* remove *n* elements from the front of *xs*

dropwhile *p xs* remove elements from the front of *xs* while the predicate *p* is satisfied

filter *p xs* return the list of all members of *xs* satisfying the predicate *p*

foldl *f a xs* accumulate the binary function *f* over all members of *xs*, starting from the initial value *a* and working from the front of the list towards its end

foldl1 *f xs* accumulate the binary function *f* over all members of *xs*, starting from the value `head xs` and working from the front of the list towards its end; *xs* must be nonempty

foldr *f a xs* accumulate the binary function *f* over all members of *xs*, starting from the initial value *a* and working from the end of the list towards its front

foldr1 *f xs* accumulate the binary function *f* over all members of *xs*, starting from the value `last xs` and working from the end of the list towards its front; *xs* must be nonempty

head *xs* return the first element of *xs*; *xs* must be nonempty

index *xs x* search for an occurrence of *x* in *xs* and return the index of the first occurrence, if any, -1 otherwise

Note: This uses equality (==) to decide whether a member of `xs` is an occurrence of `x`, so `==` must have an appropriate definition on the list members.

init xs return all but the last element of `xs`; `xs` must be nonempty

last xs return the last element of `xs`; `xs` must be nonempty

listmap f xs convenience function which works like `map`, but also deals with matrix and string arguments while ensuring that the result is always a list; this is primarily used to implement list comprehensions

map f xs apply `f` to each member of `xs`

scanl f a xs accumulate the binary function `f` over all members of `xs`, as with `foldl`, but return all intermediate results as a list

scanl1 f xs accumulate the binary function `f` over all members of `xs`, as with `foldl1`, but return all intermediate results as a list

scanr f a xs accumulate the binary function `f` over all members of `xs`, as with `foldr`, but return all intermediate results as a list

scanr1 f xs accumulate the binary function `f` over all members of `xs`, as with `foldr1`, but return all intermediate results as a list

sort p xs Sorts the elements of the list `xs` in ascending order according to the given predicate `p`, using the C `qsort` function. The predicate `p` is invoked with two arguments and should return a truth value indicating whether the first argument is “less than” the second. (An exception is raised if the result of a comparison is not a machine integer.)

tail xs return all but the first element of `xs`; `xs` must be nonempty

take n xs take `n` elements from the front of `xs`

takewhile p xs take elements from the front of `xs` while the predicate `p` is satisfied

List Generators

Some useful (infinite) list generators, as well as some finite (and eager) variations of these. The latter work like a combination of `take` or `takewhile` and the former, but are implemented directly for better efficiency.

cycle xs cycles through the elements of the nonempty list `xs`, ad infinitum

cyclen n xs eager version of `cycle`, returns the first `n` elements of `cycle xs`

iterate f x returns the stream containing `x`, `f x`, `f (f x)`, etc., ad infinitum

iteraten n f x eager version of `iterate`, returns the first `n` elements of `iterate f x`

iterwhile p f x another eager version of `iterate`, returns the list of all elements from the front of `iterate f x` for which the predicate `p` holds

repeat x returns an infinite stream of `xs`

repeatn *n* *x* eager version of `repeat`, returns a list with *n* *xs*

Zip and Friends

unzip *xy*s takes a list of pairs to a pair of lists of corresponding elements

unzip3 *xyz*s unzip with triples

zip *xs ys* return the list of corresponding pairs (*x*, *y*) where *x* runs through the elements of *xs* and *y* runs through the elements of *y*

zip3 *xs ys zs* zip with three lists, returns a list of triples

zipwith *f xs ys* apply the binary function *f* to corresponding elements of *xs* and *ys*

zipwith3 *f xs ys zs* apply the ternary function *f* to corresponding elements of *xs*, *ys* and *zs*

Pure also has the following variations of `zipwith`/`zipwith3` which throw away all intermediate results and return `()`. That is, these work like [do](#) but pull arguments from two or three lists, respectively:

dowith *f xs ys* apply the binary function *f* to corresponding elements of *xs* and *ys*, return `()`

dowith3 *f xs ys zs* apply the ternary function *f* to corresponding elements of *xs*, *ys* and *zs*, return `()`

1.4 String Functions

Pure strings are null-terminated character strings encoded in UTF-8, see the Pure Manual for details. The prelude provides various operations on strings, including a complete set of list-like operations, so that strings can be used mostly as if they were lists, although they are really implemented as C character arrays for reasons of efficiency. Pure also has some powerful operations to convert between Pure expressions and their string representation, see [Eval and Friends](#) for those.

Basic String Functions

Concatenation, indexing and slicing works just like with lists:

```
> "abc"+"xyz";
"abcxyz"
> let s = "The quick brown fox jumps over the lazy dog.";
> s!5;
"u"
> s!!(20..24);
"jumps"
```

Checking for empty strings and determining the size of a string also works as expected:

```
> null "";
1
> null s;
0
> #s;
44
```

You can search for the location of a substring in a string, and extract a substring of a given length:

index s u Returns the (zero-based) index of the first occurrence of the substring *u* in *s*, or -1 if *u* is not found in *s*.

substr s i n Extracts a substring of (at most) *n* characters at position *i* in *s*. This takes care of all corner cases, adjusting index and number of characters so that the index range stays confined to the source string.

Example:

```
> index s "jumps";
20
> substr s 20 10;
"jumps over"
```

Note that Pure doesn't have a separate type for individual characters. Instead, these are represented as strings *c* containing exactly one (UTF-8) character (i.e., `#c==1`). It is possible to convert such single character strings to the corresponding integer character codes, and vice versa:

ord c Ordinal number of a single character string *c*. This is the character's code point in the Unicode character set.

chr n Converts an integer back to the character with the corresponding code point.

In addition, the usual character arithmetic works, including arithmetic sequences of characters, so that you can write stuff like the following:

```
> "a" - "A";
32
> "u" - 32;
"U"
> "a" .. "k";
["a", "b", "c", "d", "e", "f", "g", "h", "i", "j", "k"]
```

Strings are also ordered lexicographically based on their character codes:

```
> "awe" > "awesome";
0
> "foo" >= "bar";
1
```

For convenience, the prelude provides the following functions to convert between strings and lists (or other aggregates) of characters.

chars s, list s Convert a string *s* to a list of characters.

tuple s, matrix s Convert a string *s* to a tuple or (symbolic) matrix of characters, respectively.

strcat xs Concatenate a list *xs* of strings (in particular, this converts a list of characters back to a string).

string xs Convert a list, tuple or (symbolic) matrix of strings to a string. In the case of a list, this is synonymous with `strcat`, but it also works with the other types of aggregates.

For instance:

```
> list "abc";  
["a", "b", "c"]  
> string ("a".."z");  
"abcdefghijklmnopqrstuvwxyz"
```

The following functions are provided to deal with strings of “tokens” separated by a given delimiter string.

split delim s Splits *s* into a list of substrings delimited by *delim*.

join delim xs Joins the list of strings *xs* to a single string, interpolating the given *delim* string.

Example:

```
> let xs = split " " s; xs;  
["The", "quick", "brown", "fox", "jumps", "over", "the", "lazy", "dog."]  
> join ":" xs;  
"The:quick:brown:fox:jumps:over:the:lazy:dog."
```

We mention in passing here that more elaborate string matching, splitting and replacement operations based on regular expressions are provided by the system module, see [System Interface](#).

If that isn’t enough already, most generic list operations carry over to strings in the obvious way, treating the string like a list of characters. (Some operations will actually return character lists in that case, so you might have to apply `string` explicitly to convert these back to a string.) For instance:

```
> filter (>= "k") s;  
"qukrownoxumpsovr t lzyo"  
> string $ map pred "ibm";  
"hal"
```

List comprehensions can draw values from strings, too:

```
> string [x+1 | x="HAL"];  
"IBM"
```

Low-Level Operations

The following routines are provided by the runtime to turn raw C `char*` pointers (also called **byte strings** in Pure parlance, to distinguish them from Pure's "cooked" UTF-8 string values) into corresponding Pure strings. Normally you don't have to worry about this, because the C interface already takes care of the necessary marshalling, but in some low-level code these operations are useful. Also note that here and in the following, the `cstring` routines also convert the string between the system encoding and Pure's internal UTF-8 representation.

string s, cstring s Convert a pointer `s` to a Pure string. `s` must point to a null-terminated C string. These routines take ownership of the original string value, assuming it to be malloced, so you should only use these for C strings which are specifically intended to be freed by the user.

string_dup s, cstring_dup s Convert a pointer `s` to a Pure string. Like above, but these functions take a copy of the string, leaving the original C string untouched.

The reverse transformations are also provided. These take a Pure string to a byte string (raw `char*`).

byte_string s, byte_cstring s Construct a byte string from a Pure string `s`. The result is a raw pointer object pointing to the converted string. The original Pure string is always copied (and, in the case of `byte_cstring`, converted to the system encoding). The resulting byte string is a malloced pointer which can be used like a C `char*`, and has to be freed explicitly by the caller when no longer needed.

It is also possible to convert Pure string lists or symbolic vectors of strings to byte string vectors and vice versa. These are useful if you need to pass an `argv`-like string vector (i.e., a `char**` or `char*[]`) to C routines. The computed C vectors are malloced pointers which have an extra NULL pointer as the last entry, and should thus be usable for almost any purpose which requires such a string vector in C. They also take care of garbage-collecting themselves. The original string data is always copied. As usual, the `cstring` variants do automatic conversions to the system encoding.

byte_string_pointer xs, byte_cstring_pointer xs Convert a list or vector of Pure strings to a C `char**`.

string_list n p, cstring_list n p Convert a C `char**` to a list of Pure strings.

string_vector n p, cstring_vector n p Convert a C `char**` to a symbolic vector of Pure strings.

Note that the back conversions take an additional first argument which denotes the number of strings to retrieve. If you know that the vector is NULL-terminated then this can also be an infinite value (`inf`) in which case the number of elements will be figured out automatically. Processing always stops at the first NULL pointer encountered.

Also note that, as of version 0.45, Pure has built-in support for passing `argv`-style vectors as arguments by means of the `char**` and `void**` pointer types. However, the operations provided here are more general in that they allow you to both encode and decode such

values in an explicit fashion. This is useful, e.g., for operations like `getopt(3)` which may mutate the given `char**` vector.

If you have `getopt` in your C library, you can try the following example. First enter these definitions:

```
extern int getopt(int argc, char **argv, char *optstring);
optind = get_int $ addr "optind";
optarg = cstring_dup $ get_pointer $ addr "optarg";
```

Now let's run `getopt` on a byte string vector constructed from an argument vector (which includes the "program name" in the first element):

```
> let args = byte_cstring_pointer {"programe", "boo", "-n", "-tfoo", "bar"};
> getopt 5 args "nt:", optarg;
110, #<pointer 0>
> getopt 5 args "nt:", optarg;
116, "foo"
> getopt 5 args "nt:", optarg;
-1, #<pointer 0>
```

Note that 110 and 116 are the character codes of the option characters `n` and `t`, where the latter option takes an argument, as returned by `optarg`. Finally, `getopt` returns -1 to indicate that there are no more options, and we can retrieve the current `optindex` value and the mutated argument vector to see which non-option arguments remain to be processed, as follows:

```
> optind, cstring_vector 5 args;
3, {"programe", "-n", "-tfoo", "boo", "bar"}
```

It is now an easy exercise to design your own high-level wrapper around `getopt` to process command line arguments in Pure. However, this isn't really necessary since the Pure library already offers such an operation which doesn't rely on any special system functions, see [Option Parsing](#) in the [System Interface](#) section.

1.5 Matrix Functions

The size of a matrix (number of elements) can be obtained using `#`, and the `dim` function can be used to return its dimensions (number of rows and columns):

```
> let x = {1,2,3;4,5,6}; #x;
6
> dim x;
2,3
```

`null x` can be used to check for empty matrices. Note that there are various kinds of these, as a matrix may have zero rows or columns, or both. Indexing and slicing works pretty much like in MATLAB and Octave, except that the Pure operators `!` and `!!` are used (and indices are zero-based). It is possible to access elements with a one-dimensional index (in row-major order):

```
> x!3;  
4
```

Or you can specify a pair of row and column index:

```
> x!(1,0);  
4
```

Slicing works accordingly. You can either specify a list of (one- or two-dimensional) indices, in which case the result is always a row vector:

```
> x!!(2..5);  
{3,4,5,6}
```

Or you can specify a pair of row and column index lists:

```
> x!!(0..1,1..2);  
{2,3;5,6}
```

The following abbreviations are provided to grab a slice from a row or column:

```
> x!!(1,1..2);  
{5,6}  
> x!!(0..1,1);  
{2;5}
```

Matrices can be compared with the `==` and `~=` operators which check the dimensions and the matrix elements for equality:

```
> x == transpose x;  
0
```

Most of the generic list operations are implemented on matrices, see [Common List Functions](#). Hence operations like `map` and `zipwith` work as expected:

```
> map succ {1,2,3;4,5,6};  
{2,3,4;5,6,7}  
> zipwith (+) {1,2,3;4,5,6} {1,0,1;0,2,0};  
{2,2,4;4,7,6}
```

The matrix module also provides a bunch of other specialized matrix operations, including all the necessary operations for matrix comprehensions. We briefly summarize the most important operations below; please refer to `matrices.pure` for all the gory details. Also make sure you check [Matrix Computations](#) in the Pure Manual for some more examples, and the [Record Functions](#) section for an implementation of records using symbolic vectors.

Matrix Construction and Conversions

The `matrix` function converts a list or tuple to a corresponding matrix. `matrix` also turns a list of lists or matrices specifying the rows of the matrix to the corresponding rectangular matrix; otherwise, the result is a row vector. (In the former case, the `matrix` function may

throw a `bad_matrix_value` x exception in case of dimension mismatch, where x denotes the offending submatrix.)

```
> matrix [1,2,3];
{1,2,3}
> matrix [[1,2,3],[4,5,6]];
{1,2,3;4,5,6}
```

The `rowvector` and `colvector` functions work in a similar fashion, but expect a list, tuple or matrix of elements and always return a row or column vector, respectively (i.e., a $1 \times n$ or $n \times 1$ matrix, where n is the size of the converted aggregate). Also, the `vector` function is a synonym for `rowvector`. These functions can also be used to create recursive (symbolic) matrix structures of arbitrary depth, which provide a nested array data structure with efficient (constant time) element access.

```
> rowvector [1,2,3];
{1,2,3}
> colvector [1,2,3];
{1;2;3}
> vector [rowvector [1,2,3],colvector [4,5,6]];
{{1,2,3},{4;5;6}}
```

It's also possible to create a row or column vector from an arithmetic sequence using the `rowvectorseq` and `colvectorseq` functions. Again, `vectorseq` is provided as a synonym for `rowvectorseq`. These operations are optimized for the case of int and double ranges.

```
> rowvectorseq 0 10 1;
{0,1,2,3,4,5,6,7,8,9,10}
> colvectorseq 0 10 1;
{0;1;2;3;4;5;6;7;8;9;10}
> vectorseq 0.0 0.9 0.1;
{0.0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9}
```

These operations are also handy to define macros which optimize the construction of vectors from list ranges, so that the intermediate list values are actually never constructed. For instance:

```
> def vector (n1:n2..m) = vectorseq n1 m (n2-n1);
> def vector (n..m) = vectorseq n m 1;
> foo = vector (1..10);
> bar = vector (0.0:0.1..0.9);
> show foo bar
bar = vectorseq 0.0 0.9 0.1;
foo = vectorseq 1 10 1;
> foo; bar;
{1,2,3,4,5,6,7,8,9,10}
{0.0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9}
```

(Note that the library doesn't contain any such macro rules, so you'll have to add them yourself if you want to take advantage of these optimizations.) The `dmatrix`, `cmatrix`, `imatrix` and `smatrix` functions convert a list or matrix to a matrix of the corresponding type (integer, double, complex or symbolic). If the input is a list, the result is always a row

vector; this is usually faster than the `matrix` and `rowvector` operations, but requires that the elements already are of the appropriate type.

```
> imatrix [1,2,3];
{1,2,3}
> dmatrix {1,2,3;4,5,6};
{1.0,2.0,3.0;4.0,5.0,6.0}
```

The `dmatrix`, `cmatrix` and `imatrix` functions can also be invoked with either an int `n` or a pair `(n,m)` of ints as argument, in which case they construct a zero rowvector or matrix with the corresponding dimensions.

```
> imatrix 3;
{0,0,0}
> imatrix (2,3);
{0,0,0;0,0,0}
```

There's also a number of more specialized conversions operations which are discussed under [Matrix Inspection and Manipulation](#) below. Conversely, a matrix can be converted back to a flat list or tuple with the `list` and `tuple` functions. In addition, the `list2` function converts a matrix to a list of lists (one sublist for each row of the matrix).

```
> tuple {1,2,3;4,5,6};
1,2,3,4,5,6
> list {1,2,3;4,5,6};
[1,2,3,4,5,6]
> list2 {1,2,3;4,5,6};
[[1,2,3],[4,5,6]]
> list2 {1,2,3};
[[1,2,3]]
```

Matrix Inspection and Manipulation

dmatrixp x, cmatrixp x, imatrixp x, smatrixp x, nmatrixp x Check for different kinds of matrices (double, complex, int, symbolic and numeric, i.e., non-symbolic).

vectorp x, rowvectorp x, colvectorp x Check for different kinds of vectors (these are just matrices with one row or column).

stride x The stride of a matrix denotes the real row size of the underlying C array, see the description of the [pack](#) function below for further details. There's little use for this value in Pure, but it may be needed when interfacing to C.

row x i, col x i Extract the `i`th row or column of a matrix.

rows x, cols x Return the list of all rows or columns of a matrix.

diag x, subdiag x k, supdiag x k: Extract (sub-,super-) diagonals from a matrix. Sub- and super-diagonals for `k=0` return the main diagonal. Indices for sub- and super-diagonals can also be negative, in which case the corresponding super- or sub-diagonal is returned instead. In each case the result is a row vector.

submat *x* (*i,j*) (*n,m*) Extract a submatrix of a given size at a given offset. The result shares the underlying storage with the input matrix (i.e., matrix elements are *not* copied) and so this is a comparatively cheap operation.

rowcat *xs*, **colcat** *xs* Construct matrices from lists of rows and columns. These take either scalars or submatrices as inputs; corresponding dimensions must match. **rowcat** combines submatrices vertically, like $\{x;y\}$; **colcat** combines them horizontally, like $\{x,y\}$. Note: Like the built-in matrix constructs, these operations may throw a `bad_matrix_value` exception in case of dimension mismatch, where *x* denotes the offending submatrix.

matcat *xs* Construct a matrix from a (symbolic) matrix of other matrices and/or scalars. This works like a combination of **rowcat** and **colcat**, but draws its input from a matrix instead of a list of matrices, and preserves the overall layout of the “host” matrix. The net effect is that the host matrix is flattened out. If all elements of the input matrix are scalars already, the input matrix is returned unchanged.

rowcatmap *f xs*, **colcatmap** *f xs* Combinations of **rowcat**, **colcat** and **map**. These are used, in particular, for implementing matrix comprehensions.

diagmat *x*, **subdiagmat** *x k*, **supdiagmat** *x k* Create a (sub-,super-) diagonal matrix from a row vector *x* of size *n*. The result is always a square matrix with dimension $(n+k, n+k)$, which is of the same matrix type (double, complex, int, symbolic) as the input and has the elements of the vector on its *k*th sub- or super-diagonal, with all other elements zero. A negative value for *k* turns a sub- into a super-diagonal matrix and vice versa.

re *x*, **im** *x*, **conj** *x* Extract the real and imaginary parts and compute the conjugate of a numeric matrix.

pack *x*, **packed** *x* Pack a matrix. This creates a copy of the matrix which has the data in contiguous storage. It also frees up extra memory if the matrix was created as a slice from a bigger matrix (see **submat** above) which has since gone the way of the dodo. The packed predicate can be used to verify whether a matrix is already packed. Note that even if a matrix is already packed, **pack** will make a copy of it anyway, so **pack** also provides a quick way to copy a matrix, e.g., if you want to pass it as an input/output parameter to a GSL routine.

redim (*n,m*) *x*, **redim** *n* *x* Change the dimensions of a matrix without changing its size. The total number of elements must match that of the input matrix. Reuses the underlying storage of the input matrix if possible (i.e., if the matrix is *packed*). You can also **redim** a matrix to a given row size *n*. In this case the row size must divide the total size of the matrix.

sort *p x* Sorts the elements of a matrix (non-destructively, i.e., without changing the original matrix) according to the given predicate, using the C `qsort` function. This works exactly the same as with lists (see [Common List Functions](#)), except that it takes and returns a matrix instead of a list. Note that the function sorts *all* elements of the matrix in one go (regardless of the dimensions), as if the matrix was a single big vector. The result matrix has the same dimensions as the input matrix. Example:

```
> sort (<) {10,9;8,7;6,5};
{5,6;7,8;9,10}
```

If you'd like to sort the individual rows instead, you can do that as follows:

```
> sort_rows p = rowcat . map (sort p) . rows;
> sort_rows (<) {10,9;8,7;6,5};
{9,10;7,8;5,6}
```

Likewise, to sort the columns of a matrix:

```
> sort_cols p = colcat . map (sort p) . cols;
> sort_cols (<) {10,9;8,7;6,5};
{6,5;8,7;10,9}
```

Also note that the `pure-gsl` module provides an interface to the GSL routines for sorting numeric (int and double) vectors using the standard order. These will usually be much faster than `sort`, whereas `sort` is more flexible in that it also allows you to sort symbolic matrices and to choose the order predicate.

transpose x Transpose a matrix. Example:

```
> transpose {1,2,3;4,5,6};
{1,4;2,5;3,6}
```

rowrev x, colrev x, reverse x Reverse a matrix. `rowrev` reverses the rows, `colrev` the columns, `reverse` both dimensions.

Pointers and Matrices

Last but not least, the `matrix` module also offers a bunch of low-level operations for converting between matrices and raw pointers. These are typically used to shovel around massive amounts of numeric data between Pure and external C routines, when performance and throughput is an important consideration (e.g., graphics, video and audio applications). The usual caveats concerning direct pointer manipulations apply.

pointer x Get a pointer to the underlying C array of a matrix. The data is *not* copied. Hence you have to be careful when passing such a pointer to C functions if the underlying data is non-contiguous; when in doubt, first use the [pack](#) function to place the data in contiguous storage, or use one of the matrix-pointer conversion routines below.

The following operations copy the contents of a matrix to a given pointer and return that pointer, converting to the target data type on the fly if necessary. The given pointer may also be `NULL`, in which case suitable memory is malloced and returned; otherwise the caller must ensure that the memory pointed to by `p` is big enough for the contents of the given matrix.

- `double_pointer p x`
- `float_pointer p x`
- `complex_pointer p x`

- `complex_float_pointer p x`
- `int_pointer p x`
- `short_pointer p x`
- `byte_pointer p x`

Conversely, the following functions allow you to create a numeric matrix from a pointer, copying the data and converting it from the source type on the fly if necessary. The source pointer `p` may also be `NULL`, in which case the new matrix is filled with zeros instead. Otherwise the caller must ensure that the pointer points to properly initialized memory big enough for the requested dimensions. The given dimension may also be just an integer `n` if a row vector is to be created.

- `double_matrix (n,m) p`
- `float_matrix (n,m) p`
- `complex_matrix (n,m) p`
- `complex_float_matrix (n,m) p`
- `int_matrix (n,m) p`
- `short_matrix (n,m) p`
- `byte_matrix (n,m) p`

Finally, you can use the following operations to create a numeric matrix view of existing data, without copying the data. The data must be double, complex or int, the pointer must not be `NULL` and the caller must also ensure that the memory persists for the entire lifetime of the matrix object. The given dimension may also be just an integer `n` if a row vector view is to be created.

- `double_matrix_view (n,m) p`
- `complex_matrix_view (n,m) p`
- `int_matrix_view (n,m) p`

1.6 Record Functions

As of Pure 0.41, the prelude also provides a basic record data structure, implemented as symbolic vectors of `key=>value` pairs which support a few dictionary-like operations such as `member`, `insert` and `indexing`. Records may be represented as row, column or empty vectors (i.e., the number of rows or columns must be zero or one). They must be symbolic matrices consisting only of “hash pairs” `key=>value`, where the keys can be either symbols or strings. The values can be any kind of Pure data; in particular, they may themselves be records, so records can be nested.

The following operations are provided. Please note that all updates of record members are non-destructive and thus involve copying, which takes linear time (and space) and thus might be slow for large record values; if this is a problem then you should use dictionaries

instead (cf. [Dictionaries](#)). Or you can create mutable records by using expression references (cf. [Expression References](#)) as values, which allow you to modify the data in-place. Element lookup (indexing) uses binary search on an internal index data structure and thus takes logarithmic time once the index has been constructed (which is done automatically when needed, or when calling `recordp` on a fresh record value).

Also note that records with duplicate keys are permitted; in such a case the following operations will always operate on the *last* entry for a given key.

recordp *x* Check for record values.

record *x* Normalizes a record. This removes duplicate keys and orders the record by keys (using an apparently random but well-defined order of the key values), so that normalized records are syntactically equal (`===`) if and only if they contain the same hash pairs. For convenience, this function can also be used directly on lists and tuples of hash pairs to convert them to a normalized record value.

#x The size of a record (number of entries it contains). Duplicate entries are counted. (This is in fact just the standard matrix size operation.)

member *x y* Check whether *x* contains the key *y*.

x!y Retrieves the (last) value associated with the key *y* in *x*, if any, otherwise throws an `out_of_bound` exception.

x!!ys Slicing also works as expected, by virtue of the generic definition of slicing provided by the matrix data structure.

insert *x (y=>z)*, **update** *x y z* Associate the key *y* with the value *z* in *x*. If *x* already contains the key *y* then the corresponding value is updated (the last such value if *x* contains more than one association for *y*), otherwise a new member is inserted at the end of the record.

delete *x y* Delete the key *y* (and its associated value) from *x*. If *x* contains more than one entry for *y* then the last such entry is removed.

keys *x*, **vals** *x* List the keys and associated values of *x*. If the record contains duplicate keys, they are all listed in the order in which they are stored in the record.

Here are a few basic examples:

```
> let r = {x=>5, y=>12};
> r!y; r!![y,x];           // indexing and slicing
12
{12,5}
> keys r; vals r;          // keys and values of a record
{x,y}
{5,12}
> insert r (x=>99);         // update an existing entry
{x=>99,y=>12}
> insert ans (z=>77);       // add a new entry
{x=>99,y=>12,z=>77}
> delete ans z;            // delete an existing entry
{x=>99,y=>12}
```

```

> let r = {r,x=>7,z=>3}; r;    // duplicate key x
{x=>5,y=>12,x=>7,z=>3}
> r!x, r!z;                  // indexing returns the last value of x
7,3
> delete r x;                // delete removes the last entry for x
{x=>5,y=>12,z=>3}
> record r;                  // normalize (remove dups and sort)
{x=>7,y=>12,z=>3}
> record [x=>5, x=>7, y=>12]; // construct a normalized record from a list
{x=>7,y=>12}
> record (x=>5, x=>7, y=>12); // ... or a tuple
{x=>7,y=>12}

```

More examples can be found in the [Record Data](#) section in the Pure Manual.

1.7 Primitives

This prelude module is a collection of various lowlevel operations, which are implemented either directly by machine instructions or by C functions provided in the runtime. In particular, this module defines the basic arithmetic and logic operations on machine integers, bigints and floating point numbers, as well as various type checking predicates and conversions between different types. Some low-level pointer operations are also provided, as well as “sentries” (Pure’s flavour of object finalizers) and “references” (mutable expression pointers).

Arithmetic

The basic arithmetic and logic operations provided by this module are summarized in the following table:

Kind	Operator	Meaning
Arithmetic	+ -	addition, subtraction (also unary minus)
	* /	multiplication, division (inexact)
	div mod	exact int/bigint division/modulus
	^	exponentiation (inexact)
Comparisons	== ~=	equality, inequality
	< >	less than, greater than
	<= >=	less than or equal, greater than or equal
Logic	~	logical not
	&&	and, or (short-circuit)
Bitwise	not	bitwise not
	and or	and, or
	<< >>	bit shifts

Precedence and associativity of the operators can be found in the *operators* table at the beginning of this section.

The names of some operations are at odds with C. Note, in particular, that logical negation is denoted `~` instead of `!` (and, consequently, `~=` denotes inequality, rather than `!=`), and the bitwise operations are named differently. This is necessary because Pure uses `!`, `&` and `|` for other purposes. Also, `/` always denotes inexact (double) division in Pure, whereas the integer division operators are called `div` and `mod`. (`%`, which is not defined by this module, also has a different meaning in Pure; it's the exact division operator, see [Rational Numbers](#).)

The above operations are implemented for `int`, `bigint` and, where appropriate, `double` and `pointer` operands. (Pointer arithmetic comprises `+` and `-` and works in the usual way, i.e., `p-q` returns the byte offset between two pointers `p` and `q`, and `p+n` or `p-n` offsets a pointer `p` by the given integer `n` denoting the amount of bytes.) The `math` module (see [Mathematical Functions](#)) also provides implementations of the arithmetic and comparison operators for `rational`, `complex` and `complex rational` numbers.

Note that the logical operations are actually implemented as special forms in order to provide for short-circuit evaluation. This needs special support from the compiler to work. The `primitives` module still provides definitions for these, as well as other special forms like `quote` and the thunking operator `&` so that they may be used as function values and in partial applications, but when used in this manner they lose all their special call-by-name properties; see [Special Forms](#) in the Pure Manual for details.

The constants `inf` and `nan` are defined as the usual IEEE floating point infinities and NaNs, and the predicates `infp` and `nanp` are provided to check for these kinds of values.

In addition, the following arithmetic and numeric functions are provided:

abs `x`, **sgn** `x` Absolute value and sign of a number.

min `x y`, **max** `x y` Minimum and maximum of two values. This works with any kind of values which have the ordering relations defined on them.

succ `x`, **pred** `x` Successor (+1) and predecessor (-1) functions.

gcd `x y`, **lcm** `x y` The greatest common divisor and least common multiple functions from the GMP library. These return a `bigint` if at least one of the arguments is a `bigint`, a machine `int` otherwise.

pow `x y` Computes exact powers of ints and bigints. The result is always a `bigint`. Note that `y` must always be nonnegative here, but see the `math` module ([Mathematical Functions](#)) which deals with the case `y<0` using rational numbers.

Conversions

These operations convert between various types of Pure values.

hash `x` Compute a 32 bit hash code of a Pure expression.

int `x`, **bigint** `x`, **double** `x`, **pointer** `x` Conversions between the different numeric and pointer types.

ubyte `x`, **ushort** `x`, **uint** `x`, **uint64** `x`, **ulong** `x` Convert signed (8/16/32/64) bit integers to the corresponding unsigned quantities. These functions behave as if the value was

“cast” to the corresponding unsigned C type, and are most useful for dealing with unsigned integers returned by external C routines. The routines always use the smallest Pure int type capable of holding the result: int for ubyte and ushort, bigint for uint, uint64 and ulong. All routines take int parameters. In the case of uint64, a bigint parameter is also permitted (which is what the C interface returns for 64 bit values). Also note that ulong reduces to either uint or uint64, depending on the size of long for the host architecture.

The following rounding functions work with all kinds of numbers:

floor x, ceil x Floor and ceil.

round x, trunc x Round or truncate to an integer.

frac x Fractional part (x - trunc x).

Note that all these functions return double values for double arguments, so if you need an integer result then you’ll have to apply a suitable conversion, as in int (floor x).

Predicates

A syntactic equality test is provided, as well as various type checking predicates.

x===y, x~==y, same x y Syntactic equality. In difference to == and ~= this is defined on all Pure expressions. Basically, two expressions are syntactically equal if they print out the same in the interpreter. In the special case of pointer objects and closures, which do not have a syntactic representation in Pure, x and y must be the same object (same pointer value or function).

intp x, bigintp x, doublep x, stringp x, pointerp x, matrixp x Predicates to check for the built-in types.

charp x Single character string predicate.

numberp x, complexp x, realp x, rationalp x, integerp x Additional number predicates.

exactp x, inexactp x Check whether a number is exact (i.e., doesn’t contain any double components).

applp x, listp x, listnp x, tuplep x Predicates to check for function applications, proper lists, list nodes and proper tuples.

funp x, lambdap x, thunkp x, varp x Predicates to check for various kinds of function objects (named, anonymous or thunk), and free variable symbols.

atomp x, closurep x Convenience functions to check for atoms (named function or variable) and closures (named function or lambda).

Inspection

The following operations let you peek at various internal information that the interpreter provides to Pure programs either for convenience or for metaprogramming purposes. They

are complemented by the evaluation primitives discussed below, see [Eval and Friends](#).

ans Retrieve the most recently printed result of a toplevel expression evaluated in the read-eval-print loop. This is just a convenience for interactive usage. Note that the `ans` value will stick around until a new expression is computed. (It is possible to clear the `ans` value with the interactive command `clear ans`, however.) Example:

```
> 1/3;
0.3333333333333333
> ans/2;
0.1666666666666667
```

__locals__ Return a list with the local function bindings (with clauses) visible at this point in the program. This is actually implemented as a built-in. The return value is a list of “hash” pairs `x=>f` where `x` is the (quoted) symbol denoting the function and `f` is the function itself (or its value, if `f` is a parameterless function). The `__locals__` function is useful for debugging purposes, as well as to implement dynamic environments. It is also used internally to implement the [reduce](#) macro, see [Eval and Friends](#). Example:

```
> __locals__ with foo x = x+1; x = a+b end;
[x=>a+b,foo=>foo]
> f 99 when _=>f = ans!1 end;
100
```

The prelude also provides some functions to retrieve various attributes of a function symbol which determine how the operation is applied to its operands or arguments. These functions all take a single argument, the symbol or function object to be inspected, and return an integer value.

nargs x Get the argument count of a function object, i.e., the number of arguments it expects. Returns 0 for thunks and saturated applications, -1 for over-saturated applications and non-functions.

arity x Determine the arity of an operator symbol. The returned value is 0, 1 or 2 for nullary, unary and binary symbols, respectively, -1 for symbols without a fixity declaration or other kinds of objects.

fixity f Determine the fixity of an operator symbol. The fixity is encoded as an integer $10 \cdot n + m$ where n is the precedence level (ranging from 0 to `PREC_MAX`, where `PREC_MAX` denotes the precedence of primary expressions, 16777216 in the current implementation) and m indicates the actual fixity at each level, in the order of increasing precedence (0 = infix, 1 = infixl, 2 = infixr, 3 = prefix, 4 = postfix). The fixity value of nonfix and outfix symbols, as well as symbols without a fixity declaration, is always given as $10 \cdot \text{PREC_MAX}$, and the same value is also reported for non-symbol objects. Infix, prefix and postfix symbols always have a fixity value less than $10 \cdot \text{PREC_MAX}$. (`PREC_MAX` isn’t actually defined as a constant anywhere, but you can easily do that yourself by setting `PREC_MAX` to the fixity value of any nonfix symbol or non-symbol value, e.g.: `const PREC_MAX = fixity [];`)

Note that only closures (i.e., named and anonymous functions and thunks) have a defined argument count in Pure, otherwise `nargs` returns -1 indicating an unknown argument count.

Partial applications of closures return the number of remaining arguments, which may be zero to indicate a **saturated** (but unevaluated) application, or -1 for **over-saturated** and constructor applications. (Note that in Pure a saturated application may also remain unevaluated because there is no definition for the given combination of arguments and thus the expression is in normal form, or because the application was quoted. If such a normal form application is then applied to some “extra” arguments it becomes over-saturated.)

The value returned by `nargs` always denotes the actual argument count of the given function, regardless of the declared arity if the function also happens to be an operator symbol. Often these will coincide (as, e.g., in the case of `+` which is a binary operator and also expects two arguments). But this is not necessarily the case, as shown in the following example of a binary operator which actually takes *three* arguments:

```
> infix 0 oops;  
> (oops) x y z = x*z+y;  
> arity (oops);  
2  
> nargs (oops);  
3  
> nargs (5 oops 8);  
1  
> map (5 oops 8) (1..5);  
[13,18,23,28,33]
```

Eval and Friends

Pure provides some rather powerful operations to convert between Pure expressions and their string representation, and to evaluate quoted expressions (`'x`). The string conversions `str`, `val` and `eval` also provide a convenient means to serialize Pure expressions, e.g., when terms are to be transferred to/from persistent storage. (Note, however, that this has its limitations. Specifically, some objects like pointers and anonymous functions do not have a parsable string representation. Also see the [Expression Serialization](#) section for some dedicated serialization operations which provide a more compact binary serialization format.)

str x Yields the print representation of an expression in Pure syntax, as a string.

val s Parses a single simple expression, specified as a string in Pure syntax, and returns the result as is, without evaluating it. Note that this is much more limited than the `eval` operation below, as the expression must not contain any of the special constructs (conditional expressions, `when`, `with`, etc.).

eval x Parses any expression, specified as a string in Pure syntax, and returns its value. In fact, `eval` can also parse and execute arbitrary Pure code. In that case it will return the last computed expression, if any. Alternatively, `eval` can also be invoked on a (quoted) Pure expression, which is recompiled and then evaluated. Exceptions during evaluation are reported back to the caller.

evalcmd x Like `eval`, but allows execution of interactive commands and returns their captured output as a string. No other results are returned, so this operation is most useful

for executing Pure definitions and interactive commands for their side-effects. (At this time, only the regular output of a few commands can be captured, most notably `bt`, `clear`, `mem`, `save` and `show`; otherwise the result string will be empty.)

lasterr Reports errors in `eval` and `evalcmd`. This string value will be nonempty iff a compilation or execution error was encountered during the most recent invocation of `eval` and `evalcmd`. In that case each reported error message is terminated with a newline character.

reduce x Reevaluates an expression in a local environment. This dynamically rebinds function symbols in the given expression to whatever local definitions are in effect at the point of the `reduce` call. Note that `reduce` is actually implemented as a macro which expands to the `reduce_with` primitive (see below), using the `__locals__` builtin to enumerate the bindings which are in effect at the call site.

reduce_with env x Like `reduce` above, but takes a list of replacements (given as hash pairs `u=>v`) as the first argument. The `reduce` macro expands to `reduce_with __locals__`.

Examples:

```
> str (1/3);
"0.3333333333333333"
> val "1/3";
1/3
> eval "1/3";
0.3333333333333333
> eval ('(1/3));
0.3333333333333333
> evalcmd "show evalcmd";
"extern expr* evalcmd(expr*);\n"

> eval "1/3)";
eval "1/3)"
> lasterr;
"<stdin>, line 1: syntax error, unexpected ')', expecting '=' or '|'\n"
```

The `reduce` macro provides a restricted form of dynamic binding which is useful to implement local rewriting rules. It is invoked without parameters and expands to the curried call `reduce_with __locals__` of the `reduce_with` primitive, which takes one additional argument, the expression to be rewritten. The following example shows how to expand or factorize an expression using local rules for the corresponding laws of distributivity:

```
expand = reduce with
  (a+b)*c = a*c+b*c;
  a*(b+c) = a*b+a*c;
end;

factor = reduce with
  a*c+b*c = (a+b)*c;
  a*b+a*c = a*(b+c);
end;
```

```
expand ((a+b)*2); // yields a*2+b*2
factor (a*2+b*2); // yields (a+b)*2
```

Note that instances of locally bound functions are substituted back in the computed result, thus the instances of `*` and `+` in the results `a*2+b*2` and `(a+b)*2` shown above denote the corresponding globals, not the local incarnations of `*` and `+` defined in `expand` and `factor`, respectively.

`reduce` also adjusts to quoted arguments. In this case, the local rules are applied as usual, but back-substituted globals are *not* evaluated in the result:

```
> expand ((a+1)*2);
a*2+2
> expand ('((a+1)*2));
a*2+1*2
```

Note that `reduce` only takes into account local *function* bindings from `with` clauses, local *variable* bindings do not affect its operation in any way:

```
> let y = [x, x^2, x^3];
> reduce y when x = u+v end;
[x, x^2, x^3]
```

However, in such cases you can perform the desired substitution by turning the `when` into a `with` clause:

```
> reduce y with x = u+v end;
[u+v, (u+v)^2, (u+v)^3]
```

Or you can just invoke the underlying `reduce_with` builtin directly, with the desired substitutions given as hash pairs in the first argument:

```
> reduce_with [x=>u+v] y;
[u+v, (u+v)^2, (u+v)^3]
```

Expression Serialization

Like `str` and `eval`, the following `blob` and `val` operations can be used to safely transfer expression data to/from persistent storage and between different processes (using, e.g., POSIX shared memory, pipes or sockets). However, `blob` and `val` use a binary format which is usually much more compact and gets processed much faster than the string representations used by `str` and `eval`. Also, `val` offers some additional protection against transmission errors through a crc check. (The advantage of the string representation, however, is that it's readable plain text in Pure syntax.)

blob *x* Stores the contents of the given expression as a binary object. The return value is a cooked pointer which frees itself when garbage-collected.

val *p* Reconstructs a serialized expression from the result of a previous invocation of the `blob` function.

blobp p Checks for a valid blob object. (Note that `val` may fail even if `blobp` returns `true`, because for performance reasons `blobp` only does a quick plausibility check on the header information of the blob, whereas `val` also performs a crc check and verifies data integrity.)

blob_size p, blob_crc p Determines the size (in bytes) and crc checksum of a blob, respectively. `blob_size` always returns a bigint, `blob_crc` a machine int (use `uint` on the latter to get a proper unsigned 32 bit value). For convenience, `#p` is defined as an alias for `blob_size p` on blob pointers.

Example:

```
> let b = blob {"Hello, world!", 1/3, 4711, NULL};
> b; #b; uint $ blob_crc b;
#<pointer 0x141dca0>
148L
3249898239L
> val b;
{"Hello, world!",0.3333333333333333,4711,#<pointer 0>}
```

Please note that the current implementation has some limitations:

- Just as with `str` and `eval`, runtime data (local closures and pointers other than the `NULL` pointer) can't be serialized, causing `blob` to fail. However, it *is* possible to transfer a global function, provided that the function exists (and is the same) in both the sending and the receiving process. (This condition can't be verified by `val` and thus is at the programmer's responsibility.)
- Sharing of subexpressions will in general be preserved, but sharing of list and tuple *tails* will be lost (unless the entire list or tuple is shared).
- The `val` function may fail to reconstruct the serialized expression even for valid blobs, if there is a conflict in symbol fixities between the symbol tables of the sending and the receiving process. To avoid this, make sure that symbol declarations in the sending and the receiving script match up.

Other Special Primitives

exit status Terminate the program with the given status code.

throw x Throw an exception, cf. [Exception Handling](#).

force x Force a thunk (`x&`), cf. [Special Forms](#). This usually happens automagically when the value of a thunk is needed.

Pointer Operations

These are lowlevel operations dealing with pointer values. The usual caveats apply, so *only* use these directly if you know what you're doing!

addr symbol Get the address of a C symbol (given as a string) at runtime. The library containing the symbol must already be loaded. Note that this can in fact be any kind of externally visible C symbol, so it's also possible to get the addresses of global variables. The result is returned as a pointer. The function fails if the symbol was not found.

calloc nmembers size, malloc size, realloc ptr size, free ptr Interface to malloc, free and friends. These let you allocate dynamic buffers (represented as Pure pointer values) for various nasty purposes.

The following operations perform direct memory accesses. Use with care ... or else!

- `get_byte ptr, get_short ptr, get_int ptr, get_int64 ptr, get_long ptr, get_float ptr, get_double ptr, get_string ptr, get_pointer ptr`
- `put_byte ptr x, put_short ptr x, put_int ptr x, put_int64 ptr x, put_long ptr x, put_float ptr x, put_double ptr x, put_string ptr x, put_pointer ptr x`

Sentries

Sentries are Pure's flavour of object **finalizers**. A sentry is simply an object (usually a function) which gets applied to the target expression when it is garbage-collected. This is useful to perform automatic cleanup actions on objects with internal state, such as files. Pure's sentries are *much* more useful than finalizers in other garbage-collected languages, since it is guaranteed that they are called as soon as an object "goes out of scope", i.e., becomes inaccessible.

sentry f x Places a sentry `f` at an expression `x` and returns the modified expression.

clear_sentry x Removes the sentry from an expression `x`.

get_sentry x Returns the sentry of an expression `x` (if any, fails otherwise).

As of Pure 0.45, sentries can be placed on any Pure expression. The sentry itself can also be any type of object (but usually it's a function). Example:

```
> using system;
> sentry (\->puts "I'm done for!") (1..3);
[1,2,3]
> clear ans
I'm done for!
```

Note that setting a finalizer on a global symbol won't usually be of much use since such values are cached by the interpreter. (However, the sentry *will* be invoked if the symbol gets recompiled because its definition has changed. This may be useful for some purposes.)

In Pure parlance, we call an expression **cooked** if a sentry has been attached to it. The following predicate can be used to check for this condition. Also, there is a convenience function to create cooked pointers which take care of freeing themselves when they are no longer needed.

cookedp x Check whether a given object has a sentry set on it.

cooked ptr Create a pointer which disposes itself after use. This is just a shorthand for `sentry free`. The given pointer `ptr` must be *malloced* to make this work.

Example:

```
> using system;
> let p = cooked (malloc 1024);
> cookedp p;
1
> get_sentry p;
free
> clear p
```

Besides their use as finalizers, sentries can also be handy in other circumstances, when you need to associate an expression with another, “invisible” value. In this case the sentry is usually some kind of data structure instead of a function to be executed at finalization time. For instance, here’s how we can employ sentries to implement hashing of function values:

```
using dict;
hashed f x = case get_sentry f of
  h::HDict = h!x if member h x;
  _ = y when y = f x; sentry (update h x y) f
    when h = case get_sentry f of
      h::HDict = h; _ = emptyhdict
    end;
  end;
end;
end;
```

E.g., consider the naive recursive definition of the Fibonacci function:

```
fib n::int = if n<=1 then 1 else fib (n-1)+fib (n-2);
```

A hashed version of the Fibonacci function can be defined as follows:

```
let hfib = hashed f with
  f n::int = if n<=1 then 1 else hfib (n-1)+hfib (n-2)
end;
```

This turns the naive definition of the Fibonacci function (which has exponential time complexity) into a linear time operation:

```
> stats
> fib 35;
14930352
4.53s
> hfib 35;
14930352
0.25s
```

Finally, note that there can be only one sentry per expression but, building on the operations provided here, it’s easy to design a scheme where sentries are chained. For instance:

```
chain_sentry f x = sentry (h (get_sentry x)) x with
  h g x = g x $$ f x;
end;
```

This invokes the original sentry before the chained one:

```
> using system;
> f _ = puts "sentry#1"; g _ = puts "sentry#2";
> let p = chain_sentry g $ sentry f $ malloc 10;
> clear p
sentry#1
sentry#2
```

You can chain any number of sentries that way. This scheme should work in most cases in which sentries are used just as finalizers. However, there are other uses, like the “hashed function” example above, where you’d like the original sentry to stay intact. This can be achieved by placing the new sentry as a sentry on the *original sentry* rather than the expression itself:

```
attach_sentry f x = sentry (sentry f (get_sentry x)) x;
```

This requires that the sentry will actually be garbage-collected when its hosting expression gets freed, so it will *not* work if the original sentry is a global:

```
> let p = attach_sentry g $ sentry f $ malloc 10;
> clear p
sentry#1
```

However, the attached sentry will work ok if you can ensure that the original sentry is a (partial or constructor) application. E.g.:

```
> let p = attach_sentry g $ sentry (f$) $ malloc 10;
> clear p
sentry#1
sentry#2
```

Tagged Pointers

As of Pure 0.45, the C interface now fully checks pointer parameter types at runtime (see the [C Types](#) section in the Pure Manual for details). To these ends, pointer values are internally tagged to keep track of the pointer types. The operations described in this section give you access to these tags in Pure programs. At the lowest level, a pointer tag is simply a machine int associated with a pointer value. The default tag is 0, which denotes a generic pointer value, i.e., `void*` in C. The following operations are provided to create such tags, and set, get or verify the tag of a pointer value.

ptrtag t x Places an integer tag `t` at an expression `x` and returns the modified expression. `x` must be a pointer value.

get_ptrtag x Retrieves the tag associated with `x`.

check_ptrtag t x Compares the tag associated with *x* against *t* and returns true iff the tags match. If *x* is a pointer value, this is equivalent to `get_ptrtag x==t || null x && get_ptrtag x==0`.

make_ptrtag Returns a new, unique tag each time it is invoked.

Examples:

```
> let p = malloc 10;
> get_ptrtag p; // zero by default
0
> let t = make_ptrtag; t;
12
> ptrtag t p;
#<pointer 0xc42da0>
> get_ptrtag p;
12
> check_ptrtag t p;
1
> check_ptrtag 0 p;
0
```

Note that in the case of a non-NULL pointer, `check_ptrtag` just tests the tags for equality. On the other hand, a generic NULL pointer, like in C, is considered compatible with all pointer types:

```
> let t1 = make_ptrtag; t1;
13
> check_ptrtag t1 p;
0
> check_ptrtag t1 NULL;
1
> get_ptrtag NULL;
0
```

The operations above are provided so that you can design your own, more elaborate type systems for pointer values if the need arises. However, you'll rarely have to deal with pointer tags at this level yourself. For most applications, it's enough to inspect the type of a Pure pointer and maybe modify it by "casting" it to a new target type. The following high-level operations provide these capabilities.

pointer_tag type, pointer_tag x Returns the pointer tag for the given type, denoted as a string, or the given pointer value *x*. In the former case, the type should be specified in the C-like syntax used in extern declarations; a new tag will be created using `make_ptrtag` if needed. In the latter case, `pointer_tag` simply acts as a frontend for `get_ptrtag` above.

pointer_type tag, pointer_type x Returns the type name associated with the given int value tag or pointer value *x*. Please note that this may be NULL in the case of an "anonymous" tag, which may have been created with `make_ptrtag` above, or if the tag is simply unknown because it hasn't been created yet.

pointer_cast tag x, pointer_cast type x Casts *x* (which must be a pointer value) to the

given pointer type, which may be specified either as a tag or a string denoting the type name. This returns a new pointer value with the appropriate type tag on it (the tag on the original pointer value `x` isn't affected by this operation).

Example:

```
> let p = malloc 10;
> let q = pointer_cast "char*" p;
> map pointer_type [p,q];
["void*", "char*"]
> map pointer_tag [p,q];
[0,1]
> map pointer_type (0..make_ptrtag-1);
["void*", "char*", "void**", "char**", "short*", "short**", "int*", "int**",
"float*", "float**", "double*", "double**"]
```

(The last command shows a quick and dirty way to retrieve the currently defined type tags in the interpreter. This won't work in batch-compiled scripts, however, since in this case the range of type tags is in general non-contiguous.)

If you have to do many casts to a given type, you can avoid the overhead of repeatedly looking up the type name by assigning the tag to a variable, which can then be passed to `pointer_cast` instead:

```
> let ty = pointer_tag "long*";
> pointer_cast ty p, pointer_cast ty q;
```

Note that you have to be careful when casting a cooked pointer, because `pointer_cast` may have to create a copy of the original pointer value in order not tolobber the original type tag. The sentry will then still be with the original cooked pointer value, thus you have to ensure that this value survives its type-cast duplicate. It's usually best to apply the cast right at the spot where the pointer gets passed to an external function, e.g.:

```
> extern char *gets(char*);
> let p = cooked $ malloc 1000;
> gets (pointer_cast "char*" p);
```

Such usage is always safe. If this approach isn't possible, you might want to use the lowlevel `ptrtag` operation instead. (This willlobber the type tag of the pointer, but you can always change it back afterwards.)

Expression References

Expression references provide a kind of mutable data cells which can hold any Pure expression. If you need these, then you're doomed. ;-) However, they can be useful as a last resort when you need to keep track of some local state or interface to the messy imperative world. Pure's references are actually implemented as expression pointers so that you can readily pass them as pointers to a C function which expects a `pure_expr**` parameter. This may even be useful at times.

ref x Create a reference pointing to `x` initially.

put *r x* Set a new value *x*, and return that value.

get *r* Retrieve the current value *r* points to.

unref *r* Purge the referenced object and turn the reference into a dangling pointer. (This is used as a sentry on reference objects and shouldn't normally be called directly.)

refp *x* Predicate to check for reference values.

Note that manually removing the `unref` sentry of a reference turns the reference into just a normal pointer object and renders it unusable as a reference. Doing this will also leak memory, so don't!

2 Mathematical Functions

The `math.pure` module provides Pure's basic math routines. It also defines complex and rational numbers.

2.1 Imports

To use the operations of this module, add the following import declaration to your program:

```
using math;
```

2.2 Basic Math Functions

The module defines Euler's number `e = 2.71828...` and Ludolph's number `pi = 3.1415...` as constants. It also provides a reasonably comprehensive (pseudo) random number generator which uses the [Mersenne twister](#) to avoid bad generators present in some C libraries.

Please note that as of Pure 0.41, the runtime library includes a newer release of the Mersenne twister which fixes issues with some kinds of seed values, and will yield different values for given seeds. Also, the `random31` and `random53` functions have been added as a convenience to compute unsigned 31 bit integers and 53 bit double values, and the `srandom` function now also accepts an `int` matrix as seed value.

random Return 32 bit pseudo random ints in the range `-0x80000000..0x7fffffff`.

random31 Return 31 bit pseudo random ints in the range `0..0x7fffffff`.

random53 Return pseudo random doubles in the range `[0,1)` with 53 bits resolution.

srandom seed Sets the seed of the generator to the given 32 bit integer. You can also specify longer seeds using a nonempty row vector, e.g.: `srandom {0x123, 0x234, 0x345, 0x456}`.

The following functions work with both double and int/bigint arguments. The result is always a double. For further explanations please see the descriptions of the corresponding functions from the C math library.

sqrt *x* The square root function.

exp *x*, **ln** *x*, **log** *x* Exponential function, natural and decadic logarithms.

sin *x*, **cos** *x*, **tan** *x* Trigonometric functions.

asin *x*, **acos** *x*, **atan** *x* Inverse trigonometric functions.

atan2 *y x* Computes the arcus tangent of y/x , using the signs of the two arguments to determine the quadrant of the result.

sinh *x*, **cosh** *x*, **tanh** *x* Hyperbolic trigonometric functions.

asinh *x*, **acosh** *x*, **atanh** *x* Inverse hyperbolic trigonometric functions.

2.3 Complex Numbers

We provide both rectangular ($x+y$) and polar ($r<:a$) representations, where (x,y) are the Cartesian coordinates and (r,t) the radius (absolute value) and angle (in radians) of a complex number, respectively. The $+$ and $<:$ constructors (declared in the prelude) bind weaker than all other arithmetic operators and are non-associative.

The polar representation $r<:t$ is normalized so that r is always nonnegative and t falls in the range $-\pi < t \leq \pi$.

The constant i is provided to denote the imaginary unit $0+1$.

The arithmetic operations $+$, $*$ etc. and the equality relations $==$ and $~=$ work as expected, and the square root, exponential, logarithms, trigonometric and hyperbolic trigonometric functions (see [Basic Math Functions](#)) are extended to complex numbers accordingly. These do *not* rely on complex number support in the C library, but should still conform to IEEE 754 and POSIX, provided that the C library provides a standards-compliant implementation of the basic math functions.

The following operations all work with both the rectangular and the polar representation, promoting real (double, int/bigint) inputs to complex where appropriate. When the result of an operation is again a complex number, it generally uses the same representation as the input (except for explicit conversions). Mixed rect/polar and polar/rect arithmetic always returns a rect result, and mixed complex/real and real/complex arithmetic yields a rect or polar result, depending on what the complex input was.

complex *x* Convert any kind of number to a complex value.

polar *z*, **rect** *z* Convert between polar and rectangular representations.

cis *t* Create complex values on the unit circle. Note: To quickly compute $\exp(x+y)$ in polar form, use $\exp x <: y$.

abs *z*, **arg** *z* Modulus (absolute value) and argument (angle, a.k.a. phase). Note that you can also find both of these in one go by converting to polar form.

re *z*, **im** *z* Real and imaginary part.

conj *z* Complex conjugate.

Examples:

```
> using math;
> let z = 2^(1/i); z;
0.769238901363972+-0.638961276313635
> let z = ln z/ln 2; z;
0.0+:-1.0
> abs z, arg z;
1.0,-1.5707963267949
> polar z;
1.0<:-1.5707963267949
```

Please note that, as the `+` and `<` constructors bind weaker than the other arithmetic operators, complex numbers *must* be parenthesized accordingly, e.g.:

```
> (1+:2)*(3+:4);
-5+:10
```

2.4 Rational Numbers

Pure's rational numbers are constructed with the **exact division** operator `%` (declared in the prelude) which has the same precedence and fixity as the other division operators.

The `%` operator returns a rational or complex rational for any combination of integer, rational and complex integer/rational arguments, provided that the denominator is nonzero (otherwise it behaves like `x div 0`, which will raise an exception). Machine int operands are always promoted to bigints, thus normalized rationals always take the form `x%y` where both the numerator `x` and the denominator `y` are bigints. For other numeric operands `%` works just like `/`. Rational results are normalized so that the sign is always in the numerator and numerator and denominator are relatively prime. In particular, a rational zero is always represented as `0L%1L`.

The usual arithmetic operations and equality/order relations are extended accordingly, as well as the [basic math functions](#) and the [rounding functions](#), and will return exact (rational or complex rational) results where appropriate. Rational operations are implemented using the GMP bigint functions where possible, and thus are reasonably fast.

In addition, the module also provides following operations:

rational *x* Converts a real or complex value *x* to a rational or complex rational. Note that the conversion from double values doesn't do any rounding, so it is guaranteed that converting the resulting rational back to a double reconstructs the original value.

Conversely, the [int](#), [bigint](#), [double](#), [complex](#), [rect](#), [polar](#) and [cis](#) conversion functions are overloaded so that they convert a rational to one of the other number types.

num x, den x Numerator and denominator of a rational x.

Examples:

```
> using math;
> 5%7 + 2%3;
29L%21L
> 3%8 - 1%3;
1L%24L
> pow (11%10) 3;
1331L%1000L
> let x = pow 3 (-3); x;
1L%27L
> num x, den x;
1L,27L
> rational (3/4);
3L%4L
```

Note that doubles can't represent most rationals exactly, so conversion from double to rational *will* yield funny results in many cases (which are still accurate up to rounding errors). For instance:

```
> let x = rational (1/17); x;
4238682002231055L%72057594037927936L
> num x/den x;
0.0588235294117647
> double (1%17);
0.0588235294117647
```

2.5 Semantic Number Predicates

In difference to the syntactic predicates in [Primitives](#), these check whether the given value can be represented as an object of the given target type (up to rounding errors). Note that if x is of syntactic type X, then it is also of semantic type X. Moreover, `intvalp x => bigintvalp x => ratvalp x => realvalp x => compvalp x <=> numberp x`.

compvalp x Check for complex values (this is the same as [numberp](#)).

realvalp x Check for real values (`im x==0`).

ratvalp x Check for rational values (same as `realvalp`, except that IEEE 754 infinities and NaNs are excluded).

bigintvalp x Check for “big” integer values which can be represented as a bigint.

intvalp x Check for “small” integer values which can be represented as a machine int.

3 Container Types

The standard library provides a variety of efficient container data structures for different purposes. These are all purely functional, i.e., immutable data structures implemented using different flavours of binary trees. This means that instead of modifying a data structure in-place, operations like insertion and deletion return a new instance of the container, keeping the previous instance intact. Nevertheless, all operations are performed efficiently, in logarithmic time where possible.

The container types are all implemented as abstract data structures, so client modules shouldn't rely on the internal representation. Each type provides a corresponding type tag (cf. *Type Tags* in the Pure Manual), as given in the "Data Structure" section of each type, which can be used to match values of the type, e.g.:

```
shift a::Array = rmfirst a;
```

All container types implement the equality predicates `==` and `~=` by recursively comparing their members. In addition, the dictionary, set and bag data structures also provide the other comparison predicates (`<`, `<=` etc.) which check whether one dictionary, set or bag is contained in another.

3.1 Arrays

The `array.pure` module implements an efficient functional array data structure which allows to access and update individual array members, as well as to add and remove elements at the beginning and end of an array. All these operations are carried out in logarithmic time.

Imports

To use the operations of this module, add the following import declaration to your program:

```
using array;
```

Data Structure

Arrays are represented as balanced tree structures of the form `Array T`, thus `Array` may be used as a type tag to restrict variables in pattern matching.

The internal data structure `T` is a binary tree built with the following constructors:

array::nil the empty tree.

array::tip value a leaf of the tree, holding the given value.

array::bin balance left right a nonempty tree with the given balance flag (0 or 1) and left and right subtrees.

(This is for informational purposes only. The tree constructors are private, and client modules must not rely on the internal representation.)

Operations

emptyarray return the empty array
array xs create an array from a list xs
array2 xs create a two-dimensional array from a list of lists
mkarray x n create an array consisting of n x's
mkarray2 x (n,m) create a two-dimensional array of n*m x's
arrayp x check whether x is an array
#a size of a
a!i return the ith member of a
a!(i,j) two-dimensional subscript
null a test whether a is the empty array
members a, list a list of values stored in a
members2 a, list2 a list of members in a two-dimensional array
first a, last a first and last member of a
rmfirst a, rmlast a remove first and last member from a
insert a x insert x at the beginning of a
append a x append x to the end of a
update a i x replace the ith member of a by x
update2 a (i,j) x update two-dimensional array

Examples

Import the module:

```
> using array;
```

A one-dimensional array:

```
> let a::Array = array (0.0:0.1..1.0);  
> #a; members a;  
11  
[0.0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0]
```

Indexing an array works in the usual way, using Pure's `!` operator. By virtue of the prelude, slicing an array with `!!` also works as expected:

```
> a!5;
0.5
> a!!(3..7);
[0.3,0.4,0.5,0.6,0.7]
```

Updating a member of an array produces a new array:

```
> let b::Array = update a 1 2.0;
> members b;
[0.0,2.0,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0]
```

Two-dimensional arrays can be created with `array2` from a list of lists:

```
> let a2::Array = array2 [[i,x | x = [u,v,w]] | i = 1..2];
> members2 a2;
[(1,u),(1,v),(1,w)],[(2,u),(2,v),(2,w)]]
> a2!(1,2);
2,w
> a2!![(0,1),(1,2)];
[(1,v),(2,w)]
> a2!!(0..1,1..2);
[(1,v),(1,w)],[(2,v),(2,w)]]
```

Here's how to convert an array to a Pure matrix:

```
> matrix $ members a;
{0.0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0}
> matrix $ members2 a2;
{(1,u),(1,v),(1,w);(2,u),(2,v),(2,w)}
```

Converting back from a matrix to an array:

```
> let b2::Array = array2 $ list2 {(1,u),(1,v),(1,w);(2,u),(2,v),(2,w)};
> members2 b2;
[(1,u),(1,v),(1,w)],[(2,u),(2,v),(2,w)]]
```

3.2 Heaps

Heaps are a kind of priority queue data structure which allows quick (constant time) access to the smallest member, and to remove the smallest member and insert new elements in logarithmic time. Our implementation does not allow quick update of arbitrary heap members; if such functionality is required, bags can be used instead (see [Sets and Bags](#)).

Heap members *must* be ordered by the `<=` predicate. Multiple instances of the same element may be stored in a heap; however, the order in which equal elements are retrieved is not specified.

Imports

To use the operations of this module, add the following import declaration to your program:

```
using heap;
```

Data Structure

Heaps are represented as balanced tree structures of the form `Heap T`, thus `Heap` may be used as a type tag to restrict variables in pattern matching.

The internal data structure `T` is a binary tree built with the following constructors:

heap::nil the empty tree.

heap::bin balance value left right a nonempty tree holding the given value at its root node, with the given balance flag (0 or 1) and left and right subtrees.

(This is for informational purposes only. The tree constructors are private, and client modules must not rely on the internal representation.)

Operations

emptyheap return the empty heap

heap xs create a heap from a list `xs`

heapp x check whether `x` is a heap

#h size of a heap

null h test whether `h` is the empty heap

members h, list h list the members of `h` in ascending order

first h the first (i.e., smallest) member of `h`

rmfirst h remove the first (i.e., smallest) member from `h`

insert h x insert `x` into `h`

Examples

```
> let h::Heap = heap [5,1,3,11,3];
> members h;
[1,3,3,5,11]
> first h;
1
> members $ rmfirst h;
[3,3,5,11]
```

3.3 Dictionaries

The `dict.pure` module provides Pure's dictionary data types based on AVL trees. There are actually four different types to choose from, depending on whether you need ordered or hashed dictionaries and whether multiple values for the same key should be allowed or not.

- `Dict` is an ordered dictionary. This assumes an ordered key type, i.e., the predicate `<` must be defined on the keys.
- `HDict` a hashed dictionary which works with any (mixture of) key types but stores members in an apparently random order.
- `MDict` is an ordered dictionary, like `Dict`, which allows multiple values to be associated with the same key.
- `HMDict` is a multi-valued dictionary, like `MDict`, but uses hashed keys like `HDict`.

`MDict` and `HMDict` are also colloquially referred to as (ordered or hashed) *multidicts*. This implementation guarantees that different members for the same key are always kept in the order in which they were inserted, and this is also the order in which they will be retrieved by the members, keys, vals and indexing operations.

The usual comparison predicates (`==`, `~=`, `<=`, `<` etc.) are defined on all dictionary types, where two dictionaries are considered "equal" (`d1==d2`) if they both contain the same `key=>value` pairs, and `d1<=d2` means that `d1` is a sub-dictionary of `d2`, i.e., all `key=>value` pairs of `d1` are also contained in `d2` (taking into account multiplicities in the multidict case). Ordered dictionaries compare keys using equality (assuming two keys `a` and `b` to be equal if neither `a<b` nor `b<a` holds), while hashed dictionaries check for syntactical equality (using `===`). The associated values are always compared using the `==` predicate, so this needs to be defined if you want to use any of the comparison operations.

The underlying AVL tree data structure can be found in the `avltrees.pure` module which is included in the library, but not to be invoked directly.

The AVL tree algorithm has its origin in the SWI-Prolog implementation of association lists. The original implementation was created by R. A. O'Keefe and updated for SWI-Prolog by Jan Wielemaker. For the original source see <http://www.swi-prolog.org>.

The port from SWI-Prolog and the deletion stuff (`rmfirst`, `rmlast`, `delete`) missing in the Prolog implementation was provided by Jiri Spitz. The generalization of the code to arbitrary combinations of ordered/hashed and single-/multi-valued keys was done by Albert Graef.

Imports

To use the operations of this module, add the following import declaration to your program:

```
using dict;
```

Data Structure

Dictionaries are represented as balanced tree structures of the form $D \rightarrow T$ where D may be any of `Dict`, `HDict`, `MDict` and `HMDict`. These constructors may be used as type tags to restrict variables in pattern matching.

In any case, the internal data structure T is an AVL tree built with the following constructors (from the `avltrees` module):

avl::nil the empty tree.

avl::bin (key=>value) balance left right a nonempty tree with given key and value in the root node, where `left` and `right` are the left and right subtree, and `balance` is either 1, 0 or -1, denoting $|left| - |right| = 1, 0$, or -1 , respectively.

(This is for informational purposes only. The tree constructors are private, and client modules must not rely on the internal representation.)

Operations

emptydict, emptyhdict, emptymdict, emptyhmdict return an empty dictionary

dict xs, hdict xs, mdict xs, hmdict xs create a dictionary of the corresponding type either from a list `xs` of key-value pairs in the form `key=>value`, or from another dictionary; in the latter case the argument is converted to a dictionary of the desired target type

dictp d, hdictp d, mdictp d, hmdictp d check whether `x` is a dictionary of the corresponding type

mkdict y xs, mkhdict y xs, mkmdict y xs, mkhmdict y xs create a dictionary from a list of keys and a constant value

d1+d2 union: `d1+d2` adds the members of `d2` to `d1`

d1-d2 difference: `d1-d2` removes the members of `d2` from `d1`

d1*d2 intersection: `d1*d2` removes the members *not* in `d2` from `d1`

#d size of a dictionary (the number of members it contains)

d!x get the value from `d` by key `x`; in the case of a multidict this actually returns a list of values (which may be empty if `d` doesn't contain `x`)

null d test whether `d` is an empty dictionary

member d x test whether `d` contains a member with key `x`

members d, list d list the members of `d` (in ascending order for ordered dictionaries)

keys d list the keys of `d` (in ascending order for ordered dictionaries)

vals d list the values of `d`

first d, last d return the first and the last member of `d`, respectively

rmfirst d, rmlast d remove the first and the last member from d, respectively

insert d (x=>y), update d x y insert x=>y into d (this always adds a new member in a multidict, otherwise it replaces an existing value if there is one); note that update is just a fully curried version of insert, so update d x y behaves exactly like insert d (x=>y)

delete d x remove x from d if present (in the multidict case, only the first member with the given key x is removed)

delete_val d (x=>y) remove a specific key-value pair x=>y from d if present (in the multidict case, only the first instance of x=>y is removed); please also see the notes below regarding this operation

delete_all d x remove all instances of x from d (in the non-multidict case, this is just the same as delete)

Notes:

- The infix operators +, - and * work like the corresponding set and bag operations (see [Sets and Bags](#)), treating dictionaries as collections of key=>val pairs. You can mix arbitrary operand types with these operations, as well as with the comparison operations; the necessary conversions from less general dictionary types (ordered, single-valued) to more general types (hashed, multi-valued) are handled automatically.
- The delete_val function always compares values using equality (==), so this predicate must be defined on the values stored in the dictionary in order to make this operation work. If there is more than one instance of the given value under the given key, the first such instance will be removed (which may be any instance that compares equal, not necessarily an exact match).
- In the multidict case, delete_val may require linear time with respect to the number of different values stored under the given key. Since this operation is also needed to implement some other multidict operations like comparisons, difference and intersection, these may end up requiring quadratic running times in degenerate cases (i.e., if the majority of members happens to be associated with only very few keys).

Examples

A normal (ordered) dictionary:

```
> using dict;
> let d::Dict = dict ["foo"=>77,"bar"=>99.1];
> keys d; vals d; members d;
["bar", "foo"]
[99.1, 77]
["bar"=>99.1, "foo"=>77]
```

Indexing a dictionary works in the usual way, using Pure's ! operator. An out_of_bound exception is thrown if the key is not in the dictionary:

```
> d!"foo";
77
> d!"baz";
<stdin>, line 5: unhandled exception 'out_of_bounds' while evaluating
'd!"baz"'
```

By virtue of the prelude, slicing a dictionary with !! also works as expected:

```
> d!!["foo", "bar", "baz"];
[77, 99.1]
```

A hashed dictionary can be used with any key values, which are stored in a seemingly random order:

```
> let h::HDict = hdict [foo=>77, 42=>99.1];
> keys h; vals h; members h;
[42, foo]
[99.1, 77]
[42=>99.1, foo=>77]
> h!foo;
77
> h!!keys h;
[99.1, 77]
```

Multidicts work in pretty much the same fashion, but allow more than one value for a given key to be stored in the dictionary. In this case, the indexing operation returns a list of all values for the given key, which may be empty if the key is not in the dictionary (rather than throwing an out_of_bound exception):

```
> let d::MDict = mdict ["foo"=>77, "bar"=>99.1, "foo"=>99];
> d!"foo"; d!"baz";
[77, 99]
[]
```

Slicing thus returns a list of lists of values here:

```
> d!!["foo", "bar", "baz"];
[[77, 99], [99.1], []]
```

To obtain a flat list you can just concatenate the results:

```
> cat $ d!!["foo", "bar", "baz"];
[77, 99, 99.1]
```

Hashed multidicts provide both key hashing and multiple values per key:

```
> let h::HMDict = hmdict [foo=>77, 42=>99.1, 42=>77];
> keys h; vals h; members h;
[42, 42, foo]
[99.1, 77, 77]
[42=>99.1, 42=>77, foo=>77]
> h!42;
[99.1, 77]
```

There are also some set-like operations which allow you to add/remove the members (key=>val pairs) of one dictionary to/from another dictionary, and to compute the intersection of two dictionaries. For instance:

```
> let h1 = hmdict [a=>1,b=>2];
> let h2 = hmdict [b=>2,c=>3];
> members (h1+h2);
[a=>1,c=>3,b=>2,b=>2]
> members (h1-h2);
[a=>1]
> members (h1*h2);
[b=>2]
```

It's possible to mix dictionaries of different types in these operations. The necessary conversions are handled automatically:

```
> let h1 = hmdict [a=>1,b=>2];
> let h2 = hdict [b=>3,c=>4];
> members (h1+h2);
[a=>1,c=>4,b=>2,b=>3]
```

Note that the result will always be promoted to the most general operand type in such cases (a hashed multidict in the above example). If this is not what you want, you'll have to apply the necessary conversions manually:

```
> members (hdict h1+h2);
[a=>1,c=>4,b=>3]
```

3.4 Sets and Bags

The `set.pure` module implements Pure's set data types based on AVL trees. These work pretty much like dictionaries (cf. [Dictionaries](#)) but only store keys (called "elements" or "members" here) without any associated data values. Hence sets provide membership tests like dictionaries, but no indexing operations.

There are four variations of this data structure to choose from, depending on whether the set members are ordered or hashed, and whether multiple instances of the same element are allowed (in this case the set is actually called a *multiset* or a *bag*).

- Set and Bag implement the ordered set types. They require that members be ordered, i.e., the predicate `<` must be defined on them.
- HSet and HBag implement the hashed set types which don't require an order of the members. Distinct members are stored in an apparently random order.

The usual comparison predicates (`==`, `~=`, `<=`, `<` etc.) are defined on all set and bag types, where two sets or bags are considered "equal" (`m1==m2`) if they both contain the same elements, and `m1<=m2` means that `m1` is a subset or subbag of `m2`, i.e., all elements of `m1` are also contained in `m2` (taking into account multiplicities in the multiset case). Ordered sets and bags compare elements using equality (considering two elements `a` and `b` to be equal if

neither $a < b$ nor $b < a$ holds), while hashed sets and bags check for syntactical equality (using `==`).

The underlying AVL tree data structure can be found in the `avltrees.pure` module which is included in the library, but not to be invoked directly. The AVL tree algorithm has its origin in the SWI-Prolog implementation of association lists and was ported to Pure by Jiri Spitz, see [Dictionaries](#) for details.

Imports

To use the operations of this module, add the following import declaration to your program:

```
using set;
```

Data Structure

Sets and bags are represented as balanced tree structures of the form $S\ T$, where S may be any of `Set`, `Bag`, `HSet` and `HBag`. These constructors may be used as type tags to restrict variables in pattern matching.

In any case, the internal data structure T is an AVL tree built with the following constructors (from the `avltrees` module):

`avl::nil` the empty tree.

`avl::bin key balance left right` a nonempty tree with given key (set element) in the root node, where `left` and `right` are the left and right subtree, and `balance` is either 1, 0 or -1, denoting $|left| - |right| = 1, 0$, or -1 , respectively.

(This is for informational purposes only. The tree constructors are private, and client modules must not rely on the internal representation.)

Operations

`emptyset, emptybag, emptyhset, emptyhbag` return an empty set or bag

`set xs, bag xs, hset xs, hbag xs` create a set or bag of the corresponding type from a list or another set or bag `xs`; in the latter case the argument is converted to a set or bag of the desired target type

`setp x, bagp x, hsetp x, hbagp x` check whether `x` is a set or bag of the corresponding type

`m1+m2` set and bag union: `m1+m2` adds the members of `m2` to `m1`

`m1-m2` set and bag difference: `m1-m2` removes the members of `m2` from `m1`

`m1*m2` set and bag intersection: `m1*m2` removes the members *not* in `m2` from `m1`

`#m` size of a set or bag `m`

null *m* test whether *m* is an empty set or bag

member *m x* test whether *m* contains *x*

members *m*, **list** *m* list the members of *m* (in ascending order for ordered sets and bags)

first *m*, **last** *m* return the first and the last member of *m*, respectively

rmfirst *m*, **rmlast** *m* remove the first and the last member from *m*, respectively

insert *m x* insert *x* into *m* (replaces an existing element in the case of a set)

delete *m x* remove *x* from *m* (in the bag case, only the first instance of *x* is removed)

delete_all *m x* remove all instances of *x* from *m* (in the set case, this is just the same as **delete**)

Note that the infix operators (+, -, *, as well as the comparison operations) allow you to mix arbitrary operand types; the necessary conversions from less general set types (ordered, set) to more general types (hashed, multiset) are handled automatically.

Examples

Some basic set operations:

```
> let m::Set = set [5,1,3,11,3];
> members m;
[1,3,5,11]
> map (member m) (1..5);
[1,0,1,0,1]
> members $ m+set (3..6);
[1,3,4,5,6,11]
> members $ m-set (3..6);
[1,11]
> members $ m*set (3..6);
[3,5]
```

The bag operations work in a similar fashion, but multiple instances are permitted in this case, and each instance counts as a separate member:

```
> let m::Bag = bag [5,1,3,11,3];
> members m;
[1,3,3,5,11]
> members $ delete m 3;
[1,3,5,11]
> members $ insert m 1;
[1,1,3,3,5,11]
> members $ m+bag (3..6);
[1,3,3,3,4,5,5,6,11]
> members $ m-bag (3..6);
[1,3,11]
> members $ m*bag (3..6);
[3,5]
```


As already mentioned, operands of different types can be mixed with the infix operators; the necessary conversions are handled automatically. E.g., here's how you add a set to a bag:

```
> let m1::Bag = bag [5,1,3,11,3];
> let m2::Set = set (3..6);
> members (m1+m2);
[1,3,3,3,4,5,5,6,11]
```

Note that the result will always be promoted to the most general operand type in such cases (a bag in the above example). If this is not what you want, you'll have to apply the necessary conversions manually:

```
> members (set m1+m2);
[1,3,4,5,6,11]
```

If set members aren't ordered then you'll get an exception when trying to create an ordered set or bag from them:

```
> set [a,b,c];
<stdin>, line 5: unhandled exception 'failed_cond' while evaluating
'set [a,b,c]'
```

In such a case hashed sets and bags must be used instead. These work analogously to the ordered sets and bags, but distinct members are stored in an apparently random order:

```
> members $ hset [a,b,c] * hset [c,d,e];
[c]
> members $ hbag [a,b,c] + hbag [c,d,e];
[a,c,c,b,d,e]
```

4 System Interface

This module offers some useful system routines, straight from the C library, as well as some convenience functions for wrapping these up in Pure. Even the "purest" program needs to do some basic I/O every once in a while, and this module provides the necessary stuff to do just that. The operations provided in this module should work (if necessary by a suitable emulation) on all supported systems. Most of the following functions are extensively documented in the C library manual pages, so we concentrate on the Pure-specific aspects here.

4.1 Imports

To use the operations of this module, add the following import declaration to your program:

```
using system;
```

Some functions of the system interface are provided in separate modules; see [Additional POSIX Functions](#) and [Option Parsing](#).

4.2 Errno and Friends

errno, **set_errno n**, **perror msg**, **strerror n** This value and the related routines are indispensable to give proper diagnostics when system calls fail for some reason. Note that, by its very nature, `errno` is a fairly volatile value, don't expect it to survive a return to the command line in interactive sessions.

Example:

```
> using system;
> fopen "junk" "r", perror "junk";
junk: No such file or directory
fopen "junk" "r"
```

4.3 POSIX Locale

setlocale category locale Set or retrieve the current locale.

Details are platform-specific, but you can expect that at least the categories `LC_ALL`, `LC_COLLATE`, `LC_CTYPE`, `LC_MONETARY`, `LC_NUMERIC` and `LC_TIME` are defined, as well as the following values for the locale parameter: `"C"` or `"POSIX"` (the default POSIX locale), `"` (the system default locale), and `NULL`, to just query the current locale.

Other string values which can be passed as the locale argument depend on the implementation, please check your local `setlocale(3)` documentation for details. If locale is not `NULL`, the current locale is changed accordingly. The return value is the new locale, or the current locale when passing `NULL` for the locale parameter. In either case, the string returned by `setlocale` is such that it can be passed to `setlocale` to restore the same locale again. In case of an error, `setlocale` fails (rather than returning a null pointer).

Please note that calling this function alters the Pure interpreter's idea of what the current locale is. When the interpreter starts up, it always sets the default system locale. Unless your scripts rely on a specific encoding, setting the locale to either `"C"` or `"` should always be safe.

Example:

```
> setlocale LC_ALL NULL;
"en_US.UTF-8"
```

4.4 Signal Handling

trap action sig Establish or remove Pure signal handlers.

The action parameter of `trap` can be one of the predefined integer values `SIG_TRAP`, `SIG_IGN` and `SIG_DFL`. `SIG_TRAP` causes the given signal to be handled by mapping it to a Pure exception of the form `signal sig`. `SIG_IGN` ignores the signal, `SIG_DFL` reverts to the system's default handling. See `show -g SIG*` for a list of known signal values on your system.

Note: When the interpreter runs interactively, most standard termination signals (SIGINT, SIGTERM, etc.) are already set up to report corresponding Pure exceptions; if this is not desired, you can use `trap` to either ignore these or revert to the default handlers instead.

See [Exception Handling](#) in the Pure Manual for details and examples.

4.5 Time Functions

The usual date/time functions from the C library are all provided. This includes some functions to retrieve wallclock and cpu time which usually offer much better resolution than the venerable `time` function.

time Reports the current time in seconds since the **epoch**, 00:00:00 UTC, Jan 1 1970. The result is always a bigint (in fact, the time value is already 64 bit on many OSes nowadays).

gettimeofday Returns wallclock time as seconds since the epoch, like `time`, but theoretically offers resolutions in the microsec range (actual resolutions vary, but are usually in the msec range for contemporary systems). The result is returned as a double value (which also limits precision). This function may actually be implemented through different system calls, depending on what's available on the host OS.

clock Returns the current CPU (not wallclock) time since an arbitrary point in the past, as a machine int. The number of "ticks" per second is given by the `CLOCKS_PER_SEC` constant. Note that this value will wrap around approximately every 72 minutes.

sleep t, nanosleep t Suspend execution for a given time interval in seconds. `sleep` takes integer (int/bigint) arguments only and uses the `sleep()` system function. `nanosleep` also accepts double arguments and theoretically supports resolutions down to 1 nanosecond (again, actual resolutions vary). This function may actually be implemented through different system calls, depending on what's available on the host OS. Both functions usually return zero, unless the sleep was interrupted by a signal, in which case the time remaining to be slept is returned.

Examples:

```
> time,sleep 1,time;
1270241703L,0,1270241704L
> gettimeofday,nanosleep 0.1,gettimeofday;
1270241709.06338,0.0,1270241709.16341
```

Here's a little macro which lets you time evaluations:

```
def timex x = y,(t2-t1)/CLOCKS_PER_SEC when
  t1 = clock; y = x; t2 = clock;
end;
```

Example:

```
> timex (foldl (+) 0 (1..100000));
705082704,0.07
```

The `tzset` function calls the corresponding routine from the C library and initializes the (Pure) variables `tzname`, `timezone` and `daylight` accordingly. See the `tzset(3)` manual page for details. This routine is also called automatically when the system module is loaded, so you only have to invoke it to get up-to-date information after changes to the locale or the timezone. Example:

```
> tzset;
()
> tzname, timezone, daylight;
["CET", "CEST"], -3600, 1
> tzname!daylight;
"CEST"
```

The following functions deal with date/time values in string and “broken-down” time format. See the `ctime(3)`, `gmtime(3)`, `localtime(3)`, `mktime(3)`, `asctime(3)`, `strftime(3)` and `strptime(3)` manual pages for details.

ctime t Convert a time value as returned by the `time` function to a string in local time.

gmtime t, localtime t Convert a time value to UTC or local time in “broken-down” form (a static pointer to a `tm` struct containing a bunch of `int` fields) which can then be passed to the `asctime` and `strftime` functions, or to *int_matrix* if you want to convert the data to a matrix; see the example below.

mktime tm Converts broken-down time to a time value (seconds since the epoch). As with `time`, the result is always a `bigint`.

asctime tm, strftime format tm Format broken-down time as a string. `strftime` also uses a format string supplied by the user, see below for a list of the most important conversion specifiers.

strptime s format tm Parse a date/time string `s` according to the given format (using more or less the same format specifiers as the `strftime` function) and store the broken-down time result in the given `tm` struct. This function may fail, e.g., if `strptime` finds an error in the format string. Otherwise it returns the part of the string which wasn’t processed, see the example below.

Examples:

```
> let t = time; t;
1270239790L
> let tm = localtime t; tm;
#<pointer 0x7ff97ecbdde0>
> mktime tm;
1270239790L
> asctime tm;
"Fri Apr  2 22:23:10 2010\n"
> int_matrix 9 tm;
{10,23,22,2,3,110,5,91,1}
> strftime "%c" tm;
"Fri 02 Apr 2010 10:23:10 PM CEST"
> strptime ans "%c" tm, int_matrix 9 tm;
"CEST",{10,23,22,2,3,110,5,91,1}
```

In the above example, `strptime` was given a static pointer to a `tm` struct returned by `localtime`. This always works, but in some situations it may be preferable to allocate dynamic storage instead. This storage should be properly initialized (zeroed out) before passing it to `strptime`, since `strptime` only stores the values specified (at least in principle; please consult your local C library documentation for details). Also note that while POSIX only specifies nine `int` fields in a `tm` struct, depending on the host operating system the struct may contain additional public and private fields. The actual size of a `tm` struct is given by the `sizeof_TM` constant, so a safe way to allocate suitable dynamic storage for the `strptime` function is as follows:

```
> let tm = pointer_cast "int*" $ calloc 1 sizeof_TM;
> strptime "4/2/10" "%D" tm, int_matrix 9 tm;
"", {0,0,0,2,3,110,5,91,0}
```

Instead of explicitly allocating dynamic storage and converting it to a Pure matrix later, you can also invoke `strptime` directly with an `int` matrix of sufficient size:

```
> let tm = imatrix (sizeof_TM div sizeof_INT + 1);
> strptime "4/2/10" "%D" tm, take 9 tm;
"", {0,0,0,2,3,110,5,91,0}
```

Last but not least, to make calling `strptime` more convenient, you can supply your own little wrapper function which takes care of allocating the storage, e.g.:

```
mystrptime s format = s, take 9 tm when
  tm = imatrix (sizeof_TM div sizeof_INT + 1);
  s = strptime s format tm;
end;

> mystrptime "4/2/10" "%D";
"", {0,0,0,2,3,110,5,91,0}
```

Here is a list of some common format specifiers which can be used with the `strftime` and `strptime` routines. These are all specified by POSIX and should thus be available on most platforms. Note that many more formats are usually supported than what is listed here, so please consult your local manual pages for the complete list.

- `%d`, `%m`, `%y`: Day of the month, month and year as decimal two-digit numbers.
- `%Y`: The year as a four-digit number which includes the century.
- `%H`, `%M`, `%S`: Hours (range 00 to 23), minutes and seconds as decimal two-digit numbers.
- `%I`: The hours on a 12-hour clock (range 01 to 12).

The following formats are locale-dependent:

- `%a`, `%A`: Abbreviated and full weekday name.
- `%b`, `%B`: Abbreviated and full month name.
- `%p`: AM or PM. `%P` is the same in lowercase (`strftime` only).

There are also some useful meta-formats which specify various combinations of the above:

- %c: The preferred date and time representation for the current locale.
- %D: The American date format (%m/%d/%y).
- %F: The ISO 8601 date format (%Y-%m-%d). (This is generally supported by `strftime` only, but `strptime` from GNU libc has it.)
- %r: The time in AM/PM notation (%I:%M:%S %p).
- %R: The time in 24-hour notation (%H:%M).
- %T: The time in 24-hour notation, including seconds (%H:%M:%S).

In addition, %% denotes a literal % character, %n newlines and %t tabs. (For `strptime` the latter two are synonymous and match arbitrary whitespace.)

Windows users should note that `strptime` isn't natively supported there. A basic emulation is provided by the Pure runtime, but at present this only supports the C locale.

4.6 Process Functions

The following process functions are available on all systems. (Some additional process-related functions such as `fork`, `kill`, `wait` and `waitpid` are available in the `posix` module, see [Additional POSIX Functions](#).)

system cmd Execute a shell command.

execv prog argv, execvp prog argv, execve prog argv envp Execute a new process. `prog` denotes the name of the executable to be run, `argv` the argument vector (which repeats the program name in the first component), and `envp` a vector of environment strings of the form "var=value". The `execv` function executes the program `prog` exactly as given, while `execvp` also performs a path search. The `execve` function is like `execv`, but also specifies an environment to be passed to the process. In either case, the new process replaces the current process. For convenience, both `argv` and `envp` can be specified as a Pure string vector or a list, which is automatically translated to the raw, NULL-terminated C string vectors (i.e., `char**`) required by the underlying C functions.

spawnv mode prog argv, spawnvp mode prog argv, spawnve mode prog argv envp Spawn a new child process. These work like the corresponding MS Windows functions; on Unix systems this functionality is implemented using a combination of `fork` and `execv`. The arguments are the same as for the `execv` functions, except that there's an additional `mode` argument which specifies how the process is to be executed: `P_WAIT` waits for the process to finish, after which `spawnv` returns with the exit status of the terminated child process; `P_NOWAIT` makes `spawnv` return immediately, returning the process id; and `P_OVERLAY` causes the child process to replace its parent, just like with `execv`. (On Windows, there's an additional `P_DETACH` flag which works like `P_NOWAIT` but also turns the child process into a background task.)

Note that, in addition, the prelude provides the `exit` function which terminates the program with a given exit code, cf. [Other Special Primitives](#).

Examples:

```
> system "pwd";  
/home/ag/svn/pure-lang/trunk/pure/lib  
0  
> spawnvp P_WAIT "pwd" ["pwd"];  
/home/ag/svn/pure-lang/trunk/pure/lib  
0  
> spawnvp P_WAIT "/bin/sh" ["/bin/sh", "-c", "pwd"];  
/home/ag/svn/pure-lang/trunk/pure/lib  
0
```

4.7 Basic I/O Interface

Note that this module also defines the standard I/O streams `stdin`, `stderr` and `stdout` as variables on startup. These are ready to be used with the operations described below. Also note that for convenience some of the following routines are actually Pure wrappers, rather than just providing the raw C library routines.

fopen name mode, popen cmd mode Open a file or a pipe. These take care of closing a file object automatically when it's garbage-collected, and fail (instead of returning a null pointer) in case of error, so that you can provide any desired error handling simply by adding suitable equations.

fdopen fd mode Associates a file object with a given existing file descriptor. Otherwise works like `fopen`, so the resulting file is closed automatically when it's garbage-collected.

freopen path mode fp Reopens a file object. The existing file object is closed. Otherwise works like `fopen`, so the resulting file is closed automatically when it's garbage-collected.

fclose fp, pclose fp Close a file or a pipe.

tmpfile Creates a unique temporary file (opened in "w+b" mode) which gets deleted automatically when it is closed or the file object gets garbage-collected.

feof fp, ferror fp, clearerr fp Check the end-of-file and error bits. `clearerr` clears the error bit.

fileno fp Returns the file descriptor associated with the given file.

fflush fp Flushes the given file (or all open files if `fp` is `NULL`).

fgets fp, gets Pure wrappers for the C `fgets` and `gets` functions which handle the necessary buffering automatically.

fget fp A variation of `fgets` which slurps in an entire text file at once.

fputs s fp, puts s Output a string to the given file or `stdout`, respectively. These are just the plain C functions. Note that `puts` automatically adds a newline, while `fputs` doesn't. Hmm.

fread ptr size nmemb fp, fwrite ptr size nmemb fp Binary read/writes. Here you'll have to manage the buffers yourself. See the corresponding manual pages for details.

fseek fp offset whence, ftell fp, rewind fp Reposition the file pointer and retrieve its current value. The constants `SEEK_SET`, `SEEK_CUR` and `SEEK_END` can be used for the whence argument of `fseek`. The call `rewind fp` is equivalent to `fseek fp 0 SEEK_SET` (except that the latter also returns a result code). See the corresponding manual pages for details.

setbuf fp buf, setvbuf fp buf mode size Set the buffering of a file object, given as the first argument. The second argument specifies the buffer, which must be a pointer to suitably allocated memory or `NULL`. The mode argument of `setvbuf` specifies the buffering mode, which may be one of the predefined constants `_IONBF`, `_IOLBF` and `_IOFBF` denoting no buffering, line buffering and full (a.k.a. block) buffering, respectively; the size argument denotes the buffer size.

For `setbuf`, the given buffer must be able to hold `BUFSIZ` characters, where `BUFSIZ` is a constant defined by this module. `setbuf fp buf` is actually equivalent to the following call (except that `setvbuf` also returns an integer return value):

```
setvbuf fp buf (if null buf then _IONBF else _IOFBF) BUFSIZ
```

Please see the `setbuf(3)` manual page for details.

Examples:

```
> puts "Hello, world!";
Hello, world!
14

> map fileno [stdin, stdout, stderr];
[0,1,2]

> let fp = fopen "/etc/passwd" "r";
> fgets fp;
"at:x:25:25:Batch jobs daemon:/var/spool/atjobs:/bin/bash\n"
> fgets fp;
"avahi:x:103:104:User for Avahi:/var/run/avahi-daemon:/bin/false\n"
> ftell fp;
121L
> rewind fp;
()
> fgets fp;
"at:x:25:25:Batch jobs daemon:/var/spool/atjobs:/bin/bash\n"

> split "\n" $ fget $ popen "ls *.pure" "r";
["array.pure", "dict.pure", "getopt.pure", "heap.pure", "math.pure",
"matrices.pure", "prelude.pure", "primitives.pure", "quasiquote2.pure",
"quasiquote.pure", "set.pure", "strings.pure", "system.pure", ""]
```

C-style formatted I/O is provided through the following wrappers for the C `printf` and `scanf` functions. Our wrapper functions take or return a tuple of values, and check these

against the format specifiers, so they shouldn't segfault. However, only simple formats derived from %cdioux, %efg, %s and %p are supported right now.

printf format args, fprintf fp format args Print a formatted string to stdout or the given file, respectively. Normally, these functions return the result of the underlying C routines (number of characters written, or negative on error). However, in case of an abnormal condition in the wrapper function (error in format string, argument mismatch), they will throw an exception.

sprintf format args Print a formatted string to a buffer and return the result as a string. Unlike the C routine, this wrapper just returns the string result, or a null pointer in case of an error; otherwise, the error handling is the same as with printf and fprintf. The implementation actually uses the C routine snprintf for safety, and a suitable output buffer is provided automatically.

scanf format, fscanf fp format Read formatted input from stdin or the given file, respectively. These normally return a tuple (or singleton) with the converted values. An exception of the form `scanf_error ret`, where `ret` is the tuple of successfully converted values (which may be less than the number of requested input items), is thrown if end-of-file was met or another error occurred while still reading. The handling of other abnormal conditions (e.g., error in format string) is analogous to printf et al. Also note that our implementation here doesn't accept any of the length modifiers used by the C routines. Floating point values will *always* be read in double precision, so you just specify "e", "g" etc. for these. However, the "assignment suppression" flag "*" is understood; the corresponding items will not be returned.

sscanf s format This works exactly like fscanf, but input comes from a string (first argument) rather than a file.

Examples:

```
> do (printf "%s%d\n") [("foo",5),("catch",22)];
foo5
catch22
()
> sscanf "foo 5 22" "%s %d %g";
"foo",5,22.0
```

4.8 Stat and Friends

stat path Return information about the given file. This is a simple wrapper around the corresponding system call, see the stat(2) manual page for details. The function returns a tuple with the most important fields from the stat structure, in this order: `st_dev`, `st_ino`, `st_mode`, `st_nlink`, `st_uid`, `st_gid`, `st_rdev`, `st_size`, `st_atime`, `st_mtime`, `st_ctime`. Among these, `st_mode`, `st_nlink`, `st_uid` and `st_gid` are simple machine integers, the rest is encoded as bigints (even on 32 bit platforms).

lstat path Return information about the given symbolic link (rather than the file it points to). On systems where this function isn't supported (e.g., Windows), `lstat` is identical to `stat`.

fstat fp Return information about the given file object. Same as `stat`, but here the file is given as a file pointer created with `fopen` (see [Basic I/O Interface](#) above). Note that the corresponding system function actually takes a file descriptor, so the Pure implementation is equivalent to the C call `fstat(fileno(fp))`. This function might not be supported on all platforms.

For average applications, the most interesting fields are `st_mode` and `st_size`, which can be retrieved with `stat filename!! [2,7]`. Note that to facilitate access to the `st_mode` field, the usual masks and bits for file types (`S_IFMT`, `S_IFREG`, etc.) and permissions (`S_ISUID`, `S_ISGID`, `S_IRWXU`, etc.) are defined as constants by this module. Use the command `show -g S_*` in the interpreter to get a full list of these. Other interesting fields are `st_atime`, `st_mtime` and `st_ctime`, which can be accessed using `stat filename!! (8..10)`. The values of these fields are the times of last access, last modification and creation, respectively, which can be decoded using the appropriate time functions like `ctime` or `strftime`, see [Time Functions](#).

Examples:

```
> stat "/etc/passwd";
64773L,9726294L,33188,1,0,0,0L,1623L,1250373163L,1242692339L,1242692339L
> stat "/etc/passwd"!7;                                // file size
1623L
> strftime "%c" $ localtime $ stat "/etc/passwd"!10;    // creation time
"Tue 19 May 2009 02:18:59 AM CEST"
> sprintf "%o" $ stat "/etc/passwd"!2 and not S_IFMT;    // permissions
"0644"
> stat "/etc/passwd"!2 and S_IFMT == S_IFREG;            // this is a regular file
1
> stat "/etc"!2 and S_IFMT == S_IFDIR;                   // this is a directory
1
```

4.9 Readdir

readdir name Read the contents of the given directory and return the names of all its entries as a list.

Example:

```
> readdir "/home";
["ag", ".", ".."]
```

4.10 Shell Globbing

fnmatch pat s flags Returns a simple truth value (1 if `pat` matches `s`, 0 if it doesn't), instead of an error code like the C function.

glob pat flags Returns a Pure list with the matches (unless there is an error in which case the integer result code of the underlying C routine is returned).

The available flag values and glob error codes are available as symbolic `FNM_*` and `GLOB_*` constants defined as variables in the global environment. See the `fnmatch(3)` and `glob(3)` manpages for the meaning of these.

Example:

```
> glob "*.pure" 0;
["array.pure", "dict.pure", "getopt.pure", "heap.pure", "math.pure",
"matrices.pure", "prelude.pure", "primitives.pure", "set.pure",
"strings.pure", "system.pure"]
```

4.11 Regex Matching

The POSIX regex functions (`regcomp` and `regexexec`) have a somewhat difficult calling sequence, hence we provide a couple of rather elaborate high-level wrapper functions for use in Pure programs. These are implemented in terms of a low-level interface provided in the runtime. (The low-level interface isn't documented here, but these functions are also callable if you want to create your own regular expression engines in Pure. You might wish to take a look at the implementation of the high-level functions in `system.pure` to see how this can be done.)

regex pat cflags s eflags Compiles and matches a regex in one go, and returns the list of submatches (if any).

The arguments are:

- `pat::string`, the regular expression pattern;
- `cflags::int`, the compilation flags (bitwise or of any of the flags accepted by `regcomp(3)`);
- `s::string`, the subject string to be matched;
- `eflags::int`, the matching execution flags (bitwise or of any of the flags accepted by `regexexec(3)`).

Symbolic `REG_*` constants are provided for the different flag values, see the `regcomp(3)` manpage for an explanation of these. Two particularly important compilation flags (to be included in the `cflags` argument) are `REG_NOSUB`, which prevents submatches to be computed, and `REG_EXTENDED`, which switches regex from “basic” to “extended” regular expressions so that it understands all the regular expression elements of `egrep(1)` in the pattern argument.

Depending on the flags and the outcome of the operation, the result of this function can take one of the following forms:

- `regerr code msg`: This indicates an error during compilation of the pattern (e.g., if there was a syntax error in the pattern). `code` is the nonzero integer code returned by `regcomp`, and `msg` is the corresponding error message string, as returned by `regerror`. You can redefine the `regerr` function as appropriate for your application (e.g., if you'd like to print an error message or throw an exception).

- 0 or 1: Just a truth value indicates whether the pattern matched or not. This will be the form of the result if the REG_NOSUB flag was specified for compilation, indicating that no submatch information is to be computed.
- 0 (indicating no match), or 1 (indicating a successful match), where the latter value is followed by a tuple of (pos, substr) pairs for each submatch. This will be the form of the result only if the REG_NOSUB flag was *not* specified for compilation, so that submatch information is available.

Note that, according to POSIX semantics, a return value of 1 does *not* generally mean that the entire subject string was matched, unless you explicitly tie the pattern to the beginning (^) and end (\$) of the string.

If the result takes the latter form, each (pos, substr) pair indicates a portion of the subject string which was matched; pos is the position at which the match starts, and substr is the substring (starting at position pos) which was matched. The first (pos, substr) pair always indicates which portion of the string was matched by the entire pattern, the remaining pairs represent submatches for the parenthesized subpatterns of the pattern, as described on the regcomp(3) manual page. Note that some submatches may be empty (if they matched the empty string), in which case a pair (pos, "") indicates the (nonnegative) position pos where the subpattern matched the empty string. Other submatches may not participate in the match at all, in which case the pair (-1, "") is returned.

The following helper functions are provided to analyze the result returned by regex.

reg_result res Returns the result of a regex call, i.e., a regerr term if compilation failed, and a flag indicating whether the match was successful otherwise.

reg_info res Returns the submatch info if any, otherwise it returns ().

reg_n_info Returns the nth submatch of the given submatch info, where info is the result of a reg_info call.

regs_info Returns all valid submatches, i.e., the list of all triples (n,p,s) for which reg_n == (p,s) with p>=0.

In addition, the following convenience functions are provided to perform global regex searches, to perform substitutions, and to tokenize a string according to a given delimiter regex.

regexg f pat cflags s eflags Perform a global regular expression search. This routine will scan the entire string for (non-overlapping) instances of the pattern, applies the given function f to the reg_info for each match, and collects all results in a list. Note: Never specify the REG_NOSUB flag with this function, it needs the submatch info.

regexgg f pat cflags s eflags This works like regexg, but allows overlapping matches.

regsub f pat cflags s eflags Replaces all non-overlapping instances of a pattern with a computed substitution string. To these ends, the given function f is applied to the reg_info for each match. The result string is then obtained by concatenating f info for all matches, with the unmatched portions of the string in between. To make this work, f must always return a string value; otherwise, regsub throws a bad_string_value exception.

regsplit pat cflags s eflags Splits a string into constituents delimited by substrings matching the given pattern.

Please note that these operations all operate in an eager fashion, i.e., they process the entire input string in one go. This may be unwieldy or at least inefficient for huge amounts of text. As a remedy, the following lazy alternatives are available:

- `regexgs f pat cflags s eflags`
- `regexggs f pat cflags s eflags`
- `regsplits pat cflags s eflags`

These work like `regexg`, `regexgg` and `regsplit` above, but return a stream result which enables you to process the matches one by one, using “call by need” evaluation.

Basic Examples

Let’s have a look at some simple examples:

```
> let pat = "[[:alpha:]]+[[:alnum:]]*";  
> let s = "1var foo 99 BAR $%&";
```

Simple match:

```
> regex pat 0 s 0;  
1,1,"var"
```

Same without match info:

```
> regex pat REG_NOSUB s 0;  
1
```

Global match, return the list of all matches:

```
> regexg id pat 0 s 0;  
[(1,"var"),(5,"foo"),(12,"BAR")]
```

Same with overlapping matches:

```
> regexgg id pat 0 s 0;  
[(1,"var"),(2,"ar"),(3,"r"),(5,"foo"),(6,"oo"),(7,"o"),(12,"BAR"),  
(13,"AR"),(14,"R")]
```

Note that `id` (the identity function) in the examples above can be replaced with an arbitrary function which processes the matches. For instance, if we only want the matched strings instead of the full match info:

```
> regexg (!1) pat 0 s 0;  
["var","foo","BAR"]
```

Lazy versions of both `regexg` and `regexgg` are provided which return the result as a stream instead. These can be processed in a “call by need” fashion:

```
> regexgs id pat 0 s 0;
(1,"var"):##<thunk 0x7fb1b7976750>
> last ans;
12,"BAR"
```

Let's verify that the processing is really done lazily:

```
> test x = printf "got: %s\n" (str x) $$ x;
> let xs = regexgs test pat 0 s 0;
got: 1,"var"
> xs!1;
got: 5,"foo"
5,"foo"
> last xs;
got: 12,"BAR"
12,"BAR"
```

As you can see, the first match is produced immediately, while the remaining matches are processed as the result stream is traversed. This is most useful if you have to deal with bigger amounts of text. By processing the result stream in a piecemeal fashion, you can avoid keeping the entire result list in memory. For instance, compare the following:

```
> let s2 = fget $ fopen "system.pure" "r";
> stats -m
> #regexg id pat 0 s2 0;
7977
0.18s, 55847 cells
> #regexgs id pat 0 s2 0;
7977
0.12s, 20 cells
```

Regex Substitutions and Splitting

We can also perform substitutions on matches:

```
> regsub (sprintf "<%d:%s>") pat 0 s 0;
"1<1:var> <5:foo> 99 <12:BAR> $%&"
```

Or split a string using a delimiter pattern (this uses an egrep pattern):

```
> let delim = "[[:space:]]+";
> regsplit delim REG_EXTENDED s 0;
["1var","foo","99","BAR","$%&"]
> regsplit delim REG_EXTENDED "The quick brown fox" 0;
["The","quick","brown","fox"]
```

The regsplit operation also has a lazy variation:

```
> regsplits "[[:space:]]+" REG_EXTENDED "The quick brown fox" 0;
"The":##<thunk 0x7fb1b79775b0>
```

```
> last ans;  
"fox"
```

Empty Matches

Empty matches are permitted, too, subject to the constraint that at most one match is reported for each position (which also prevents looping). And of course an empty match will only be reported if nothing else matches. For instance:

```
> regexg id "" REG_EXTENDED "foo" 0;  
[(0,""),(1,""),(2,""),(3,"")]  
> regexg id "o*" REG_EXTENDED "foo" 0;  
[(0,""),(1,"oo"),(3,"")]  
> regexg id "o*" REG_EXTENDED "foo" 0;  
[(0,""),(1,"oo"),(2,"o"),(3,"")]
```

This also works when substituting or splitting:

```
> regsub (cst " ") "" REG_EXTENDED "some text" 0;  
" s o m e   t e x t "  
> regsub (cst " ") " ?" REG_EXTENDED "some text" 0;  
" s o m e   t e x t "  
> regsplit "" REG_EXTENDED "some text" 0;  
["","s","o","m","e"," ","t","e","x","t",""]  
> regsplit " ?" REG_EXTENDED "some text" 0;  
["","s","o","m","e","","t","e","x","t",""]
```

Submatches

Parenthesized subexpressions in a pattern yield corresponding submatch information, which is useful if we need to retrieve the text matched by a given subexpression. For instance, suppose we want to parse environment lines, such as those returned by the shell's set command. These can be dissected using the following regex:

```
> const env_pat = "^( [=]=+)(.*)$";  
> const env_flags = REG_EXTENDED or REG_NEWLINE;  
> regex env_pat env_flags "SHELL=/bin/sh" 0;  
1,0,"SHELL=/bin/sh",0,"SHELL",6,"/bin/sh"
```

Note that we again used an extended regex here, and we also added the REG_NEWLINE flag so that we properly deal with multiline input. The desired information is in the 4th and 6th element of the submatch info, we can retrieve that as follows:

```
> parse_env s = regexg (\info -> info!3 => info!5) env_pat env_flags s 0;  
> parse_env "SHELL=/bin/sh\nHOME=/home/bar\n";  
["SHELL"=>"/bin/sh","HOME"=>"/home/bar"]
```

We can get hold of the real process environment as follows:

```
> let env = parse_env $ fget $ popen "set" "r";
> #env;
109
> head env;
"BASH"=>"/usr/bin/sh"
```

Just for the fun of it, let's convert this to a record, providing easy random access to the environment variables:

```
> let env = record env;
> env!!["SHELL", "HOME"];
{"/bin/bash", "/home/ag"}
```

4.12 Additional POSIX Functions

The `posix` module provides some additional POSIX functions not available on all supported systems. (In particular, none of these functions are provided on MS Windows.) You can load this module in addition to the `system` module if you need the additional functionality. To use the operations of this module, add the following import declaration to your program:

```
using posix;
```

The following operations are provided. Please see the appropriate POSIX manual pages for a closer description of these functions.

fork Fork a new process.

getpid, getppid Get the process id of the current process and its parent process, respectively.

wait status, waitpid pid status options Wait for any child process, or the given one. The `status` argument must be a pointer to an `int` value, which is used to return the status of the child process.

kill pid sig Send the given signal to the given process.

raise sig Raise the given signal in the current process.

pause Sleep until a signal is caught.

4.13 Option Parsing

This is a quick-and-dirty replacement for the GNU `getopt` functions, ported from the `Q` library. To use the operations of this module, add the following import declaration to your program:

```
using getopt;
```

The following operation is provided:

getopt opts args Parse options as given by `opts` in the command line arguments `args`, return the parsed options along with a list of the remaining (non-option) command line arguments.

The `getopt` function takes two arguments: `opts`, a list of option descriptions in the format described below, and `args`, a list of strings containing the command line parameters to be parsed for options. The result is a pair (`opts_return`, `args_return`) where `opts_return` is a list of options and their values, and `args_return` is the list of remaining (non-option) arguments. Options are parsed using the rules of GNU `getopt(1)`. If an invalid option is encountered (unrecognized option, missing or extra argument, etc.), `getopt` throws the offending option string as an exception.

The `opts_return` value is a list of “hash pairs” `opt=>val` where `opt` is the (long) option name (as given by the `long_opt` field given in the `opts` argument, see below) and `val` is the corresponding value (`()` if none). Note that this format is ready to be passed to the [dict](#) or [hdict](#) function, cf. [Dictionaries](#), which makes it easy to retrieve option values or check for the presence of options. (As of Pure 0.41, you can also just convert the list to a record and employ the record functions to access the option data, cf. [Record Functions](#).)

The `opts` argument of `getopt` must be a list of triples (`long_opt`, `short_opt`, `flag`), where `long_opt` denotes the long option, `short_opt` the equivalent short option, and `flag` is one of the symbolic integer values `NOARG`, `OPTARG` and `REQARG` which specifies whether the option has no argument, an optional argument or a required argument, respectively. Either `long_opt` or `short_opt` should be a string value of the form “- -abc” or “-x”, respectively. Note that since the `long_opt` value is always used to denote the corresponding option in the `opts_return` list, you always have to specify a sensible value for that field. If no separate long option name is needed, you can specify the same value as in the `short_opt` field, or some other convenient value (e.g., an integer) which designates the option. Conversely, to indicate that an option has no short option equivalent, simply specify an empty option string for the `short_opt` field.

Examples:

```
> let opts = [("--help", "-h", NOARG),           // no argument
>             ("--version", "", NOARG),         // no short option
>             ("--filename", "-f", REQARG),      // required argument
>             ("--count", "-n", OPTARG)];        // optional argument
> getopt opts ["foo", "-h", "--filename", "bar", "-n0", "baz"];
["--help"=>(), "--filename"=>"bar", "--count"=>"0"], ["foo", "baz"]
> catch invalid_option $ getopt opts ["-h", "-v"];
invalid_option "-v"
> getopt opts [foo, "-h", bar];
["--help"=>()], [foo, bar]
```

As the last example shows, non-option arguments (as well as option values specified as separate arguments) can actually be any values which are just copied to the result lists as is.

pure-doc

Version 0.6, 2010-11-14

Albert Graef <Dr.Graef@t-online.de>

pure-doc is a simple utility for literate programming and documenting source code written in the Pure programming language. It is designed to be used with the excellent [docutils](#) tools and the gentle markup format supported by these, called [RST](#) a.k.a. “reStructuredText”, usually pronounced “rest”.

The basic idea is that you just comment your code as usual, but using RST markup instead of plain text. In addition, you can also designate literate programming fragments in your code, which will be translated to RST literal blocks automatically. You then run pure-doc on your source files to extract all marked up comments and the literate code blocks. The resulting RST source can then be processed with the docutils utilities like [rst2html.py](#) and [rst2latex.py](#) to create the documentation in a variety of formats.

1 Copying

Copyright (c) 2009-2010 by Albert Graef.

pure-doc is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

pure-doc is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

2 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-doc-0.6.tar.gz>.

Unpack and do the customary `&& sudo make install`. This only needs flex and a standards-compliant C++ compiler.

3 Usage

First, see the description of the [RST](#) format. RST is a very simple markup format, almost like plain text (in fact, you're looking at RST right now, this document is written in it!). You can learn enough of it to start marking up your source in about five minutes.

Second, you'll have to mark up your source comments. pure-doc recognizes comments in RST format by looking at the first non-empty line of the comment. A comment (either `/* ... */` or a contiguous sequence of `//` line comments) is assumed to contain RST format if the first non-empty line starts with `:`, `..` or `---`. Other comments are taken to be plain text and are ignored by pure-doc.

Notes:

- pure-doc makes no other assumption about the contents of marked up comments, so you can include whatever you want: titles, section headers, fields, admonitions, plain text, whatever. Just make sure that the comment starts with one of the special tokens listed above. (You can always put just `..` at the beginning of the comment to force it to be recognized, this will be treated as a comment by the docutils tools.)
- Also, pure-doc makes very few assumptions about the source; in fact, any source files with a C/C++-like comment and string syntax should work. So you could also use it to document your C/C++ programs, or even plain text files like this one, as long as they adhere to these standards.
- Indentation in extracted comments is preserved (assuming tabs = 8 spaces by default, you can change this with the `-t` option). This is important because indentation conveys document structure in RST.

For instance, here is a sample RST-formatted comment:

```
/* :Name: ``rand`` - compute random numbers
   :Synopsis: ``rand``
   :Description: Computes a (pseudo) random number. Takes no parameters.
   :Example: Here is how you can call ``rand`` in Pure:
   ::

       > extern int rand();
       > rand;
       1804289383

   :See Also: rand(3) */
```

This will be rendered as follows:

```
Name rand - compute random numbers
Synopsis rand
```

Description Computes a (pseudo) random number. Takes no parameters.

Example Here is how you can call `rand` in Pure:

```
> extern int rand();
> rand;
1804289383
```

See Also `rand(3)`

Finally, to extract the documentation you run `pure-doc` on your source files as follows:

```
pure-doc source-files ...
```

If no input files are specified then the source is read from standard input. Otherwise all input files are read and processed in the indicated order. The output is written to `stdout`, so that you can directly pipe it into one of the `docutils` programs:

```
pure-doc source-files ... | rst2html.py
```

If you prefer to write the output to a file, you can do that as follows:

```
pure-doc source-files ... > rst-file
```

`pure-doc` also understands the following options. These must come before any file arguments.

- h** Print a short help message.
- i** Automatic index creation (see below).
- s** Generate Sphinx-compatible output (see below).
- twidth** Set the tab width to the given number of spaces.

There are no other options. By its design `pure-doc` is just a plain simple “docstring scraping” utility with no formatting knowledge of its own. All actual formatting is handled by the `docutils` programs which offer plenty of options to change the appearance of the generated output; please refer to the [docutils](#) documentation for details.

Note that since Pure 0.46, all Pure documentation is usually formatted using [Sphinx](#), the RST formatter used by the Python project which provides cross-document indexing and referencing, and even more elaborate formatting options and prettier output than `docutils`. `pure-doc` versions since 0.6 support this by adding the `-s` option which makes its output compatible with Sphinx. (At present this option actually has any effect only when combined with the `-i` index generation option, see [Hyperlink Targets and Index Generation](#) below.)

4 Literate Programming

`pure-doc` also recognizes literate code delimited by comments which, besides the comment delimiters and whitespace, contain nothing but the special start and end “tags” `>>>` and `<<<`.

Code between these delimiters (including all comments) is extracted from the source and output as a RST literal code block.

For instance:

```
/* ..

    pure-doc supports literate programming, too. */

// >>>

// This is a literate comment.
/* .. This too! */

extern int rand();
rand;

// <<<
```

This will be rendered as follows:

```
pure-doc supports literate programming, too.

    // This is a literate comment.
    /* .. This too! */

    extern int rand();
    rand;
```

Try it now! You can scrape all the sample “documentation” from this file and format it as html, as follows:

```
pure-doc README | rst2html.py --no-doc-title --no-doc-info > test.html
```

5 Hyperlink Targets and Index Generation

Note: This feature is now largely obsolete as Pure uses Sphinx for formatting its documentation these days. Thus, as of version 0.6, the indexing feature must be enabled explicitly with the `-i` option.

When run with the `-i` option, pure-doc supplements the normal hyperlink target processing by the docutils tools, by recognizing explicit hyperlink targets of the form `.. _target:` and automatically creating raw html targets (`<a name=...`) for them. This works around the docutils name mangling (which is undesirable if you’re indexing, say, function names). It also resolves a quirk with some w3m versions which don’t pick up all id attributes in the docutils-generated html source.

In addition, you can also have pure-doc generate an index from all explicit targets. To these ends, just add the following special directive at the place where you want the index to appear:

```
.. makeindex::
```

The directive will be replaced with a list of references to all targets collected *up to that point*, sorted alphabetically. This also resets the list of collected targets, so that you can have multiple smaller indices in your document instead of one big one.

It goes without saying that this facility is rather simplistic, but it may be useful when you are working with plain docutils which does not provide its own indexing facility. Note, however, that docutils doesn't allow multiple explicit targets with the same name, so you should take that into consideration when devising your index terms.

Also note that in Sphinx compatibility mode (-s), pure-doc will generate the appropriate Sphinx markup for index entries (index:;) instead, and the makeindex:: directive will be ignored. You should then use Sphinx to generate the index.

Finally, if the -i option isn't specified, then all this special processing is disabled and the makeindex:: directive won't be recognized at all. This is the recommended way to process Pure documentation files which have been fully converted to Sphinx.

6 Generating and Installing Local Documentation

Note: This section only applies to 3rd party packages with their own bundled documentation which isn't part of the "official" Pure documentation. In this case it is possible to use docutils or some other RST formatting software to generate additional documentation files for use with the Pure interpreter. Please note that the method sketched out in this section doesn't provide full integration with the rest of Pure's documentation, but at least it makes it possible to read the local documentation in the interpreter.

If you're generating some library documentation for which you have to process a bigger collection of source files, then it is often convenient to have a few Makefile rules to automatize the process. To these ends, simply add rules similar to the following to your Makefile (the following assumes GNU make and that you're using docutils to format the documentation):

```
# The sources. Order matters here. The generated documentation will have the
# comments from each source file in the indicated order.
sources = foo.pure bar.pure
```

```
# The basename of the documentation files to be generated.
target = foo
```

```
.PHONY: html tex pdf
```

```
html: $(target).html
tex: $(target).tex
pdf: $(target).pdf
```

```
$(target).txt: $(sources)
    pure-doc $(sources) > $@
```

```
# This requires that you have docutils installed.
```

```
%.html: %.txt
    rst2html.py $< $@
```

```
%.tex: %.txt
    rst2latex.py $< $@
```

```
# This also requires that you have TeX installed.
```

```
%.pdf: %.tex
    pdflatex $<
    rm -f *.aux *.log *.out
```

```
clean:
    rm -f *.html *.tex *.pdf
```

You might want to add `-i` to the `pure-doc` command line if you want to enable the indexing feature described in the previous section. If you want to use some other RST formatting software, please check the corresponding documentation for information on how to format your documents and adjust the above rules for the `html`, `tex` and `pdf` targets accordingly.

Now you can just type `make html` to generate the documentation in `html` format, and `make tex` or `make pdf` to generate the other formats. The `clean` target removes the generated files.

Having generated the documentation files in `html` format, you can install them in the `docs` subdirectory of the Pure library directory to make it known to the Pure interpreter, so that you can read your documentation with the `help` command of the interpreter. (When doing this, name your documentation files in such a manner that you don't overwrite any of the Pure documentation files there.) The following Makefile rule automatizes this process. Add this to the Makefile in the previous section:

```
# Try to guess the installation prefix (this needs GNU make):
prefix = $(patsubst %/bin/pure,%, $(shell which pure 2>/dev/null))
ifeq ($(strip $(prefix)),)
# Fall back to /usr/local.
prefix = /usr/local
endif

libdir = $(prefix)/lib
docsdir = $(libdir)/pure/docs

install:
    test -d "$(DESTDIR)$(docsdir)" || mkdir -p "$(DESTDIR)$(docsdir)"
    cp $(target).html "$(DESTDIR)$(docsdir)"
```

After a `make install` your documentation should now end up in the appropriate place in the Pure library directory and you can read it in the Pure interpreter using a command like the following:


```
> help foo#
```

Note the hash character. This tells the `help` command that this is an auxiliary documentation file, rather than a search term to be looked up in the Pure documentation. You can also look up a specific section in your manual as follows:

```
> help foo#section-name
```

Please also refer to the Pure Manual for more information on how to use the interpreter's online help.

7 Formatting Tips

If you're generating documentation in pdf format via LaTeX, you might have to fiddle with the formatting to get results suitable for publication purposes. Newer versions of the `rsts2latex.py` program provide some options which let you adjust the formatting of various document elements. Here are the options that the author found particularly helpful when working with plain docutils:

- The table of contents that RST produces isn't all that useful in printed documentation, since it lacks page numbers. As a remedy, you can invoke `rsts2latex` with `--use-latex-toc` to have LaTeX handle the formatting of the table of contents, which looks much nicer.
- Similarly, `--use-latex-docinfo` can be used to tell `rsts2latex` that you want the title information (author and date) to be formatted the LaTeX way.
- If you need specific LaTeX document options, these can be specified with `--documentoptions`, e.g.: `--documentoptions="11pt"`.
- For more comprehensive formatting changes which require special LaTeX code and/or packages, you can use the `--stylesheet` option. E.g., `--stylesheet=preamble.tex` will cause a `preamble.tex` file with your own definitions to be included in the preamble of the generated document.
- To format literal code blocks using an alternative environment instead of the default `verbatim` environment, use the `--literal-block-env` option. E.g., `--literal-block-env=lstlisting` will use the highlighted code environment from the `listings` package. (Note that in this case you'll also need a preamble which loads the corresponding package.).

To learn more about this, please consult the `rsts2latex.py` documentation at the docutils website.

In addition, the `pure-doc` package contains a little GNU awk script called `fixdoc`, which attempts to improve the LaTeX output produced by older svn versions of `rsts2latex` in various ways. (This shouldn't be necessary for the latest `rsts2latex` versions.)

pure-ffi

Version 0.10, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

The libffi library provides a portable, high level programming interface to various calling conventions. This allows a programmer to call any function specified by a call interface description at run time. libffi should be present on most gcc-based systems, but it is also available as a standalone package at <http://sourceware.org/libffi/>.

This module provides an interface to libffi which enables you to call C functions from Pure and vice versa. It goes beyond Pure's built-in C interface in that it also handles C structs and makes Pure functions callable from C. Moreover, depending on the libffi implementation, it may also be possible to call foreign languages other than C.

1 Copying

Copyright (c) 2008, 2009 by Albert Graef.

pure-ffi is free software: you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

pure-ffi is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

2 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-ffi-0.10.tar.gz>.

Run `make` to compile the module and `make install` (as root) to install it in the Pure library directory. This requires GNU make, and of course you need to have Pure and libffi installed.

make tries to guess your Pure installation directory and platform-specific setup. If it gets this wrong, you can set some variables manually. In particular, `make install prefix=/usr` sets the installation prefix, and `make PIC=-fPIC` or some similar flag might be needed for compilation on 64 bit systems. Please see the Makefile for details.

NOTE: This module requires libffi 3.x (3.0.8 has been tested). Old libffi versions (2.x) do not appear to work (closures are broken). Patches are welcome.

3 Usage

The module exposes a simplified interface to libffi tailored to the Pure language. Call interfaces are described using the desired ABI, return type and tuple of argument types. The ABI is specified using one of the `FFI_*` constants defined by the module; for most purposes, `FFI_DEFAULT_ABI` is all that's needed. C types are specified using special descriptors `void_t`, `uint_t` etc., see `ffi.pure` for details. You can also get a list of these values using `show -g FFI_* *_t` after importing the `ffi` module.

The primary interface for calling C from Pure and vice versa is as follows:

- `fcall name abi rtype atypes`: Creates a Pure function from a C function with the given name, specified as a string. This makes the C function callable in Pure, no matter whether it is already declared as an extern or not. But note that if the function resides in a shared library, you still have to import that library using a Pure using declaration, see the Pure manual for details.
- `fclos fn abi rtype atypes`: Creates a pointer to a C function from the given Pure function `fn`. The resulting pointer can then be passed to other C functions expecting functions as arguments. This allows you to create C callbacks from Pure functions without writing a single line of C code. (This functionality might not be available on some platforms.)

Note that in difference to extern functions, arguments to functions created with libffi are always passed in uncurried form, as a Pure tuple. E.g.:

```
> using ffi;
> let fmod = fcall "fmod" FFI_DEFAULT_ABI double_t (double_t,double_t);
> fmod (5.3,0.7);
0.4
```

C structs are fully supported and are passed in a type-safe manner, see `ffi.pure` for details. Note that these are to be used for passing structs by value. (When passing a pointer to a struct, you must use `pointer_t` instead.) For instance:

```
> let complex_t = struct_t (double_t,double_t);
> let cexp = fcall "cexp" FFI_DEFAULT_ABI complex_t (complex_t);
> members (cexp (struct complex_t (0.0,1.0)));
0.54030230586814,0.841470984807897
```

See the examples folder in the sources for more examples.

4 TODO

The API isn't perfect yet. In particular, one might consider to implement type descriptors as structs instead of raw pointers, and support for typed pointers would be useful. Contributions and suggestions are welcome.

pure-gen: Pure interface generator

Version 0.9, November 14, 2010

Albert Gräf <Dr.Graef@t-online.de>

pure-gen is a C interface generator for the Pure language. It takes a C header file as input and generates a corresponding Pure module with the constant definitions and extern declarations needed to use the C module from Pure. pure-gen can also generate FFI interfaces rather than externs (using the pure-ffi module), and it can optionally create a C wrapper module which allows you to create interfaces to pretty much any code which can be called via C.

1 Synopsis

```
pure-gen [options ...] input-file
```

2 Options

2.1 General Options

- h, --help** Print a brief help message and exit.
- V, --version** Print version number and exit.
- e, --echo** Echo preprocessor lines. Prints all processed `#defines`, useful for debugging purposes.
- v, --verbose** Show parameters and progress information. Gives useful information about the conversion process.
- w[level], --warnings[=level]** Display warnings, `level = 0` (disable most warnings), `1` (default, shows important warnings only) or `2` (lots of additional warnings useful for debugging purposes).

2.2 Preprocessor Options

- I path, --include path** Add include path. Passed to the C preprocessor.
- D name[=value], --define name[=value]** Define symbol. Passed to the C preprocessor.
- U name, --undefine name** Undefine symbol. Passed to the C preprocessor.
- C option, --cpp option** Pass through other preprocessor options and arguments.

2.3 Generator Options

- f iface, --interface iface** Interface type (extern, c, ffi or c-ffi). Default is extern. The extern and c types generate Pure extern declarations, which is what you want in most cases. ffi and c-ffi employ Pure's libffi interface instead. The c and c-ffi types cause an additional C wrapper module to be created (see [Generating C Code](#)). These can also be combined with the -auto suffix which creates C wrappers only when needed to get C struct arguments and returns working, see [Dealing with C Structs](#) for details.
- l lib, --lib-name lib** Add dynamic library module to be imported in the Pure output file. Default is -l c-file (the filename specified with -c, see below, without filename extension) if one of the -fc options was specified, none otherwise.
- m name, --namespace name** Module namespace in which symbols should be declared.
- p prefix, --prefix prefix** Module name prefix to be removed from C symbols.
- P prefix, --wrap prefix** Prefix to be prepended to C wrapper symbols (-fc and friends). Default is Pure_.
- a, --all** Include "hidden" symbols in the output. Built-in preprocessor symbols and symbols starting with an underscore are excluded unless this option is specified.
- s pattern, --select pattern** Selection of C symbols to be included in the output. pattern takes the form [glob-patterns::][regex-pattern], designating a comma separated list of glob patterns matching the source filenames, and an extended regular expression matching the symbols to be processed. See glob(7) and regex(7). The default pattern is empty which matches all symbols in all source modules.
- x pattern, --exclude pattern** Like -s, but *excludes* all matching C symbols from the selection.
- t file, --template file** Specify a C template file to be used with C wrapper generation (-fc). See [Generating C Code](#) for details.
- T file, --alt-template file** Specify an alternate C template file to be used with C wrapper generation (-fc). See [Generating C Code](#) for details.

2.4 Output Options

- n, --dry-run** Only parse without generating any output.
- N, --noclobber** Append output to existing files.
- o file, --output file** Pure output (.pure) filename. Default is input-file with new extension .pure.
- c file, --c-output file** C wrapper (.c) filename (-fc). Default is input-file with new extension .c.

3 Description

pure-gen generates Pure bindings for C functions from a C header file. For instance, the command

```
pure-gen foo.h
```

creates a Pure module foo.pure with extern declarations for the constants (#defines and enums) and C routines declared in the given C header file and (recursively) its includes.

pure-gen only accepts a single header file on the command line. If you need to parse more than one header in a single run, you can just create a dummy header with all the necessary #includes in it and pass that to pure-gen instead.

When invoked with the -n option, pure-gen performs a dry run in which it only parses the input without actually generating any output files. This is useful for checking the input (possibly in combination with the -e, -v and/or -w options) before generating output. A particularly useful example is

```
pure-gen -ne foo.h |
| awk ' $1=="#" && $2~/^[0-9]+$ / && $3!~/^"<.*>"$ / { print $3 }' |
| sort | uniq
```

which prints on standard output all headers which are included in the source. This helps to decide which headers you want to be included in the output, so that you can set up a corresponding filter patterns (-s and -x options, see below).

The -I, -D and -U options are simply passed to the C preprocessor, as well as any other option or argument escaped with the -C flag. This is handy if you need to define additional preprocessor symbols, add directories to the include search path, etc., see cpp(1) for details.

There are some other options which affect the generated output. In particular, -f c generates a C wrapper module along with the Pure module (see [Generating C Code](#) below), and -f ffi generates a wrapper using Pure's ffi module. Moreover, -l libfoo generates a using "lib:libfoo" declaration in the Pure source, for modules which require a shared library to be loaded. Any number of -l options can be specified.

Other options for more advanced uses are explained in the following sections.

4 Filtering

Note that pure-gen always parses the given header file as well as *all* its includes. If the header file includes system headers, by default you will get those declarations as well. This is often undesirable. As a remedy, pure-gen normally excludes built-in #defines of the C preprocessor, as well as identifiers with a leading underscore (which are often found in system headers) from processing. You can use the -a option to disable this, so that all these symbols are included as well.

In addition, the -s and -x options enable you to filter C symbols using the source filename and the symbol as search criteria. For instance, to just generate code for a single header foo.h and none of the other headers included in foo.h, you can invoke pure-gen as follows:

```
pure-gen -s foo.h:: foo.h
```

Note that even in this case all included headers will be parsed so that #defined constants and enum values can be resolved, but the generated output will only contain definitions and declarations from the given header file.

In general, the -s option takes an argument of the form glob-patterns::regex-pattern denoting a comma-separated list of glob patterns to be matched against the source filename in which the symbol resides, and an extended regex to be matched against the symbol itself. The glob-patterns:: part can also be omitted in which case it defaults to :: which matches any source file. The regex can also be empty, in which case it matches any symbol. The generated output will contain only the constant and function symbols matching the given regex, from source files matching any of the the glob patterns. Thus, for instance, the option -s foo.h,bar.h:^(foo|bar)_ pulls all symbols prefixed with either foo_ or bar_ from the files foo.h and bar.h in the current directory.

Instead of :: you can also use a single semicolon ; to separate glob and regex pattern. This is mainly for Windows compatibility, where the msys shell sometimes eats the colons or changes them to ;.

The -x option works exactly the same, but *excludes* all matching symbols from the selection. Thus, e.g., the option -x ^bar_ causes all symbols with the prefix bar_ to *not* be included in the output module.

Processing of glob patterns is performed using the customary rules for filename matching, see glob(7) for details. Note that some include files may be specified using a full pathname. This is the case, in particular, for system includes such as #include <stdio.h>, which are resolved by the C preprocessor employing a search of the system include directories (as well as any directories named with the -I option).

Since the * and ? wildcards never match the pathname separator /, you have to specify the path in the glob patterns in such cases. Thus, e.g., if the foo.h file actually lives in either /usr/include or /usr/local/include, then it must be matched using a pattern like /usr/include/*.h,/usr/local/include/*.h::. Just foo.h:: will not work in this case. On the other hand, if you have set up your C sources in some local directory then specifying a relative pathname is ok.

5 Name Mangling

The `-s` option is often used in conjunction with the `-p` option, which lets you specify a “module name prefix” which should be stripped off from C symbols. Case is insignificant and a trailing underscore will be removed as well, so `-p foo` turns `fooBar` into `Bar` and `F00_BAR` into `BAR`. Moreover, the `-m` option allows you to specify the name of a Pure namespace in which the resulting constants and functions are to be declared. So, for instance, `-s "^(foo|F00)" -p foo -m foo` will select all symbols starting with the `foo` or `F00` prefix, stripping the prefix from the selected symbols and finally adding a `foo::` namespace qualifier to them instead.

6 Generating C Code

As already mentioned, `pure-gen` can be invoked with the `-fc` or `-fc-ffi` option to create a C wrapper module along with the Pure module it generates. There are various situations in which this is preferable, e.g.:

- You are about to create a new module for which you want to generate some boilerplate code.
- The C routines to be wrapped aren’t available in a shared library, but in some other form (e.g., object file or static library).
- You need to inject some custom code into the wrapper functions (e.g., to implement custom argument preprocessing or lazy dynamic loading of functions from a shared library).
- The C routines can’t be called directly through Pure externs.

The latter case might arise, e.g., if the module uses non-C linkage or calling conventions, or if some of the operations to be wrapped are actually implemented as C macros. (Note that in order to wrap macros as functions you’ll have to create a staged header which declares the macros as C functions, so that they are wrapped in the C module. `pure-gen` doesn’t do this automatically.)

Another important case is that some of the C routines pass C structs by value or return them as results. This is discussed in more detail in the following section.

For instance, let’s say that we want to generate a wrapper `foo.c` from the `foo.h` header file whose operations are implemented in some library `libfoo.a` or `libfoo.so`. A command like the following generates both the C wrapper and the corresponding Pure module:

```
pure-gen -fc foo.h
```

This creates `foo.pure` and `foo.c`, with an import clause for `"lib:foo"` at the beginning of the Pure module. (You can also change the name of the Pure and C output files using the `-o` and `-c` options, respectively.)

The generated wrapper is just an ordinary C file which should be compiled to a shared object (dll on Windows) as usual. E.g., using gcc on Linux:

```
gcc -shared -o foo.so foo.c -lfoo
```

That's all. You should now be able to use the foo module by just putting the declaration using `foo`; into your programs. The same approach also works with the ffi interface if you replace the `-fc` option with `-fc-ffi`.

You can also adjust the C wrapper code to some extent by providing your own template file, which has the following format:

```
/* frontmatter here */
#include %h
%%

/* wrapper here */
%r %w(%p)
{
    return %n(%a);
}
```

Note that the code up to the symbol `%%` on a line by itself denotes “frontmatter” which gets inserted at the beginning of the C file. (The frontmatter section can also be empty or missing altogether if you don't need it, but usually it will contain at least an `#include` for the input header file.)

The rest of the template is the code for each wrapper function. Substitutions of various syntactical fragments of the function definition is performed using the following placeholders:

`%h` input header file

`%r` return type of the function

`%w` the name of the wrapper function

`%p` declaration of the formal parameters of the wrapper function

`%n` the real function name (i.e., the name of the target C function to be called)

`%a` the arguments of the function call (formal parameters with types stripped off)

`%%` escapes a literal `%`

A default template is provided if you don't specify one (which looks pretty much like the template above, minus the comments). A custom template is specified with the `-t` option. (There's also a `-T` option to specify an “alternate” template for dealing with routines returning struct values, see [Dealing with C Structs](#).)

For instance, suppose that we place the sample template above into a file `foo.templ` and invoke `pure-gen` on the `foo.h` header file as follows:

```
pure-gen -fc -t foo.templ foo.h
```

Then in `foo.c` you'd get C output code like the following:

```
/* frontmatter here */
#include "foo.h"

/* wrapper here */
void Pure_foo(int arg0, void* arg1)
{
    return foo(arg0, arg1);
}

/* wrapper here */
int Pure_bar(int arg0)
{
    return foo(arg0);
}
```

As indicated, the wrapper function names are usually stropped with the `Pure_` prefix. You can change this with the `-P` option.

This also works great to create boilerplate code for new modules. For this purpose the following template will do the trick:

```
/* Add #includes etc. here. */
%%

%r %n(%p)
{
    /* Enter code of %n here. */
}
```

7 Dealing with C Structs

Modern C compilers allow you to pass C structs by value or return them as results from a C function. This represents a problem, because Pure doesn't provide any support for that in its extern declarations. Even Pure's libffi interface only has limited support for C structs (no unions, no bit fields), and at present pure-gen itself does not keep track of the internal structure of C structs either.

Hence pure-gen will bark if you try to wrap an operation which passes or returns a C struct, printing a warning message like the following which indicates that the given function could not be wrapped:

```
Warning: foo: struct argument or return type, try -fc-auto
```

What Pure *does* know is how to pass and return *pointers* to C structs in its C interface. This makes it possible to deal with struct arguments and return values in the C wrapper. To make this work, you need to create a C wrapper module as explained in the previous section. However, as C wrappers are only needed for functions which actually have struct arguments or return values, you can also use the `-fc-auto` option (or `-fc-ffi-auto` if you prefer the

ffi interface) to only generate the C wrapper when required. This saves the overhead of an extra function call if it's not actually needed.

Struct arguments in the original C function then become struct pointers in the wrapper function. E.g., if the function is declared in the header as follows:

```
typedef struct { double x, y; } point;
extern double foo(point p);
```

Then the generated wrapper code becomes:

```
double Pure_foo(point* arg0)
{
    return foo(*arg0);
}
```

Which is declared in the Pure interface as:

```
extern double Pure_foo(point*) = foo;
```

Struct return values are handled by returning a pointer to a static variable holding the return value. E.g.,

```
extern point bar(double x, double y);
```

becomes:

```
point* Pure_bar(double arg0, double arg1)
{
    static point ret;
    ret = bar(arg0, arg1); return &ret;
}
```

Which is declared in the Pure interface as:

```
extern point* Pure_bar(double, double) = bar;
```

(Note that the generated code in this case comes from an alternate template. It's possible to configure the alternate template just like the normal one, using the `-T` option instead of `-t`. See the [Generating C Code](#) section above for details about code templates.)

In a Pure script you can now call `foo` and `bar` as:

```
> foo (bar 0.0 1.0);
```

Note, however, that the pointer returned by `bar` points to static storage which will be overwritten each time you invoke the `bar` function. Thus in the following example *both* `u` and `v` will point to the same point struct, namely that defined by the latter call to `bar`:

```
> let u = bar 1.0 0.0; let v = bar 0.0 1.0;
```

Which most likely is *not* what you want. To avoid this, you'll have to take dynamic copies of returned structs. It's possible to do this manually by fiddling around with `malloc` and

memcpy, but the most convenient way is to employ the struct functions provided by Pure's ffi module:

```
> using ffi;
> let point_t = struct_t (double_t, double_t);
> let u = copy_struct point_t (bar 1.0 0.0);
> let v = copy_struct point_t (bar 0.0 1.0);
```

Now `u` and `v` point to different, malloc'd structs which even take care of freeing themselves when they are no longer needed. Moreover, the ffi module also allows you to access the members of the structs in a direct fashion. Please refer to the pure-ffi documentation for further details.

8 Notes

pure-gen currently requires gcc (-E) as the C preprocessor. It also needs a version of gcc which understands the `-fdirectives-only` option, which means gcc 4.3 or later. It will run with older versions of gcc, but then you'll get an error message from gcc indicating that it doesn't understand the `-fdirectives-only` option. pure-gen then won't be able to extract any `#defined` constants from the header files.

pure-gen itself is written in Pure, but uses a C parser implemented in Haskell, based on the Language.C library written by Manuel Chakravarty and others.

pure-gen can only generate C bindings at this time. Other languages may have their own calling conventions which make it hard or even impossible to call them directly through Pure's extern interface. However, if your C compiler knows how to call the other language, then it may be possible to interface to modules written in that language by faking a C header for the module and generating a C wrapper with a custom code template, as described in [Generating C Code](#). In principle, this approach should even work with behemoths like C++, although it might be easier to use third-party tools like SWIG for that purpose.

In difference to SWIG and similar tools, pure-gen doesn't require you to write any special "interface files", is controlled entirely by command line options, and the amount of marshalling overhead in C wrappers is negligible. This is possible since pure-gen targets only the Pure-C interface and Pure has good support for interfacing to C built into the language already.

pure-gen usually works pretty well if the processed header files are written in a fairly clean fashion. Nevertheless, some libraries defy fully automatic wrapper generation and may thus require staged headers and/or manual editing of the generated output to get a nice wrapper module.

In complex cases it may also be necessary to assemble the output of several runs of pure-gen for different combinations of header files, symbol selections and/or namespace/prefix settings. In such a situation it is usually possible to just concatenate the various output files produced by pure-gen to consolidate them into a single wrapper module. To make this easier, pure-gen provides the `-N` a.k.a. `--noclobber` option which appends the output to existing files instead of overwriting them. See the example below.

9 Example

For the sake of a substantial, real-world example, here is how you can wrap the entire GNU Scientific Library in a single Pure module `mygsl.pure`, with the accompanying C module in `mygsl.c`:

```
rm -f mygsl.pure mygsl.c
DEFS=-DGSL_DISABLE_DEPRECATED
for x in /usr/include/gsl/gsl*.h; do
    pure-gen $DEFS -N -fc-auto -s "$x:." $x -o mygsl.pure -c mygsl.c
done
```

The C module can then be compiled with:

```
gcc $DEFS -shared -o mygsl.so mygsl.c
```

Note that the `GSL_DISABLE_DEPRECATED` symbol must be defined here to avoid some botches with constants being defined in incompatible ways in different GSL headers. Also, some GSL versions have broken headers lacking some system includes which causes hiccups in pure-gen's C parser. Fixing those errors or working around them through some appropriate cpp options should be a piece of cake, though. ;-)

10 License

BSD-like. See the accompanying COPYING file for details.

11 Authors

Scott E. Dillard (University of California at Davis), Albert Graef (Johannes Gutenberg University at Mainz, Germany).

12 See Also

Language.C A C parser written in Haskell by Manuel Chakravarty et al, <http://www.sivity.net/projects/language.c>.

SWIG The Simplified Wrapper and Interface Generator, <http://www.swig.org>.

pure-readline

Version 0.1, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

Get the latest source from <http://pure-lang.googlecode.com/files/pure-readline-0.1.tar.gz>.

This is a trivial wrapper around GNU readline, which gives Pure scripts access to the `readline` and `add_history` functions. The wrapper can also be used with the BSD editline a.k.a. libedit library, a readline replacement licensed under the 3-clause BSD license. You can find these at:

- GNU readline: <http://tiswww.tis.case.edu/~chet/readline/rltop.html>
- BSD editline/libedit: <http://www.thrysoee.dk/editline>

We recommend GNU readline because it's easier to use and has full UTF-8 support, but in some situations BSD editline/libedit may be preferable for license reasons or because it's what the operating system provides. Note that in either case Pure programs using this module are subject to the license terms of the library that you use (GPLv3+ in case of GNU readline, BSD license in the case of BSD editline/libedit).

Normally, you should choose the same library that you use with the Pure interpreter, to avoid having two different versions of the library linked into your program. (This doesn't matter if you only use this module with batch-compiled scripts, though, since the Pure runtime doesn't depend on readline in any way.) By default, the module will be built with GNU readline. To select editline/libedit instead, you only have to uncomment a line at the beginning of the Makefile. Also, you might want to check the beginning of `readline.c` for the proper location of the corresponding header files.

The module provides two functions:

readline prompt Read a line of input from the user, with prompting and command line editing. Returns the input line (with the trailing newline removed), or `NULL` when reaching end of file.

add_history line Adds the given line (a string) to the command history.

Example:

```
> readline "input> ";
input> Hello, world!
"Hello, world!"
```

```
> add_history ans;
()
> readline "input> ";
input> <EOF>
#<pointer 0>
```

pure-sockets: Pure Sockets Interface

Version 0.1, November 14, 2010

Albert Gräf <Dr.Graef@t-online.de>

This is an interface to the Berkeley socket functions. It provides most of the core functionality, so you can create sockets for both stream and datagram based protocols and use these to transmit messages. Unix-style file sockets are also available if the host system supports them.

1 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-sockets-0.1.tar.gz>.

This package uses Pure's new bitcode loader to ease porting, so you'll need Pure 0.45 or later and an LLVM-based C compiler installed. The Makefile is set up for clang (<http://clang.llvm.org/>) by default, because it's easier to install, but you can also use llvm-gcc if you set the CC make variable accordingly.

For clang compilation, do the usual `make` && `sudo make install` and you should be set. To use llvm-gcc, replace `make` with `make CC=llvm-gcc`. You can also set custom compilation options with the CFLAGS variable, e.g.: `make CFLAGS=-O3`.

To uninstall the module, use `sudo make uninstall`. There are a number of other targets (mostly for maintainers), please see the Makefile for details.

2 Usage

To use the operations of this module, put the following in your Pure script:

```
using sockets;
```

NOTE: On some systems you may also have to pull in a special system library to get the socket functions. E.g.:

```
using "lib:wsock32", sockets;
```

With the `sockets` module loaded, all the standard socket functions are available and work pretty much like in C. The only real difference is that, for convenience, functions taking socket addresses as parameters (`struct_sockaddr*` pointers in Pure), are called without the `addrlen` parameter; the size of the socket address structure will be inferred automatically and passed to the underlying C functions. Also, there are some convenience functions which act as wrappers around `getaddrinfo` and `getnameinfo` to create socket addresses from symbolic information (hostname or ip, port names or numbers) and return information about existing address pointers, see [Creating and Inspecting Socket Addresses](#) below.

Below is a list of the provided functions. Please see the corresponding manual pages for details, and check the Pure scripts in the examples subdirectory for some examples.

2.1 Creating and Inspecting Socket Addresses

These functions are Pure-specific. The created socket addresses are malloc'ed and free themselves automatically when garbage-collected.

```
struct_sockaddr *sockaddr();
```

Create a pointer to an empty socket address suitable to hold the socket address result of routines like `accept`, `getsockname`, `recvfrom`, etc. which return a socket address.

```
struct_sockaddr *sockaddr([int family,] char *path);
```

Create a local (a.k.a. file) socket for the given pathname. The `family` parameter, if specified, must be `AF_UNIX` here. Please note that `AF_UNIX` is not supported on all platforms. You can check for this by testing the `HAVE_AF_UNIX` constant, which is a truth value specifying whether `AF_UNIX` is available on your system.

```
struct_sockaddr *sockaddr([int family,] char *host, char *port);
```

```
struct_sockaddr *sockaddr([int family,] char *host, int port);
```

This uses `getaddrinfo` to retrieve an `AF_INET` or `AF_INET6` address for the given hostname (or numeric IP address in string form) and port (specified either as an int or a string). If `family` is omitted, it defaults to `AF_UNSPEC` which matches both `AF_INET` and `AF_INET6` addresses.

```
struct_sockaddr *sockaddrs([int family,] char *host, char *port);
```

```
struct_sockaddr *sockaddrs([int family,] char *host, int port);
```

This works like `sockaddr` above, but returns a list with *all* matching addresses.

```
int sockaddr_family(struct_sockaddr *addr);
```

Returns the address family of the given address.

```
char *sockaddr_path(struct_sockaddr *addr);
```

Returns the pathname for `AF_UNIX` addresses.

```
char *sockaddr_hostname(struct_sockaddr *addr);
```

Returns the hostname if available, the IP address otherwise.

```
char *sockaddr_ip(struct_sockaddr *addr);
```

Returns the IP address.

```
char *sockaddr_service(struct_sockaddr *addr);
```

Returns the service (a.k.a. port) name.

```
int sockaddr_port(struct_sockaddr *addr);
```

Returns the port number.

```
expr *sockaddr_info(struct_sockaddr *addr);
```

Returns a readable description of a socket address, as a (family, hostname, port) tuple. You should be able to pass this into `sockaddr` again to get the original address.

2.2 Creating and Closing Sockets

```
int socket(int domain, int type, int protocol);
```

Creates a socket for the given protocol family (`AF_UNIX`, `AF_INET` or `AF_INET6`), socket type (`SOCK_STREAM`, `SOCK_DGRAM`, etc.) and protocol. Note that on Linux we also support the `SOCK_NONBLOCK` (non-blocking) and `SOCK_CLOEXEC` (close-on-exec) flags which can be or'ed with the socket type to get sockets with the corresponding features. The protocol number is usually 0, denoting the default protocol, but it can also be any of the prescribed `IPPROTO` constants (a few common ones are predefined by this module, try `show -g IPPROTO_*` for a list of those).

```
int socketpair(int domain, int type, int protocol, int *sv);
```

Create a pair of sockets. The descriptors are returned in the integer vector passed in the last argument.

```
int shutdown(int fd, int how);
```

Perform shutdown on a socket. The second argument should be one of `SHUT_RD`, `SHUT_WR` and `SHUT_RDWR`.

```
closesocket(int fd);
```

This is provided for Windows compatibility. On POSIX systems this is just `close`.

2.3 Establishing Connections

```
int accept(int sockfd, struct_sockaddr *addr);
```

```
int bind(int sockfd, struct_sockaddr *addr);
```

```
int connect(int sockfd, struct_sockaddr *addr);
```

```
int listen(int sockfd, int backlog);
```

2.4 Socket I/O

```
size_t recv(int fd, void *buf, size_t len, int flags);
size_t send(int fd, void *buf, size_t len, int flags);
size_t recvfrom(int fd, void *buf, size_t len, int flags, struct_sockaddr
*addr);
size_t sendto(int fd, void *buf, size_t len, int flags, struct_sockaddr
*addr);
```

The usual send/recv flags specified by POSIX (MSG_EOR, MSG_OOB, MSG_PEEK, MSG_WAITALL) are provided. On Linux we also support MSG_DONTWAIT. Note that on POSIX systems you can also just fdopen the socket descriptor and use the standard file I/O operations from the system module instead.

2.5 Socket Information

```
int getsockname(int fd, struct_sockaddr *addr);
int getpeername(int fd, struct_sockaddr *addr);
int getsockopt(int fd, int level, int name, void *val, int *len);
int setsockopt(int fd, int level, int name, void *val, int len);
```

For getsockopt and setsockopt, currently only the SOL_SOCKET level is defined (level argument) along with the available POSIX socket options (name argument). Try `show -g SO_*` to get a list of those. Also note that for most socket level options the val argument is actually an int*, so you can pass a Pure int vector (with len = sizeof_int) for that parameter.

3 Example

Here is a fairly minimal example using Unix stream sockets. To keep things simple, this does no error checking whatsoever and just keeps sending strings back and forth. More elaborate examples can be found in the examples directory in the sources.

```
using sockets, system;

const path = "server_socket";
extern int unlink(char *name);

server = loop with
  loop = loop if ~null s && ~response fp s when
    // Connect to a client.
    cfd = accept fd $ sockaddr ();
```

```
    // Open the client socket as a FILE* and read a request.
    fp = fdopen cfd "r+"; s = fgets fp;
end;
loop = puts "server is exiting" $$ closesocket fd $$
      unlink path $$ () otherwise;
response fp s::string = s=="quit\n" when
  // Process the request. (Here we just print the received
  // message and echo it back to the client.)
  printf "server> %s" s;
  fputs s fp;
end;
end when
  // Create the server socket and start listening.
  unlink path;
  fd = socket AF_UNIX SOCK_STREAM 0;
  bind fd (sockaddr path); listen fd 5;
  printf "server listening at '%s'\n" path;
end;

client = loop with
  // Keep reading requests from stdin.
  loop = loop if ~null s && ~request s when
    fputs "client> " stdout; s = fgets stdin;
  end;
  loop = puts "client is exiting" $$ () otherwise;
  request s::string = s=="quit\n" when
    fd = socket AF_UNIX SOCK_STREAM 0;
    connect fd (sockaddr path);
    // Send the request to the server.
    fp = fdopen fd "r+"; fputs s fp;
    // Get the reply.
    s = fgets fp;
  end;
end;
```

To use this example, run the server function in one instance of the Pure interpreter and the client function in another. Enter a line when the client prompts you for input; it will be printed by the server. Behind the scenes, the server also sends the line back to the client. After receiving the reply, the client prompts for the next input line. Entering end-of-file at the client prompt terminates the client but keeps the server running, so that you can start another client and reconnect to the server. Entering just quit in the client terminates both server and client. Here is how a typical interaction may look like:

```
> client;
client> 1+1
client> foo bar
client> quit
client is exiting
()

> server;
server listening at 'server_socket'
```

```
server> 1+1
server> foo bar
server> quit
server is exiting
()
```

Note that while the server processes requests sequentially, it accepts connections from a new client after each request, so that you can run as many clients as you like.

Gnumeric/Pure: A Pure Plugin for Gnumeric

Version 0.8, November 14, 2010

Albert Gräf <Dr.Graef@t-online.de>

Gnumeric/Pure is a [Gnumeric](#) extension which lets you use [Pure](#) functions in Gnumeric, the Gnome spreadsheet. It offers better execution speed than the existing Perl and Python plugins, and provides some powerful features not found in other Gnumeric scripting plugins, such as asynchronous data sources created from Pure streams and OpenGL rendering in Gnumeric frame widgets via Pure's OpenGL module.

1 Introduction

This package provides a [Gnumeric](#) extension which gives you access to the [Pure](#) programming language in Gnumeric. It works pretty much like the Perl and Python plugin loaders which are distributed with Gnumeric, but Gnumeric/Pure offers some powerful features which aren't found in other Gnumeric scripting plugins:

- Pure is a functional programming language which fits the computational model of spreadsheet programs very well.
- Pure is based on term rewriting and thus enables you to do symbolic computations in addition to the usual numeric calculations.
- Pure has a built-in [MATLAB/Octave](#)-like matrix data structure which makes it easy to deal with cell ranges in a spreadsheet in an efficient manner.
- Pure also provides a bridge to [Octave](#) so that you can call arbitrary Octave functions using this extension.
- Gnumeric/Pure offers support for rendering [OpenGL](#) scenes in Gnumeric frame widgets, via Pure's own OpenGL interface.
- Pure also has built-in support for lazy data structures and thus allows you to handle potentially infinite amounts of data such as the list of all prime numbers. Gnumeric/Pure lets you turn such lazy values into asynchronous data sources computed in the background, which update the spreadsheet automatically as results become available.

- Last but not least, Pure is compiled to native code on the fly. This means that, while startup times are a bit longer due to Pure's JIT compiler kicking in (you'll notice this if you open a spreadsheet with Pure functions), the resulting compiled code then typically executes *much* faster than equivalent interpreted Perl and Python code.

Adding Pure functions to Gnumeric is quite easy. Once the plugin is installed and enabled, you can simply start adding Pure functions to the provided `pure_func.pure` script, or you can create your own Pure plugin folders. Gnumeric/Pure also provides a programming interface for use in Pure which gives you access to various advanced features, such as modifying entire ranges of cells with one Pure call, calling Gnumeric functions from Pure, and setting up asynchronous data sources and OpenGL frames. This is all explained in detail below.

2 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/gnumeric-pure-0.8.tar.gz>.

Obviously, you need to have both Pure and Gnumeric installed. Pure 0.36 and Gnumeric 1.9.13 are known to work. We recommend Gnumeric 1.9.14 or later since it has improved support for GUI widgets. (Older Gnumeric versions probably work as well if you're willing to fiddle with the Makefile and/or the sources. See the beginning of the Makefile for related information.)

As shipped, the Makefile is set up to build Gnumeric/Pure with OpenGL support, which requires that you have the OpenGL libraries as well as `GtkGLExt` (the Gtk OpenGL extension) installed. These should be readily available on most systems, but you can also disable this feature by invoking `make` as `make GLDEPS=`.

Run `make` to compile the software. You might have to adjust the settings at the beginning of the Makefile to make this work. If you're lucky and the compile goes through, you should now have a `pure_loader.so` file in the `pure-loader` subdirectory. You can install the plugin and related stuff with `sudo make install` in the global Gnumeric plugin directory, or if you prefer to install it into your personal plugin directory then run `make install-local` instead. (We recommend the latter since it lets you adjust `pure_func.pure` for your purposes more easily.) Optionally, you might also want to copy the `gnumeric-pure.html` file to your Pure library directory so that you can read it with the `help` command of the Pure interpreter or in Emacs Pure mode.

Typically, `make install` and `make install-local` will install the plugins into the following directories by default:

- System-wide installations go into `/usr/local/gnumeric/1.9.13/plugins` or similar, depending on where Gnumeric is installed.
- User-specific installations go into `~/.gnumeric/1.9.13/plugins`.

The Makefile tries to guess the installation path and version number of Gnumeric on its own. If it guesses wrong, you can change these using the Makefile variables `prefix` and

gnmversion, respectively. For instance:

```
$ make prefix=/usr gnmversion=1.9.13
```

If `make install` doesn't work for some reason, you can also just copy the `pure-func`, `pure-glfw` and `pure-loader` directories manually to your Gnumeric plugin directory.

3 Setup

Once Gnumeric/Pure has been properly installed, you should see it in Gnumeric's Tools/Plug-ins dialog. There are actually two main entries, one labelled "Pure functions" and the other one labelled "Pure plugin loader". You need to enable both before you can start using Pure functions in your Gnumeric spreadsheets. There's also a third entry labelled "Pure OpenGL functions" which you might want to enable if you want to try the OpenGL capabilities (this will only work if you built Gnumeric/Pure with OpenGL support and have Pure's OpenGL module installed; see [OpenGL Interface](#) for details).

Gnumeric doesn't provide much in the way of GUI customization options right now, but at least it's possible for plugins to install and configure additional menu and toolbar options. Gnumeric/Pure adds three additional options to the Tools menu which allow you to stop asynchronous data sources, reload Pure scripts and edit them. After installation, the definitions of these items can be found in the `pure-loader/pure-ui.xml` file in your Gnumeric plugin directory. Have a look at this file and edit it as desired. E.g., if you want to put the Pure-related options into a submenu and enable toolbar buttons for these options, then your `pure-ui.xml` file should look as follows:

```
<ui>
  <menubar>
    <menu name="Tools" action="MenuTools">
      <separator/>
      <menu name="Pure" action="PureMenu">
        <menuitem action="PureStop"/>
        <menuitem action="PureReload"/>
        <menuitem action="PureEdit"/>
      </menu>
    </menu>
  </menubar>
  <toolbar name="StandardToolbar">
    <separator/>
    <toolitem action="PureStop"/>
    <toolitem action="PureReload"/>
    <toolitem action="PureEdit"/>
  </toolbar>
</ui>
```

4 Basic Usage

With Pure/Gnumeric installed and enabled, you should be ready to join the fun now. Start up Gnumeric, click on a cell and invoke the “f(x)” dialog. The Pure functions available for use are shown in the “Pure” category. E.g., click on `pure_hello`. Now the Pure interpreter will be loaded and the function description displayed. Click “Insert” and then “Ok”. You should now be able to read the friendly greeting returned by the `pure_hello` function.

Of course, you can also enter the function call directly as a formula into a cell as usual. Click on a cell, then enter the following:

```
=pure_hello(getenv("USER"))
```

The greeting should now be displayed with your login name in it.

Play around a bit with the other Pure functions. These functions are nothing special; they are just ordinary Pure functions which are defined by the `pure_func.pure` script in the `pure-func` subdirectory of your Gnumeric plugin directory. You can have a look at them by invoking the “Edit Pure Script” option which gets added to the Tools/Pure menu once the Pure plugin loader is enabled. (This will invoke the emacs editor by default, or the editor named by the `EDITOR` environment variable. You can set this environment variable in your shell’s startup files.) The Tools/Pure menu contains a second Pure-related option, “Reload Pure Scripts” which can be used to quickly reload all loaded Pure scripts after edits; more about that later.

Please note that most of the functions in `pure_func.pure` are rather useless, they are only provided for illustrative purposes. However, there are some useful examples in there, too, in particular:

- `pure_eval` lets you evaluate any Pure expression, given as a string in its first argument. E.g., try something like `=pure_eval("foldl (+) 0 (1..100)")`. Additional parameters are accessible as `x!0`, `x!1`, etc. For instance: `=pure_eval("x!0+x!1",A1,B1)`.
- `pure_echo` just displays its arguments as a string in Pure syntax, as the interpreter sees them. This is useful for debugging purposes. E.g., `=pure_echo(A1:B10)` shows the given range as a Pure matrix.
- `pure_shell` is a variation of `pure_eval` which executes arbitrary Pure code and returns the last evaluated expression (if any) as a string. This is mainly provided as a convenience to create an “interactive Pure shell” which lets you evaluate Pure code inside Gnumeric. To these ends, simply prepare a text cell for entering the code to be evaluated, and then apply `pure_shell` on this text cell in another cell to display the result.

A spreadsheet showing most of the predefined functions in action can be found in `pure-examples.gnumeric` example distributed with Gnumeric/Pure.

5 Interactive Pure Shell

The `pure-examples.gnumeric` spreadsheet also includes an instance of `pure_shell` which lets you evaluate arbitrary Pure code in the same interpreter instance that executes Gnumeric/Pure functions. This is very helpful if you're developing new Pure functions to be used in Gnumeric. It also lets you use Gnumeric as a kind of GUI frontend to the Pure interpreter. You can try this now. Open the `pure-examples` spreadsheet in Gnumeric and enter the following into the input cell of the Pure shell:

```
> scanl (+) 0 (1..20)
[0,1,3,6,10,15,21,28,36,45,55,66,78,91,105,120,136,153,171,190,210]
```

Note that here and in the following the prompt `>` indicates a Pure expression to be evaluated in *Gnumeric* (rather than the standalone Pure interpreter), which is followed by another line indicating the result (printed in the output cell below the input cell of the Pure shell). You can find the Pure shell at the bottom of the first sheet in `pure-examples`, see the screenshot below. For your convenience, there's also a second, bigger one on the second sheet. You might want to copy this over to a separate spreadsheet which you can use as a scratchpad for experimentation purposes.




Figure .1: The Pure shell.

Also note that this is in fact *Pure code* (not a Gnumeric formula) being evaluated there. You can execute any Pure code, including Pure declarations, so you can type:

```
> using system; puts "Hello, world!";
14
```

This prints the string `"Hello, world!"` on standard output, visible in the terminal window where you launched Gnumeric. Here is another example, showing how you can invoke any function from the C library, by declaring it as a Pure extern function:

```
> extern int rand(); [rand | i = 1..5];
[1810821799,2106746672,1436605662,1363610028,695042099]
```

All functions in the Pure prelude are readily available in the Gnumeric Pure shell, as well as the functions defined in `pure_func.pure` and its imports, including the programming interface described in [Advanced Features](#). For instance, here's how you can retrieve a cell value from the current sheet:

```
> get_cell "A1"
"Gnumeric/Pure Examples"
```

Using `call` (see [Calling Gnumeric from Pure](#)), you can also invoke any Gnumeric function:

```
> call "product" (1..10)
3628800.0
```

6 Defining Your Own Functions

After playing around with `pure_func.pure` and the interactive Pure shell for a while, of course you will want to write your own functions, that's what this plugin is about after all! For the beginning, you can just add your definitions to the existing `pure_func.pure` script. Use the "Edit Pure Script" option to edit the script in your favourite editor, and see the comments and the examples in the script for guidance. (This document assumes that you're already familiar with Pure, if not then you should consult the available Pure documentation.)

Note that if you delete or rename any functions in this file, or add new ones to it, then you also have to change the list of function names in the `plugin.xml` file in the same directory accordingly. This file tells Gnumeric which functions are provided by the script. Unfortunately, you'll have to restart Gnumeric to make changes in this file take effect. If you only change the definition of an existing function then it's usually sufficient to just invoke "Reload Pure Scripts" afterwards, and maybe run "Recalculate" (F9) to recompute the spreadsheet. However, if you also made changes to the function descriptions provided via `gnm_info` (see the following section for explanation), then you'll also have to restart Gnumeric so that it picks up the changes.

Once you understand how this works, you can also create your own plugin directories with your personal collections of Gnumeric/Pure functions, using the `pure_func` directory as a template. For instance, let's assume that your Gnumeric/Pure stuff is in a script named `gnumeric.pure` under `/some/path/pure/gnumeric`. The `plugin.xml` file in that directory might look as follows:

```
<?xml version="1.0" encoding="UTF-8"?>
<plugin id="Gnumeric_MyPureFunc">
  <information>
    <name>My Pure functions</name>
    <description>My Pure functions for use in Gnumeric.</description>
    <require_explicit_enabling/>
  </information>
  <loader type="Gnumeric_PureLoader:pure">
    <attribute value="gnumeric" name="module_name"/>
  </loader>
  <services>
    <service type="function_group" id="my_pure_func">
      <category>Pure</category>
      <functions>
        <!-- My Pure functions go here, e.g.: -->
        <function name="my_pure_func"/>
      </functions>
    </service>
  </services>
</plugin>
```

```
</services>
</plugin>
```

The following steps are needed to tell Gnumeric about your new Pure plugin:

- Open the Directories tab in the Tools/Plug-ins dialog and check that the `/some/path/pure` directory is in your plugin search path. Add it if necessary, and restart Gnumeric after that. (Note that you really have to add the *parent* directory `/some/path/pure`, not `/some/path/pure/gnumeric` itself.)
- The new plugin should now be listed as “My Pure functions” on the Plugin List tab in the Tools/Plug-ins dialog. Check it to enable it.

The Pure loader can load multiple Pure plugins. You only need to tell Gnumeric about them after creating the scripts and `plugin.xml` files and placing them into corresponding plugin directories. Just enable the ones that you want in Tools/Plug-ins. All scripts are loaded in the same Pure interpreter (and thus are treated like one big script) so that functions in one script can use the function and variable definitions in another. If you need to access the definitions in the `pure_func.pure` “mother script”, you can also just import it into your scripts with a `using` clause, i.e.: `using pure_func;`

7 Gnumeric/Pure Interface

By default, when a Pure function is called from Gnumeric, it receives its arguments in a list. However, it is possible to tell Gnumeric about the expected arguments of the function and also specify a help text to be displayed in the “f(x)” dialog, by giving a definition of the special `gnm_info` function as explained below.

7.1 Function Descriptions

To describe a given function to Gnumeric, define `gnm_info` “<name>” (where <name> is the name of the function) as a pair with the following elements:

- The first element, a string, gives the signature of the function. E.g., “” denotes a function without arguments, “f” a function taking a single float parameter, “fs” a function taking a float and a string argument (in that order), etc. Optional parameters can be indicated using `|`, as in “ff|s” (two non-optional floats, followed by an optional string). See below for a complete list of the supported parameter types.
- The second element is a list of hash pairs `key=>text` which together make up the help text shown in Gnumeric’s “f(x)” dialog. You should at least specify the function name along with a short synopsis here, e.g. `GNM_FUNC_HELP_NAME => "frob: the frob function"`. Parameter descriptions take the form `GNM_FUNC_HELP_ARG => "x: integer"`. There are a number of other useful elements, see below for details.

Both the signature and the function description are optional. That is, `gnm_info` may return either just a signature string, or a list of hash pairs with the function description, or both.

The signature defaults to a variadic function which takes any number of parameters of any type (see below), and the description defaults to some boilerplate text which says that the function hasn't been documented yet.

Note that if no signature is given, then the function accepts any number of parameters of any type. In that case, or if there are optional parameters, the function becomes variadic and the (optional) parameters are passed as a Pure list (in addition to the non-optional parameters).

Here's the list of valid parameter types, as they are documented in the Gnumeric sources:

<code>f</code>	: float	no errors, string conversion attempted
<code>b</code>	: boolean	identical to <code>f</code>
<code>s</code>	: string	no errors
<code>S</code>	: scalar	any non-error scalar
<code>E</code>	: scalar	any scalar, including errors
<code>r</code>	: cell range	content may not be evaluated yet
<code>A</code>	: area	array, range (as above), or scalar
<code>?</code>	: anything	any value (scalars, non-scalars, errors, whatever)

The keys used in the function description may be any of the following, along with sample text for each type of field:

<code>GNM_FUNC_HELP_NAME</code>	=> <code>"name:synopsis"</code>
<code>GNM_FUNC_HELP_ARG</code>	=> <code>"name:parameter description"</code>
<code>GNM_FUNC_HELP_DESCRIPTION</code>	=> <code>"Long description."</code>
<code>GNM_FUNC_HELP_NOTE</code>	=> <code>"Note."</code>
<code>GNM_FUNC_HELP_EXAMPLES</code>	=> <code>"=sample_formula()"</code>
<code>GNM_FUNC_HELP_SEEALSO</code>	=> <code>"foo,bar,..."</code>

The following keys are only supported in the latest Gnumeric versions:

<code>GNM_FUNC_HELP_EXTREF</code>	=> <code>"wiki:en:Trigonometric_functions"</code>
<code>GNM_FUNC_HELP_EXCEL</code>	=> <code>"Excel compatibility information."</code>
<code>GNM_FUNC_HELP_ODF</code>	=> <code>"OpenOffice compatibility information."</code>

Note that inside the descriptions, the notation `@{arg}` (`@arg` in older Gnumeric versions) can be used to refer to a parameter value. For instance, here's a sample description for a binary function which also includes a help text:

```
gnm_info "pure_max" = "ff",
[GNM_FUNC_HELP_NAME    => "pure_max:maximum of two numbers",
 GNM_FUNC_HELP_ARG     => "x:number",
 GNM_FUNC_HELP_ARG     => "y:number",
 GNM_FUNC_HELP_DESCRIPTION =>
  "Computes the maximum of two numbers @{x} and @{y}.",
 GNM_FUNC_HELP_EXAMPLES => "=pure_max(17,22)"];
```

As you can see, the function descriptions are a bit unwieldy, so it's convenient to construct them using this little helper function defined in `pure_func.pure`:

```
gnm_help name::string args descr::string notes examples see_also =
  [GNM_FUNC_HELP_NAME    => name] +
```



```
[GNM_FUNC_HELP_ARG          => x | x::string = args ] +
[GNM_FUNC_HELP_DESCRIPTION  => descr ] +
[GNM_FUNC_HELP_NOTE         => x | x::string = notes ] +
[GNM_FUNC_HELP_EXAMPLES     => x | x::string = examples ] +
(if null see_also then [] else
 [GNM_FUNC_HELP_SEEALSO     => join ", " see_also]);
```

Now the description can be written simply as follows:

```
gnm_info "pure_max" = "ff", gnm_help "pure_max:maximum of two numbers"
  ["x:number", "y:number"]
  "Computes the maximum of two numbers @{x} and @{y}."
  [] ["=pure_max(17,22)"] [];
```

Since this function only has fixed arguments, it will be called in curried form, i.e., as `pure_max x y`. For instance, the actual definition of `pure_max` may look as follows:

```
pure_max x y = max x y;
```

Conversely, if no signature is given, then the function accepts any number of parameters of any type, which are passed as a list. For instance:

```
gnm_info "pure_sum" = gnm_help "pure_sum:sum of a collection of numbers"
  [] "Computes the sum of a collection of numbers."
  [] ["=pure_sum(1,2,3,4,5,6)"] ["pure_sums"];
```

Here the function will be called as `pure_sum [x1,x2,...]`, where `x1`, `x2`, etc. are the arguments the function is invoked with. Note that in this case there may be any number of arguments (including zero) of any type, so your function definition must be prepared to handle this. If a function does not have a `gnm_info` description at all then it is treated in the same fashion. The `pure_func.pure` script contains some examples showing how to write functions which can deal with any numbers of scalars, arrays or ranges, see the `pure_sum` and `pure_sums` examples. These employ the following `ranges` function to “flatten” a parameter list to a list holding all denoted values:

```
ranges xs = cat [ case x of _::matrix = list x; _ = [x] end | x = xs ];
```

E.g., the `pure_sum` function can now be defined as follows:

```
pure_sum xs = foldl (+) 0 (ranges xs);
```

A function may also have both fixed and optional arguments (note that in what follows we’re going to omit the detailed function descriptions for brevity):

```
gnm_info "foo" = "ff|ff";
```

In this case the fixed arguments are passed in curried form as usual, while the optional parameters are passed as a list. That is, `foo` may be called as `foo x y []`, `foo x y [z]` or `foo x y [z,t]`, depending on whether it is invoked with two, three or four arguments.

7.2 Conversions Between Pure and Gnumeric Values

The marshalling of types between Gnumeric and Pure is pretty straightforward; basically, Pure numbers, strings and matrices map to Gnumeric numbers, strings and arrays, respectively. The following table summarizes the available conversions:

Pure	Gnumeric
<code>gnm_error "#N/A"</code>	error
<code>4711, 4711L, 4711.0</code>	scalar (number)
<code>"Hello world"</code>	string
<code>()</code>	empty
<code>(1,2,3)</code>	array
<code>[1,2,3]</code>	array
<code>{1,2,3;4,5,6}</code>	array (or cell range)
<code>"A1:B10"</code>	cell range ("r" conversion)

These conversions mostly work both ways. Note that on input, cell ranges are usually passed as matrices to Pure functions (i.e., they are passed “by value”), unless the function signature specifies a “r” conversion in which case the cell ranges themselves are passed to the function in string form. (Such values can also be passed on to Gnumeric functions which expect a cell range (“r”) parameter, see [Calling Gnumeric from Pure](#) below.)

Conversely, matrices, lists and tuples all become Gnumeric arrays on output, so usually you’ll want to enter these as array functions (Ctrl-Shift-Enter in Gnumeric). As a special case, the empty tuple can be used to denote empty cell values (but note that empty Gnumeric values may become zeros when passed as float or array arguments to Pure functions).

If a Pure function returns a value that doesn’t match any of the above then it is converted to a string in Pure expression syntax and that string is returned as the result of the function invocation in Gnumeric. This makes it possible to return any kind of symbolic Pure value, but note that if such a value is then fed into another Pure function, that function will have to convert the string value back to the internal representation if needed; this can be done very conveniently using Pure’s `eval` function, see the Pure documentation for details.

8 Advanced Features

This section explains various additional features provided by the Gnumeric/Pure interface that should be useful for writing your own functions. Note that, for your convenience all functions discussed in this section are declared in `pure_func.pure`.

8.1 Calling Gnumeric from Pure

It is possible to call Gnumeric functions from Pure using the `call` function which takes the name of the function (a string) as its first, and the parameters as the second (list) argument. For instance:

```
gnm_info "gnm_bitand" = "ff";
gnm_bitand x y = call "bitand" [x,y];
```

Note that `call` is an external C function provided by Gnumeric/Pure. If you want to use it, it must be declared in your Pure script as follows:

```
extern expr* pure_gnmcall(char* name, expr* args) = call;
```

However, `pure_func.pure` already contains the above declaration, so you don't have to do this yourself if you import `pure_func.pure` in your scripts.

Also note that `call` doesn't do any of Gnumeric's automatic conversions on the parameters, so you have to pass the proper types of arguments as required by the function.

8.2 Accessing Spreadsheet Cells

Gnumeric/Pure provides the following functions to retrieve and modify the contents of spreadsheet cells and ranges of such cells:

```
extern expr* pure_get_cell(char* s) = get_cell;
extern expr* pure_get_cell_text(char* s) = get_cell_text;
extern expr* pure_get_cell_format(char* s) = get_cell_format;
extern expr* pure_set_cell(char* s, expr *x) = set_cell;
extern expr* pure_set_cell_text(char* s, expr *x) = set_cell_text;
extern expr* pure_set_cell_format(char* s, expr *x) = set_cell_format;
extern expr* pure_get_range(char* s) = get_range;
extern expr* pure_get_range_text(char* s) = get_range_text;
extern expr* pure_get_range_format(char* s) = get_range_format;
extern expr* pure_set_range(char* s, expr *x) = set_range;
extern expr* pure_set_range_text(char* s, expr *x) = set_range_text;
extern expr* pure_set_range_format(char* s, expr *x) = set_range_format;
```

For instance, here's how you use these functions to write and then read some cell values (try this in the interactive Pure shell):

```
> set_cell "A14" 42
()
> get_cell "A14"
42.0
> set_range "A14:G14" $ scanl (*) 1 (1..6)
()
> get_range "A14:G14"
{1.0,1.0,2.0,6.0,24.0,120.0,720.0}
> set_cell_text "A14" "=sum(B14:G14)"
()
> get_cell "A14"
873.0
> get_cell_text "A14"
"=sum(B14:G14)"
> get_range_text "A14:G14"
{"=sum(B14:G14)","1","2","6","24","120","720"}
```

Note that while the `set_cell` function sets the given cell to a constant value, `set_cell_text` also allows you to store a formula in a cell which will then be evaluated as usual. Similarly, `get_cell` retrieves the cell value, while `get_cell_text` yields the text in the cell, as entered by the user (which will either be a formula or the textual representation of a constant value). The `set_range`, `set_range_text`, `get_range` and `get_range_text` functions work analogously, but are used to manipulate entire ranges of cells, which can be set from Pure tuples, lists or matrices, and retrieved as Pure matrices.

Functions to retrieve and change the cell format are also provided (watch the contents of the cell A14 change its color to blue on entering the first expression):

```
> set_cell_format "A14" "[Blue]0.00"
()
> get_range_format "A14:C14"
{"[Blue]0.00", "General", "General"}
```

There are also functions to get the position of the “current” cell (i.e., the cell from which a Pure function was called), and to translate between cell ranges in Gnumeric syntax and the corresponding internal representation consisting of a pointer to a Gnumeric sheet and the cell or range coordinates:

```
extern expr* pure_this_cell() = this_cell;
extern expr* pure_parse_range(char* s) = parse_range;
extern expr* pure_make_range(expr* x) = make_range;
```

Examples:

```
> this_cell
"B4"
> parse_range this_cell
#<pointer 0x875220>,1,3
> make_range (NULL,0,0,10,10)
"Sheet2!A1:K11"
```

8.3 Asynchronous Data Sources

Gnumeric/Pure makes it easy to set up asynchronous data sources which draw values from a Pure computation executed in a background process. This facility is useful to carry out lengthy computations in the background while you can continue to work with your spreadsheet. It also allows you to process incoming data and asynchronous events from special devices (MIDI, sensors, stock tickers, etc.) in (soft) realtime.

To do this, you simply pass an expression to the `datasource` function. This is another external C function provided by Gnumeric/Pure, which is declared in `pure_func.pure` as follows:

```
extern expr* pure_datasource(expr* x) = datasource;
```

The argument to `datasource` is typically a thunk or stream (lazy list) which is to be evaluated in the background. The call to `datasource` initially returns a #N/A value (`gnm_error "#N/A"`)

while the computation is still in progress. The cell containing the data source then gets updated automatically as soon as the value becomes available, at which point the `datasource` call now returns the computed value. E.g., here's how you would wrap up a lengthy calculation as a thunk and submit it to `datasource` which carries out the computation as a background task:

```
gnm_info "pure_frob" = "f";
pure_frob x = datasource (lengthy_calculation x&);
lengthy_calculation x = sleep 3 $$ foldl (*) 1 (1..x);
```

Note that a cell value may draw values from as many independent data sources as you want, so the definition of a cell may also involve multiple invocations of `datasource`:

```
gnm_info "pure_frob2" = "ff";
pure_frob2 x y = datasource (lengthy_calculation x&),
  datasource (lengthy_calculation y&);
```

Special treatment is given to (lazy) lists, in this case `datasource` returns a new value each time a list element becomes available. For instance, the following function uses an infinite stream to count off the seconds starting from a given initial value:

```
gnm_info "pure_counter" = "f";
pure_counter x = datasource [sleep (i>x) $$ i | i = x..inf];
```

You can also try this interactively in the Pure shell:

```
> datasource [sleep (i>0) $$ i | i = 0..inf]
0
1
...
```

Here's another example for the Pure shell which prints the prime numbers:

```
> datasource primes with primes = sieve (2..inf);
  sieve (p:qs) = p : (sleep 1 $$ sieve [q | q = qs; q mod p])& end
2
3
5
...
```

Note that when processing a lazy list, the cell containing the call will keep changing as long as new values are produced (i.e., forever in this example). The "Stop Data Sources" option in the Tools/Pure menu can be used to stop all active data sources. "Reload Pure Scripts" also does this. You can then restart the data sources at any time by using "Recalculate" (F9) to recompute the spreadsheet.

Also note that because of the special way that `datasource` handles list values, you cannot return a list directly as the result of `datasource`, if it is to be treated as a single result. Instead, you'll have to wrap the result in a singleton list (e.g., `datasource [[lengthy_calculation x, lengthy_calculation y]&]]`), or return another aggregate (i.e., a matrix or a tuple).

Finally, note that when the arguments of a call involving `datasource` change (because they depend on other cells which may have been updated), the computation is automatically

restarted with the new parameters. The default behaviour in this case is that the entire computation will be redone from scratch, but it's also possible to wrap up calls to `datasource` in a manner which enables more elaborate communication between Gnumeric and background tasks initiated with `datasource`. This is beyond the scope of this manual, however, so we leave this as an exercise to the interested reader.

8.4 Triggers

In addition to asynchronous data sources, the `trigger` function is provided to compute values or take actions depending on some external condition, such as the availability of data on a special device or the creation of some widget (see the next section):

```
extern expr* pure_trigger(int timeout, expr* cond, expr *val, expr *data)
    = trigger;
```

Thus a typical invocation of the function looks as follows:

```
trigger timeout condition value data
```

The condition and value arguments are callback functions which get invoked by `trigger`, passing them the given data argument. The trigger reevaluates the given condition in regular intervals (1 second in the current implementation) and, as soon as it becomes true, computes the given value and returns that value as the result of the trigger call. As long as the condition doesn't hold, trigger returns a #N/A value (`gnm_error "#N/A"`). Note that, in difference to `datasource`, both the condition and the value are computed in Gnumeric (rather than a child process), so that it is possible to access the current information in the loaded spreadsheet.

The timeout value determines how often the condition is checked. If it is positive, the condition will be reevaluated `timeout+1` times (once initially, and then once per second for a total duration of timeout seconds). If it is negative, the trigger never times out and the condition will be checked repeatedly until the trigger expression is removed (or Gnumeric is exited). In either case value data will be recomputed each time condition data yields true. (This is most useful if the computed value, as a side effect, arranges for the condition to become false again afterwards.) Finally, if timeout is zero then the trigger fires at most once, as soon as the condition becomes true, at which point value data is computed just once.

Here's a (rather useless) example of a trigger which fires exactly once, as soon as a certain cell goes to a certain value, and then modifies another cell value accordingly:

```
> trigger 0 (\->get_cell "A14"=="Hello") (\->set_cell "A15" "World") ()
```

Now, as soon as you type Hello in the cell A14, the trigger will print World in cell A15. Note that the data argument isn't used here. A more useful example will be discussed in the following section.

8.5 Sheet Objects

Gnumeric offers some kinds of special objects which can be placed on a sheet. This comprises the chart and image objects which can be found in the “Insert” menu, as well as a number of useful graphical elements and GUI widgets on the “Object” toolbar, accessible via “View/Toolbars”. The latter are also useful for providing control input to Pure functions.

Gnumeric/Pure provides the following function to retrieve information about the special objects in a spreadsheet:

```
extern expr* pure_sheet_objects() = sheet_objects;
```

For instance, with one button object in your spreadsheet, the output of `sheet_objects` might look like this:

```
> sheet_objects
[("Sheet1", "button", "Push Me!", "A11", [#<pointer 0x2a1dcd0>])]
```

Each object is described by a tuple which lists the name of the sheet on which the object is located, the type of object, the object’s content or label (if applicable), the cell which the object is linked to (if applicable), and a list of pointers to the corresponding `GtkWidgets` (if any). Note that in general a GUI object may be associated with several widgets, as Gnumeric allows you to have multiple views on the same spreadsheet, so there will be one widget for each view an object is visible in. Also note that the content/label information depends on the particular type of object:

- List and combo widgets return the content link (referring to the cells in the spreadsheet holding the items shown in the list).
- Frame and button widgets return the label shown on the widget.
- Graphic objects like rectangles and ellipses return the text content of the object.
- Image objects (type “image/xyz”, where xyz is the type of image, such as svg or png) return a pointer to the image data in this field.

The `sheet_objects` function is a bit tricky to use, since some of the objects or their associated widgets might not have been created yet when the spreadsheet is loaded. Therefore it is necessary to use a trigger to make sure that the information is updated once all objects are fully displayed. The `pure_func.pure` script contains the following little wrapper around `sheet_objects` which does this:

```
pure_objects = trigger 0 (\_ ->all realized sheet_objects)
(\_ ->matrix$map list sheet_objects) ()
with realized (_,_,_,_,w) = ~listp w || ~null w && ~any null w end;
```

See the `widgets.gnumeric` spreadsheet in the distribution for an example.

Possible uses of this facility are left to your imagination. Using Gnumeric’s internal APIs and Pure’s Gtk interface, you might manipulate the GUI widgets in various ways (add icons to buttons or custom child widgets to frames, etc.). One particularly useful case, for which

Gnumeric/Pure has built-in support, is rendering an OpenGL scene in a Gnumeric frame widget, see below.

8.6 OpenGL Interface

Gnumeric/Pure provides special support for rendering OpenGL scenes into Gnumeric frame widgets. To actually use this, you must have Pure's OpenGL module installed. The following function is provided to equip a Gnumeric frame with the OpenGL rendering capability:

```
extern expr *pure_gl_window(char *name, int timeout,  
                             expr *setup_cb, expr *config_cb,  
                             expr *display_cb, expr *timer_cb,  
                             expr *user_data) = gl_window;
```

The meaning of the parameters is as follows:

- `name` is a string which specifies the label of the frame widget into which the scene is to be rendered.
- `timeout` is a time value in milliseconds (an integer) which specifies the period for invocations of the `timer_cb` callback, see below. If this value is zero or negative then the timer callback is disabled.
- `setup_cb`, `config_cb`, `display_cb` and `timer_cb` are the Pure callback functions which are invoked by `gl_window` to actually render the scene. The callbacks are all invoked with two arguments: the `user_data` parameter that `gl_window` was invoked with, and a second parameter with callback-specific information as described below.
- `user_data` is any information that the caller wants to be passed as the first argument to the callback functions.

The different callbacks are:

- `setup_cb` is called for initializing the scene. It receives the `GtkDrawingArea` widget as the second argument. Typically this is used to set the initial projection and modelview matrices, enable lighting, etc.
- `config_cb` is called when width or height of the frame widget changes so that the rendering parameters (typically the viewport) can be adjusted accordingly. It is invoked with a pair of integers (`w,h`) as the second argument, which denotes the new dimension allocated to the drawing area.
- `display_cb` is called whenever the contents of the drawing area needs to be redrawn. This typically does most of the work necessary to render the scene. The second callback argument is always `()`.
- `timer_cb` is called at regular intervals as specified with the `timeout` parameter (see above), unless the `timeout` value is zero or negative in which case this callback is disabled. The second callback argument is always `()`. The timer callback is typically used to incrementally adjust some parameters of the scene in order to render animations.

When invoked, the callback automatically arranges for `display_cb` to be called afterwards, so you don't have to do that manually in your callback definition.

You'll need at least either the `display_cb` or the `timer_cb` function to render anything, but typically all of these callbacks will be needed for animated scenes. Callback functions which aren't needed can be specified as `()`.

There's also a related helper function which can be used as a trigger condition to defer rendering until the target frame widget has been realized:

```
extern bool pure_check_window(char *name) = check_window;
```

This function returns true as soon as the named frame widget is ready to go, at which time `gl_windows` can be called on the widget. So your call to `gl_windows` should usually be wrapped up like this:

```
trigger 0 check_window
(\frame->gl_window frame timeout setup config display timer user_data) frame
```

Here's an example from `pure_glfunc.pure` which shows how these functions are to be used:

```
using pure_func, GL, GLU;
extern void gdk_gl_draw_teapot(bool solid, double scale);
gnm_info "gltest" = "sbfff";

gltest frame m a b c = trigger 0 check_window
  (\frame->gl_window frame (m*40) setup config display timer ()) frame
with
  setup _ _ = () when
    // Initialize.
    GL::ClearColor 0.1 0.1 0.3 1.0;
    GL::ShadeModel GL::SMOOTH;
    GL::Enable GL::DEPTH_TEST;
    // Initial projection and modelview matrices.
    GL::MatrixMode GL::PROJECTION;
    GL::LoadIdentity;
    GL::Rotatef 20.0 (-1.0) 0.0 0.0;
    GL::MatrixMode GL::MODELVIEW;
    GL::LoadIdentity;
    // Lighting.
    GL::Lightfv GL::LIGHT0 GL::DIFFUSE {1.0,0.0,0.0,1.0};
    GL::Lightfv GL::LIGHT0 GL::POSITION {2.0,2.0,-5.0,1.0};
    GL::Enable GL::LIGHTING;
    GL::Enable GL::LIGHT0;
  end;
  config _ (w,h) = GL::Viewport 0 0 w h;
  display _ _ = () when
    GL::Clear (GL::DEPTH_BUFFER_BIT or GL::COLOR_BUFFER_BIT);
    gdk_gl_draw_teapot true 0.5;
  end if m;
  display _ _ = () when
    GL::Clear (GL::DEPTH_BUFFER_BIT or GL::COLOR_BUFFER_BIT);
    GL::LoadIdentity;
```

```

GL::Rotatef (scale 360 a) 0.0 1.0 0.0;
GL::Rotatef (scale 360 b) 1.0 0.0 0.0;
GL::Rotatef (scale 360 c) 0.0 0.0 1.0;
gdk_gl_draw_teapot true 0.5;
end;
timer _ _ = () when
  GL::Rotatef (scale 36 a) 0.0 1.0 0.0;
  GL::Rotatef (scale 36 b) 1.0 0.0 0.0;
  GL::Rotatef (scale 36 c) 0.0 0.0 1.0;
end;
scale step x = (x/100*step);
end;

```

Have a look at the `gl-example.gnumeric` spreadsheet included in the distribution to see this example in action. (You first need to enable the “Pure OpenGL functions” in the Plugin Manager to make this work.) The screenshot below shows how the example looks like in Gnumeric.



Figure .2: Gnumeric/Pure OpenGL example.

9 Copying

Copyright (c) 2009 by Albert Graef.

Gnumeric/Pure is free software: you can redistribute it and/or modify it under the terms

of the GNU General Public License as published by the Free Software Foundation, either version 2 of the License, or (at your option) any later version.

Gnumeric/Pure is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

Pure-GLPK - GLPK interface for the Pure programming language

Version 0.1, November 14, 2010

Jiri Spitz <jiri.spitz@bluetone.cz>

This module provides a feature complete GLPK interface for the Pure programming language, which lets you use all capabilities of the GNU Linear Programming Kit (GLPK) directly from Pure.

GLPK (see <http://www.gnu.org/software/glpk>) contains an efficient simplex LP solver, a simplex LP solver in exact arithmetics, an interior-point solver, a branch-and-cut solver for mixed integer programming and some specialized algorithms for net/grid problems. Using this interface you can build, modify and solve the problem, retrieve the solution, load and save the problem and solution data in standard formats and use any of advanced GLPK features.

The interface uses native Pure data types - lists and tuples - so that you need not perform any data conversions to/from GLPK internal data structures.

To make this module work, you must have a GLPK installation on your system, the version 4.38 is required.

1 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-glpk-0.1.tar.gz>.

Run `make` to compile the module and `make install` (as root) to install it in the Pure library directory. This requires GNU `make`, and of course you need to have Pure installed.

The default make options suppose that GLPK was configured with the following options:
`--enable-dl --enable-odbc --enable-mysql --with-gmp --with-zlib`

Using the given options the dependencies are:

- GNU Multiprecision Library (GMP) - serves for the exact simplex solver. When disabled, the exact solver still works but it is much slower.

- ODBC library - serves for reading data directly from database tables within the GNU MathProg language translator through the ODBC interface.
- zlib compression library - enables reading and writing gzip compressed problem and solution files.
- MySQL client library - serves for reading data directly from MySQL tables within the GNU MathProg language translator.
- ltdl dlopen library - must be enabled together with any of ODBC, zlib or MySQL.

make tries to guess your Pure installation directory and platform-specific setup. If it gets this wrong, you can set some variables manually. In particular, `make install prefix=/usr` sets the installation prefix, and `make PIC=-fPIC` or some similar flag might be needed for compilation on 64 bit systems. The variable `ODBCLIB` specifies the ODBC library to be linked with. The default value is `ODBCLIB=-lodbc`. Please see the Makefile for details.

2 Error Handling

When an error condition occurs, the GLPK library itself prints an error message and terminates the application. This behaviour is not pleasant when working within an interpreter. Therefore, the Pure - GLPK bindings catches at least the most common errors like indices out of bounds. On such an error an appropriate message is returned to the interpreter. The less common errors are still trapped by the GLPK library.

When one of the most common errors occurs, an error term of the form `glp::error message` will be returned, which specifies what kind of error happened. For instance, an index out of bounds will cause a report like the following:

```
glp::error "[Pure GLPK error] row index out of bounds"
```

You can check for such return values and take some appropriate action. By redefining `glp::error` accordingly, you can also have it generate exceptions or print an error message. For instance:

```
glp::error message = fprintf stderr "%s\n" message $$ ();
```

NOTE: When redefining `glp::error` in this manner, you should be aware that the return value of `glp::error` is what will be returned by the other operations of this module in case of an error condition. These return values are checked by other functions. Thus the return value should still indicate that an error has happened, and not be something that might be interpreted as a legal return value, such as an integer or a nonempty tuple. It is usually safe to have `glp::error` return an empty tuple or throw an exception, but other types of return values should be avoided.

IMPORTANT: It is really good to define a `glp::error` function, otherwise the errors might remain unnoticed.

3 Further Information and Examples

For further details about the operations provided by this module please see the GLPK Reference Manual. Sample scripts illustrating the usage of the module can be found in the examples directory.

4 Interface description

Most GLPK functions and symbols live in the namespace `glp`. There are a few functions and symbols in the namespace `lpx`. These functions and symbols are likely to be removed and replaced by new ones in the future.

In general, when you replace the `glp_` prefix from the GLPK Reference Manual with the namespace specification `glp::` then you receive the function name in this module. The same is valid for `lpx_` and `lpx::`. The symbolic constants are converted into lower case in this module, again obeying the same prefix rules.

5 Descriptions of interface functions

5.1 Basic API routines

Problem creating and modifying routines

Create the GLPK problem object Synopsis:

```
glp::create_prob
```

Parameters:

none

Returns:

pointer to the LP problem object

Example:

```
> let lp = glp::create_prob;  
> lp;  
#<pointer 0x9de7168>
```

Set the problem name Synopsis:

```
glp::set_prob_name lp name
```

Parameters:

lp pointer to the LP problem object
name problem name

Returns:

()

Example:

```
> glp::set_prob_name lp "Testing problem";  
()
```

Set objective name Synopsis:

```
glp::set_obj_name lp name
```

Parameters:

lp pointer to the LP problem object
name objective name

Returns:

()

Example:

```
> glp::set_obj_name lp "Total costs";  
()
```

Set the objective direction Synopsis:

```
glp::set_obj_dir lp direction
```

Parameters:

lp pointer to the LP problem object
direction one of the following:
 glp::min minimize
 glp::max maximize

Returns:

()

Example:

```
> glp::set_obj_dir lp glp::min;  
()
```


Add new rows to the problem Synopsis:

```
glp::add_rows lp count
```

Parameters:

lp pointer to the LP problem object

count number of rows to add

Returns:

index of the first row added

Example:

```
> let first_added_row = glp_add_rows lp 3;  
> first_added_row;  
6
```

Add new columns to the problem Synopsis:

```
glp::add_cols lp count
```

Parameters:

lp pointer to the LP problem object

count number of columns to add

Returns:

index of the first column added

Example:

```
> let first_added_col = glp_add_cols lp 3;  
> first_added_col;  
5
```

Set the row name Synopsis:

```
glp::set_row_name lp (rowindex, rowname)
```

Parameters:

lp pointer to the LP problem object

rowindex row index

rowname row name

Returns:

()

Example:

```
> glp::set_row_name lp (3, "The third row");  
()
```

Set the column name Synopsis:

```
glp::set_col_name lp (colindex, colname)
```

Parameters:

lp pointer to the LP problem object

colindex column index

colname column name

Returns:

()

Example:

```
> glp::set_col_name lp (3, "The third column");  
()
```

Set (change) row bounds Synopsis:

```
glp::set_row_bnds lp (rowindex, rowtype, lowerbound, upperbound)
```

Parameters:

lp pointer to the LP problem object

rowindex row index

rowtype one of the following:

glp::fr free variable (both bounds are ignored)

glp::lo variable with lower bound (upper bound is ignored)

glp::up variable with upper bound (lower bound is ignored)

glp::db double bounded variable

glp::fx fixed variable (lower bound applies, upper bound is ignored)

lowerbound lower row bound

upperbound upper row bound

Returns: ()

Example:: `glp::set_row_bnds lp (3, glp::up, 0.0, 150.0);`

Set (change) column bounds **Synopsis:**

`glp::set_col_bnds lp (colindex, coltype, lowerbound, upperbound)`

Parameters:

lp pointer to the LP problem object

colindex column index

coltype one of the following:

glp::fr free variable (both bounds are ignored)

glp::lo variable with lower bound (upper bound is ignored)

glp::up variable with upper bound (lower bound is ignored)

glp::db double bounded variable

glp::fx fixed variable (lower bound applies, upper bound is ignored)

lowerbound lower column bound

upperbound upper column bound

Returns:

()

Example:

```
> glp::set_col_bnds lp (3, glp::db, 100.0, 150.0);  
( )
```

Set (change) objective coefficient or constant term **Synopsis:**

`glp::set_obj_coef lp (colindex, coefficient)`

Parameters:

lp pointer to the LP problem object

colindex column index, zero index denotes the constant term (objective shift)

Returns:

()

Example:

```
> glp::set_obj_coef lp (3, 15.8);  
()
```

Load or replace matrix row Synopsis:

```
glp::set_mat_row lp (rowindex, rowvector)
```

Parameters:

lp pointer to the LP problem object

rowindex row index

rowvector list of tuples (colindex, coefficient); only non-zero coefficients have to be specified, the order of column indices is not important, duplicates are **not** allowed

Returns:

()

Example:

```
> glp::set_mat_row lp (3, [(1, 3.0), (4, 5.2)]);  
()
```

Load or replace matrix column Synopsis:

```
glp::set_mat_col lp (colindex, colvector)
```

Parameters:

lp pointer to the LP problem object

colindex column index

colvector list of tuples (rowindex, coefficient); only non-zero coefficients have to be specified, the order of row indices is not important, duplicates are **not** allowed

Returns:

()

Example:

```
> glp::set_mat_col lp (2, [(4, 2.0), (2, 1.5)]);  
()
```

Load or replace the whole problem matrix Synopsis:

```
glp::load_matrix lp matrix
```

Parameters:

lp pointer to the LP problem object

matrix list of tuples (rowindex, colindex, coefficient); only non-zero coefficients have to be specified, the order of indices is not important, duplicates are **not** allowed

Returns:

()

Example:

```
> glp::load_matrix lp [(1, 3, 5.0), (2, 2, 3.5), (3, 1, -2.0), (3, 2, 1.0)];
()
```

Check for duplicate elements in sparse matrix Synopsis:

```
glp::check_dup numrows numcols indices
```

Parameters:

numrows number of rows

numcols number of columns

indices list of tuples (rowindex, colindex); indices of only non-zero coefficients have to be specified, the order of indices is not important

Returns:

returns one of the following:

0 the matrix has no duplicate elements

-k rowindex or colindex of the k-th element in indices is out of range

+k the k-th element in indices is duplicate

Remark:

Notice, that k counts from 1, whereas list members are counted from 0.

Example:

```
> glp::check_dup 3 3 [(1, 3), (2, 2), (3, 1), (2, 2)];
4
```

Sort elements of the constraint matrix Synopsis:

```
glp::sort_matrix lp
```

Parameters:

lp pointer to the LP problem object

Returns:

()

Example:

```
> glp::sort_matrix lp;  
()
```

Delete rows from the matrix Synopsis:

```
glp::del_rows lp rows
```

Parameters:

lp pointer to the LP problem object

rows list of indices of rows to be deleted; the order of indices is not important, duplicates are **not** allowed

Returns:

()

Remark:

Deleting rows involves changing ordinal numbers of other rows remaining in the problem object. New ordinal numbers of the remaining rows are assigned under the assumption that the original order of rows is not changed.

Example:

```
> glp::del_rows lp [3, 4, 7];  
()
```

Delete columns from the matrix Synopsis:

```
glp::del_cols lp cols
```

Parameters:

lp pointer to the LP problem object

cols list of indices of columns to be deleted; the order of indices is not important, duplicates are **not** allowed

Returns:

()

Remark:

Deleting columns involves changing ordinal numbers of other columns remaining in the problem object. New ordinal numbers of the remaining columns are assigned under the assumption that the original order of columns is not changed.

Example:

```
> glp::del_cols lp [6, 4, 5];
()
```

Copy the whole content of the GLPK problem object to another one **Synopsis:**

```
glp::copy_prob destination source names
```

Parameters:

destination pointer to the destination LP problem object (must already exist)

source pointer to the source LP problem object

names one of the following:

glp::on copy all symbolic names as well

glp::off do not copy the symbolic names

Returns:

()

Example:

```
> glp::copy_prob lp_dest lp_src glp::on;
()
```

Erase all data from the GLPK problem object **Synopsis:**

```
glp::erase_prob lp
```

Parameters:

lp pointer to the LP problem object, it remains still valid after the function call

Returns:

()

Example:

```
> glp::erase_prob lp;  
()
```

Delete the GLPK problem object Synopsis:

```
glp::delete_prob lp
```

Parameters:

lp pointer to the LP problem object, it is not valid any more after the function call

Returns:

()

Example:

```
> glp::delete_prob lp;  
()
```

Problem retrieving routines

Get the problem name Synopsis:

```
glp::get_prob_name lp
```

Parameters:

lp pointer to the LP problem object

Returns:

name of the problem

Example:

```
> glp::get_prob_name lp;  
"Testing problem"
```

Get the objective name Synopsis:

```
glp::get_obj_name lp
```

Parameters:

lp pointer to the LP problem object

Returns:

name of the objective

Example:

```
> glp::get_obj_name lp;  
"Total costs"
```

Get the objective direction **Synopsis:**

```
glp::get_obj_dir lp
```

Parameters:

`lp` pointer to the LP problem object

Returns:

returns one of the following:

glp::min minimize

glp::max maximize

Example:

```
> glp::get_obj_dir lp;  
glp::min
```

Get number of rows **Synopsis:**

```
glp::get_num_rows lp
```

Parameters:

`lp` pointer to the LP problem object

Returns:

number of rows (constraints)

Example:

```
> glp::get_num_rows lp;  
58
```

Get number of columns **Synopsis:**

```
glp::get_num_cols lp
```

Parameters:

`lp` pointer to the LP problem object

Returns:

number of columns (structural variables)

Example:

```
> glp::get_num_cols lp;  
65
```

Get name of a row Synopsis:

```
glp::get_row_name lp rowindex
```

Parameters:

lp pointer to the LP problem object
rowindex row index

Returns:

name of the given row

Example:

```
> glp::get_row_name lp 3;  
"The third row"
```

Get name of a column Synopsis:

```
glp::get_col_name lp colindex
```

Parameters:

lp pointer to the LP problem object
colindex column index

Returns:

name of the given column

Example:

```
> glp::get_col_name lp 2;  
"The second column"
```

Get row type Synopsis:

```
glp::get_row_type lp rowindex
```

Parameters:

lp pointer to the LP problem object
rowindex row index

Returns:

returns one of the following:

- glp::fr** free variable
- glp::lo** variable with lower bound
- glp::up** variable with upper bound
- glp::db** double bounded variable
- glp::fx** fixed variable

Example:

```
> glp::get_row_type lp 3;
glp::db
```

Get row lower bound Synopsis:

```
glp::get_row_lb lp rowindex
```

Parameters:

- lp** pointer to the LP problem object
- rowindex** row index

Returns:

the row lower bound; if the row has no lower bound then it returns the smallest double number

Example:

```
> glp::get_row_lb lp 3;
50.0
```

Get row upper bound Synopsis:

```
glp::get_row_ub lp rowindex
```

Parameters:

- lp** pointer to the LP problem object
- rowindex** row index

Returns:

the row upper bound; if the row has no upper bound then it returns the biggest double number

Example:

```
> glp::get_row_ub lp 3;
150.0
```

Get column type Synopsis:

```
glp::get_col_type lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

returns one of the following:

glp::fr free variable

glp::lo variable with lower bound

glp::up variable with upper bound

glp::db double bounded variable

glp::fx fixed variable

Example:

```
> glp::get_col_type lp 2;
glp::up
```

Get column lower bound Synopsis:

```
glp::get_col_lb lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

the column lower bound; if the column has no lower bound then it returns the smallest double number

Example:

```
> glp::get_col_lb lp 3;
-1.79769313486232e+308
```

Get column upper bound Synopsis:

```
glp::get_col_ub lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

the column upper bound; if the column has no upper bound then it returns the biggest double number

Example:

```
> glp::get_col_lb lp 3;
150.0
```

Get objective coefficient Synopsis:

```
glp::get_obj_coef lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index; zero index denotes the constant term (objective shift)

Returns:

the coefficient of given column in the objective

Example:

```
> glp::get_obj_coef lp 3;
5.8
```

Get number of nonzero coefficients Synopsis:

```
glp::get_num_nz lp
```

Parameters:

lp pointer to the LP problem object

Returns:

number of non-zero coefficients in the problem matrix

Example:

```
> glp::get_num_nz lp;  
158
```

Retrieve a row from the problem matrix **Synopsis:**

```
glp::get_mat_row lp rowindex
```

Parameters:

lp pointer to the LP problem object

rowindex row index

Returns:

non-zero coefficients of the given row in a list form of tuples (colindex, coefficient)

Example:

```
> get_mat_row lp 3;  
[(3,6.0),(2,2.0),(1,2.0)]
```

Retrieve a column from the problem matrix **Synopsis:**

```
glp::get_mat_col lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

non-zero coefficients of the given column in a list form of tuples (rowindex, coefficient)

Example:

```
> get_mat_col lp 2;  
[(3,2.0),(2,4.0),(1,1.0)]
```

Row and column searching routines

Create index for searching rows and columns by their names **Synopsis:**

```
glp::create_index lp
```

Parameters:

lp pointer to the LP problem object

Returns:

()

Example:

```
> glp::create_index lp;  
()
```

Find a row number by name Synopsis:

```
glp::find_row lp rowname
```

Parameters:

lp pointer to the LP problem object

rowname row name

Returns:

ordinal number (index) of the row

Remark:

The search index is automatically created if it does not already exists.

Example:

```
> glp::find_row lp "The third row";  
3
```

Find a column number by name Synopsis:

```
glp::find_col lp colname
```

Parameters:

lp pointer to the LP problem object

colname column name

Returns:

ordinal number (index) of the column

Remark:

The search index is automatically created if it does not already exists.

Example:

```
> glp::find_col lp "The second row";  
2
```

Delete index for searching rows and columns by their names Synopsis:

```
glp::delete_index lp
```

Parameters:

lp pointer to the LP problem object

Returns:

()

Example:

```
> glp::delete:index lp;  
()
```

Problem scaling routines

Set the row scale factor Synopsis:

```
glp::set_rii lp (rowindex, coefficient)
```

Parameters:

lp pointer to the LP problem object

rowindex row index

coefficient scaling coefficient

Returns:

()

Example:

```
> glp::set_rii lp (3, 258.6);  
()
```

Set the column scale factor Synopsis:

```
glp::set_sjj lp (colindex, coefficient)
```

Parameters:

lp pointer to the LP problem object

colindex column index

coefficient scaling coefficient

Returns:

()

Example:

```
> glp::set_sjj lp (2, 12.8);  
()
```

Retrieve the row scale factor Synopsis:

```
glp::get_rii lp rowindex
```

Parameters:

lp pointer to the LP problem object

rowindex row index

Returns:

scaling coefficient of given row

Example:

```
> glp::get_rii lp 3;  
258.6
```

Retrieve the column scale factor Synopsis:

```
glp::get_sjj lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

scaling coefficient of given column

Example:

```
> glp::get_sjj lp 2;  
12.8
```

Scale the problem data according to supplied flags Synopsis:

```
glp::scale_prob lp flags
```

Parameters:

lp pointer to the LP problem object

flags symbolic integer constants which can be combined together by arithmetic **or**; the possible constants are:

glp::sf_gm perform geometric mean scaling
glp::sf_eq perform equilibration scaling
glp::sf_2n round scale factors to power of two
glp::sf_skip skip if problem is well scaled
glp::sf_auto choose scaling options automatically

Returns:

()

Example:

```
> glp::scale_prob lp (glp::sf_gm || glp::sf_2n);  
()
```

Unscale the problem data Synopsis:

```
glp::unscale_prob lp
```

Parameters:

lp pointer to the LP problem object

Returns:

()

Example:

```
> glp::unscale_prob lp;  
()
```

LP basis constructing routines

Set the row status Synopsis:

```
glp::set_row_stat lp (rowindex, status)
```

Parameters:

lp pointer to the LP problem object

rowindex row index

status one of the following:

glp::bs make the row basic (make the constraint inactive)

glp::nl make the row non-basic (make the constraint active)

glp::nu make the row non-basic and set it to the upper bound; if the row is not double-bounded, this status is equivalent to **glp::nl** (only in the case of this routine)

glp::nf the same as **glp::nl** (only in the case of this routine)

glp::ns the same as **glp::nl** (only in the case of this routine)

Returns:

()

Example:

```
> glp::set_row_stat lp (3, glp::nu);
()
```

Set the column status Synopsis:

```
glp::set_col_stat lp (colindex, status)
```

Parameters:

lp pointer to the LP problem object

colindex column index

status one of the following:

glp::bs make the column basic

glp::nl make the column non-basic

glp::nu make the column non-basic and set it to the upper bound; if the column is not double-bounded, this status is equivalent to **glp::nl** (only in the case of this routine)

glp::nf the same as **glp::nl** (only in the case of this routine)

glp::ns the same as **glp::nl** (only in the case of this routine)

Returns:

()

Example:

```
> glp::set_col_stat lp (2, glp::bs);
()
```

Construct standard problem basis Synopsis:

```
glp::std_basis lp
```

Parameters:

`lp` pointer to the LP problem object

Returns:

`()`

Example:

```
> glp::std_basis lp;  
()
```

Construct advanced problem basis **Synopsis:**

```
glp::adv_basis lp
```

Parameters:

`lp` pointer to the LP problem object

Returns:

`()`

Example:

```
> glp::adv_basis lp;  
()
```

Construct Bixby's problem basis **Synopsis:**

```
glp::cpx_basis lp
```

Parameters:

`lp` pointer to the LP problem object

Returns:

`()`

Example:

```
> glp::cpx_basis lp;  
()
```

Simplex method routines

Solve the LP problem using simplex method **Synopsis:**

```
glp::simplex lp options
```

Parameters:

lp pointer to the LP problem object

options list of solver options in the form of tuples (option_name, value):

glp::msg_lev

(default: **glp::msg_all**) - message level for terminal output:

glp::msg_off: no output

glp::msg_err: error and warning messages only

glp::msg_on: normal output;

glp::msg_all: full output (including informational messages)

glp::meth (default: **glp::primal**) - simplex method option

glp::primal: use two-phase primal simplex

glp::dual: use two-phase dual simplex;

glp::dualp: use two-phase dual simplex, and if it fails, switch to the primal simplex

glp::pricing (default: **glp::pt_pse**) - pricing technique

glp::pt_std: standard (textbook)

glp::pt_pse: projected steepest edge

glp::r_test (default: **glp::rt_har**) - ratio test technique

glp::rt_std: standard (textbook)

glp::rt_har: Harris' two-pass ratio test

glp::tol_bnd (default: 1e-7) - tolerance used to check if the basic solution is primal feasible

glp::tol_dj (default: 1e-7) - tolerance used to check if the basic solution is dual feasible

glp::tol_piv (default: 1e-10) - tolerance used to choose eligible pivotal elements of the simplex table

glp::obj_ll (default: -DBL_MAX) - lower limit of the objective function - if the objective function reaches this limit and continues decreasing, the solver terminates the search - used in the dual simplex only

glp::obj_ul (default: +DBL_MAX) - upper limit of the objective function. If the objective function reaches this limit and continues increasing, the solver terminates the search - used in the dual simplex only

glp::it_lim (default: INT_MAX) - simplex iteration limit

glp::tm_lim (default: INT_MAX) - searching time limit, in milliseconds

glp::out_frq (default: 200) - output frequency, in iterations - this parameter specifies how frequently the solver sends information about the solution process to the terminal

glp::out_dly (default: 0) - output delay, in milliseconds - this parameter specifies how long the solver should delay sending information about the solution process to the terminal

glp::presolve (default: glp::off) - LP presolver option:

glp::on: enable using the LP presolver

glp::off: disable using the LP presolver

Returns:

one of the following:

glp::ok the LP problem instance has been successfully solved; this code does not necessarily mean that the solver has found optimal solution, it only means that the solution process was successful

glp::ebadb unable to start the search, because the initial basis specified in the problem object is invalid - the number of basic (auxiliary and structural) variables is not the same as the number of rows in the problem object

glp::esing unable to start the search, because the basis matrix corresponding to the initial basis is singular within the working precision

glp::econd unable to start the search, because the basis matrix corresponding to the initial basis is ill-conditioned, i.e. its condition number is too large

glp::ebound unable to start the search, because some double-bounded (auxiliary or structural) variables have incorrect bounds

glp::efail the search was prematurely terminated due to the solver failure

glp::eobjll the search was prematurely terminated, because the objective function being maximized has reached its lower limit and continues decreasing (the dual simplex only)

glp::eobjul the search was prematurely terminated, because the objective function being minimized has reached its upper limit and continues increasing (the dual simplex only)

glp::eitlim the search was prematurely terminated, because the simplex iteration limit has been exceeded

glp::etmlim the search was prematurely terminated, because the time limit has been exceeded

glp::enopfs the LP problem instance has no primal feasible solution (only if the LP presolver is used)

glp::enodfs the LP problem instance has no dual feasible solution (only if the LP presolver is used)

When the list of options contains some bad option(s) then a list of bad options is returned instead.

Remark:

Options not mentioned in the option list are set to their default values.

Example:

```
> glp::simplex lp [(glp::presolve, glp::on), (glp::msg_lev, glp::msg_all)];
glp_simplex: original LP has 3 rows, 3 columns, 9 non-zeros
glp_simplex: presolved LP has 3 rows, 3 columns, 9 non-zeros
Scaling...
  A: min|aij| = 1,000e+000  max|aij| = 1,000e+001  ratio = 1,000e+001
Problem data seem to be well scaled
Crashing...
Size of triangular part = 3
*    0: obj = 0,000000000e+000  infeas = 0,000e+000 (0)
*    2: obj = 7,333333333e+002  infeas = 0,000e+000 (0)
OPTIMAL SOLUTION FOUND
glp::ok
```

Solve the LP problem using simplex method in exact arithmetics **Synopsis:**

`glp::exact lp options`

Parameters:

lp pointer to the LP problem object

options list of solver options in the form of tuples (option_name, value):

glp::it_lim (default: INT_MAX) - simplex iteration limit

glp::tm_lim (default: INT_MAX) - searching time limit, in milliseconds

Returns:

one of the following:

glp::ok the LP problem instance has been successfully solved; this code does not necessarily mean that the solver has found optimal solution, it only means that the solution process was successful

glp::ebadb unable to start the search, because the initial basis specified in the problem object is invalid - the number of basic (auxiliary and structural) variables is not the same as the number of rows in the problem object

glp::esing unable to start the search, because the basis matrix corresponding to the initial basis is singular within the working precision

glp::ebound unable to start the search, because some double-bounded (auxiliary or structural) variables have incorrect bounds

glp::efail the search was prematurely terminated due to the solver failure

glp::eitlim the search was prematurely terminated, because the simplex iteration limit has been exceeded

glp::etmlim the search was prematurely terminated, because the time limit has been exceeded

When the list of options contains some bad option(s) then a list of bad options is returned instead.

Remark:

Options not mentioned in the option list are set to their default values.

Example:

```
> glp::exact lp [];  
glp_exact: 3 rows, 3 columns, 9 non-zeros  
GNU MP bignum library is being used
```



```
*      2:  objval =                0      (0)
*      4:  objval =      733,333333333333 (0)
OPTIMAL SOLUTION FOUND
glp::ok
```

Retrieve generic status of basic solution Synopsis:

```
glp::get_status lp
```

Parameters:

lp pointer to the LP problem object

Returns:

one of the following:

- glp::undef** solution is undefined
- glp::feas** solution is feasible
- glp::infeas** solution is infeasible
- glp::nofeas** no feasible solution exists
- glp::opt** solution is optimal
- glp::unbnd** solution is unbounded

Example:

```
> glp::get_status lp;
glp::opt
```

Retrieve generic status of primal solution Synopsis:

```
glp::get_prim_stat lp
```

Parameters:

lp pointer to the LP problem object

Returns:

one of the following:

- glp::undef** primal solution is undefined
- glp::feas** primal solution is feasible
- glp::infeas** primal solution is infeasible
- glp::nofeas** no primal feasible solution exists

Example:

```
> glp::get_prim_stat lp;  
glp::feas
```

Retrieve generic status of dual solution Synopsis:

```
glp::get_dual_stat lp
```

Parameters:

lp pointer to the LP problem object

Returns:

one of the following:

glp::undef dual solution is undefined

glp::feas dual solution is feasible

glp::infeas dual solution is infeasible

glp::nofeas no dual feasible solution exists

Example:

```
> glp::get_dual_stat lp;  
glp::feas
```

Retrieve value of the objective function Synopsis:

```
glp::get_obj_val lp
```

Parameters:

lp pointer to the LP problem object

Returns:

value of the objective function

Example:

```
> glp::get_obj_val lp  
733.333333333333
```

Retrieve generic status of a row variable Synopsis:

```
glp::get_row_stat lp rowindex
```

Parameters:

lp pointer to the LP problem object

rowindex row index

Returns:

one of the following:

glp::bs basic variable

glp::nl non-basic variable on its lower bound

glp::nu non-basic variable on its upper bound

glp::nf non-basic free (unbounded) variable

glp::ns non-basic fixed variable

Example:

```
> glp::get_row_stat lp 3;
glp::bs
```

Retrieve row primal value

Synopsis:: `glp::get_row_prim lp rowindex`

Parameters:

lp pointer to the LP problem object

rowindex row index

Returns:

primal value of the row (auxiliary) variable

Example:

```
> glp::get_row_prim lp 3;
200.0
```

Retrieve row dual value Synopsis:

`glp::get_row_dual lp rowindex`

Parameters:

lp pointer to the LP problem object

rowindex row index

Returns:

dual value of the row (auxiliary) variable

Example:

```
> glp::get_row_dual lp 3;  
0.0
```

Retrieve generic status of a column variable **Synopsis:**

```
glp::get_col_stat lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

one of the following:

glp::bs basic variable

glp::nl non-basic variable on its lower bound

glp::nu non-basic variable on its upper bound

glp::nf non-basic free (unbounded) variable

glp::ns non-basic fixed variable

Example:

```
> glp::get_col_stat lp 2;  
glp::bs
```

Retrieve column primal value **Synopsis:**

```
glp::get_col_prim lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

primal value of the column (structural) variable

Example:

```
> glp::get_col_prim lp 2;  
66.6666666666667
```

Retrieve column dual value Synopsis:

```
glp::get_col_dual lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

dual value of the column (structural) variable

Example:

```
> glp::get_col_dual lp 2;
0.0
```

Determine variable causing unboundedness Synopsis:

```
glp::get_unbnd_ray lp
```

Parameters:

lp pointer to the LP problem object

Returns:

The routine `glp_get_unbnd_ray` returns the number k of a variable, which causes primal or dual unboundedness. If $1 \leq k \leq m$, it is k -th auxiliary variable, and if $m + 1 \leq k \leq m + n$, it is $(k - m)$ -th structural variable, where m is the number of rows, n is the number of columns in the problem object. If such variable is not defined, the routine returns 0.

Remark:

If it is not exactly known which version of the simplex solver detected unboundedness, i.e. whether the unboundedness is primal or dual, it is sufficient to check the status of the variable with the routine `glp::get_row_stat` or `glp::get_col_stat`. If the variable is non-basic, the unboundedness is primal, otherwise, if the variable is basic, the unboundedness is dual (the latter case means that the problem has no primal feasible dolution).

Example:

```
> glp::get_unbnd_ray lp;
0
```

Interior-point method routines

Solve the LP problem using interior-point method Synopsis:

`glp::interior lp options`

Parameters:

lp pointer to the LP problem object

options list of solver options in the form of tuples (option_name, value):

glp::msg_lev

(default: **glp::msg_all**) - message level for terminal output:

glp::msg_off: no output

glp::msg_err: error and warning messages only

glp::msg_on: normal output;

glp::msg_all: full output (including informational messages)

glp::ord_alg (default: **glp::ord_amd**) - ordering algorithm option

glp::ord_none: use natural (original) ordering

glp::ord_qmd: quotient minimum degree (QMD)

glp::ord_amd: approximate minimum degree (AMD)

glp::ord_sysamd: approximate minimum degree (SYSAMD)

Returns:

one of the following:

glp::ok the LP problem instance has been successfully solved; this code does not necessarily mean that the solver has found optimal solution, it only means that the solution process was successful

glp::efail the problem has no rows/columns

glp::enocvg very slow convergence or divergence

glp::eitlim iteration limit exceeded

glp::einstab numerical instability on solving Newtonian system

Example:

```
> glp::interior lp [(glp::ord_alg, glp::ord_amd)];
Original LP has 3 row(s), 3 column(s), and 9 non-zero(s)
Working LP has 3 row(s), 6 column(s), and 12 non-zero(s)
Matrix A has 12 non-zeros
Matrix S = A*A' has 6 non-zeros (upper triangle)
Approximate minimum degree ordering (AMD)...
Computing Cholesky factorization S = L*L'...
Matrix L has 6 non-zeros
Guessing initial point...
Optimization begins...
 0: obj = -8,218489503e+002; rpi = 3,6e-001; rdi = 6,8e-001; gap = 2,5e-001
 1: obj = -6,719060895e+002; rpi = 3,6e-002; rdi = 1,9e-001; gap = 1,4e-002
 2: obj = -6,917210389e+002; rpi = 3,6e-003; rdi = 9,3e-002; gap = 3,0e-002
 3: obj = -7,267557732e+002; rpi = 2,1e-003; rdi = 9,3e-003; gap = 4,4e-002
 4: obj = -7,323038146e+002; rpi = 2,1e-004; rdi = 1,1e-003; gap = 4,8e-003
 5: obj = -7,332295932e+002; rpi = 2,1e-005; rdi = 1,1e-004; gap = 4,8e-004
 6: obj = -7,333229585e+002; rpi = 2,1e-006; rdi = 1,1e-005; gap = 4,8e-005
 7: obj = -7,333322959e+002; rpi = 2,1e-007; rdi = 1,1e-006; gap = 4,8e-006
 8: obj = -7,333332296e+002; rpi = 2,1e-008; rdi = 1,1e-007; gap = 4,8e-007
 9: obj = -7,333333230e+002; rpi = 2,1e-009; rdi = 1,1e-008; gap = 4,8e-008
10: obj = -7,333333323e+002; rpi = 2,1e-010; rdi = 1,1e-009; gap = 4,8e-009
OPTIMAL SOLUTION FOUND
glp::ok
```

Retrieve status of interior-point solution Synopsis:

```
glp::ipt_status lp
```

Parameters:

lp pointer to the LP problem object

Returns:

one of the following

glp::undef interior-point solution is undefined

glp::opt interior-point solution is optimal

glp::infeas interior-point solution is infeasible

glp::nofeas no feasible primal-dual solution exists

Example:

```
> glp::ipt_status lp;
glp::opt
```

Retrieve the objective function value of interior-point solution Synopsis:

```
glp::ipt_obj_val lp
```

Parameters:

lp pointer to the LP problem object

Returns:

objective function value of interior-point solution

Example:

```
> glp::ipt_obj_val lp;  
733.333332295849
```

Retrieve row primal value of interior-point solution **Synopsis:**

```
glp::ipt_row_prim lp rowindex
```

Parameters:

lp pointer to the LP problem object

rowindex row index

Returns:

primal value of the row (auxiliary) variable

Example:

```
> glp::ipt_row_prim lp 3;  
200.000000920688
```

Retrieve row dual value of interior-point solution **Synopsis:**

```
glp::ipt_row_dual lp rowindex
```

Parameters:

lp pointer to the LP problem object

rowindex row index

Returns:

dual value of the row (auxiliary) variable

Example:

```
> glp::ipt_row_dual lp 3;  
2.50607466186742e-008
```


Retrieve column primal value of interior-point solution Synopsis:

```
glp::ipt_col_prim lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

primal value of the column (structural) variable

Example:

```
> glp::ipt_col_prim lp 2;  
66.666666406779
```

Retrieve column dual value of interior-point solution Synopsis:

```
glp::ipt_col_dual lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

dual value of the column (structural) variable

Example:

```
> glp::ipt_col_dual lp 2;  
2.00019467655466e-009
```

Mixed integer programming routines

Set column kind Synopsis:

```
glp::set_col_kind lp (colindex, colkind)
```

Parameters:

lp pointer to the LP problem object

colindex column index

colkind column kind - one of the following:

glp::cv continuous variable

glp::iv integer variable

glp::bv binary variable

Returns:

()

Example:

```
> glp::set_col_kind lp (1, glp::iv);  
()
```

Retrieve column kind Synopsis:

`glp::get_col_kind lp colindex`

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

one of the following:

glp::cv continuous variable

glp::iv integer variable

glp::bv binary variable

Example:

```
> glp::get_col_kind lp 1;  
glp::iv
```

Retrieve number of integer columns Synopsis:

`glp::get_num_int lp`

Parameters:

lp pointer to the LP problem object

Returns:

number of integer columns (including binary columns)

Example:

```
> glp_get_num_int lp;  
1
```

Retrieve number of binary columns Synopsis:

```
glp::get_num_bin lp
```

Parameters:

lp pointer to the LP problem object

Returns:

number of binary columns

Example:

```
> glp::get_num_bin lp
0
```

Solve the MIP problem using branch-and-cut method Synopsis:

```
glp::intopt lp options
```

Parameters:

lp pointer to the LP problem object

options list of solver options in the form of tuples (option_name, value):

glp::msg_lev

(default: **glp::msg_all**) - message level for terminal output:

glp::msg_off: no output

glp::msg_err: error and warning messages only

glp::msg_on: normal output;

glp::msg_all: full output (including informational messages)

glp::br_tech (default: **glp::bt::blb**) - branching technique

glp::br_ffv: first fractional variable

glp::br_lfv: last fractional variable

glp::br_mfv: most fractional variable

glp::br_dth: heuristic by Driebeck and Tomlin

glp::br_pch: hybrid pseudocost heuristic

glp::bt_tech (default: **glp::pt_pse**) - backtracking technique

glp::bt_dfs: depth first search;
glp::bt_bfs: breadth first search;
glp::bt_blb: best local bound;
glp::bt_bph: best projection heuristic.

glp::pp_tech (default: **glp::pp_all**) - preprocessing technique

glp::pp_none: disable preprocessing;
glp::pp_root: perform preprocessing only on the root level
glp::pp_all: perform preprocessing on all levels

glp::fp_heur (default: **glp::off**) - feasibility pump heuristic:

glp::on: enable applying the feasibility pump heuristic
glp::off: disable applying the feasibility pump heuristic

glp::gmi_cuts

(default: **glp::off**) - **Gomory's mixed integer** cuts:

glp::on: enable generating Gomory's cuts;
glp::off: disable generating Gomory's cuts.

glp::mir_cuts

(default: **glp::off**) - **mixed integer rounding** (MIR) cuts:

glp::on: enable generating MIR cuts;
glp::off: disable generating MIR cuts.

glp::cov_cuts (default: `glp::off`) - mixed cover cuts:

glp::on: enable generating mixed cover cuts;

glp::off: disable generating mixed cover cuts.

glp::clq_cuts (default `glp::off`) - clique cuts:

glp::on: enable generating clique cuts;

glp::off: disable generating clique cuts.

glp::tol_int (default: `1e-5`) - absolute tolerance used to check if optimal solution to the current LP relaxation is integer feasible

glp::tol_obj (default: `1e-7`) - relative tolerance used to check if the objective value in optimal solution to the current LP relaxation is not better than in the best known integer feasible solution

glp::mip_gap (default: `0.0`) - the relative mip gap tolerance; if the relative mip gap for currently known best integer feasible solution falls below this tolerance, the solver terminates the search - this allows obtaining sub-optimal integer feasible solutions if solving the problem to optimality takes too long time

glp::tm_lim (default: `INT_MAX`) - searching time limit, in milliseconds

glp::out_frq (default: `5000`) - output frequency, in milliseconds - this parameter specifies how frequently the solver sends information about the solution process to the terminal

glp::out_dly (default: `10000`) - output delay, in milliseconds - this parameter specifies how long the solver should delay sending information about the solution of the current LP relaxation with the simplex method to the terminal

glp::cb_func

(default: `glp::off`) - specifies whether to use the user-defined callback routine

glp::on: use user-defined callback function - the function `glp::mip_cb tree info` **must** be defined by the user

glp::off: do not use user-defined callback function

glp::cb_info (default: `NULL`) - transit pointer passed to the routine `glp::mip_cb tree info` (see above)

glp::cb_size (default: 0) - the number of extra (up to 256) bytes allocated for each node of the branch-and-bound tree to store application-specific data - on creating a node these bytes are initialized by binary zeros

glp::presolve (default: `glp::off`) - LP presolver option:

glp::on: enable using the MIP presolver

glp::off: disable using the MIP presolver

glp::binarize

(default: `glp::off`) - **binarization** (used only if the presolver is enabled):

glp::on: replace general integer variables by binary ones

glp::off: do not use binarization

Returns:

one of the following:

glp::ok the MIP problem instance has been successfully solved; this code does not necessarily mean that the solver has found optimal solution, it only means that the solution process was successful

glp::ebound unable to start the search, because some double-bounded (auxiliary or structural) variables have incorrect bounds or some integer variables have non-integer (fractional) bounds

glp::eroot unable to start the search, because optimal basis for initial LP relaxation is not provided - this code may appear only if the presolver is disabled

glp::enopfs unable to start the search, because LP relaxation of the MIP problem instance has no primal feasible solution - this code may appear only if the presolver is enabled

glp::enodfs unable to start the search, because LP relaxation of the MIP problem instance has no dual feasible solution; in other word, this code means that if the LP relaxation has at least one primal feasible solution, its optimal solution is unbounded, so if the MIP problem has at least one integer feasible solution, its (integer) optimal solution is also unbounded - this code may appear only if the presolver is enabled

glp::efail the search was prematurely terminated due to the solver failure

glp::emipgap the search was prematurely terminated, because the relative mip gap tolerance has been reached

glp::etmlim the search was prematurely terminated, because the time limit has been exceeded

glp::estop the search was prematurely terminated by application - this code may appear only if the advanced solver interface is used

When the list of options contains some bad option(s) then a list of bad options is returned instead.

Remark:

Options not mentioned in the option list are set to their default values.

Example:

```
> glp::intopt lp [(glp::presolve, glp::on)];
ipp_basic_tech: 0 row(s) and 0 column(s) removed
ipp_reduce_bnds: 2 pass(es) made, 3 bound(s) reduced
ipp_basic_tech: 0 row(s) and 0 column(s) removed
ipp_reduce_coef: 1 pass(es) made, 0 coefficient(s) reduced
glp_intopt: presolved MIP has 3 rows, 3 columns, 9 non-zeros
glp_intopt: 3 integer columns, none of which are binary
Scaling...
  A: min|aij| = 1,000e+00 max|aij| = 1,000e+01 ratio = 1,000e+01
Problem data seem to be well scaled
Crashing...
Size of triangular part = 3
Solving LP relaxation...
*   2: obj = 0,000000000e+00 infeas = 0,000e+00 (0)
*   5: obj = 7,333333333e+02 infeas = 0,000e+00 (0)
OPTIMAL SOLUTION FOUND
Integer optimization begins...
+   5: mip = not found yet <= +inf (1; 0)
+   6: >>>> 7,320000000e+02 <= 7,320000000e+02 0.0% (2; 0)
+   6: mip = 7,320000000e+02 <= tree is empty 0.0% (0; 3)
INTEGER OPTIMAL SOLUTION FOUND
glp::ok
```

Retrieve status of mip solution Synopsis:

```
glp::mip_status lp
```

Parameters:

lp pointer to the LP problem object

Returns:

one of the following:

glp::undef MIP solution is undefined

glp::opt MIP solution is integer optimal

glp::feas MIP solution is integer feasible, however, its optimality (or non-optimality) has not been proven, perhaps due to premature termination of the search

glp::nofeas problem has no integer feasible solution (proven by the solver)

Example:

```
> glp::mip_status lp;  
glp::opt
```

Retrieve the objective function value of mip solution Synopsis:

```
glp::mip_obj_val lp
```

Parameters:

lp pointer to the LP problem object

Returns:

objective function value of mip solution

Example:

```
> glp::mip_obj_val lp;  
732.0
```

Retrieve row value of mip solution Synopsis:

```
glp::mip_row_val lp rowindex
```

Parameters:

lp pointer to the LP problem object

rowindex row index

Returns:

row value (value of auxiliary variable)

Example:

```
> glp::mip_row_val lp 3;
200.0
```

Retrieve column value of mip solution **Synopsis:**

```
glp::mip_col_val lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

column value (value of structural variable)

Example:

```
> glp::mip_col_val lp 2;
67.0
```

Additional routines

Check Karush-Kuhn-Tucker conditions **Synopsis:**

```
lpx::check_kkt lp scaled
```

Parameters:

lp pointer to the LP problem object

scaled one of the following:

true test the scaled problem

false test the unscaled problem

Returns:

list of four tuples with five members (see GLPK reference manual):

Condition	Member	Comment
(KKT.PE)	pe_ae_max	Largest absolute error
	pe_ae_row	Number of row with largest absolute error
	pe_re_max	Largest relative error
	pe_re_row	Number of row with largest relative error
(KKT.PB)	pe_quality	Quality of primal solution
	pb_ae_max	Largest absolute error
	pb_ae_ind	Number of variable with largest absolute error
	pb_re_max	Largest relative error
(KKT.DE)	pb_re_ind	Number of variable with largest relative error
	pb_quality	Quality of primal feasibility
	de_ae_max	Largest absolute error
	de_ae_col	Number of column with largest absolute error
(KKT.DB)	de_re_max	Largest relative error
	de_re_col	Number of column with largest relative error
	de_quality	Quality of dual solution
	db_ae_max	Largest absolute error
	db_ae_ind	Number of variable with largest absolute error
	db_re_max	Largest relative error
	db_re_ind	Number of variable with largest relative error
	db_quality	Quality of dual feasibility

where number of variable is $(1 \leq k \leq m)$ for auxiliary variable and $(m+1 \leq k \leq m+n)$ for structural variable

Example:

```
> lpx::check_kkt lp true;
[(1.4210854715202e-14,2,3.54385404369127e-17,3,"H"),(0.0,0,0.0,0,"H"),
(4.44089209850063e-16,1,2.11471052309554e-17,1,"H"),(0.0,0,0.0,0,"H")]
```

5.2 Utility API routines

Problem data reading/writing routines

Read LP problem data from a MPS file Synopsis:

```
glp::read_mps lp format filename
```

Parameters:

lp pointer to the LP problem object

format one of the following

glp::mps_deck fixed (ancient) MPS file format

glp::mps_file free (modern) MPS file format

filename file name - if the file name ends with suffix **.gz**, the file is assumed to be compressed, in which case the routine `glp::read_mps` decompresses it “on the fly”

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> glp::read_mps lp glp::mps_deck "examples/plan.mps";
Reading problem data from 'examples/plan.mps'...
Problem PLAN
Objective R00000000
8 rows, 7 columns, 55 non-zeros
63 records were read
0
```

Write LP problem data into a MPS file Synopsis:

```
glp::write_mps lp format filename
```

Parameters:

lp pointer to the LP problem object

format one of the following

glp::mps_deck fixed (ancient) MPS file format

glp::mps_file free (modern) MPS file format

filename file name - if the file name ends with suffix **.gz**, the file is assumed to be compressed, in which case the routine `glp_write_mps` performs automatic compression on writing it

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::write_mps lp glp::mps_file "examples/plan1.mps";
Writing problem data to 'examples/plan1.mps'...
63 records were written
0
```

Read LP problem data from a CPLEX file Synopsis:

```
glp::read_lp lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name - if the file name ends with suffix **.gz**, the file is assumed to be compressed, in which case the routine `glp::read_lp` decompresses it “on the fly”

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::read_lp lp "examples/plan.lp";  
reading problem data from 'examples/plan.lp'...  
8 rows, 7 columns, 48 non-zeros  
39 lines were read  
0
```

Write LP problem data into a CPLEX file Synopsis:

```
glp::write_lp lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name - if the file name ends with suffix **.gz**, the file is assumed to be compressed, in which case the routine `glp::write_lp` performs automatic compression on writing it

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::write_lp lp "examples/plan1.lp";  
writing problem data to 'examples/plan1.lp'...  
29 lines were written  
0
```

Read LP problem data in GLPK format Synopsis:

```
glp::read_prob lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name - if the file name ends with suffix **.gz**, the file is assumed to be compressed, in which case the routine `glp::read_prob` decompresses it “on the fly”

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::read_prob lp "examples/plan.glpk";
reading problem data from 'examples/plan.glpk'...
8 rows, 7 columns, 48 non-zeros
86 lines were read
0
```

Write LP problem data in GLPK format Synopsis:

```
glp::write_prob lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name - if the file name ends with suffix **.gz**, the file is assumed to be compressed, in which case the routine `glp::write_prob` performs automatic compression on writing it

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::write_prob lp "examples/plan1.glpk";
writing problem data to 'examples/plan1.glpk'...
86 lines were written
0
```

Routines for MathProg models

Create the MathProg translator object Synopsis:

```
glp::mpl_alloc_wksp
```

Parameters:

none

Returns:

pointer to the MathProg translator object

Example:

```
> let mpt = glp::mpl_alloc_wksp;
> mpt;
#<pointer 0xa0d0180>
```

Read and translate model section Synopsis:

```
glp::mpl_read_model tranobject filename skip
```

Parameters:

tranobject pointer to the MathProg translator object

filename file name

skip if 0 then the data section from the model file is read; if non-zero,
the data section in the data model is skipped

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> mpl_read_model mpt "examples/sudoku.mod" 1;
Reading model section from examples/sudoku.mod...
examples/sudoku.mod:69: warning: data section ignored
69 lines were read
0
```

Read and translate data section Synopsis:

```
glp::mpl_read_data tranobject filename
```

Parameters:

tranobject pointer to the MathProg translator object

filename file name

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> glp::mpl_read_data mpt "examples/sudoku.dat";
Reading data section from examples/sudoku.dat...
16 lines were read
0
```

Generate the model Synopsis:

```
glp::mpl_generate tranobject filename
```

Parameters:

tranobject pointer to the MathProg translator object

filename file name

Returns:

0 if generating went OK; non-zero in case of an error

Example:

```
> glp::mpl_generate mpt "examples/sudoku.lst";
Generating fa...
Generating fb...
Generating fc...
Generating fd...
Generating fe...
Model has been successfully generated
0
```

Build problem instance from the model Synopsis:

```
glp::mpl_build_prob tranobject lp
```

Parameters:

tranobject pointer to the MathProg translator object

lp pointer to the LP problem object

Returns:

()

Example:

```
> glp::mpl_build_prob mpt lp;
()
```

Postsolve the model Synopsis:

```
glp::mpl_postsolve tran lp solution
```

Parameters:

tranobject pointer to the MathProg translator object

lp pointer to the LP problem object

solution one of the following:

glp::sol use the basic solution

glp::ipt use the interior-point solution

glp::mip use mixed integer solution

Returns:

0 if postsolve went OK; non-zero in case of an error

Example:

```
> glp::mpl_postsolve mpt lp glp::sol;
Model has been successfully processed
0
```

Delete the MathProg translator object Synopsis:

```
glp::mpl_free_wksp tranobject
```

Parameters:

tranobject pointer to the MathProg translator object

Returns:

()

Example:

```
> glp::mpl_free_wksp mpt;
()
```

Problem solution reading/writing routines

Write basic solution in printable format Synopsis:

```
glp::print_sol lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::print_sol lp "examples/test.txt";
Writing basic solution to 'examples/test.txt'...
0
```

Read basic solution from a text file Synopsis:

```
glp::read_sol lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> glp::read_sol lp "examples/test.txt";
Reading basic solution from 'examples/test.txt'...
1235 lines were read
0
```

Write basic solution into a text file Synopsis:

```
glp::write_sol lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::write_sol lp "examples/test.txt";
Writing basic solution to 'examples/test.txt'...
1235 lines were written
0
```

Print sensitivity analysis report Synopsis:

```
glp::print_ranges lp indices filename
```

Parameters:

lp pointer to the LP problem object

indices list indices k of rows and columns to be included in the report. If $1 \leq k \leq m$, the basic variable is k -th auxiliary variable, and if $m + 1 \leq k \leq m + n$, the non-basic variable is $(k - m)$ -th structural variable, where m is the number of rows and n is the number of columns in the specified problem object. An empty lists means printing report for all rows and columns.

filename file name

Returns:

0 if the operation was successful

non-zero if the operation failed

Example:

```
> glp::print_ranges lp [] "sensitivity.rpt";
Write sensitivity analysis report to 'sensitivity.rpt'...
0
```

Write interior-point solution in printable format Synopsis:

```
glp::print_ipt lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::print_ipt lp "examples/test.txt";
Writing interior-point solution to 'examples/test.txt'...
0
```

Read interior-point solution from a text file Synopsis:

```
glp::read_ipt lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> glp::read_ipt lp "examples/test.txt";
Reading interior-point solution from 'examples/test.txt'...
1235 lines were read
0
```

Write interior-point solution into a text file Synopsis:

```
glp::write_ipt lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::write_ipt lp "examples/test.txt";
Writing interior-point solution to 'examples/test.txt'...
1235 lines were written
0
```

Write MIP solution in printable format Synopsis:

```
glp::print_mip lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::print_mip lp "examples/test.txt";
Writing MIP solution to 'examples/test.txt'...
0
```

Read MIP solution from a text file Synopsis:

```
glp::read_mip lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> glp::read_mip lp "examples/test.txt";
Reading MIP solution from 'examples/test.txt'...
1235 lines were read
0
```

Write MIP solution into a text file Synopsis:

```
glp::write_mip lp filename
```

Parameters:

lp pointer to the LP problem object

filename file name

Returns:

0 if writing went OK; non-zero in case of an error

Example:

```
> glp::write_mip lp "examples/test.txt";
Writing MIP solution to 'examples/test.txt'...
1235 lines were written
0
```

5.3 Advanced API routines

LP basis routines

Check whether basis factorization exists Synopsis:

```
glp::bf_exists lp
```

Parameters:

lp pointer to the LP problem object

Returns:

non-zero the basis factorization exists and can be used for calculations

0 the basis factorization does not exist

Example:

```
> glp::bf_exists lp;
1
```

Compute the basis factorization Synopsis:

```
glp::factorize lp
```

Parameters:

lp pointer to the LP problem object

Returns:

one of the following:

glp::ok the basis factorization has been successfully computed

glp::ebadb the basis matrix is invalid, because the number of basic (auxiliary and structural) variables is not the same as the number of rows in the problem object

glp::esing the basis matrix is singular within the working precision

glp::exond the basis matrix is ill-conditioned, i.e. its condition number is too large

Example:

```
> glp::factorize lp;  
glp::ok
```

Check whether basis factorization has been updated Synopsis:

```
glp::bf_updated lp
```

Parameters:

lp pointer to the LP problem object

Returns:

0 if the basis factorization has been just computed from “scratch”

non-zero if the factorization has been updated at least once

Example:

```
> glp::bf_updated lp;  
0
```

Get basis factorization parameters Synopsis:

```
glp::get_bfcp lp
```

Parameters:

lp pointer to the LP problem object

Returns:

complete list of options in a form of tuples (option_name, value):

glp::type basis factorization type:

glp::bf_ft LU + Forrest–Tomlin update

glp::bf_bg LU + Schur complement + Bartels–Golub update

glp::bf_gr LU + Schur complement + Givens rotation update

glp::lu_size the initial size of the Sparse Vector Area, in non-zeros, used on computing LU-factorization of the basis matrix for the first time - if this parameter is set to 0, the initial SVA size is determined automatically

glp::piv_tol threshold pivoting (Markowitz) tolerance, $0 < \text{piv_tol} < 1$, used on computing LU-factorization of the basis matrix

glp::piv_lim this parameter is used on computing LU-factorization of the basis matrix and specifies how many pivot candidates needs to be considered on choosing a pivot element, $\text{piv_lim} \geq 1$

glp::suhl this parameter is used on computing LU-factorization of the basis matrix

glp::on enables applying the heuristic proposed by Uwe Suhl

glp::off disables this heuristic

glp::eps_tol epsilon tolerance, $\text{eps_tol} \geq 0$, used on computing LU-factorization of the basis matrix

glp::max_gro maximal growth of elements of factor U, $\text{max_gro} \geq 1$, allowable on computing LU-factorization of the basis matrix

glp::nfs_max maximal number of additional row-like factors (entries of the eta file), $\text{nfs_max} \geq 1$, which can be added to LU-factorization of the basis matrix on updating it with the Forrest–Tomlin technique

glp::upd_tol update tolerance, $0 < \text{upd_tol} < 1$, used on updating LU-factorization of the basis matrix with the Forrest–Tomlin technique

glp::nrs_max maximal number of additional rows and columns, $\text{nrs_max} \geq 1$, which can be added to LU-factorization of the basis matrix on updating it with the Schur complement technique

glp::rs_size the initial size of the Sparse Vector Area, in non-zeros, used to store non-zero elements of additional rows and columns introduced on updating LU-factorization of the basis matrix with

the Schur complement technique - if this parameter is set to 0, the initial SVA size is determined automatically

Example:

```
> glp::get_bfcp lp;
[(glp::type,glp::bf_ft),(glp::lu_size,0),(glp::piv_tol,0.1),(glp::piv_lim,4),
 (glp::suhl,glp::on),(glp::eps_tol,1e-15),(glp::max_gro,10000000000.0),
 (glp::nfs_max,50),(glp::upd_tol,1e-06),(glp::nrs_max,50),(glp::rs_size,0)]
```

Change basis factorization parameters Synopsis:

```
glp::set_bfcp lp options
```

Parameters:

lp pointer to the LP problem object

options list of options in a form of tuples (option_name, value):

glp::type (default: `glp::bf_ft`) - basis factorization type:

glp::bf_ft LU + Forrest–Tomlin update

glp::bf_bg LU + Schur complement + Bartels–Golub update

glp::bf_gr LU + Schur complement + Givens rotation update

glp::lu_size (default: 0) - the initial size of the Sparse Vector Area, in non-zeros, used on computing LU-factorization of the basis matrix for the first time - if this parameter is set to 0, the initial SVA size is determined automatically

glp::piv_tol (default: 0.10) - threshold pivoting (Markowitz) tolerance, $0 < \text{piv_tol} < 1$, used on computing LU-factorization of the basis matrix.

glp::piv_lim (default: 4) - this parameter is used on computing LU-factorization of the basis matrix and specifies how many pivot candidates needs to be considered on choosing a pivot element, $\text{piv_lim} \geq 1$

glp::suhl (default: `glp::on`) - this parameter is used on computing LU-factorization of the basis matrix.

glp::on enables applying the heuristic proposed by Uwe Suhl

glp::off disables this heuristic

glp::eps_tol (default: 1e-15) - epsilon tolerance, $\text{eps_tol} \geq 0$, used on computing LU -factorization of the basis matrix.

glp::max_gro (default: 1e+10) - maximal growth of elements of factor U , $\text{max_gro} \geq 1$, allowable on computing LU-factorization of the basis matrix.

glp::nfs_max (default: 50) - maximal number of additional row-like factors (entries of the eta file), $\text{nfs_max} \geq 1$, which can be added to LU-factorization of the basis matrix on updating it with the Forrest-Tomlin technique.

glp::upd_tol (default: 1e-6) - update tolerance, $0 < \text{upd_tol} < 1$, used on updating LU -factorization of the basis matrix with the Forrest-Tomlin technique.

glp::nrs_max (default: 50) - maximal number of additional rows and columns, $\text{nrs_max} \geq 1$, which can be added to LU-factorization of the basis matrix on updating it with the Schur complement technique.

glp::rs_size (default: 0) - the initial size of the Sparse Vector Area, in non-zeros, used to store non-zero elements of additional rows and columns introduced on updating LU-factorization of the basis matrix with the Schur complement technique - if this parameter is set to 0, the initial SVA size is determined automatically

Remarks:

Options not mentioned in the option list are left unchanged.

All options will be reset to their default values when an empty option list is supplied.

Returns:

() if all options are OK, otherwise returns a list of bad options

Example:

```
> glp_set_bfcp lp [(glp::type, glp::bf_ft), (glp::piv_tol, 0.15)];
()
```

Retrieve the basis header information Synopsis:

```
glp::get_bhead lp k
```

Parameters:

lp pointer to the LP problem object

k variable index in the basis matrix

Returns:

If basic variable $(xB)_k$, $1 \leq k \leq m$, is i -th auxiliary variable ($1 \leq i \leq m$), the routine returns i . Otherwise, if $(xB)_k$ is j -th structural variable ($1 \leq j \leq n$), the routine

returns $m+j$. Here m is the number of rows and n is the number of columns in the problem object.

Example:

```
> glp::get_bhead lp 3;  
5
```

Retrieve row index in the basis header Synopsis:

```
glp::get_row_bind lp rowindex
```

Parameters:

lp pointer to the LP problem object

rowindex row index

Returns:

This routine returns the index k of basic variable $(xB)_k$, $1 \leq k \leq m$, which is i -th auxiliary variable (that is, the auxiliary variable corresponding to i -th row), $1 \leq i \leq m$, in the current basis associated with the specified problem object, where m is the number of rows. However, if i -th auxiliary variable is non-basic, the routine returns zero.

Example:

```
> glp::get_row_bind lp 3;  
1
```

Retrieve column index in the basis header Synopsis:

```
glp::get_col_bind lp colindex
```

Parameters:

lp pointer to the LP problem object

colindex column index

Returns:

This routine returns the index k of basic variable $(xB)_k$, $1 \leq k \leq m$, which is j -th structural variable (that is, the structural variable corresponding to j -th column), $1 \leq j \leq n$, in the current basis associated with the specified problem object, where m is the number of rows, n is the number of columns. However, if j -th structural variable is non-basic, the routine returns zero.

Example:

```
> glp::get_col_bind lp 2;  
3
```

Perform forward transformation Synopsis:

```
glp::ftran lp vector
```

Parameters:

lp pointer to the LP problem object

vector vector to be transformed - a dense vector in a form of a list of double numbers has to be supplied and the number of its members must exactly correspond to the number of LP problem constraints

Returns:

the transformed vector in the same format

Example:

```
> glp::ftran lp [1.5, 3.2, 4.8];  
[1.8,0.466666666666667,-1.96666666666667]
```

Perform backward transformation Synopsis:

```
glp::btran lp vector
```

Parameters:

lp pointer to the LP problem object

vector vector to be transformed - a dense vector in a form of a list of double numbers has to be supplied and the number of its members must exactly correspond to the number of LP problem constraints

Returns:

the transformed vector in the same format

Example:

```
> glp::btran lp [1.5, 3.2, 4.8];  
[-8.86666666666667,0.26666666666667,1.5]
```

Warm up LP basis Synopsis:

```
glp::warm_up lp
```

Parameters:

lp pointer to the LP problem object

Returns:

one of the following:

glp::ok the LP basis has been successfully “warmed up”

glp::ebadb the LP basis is invalid, because the number of basic variables is not the same as the number of rows

glp::esing the basis matrix is singular within the working precision

glp::econd the basis matrix is ill-conditioned, i.e. its condition number is too large

Example:

```
> glp::warm_up lp;
glp::e_ok
```

Simplex tableau routines

Compute row of the tableau Synopsis:

```
glp::eval_tab_row lp k
```

Parameters:

lp pointer to the LP problem object

k variable index such that it corresponds to some basic variable: if $1 \leq k \leq m$, the basic variable is k -th auxiliary variable, and if $m + 1 \leq k \leq m + n$, the basic variable is $(k - m)$ -th structural variable, where m is the number of rows and n is the number of columns in the specified problem object (the basis factorization must exist)

Returns:

simplex tableau row in a sparse form as a list of tuples (index, value), where index has the same meaning as k in parameters

Example:

```
> glp::eval_tab_row lp 3;
[(1,2.0),(6,4.0)]
```

Compute column of the tableau Synopsis:

```
glp::eval_tab_col lp k
```

Parameters:

lp pointer to the LP problem object

k variable index such that it corresponds to some non-basic variable: if $1 \leq k \leq m$, the non-basic variable is k -th auxiliary variable, and if $m + 1 \leq k \leq m + n$, the non-basic variable is $(k - m)$ -th structural variable, where m is the number of rows and n is the number

of columns in the specified problem object (the basis factorization must exist)

Returns:

simplex tableau column in a sparse form as a list of tuples (index, value), where index has the same meaning as *k* in parameters

Example:

```
> glp::eval_tab_col lp 1;
[(3,2.0),(4,-0.666666666666667),(5,1.666666666666667)]
```

Transform explicitly specified row Synopsis:

```
glp::transform_row lp rowvector
```

Parameters:

lp pointer to the LP problem object

rowvector row vector to be transformed in a sparse form as a list of tuples (*k*, value): if $1 \leq k \leq m$, the non-basic variable is *k*-th auxiliary variable, and if $m + 1 \leq k \leq m + n$, the non-basic variable is (*k* − *m*)-th structural variable, where *m* is the number of rows and *n* is the number of columns in the specified problem object (the basis factorization must exist)

Returns:

the transformed row in a sparse form as a list of tuples (index, value), where index has the same meaning as *k* in parameters

Example:

```
> glp::transform_row lp [(1, 3.0), (2, 3.5)];
[(1,3.833333333333333),(2,-0.0833333333333333),(6,-3.41666666666667)]
```

Transform explicitly specified column Synopsis:

```
glp::transform_col lp colvector
```

Parameters:

lp pointer to the LP problem object

colvector column vector to be transformed in a sparse form as a list of tuples (*k*, value): if $1 \leq k \leq m$, the non-basic variable is *k*-th auxiliary variable, and if $m + 1 \leq k \leq m + n$, the non-basic variable is (*k* − *m*)-th structural variable, where *m* is the number of rows and *n* is the number of columns in the specified problem object (the basis factorization must exist)

Returns:

the transformed column in a sparse form as a list of tuples (index, value), where index has the same meaning as k in parameters

Example:

```
> glp::transform_col lp [(2, 1.0), (3, 2.3)];
[(3,2.3), (4, -0.1666666666666667), (5, 0.1666666666666667)]
```

Perform primal ratio test Synopsis:

```
glp::prim_rtest lp colvector dir eps
```

Parameters:

lp pointer to the LP problem object

colvector simplex tableau column in a sparse form as a list of tuples (k , value): if $1 \leq k \leq m$, the basic variable is k -th auxiliary variable, and if $m + 1 \leq k \leq m + n$, the basic variable is $(k - m)$ -th structural variable, where m is the number of rows and n is the number of columns in the specified problem object (the basis factorization must exist and the primal solution must be feasible)

dir specifies in which direction the variable y changes on entering the basis: $+1$ means increasing, -1 means decreasing

eps relative tolerance (small positive number) used to skip small values in the column

Returns:

The routine returns the index, piv , in the colvector corresponding to the pivot element chosen, $1 \leq piv \leq len$. If the adjacent basic solution is primal unbounded, and therefore the choice cannot be made, the routine returns zero.

Example:

```
> glp::prim_rtest lp [(3, 2.5), (5, 7.0)] 1 1.0e-5;
3
```

Perform dual ratio test Synopsis:

```
glp::dual_rtest lp rowvector dir eps
```

Parameters:

lp pointer to the LP problem object

rowvector simplex tableau row in a sparse form as a list of tuples (k , value): if $1 \leq k \leq m$, the non-basic variable is k -th auxiliary variable, and if $m + 1 \leq k \leq m + n$, the non-basic variable is $(k -$

m)-th structural variable, where m is the number of rows and n is the number of columns in the specified problem object (the basis factorization must exist and the dual solution must be feasible)

dir specifies in which direction the variable y changes on leaving the basis: +1 means increasing, -1 means decreasing

eps relative tolerance (small positive number) used to skip small values in the row

Returns:

The routine returns the index, piv, in the rowvector corresponding to the pivot element chosen, $1 \leq \text{piv} \leq \text{len}$. If the adjacent basic solution is dual unbounded, and therefore the choice cannot be made, the routine returns zero.

Example:

```
> glp::dual_rtest lp [(1, 1.5), (6, 4.0)] 1 1.0e-5;
6
```

Analyze active bound of non-basic variable Synopsis:

```
glp::analyze_bound lp k
```

Parameters:

lp pointer to the LP problem object

k if $1 \leq k \leq m$, the non-basic variable is k-th auxiliary variable, and if $m + 1 \leq k \leq m + n$, the non-basic variable is (k - m)-th structural variable, where m is the number of rows and n is the number of columns in the specified problem object (the basis factorization must exist and the solution must be optimal)

Returns:

The routine returns a tuple (limit1, var1, limit2 var2) where:

value1 the minimal value of the active bound, at which the basis still remains primal feasible and thus optimal. -DBL_MAX means that the active bound has no lower limit.

var1 the ordinal number of an auxiliary (1 to m) or structural (m + 1 to m + n) basic variable, which reaches its bound first and thereby limits further decreasing the active bound being analyzed. If value1 = -DBL_MAX, var1 is set to 0.

value2 the maximal value of the active bound, at which the basis still remains primal feasible and thus optimal. +DBL_MAX means that the active bound has no upper limit.

var2 the ordinal number of an auxiliary (1 to m) or structural (m + 1 to m + n) basic variable, which reaches its bound first and thereby limits further increasing the active bound being analyzed. If value2 = +DBL_MAX, var2 is set to 0.

Example:

```
> analyze_bound lp 2;
1995.06864446899,12,2014.03478832467,4
```

Analyze objective coefficient at basic variable Synopsis:

```
glp::analyze_coef lp k
```

Parameters:

lp pointer to the LP problem object

k if $1 \leq k \leq m$, the basic variable is k-th auxiliary variable, and if $m + 1 \leq k \leq m + n$, the non-basic variable is (k - m)-th structural variable, where m is the number of rows and n is the number of columns in the specified problem object (the basis factorization must exist and the solution must be optimal)

Returns:

The routine returns a tuple (coef1, var1, value1, coef2 var2, value2) where:

coef1 the minimal value of the objective coefficient, at which the basis still remains dual feasible and thus optimal. -DBL_MAX means that the objective coefficient has no lower limit.

var1 is the ordinal number of an auxiliary (1 to m) or structural (m + 1 to m + n) non-basic variable, whose reduced cost reaches its zero bound first and thereby limits further decreasing the objective coefficient being analyzed. If coef1 = -DBL_MAX, var1 is set to 0.

value1 value of the basic variable being analyzed in an adjacent basis, which is defined as follows. Let the objective coefficient reaches its minimal value (coef1) and continues decreasing. Then the reduced cost of the limiting non-basic variable (var1) becomes dual infeasible and the current basis becomes non-optimal that forces the limiting non-basic variable to enter the basis replacing there some basic variable that leaves the basis to keep primal feasibility. Should note that on determining the adjacent basis current bounds of the basic variable being analyzed are ignored as if it were free (unbounded) variable, so it cannot leave the basis. It may happen that no dual feasible adjacent basis exists, in which case value1 is set to -DBL_MAX or +DBL_MAX.

coef2 the maximal value of the objective coefficient, at which the basis still remains dual feasible and thus optimal. +DBL_MAX means that the objective coefficient has no upper limit.

var2 the ordinal number of an auxiliary (1 to m) or structural (m + 1 to m + n) non-basic variable, whose reduced cost reaches its zero bound first and thereby limits further increasing the objective coefficient being analyzed. If coef2 = +DBL_MAX, var2 is set to 0.

value2 value of the basic variable being analyzed in an adjacent basis, which is defined exactly in the same way as value1 above with exception that now the objective coefficient is increasing.

Example:

```
> analyze_coef lp 1;  
-1.0,3,306.771624713959,1.79769313486232e+308,0,296.216606498195
```

5.4 Branch-and-cut API routines

All branch-and-cut API routines are supposed to be called from the callback routine. They cannot be called directly.

Basic routines

Determine reason for calling the callback routine Synopsis:

```
glp::ios_reason tree
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

one of the following:

glp::irowgen request for row generation

glp::ibingo better integer solution found

glp::iheur request for heuristic solution

glp::icutgen request for cut generation

glp::ibbranch request for branching

glp::iselect request for subproblem selection

glp::iprepro request for preprocessing

Example:


```
glp::ios:reason tree;
```

Access the problem object Synopsis:

```
glp::ios_get_prob tree
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

The routine returns a pointer to the problem object used by the MIP solver.

Example:

```
glp::ios_get_prob tree;
```

Determine additional row attributes Synopsis:

```
glp::ios_row_attr tree rowindex
```

Parameters:

tree pointer to the branch-and-cut search tree

rowindex row index

Returns:

The routine returns a tuple consisting of three values (level, origin, klass):

level subproblem level at which the row was created

origin the row origin flag - one of the following:

glp::rf_reg regular constraint

glp::rf_lazy "lazy" constraint

glp::rf_cut cutting plane constraint

klass the row class descriptor, which is a number passed to the routine `glp_ios_add_row` as its third parameter - if the row is a cutting plane constraint generated by the solver, its class may be the following:

glp::rf_gmi Gomory's mixed integer cut

glp::rf_mir mixed integer rounding cut

glp::rf_cov mixed cover cut

glp::rf_clq clique cut

Example:

```
glp::ios_row_attr tree 3;
```

Compute relative MIP gap Synopsis:

```
glp::ios_mip_gap tree
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

The routine returns the relative MIP gap.

Example:

```
> glp::ios_mip_gap tree;
```

Access application-specific data Synopsis:

```
glp::ios_node_data tree node
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

The routine `glp_ios_node_data` returns a pointer to the memory block for the specified subproblem. Note that if `cb_size = 0` was specified in the call of the **intopt** function, the routine returns a null pointer.

Example:

```
> glp::ios_node_data tree 23;
```

Select subproblem to continue the search Synopsis:

```
glp::ios_select_node tree node
```

Parameters:

tree pointer to the branch-and-cut search tree

node reference number of the subproblem from which the search will continue

Returns:

()

Example:

```
> glp::ios_select_node tree 23;
```

Provide solution found by heuristic Synopsis:

```
glp::ios_heur_sol tree colvector
```

Parameters:

tree pointer to the branch-and-cut search tree

colvector solution found by a primal heuristic. Primal values of all variables (columns) found by the heuristic should be placed in the list, i. e. the list must contain n numbers where n is the number of columns in the original problem object. Note that the routine does not check primal feasibility of the solution provided.

Returns:

If the provided solution is accepted, the routine returns zero. Otherwise, if the provided solution is rejected, the routine returns non-zero.

Example:

```
> glp::ios_heur_sol tree [15.7, (-3.1), 2.2];
```

Check whether can branch upon specified variable Synopsis:

```
glp::ios_can_branch tree j
```

Parameters:

tree pointer to the branch-and-cut search tree

j variable (column) index

Returns:

The function returns non-zero if j -th variable can be used for branching. Otherwise, it returns zero.

Example:

```
> glp::ios_can_branch tree 23;
```

Choose variable to branch upon Synopsis:

```
glp::ios_branch_upon tree j selection
```

Parameters:

tree pointer to the branch-and-cut search tree
j ordinal number of the selected branching variable
selection one of the following:
 glp::dn_brnch select down-branch
 glp::up_brnch select up-branch
 glp::no_brnch use general selection technique

Returns:

()

Example:

```
> glp::ios_branch_upon tree 23 glp::up_brnch;
```

Terminate the solution process Synopsis:

```
glp::ios_terminate tree
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

()

Example:

```
> glp::ios_terminate tree;
```

The search tree exploring routines

Determine the search tree size Synopsis:

```
glp::ios_tree_size tree
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

The routine returns a tuple (a_cnt, n_cnt, t_cnt), where

a_cnt the current number of active nodes

n_cnt the current number of all (active and inactive) nodes

t_cnt the total number of nodes including those which have been already removed from the tree. This count is increased whenever a new node appears in the tree and never decreased.

Example:

```
> glp::ios_tree_size tree;
```

Determine current active subproblem Synopsis:

```
glp::ios_curr_node tree
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

The routine returns the reference number of the current active subproblem. If the current subproblem does not exist, the routine returns zero.

Example:

```
> glp::ios_curr_node tree;
```

Determine next active subproblem Synopsis:

```
glp::ios_next_node tree node
```

Parameters:

tree pointer to the branch-and-cut search tree

node reference number of an active subproblem or zero

Returns:

If the parameter *p* is zero, the routine returns the reference number of the first active subproblem. If the tree is empty, zero is returned. If the parameter *p* is not zero, it must specify the reference number of some active subproblem, in which case the routine returns the reference number of the next active subproblem. If there is no next active subproblem in the list, zero is returned. All subproblems in the active list are ordered chronologically, i.e. subproblem A precedes subproblem B if A was created before B.

Example:

```
> glp::ios_next_node tree 23;
```

Determine previous active subproblem Synopsis:

`glp::ios_prev_node tree node`

Parameters:

tree pointer to the branch-and-cut search tree

node reference number of an active subproblem or zero

Returns:

If the parameter `p` is zero, the routine returns the reference number of the last active subproblem. If the tree is empty, zero is returned. If the parameter `p` is not zero, it must specify the reference number of some active subproblem, in which case the routine returns the reference number of the previous active subproblem. If there is no previous active subproblem in the list, zero is returned. All subproblems in the active list are ordered chronologically, i.e. subproblem A precedes subproblem B if A was created before B.

Example:

```
> glp::ios_prev_node tree 23;
```

Determine parent active subproblem Synopsis:

`glp::ios_up_node tree node`

Parameters:

tree pointer to the branch-and-cut search tree

node reference number of an active or inactive subproblem

Returns:

The routine returns the reference number of its parent subproblem. If the specified subproblem is the root of the tree, the routine returns zero.

Example:

```
> glp::ios_up_node tree 23;
```

Determine subproblem level Synopsis:

`glp::ios_node_level tree node`

Parameters:

tree pointer to the branch-and-cut search tree

node reference number of an active or inactive subproblem

Returns:

The routine returns the level of the given subproblem in the branch-and-bound tree. (The root subproblem has level 0.)

Example:

```
> glp::ios_node_level tree 23;
```

Determine subproblem local bound Synopsis:

```
glp::ios_node_bound tree node
```

Parameters:

tree pointer to the branch-and-cut search tree

node reference number of an active or inactive subproblem

Returns:

The routine returns the local bound for the given subproblem.

Example:

```
> glp::ios_node_bound tree 23;
```

Find active subproblem with the best local bound Synopsis:

```
glp::ios_best_node tree
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

The routine returns the reference number of the active subproblem, whose local bound is best (i.e. smallest in case of minimization or largest in case of maximization). If the tree is empty, the routine returns zero.

Example:

```
> glp::ios_best_node tree;
```

The cut pool routines**Determine current size of the cut pool Synopsis:**

```
glp::ios_pool_size tree
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

The routine returns the current size of the cut pool, that is, the number of cutting plane constraints currently added to it.

Example:

```
> glp::ios_pool_size tree;
```

Add constraint to the cut pool Synopsis:

```
glp::ios_add_row tree (name, klass, flags, row, type, rhs)
```

Parameters:

tree pointer to the branch-and-cut search tree

name symbolic name of the constraint

klass specifies the constraint class, which must be either zero or a number in the range from 101 to 200. The application may use this attribute to distinguish between cutting plane constraints of different classes.

flags currently is not used and must be zero

row list of pairs (colindex, coefficient)

type one of the following:

glp::lo $\sum(a_j.x_j) \geq \text{RHS constraint}$

glp::up $\sum(a_j.x_j) \leq \text{RHS constraint}$

rhs right hand side of the constraint

Returns:

The routine returns the ordinal number of the cutting plane constraint added, which is the new size of the cut pool.

Example:

```
> glp::ios_add_row tree ("new_constraint", 101, 0,
                        [(3, 15.0), (4, 6.7), (8, 1.25)], glp::up, 152.7);
```

Remove constraint from the cut pool Synopsis:

```
glp::ios_del_row tree rowindex
```

Parameters:

tree pointer to the branch-and-cut search tree

rowindex index of row to be deleted from the cut pool

Returns:

()

Remark:

Note that deleting a constraint from the cut pool leads to changing ordinal numbers of other constraints remaining in the pool. New ordinal numbers of the remaining constraints are assigned under assumption that the original order of constraints is not changed.

Example:

```
> glp::ios_del_row tree 5;
```

Remove all constraints from the cut pool Synopsis:

```
glp::ios_clear_pool tree
```

Parameters:

tree pointer to the branch-and-cut search tree

Returns:

()

Example:

```
> glp::ios_clear_pool tree;
```

5.5 Graph and network API routines

Basic graph routines

Create the GLPK graph object Synopsis:

```
glp::create_graph v_size a_size
```

Parameters:

v_size size of vertex data blocks, in bytes, $0 \leq \text{v_size} \leq 256$

a_size size of arc data blocks, in bytes, $0 \leq \text{a_size} \leq 256$.

Returns:

The routine returns a pointer to the graph created.

Example:

```
> let g = glp::create_graph 32 64;  
> g;  
#<pointer 0x9de7168>
```

Set the graph name Synopsis:

```
glp::set_graph_name graph name
```

Parameters:

graph pointer to the graph object

name the graph name, an empty string erases the current name

Returns:

()

Example:

```
> glp::set_graph_name graph "MyGraph";  
()
```

Add vertices to a graph Synopsis:

```
glp::add_vertices graph count
```

Parameters:

graph pointer to the graph object

count number of vertices to add

Returns:

The routine returns the ordinal number of the first new vertex added to the graph.

Example:

```
> glp::add_vertices graph 5;  
18
```

Add arc to a graph Synopsis:

```
glp::add_arc graph i j
```

Parameters:

graph pointer to the graph object

i index of the tail vertex

j index of the head vertex

Returns:

()

Example:

```
> glp::add_arc graph 7 12;  
()
```

Erase content of the GLPK graph object Synopsis:

```
glp::erase_graph graph v_size a_size
```

Parameters:

graph pointer to the graph object

v_size size of vertex data blocks, in bytes, $0 \leq \text{v_size} \leq 256$

a_size size of arc data blocks, in bytes, $0 \leq \text{a_size} \leq 256$.

Returns:

()

Remark:

The routine reinitialises the graph object. Its effect is equivalent to calling `delete_graph` followed by a call to `create_graph`.

Example:

```
> glp::erase_graph graph 16 34;  
()
```

Delete the GLPK graph object Synopsis:

```
glp::delete_graph graph
```

Parameters:

graph pointer to the graph object

Returns:

()

Remark:

The routine destroys the graph object and invalidates the pointer. This is done automatically when the graph is not needed anymore, the routine need not be usually called.

Example:

```
> glp::delete_graph graph  
()
```

Read graph in a plain text format Synopsis:

`glp::read_graph graph filename`

Parameters:

graph pointer to the graph object

filename file name

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> glp::read_graph graph "graph_data.txt";  
0
```

Write graph in a plain text format Synopsis:

`glp::write_graph graph filename`

Parameters:

graph pointer to the graph object

filename file name

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> glp::write_graph graph "graph_data.txt";  
0
```

Graph analysis routines

Find all weakly connected components of a graph Synopsis:

`glp::weak_comp graph v_num`

Parameters:

graph pointer to the graph object

v_num offset of the field of type `int` in the vertex data block, to which the routine stores the number of a weakly connected component containing that vertex - if `v_num < 0`, no component numbers are stored

Returns:

The routine returns the total number of components found.

Example:

```
> glp::weak_comp graph 16;
3
```

Find all strongly connected components of a graph Synopsis:

```
glp::strong_comp graph v_num
```

Parameters:

graph pointer to the graph object

v_num offset of the field of type int in the vertex data block, to which the routine stores the number of a strongly connected component containing that vertex - if $v_num < 0$, no component numbers are stored

Returns:

The routine returns the total number of components found.

Example:

```
> glp::strong_comp graph 16;
4
```

Minimum cost flow problem

Read minimum cost flow problem data in DIMACS format Synopsis:

```
glp::read_mincost graph v_rhs a_low a_cap a_cost filename
```

Parameters:

graph pointer to the graph object

v_rhs offset of the field of type double in the vertex data block, to which the routine stores b_i , the supply/demand value - if $v_rhs < 0$, the value is not stored

a_low offset of the field of type double in the arc data block, to which the routine stores l_{ij} , the lower bound to the arc flow - if $a_low < 0$, the lower bound is not stored

a_cap offset of the field of type double in the arc data block, to which the routine stores u_{ij} , the upper bound to the arc flow (the arc capacity) - if $a_cap < 0$, the upper bound is not stored

a_cost offset of the field of type double in the arc data block, to which the routine stores cij, the per-unit cost of the arc flow - if a_cost < 0, the cost is not stored

fname the name of a text file to be read in - if the file name ends with the suffix '.gz', the file is assumed to be compressed, in which case the routine decompresses it "on the fly"

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> glp::read_mincost graph 0 8 16 24 "graphdata.txt";  
0
```

Write minimum cost flow problem data in DIMACS format Synopsis:

```
glp::write_mincost graph v_rhs a_low a_cap a_cost fname
```

Parameters:

graph pointer to the graph object

v_rhs offset of the field of type double in the vertex data block, to which the routine stores bi, the supply/demand value - if v_rhs < 0, the value is not stored

a_low offset of the field of type double in the arc data block, to which the routine stores lij, the lower bound to the arc flow - if a_low < 0, the lower bound is not stored

a_cap offset of the field of type double in the arc data block, to which the routine stores uij, the upper bound to the arc flow (the arc capacity) - if a_cap < 0, the upper bound is not stored

a_cost offset of the field of type double in the arc data block, to which the routine stores cij, the per-unit cost of the arc flow - if a_cost < 0, the cost is not stored

fname the name of a text file to be written out - if the file name ends with the suffix '.gz', the file is assumed to be compressed, in which case the routine compresses it "on the fly"

Returns:

0 if reading went OK; non-zero in case of an error

Example:

```
> glp::write_mincost graph 0 8 16 24 "graphdata.txt";  
0
```

Convert minimum cost flow problem to LP Synopsis:

```
glp::mincost_lp lp graph names v_rhs a_low a_cap a_cost
```

Parameters:

lp pointer to the LP problem object

graph pointer to the graph object

names one of the following:

glp::on assign symbolic names of the graph object components to symbolic names of the LP problem object components

glp::off no symbolic names are assigned

v_rhs offset of the field of type double in the vertex data block, to which the routine stores b_i , the supply/demand value - if $v_rhs < 0$, it is assumed $b_i = 0$ for all nodes

a_low offset of the field of type double in the arc data block, to which the routine stores l_{ij} , the lower bound to the arc flow - if $a_low < 0$, it is assumed $l_{ij} = 0$ for all arcs

a_cap offset of the field of type double in the arc data block, to which the routine stores u_{ij} , the upper bound to the arc flow (the arc capacity) - if $a_cap < 0$, it is assumed $u_{ij} = 1$ for all arcs, value of DBL_MAX means an uncapacitated arc

a_cost offset of the field of type double in the arc data block, to which the routine stores c_{ij} , the per-unit cost of the arc flow - if $a_cost < 0$, it is assumed $c_{ij} = 0$ for all arcs

Returns:

()

Example:

```
> glp::mincost_lp lp graph glp::on 0 8 16 24;
()
```

Solve minimum cost flow problem with out-of-kilter algorithm Synopsis:

```
glp::mincost_okalg graph v_rhs a_low a_cap a_cost a_x v_pi
```

Parameters:

graph pointer to the graph object

v_rhs offset of the field of type double in the vertex data block, to which the routine stores b_i , the supply/demand value - if $v_rhs < 0$, it is assumed $b_i = 0$ for all nodes

- a_low** offset of the field of type double in the arc data block, to which the routine stores lij , the lower bound to the arc flow - if $a_low < 0$, it is assumed $lij = 0$ for all arcs
- a_cap** offset of the field of type double in the arc data block, to which the routine stores uij , the upper bound to the arc flow (the arc capacity) - if $a_cap < 0$, it is assumed $uij = 1$ for all arcs, value of `DBL_MAX` means an uncapacitated arc
- a_cost** offset of the field of type double in the arc data block, to which the routine stores cij , the per-unit cost of the arc flow - if $a_cost < 0$, it is assumed $cij = 0$ for all arcs
- a_x** offset of the field of type double in the arc data block, to which the routine stores xij , the arc flow found - if $a_x < 0$, the arc flow value is not stored
- v_pi** specifies an offset of the field of type double in the vertex data block, to which the routine stores pi , the node potential, which is the Lagrange multiplier for the corresponding flow conservation equality constraint

Remark:

Note that all solution components (the objective value, arc flows, and node potentials) computed by the routine are always integer-valued.

Returns:

The function returns a tuple in the form **(code, obj)**, where **code** is one of the following

- glp::ok** optimal solution found
- glp::enopfs** no (primal) feasible solution exists
- glp::edata** unable to start the search, because some problem data are either not integer-valued or out of range; this code is also returned if the total supply, which is the sum of b_i over all source nodes (nodes with $b_i > 0$), exceeds `INT_MAX`
- glp::erange** the search was prematurely terminated because of integer overflow
- glp::efail** an error has been detected in the program logic - if this code is returned for your problem instance, please report to [<bug-glpk@gnu.org>](mailto:bug-glpk@gnu.org)

and **obj** is value of the objective function.

Example:

```
> glp::mincost_okalg graph 0 8 16 24 32 40;
(glp::ok, 15)
```


Klingman's network problem generator Synopsis:

`glp::netgen graph v_rhs a_cap a_cost parameters`

Parameters:

graph pointer to the graph object

v_rhs offset of the field of type double in the vertex data block, to which the routine stores b_i , the supply/demand value - if $v_rhs < 0$, it is assumed $b_i = 0$ for all nodes

a_cap offset of the field of type double in the arc data block, to which the routine stores u_{ij} , the upper bound to the arc flow (the arc capacity) - if $a_cap < 0$, it is assumed $u_{ij} = 1$ for all arcs, value of DBL_MAX means an uncapacitated arc

a_cost offset of the field of type double in the arc data block, to which the routine stores c_{ij} , the per-unit cost of the arc flow - if $a_cost < 0$, it is assumed $c_{ij} = 0$ for all arcs

parameters tuple of exactly 15 integer numbers with the following meaning:

parm[1] iseed 8-digit positive random number seed

parm[2] nprob 8-digit problem id number

parm[3] nodes total number of nodes

parm[4] nsorc total number of source nodes (including transshipment nodes)

parm[5] nsink total number of sink nodes (including transshipment nodes)

parm[6] iarcs number of arc

parm[7] mincst minimum cost for arcs

parm[8] maxcst maximum cost for arcs

parm[9] itsup total supply

parm[10] ntsorc number of transshipment source nodes

parm[11] ntsink number of transshipment sink nodes

parm[12] iphic percentage of skeleton arcs to be given the maximum cost

parm[13] ipcap percentage of arcs to be capacitated

parm[14] mincap minimum upper bound for capacitated arcs

parm[15] maxcap maximum upper bound for capacitated arcs

Returns:

0 if the instance was successfully generated, nonzero otherwise

Example:

```
> glp::netgen graph 0 8 16 (12345678, 87654321, 20, 12, 8,  
                           25, 5, 20, 300, 6, 5, 15, 100, 1, 30);  
0
```

Grid-like network problem generator Synopsis:

```
glp::gridgen graph v_rhs a_cap a_cost parameters
```

Parameters:

graph pointer to the graph object

v_rhs offset of the field of type double in the vertex data block, to which the routine stores b_i , the supply/demand value - if $v_rhs < 0$, it is assumed $b_i = 0$ for all nodes

a_cap offset of the field of type double in the arc data block, to which the routine stores u_{ij} , the upper bound to the arc flow (the arc capacity) - if $a_cap < 0$, it is assumed $u_{ij} = 1$ for all arcs, value of DBL_MAX means an uncapacitated arc

a_cost offset of the field of type double in the arc data block, to which the routine stores c_{ij} , the per-unit cost of the arc flow - if $a_cost < 0$, it is assumed $c_{ij} = 0$ for all arcs

parameters tuple of exactly 14 integer numbers with the following meaning:

parm[1] two-ways arcs indicator:

1: if links in both direction should be generated

0: otherwise

parm[2] random number seed (a positive integer)

parm[3] number of nodes (the number of nodes generated might be slightly different to make the network a grid)

parm[4] grid width

parm[5] number of sources

parm[6] number of sinks
parm[7] average degree
parm[8] total flow
parm[9] distribution of arc costs:

1: uniform
 2: exponential

parm[10] lower bound for arc cost (uniform), 100 lambda,
 (exponential)
parm[11] upper bound for arc cost (uniform), not used (ex-
 ponential)
parm[12] distribution of arc capacities:

1: uniform
 2: exponential

parm[13] lower bound for arc capacity (uniform), 100
 lambda (exponential)
parm[14] upper bound for arc capacity (uniform), not
 used (exponential)

Returns:

0 if the instance was successfully generated, nonzero otherwise

Example:

```
> glp::gridgen graph 0 8 16 (1, 123, 20, 4, 7, 5, 3, 300, 1, 1, 5, 1, 5, 30);
0
```

Maximum flow problem

Read maximum cost flow problem data in DIMACS format Synopsis:

```
glp::read_maxflow graph a_cap filename
```

Parameters:

graph pointer to the graph object

Returns:

Example:

>

Write maximum cost flow problem data in DIMACS format **Synopsis:**

`glp::write_maxflow graph s t a_cap filename`

Parameters:

`graph` pointer to the graph object

Returns:

Example:

>

Convert maximum flow problem to LP **Synopsis:**

`glp::maxflow_lp lp graph names s t a_cap`

Parameters:

`graph` pointer to the graph object

Returns:

Example:

>

Solve maximum flow problem with Ford-Fulkerson algorithm **Synopsis:**

`glp::maxflow_ffalg graph s t a_cap a_x v_cut`

Parameters:

`graph` pointer to the graph object

Returns:

Example:

>

Goldfarb's maximum flow problem generator **Synopsis:**

`glp::rmfgen graph a_cap parameters`

Parameters:

graph pointer to the graph object

Returns:

Example:

>

5.6 Miscellaneous routines

Library environment routines

Determine library version Synopsis:

`glp::version`

Parameters:

none

Returns:

GLPK library version

Example:

```
> glp::version;  
"4.38"
```

Enable/disable terminal output Synopsis:

`glp::term_out` switch

Parameters:

switch one of the following:

glp::on enable terminal output from GLPK routines

glp::off disable terminal output from GLPK routines

Returns:

()

Example:

```
> glp::term_out glp:off;  
()
```

Enable/disable the terminal hook routine Synopsis:

```
glp::term_hook switch info
```

Parameters:

switch one of the following:

glp::on use the terminal callback function

glp::off don't use the terminal callback function

info pointer to a memory block which can be used for passing additional information to the terminal callback function

Returns:

()

Example:

```
> glp::term_hook glp::on NULL;  
()
```

Get memory usage information Synopsis:

```
glp::mem_usage
```

Parameters:

none

Returns:

tuple consisting of four numbers:

- count (int) - the number of currently allocated memory blocks
- cpeak (int) - the peak value of count reached since the initialization of the GLPK library environment
- total (bigint) - the total amount, in bytes, of currently allocated memory blocks
- tpeak (bigint) - the peak value of total reached since the initialization of the GLPK library environment

Example:

```
> glp::mem_usage;  
7,84,10172L,45304L
```

Set memory usage limit Synopsis:

```
glp::mem_limit limit
```

Parameters:

limit memory limit in megabytes

Returns:

()

Example:

```
> glp::mem_limit 200;  
()
```

Free GLPK library environment Synopsis:

```
glp::free_env
```

Parameters:

none

Returns:

()

Example:

```
> glp_free_env;  
()
```


Gnuplot bindings

Kay-Uwe Kirstein

1 Copying

Copyright (c) 2009, 2010 by Kay-Uwe Kirstein.

pure-gplot is free software: you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

2 Introduction

This module contains a pure binding to gnuplot. Communication to gnuplot is performed via pipes. The usual work flow to generate plot via gnuplot is the following:

1. open pipe via `open`
2. send plot commands, e.g., with `plot`
3. close pipe with `close`

3 Function Reference

3.1 Open / Closing Functions

```
gplot::open cmd;
```

opens a pipe to gnuplot, using *cmd*. *cmd* usually is something like `gnuplot` or `/path/to/gnuplot/bin/gnuplot` depending on your path configuration. `open` returns a pointer to the actual pipe for later usage, so a typical call to `open` might look like this:

```
let gp = gplot::open "/path_to_gnuplot/gnuplot";
```

`gplot::GPLOT_EXE` is a predefined variable with the standard Gnuplot executable. It is set to `pgnuplot` on Windows and to `gnuplot` otherwise and can be overridden by the `GPLOT_EXE` environment variable. (`pgnuplot.exe` is a special executable for Windows, which is capable of stdin pipes in contrast to the normal `gnuplot.exe`). Usage of `gplot::GPLOT_EXE` might look like this:

```
let gp = gplot::open gplot::GPLOT_EXE;
gplot::close gp;
```

closes a gnuplot session, given by the handle *gp*.

3.2 Low-Level Commands

```
gplot::puts_no_echo string gp;
```

sends the string to the gnuplot session *gp* points to. As the name states, there is no echo read back from gnuplot (Don't know whether *gnuplot* or *pgnuplot.exe* supports reading/bidirectional pipes at all).

```
gplot::puts string gp;
```

is a convenience wrapper to `gplot::puts_no_echo`.

3.3 Plot Commands

The main (versatile) function to generate plots is the simple plot command, which expects a list of the data to be plotted.

```
gplot::plot gp data opt;
```

where *gp* is the pointer to the gnuplot session, *data* is a list containing the data to be plotted and *opt* is a tuple, containing options for the plot. *opt* might be empty () or `DEFAULT` for default options (refer to gnuplot for them).

If data for the x-axis (ordinate) should be explicitly given *plotxy* should be used instead:

```
gplot::plotxy_deprecated gp (xdata, ydata) opt;
gplot::plotxy gp (xdata, ydata) opt [];
```

Multiple datasets can be plotted into a single graph by combining them to tuples of lists:

```
gplot::plotxy gp (xdata, y1data, y2data, ..) opt;
gplot::plotxy gp (xdata, y1data, y2data, ..) opt [];
gplot::plotxy gp (xdata, y1data, y2data, ..) opt titles;
```

where the latter form gives additional titles for each y-data set.

3.4 Plot Options

```
gplot::xtics gp list_of_tic_labels;
```

Sets the tic labels of the x-axis to the given text labels. The labels can be given as a simple list of strings, which are taken as successive labels or as a list of tuples with the form (value, label), in which case each label is placed at its value position.

```
gplot::xtics gp () or gplot::xtics gp "default";
```

This restores the default tics on the y-axis.

```
gplot::title t;
```

Sets a title string on top of the plot (default location)

```
gplot::output gp terminal name;
```

Sets the terminal and output name for the successive plots. For some terminal additional options might be given:

```
gplot::output gp (terminal, options) name.
```

For terminals like x11 or windows, name can be empty ().

```
gplot::xlabel gp name or gplot::ylabel gp name
```

Adds labels to the x- or y-axis, respectively. An empty name removes the label for successive plots, e.g., `gplot::xlabel gp ""`.

3.5 Private Functions

```
gpdata data, gpxydata (xdata, yldata, ..)
```

Internal functions to handle lists of data point (gpdata) or tuples of lists of data points (gpxydata) and convert them to be understood by Gnuplot.

```
gpxycmd, gpxycmdtitle
```

Internal function to generate the plotting command for multiple datasets. gpxycmdtitle adds titles to each dataset, a.k.a plot legend.

```
gplot::gpopt ("style", style, args);
```

Internal function to convert a plot style to the respective gnuplot syntax

```
gplot::gptitle t;
```

Internal function to generate title information for individual datasets

pure-gsl - GNU Scientific Library Interface for Pure

Version 0.9, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

Eddie Rucker <erucker@bmc.edu>

License: GPL V3 or later, see the accompanying COPYING file

Building on Pure's GSL-compatible matrix support, this module aims to provide a complete wrapper for the GNU Scientific Library which provides a wide range of mathematical routines useful for scientific programming, number crunching and signal processing applications.

This is still work in progress, only a small part of the interface is finished right now. Here is a brief summary of the operations which are implemented:

- Matrix-scalar and matrix-matrix arithmetic. This is fairly complete and includes matrix multiplication, as well as element-wise exponentiation (^) and integer operations (div, mod, bit shifts and bitwise logical operations) which aren't actually in the GSL API.
- SVD (singular value decomposition), as well as the corresponding solvers, pseudo inverses and left and right matrix division. This is only available for real matrices right now, as GSL doesn't implement complex SVD.
- Random distributions (p.d.f. and c.d.f.) and statistic functions.
- Polynomial evaluation and roots.
- Linear least-squares fitting. Multi-fitting is not available yet.

Installation instructions: Get the latest source from <http://pure-lang.googlecode.com/files/pure-gsl-0.9.tar.gz>. Run `make` to compile the module and `make install` (as root) to install it in the Pure library directory. This requires GNU make, and of course you need to have Pure and GSL installed. The `make install` step is only necessary for system-wide installation.

make tries to guess your Pure installation directory and platform-specific setup. If it gets this wrong, you can set some variables manually. In particular, make `install prefix=/usr` sets the installation prefix, and make `PIC=-fPIC` or some similar flag might be needed for compilation on 64 bit systems. Please see the Makefile for details.

The current release requires GSL 1.11 or later and Pure 0.45 or later. Older GSL versions might still work, but then some operations may be missing. The latest and greatest GSL version is always available from <http://www.gnu.org/software/gsl>.

After installation, you can import the entire GSL interface as follows:

```
using gsl;
```

For convenience, the different parts of the GSL interface are also available as separate modules. E.g., if you only need the matrix operations:

```
using gsl::matrix;
```

In either case, the global `gsl_version` variable reports the installed GSL version:

```
> show gsl_version  
let gsl_version = "1.11";
```

(This variable used to be defined by the Pure runtime but has been moved into pure-gsl as of Pure 0.37.)

Most other operations are declared in separate namespaces which are in 1-1 correspondence with the module names. Thus, e.g., the `gsl_poly_eval` routine is named `gsl::poly::eval` in Pure and can be found in the `gsl::poly` module and namespace. The `using namespace` declaration can be used to facilitate access to the operations in a given namespace, e.g.:

```
> using gsl::poly;  
> using namespace gsl::poly;  
> eval {1,2,3} 2;  
17
```

See the examples folder in the sources for some examples.

If you'd like to contribute, please mail the authors or contact us at <http://groups.google.com/group/pure-lang>.

1 Polynomials

This module provides Pure wrappers for the GSL polynomial routines. For detail about the routines, see Chapter 6 of the GSL manual,

http://www.gnu.org/software/gsl/manual/html_node/Polynomials.html.

Polynomials are represented by vectors (one row matrices).

1.1 Routines

eval **c::matrix** **x** implements `gsl_poly_eval`, `gsl_poly_complex_eval`, and `gsl_complex_poly_eval` without the `len` parameter.

GSL does not supply an integer routine for evaluating polynomials with `int` or `bigint` coefficients. Therefore, an integer routine has been provided in `pure-gsl` using the Chinese Remainder Theorem.

dd_init **x::matrix** **y::matrix** implements `gsl_poly_dd_init` without the `size` parameter.

dd_eval **dd::matrix** **xa::matrix** **x::double** implements `gsl_poly_dd_eval` without the `size` parameter.

dd_taylor **xp::double** **dd::matrix** **xa::matrix** implements `gsl_poly_dd_taylor` without the `size` and `workspace w` arguments.

solve_quadratic **a b c** implements `gsl_poly_solve_quadratic`. This function returns a list of roots instead of passing them through the parameters `x0` and `x1`.

complex_solve_quadratic **a b c** implements `gsl_poly_complex_solve_quadratic`. This function returns a list of roots instead of passing through the parameters `z0` and `z1`.

solve_cubic **a b c** implements `gsl_poly_solve_cubic`. This function returns a list of roots instead of passing them through the parameters `x0`, `x1`, and `x2`.

complex_solve_cubic **a b c** implements `gsl_poly_complex_solve_cubic`. This function returns a list of roots instead of passing them through the parameters `z0`, `z1`, and `z2`.

complex_solve **c::matrix** implements `gsl_poly_complex_solve` omitting the parameters `n` and `w`. The GSL routines for creating and freeing the workspace are handled automatically.

1.2 Examples

Usage of each library routine is illustrated below.

```
> using gsl::poly;
> using namespace gsl::poly;
> eval {1,2,3} 2;
17
> eval {1.0,2.0,3.0} (-2.0);
9.0
> eval {1, 2, 2} (1+:1);
3.0+:6.0
> eval {1+:2, 2+:3, 2+:3} (1+:1);
-6.0+:11.0
> let dd = dd_init {1,2,3} {2,4,6};
> dd;
{2.0,2.0,0.0}
> dd_eval dd {1,2,3} 2;
```

```
4.0
> dd_taylor 0.0 dd {1,2,3};
{0.0,2.0,0.0}
> solve_quadratic 2 4 1;
[-1.70710678118655,-0.292893218813452]
> solve_quadratic 1 4 4;
[-2.0,-2.0]
> solve_quadratic 0 2 1;
[-0.5]
> solve_quadratic 1 2 8;
[]
> complex_solve_quadratic 0 2 1;
[-0.5+:0.0]
> complex_solve_quadratic 2 2 3;
[-0.5+:-1.11803398874989,-0.5+:1.11803398874989]
> solve_cubic 3 3 1;
[-1.0,-1.0,-1.0]
> solve_cubic 3 2 1;
[-2.32471795724475]
> complex_solve_cubic 2 2 1;
[-1.0+:0.0,-0.5+:-0.866025403784439,-0.5+:0.866025403784439]
> complex_solve {6,1,-7,-1,1};
[1.0+:0.0,-1.0+:0.0,-2.0+:0.0,3.0+:0.0]
```

2 Special Functions

This module is loaded via the command using `gsl::sf` and provides Pure wrappers for the GSL Special Functions. For details, see Chapter 7 of the GSL manual,

http://www.gnu.org/software/gsl/manual/html_node/Special-Functions.html.

To load the library, use the Pure command using `gsl::sf`. Modes for the functions must be one of:

```
GSL_PREC_DOUBLE
GSL_PREC_SINGLE
GSL_PREC_APPROX
```

Results for some of the functions are returned as a Pure list instead of the `gsl_sf_result` or `gsl_sf_result_e10` structures in C. In these cases, the resulting list is one of the following forms.

- `[val, err]` for the `gsl_sf_result` struct and
- `[val, err, e10]` for the `gsl_sf_result_e10` struct.

2.1 Airy Functions

airy_Ai x and **airy_Ai (x, mode::int)** implements `gsl_sf_airy_Ai`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Ai_e x and **airy_Ai_e (x, mode::int)** implements `gsl_sf_airy_Ai_e`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Ai_scaled x and **airy_Ai_scaled (x, mode::int)** implements `gsl_sf_airy_Ai_scaled`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Ai_scaled_e x and **airy_Ai_scaled_e (x, mode::int)** implements `gsl_sf_airy_Ai_scaled_e`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Bi x and **airy_Bi (x, mode::int)** implements `gsl_sf_airy_Bi`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Bi_e x and **airy_Bi_e (x, mode::int)** implements `gsl_sf_airy_Bi_e`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Bi_scaled x and **airy_Bi_scaled (x, mode::int)** implements `gsl_sf_airy_Bi_scaled`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Bi_scaled_e x and **airy_Bi_scaled_e (x, mode::int)** implements `gsl_sf_airy_Bi_scaled_e`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Ai_deriv x and **airy_Ai_deriv (x, mode::int)** implements `gsl_sf_airy_Ai_deriv`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Ai_deriv_e x and **airy_Ai_deriv_e (x, mode::int)** implements `gsl_sf_airy_Ai_deriv_e`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Ai_deriv_scaled x and **airy_Ai_deriv_scaled (x, mode::int)** implements `gsl_sf_airy_Ai_deriv_scaled`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Ai_deriv_scaled_e x and **airy_Ai_deriv_scaled_e (x, mode::int)** implements `gsl_sf_airy_Ai_deriv_scaled_e`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Bi_deriv x and **airy_Bi_deriv (x, mode::int)** implements `gsl_sf_airy_Bi_deriv`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Bi_deriv_e x and **airy_Bi_deriv_e (x, mode::int)** implements `gsl_sf_airy_Bi_deriv_e`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Bi_deriv_scaled *x* and **airy_Bi_deriv_scaled** (*x*, *mode::int*) implements `gsl_sf_airy_Bi_deriv_scaled`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_Bi_deriv_scaled_e *x* and **airy_Bi_deriv_scaled_e** (*x*, *mode::int*) implements `gsl_sf_airy_Bi_deriv_scaled_e`. The first form computes the function with `mode = GSL_PREC_DOUBLE`.

airy_zero_Ai *s* implements `gsl_sf_airy_zero_Ai`.

airy_zero_Ai_e *s* implements `gsl_sf_airy_zero_Ai_e`.

airy_zero_Bi *s* implements `gsl_sf_airy_zero_Bi`.

airy_zero_Bi_e *s* implements `gsl_sf_airy_zero_Bi_e`.

airy_zero_Ai_deriv *s* implements `gsl_sf_airy_zero_Ai_deriv`.

airy_zero_Ai_deriv_e *s* implements `gsl_sf_airy_zero_Ai_deriv_e`.

airy_zero_Bi_deriv *s* implements `gsl_sf_airy_zero_Bi_deriv`.

airy_zero_Bi_deriv_e *s* implements `gsl_sf_airy_zero_Bi_deriv_e`.

2.2 Examples

The following illustrate the Airy functions.

```
> using gsl::sf;
> using namespace gsl::sf;
> airy_Ai (-1.2); // defaults to GSL_PREC_DOUBLE
0.52619437480212
> airy_Ai_scaled (-1.2);
0.52619437480212
> airy_Ai (-1.2, GSL_PREC_APPROX);
0.526194374771687
> airy_Ai_scaled (-1.2, GSL_PREC_SINGLE);
0.526194374771687
> airy_Ai_e (-1.2);
[0.52619437480212, 1.88330586480371e-15]
> airy_Ai_e (-1.2, GSL_PREC_APPROX);
[0.526194374771687, 1.01942940819652e-08]
> airy_Ai_scaled_e (-1.2);
[0.52619437480212, 1.88330586480371e-15]
> airy_Ai_scaled_e (-1.2, GSL_PREC_APPROX);
[0.526194374771687, 1.01942940819652e-08]
> airy_Bi (-1.2);
-0.015821370184632
> airy_Bi_scaled (-1.2);
-0.015821370184632
> airy_Bi (-1.2, GSL_PREC_APPROX);
-0.0158213701898015
> airy_Bi_scaled (-1.2, GSL_PREC_SINGLE);
```

```
-0.0158213701898015
> airy_Bi_e (-1.2);
[-0.015821370184632,1.31448899295896e-16]
> airy_Bi_e (-1.2,GSL_PREC_APPROX);
[-0.0158213701898015,4.10638404843775e-10]
> airy_Bi_scaled_e (-1.2);
[-0.015821370184632,1.31448899295896e-16]
> airy_Bi_scaled_e (-1.2,GSL_PREC_APPROX);
[-0.0158213701898015,4.10638404843775e-10]
> airy_Ai_deriv (-1.2); // defaults to GSL_PREC_DOUBLE
0.107031569272281
> airy_Ai_deriv_scaled (-1.2);
0.107031569272281
> airy_Ai_deriv (-1.2,GSL_PREC_APPROX);
0.107031569264504
> airy_Ai_deriv_scaled (-1.2, GSL_PREC_SINGLE);
0.107031569264504
> airy_Ai_deriv_e (-1.2);
[0.107031569272281,3.02919983680384e-16]
> airy_Ai_deriv_e (-1.2,GSL_PREC_APPROX);
[0.107031569264504,9.25921017197604e-11]
> airy_Ai_deriv_scaled_e (-1.2);
[0.107031569272281,3.02919983680384e-16]
> airy_Ai_deriv_scaled_e (-1.2,GSL_PREC_APPROX);
[0.107031569264504,9.25921017197604e-11]
> airy_Bi_deriv (-1.2);
0.601710157437464
> airy_Bi_deriv_scaled (-1.2);
0.601710157437464
> airy_Bi_deriv (-1.2,GSL_PREC_APPROX);
0.601710157441937
> airy_Bi_deriv_scaled (-1.2, GSL_PREC_SINGLE);
0.601710157441937
> airy_Bi_deriv_e (-1.2);
[0.601710157437464,1.7029557943563e-15]
> airy_Bi_deriv_e (-1.2,GSL_PREC_APPROX);
[0.601710157441937,5.20534347823991e-10]
> airy_Bi_deriv_scaled_e (-1.2);
[0.601710157437464,1.7029557943563e-15]
> airy_Bi_deriv_scaled_e (-1.2,GSL_PREC_APPROX);
[0.601710157441937,5.20534347823991e-10]
> airy_zero_Ai 2;
-4.08794944413097
> airy_zero_Ai_e 3;
[-5.52055982809555,1.22581052599448e-15]
> airy_zero_Bi 2;
-3.27109330283635
> airy_zero_Bi_e 3;
[-4.83073784166202,1.07263927554824e-15]
> airy_zero_Ai_deriv 2;
-3.24819758217984
> airy_zero_Ai_deriv_e 3;
```

```
[-4.82009921117874,1.07027702504564e-15]
> airy_zero_Bi_deriv 2;
-4.07315508907183
> airy_zero_Bi_deriv_e 3;
[-5.5123957296636,1.22399773198358e-15]
```

2.3 Bessel Functions

bessel_J0 *x* implements `gsl_sf_bessel_J0`.

bessel_J0_e *x* implements `gsl_sf_bessel_J0_e`.

bessel_J1 *x* implements `gsl_sf_bessel_J1`.

bessel_J1_e *x* implements `gsl_sf_bessel_J1_e`.

bessel_Jn *n x* implements `gsl_sf_bessel_Jn`.

bessel_Jn_e *n x* implements `gsl_sf_bessel_Jn_e`.

bessel_Jn_array *nmin::int nmax::int x* implements `gsl_sf_bessel_Jn_array`.

bessel_Y0 *x* implements `gsl_sf_bessel_Y0`.

bessel_Y0_e *x* implements `gsl_sf_bessel_Y0_e`.

bessel_Y1 *x* implements `gsl_sf_bessel_Y1`.

bessel_Y1_e *x* implements `gsl_sf_bessel_Y1_e`.

bessel_Yn *x* implements `gsl_sf_bessel_Yn`.

bessel_Yn_e *x* implements `gsl_sf_bessel_Yn_e`.

bessel_Yn_array *nmin::int nmax::int x* implements `gsl_sf_bessel_Yn_array`.

bessel_I0 *x* implements `gsl_sf_bessel_I0`.

bessel_I0_e *x* implements `gsl_sf_bessel_I0_e`.

bessel_I1 *x* implements `gsl_sf_bessel_I1`.

bessel_I1_e *x* implements `gsl_sf_bessel_I1_e`.

bessel_In *n::int x* implements `gsl_sf_bessel_In`.

bessel_In_e *n::int x* implements `gsl_sf_bessel_In_e`.

bessel_In_array *nmin::int nmax::int x* implements `gsl_sf_bessel_In_array`.

bessel_I0_scaled *x* implements `gsl_sf_bessel_I0_scaled`.

bessel_I0_scaled_e *x* implements `gsl_sf_bessel_I0_scaled_e`.

bessel_I1_scaled *x* implements `gsl_sf_bessel_I1_scaled`.

bessel_I1_scaled_e *x* implements `gsl_sf_bessel_I1_scaled_e`.

`bessel_In_scaled n::int x` implements `gsl_sf_bessel_In_scaled`.

`bessel_In_scaled_e n::int x` implements `gsl_sf_bessel_In_scaled_e`.

`bessel_In_scaled_array nmin::int nmax::int x` implements `gsl_sf_bessel_In_array`.

`bessel_K0 x` implements `gsl_sf_bessel_K0`.

`bessel_K0_e x` implements `gsl_sf_bessel_K0_e`.

`bessel_K1 x` implements `gsl_sf_bessel_K1`.

`bessel_K1_e x` implements `gsl_sf_bessel_K1_e`.

`bessel_Kn n::int x` implements `gsl_sf_bessel_Kn`.

`bessel_Kn_e n::int x` implements `gsl_sf_bessel_Kn_e`.

`bessel_Kn_array nmin::int nmax::int x` implements `gsl_sf_bessel_Kn_array`.

`bessel_K0_scaled x` implements `gsl_sf_bessel_K0_scaled`.

`bessel_K0_scaled_e x` implements `gsl_sf_bessel_K0_scaled_e`.

`bessel_K1_scaled x` implements `gsl_sf_bessel_K1_scaled`.

`bessel_K1_scaled_e x` implements `gsl_sf_bessel_K1_scaled_e`.

`bessel_Kn_scaled n::int x` implements `gsl_sf_bessel_Kn_scaled`.

`bessel_Kn_scaled_e n::int x` implements `gsl_sf_bessel_Kn_scaled_e`.

`bessel_Kn_scaled_array nmin::int nmax::int x` implements `gsl_sf_bessel_Kn_array`.

`bessel_j0 x` implements `gsl_sf_bessel_j0`.

`bessel_j0_e x` implements `gsl_sf_bessel_j0_e`.

`bessel_j1 x` implements `gsl_sf_bessel_j1`.

`bessel_j1_e x` implements `gsl_sf_bessel_j1_e`.

`bessel_j2 x` implements `gsl_sf_bessel_j2`.

`bessel_j2_e x` implements `gsl_sf_bessel_j2_e`.

`bessel_jl l::int x` implements `gsl_sf_bessel_jl`.

`bessel_jl_e l::int x` implements `gsl_sf_bessel_jl_e`.

`bessel_jl_array lmax::int x` implements `gsl_sf_bessel_jl_array`.

`bessel_jl_stepped_array lmax::int x` implements `gsl_sf_bessel_jl_stepped_array`.

`bessel_y0 x` implements `gsl_sf_bessel_y0`.

`bessel_y0_e x` implements `gsl_sf_bessel_y0_e`.

`bessel_y1 x` implements `gsl_sf_bessel_y1`.

`bessel_y1_e x` implements `gsl_sf_bessel_y1_e`.

`bessel_y2 x` implements `gsl_sf_bessel_y2`.
`bessel_y2_e x` implements `gsl_sf_bessel_y2_e`.
`bessel_yl l::int x` implements `gsl_sf_bessel_yl`.
`bessel_yl_e l::int x` implements `gsl_sf_bessel_yl_e`.
`bessel_yl_array lmax::int x` implements `gsl_sf_bessel_yl_array`.
`bessel_i0_scaled x` implements `gsl_sf_bessel_i0_scaled`.
`bessel_i0_scaled_e x` implements `gsl_sf_bessel_i0_scaled_e`.
`bessel_i1_scaled x` implements `gsl_sf_bessel_i1_scaled`.
`bessel_i1_scaled_e x` implements `gsl_sf_bessel_i1_scaled_e`.
`bessel_i2_scaled x` implements `gsl_sf_bessel_i2_scaled`.
`bessel_i2_scaled_e x` implements `gsl_sf_bessel_i2_scaled_e`.
`bessel_il_scaled l::int x` implements `gsl_sf_bessel_il_scaled`.
`bessel_il_scaled_e l::int x` implements `gsl_sf_bessel_il_scaled_e`.
`bessel_il_scaled_array lmax::int x` implements `gsl_sf_bessel_il_scaled_array`.
`bessel_k0_scaled x` implements `gsl_sf_bessel_k0_scaled`.
`bessel_k0_scaled_e x` implements `gsl_sf_bessel_k0_scaled_e`.
`bessel_k1_scaled x` implements `gsl_sf_bessel_k1_scaled`.
`bessel_k1_scaled_e x` implements `gsl_sf_bessel_k1_scaled_e`.
`bessel_k2_scaled x` implements `gsl_sf_bessel_k2_scaled`.
`bessel_k2_scaled_e x` implements `gsl_sf_bessel_k2_scaled_e`.
`bessel_kl_scaled l::int x` implements `gsl_sf_bessel_kl_scaled`.
`bessel_kl_scaled_e l::int x` implements `gsl_sf_bessel_kl_scaled_e`.
`bessel_kl_scaled_array lmax::int x` implements `gsl_sf_bessel_kl_scaled_array`.
`bessel_Jnu nu x` implements `gsl_sf_bessel_Jnu`.
`bessel_Jnu_e nu x` implements `gsl_sf_bessel_Jnu_e`.
`bessel_sequence_Jnu_e nu v::matrix` implements `gsl_sf_bessel_sequence_Jnu_e`.
`bessel_Ynu nu x` implements `gsl_sf_bessel_Ynu`.
`bessel_Ynu_e nu x` implements `gsl_sf_bessel_Ynu_e`.
`bessel_Inu nu x` implements `gsl_sf_bessel_Inu`.
`bessel_Inu_e nu x` implements `gsl_sf_bessel_Inu_e`.
`bessel_Inu_scaled nu x` implements `gsl_sf_bessel_Inu_scaled`.

`bessel_Inu_scaled_e nu x` implements `gsl_sf_bessel_Inu_scaled_e`.
`bessel_Knu nu x` implements `gsl_sf_bessel_Knu`.
`bessel_Knu_e nu x` implements `gsl_sf_bessel_Knu_e`.
`bessel_lnKnu nu x` implements `gsl_sf_bessel_lnKnu`.
`bessel_lnKnu_e nu x` implements `gsl_sf_bessel_lnKnu_e`.
`bessel_Knu_scaled nu x` implements `gsl_sf_bessel_Knu_scaled`.
`bessel_Knu_scaled_e nu x` implements `gsl_sf_bessel_Knu_scaled_e`.
`bessel_zero_J0 s::int` implements `gsl_sf_bessel_zero_J0`.
`bessel_zero_J0_e s::int` implements `gsl_sf_bessel_zero_J0_e`.
`bessel_zero_J1 s::int` implements `gsl_sf_bessel_zero_J1`.
`bessel_zero_J1_e s::int` implements `gsl_sf_bessel_zero_J1_e`.
`bessel_zero_Jnu nu s::int` implements `gsl_sf_bessel_zero_Jnu`.
`bessel_zero_Jnu_e nu s::int` implements `gsl_sf_bessel_zero_Jnu_e`.

2.4 Examples

The following illustrate the Bessel functions.

```
> using gsl::sf;
> using namespace gsl::sf;
> bessel_J0 (-1.2);
0.671132744264363
> bessel_J0_e 0.75;
[0.864242275166649,7.07329111491049e-16]
> bessel_J1 1.2;
0.498289057567216
> bessel_J1_e (-0.2);
[-0.099500832639236,5.00768737808415e-17]
> bessel_Jn 0 (-1.2);
0.671132744264363
> bessel_Jn_e 2 0.75;
[0.0670739972996506,5.48959386474892e-17]
> bessel_Jn_array 0 4 0.5;
[0.938469807240813,0.242268457674874,0.0306040234586826,
 0.00256372999458724,0.000160736476364288]
> bessel_Y0 0.25;
-0.931573024930059
> bessel_Y0_e 0.25;
[-0.931573024930059,6.4279898430593e-16]
> bessel_Y1 0.125;
-5.19993611253477
> bessel_Y1_e 4.325;
[0.343041276811844,2.74577716760089e-16]
```

```
> bessell_Yn 3 4.325;
-0.0684784962694202
> bessell_Yn_e 3 4.325;
[-0.0684784962694202, 3.37764590906247e-16]
> bessell_Yn_array 2 4 1.35;
[-1.07379345815726, -2.66813016175689, -10.7845628163178]
> bessell_I0 1.35;
1.51022709775726
> bessell_I0_e 1.35;
[1.51022709775726, 2.37852166449918e-15]
> bessell_I1 0.35;
0.177693400031422
> bessell_I1_e 0.35;
[0.177693400031422, 1.55520651386126e-16]
> bessell_In 2 3.0;
2.24521244092995
> bessell_In_e 2 3.0;
2.24521244092995, 5.98244771302867e-15]
> bessell_In_array 3 5 (-0.1);
[-2.08463574223272e-05, 2.60546902129966e-07, -2.6052519298937e-09]
> bessell_I0_scaled 1.05;
0.453242541279856
> bessell_I0_scaled_e 1.05;
[0.453242541279856, 4.10118141697477e-16]
> bessell_I1_scaled 1.05;
0.210226017612868
> bessell_I1_scaled_e 1.05;
[0.210226017612868, 2.12903131803686e-16]
> bessell_In_scaled 3 1.05;
0.00903732602788281
> bessell_In_scaled_e 3 1.05;
[0.00903732602788281, 2.00668948743994e-17]
> bessell_In_scaled_array 3 5 1.05;
[0.00903732602788281, 0.0011701685245855, 0.000121756316755217]
> bessell_K0 2.3;
0.0791399330020936
> bessell_K0_e 2.3;
[0.0791399330020936, 1.15144454318261e-16]
> bessell_K1 2.3;
0.0949824438453627
> bessell_K1_e 2.3;
[0.0949824438453627, 9.85583638959967e-17]
> bessell_Kn 2 3.4;
0.0366633035851529
> bessell_Kn_e 2 3.4;
[0.0366633035851529, 2.01761856558251e-16]
> bessell_Kn_array 1 3 2.5;
[0.0738908163477471, 0.121460206278564, 0.268227146393449]
> bessell_K0_scaled 1.5;
0.367433609054158
> bessell_K0_scaled_e 1.5;
[0.958210053294896, 1.25816573186951e-14]
```



```
> bessel_K1_scaled 1.5;
1.24316587355255
> bessel_K1_scaled_e 1.5;
[1.24316587355255, 2.32370553362606e-15]
> bessel_Kn_scaled 4 1.5;
35.4899165934682
> bessel_Kn_scaled_e 4 1.5;
[35.4899165934682, 3.89252285021454e-14]
> bessel_Kn_scaled_array 4 6 1.5;
[35.4899165934682, 197.498093175689, 1352.14387109806]
> bessel_j0 0.01;
0.999983333416666
> bessel_j0_e 0.01;
[0.999983333416666, 4.44081808400239e-16]
> bessel_j1 0.2;
0.0664003806703222
> bessel_j1_e 0.2;
[0.0664003806703222, 2.94876925856268e-17]
> bessel_j2 0.3;
0.00596152486862022
> bessel_j2_e 0.3;
[0.00596152486862022, 2.64744886840705e-18]
> bessel_jl 4 0.3;
8.53642426502516e-06
> bessel_jl_e 4 0.3;
[8.53642426502516e-06, 1.02355215483598e-19]
> bessel_jl_array 2 1.2;
[0.776699238306022, 0.34528456985779, 0.0865121863384538]
> bessel_jl_stepped_array 2 1.2;
[0.776699238306022, 0.34528456985779, 0.0865121863384538]
> bessel_y0 1;
-0.54030230586814
> bessel_y0_e 3;
[0.329997498866815, 2.93096657048522e-16]
> bessel_y1 3;
0.062959163602316
> bessel_y1_e 3.0;
[0.062959163602316, 1.04609100698801e-16]
> bessel_yl 3 5;
-0.0154429099129942
> bessel_yl_e 3 5;
[-0.0154429099129942, 2.87258769784673e-17]
> bessel_i0_scaled 3;
0.166253541303889
> bessel_i0_scaled_e 3;
[0.166253541303889, 7.38314037924188e-17]
> bessel_i1_scaled 3;
0.111661944928148
> bessel_i1_scaled_e 3;
[0.111661944928148, 4.95878648934625e-17]
> bessel_i2_scaled 3;
0.0545915963757409
```

```
> besseli2_scaled_e 3;
[0.0545915963757409, 2.42435388989563e-17]
> besseli_scaled 3 1;
0.0037027398773348
> besseli_scaled_e 3 1;
[0.0037027398773348, 8.46838615599053e-17]
> besseli_scaled_array 3 1;
[0.432332358381693, 0.135335283236613, 0.0263265086718556, 0.0037027398773348]
> besserk0_scaled 3;
0.523598775598299
> besserk0_scaled_e 3;
[0.523598775598299, 2.32524566533909e-16]
> besserk1_scaled 4;
0.490873852123405
> besserk1_scaled_e 4;
[0.490873852123405, 2.17991781125539e-16]
> besserk2_scaled 4;
0.760854470791278
> besserk2_scaled_e 4;
[0.760854470791278, 3.37887260744586e-16]
> besserk_scaled 2 4;
0.760854470791278
> besserk_scaled_e 2 4;
[0.760854470791278, 3.37887260744586e-16]
> besserk_scaled_array 2 4;
[0.392699081698724, 0.490873852123405, 0.760854470791278]
> besseljnu 2 2.3;
0.413914591732062
> besseljnu_e 2 2.3;
[0.413914591732062, 6.43352513956959e-16]
> besseljnu_sequence 2 {.1,.2,.3};
[0.00124895865879992, 0.00498335415278356, 0.011165861949064]
> besselynu 1 0.5;
-1.47147239267024
> besselynu_e 1 0.5;
[-1.47147239267024, 8.49504515830242e-15]
> besseljnu 1.2 3.4;
5.25626563437082
> besseljnu_e 1.2 3.4;
[5.25626563437082, 1.00839636820646e-13]
> besseljnu_scaled 1.2 3.4;
0.175418771999042
> besseljnu_scaled_e 1.2 3.4;
[0.175418771999042, 3.15501414592188e-15]
> besserknu 3 3;
0.122170375757184
> besserknu_e 3 3;
[0.122170375757184, 4.34036365096743e-16]
> besserklnknu 3 3;
-2.10233868587978
> besserklnknu_e 3 3;
[-2.10233868587978, 4.24157124665032e-15]
```

```
> bessel_Knu_scaled 3 3;  
2.45385759319062  
> bessel_Knu_scaled_e 3 3;  
[2.45385759319062,7.6281217575122e-15]  
> bessel_zero_J0 3;  
8.65372791291102  
> bessel_zero_J0_e 3;  
[8.65372791291102,2.59611837387331e-14]  
> bessel_zero_J1 3;  
10.1734681350627  
> bessel_zero_J1_e 3;  
[10.1734681350627,2.03469362701254e-13]  
> bessel_zero_Jnu 1.2 3;  
10.46769  
> bessel_zero_Jnu_e 1.2 3;  
[10.4676986203553,2.09353972407105e-14]86203553
```

2.5 Clausen Functions

clausen x implements `gsl_sf_clausen`.

clausen_e x implements `gsl_sf_clausen_e`.

2.6 Examples

The following illustrate the Clausen functions.

```
> using gsl::sf;  
> using namespace gsl::sf;  
> clausen 4.5;  
-0.831839220823219  
> clausen_e 4.5;  
[-0.831839220823219,8.60688668835964e-16]
```

2.7 Colomb Functions

The results of the Coulomb wave functions are returned as a list whose elements are ordered corresponding to the argument order of the corresponding C functions in GSL library.

hydrogenicR_1 Z r implements `gsl_sf_hydrogenicR_1`.

hydrogenicR_1_e Z r implements `gsl_sf_hydrogenicR_1_e`.

hydrogenicR n::int l::int Z r implements `gsl_sf_hydrogenicR_1`.

hydrogenicR_e n::int l::int Z r implements `gsl_sf_hydrogenicR_1_e`.

coulomb_wave_FG_e eta x L_F k::int implements `gsl_sf_coulomb_wave_FG_e`.

coulomb_wave_F_array L_min kmax::int eta x implements `gsl_sf_coulomb_wave_F_array`.

coulomb_wave_FG_array *L_min kmax::int eta x* implements
`gsl_sf_coulomb_wave_FG_array`.

coulomb_wave_FGp_array *L_min kmax::int eta x* implements
`gsl_sf_coulomb_wave_FGp_array`.

coulomb_wave_sphF_array *L_min kmax::int eta x* implements
`gsl_sf_coulomb_wave_sphF_array`.

coulomb_CL_e *L eta* implements `gsl_sf_coulomb_wave_CL_e`.

coulomb_CL_array *Lmin kmax eta* implements `gsl_sf_coulomb_wave_CL_array`.

2.8 Examples

The following illustrate the Coulomb functions.

```
> using gsl::sf;
> using namespace gsl::sf;
> hydrogenicR_1 0.2 4;
0.0803784086420537
> hydrogenicR_1_e 0.2 4;
[0.0803784086420537, 2.85561471862841e-17]
> hydrogenicR 3 1 0.25 3.2;
0.00802954301593587
> hydrogenicR_e 3 1 0.25 3.2;
[0.00802954301593587, 3.90138748076797e-17]
> coulomb_wave_F_array 1 2 0.5 0.5;
[{0.0387503306520188, 0.0038612830533923, 0.000274978904710252}, 0.0]
> coulomb_wave_FG_array 1 2 0.5 0.5;
[{0.0387503306520188, 0.0038612830533923, 0.000274978904710252},
 {4.13731494044202, 25.4479852847406, 257.269816591168}, 0.0, 0.0]
> coulomb_wave_FGp_array 1 2 0.5 0.5;
[{0.0387503306520188, 0.0038612830533923, 0.000274978904710252},
 {4.13731494044202, 25.4479852847406, 257.269816591168}, 0.0, 0.0]
> coulomb_wave_sphF_array 1 2 0.5 0.5;
[{0.0775006613040376, 0.0077225661067846, 0.000549957809420504}, 0.0]
> coulomb_CL_e (-0.5) 3;
[0.000143036170217949, 2.92195771135514e-18]
> coulomb_CL_array (-0.5) 4 1.5;
[0.0159218263353144, 0.0251746178646226, 0.00890057150292734,
 0.00172996014234001, 0.000235267570111599]
```

2.9 Coupling Coefficients

coupling_3j *m::matrix* implements `gsl_sf_coupling_3j` except the input is a 2x3 (row by column) integer matrix instead of six integer arguments.

coupling_3j_e *m::matrix* implements `gsl_sf_coupling_3j_e` except the input is a 2x3 (row by column) integer matrix instead of six integer arguments.

coupling_6j **m::matrix** implements `gsl_sf_coupling_6j` except the input is a 2x3 (row by column) integer matrix instead of six integer arguments.

coupling_6j_e **m::matrix** implements `gsl_sf_coupling_6j_e` except the input is a 2x3 (row by column) integer matrix instead of six integer arguments.

coupling_9j **m::matrix** implements `gsl_sf_coupling_9j` except the input is a 3x3 integer matrix instead of six integer arguments.

coupling_9j_e **m::matrix** implements `gsl_sf_coupling_9j_e` except the input is a 3x3 integer matrix instead of six integer arguments.

2.10 Examples

The following illustrate the coupling coefficient functions.

```
> using gsl::sf;
> using namespace gsl::sf;
> coupling_3j {6,4,2;0,0,0};
-0.29277002188456
> coupling_3j_e {6,4,2;0,0,0};
[-0.29277002188456,1.300160076865e-16]
> coupling_6j {1,2,3;2,1,2};
-0.166666666666667
> coupling_6j_e {1,2,3;2,1,2};
[-0.166666666666667,2.22044604925031e-16]
> coupling_9j {1,2,3;2,1,2;1,1,1};
-0.0962250448649376
> coupling_9j_e {1,2,3;2,1,2;1,1,1};
[-0.0962250448649376,4.84948508304183e-16]
```

2.11 Dawson Function

dawson **x** implements `gsl_sf_dawson`.

dawson_e **x** implements `gsl_sf_dawson_e`.

2.12 Examples

The following illustrate the dawson functions.

```
> dawson 3;/**-
0.178271030610558
> dawson_e 3;
[0.178271030610558,8.9920386788099e-16]
```

2.13 Debye Functions

debye_1 *x* implements `gsl_sf_debye_1`.
debye_1_e *x* implements `gsl_sf_debye_1_e`.
debye_2 *x* implements `gsl_sf_debye_2`.
debye_2_e *x* implements `gsl_sf_debye_2_e`.
debye_3 *x* implements `gsl_sf_debye_3`.
debye_3_e *x* implements `gsl_sf_debye_3_e`.
debye_4 *x* implements `gsl_sf_debye_4`.
debye_4_e *x* implements `gsl_sf_debye_4_e`.
debye_5 *x* implements `gsl_sf_debye_5`.
debye_5_e *x* implements `gsl_sf_debye_5_e`.
debye_6 *x* implements `gsl_sf_debye_6`.
debye_6_e *x* implements `gsl_sf_debye_6_e`.

2.14 Examples

The following illustrate the debye functions.

```
> debye_1 0.4;
0.904437352623294
> debye_1_e 0.4;
[0.904437352623294, 3.84040456356756e-16]
> debye_2 1.4;
0.613281386045505
> debye_2_e 1.4;
[0.613281386045505, 5.15090106564116e-16]
> debye_3 2.4;
0.370136882985216
> debye_3_e 2.4;
[0.370136882985216, 6.0792125556598e-16]
> debye_4 3.4;
0.205914922541978
> debye_4_e 3.4;
[0.205914922541978, 7.42872979584512e-16]
> debye_5 4.4;
0.107477287722471
> debye_5_e 4.4;
[0.107477287722471, 2.38647518907499e-17]
> debye_6 5.4;
0.0533132925698824
> debye_6_e 5.4;
[0.0533132925698824, 1.18379289859322e-17]
```

2.15 Dilogarithm

dilog *x* implements `gsl_sf_dilog`.

dilog (*r<:theta*) implements `gsl_sf_complex_dilog_e` except that results are returned as the complex value `re+im` and the error values are not returned.

dilog_e *x* implements `gsl_sf_dilog_e`.

dilog_e (*r<:theta*) implements `gsl_sf_complex_dilog_e` except the results are returned as the list `[re+im, re_error, im_error]`.

2.16 Examples

The following illustrate the `dilog` functions.

```
> dilog 1.0;
1.64493406684823
> dilog (1<:2);
-0.496658586741567+0.727146050863279i
> dilog_e (1%3);
[0.366213229977064, 8.22687466397711e-15]
> dilog_e (1<:3);
[-0.817454913536463+0.0980262093913011i, 3.8224192909699e-15,
 1.47247478976757e-15]
```

multiply_e *x y* implements `gsl_sf_multiply_e`.

multiply_err_e *x dx y dy* implements `gsl_sf_multiply_err_e`.

2.17 Examples

The following illustrate the `multiply` functions.

```
> multiply_e 10.0 11.0;
[110.0, 4.88498130835069e-14]
> multiply_err_e 10.0 0.04 11.0 0.002;
[110.0, 0.4600000000000049]
```

3 Matrices

This module is loaded via the command using `gsl::matrix` and provides wrappers for many of the GSL matrix, BLAS, and linear algebra routines found in Chapters 8, 12, and 13, respectively of the GSL Reference Manual:

- [Vectors and Matrices](#)
- [BLAS Support](#)

- [Linear Algebra](#)

It also contains some general utility functions for creating various types of matrices.

3.1 Matrix Creation

The utility functions `zeros` and `ones` create matrices with all elements zero or one, respectively, and `eye` creates identity matrices. These functions can be invoked either with a pair `(n,m)` denoting the desired number of rows or columns, or an integer `n` in which case a square `n x n` matrix is created. The result is always a double matrix. Analogous functions `izeros`, `czeros`, etc. are provided to create integer and complex matrices, respectively.

`zeros (n :: int, m :: int)` creates an `n x m` double matrix with all of its entries being zero.

`zeros n :: int` creates an `n x n` double matrix with all of its entries being zero.

`izeros (n :: int, m :: int)` creates an `n x m` integer matrix with all of its entries being zero.

`izeros n :: int` creates an `n'x'n` integer matrix with all of its entries being zero.

`czeros (n :: int, m :: int)` creates an `n x m` complex matrix with all of its entries being zero.

`czeros n :: int` creates an `n x n` complex matrix with all of its entries being zero.

`ones (n :: int, m :: int)` creates an `n x m` double matrix with all of its entries being one.

`ones n :: int` creates an `n x n` double matrix with all of its entries being one.

`iones (n :: int, m :: int)` creates an `n x m` integer matrix with all of its entries being one.

`iones n :: int` creates an `n x n` integer matrix with all of its entries being one.

`cones (n :: int, m :: int)` creates an `n x m` complex matrix with all of its entries being one.

`cones n :: int` creates an `n x n` complex matrix with all of its entries being one.

`eye (n :: int, m :: int)` creates an `n x m` identity matrix with double entries.

`eye n :: int` creates an `n x n` identity matrix with double entries.

`ieye (n :: int, m :: int)` creates an `n x m` identity matrix with integer entries.

`ieye n :: int` creates an `n x n` identity matrix with integer entries.

`ceye (n :: int, m :: int)` creates an `n x m` identity matrix with complex entries.

`ceye n :: int` creates an `n x n` identity matrix with complex entries.

3.2 Matrix Operators and Functions

The following operations are defined for constant a and matrices x and y . Some operators are not defined in the GSL library but are provided here for convenience.

$a + x, x + a$ returns a matrix with entries $a + x!(i, j)$.

$x + y$ adds matrix x to matrix y .

$-x$ returns a matrix with entries $-x!(i, j)$. Note that `neg x` is equivalent to $-x$.

$a - x$ returns a matrix with entries $a - x!(i, j)$.

$x - a$ returns a matrix with entries $x!(i, j) - a$.

$x - y$ subtracts matrix y from matrix x .

$a * x, x * a$ returns a matrix with entries $a * x!(i, j)$.

$x .* y$ multiplies, element-wise, matrix x to matrix y .

$x * y$ multiplies matrix x to matrix y .

a / x returns a matrix with entries $a / x!(i, j)$. Note that matrix x must not have any zero entries.

x / a returns a matrix with entries $x!(i, j) / a$. Note that a must be nonzero.

$x ./ y$ divides, element-wise, matrix x by matrix y .

x / y right divides matrix x by matrix y .

$x \backslash y$ left divides matrix x by matrix y .

$a \text{ div } x$ returns an integer matrix with entries $a \text{ div } x!(i, j)$. Note that a must be an integer and matrix x must be an integer matrix with nonzero entries.

$x \text{ div } a$ returns an integer matrix with entries $x!(i, j) \text{ div } a$. Note that a must be a nonzero integer and matrix x must have integer entries.

$x \text{ div } y$ computes the quotient integer matrix x by integer matrix y .

$a \text{ mod } x$ returns an integer matrix with entries $a \text{ mod } x!(i, j)$. Note that a must be an integer and matrix x must be an integer matrix with nonzero entries.

$x \text{ mod } a$ returns an integer matrix with entries $x!(i, j) \text{ mod } a$. Note that a must be an integer and matrix x must be an integer matrix with nonzero entries.

$x \text{ mod } y$ returns the remainder integer matrix $x \text{ mod } y$.

`not` x returns a matrix with integer entries `not` $x!(i, j)$. Note that x must be a matrix with integer entries and `not` is the bitwise negation operation.

$a ^ x$ returns a matrix with entries $a ^ x!(i, j)$. Note that 0^0 is defined as 1.

$x ^ a$ returns a matrix with entries $x!(i, j) ^ a$. Note that 0^0 is defined as 1.

$x .^ y$ returns a matrix with entries $x!(i, j) ^ y!(i, j)$.

`x ^ y` returns a matrix with entries $x!(i, j) \wedge y!(i, j)$.

`x << a` returns an integer matrix with entries $x!(i, j) << a$. Note that `a` must be an integer and matrix `x` must have integer entries.

`x << y` returns an integer matrix with entries $x!(i, j) << y!(i, j)$. Note that `x` and `y` must have integer entries.

`x >> a` returns an integer matrix with entries $x!(i, j) >> a$. Note that `a` must be an integer and matrix `x` must have integer entries.

`x >> y` returns an integer matrix with entries $x!(i, j) >> y!(i, j)$. Note that `x` and `y` must have integer entries.

`x and a, a and x` returns an integer matrix with entries `a and x!(i, j)`. Note that `a` must be an integer, matrix `x` must have integer entries, and `and` is a bitwise operator.

`x and y` returns an integer matrix with entries $x!(i, j) \text{ and } y!(i, j)$. Note that `x` and `y` must be matrices with integer entries.

`x or a, a or x` returns an integer matrix with entries `a or x!(i, j)`. Note that `a` must be an integer, matrix `x` must have integer entries, and `or` is a bitwise operator.

`x or y` returns an integer matrix with entries $x!(i, j) \text{ or } y!(i, j)$. Note that `x` and `y` must be matrices with integer entries.

The `pow` function computes powers of matrices by repeated matrix multiplication.

`pow x :: matrix k :: int, pow x :: matrix k :: bigint` Raises matrix `x` to the `k`th power. Note `x` must be a square matrix and `k` a nonnegative integer.

3.3 Singular Value Decomposition

For a given $n \times m$ matrix `x`, these functions yield a singular-value decomposition `u, s, v` of the matrix such that $x == u*s*transpose\ v$, where `u` and `v` are orthogonal matrices of dimensions $n \times m$ and $n \times n$, respectively, and `s` is a $n \times n$ diagonal matrix which has the singular values in its diagonal, in descending order. Note that GSL implements this only for double matrices right now. Also, GSL only handles the case of square or overdetermined systems, but we work around that in our wrapper functions by just adding a suitable number of zero rows in the underdetermined case.

`svd x` singular-value decomposition of matrix `x`.

`svd_mod` This uses the modified Golub-Reinsch algorithm, which is faster if $n > m$ but needs $O(m^2)$ extra memory as internal workspace.

`svd_jacobi x` This uses one-sided Jacobi orthogonalization which provides better relative accuracy but is slower.

`svd_solve (u, s, v) b` Solve the system $Ax=b$, using the SVD of `A`. `svd_solve` takes the result `(u, s, v)` of a `svd` call, and a column vector `b` of the appropriate dimension. The result is another column vector solving the system (possibly in the least-squares sense).

pinv **x** Computes the pseudo inverse of a matrix from its singular value decomposition.

4 Least-Squares Fitting

This module is loaded via the command using `gsl::fit` and provides Pure wrappers for the GSL least-squares fitting routines found in Chapter 36 of the GSL manual,

http://www.gnu.org/software/gsl/manual/html_node/Least_002dSquares-Fitting.html.

4.1 Routines

linear **x::matrix** **y::matrix** implements `gsl_fit_linear` without the `xstride`, `ystride`, and `n` parameters. Results are returned as a list [`c0`, `c1`, `cov00`, `cov01`, `cov11`, `sumsq`].

wlinear **x::matrix** **w::matrix** **y::matrix** implements `gsl_fit_wlinear` without the `xstride`, `wstride`, `ystride`, and `n` parameters. Results are given as a list [`c0`, `c1`, `cov00`, `cov01`, `cov11`, `chisq`].

linear_est **x** **c0::double** **c1::double** **cov00::double** **cov01::double** **cov11::double** implements `gsl_fit_linear_est`. Results are returned as a list [`y`, `y_err`].

mul **x::matrix** **y::matrix** implements `gsl_fit_mul` omitting the parameters `xstride`, `ystride`, and `n`. Results are returned as a list [`c1`, `cov11`, `sumsq`].

wmul **x::matrix** **w::matrix** **y::matrix** implements `gsl_fit_wmul` omitting the parameters `xstride`, `ystride`, and `n`. Results are returned as a list [`c1`, `cov11`, `sumsq`].

mul_est **x** **c1::double** **cov11::double** implements `gsl_fit_mul_est`. Results are returned as a list [`y`, `y_err`].

4.2 Examples

Usage of each implemented library routine is illustrated below.

```
> using gsl::fit;
> using namespace gsl::fit;
```

The following code determines the equation for the least-squares line through the points (1,0.01), (2,1.11), (3,1.9), (4,2.85), and (5,4.01).

```
> Y x = '(a + b * x)
> when
>   a:b:_ = linear {1,2,3,4,5} {0.01,1.11,1.9,2.85,4.01}
> end;
> Y x;
-0.946+0.974*x
```

```
> eval $ Y 2;
1.002
```

The following code illustrates estimating y-values without constructing an equation for the least-squares line determined by the points $\{x_1, x_2, x_3, \dots, x_n\}$, $\{y_1, y_2, y_3, \dots, y_n\}$. Here we estimate the y-value at $x = 1$, $x = 2$, and $x = 3$. Compare the output above at $x = 2$ to the output at $x = 2$ below.

```
> let c0:c1:cov00:cov01:cov11:_ = linear {1,2,3,4,5}
> {0.01,1.11,1.9,2.85,4.01};
> linear_est 1 c0 c1 cov00 cov01 cov11;
[0.028,0.0838570211729465]
> linear_est 2 c0 c1 cov00 cov01 cov11;
[1.002,0.0592958683214944]
> linear_est 3 c0 c1 cov00 cov01 cov11;
[1.976,0.0484148737476408]
```

Next, we determine a least-squares line through the points (1,0.01), (2,1.11), (3,1.9), (4,2.85), and (5,4.01) using weights 0.1, 0.2, 0.3, 0.4, and 0.5.

```
> W x = '(a + b * x)
> when
> a:b:_ = wlinear (matrix (1..5))
> {0.1, 0.2, 0.3, 0.4, 0.5}
> {0.01, 1.11, 1.9, 2.85, 4.01};
> end;
> W u;
-0.99+0.986*u
> eval $ W 2;
0.982
```

The least-squares slope for $Y = c_1 * X$ using the points (1,3), (2,5), and (3,7) is calculated below. Also, the y-values and standard error about $x = 1, 2$, and 3 are given.

```
> let c1:cov11:sumsq:_ = mul {1,2,3} {3,5,7};
> mul_est 1 c1 cov11;
[2.42857142857143,0.123717914826348]
> mul_est 2 c1 cov11;
[4.85714285714286,0.247435829652697]
> mul_est 3 c1 cov11;
[7.28571428571428,0.371153744479045]
```

The least-squares slope for $Y = c_1 * X$ using the points (1,3), (2,5), and (3,7), and weights 0.4, 0.9, and 0.4 is calculated below. The approximation of y-values and standard error about $x = 1, 2$, and 3 follows.

```
> let c1:cov11:sumsq:_ = wmul {1,2,3} {0.4,0.9,0.4} {3,5,7};
> mul_est 1 c1 cov11;
[2.44736842105263,0.362738125055006]
> mul_est 2 c1 cov11;
[4.89473684210526,0.725476250110012]
> mul_est 3 c1 cov11;
```

[7.34210526315789, 1.08821437516502]

5 Statistics

This module is loaded via the command using `gsl::stats` and provides Pure wrappers for the GSL Statistics routines found in Chapter 20 of the GSL manual,

http://www.gnu.org/software/gsl/manual/html_node/Statistics.html.

5.1 Routines

mean data::matrix implements `gsl_stats_mean` without stride and n arguments.

variance data::matrix implements `gsl_stats_variance` without stride and n arguments.

variance data::matrix mean implements `gsl_stats_variance_m` without stride and n arguments.

sd data::matrix implements `gsl_stats_sd` without stride and n arguments.

sd_m data::matrix mean implements `gsl_stats_sd_m` without stride and n arguments.

tss data::matrix implements `gsl_stats_tss` without stride and n arguments.

tss_m data::matrix mean implements `gsl_stats_tss_m` without stride and n arguments.

variance_with_fixed_mean data::matrix mean implements `gsl_stats_variance_with_fixed_mean` without stride and n arguments.

sd_with_fixed_mean data::matrix mean implements `gsl_stats_sd_with_fixed_mean` without stride and n arguments.

absdev data::matrix implements `gsl_stats_absdev` without stride and n arguments.

absdev_m data::matrix mean implements `gsl_stats_absdev_m` without stride and n arguments.

skew data::matrix mean implements `gsl_stats_skew` without stride and n arguments.

skew_m_sd data::matrix mean sd implements `gsl_stats_skew_m_sd` without stride and n arguments.

kurtosis data::matrix implements `gsl_stats_kurtosis` without stride and n arguments.

kurtosis_m_sd data::matrix mean sd implements `gsl_stats_kurtosis_m_sd` without stride and n arguments.

lag1_autocorrelation data::matrix implements `gsl_stats_lag1_autocorrelation` without stride and n arguments.

lag1_autocorrelation_m **data::matrix mean** implements `gsl_stats_lag1_autocorrelation_m` without stride and n arguments.

covariance **d1::matrix d2::matrix** implements `gsl_stats_covariance` without stride1, stride2, and n arguments.

covariance_m **d1::matrix d2::matrix mean1 mean2** implements `gsl_stats_covariance_m` without stride1, stride2, and n arguments.

correlation **d1::matrix d2::matrix** implements `gsl_stats_correlation` without stride1, stride2, and n arguments.

wmean **weight::matrix data::matrix** implements `gsl_stats_wmean` without stride and n arguments.

wvariance **weight::matrix data::matrix** implements `gsl_stats_wvariance` without stride and n arguments.

wvariance_m **weight::matrix data::matrix mean** implements `gsl_stats_wvariance_m` without stride and n arguments.

wsd **weight::matrix data::matrix** implements `gsl_stats_wsd` without stride and n arguments.

wsd_m **weight::matrix data::matrix mean** implements `gsl_stats_wsd_m` without stride and n arguments.

wvariance_with_fixed_mean **weight::matrix data::matrix mean** implements `gsl_stats_wvariance_with_fixed_mean` without stride and n arguments.

wsd_with_fixed_mean **weight::matrix data::matrix mean** implements `gsl_stats_wsd_with_fixed_mean` without stride and n arguments.

wtss **weight::matrix data::matrix** implements `gsl_stats_wtss` without stride and n arguments.

wtss_m **weight::matrix data::matrix mean** implements `gsl_stats_wtss_m` without stride and n arguments.

wabsdev **weight::matrix data::matrix** implements `gsl_stats_wabsdev` without stride and n arguments.

wabsdev_m **weight::matrix data::matrix mean** implements `gsl_stats_wabsdev_m` without stride and n arguments.

wskew **weight::matrix data::matrix** implements `gsl_stats_wskew` without stride and n arguments.

wskew_m_sd **weight::matrix data::matrix mean sd** implements `gsl_stats_wskew_m_sd` without stride and n arguments.

wkurtosis **weight::matrix data::matrix** implements `gsl_stats_wkurtosis` without stride and n arguments.

wkurtosis_m_sd **weight::matrix data::matrix** implements `gsl_stats_wkurtosis_m_sd` without stride and n arguments.

max data::matrix implements `gsl_stats_max` without `stride` and `n` arguments.

min data::matrix implements `gsl_stats_min` without `stride` and `n` arguments.

minmax data::matrix implements `gsl_stats_minmax` without `stride` and `n` arguments. Results are returned as a list `[min, max]`.

min_index data::matrix implements `gsl_stats_min_index` without `stride` and `n` arguments.

max_index data::matrix implements `gsl_stats_max_index` without `stride` and `n` arguments.

minmax_index data::matrix implements `gsl_stats_minmax_index` without `stride` and `n` arguments. Results are returned as a list `[min_index, max_index]`.

median_from_sorted_data data::matrix implements `gsl_stats_median_from_sorted_data` without `stride` and `n` arguments.

quantile_from_sorted_data data::matrix f::double implements `gsl_stats_quantile_from_sorted_data` without `stride` and `n` arguments.

5.2 Examples

The following illustrates the use of each function in the `stats` module.

```
> using gsl::stats;
> using namespace gsl::stats;
> mean {1,2,3,4,5};
3.0
> variance {1,2,3,4,5};
2.5
> variance_m {1,2,3,4,5} 4;
3.75
> sd {1,2,3,4,5};
1.58113883008419
> sd_m {1,2,3,4,5} 4;
1.93649167310371
> tss {1,2,3,4,5};
10.0
> tss_m {1,2,3,4,5} 4;
15.0
> variance_with_fixed_mean {0.0,1.2,3.4,5.6,6.0} 4.1;
6.314
> sd_with_fixed_mean {0.0,1.2,3.4,5.6,6.0} 4.1;
2.51276739870606
> absdev {2,2,3,4,4};
0.8
> absdev_m {2,2,3,4,4} 4;
1.0
> skew {1,1,1,1,2,2,2,2,2,2,2,3,30};
2.94796699504537
> skew_sd {1,2,2,3,3,3,3,3,3,3,4,4,5} 3 1;
```

```
0.0
> kurtosis {1,2,2,3,3,3,3,3,3,4,4,5};
-0.230769230769231
> kurtosis_m_sd {1,2,2,3,3,3,3,3,3,4,4,5} 3 1;
-0.230769230769231
> lag1_autocorrelation {1,2,3,4,5};
0.4
> lag1_autocorrelation_m {1,2,3,4,5} 2.5;
0.444444444444444
> covariance {1,2,3,4,5} {3.0,4.5,6.0,7.5,9.0};
3.75
> covariance_m {1,2,3,4,5} {3.0,4.5,6.0,7.5,9.0} 3 6;
3.75
> correlation {1,2,3,4} {2,3,4,5};
1.0
> wmean {0.4,0.2,0.3,0.3,0.3} {2,3,4,5,6};
3.93333333333333
> wvariance {0.4,0.2,0.3,0.3,0.3} {2,3,4,5,6};
2.7752808988764
> wvariance_m {0.4,0.2,0.3,0.3,0.3} {2,3,4,5,6} 3.0;
3.87640449438202
> wsd {0.4,0.2,0.3,0.3,0.3} {2,3,4,5,6};
1.66591743459164
> wsd_m {0.4,0.2,0.3,0.3,0.3} {2,3,4,5,6} 3.0;
1.96885867811329
> wvariance_with_fixed_mean {1,2,3,4} {1,2,3,4} 2.5;
1.25
> wsd_with_fixed_mean {1,2,3,4} {1,2,3,4} 2.5;
1.11803398874989
> wtss {1,1,2,2} {2,3,4,5};
6.83333333333333
> wtss_m {1,1,2,2} {2,3,4,5} 3.1;
10.06
> wabsdev {1,1,2,2} {2,3,4,5};
0.888888888888889
> wabsdev_m {1,1,2,2} {2,3,4,5} 3.1;
1.13333333333333
> wskew {1,1,2,2} {2,3,4,5};
-0.299254338484713
> wskew_m_sd {1,1,2,2} {2,3,4,5} 3.1 1.2;
1.33526234567901
> wkurtosis {1,1,2,2} {2,3,4,5};
-1.96206512878137
> wkurtosis_m_sd {1,1,2,2} {2,3,4,5} 3.1 1.2;
-0.681921939300412
> min {9,4,2,1,9};
1
> max {9.1,4.2,2.6,1.1,9.2};
9.2
> minmax {9.0,4.0,2.0,1.0,9.0};
[1.0,9.0]
> min_index {9.1,4.2,2.6,1.1,9.2};
```



```
3
> max_index {9,4,2,1,9};
0
> minmax_index {9,4,2,1,0,9};
[4,0]
> median_from_sorted_data {1.0,2.0,3.0};
2.0
> quantile_from_sorted_data {1.0,2.0,3.0} 0.25;
1.5
```

6 Random Number Distributions

This module is loaded via the command using `gsl::randist` and provides Pure wrappers for the GSL random distribution routines found in Chapter 19 of the GSL manual,

http://www.gnu.org/software/gsl/manual/html_node/Random-Number-Distributions.html.

There are two namespaces provided by `randist.pure`, `gsl::ran` for probability density functions and `gsl::cdf` for cumulative distribution functions. The two namespaces minimize typing of the prefixes `gsl_ran_` and `gsl_cdf_` respectively.

6.1 Routines

ugaussian_pdf *x* implements `gsl_ran_ugaussian`.

gaussian_pdf *x sigma* implements `gsl_ran_gaussian_pdf`.

gaussian_tail_pdf *x a sigma* implements `gsl_ran_gaussian_tail_pdf`.

ugaussian_tail_pdf *x a* implements `gsl_ran_ugaussian_tail_pdf`.

bivariate_gaussian_pdf *x a* implements `gsl_ran_bivariate_gaussian_pdf`.

exponential_pdf *x mu* implements `gsl_ran_exponential_pdf`.

laplace_pdf *x a* implements `gsl_ran_laplace_pdf`.

exppow_pdf *x a b* implements `gsl_ran_exppow_pdf`.

cauchy_pdf *x a* implements `gsl_ran_cauchy_pdf`.

rayleigh_pdf *x sigma* implements `gsl_ran_rayleigh_pdf`.

rayleigh_tail_pdf *x a sigma* implements `gsl_ran_rayleigh_tail_pdf`.

landau_pdf *x* implements `gsl_ran_landau_pdf`.

gamma_pdf *x a b* implements `gsl_ran_gamma_pdf`.

flat_pdf *x a b* implements `gsl_ran_flat_pdf`.

lognormal_pdf *x zeta sigma* implements `gsl_ran_lognormal_pdf`.

chisq_pdf *x nu* implements `gsl_ran_chisq_pdf`.

fdist_pdf *x nu1 nu2* implements `gsl_ran_fdist_pdf`.

tdist_pdf *x nu* implements `gsl_ran_tdist_pdf`.

beta_pdf *x a b* implements `gsl_ran_beta_pdf`.

logistic_pdf *x a* implements `gsl_ran_logistic_pdf`.

pareto_pdf *x a b* implements `gsl_ran_pareto_pdf`.

weibull_pdf *x a b* implements `gsl_ran_weibull_pdf`.

gumbell1_pdf *x a b* implements `gsl_ran_gumbell1_pdf`.

gumbell2_pdf *x a b* implements `gsl_ran_gumbell2_pdf`.

dirichlet_pdf *alpha::matrix theta::matrix* implements `gsl_ran_dirichlet_pdf`.

dirichlet_lnpdf *alpha::matrix theta::matrix* implements `gsl_ran_dirichlet_lnpdf`.

discrete_preproc *p::matrix* implements `gsl_ran_discrete_preproc` without the *K* parameter.

discrete_pdf *k::int p::pointer* implements `gsl_ran_discrete_pdf` without the *K* parameter.

discrete_free *p::pointer* implements `gsl_ran_discrete_free`

poisson_pdf *k::int mu* implements `gsl_ran_poisson_pdf`.

bernoulli_pdf *k::int p* implements `gsl_ran_bernoulli_pdf`.

binomial_pdf *k::int p n::int* implements `gsl_ran_binomial_pdf`.

multinomial_pdf *p::matrix n::matrix* implements `gsl_ran_multinomial_pdf`.

multinomial_lnpdf *p::matrix n::matrix* implements `gsl_ran_multinomial_lnpdf`.

negative_binomial_pdf *k::int p n* implements `gsl_ran_negative_binomial_pdf`.

pascal_pdf *k::int p n::int* implements `gsl_ran_pascal_pdf`.

geometric_pdf *k::int p* implements `gsl_ran_geometric_pdf`.

hypergeometric_pdf *k::int n1::int n2::int t::int* implements `gsl_ran_hypergeometric_pdf`.

logarithmic_pdf *k::int p* implements `gsl_ran_logarithmic_pdf`.

ugaussian_P *x* implements `gsl_cdf_ugaussian_P`.

ugaussian_Q *x* implements `gsl_cdf_ugaussian_Q`.

ugaussian_Pinv *p* implements `gsl_cdf_ugaussian_Pinv`.

ugaussian_Qinv *q* implements `gsl_cdf_ugaussian_Qinv`.

gaussian_P *x sigma* implements `gsl_cdf_gaussian_P`.

gaussian_Q *x sigma* implements `gsl_cdf_gaussian_Q`.
gaussian_Pinv *p sigma* implements `gsl_cdf_gaussian_Pinv`.
gaussian_Qinv *q sigma* implements `gsl_cdf_gaussian_Qinv`.
exponential_P *x mu* implements `gsl_cdf_exponential_P`.
exponential_Q *x mu* implements `gsl_cdf_exponential_Q`.
exponential_Pinv *p mu* implements `gsl_cdf_exponential_Pinv`.
exponential_Qinv *q mu* implements `gsl_cdf_exponential_Qinv`.
laplace_P *x a* implements `gsl_cdf_laplace_P`.
laplace_Q *x a* implements `gsl_cdf_laplace_Q`.
laplace_Pinv *p a* implements `gsl_cdf_laplace_Pinv`.
laplace_Qinv *q a* implements `gsl_cdf_laplace_Qinv`.
exppow_P *x a b* implements `gsl_cdf_exppow_P`.
exppow_Q *x a b* implements `gsl_cdf_exppow_Q`.
cauchy_P *x a* implements `gsl_cdf_cauchy_P`.
cauchy_Q *x a* implements `gsl_cdf_cauchy_Q`.
cauchy_Pinv *p a* implements `gsl_cdf_cauchy_Pinv`.
cauchy_Qinv *q a* implements `gsl_cdf_cauchy_Qinv`.
rayleigh_P *x sigma* implements `gsl_cdf_rayleigh_P`.
rayleigh_Q *x sigma* implements `gsl_cdf_rayleigh_Q`.
rayleigh_Pinv *p sigma* implements `gsl_cdf_rayleigh_Pinv`.
rayleigh_Qinv *q sigma* implements `gsl_cdf_rayleigh_Qinv`.
gamma_P *x a b* implements `gsl_cdf_gamma_P`.
gamma_Q *x a b* implements `gsl_cdf_gamma_Q`.
gamma_Pinv *p a b* implements `gsl_cdf_gamma_Pinv`.
gamma_Qinv *q a b* implements `gsl_cdf_gamma_Qinv`.
flat_P *x a b* implements `gsl_cdf_flat_P`.
flat_Q *x a b* implements `gsl_cdf_flat_Q`.
flat_Pinv *p a b* implements `gsl_cdf_flat_Pinv`.
flat_Qinv *q a b* implements `gsl_cdf_flat_Qinv`.
lognormal_P *x zeta sigma* implements `gsl_cdf_lognormal_P`.
lognormal_Q *x zeta sigma* implements `gsl_cdf_lognormal_Q`.

`lognormal_Pinv p zeta sigma` implements `gsl_cdf_lognormal_Pinv`.
`lognormal_Qinv q zeta sigma` implements `gsl_cdf_lognormal_Qinv`.
`chisq_P x nu` implements `gsl_cdf_chisq_P`.
`chisq_Q x nu` implements `gsl_cdf_chisq_Q`.
`chisq_Pinv p nu` implements `gsl_cdf_chisq_Pinv`.
`chisq_Qinv q nu` implements `gsl_cdf_chisq_Qinv`.
`fdist_P x nu1 nu2` implements `gsl_cdf_fdist_P`.
`fdist_Q x nu1 nu2` implements `gsl_cdf_fdist_Q`.
`fdist_Pinv p nu1 nu2` implements `gsl_cdf_fdist_Pinv`.
`fdist_Qinv q nu1 nu2` implements `gsl_cdf_fdist_Qinv`.
`tdist_P x nu` implements `gsl_cdf_tdist_P`.
`tdist_Q x nu` implements `gsl_cdf_tdist_Q`.
`tdist_Pinv p nu` implements `gsl_cdf_tdist_Pinv`.
`tdist_Qinv q nu` implements `gsl_cdf_tdist_Qinv`.
`beta_P x a b` implements `gsl_cdf_beta_P`.
`beta_Q x a b` implements `gsl_cdf_beta_Q`.
`beta_Pinv p a b` implements `gsl_cdf_beta_Pinv`.
`beta_Qinv q a b` implements `gsl_cdf_beta_Qinv`.
`logistic_P x a` implements `gsl_cdf_logistic_P`.
`logistic_Q x a` implements `gsl_cdf_logistic_Q`.
`logistic_Pinv p a` implements `gsl_cdf_logistic_Pinv`.
`logistic_Qinv q a` implements `gsl_cdf_logistic_Qinv`.
`pareto_P x a b` implements `gsl_cdf_pareto_P`.
`pareto_Q x a b` implements `gsl_cdf_pareto_Q`.
`pareto_Pinv p a b` implements `gsl_cdf_pareto_Pinv`.
`pareto_Qinv q a b` implements `gsl_cdf_pareto_Qinv`.
`weibull_P x a b` implements `gsl_cdf_weibull_P`.
`weibull_Q x a b` implements `gsl_cdf_weibull_Q`.
`weibull_Pinv p a b` implements `gsl_cdf_weibull_Pinv`.
`weibull_Qinv q a b` implements `gsl_cdf_weibull_Qinv`.
`gumbell_P x a b` implements `gsl_cdf_gumbell_P`.

gumbel1_Q *x a b* implements `gsl_cdf_gumbel1_Q`.
gumbel1_Pinv *p a b* implements `gsl_cdf_gumbel1_Pinv`.
gumbel1_Qinv *q a b* implements `gsl_cdf_gumbel1_Qinv`.
gumbel2_P *x a b* implements `gsl_cdf_gumbel2_P`.
gumbel2_Q *x a b* implements `gsl_cdf_gumbel2_Q`.
gumbel2_Pinv *p a b* implements `gsl_cdf_gumbel2_Pinv`.
gumbel2_Qinv *q a b* implements `gsl_cdf_gumbel2_Qinv`.
poisson_P *k::int mu* implements `gsl_cdf_poisson_P`.
poisson_Q *k::int mu* implements `gsl_cdf_poisson_Q`.
binomial_P *k::int p n::int* implements `gsl_cdf_binomial_P`.
binomial_Q *k::int q n::int* implements `gsl_cdf_binomial_Q`.
negative_binomial_P *k::int p n* implements `gsl_cdf_negative_binomial_P`.
negative_binomial_Q *k::int p n* implements `gsl_cdf_negative_binomial_Q`.
pascal_P *k::int p n::int* implements `gsl_cdf_pascal_P`.
pascal_Q *k::int p n::int* implements `gsl_cdf_pascal_Q`.
geometric_P *k::int p* implements `gsl_cdf_geometric_P`.
geometric_Q *k::int p* implements `gsl_cdf_geometric_Q`.
hypergeometric_P *k::int n1::int n2::int t::int* implements
 `gsl_cdf_hypergeometric_P`.
hypergeometric_Q *k::int n1::int n2::int t::int* implements
 `gsl_cdf_hypergeometric_Q`.

6.2 Examples

The following illustrates the use of each function in the `randist` module. The pdf functions are illustrated first.

```
> using gsl::stats;
> using namespace gsl::ran;
> ugaussian_pdf 1.2;
0.194186054983213
> gaussian_pdf (-1.3) 1.5;
0.182690978264686
> gaussian_tail_pdf 2.0 1.0 1.5;
0.433042698395299
> ugaussian_tail_pdf 2.0 1.0;
0.34030367841782
> bivariate_gaussian_pdf 1.2 0.9 1.0 1.0 0.95;
```

```
0.184646843689817
> exponential_pdf 1.0 0.5;
0.270670566473225
> laplace_pdf 1.5 2.0;
0.118091638185254
> exppow_pdf 0.0 1.0 1.5;
0.553866083716236
> cauchy_pdf (-1.0) 1.0;
0.159154943091895
> rayleigh_pdf 2.5 1.0;
0.109842334058519
> rayleigh_tail_pdf 1.5 1.0 1.0;
0.802892142778485
> landau_pdf 1.1;
0.140968737919623
> gamma_pdf 1.0 1.0 1.5;
0.342278079355061
> flat_pdf 1.0 0.5 2.5;
0.5
> lognormal_pdf 0.01 0.0 1.0;
0.000990238664959182
> chisq_pdf 1.0 2.0;
0.303265329856317
> fdist_pdf 0.5 3.0 2.0;
0.480970043785452
> tdist_pdf 0.1 10.0;
0.386975225815181
> beta_pdf 0.5 4.0 1.0;
0.499999999999999
> logistic_pdf (-1.0) 2.0;
0.117501856100797
> pareto_pdf 0.01 3.0 2.0;
0.0
> weibull_pdf 0.01 1.0 1.0;
0.990049833749168
> gumbell_pdf 0.01 1.0 1.0;
0.367861108816436
> gumbel2_pdf 0.01 1.0 1.0;
3.72007597602084e-40
> dirichlet_pdf {0.1,0.2,0.8} {2.0,2.0,2.0};
0.00501316294425874
> dirichlet_lnpdf {0.1,0.2,0.8} {2.0,2.0,2.0};
-5.29568823688856
> poisson_pdf 4 0.4;
0.000715008049104682
> bernoulli_pdf 1 0.7;
0.7
> binomial_pdf 3 0.5 9;
0.1640625
> multinomial_pdf {0.1,0.2,0.7} {2,2,2};
0.0
> multinomial_lnpdf {0.1,0.2,0.7} {2,2,2};
```

```
-1728120799.71174
> negative_binomial_pdf 10 0.5 3.5;
0.0122430486923836
> pascal_pdf 10 0.5 3;
0.00805664062499999
> geometric_pdf 5 0.4;
0.05184
> hypergeometric_pdf 1 5 20 3;
0.413043478260872
> logarithmic_pdf 10 0.7;
0.00234619293712492
> test_discrete
>   = v
>   when
>     px = discrete_preproc {0.1,0.3,0.4};
>     v = discrete_pdf 0 px +
>         discrete_pdf 1 px +
>         discrete_pdf 2 px;
>     _ = discrete_free px
>   end;
> test_discrete;
1.0
```

The cumulative distribution functions are shown.

```
> ugaussian_P (-1.3);
0.0968004845856103
> ugaussian_Q (-1.3);
0.90319951541439
> ugaussian_Pinv 0.84;
0.994457883209753
> ugaussian_Qinv 0.84;
-0.994457883209753
> gaussian_P (1.3) 1.5;
0.806937662858093
> gaussian_Q (1.3) 1.5;
0.193062337141907
> gaussian_Pinv 0.4 5.0;
-1.266735515679
> gaussian_Qinv 0.4 5.0;
1.266735515679
> exponential_P 1.0 0.5;
0.864664716763387
> exponential_Q 1.0 0.5;
0.135335283236613
> exponential_Pinv 0.6 0.5;
0.458145365937077
> exponential_Qinv 0.6 0.5;
0.255412811882995
> laplace_P 1.5 2.0;
0.763816723629493
> laplace_Q 1.5 2.0;
```

```
0.236183276370507
> laplace_Pinv 0.6 2.0;
0.446287102628419
> laplace_Qinv 0.4 2.0;
0.446287102628419
> exppow_P 0.0 1.0 2.5;
0.5
> exppow_Q 0.0 1.0 0.5;
0.5
> cauchy_P (-1.0) 1.0;
0.25
> cauchy_Q (-1.0) 1.0;
0.75
> cauchy_Pinv 0.75 1.0;
1.0
> cauchy_Qinv 0.25 1.0;
1.0
> rayleigh_P 1.5 2.0;
0.245160398010993
> rayleigh_Q 0.5 1.0;
0.882496902584595
> rayleigh_Pinv 0.5 1.0;
1.17741002251547
> rayleigh_Qinv 0.5 1.0;
1.17741002251547
> gamma_P 1.0 1.0 3.0;
0.283468689426211
> gamma_Q 1.0 1.0 3.0;
0.716531310573789
> gamma_Pinv 0.5 1.0 1.0;
0.693147180559945
> gamma_Qinv 0.5 1.0 1.0;
0.693147180559945
> flat_P 2.0 1.2 4.8;
0.222222222222222
> flat_Q 2.0 1.2 4.8;
0.777777777777778
> flat_Pinv 0.2 0.5 2.5;
0.9
> flat_Qinv 0.2 0.5 2.5;
2.1
> lognormal_P 0.01 0.0 1.0;
2.06064339597172e-06
> lognormal_Q 0.01 0.0 1.0;
0.999997939356604
> lognormal_Pinv 0.1 0.0 1.0;
0.27760624185201
> lognormal_Qinv 0.1 0.0 1.0;
3.60222447927916
> chisq_P 1.0 2.0;
0.393469340287367
> chisq_Q 1.0 2.0;
```



```
0.606530659712633
> chisq_Pinv 0.5 2.0;
0.221199216928595
> chisq_Qinv 0.5 2.0;
1.38629436111989
> fdist_P 1.0 3.0 2.0;
0.46475800154489
> fdist_Q 1.0 3.0 2.0;
0.53524199845511
> fdist_Pinv 0.5 3.0 2.0;
1.13494292261288
> fdist_Qinv 0.5 3.0 2.0;
1.13494292261288
> tdist_P 2.1 10.0;
0.968961377898891
> tdist_Q (-2.1) 10.0;
0.968961377898891
> tdist_Pinv 0.68 10.0;
0.482264205919689
> tdist_Qinv 0.68 10.0;
-0.482264205919689
> beta_P 0.75 2.0 2.0;
0.84375
> beta_Q 0.75 2.0 2.0;
0.15625
> beta_Pinv 0.75 2.0 2.0;
0.673648177666931
> beta_Qinv 0.25 2.0 2.0;
0.673648177666931
> logistic_P (-1.0) 2.0;
1
> logistic_Q (-1.0) 2.0;
0.622459331201855
> logistic_Pinv 0.75 1.0;
1.09861228866811
> logistic_Qinv 0.25 1.0;
1.09861228866811
> pareto_P 2.01 3.0 2.0;
0.0148512406901899
> pareto_Q 2.01 3.0 2.0;
0.98514875930981
> pareto_Pinv 0.1 3.0 2.0;
2.07148833730257
> pareto_Qinv 0.1 3.0 2.0;
4.30886938006377
> weibull_P 1.01 1.0 2.0;
0.639441117518024
> weibull_Q 1.01 2.0 3.0;
0.879160657465162
> weibull_Pinv 0.1 1.0 2.0;
0.324592845974501
> weibull_Qinv 0.1 1.0 2.0;
```

```
1.51742712938515
> gumbel1_P 1.01 1.0 1.0;
0.694739044426344
> gumbel1_Q 1.01 1.0 1.0;
0.305260955573656
> gumbel1_Pinv 0.1 1.0 1.0;
-0.834032445247956
> gumbel1_Qinv 0.1 1.0 1.0;
2.25036732731245
> gumbel2_P 1.01 1.0 1.0;
0.371539903071873
> gumbel2_Q 1.01 1.0 1.0;
0.628460096928127
> gumbel2_Pinv 0.1 1.0 1.0;
0.434294481903252
> gumbel2_Qinv 0.1 1.0 1.0;
9.4912215810299
> poisson_P 4 0.4;
0.999938756672898
> poisson_Q 4 0.6;
0.000394486018340255
> binomial_P 3 0.5 10;
0.171874999999999
> binomial_Q 3 0.5 10;
0.828125000000001
> negative_binomial_P 10 0.5 3.0;
0.98876953125
> negative_binomial_Q 10 0.5 3.0;
0.01123046875
> pascal_P 10 0.5 3;
0.98876953125
> pascal_Q 10 0.5 3;
0.01123046875
> geometric_P 5 0.4;
0.92224
> geometric_Q 5 0.6;
0.01024
> hypergeometric_P 1 5 20 3;
0.908695652173913
> hypergeometric_Q 1 5 20 3;
0.0913043478260873
```

7 Sorting

This module is loaded via the command using `gsl::sort` and provides Pure wrappers for the GSL sorting routines found in Chapter 11 of the GSL manual,

http://www.gnu.org/software/gsl/manual/html_node/Sorting.html.

7.1 Routines

sort_vector `m::matrix` implements `gsl_sort` and `gsl_sort_int` without `stride` and `n` parameters.

sort_vector_index `m::matrix` implements `gsl_sort_index` and `gsl_sort_int_index` without `stride` and `n` parameters.

7.2 Examples

Usage of each library routine is illustrated below.

```
> using gsl::sort;
> using namespace gsl;
> sort_vector {0,3,2,4,5};
{0,2,3,4,5}
> sort_vector_index {0.0,1.0,5.0,2.0,8.0,0.0};
{5,0,1,3,2,4}
```


pure-octave

Version 0.1, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

A Pure interface to [GNU Octave](#).

1 Introduction

This is an Octave module for the Pure programming language, based on Paul Kienzle's [octave_embed](#) which allows Octave to be embedded in other languages. It allows you to execute arbitrary Octave commands and Octave functions from Pure.

2 Copying

pure-octave is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

pure-octave is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

Please see the accompanying COPYING file for the precise license terms. The GPL can also be read online at <http://www.gnu.org/licenses/>.

3 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-octave-0.1.tar.gz>.

Run `make` to compile the module and `make install` (as root) to install it in the Pure library directory. This requires GNU `make`, and of course you need to have both Pure and Octave installed (including Octave's `mkoctfile` utility and the corresponding header files and libraries).

make tries to guess your Pure installation directory and platform-specific setup. If it gets this wrong, you can set some variables manually. In particular, `make install prefix=/usr` sets the installation prefix. Please see the Makefile for details.

NOTE: This release of pure-gen has been tested with Octave 3.2.4. Older versions might require some fiddling with the sources to get the embedded Octave interface working.

4 Basic Usage

Import this module into your Pure scripts as follows:

```
using octave;
```

This will add an embedded instance of the Octave interpreter to your program. (You can import this module as often as you want, but there's always only one instance of Octave in each process.)

Arbitrary Octave code can be executed with the `octave_eval` function:

```
> octave_eval "eig([1 2;3 4]);"
ans =

-0.37228
 5.37228

0
```

This prints the result on stdout and returns a result code (zero if everything went fine). To suppress the printing of the result, simply terminate the Octave statement with a semicolon:

```
> octave_eval "eig([1 2;3 4]);";
0
```

You can also set and get global variables in the Octave interpreter. This allows you to define values to be used when evaluating Octave code, and to transfer results back to Pure. However, before such globals can be accessed in Octave, you must explicitly declare them as globals:

```
> octave_eval "global x y ans";
0
```

Now you can use `octave_set` and `octave_get` to transfer values between Pure and Octave as follows:

```
> octave_set "x" {1.0,2.0;3.0,4.0};
{1.0,2.0;3.0,4.0}
> octave_eval "eig(x);";
0
> octave_get "ans";
{-0.372281323269014;5.37228132326901}
```

Note that the most recent result can always be accessed through Octave's `ans` variable. You can also use an explicit variable definition as follows:

```
> octave_eval "y = eig(x);";
0
> octave_get "y";
{-0.372281323269014;5.37228132326901}
```

5 Direct Function Calls

Octave functions can also be called in a direct fashion with the `octave_call` function, which gets invoked with the name of the function, the number of function results and the function arguments:

```
> let x = {1.0,2.0;3.0,4.0};
> octave_call "eig" 1 x;
{-0.372281323269014;5.37228132326901}
```

Note the second argument, which denotes the desired number of *return* values. This will usually be 1 (or 0 if you don't care about the result), but some Octave functions may return a variable number of results, depending on how they're called. Multiple values are returned as tuples in Pure:

```
> octave_call "eig" 2 x;
{-0.824564840132394,-0.415973557919284;0.565767464968992,-0.909376709132124},
{-0.372281323269014,0.0;0.0,5.37228132326901}
```

If there are multiple *arguments*, you can specify them either as a tuple or a list:

```
> octave_call "norm" 1 (x, 2);
5.46498570421904
> octave_call "norm" 1 [x, 1];
6.0
```

Instead of a function name, you can also specify the function to be called using a special kind of Octave object, a function value. These are returned, e.g., by Octave's `str2func` and `inline` builtins. For your convenience, pure-octave provides a frontend to these builtins, the `octave_func` function, which lets you specify an Octave function in one of two ways:

- `octave_func name`, where `name` is the function name as a string, returns the Octave function with the given name. This works like Octave's `str2func` builtin.
- `octave_func expr`, where `expr` is an Octave expression (as a string) describing the function value, returns an "inline" function. This works like Octave's `inline` builtin. Instead of just an Octave expression, you can also specify a tuple or a list consisting of the inline expression and the parameter names. (Otherwise the parameters are determined automatically, see the description of the `inline` function in the Octave manual for details.)

Note that inline functions allow you to call stuff that is not an Octave function and hence cannot be specified directly in `octave_call`, such as an operator. Examples:

```
> let eig = octave_func "eig";
> let mul = octave_func "x*y";
> octave_call eig 1 (octave_call mul 1 (x,x));
{0.138593383654928;28.8614066163451}
> let add = octave_func ("x+y","x","y");
> octave_call add 1 (x,x);
{2.0,4.0;6.0,8.0}
```

6 Data Conversions

As shown above, the `octave_set`, `octave_get` and `octave_call` functions convert Pure data to corresponding Octave values and vice versa. Octave scalars and matrices of boolean, integer, double, complex and character data are all supported by this interface, and are mapped to the corresponding Pure data types in a straightforward manner (scalars to scalars, matrices to matrices and strings to strings). Note that in any case these conversions create *copies* of the data, so modifying, say, an Octave matrix received via `octave_get` in Pure only affects the Pure copy of the matrix and leaves the original Octave matrix unchanged.

Any other kind of Octave object (including Octave function objects, see [Direct Function Calls](#)) is just passed through as is, in the form of a cooked pointer to an Octave value which frees itself when garbage-collected. You can check for such objects with the `octave_valuep` predicate:

```
> let eig = octave_func "eig";
> eig; octave_valuep eig;
#<pointer 0x230dba0>
1
```

Such Octave value pointers can be used freely wherever an Octave value is needed (i.e., in invocations of `octave_set` and `octave_call`).

You should also note the following:

- Octave's cell arrays and structures are not supported at this time, so they will show up as opaque Octave value pointers in Pure land. This allows these objects to be passed around freely, but you can't inspect or modify them in Pure.
- Scalars and 1x1 matrices are indistinguishable in Octave, thus any 1x1 matrix will be returned as a scalar from Octave to Pure.
- All types of boolean and integer matrices are returned from Octave to Pure as (machine) integer matrices. When converted back to Octave, these all become Octave `int32` matrices, but you can easily convert them to boolean or other types of matrices in Octave as needed. For instance:

```
> octave_set "a" {1,2;3,4};
{1,2;3,4}
```



```
> octave_eval "global a ans";
0
> octave_eval "eig(a)";
error: eig: wrong type argument 'int32 matrix'
1
> octave_eval "eig(double(a))";
ans =

    -0.37228
     5.37228

0
> octave_eval "a>2";
ans =

     0     0
     1     1

0
> octave_get "ans";
{0,0;1,1}
```

- Octave strings are mapped to Pure strings, and character matrices with more than one row are mapped to (symbolic) column vectors of Pure strings. Example:

```
> octave_set "a" "Hello, world!";
"Hello, world!"
> octave_eval "a";
a = Hello, world!
0
> octave_eval "[a;'abc']";
ans =

Hello, world!
abc

0
> octave_get "ans";
{"Hello, world!";"abc" }
```

7 Calling Back Into Pure

The embedded Octave interpreter provides one special builtin, the `pure_call` function which can be used to call any function defined in the executing Pure script from Octave. For instance:

```
> even m::matrix = {~(int x mod 2) | x=m};
> octave_eval "pure_call('even', 1:12)";
ans =
```

```
0 1 0 1 0 1 0 1 0 1 0 1
```

```
0
```

Here's the description of the `pure_call` function, as it is printed with Octave's help command:

'pure_call' is a built-in function

```
RES = pure_call(NAME, ARG, ...)
[RES, ...] = pure_call(NAME, ARG, ...)
```

Execute the Pure function named NAME (a string) with the given arguments. Arguments and result types may be scalars and matrices of boolean, integer, double, complex and character data. The Pure function may return multiple results as a tuple. Example: `pure('succ', 99) => 100`.

8 Caveats and Notes

Directly embedding Octave in Pure programs is convenient because it allows easy exchange of data between Pure and Octave, and you can also call Octave functions directly from Pure and vice versa. However, it also comes at a cost. A default build of Octave pulls in quite a few dependencies of its own which might conflict with other modules loaded in a Pure script. Specifically, we have found that building Octave with native graphics support (as is included in the latest Octave versions) may give problems with third-party graphics libraries such as [VTK](#) (which segfaults, apparently due to FreeType library compatibility issues, if used in the same program as Octave).

We're not sure whether Octave or VTK is to blame here, but until these quirks get ironed out you can work around them by disabling native graphics support in Octave. Unfortunately, Octave currently doesn't provide an option to do so, but the following patches available at the Pure website will do the trick (there are different versions of the patch, corresponding to the current stable and development versions of Octave):

- Octave 3.2.4: <http://pure-lang.googlecode.com/svn/wiki/octave-3.2.4.patch>
- Octave 3.3.51: <http://pure-lang.googlecode.com/svn/wiki/octave-3.3.51.patch>

The patch is to be applied to Octave's `configure.in` (`configure.ac` for the latest development versions) by running a command like the following in the Octave source directory:

```
patch -p0 < octave-3.2.4.patch
```

Then just rerun `autoconf` and `configure` and compile Octave as usual. The resulting version of Octave will have native graphics disabled and should work ok with third party libraries such as VTK.

(If you have an Octave version which differs considerably from what we provide here then you may have to rework the patch accordingly or apply the necessary changes to `configure.in` or `configure.ac` manually.)

Pure-Rational - Rational number library for the Pure programming language

Version 0.1, November 14, 2010

Rob Hubbard

Albert Graef <Dr.Graef@t-online.de>

Jiri Spitz <jiri.spitz@bluetone.cz>

This package provides a Pure port of [Q+Q](#), Rob Hubbard's rational number library for the Q programming language. The port was done by Jiri Spitz. It contains `rational.pure`, a collection of utility functions for rational numbers, and `rat_interval.pure`, a module for doing interval arithmetic needed by `rational.pure`. These modules are designed to work with the `math.pure` module (part of the standard Pure library), which contains the definition of Pure's rational type and implements the basic rational arithmetic.

This document is an edited version of Rob's original [Q+Q manual](#) available from the Q website, slightly adjusted to account for the Pure specifics of the implementation. In particular, note that the operations provided by `rational.pure` and `rat_interval.pure` live in their own `rational` and `interval` namespaces, respectively, so if you want to get unqualified access to the symbols of these modules (as the examples in this manual assume) then you'll have to import the modules as follows:

```
using rational, rat_interval;  
using namespace rational, interval;
```

Also note that `rational` always pulls in the `math` module, so you don't have to import the latter explicitly if you are using `rational`.

Another minor difference to the Q version of this module is that `rational` results always have Pure bigints as their numerators and denominators, hence the `L` suffix in the printed results. Also, unary minus binds weaker in Pure than the rational division operator, so a negative rational number will be printed as, e.g., `(-1L)%2L`, which looks a bit funny but is correct since Pure rationals always carry their sign in the numerator.

1 Copying

Copyright (c) 2006 - 2010 by Rob Hubbard.

Copyright (c) 2006 - 2010 by Albert Graef <Dr.Graef@t-online.de>.

Copyright (c) 2010 by Jiri Spitz <jiri.spitz@bluetone.cz>.

Pure-rational is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

Pure-rational is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

2 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-rational-0.1.tar.gz>.

Then run `make install` (as root) to install pure-rational in the Pure library directory. This requires GNU make, and of course you need to have Pure installed.

`make install` tries to guess your Pure installation directory. If it gets this wrong, you can install using `make install prefix=/usr` which sets the installation prefix. Please see the Makefile for details.

3 Introduction

3.1 The Rational Module

This module provides additional operations on the rational number type provided by the `math.pure` module in the standard library. The module is compatible with Pure version 0.43 (onwards).

3.2 The Files and the Default Prelude

The implementation of the rational type and associated utilities is distributed across various files.

math.pure and Other Files

The file `math.pure` defines the type, its constructors and ‘deconstructors’ and basic arithmetical and mathematical operators and functions. This is part of the standard Pure library. A few definitions associated with rationals are also defined in other standard library modules. In particular, the type tests are contained in `primitives.pure`.

It is also possible to create rational complex numbers (in addition to double complex numbers and integral or Gaussian complex numbers). That is, rationals play nicely with the complex number constructors provided in the `math.pure` module. This is discussed further in [Rational Complex Numbers](#).

rational.pure

Additional ‘rational utilities’, not included in the `math.pure` module, are defined in `rational.pure`. The functions include further arithmetical and mathematical operators and functions, continued fraction support, approximation routines and string formatting and evaluation.

The rational utilities include some ‘rational complex number’ functions.

rat_interval.pure

Amongst the rational utilities are some functions that return a rational interval. The file `rat_interval.pure` is a partial implementation of interval arithmetic. Intervals are discussed further in [Intervals](#).

3.3 Notation

Throughout this document, the parameters q, q_0, q_1, \dots usually denote rationals ($\in \mathbf{Q}$), parameters z, \dots usually denote integers ($\in \mathbf{Z}$), r, \dots usually denote real numbers ($\in \mathbf{R}$), c, \dots usually denote complex numbers ($\in \mathbf{C}$), n, \dots usually denote parameters of any numeric type, v, \dots usually denote parameters of any interval type, and x, \dots usually denote parameters of any type.

The reals are not just the doubles, but include rationals and integers. The term ‘rational’ usually refers to a rational number $\in \mathbf{Q} \supset \mathbf{Z}$, or an expression of type `rational` or `integer`.

4 The Rational Type

4.1 Constructors

Rationals are constructed with the `%exact` division operator, and other kinds of numbers can be converted to rationals with the `rational` function. These are both defined in `math.pure`.

n1 % n2 is the exact division operator, which may be used as a constructor (for integers *n1* and *n2*). This is described in [More on Division](#).

rational x converts the given number *x* to a rational.

Example 1 Constructing a fraction:

```
> 44%14;  
22L%7L
```

Example 2 Converting from an integer:

```
> rational 3;  
3L%1L
```

4.2 ‘Deconstructors’

A rational number is in simplest form if the numerator and denominator are coprime (i.e. do not have a factor in common) and the denominator is positive (and, specifically, non-zero). Sometimes the term ‘irreducible’ is used for a rational in simplest form. This is a property of the representation of the rational number and not of the number itself.

num q given a rational or integer *q*, returns the ‘(signed) simplest numerator’, i.e. the numerator of the normalised form of *q*.

den q given a rational or integer *q*, returns the ‘(positive) simplest denominator’, i.e. the denominator of the normalised form of *q*.

num_den q given a rational or integer *q*, returns a pair (*n*, *d*) containing the (signed) simplest numerator *n* and the (positive) simplest denominator *d*. This is the inverse (up to equivalence) of *rational* as defined on integer pairs (see [Constructors](#)).

Example 3 Using *num_den* to obtain a representation in simplest form:

```
> let q = (44%(-14));  
> num q;  
-22L  
> den q;  
7L  
> num_den q;  
-22L, 7L  
> num_den 3;  
3L, 1L  
> num_den (-3);  
-3L, 1L
```

Together, *num* and *den* are a pair of ‘decomposition’ operators, and *num_den* is also a decomposition operator. There are others (see [Decomposition](#)). The *integer* and *fraction* function (see [Integer and Fraction Parts](#)) may be used in conjunction with *num_den_gauss* to decompose a rational into integer, numerator and denominator parts.

4.3 Type and Value Tests

The functions `rationalp` and `ratvalp` and other rational variants are new for rationals and the standard functions `exactp` and `inexactp` are extended for rationals.

A value is ‘exact’, or of an exact type, if it is of a type that is able to represent the values returned by arithmetical operations exactly; in a sense, it is ‘closed’ under arithmetical operations. Otherwise, a value is ‘inexact’. Inexact types are able to store some values only approximately.

The doubles are not an exact type. The results of some operations on some values that are stored exactly, can’t be stored exactly. (Furthermore, doubles are intended to represent real numbers; no irrational number ($\in \mathbf{R} \setminus \mathbf{Q}$) can be stored exactly as a double; even some rational ($\in \mathbf{Q}$) numbers are not stored exactly.)

The rationals are an exact type. All rational numbers (subject to available resources, of course) are stored exactly. The results of the arithmetical operations on rationals are rationals represented exactly. Beware that the standard `intvalp` and `ratvalp` may return 1 even if the value is of double type. However, these functions may be combined with `exactp`.

exactp *x* returns whether *x* has an exact value.

inexactp *x* returns whether *x* has an inexact value.

rationalp *x* returns whether *x* is of rational type.

ratvalp *x* returns whether *x* has a rational value.

Example 4 Rational value tests:

```
> let l = [9, 9%1, 9%2, 4.5, sqrt 2, 1+i, inf, nan];
> map exactp l;
[1,1,1,0,0,1,0,0]
> map inexactp l;
[0,0,0,1,1,0,1,1]
> map rationalp l;
[0,1,1,0,0,0,0,0]
> map ratvalp l;
[1,1,1,1,1,0,0,0]
> map (\x -> (exactp x && ratvalp x)) l; // "has exact rational value"
[1,1,1,0,0,0,0,0]
> map intvalp l; // for comparison
[1,1,0,0,0,0,0,0]
> map (\x -> (exactp x && intvalp x)) l; // "has exact integer value"
[1,1,0,0,0,0,0,0]
```

See [Rational Complex Numbers](#) for details about rational complex numbers, and [Rational Complex Type and Value Tests](#) for details of their type and value tests.

5 Arithmetic

5.1 Operators

The standard arithmetic operators (+), (−) and (*) are overloaded to have at least one rational operand. If both operands are rational then the result is rational. If one operand is integer, then the result is rational. If one operand is double, then the result is double. The operators (/) and (%) are overloaded for division on at least one rational operand. The value returned by (/) is always inexact (in the sense of [Type and Value Tests](#)). The value returned by (%) is exact (if it exists). The standard function pow is overloaded to have a rational left operand. If pow is passed integer operands where the right operand is negative, then a rational is returned. The right operand should be an integer; negative values are permitted (because $q^{-z} = 1/q^z$). It is not overloaded to also have a rational right operand because such values are not generally rational (e.g. $q^{1/n} = \sqrt[n]{q}$). The standard arithmetic operator (^) is also overloaded, but produces a double value (as always).

Example 5 Arithmetic:

```
> 5%7 + 2%3;
29L%21L
> str_mixed ans;
"1L+8L/21L"
> 1 + 2%3;
5L%3L
> ans + 1.0;
2.666666666666667
> 3%8 - 1%3;
1L%24L
> (11%10) ^ 3;
1.331
> pow (11%10) 3;
1331L%1000L
> pow 3 5;
243L
> pow 3 (-5);
1L%243L
```

(See the function [str_mixed](#).)

Beware that (/) on integers will not produce a rational result.

Example 6 Division:

```
> 44/14;
3.14285714285714
> 44%14;
22L%7L
> str_mixed ans;
"3L+1L/7L"
```

(See the function [str_mixed](#).)

5.2 More on Division

There is a rational-aware divide operator on the numeric types:

n1 % n2 returns the quotient ($\in \mathbf{Q}$) of n1 and n2. If n1 and n2 are rational or integer then the result is rational. This operator has the precedence of division (/).

Example 7 Using % like a constructor:

```
> 44 % 14;
22L%7L
> 2 + 3%8; // "2 3/8"
19L%8L
> str_mixed ans;
"2L+3L/8L"
```

(See the function [str_mixed](#).)

reciprocal n returns the reciprocal of n: $1/n$.

Example 8 Reciprocal:

```
> reciprocal (22%7);
7L%22L
```

The following division functions are parameterised by a rounding mode [roundfun](#). The available rounding modes are described in [Rounding to Integer](#).

divide roundfun n d for rationals n and d returns a pair (q, r) of ‘quotient’ and ‘remainder’ where q is an integer and r is a rational such that $|r| < |d|$ (or better) and $n = q * d + r$. Further conditions may hold, depending on the chosen rounding mode [roundfun](#) (see [Rounding to Integer](#)). If [roundfun](#) = [floor](#) then $0 \leq r < d$. If [roundfun](#) = [ceil](#) then $-d < r \leq 0$. If [roundfun](#) = [trunc](#) then $|r| < |d|$ and $\text{sgn } r \in \{0, \text{sgn } d\}$. If [roundfun](#) = [round](#), [roundfun](#) = [round_zero_bias](#) or [roundfun](#) = [round_unbiased](#) then $|r| \leq d/2$.

quotient roundfun nN d returns just the quotient as produced by [divide roundfun n d](#).

modulus roundfun n d returns just the remainder as produced by [divide roundfun n d](#).

q1 div q2 (overload of the built-in div) q1 and q2 may be rational or integer. Returns an integer.

q1 mod q2 (overload of the built-in mod) q1 and q2 may be rational or integer. Returns a rational. If $q = q1 \text{ div } q2$ and $r = q1 \text{ mod } q2$ then $q1 = q * q2 + r$, $q \in \mathbf{Z}$, $|r| < |q2|$ and $\text{sgn } r \in \{0, \text{sgn } q2\}$.

5.3 Relations — Equality and Inequality Tests

The standard arithmetic operators (==), (~=), (<), (<=), (>), (>=) are overloaded to have at least one rational operand. The other operand may be rational, integer or double.

Example 9 Inequality:

```
> 3%8 < 1%3;  
0
```

5.4 Comparison Function

cmp *n1 n2* is the ‘comparison’ (or ‘compare’) function, and returns $\text{sgn}(n1 - n2)$; that is, it returns -1 if $n1 < n2$, 0 if $n1 = n2$, and $+1$ if $n1 > n2$.

Example 10 Compare:

```
> cmp (3%8) (1%3);  
1
```

6 Mathematical Functions

Most mathematical functions, including the elementary functions (\sin , \sin^{-1} , \sinh , \sinh^{-1} , \cos , \dots , \exp , \ln , \dots), are not closed on the set of rational numbers. That is, most mathematical functions do not yield a rational number in general when applied to a rational number. Therefore the elementary functions are not defined for rationals. To apply these functions, first apply a cast to double, or compose the function with a cast.

6.1 Absolute Value and Sign

The standard `abs` and `sgn` functions are overloaded for rationals.

abs *q* returns absolute value, or magnitude, $|q|$ of *q*; $\text{abs } q = |q| = q \times \text{sgn } q$ (see below).

sgn *q* returns the sign of *q* as an integer; returns -1 if $q < 0$, 0 if $q = 0$, $+1$ if $q > 0$.

Together, these functions satisfy the property $\forall q \bullet (\text{sgn } q) * (\text{abs } q) = q$ (i.e. $\forall q \bullet (\text{sgn } q) * |q| = q$). Thus these provide a pair of ‘decomposition’ operators; there are others (see [Decomposition](#)).

6.2 Greatest Common Divisor (GCD) and Least Common Multiple (LCM)

The standard functions `gcd` and `lcm` are overloaded for rationals, and mixtures of integer and rational.

gcd *n1 n2* The GCD is also known as the Highest Common Factor (HCF). The GCD of rationals *q1* and *q2* is the largest (therefore positive) rational *f* such that *f* divides into both *q1* and *q2* exactly, i.e. an integral number of times. This is not defined for *n1* and *n2* both zero. For integral *q1* and *q2*, this definition coincides with the usual definition of GCD for integers.

Example 11 With two rationals:

```
> let a = 7%12;
> let b = 21%32;
> let f = gcd a b;
> f;
7L%96L
> a % f;
8L%1L
> b % f;
9L%1L
```

Example 12 With a rational and an integer:

```
> let f = gcd (6%5) 4;
> f;
2L%5L
> (6%5) % f;
3L%1L
> 4 % f;
10L%1L
```

Example 13 With integral rationals and with integers:

```
> gcd (rational 18) (rational 24);
6L%1L
> gcd 18 24;
6
```

Example 14 The behaviour with negative numbers:

```
> gcd (rational (-18)) (rational 24);
6L%1L
> gcd (rational 18) (rational (-24));
6L%1L
> gcd (rational (-18)) (rational (-24));
6L%1L
```

lcm n1 n2 The LCM of rationals q_1 and q_2 is the smallest positive rational m such that both q_1 and q_2 divide m exactly. This is not defined for n_1 and n_2 both zero. For integral q_1 and q_2 , this definition coincides with the usual definition of LCM for integers.

Example 15 With two rationals:

```
> let a = 7%12;
> let bB = 21%32;
> let m = lcm a b;
> m;
21L%4L
> m % a;
9L%1L
> m % b;
8L%1L
```

Example 16 With a rational and an integer:

```
> let m = lcm (6%5) 4;
> m;
12L%1L
> m % (6%5);
10L%1L
```

Example 17 The behaviour with negative numbers:

```
> lcm (rational (-18)) (rational 24);
72L%1L
> lcm (rational 18) (rational (-24));
72L%1L
> lcm (rational (-18)) (rational (-24));
72L%1L
```

Together, the GCD and LCM have the following property when applied to two numbers: $(\text{gcd } q1 \ q2) * (\text{lcm } q1 \ q2) = |q1 * q2|$.

6.3 Extrema (Minima and Maxima)

The standard `min` and `max` functions work with rational values.

Example 18 Maximum:

```
> max (3%8) (1%3);
3L%8L
```

7 Special Rational Functions

7.1 Complexity

The ‘complexity’ (or ‘complicatedness’) of a rational is a measure of the greatness of its simplest (positive) denominator.

The complexity of a number is not itself made available, but various functions and operators are provided to allow complexities to be compared. Generally, it does not make sense to operate directly on complexity values.

The complexity functions in this section may be applied to integers (the least complex), rationals, or reals (doubles; the most complex).

Functions concerning ‘complexity’ are named with ‘`cplx`’, whereas functions concerning ‘complex numbers’ (see [Rational Complex Numbers](#)) are named with ‘`comp`’.

Complexity Relations

n1 eq_cplx n2 “[is] equally complex [to]” — returns 1 if `n1` and `n2` are equally complex; returns 0 otherwise. Equal complexity is not the same as equality; `n1` and `n2` are equally

complex if their simplest denominators are equal. Equal complexity forms an equivalence relation on rationals.

Example 19 Complexity equality test:

```
> (1%3) eq_cplx (100%3);
1
> (1%4) eq_cplx (1%5);
0
> (3%3) eq_cplx (1%3); // LHS is not in simplest form
0
```

n1 not_eq_cplx n2 “not equally complex” — returns 0 if n1 and n2 are equally complex; returns 1 otherwise.

n1 less_cplx n2 “[is] less complex [than]” (or “simpler”) — returns 1 if n1 is strictly less complex than n2; returns 0 otherwise. This forms a partial strict ordering on rationals.

Example 20 Complexity inequality test:

```
> (1%3) less_cplx (100%3);
0
> (1%4) less_cplx (1%5);
1
> (3%3) less_cplx (1%3); // LHS is not in simplest form
1
```

n1 less_eq_cplx n2 “less or equally complex” (or “not more complex”) — returns 1 if n1 is less complex than or equally complex to n2; returns 0 otherwise. This forms a partial non-strict ordering on rationals.

n1 more_cplx n2 “[is] more complex [than]” — returns 1 if n1 is strictly more complex than n2; returns 0 otherwise. This forms a partial strict ordering on rationals.

n1 more_eq_cplx n2 “more or equally complex” (or “not less complex”) — returns 1 if n1 is more complex than or equally complex to n2; returns 0 otherwise. This forms a partial non-strict ordering on rationals.

Complexity Comparison Function

cmp_complexity n1 n2 is the ‘complexity comparison’ function, and returns the sign of the difference in complexity; that is, it returns -1 if n1 is less complex than n2, 0 if n1 and n2 are equally complex (but not necessarily equal), and $+1$ if n1 is more complex than n2.

Example 21 Complexity comparison:

```
> cmp_complexity (1%3) (100%3);
0
> cmp_complexity (1%4) (1%5);
-1
```

```
> cmp_complexity (3%3) (1%3); // LHS is not in simplest form
-1
```

Complexity Extrema

least_cplx *n1 n2* returns the least complex of *n1* and *n2*; if they're equally complex, *n1* is returned.

Example 22 Complexity selection:

```
> least_cplx (100%3) (1%3);
100L%3L
> least_cplx (1%5) (1%4);
1L%4L
> least_cplx (1%3) (3%3); // second argument not in simplest form
1L%1L
```

most_cplx *n1 n2* returns the most complex of *n1* and *n2*; if they're equally complex, *n1* is returned.

Other Complexity Functions

complexity_rel *n1 op n2* returns “complexity-of *n1*” compared by operator *op* to the “complexity-of *n2*”. This is equivalent to prefix complexity *rel op n1 n2* (below), but is the more readable form.

Example 23 Complexity relations:

```
> complexity_rel (1%3) (==) (100%3);
1
> complexity_rel (1%4) (<=) (1%5);
1
> complexity_rel (1%4) (>) (1%5);
0
```

prefix_complexity_rel *op n1 n2* returns the same as **complexity_rel** *n1 op n2*, but this form is more convenient for currying.

7.2 Mediants and Farey Sequences

mediant *q1 q2* returns the canonical mediant of the rationals *q1* and *q2*, a form of (nonarithmetic) average on rationals. The mediant of the representations $n_1/d_1 = q_1$ and $n_2/d_2 = q_2$, where d_1 and d_2 must be positive, is defined as $(n_1 + n_2)/(d_1 + d_2)$. A mediant of the rationals *q1* and *q2* is a mediant of some representation of each of *q1* and *q2*. That is, the mediant is dependent upon the representations and therefore is not well-defined as a function on pairs of rationals. The canonical mediant always assumes the simplest representation, and therefore is well-defined as a function on pairs of rationals.

By the phrase “the mediant” (as opposed to just “a mediant”) we always mean “the canonical mediant”.

If $q_1 < q_2$, then any mediant q is always such that $q_1 < q < q_2$.

The (canonical) mediant has some special properties. If q_1 and q_2 are integers, then the mediant is the arithmetic mean. If q_1 and q_2 are unit fractions (reciprocals of integers), then the mediant is the harmonic mean. The mediant of q and $1/q$ is ± 1 , (which happens to be a geometric mean with the correct sign, although this is a somewhat uninteresting and degenerate case).

Example 24 Mediants:

```
> mediant (1%4) (3%10);
2L%7L
> mediant 3 7; // both integers
5L%1L
> mediant 3 8; // both integers again
11L%2L
> mediant (1%3) (1%7); // both unit fractions
1L%5L
> mediant (1%3) (1%8); // both unit fractions again
2L%11L
> mediant (-10) (-1%10);
(-1L)%1L
```

farey -k for an integer k , farey returns the ordered list containing the order- k Farey sequence, which is the ordered list of all rational numbers between 0 and 1 inclusive with (simplest) denominator at most k .

Example 25 A Farey sequence:

```
> map str_mixed (farey 6);
["0L", "1L/6L", "1L/5L", "1L/4L", "1L/3L", "2L/5L", "1L/2L", "3L/5L", "2L/3L",
"3L/4L", "4L/5L", "5L/6L", "1L"]
```

(See the function [str_mixed](#).)

Farey sequences and mediants are closely related. Three rationals $q_1 < q_2 < q_3$ are consecutive members of a Farey sequence if and only if q_2 is the mediant of q_1 and q_3 . If rationals $q_1 = n_1/d_1 < q_2 = n_2/d_2$ are consecutive members of a Farey sequence, then $n_2d_1 - n_1d_2 = 1$.

7.3 Rational Type Simplification

rat_simplify q returns q with rationals simplified to integers, if possible.

Example 26 Rational type simplification:

```
> let l = [9, 9%1, 9%2, 4.5, 9%1+i, 9%2+i]; l;
[9, 9L%1L, 9L%2L, 4.5, 9L%1L+:1, 9L%2L+:1]
```

```
> map rat_simplify l;  
[9,9,9L%2L,4.5,9+:1,9L%2L+:1]
```

See [Rational Complex Numbers](#) for details about rational complex numbers, and [Rational Complex Type Simplification](#) for details of their type simplification.

8 Q -> Z — Rounding

8.1 Rounding to Integer

Some of these are new functions, and some are overloads of standard functions. The behaviour of the overloads is consistent with that of the standard functions.

floor q (overload of standard function) returns q rounded downwards, i.e. towards -1 , to an integer, usually denoted bQc .

ceil q (overload of standard function) returns q rounded upwards, i.e. towards $+1$, to an integer, usually denoted dQe .

trunc q (overload of standard function) returns q truncated, i.e. rounded towards 0, to an integer.

round q (overload of standard function) returns q ‘rounded off’, i.e. rounded to the nearest integer, with ‘half-integers’ (values that are an integer plus a half) rounded away from zero.

round_zero_bias q (new function) returns q ‘rounded off’, i.e. rounded to the nearest integer, but with ‘half-integers’ rounded towards zero.

round_unbiased q (new function) returns q rounded to the nearest integer, with ‘half-integers’ rounded to the nearest even integer.

Example 27 Illustration of the different rounding modes:

```
> let l = iterwhile (<= 3) (+(1%2)) (- rational 3);  
> map double l; // (just to show the values in a familiar format)  
[-3.0,-2.5,-2.0,-1.5,-1.0,-0.5,0.0,0.5,1.0,1.5,2.0,2.5,3.0]  
> map floor l;  
[-3L,-3L,-2L,-2L,-1L,-1L,0L,0L,1L,1L,2L,2L,3L]  
> map ceil l;  
[-3L,-2L,-2L,-1L,-1L,0L,0L,1L,1L,2L,2L,3L,3L]  
> map trunc l;  
[-3L,-2L,-2L,-1L,-1L,0L,0L,0L,1L,1L,2L,2L,3L]  
> map round l;  
[-3L,-3L,-2L,-2L,-1L,-1L,0L,1L,1L,2L,2L,3L,3L]  
> map round_zero_bias l;  
[-3L,-2L,-2L,-1L,-1L,0L,0L,0L,1L,1L,2L,2L,3L]  
> map round_unbiased l;  
[-3L,-2L,-2L,-2L,-1L,0L,0L,0L,1L,2L,2L,2L,3L]
```

(See the function [double](#).)

8.2 Integer and Fraction Parts

integer_and_fraction roundfun q returns a pair (z, f) where z is the ‘integer part’ as an integer, f is the ‘fraction part’ as a rational, where the rounding operations are performed using rounding mode **roundfun** (see [Rounding to Integer](#)).

Example 28 Integer and fraction parts with the different rounding modes:

```
> let nc = -22%7;
> integer_and_fraction floor nc;
-4L,6L%7L
> integer_and_fraction trunc nc;
-3L,(-1L)%7L
> integer_and_fraction round nc;
-3L,(-1L)%7L
```

It is always the case that z and f have the property that $q = z + f$. However, the remaining properties depend upon the choice of **roundfun**. Thus this provides a ‘decomposition’ operator; there are others (see [Decomposition](#)). If **roundfun** = **floor** then $0 \leq f < 1$. If **Round** = **ceil** then $-1 < f \leq 0$. If **roundfun** = **trunc** then $|f| < 1$ and $\text{sgn } f \in \{0, \text{sgn } q\}$. If **roundfun** = **round**, **roundfun** = **round_zero_bias** or **roundfun** = **round_unbiased** then $|f| \leq 1/2$.

fraction roundfun q returns just the ‘fraction part’ as a rational, where the rounding operations are performed using **roundfun**. The corresponding function ‘integer’ is not provided, as integer **roundfun q** would be just **roundfun q**. The integer and fraction function (probably with **trunc** or **floor** rounding mode) may be used in conjunction with **num_den** (see [‘Deconstructors’](#)) to decompose a rational into integer, numerator and denominator parts.

int q overloads the built-in int and returns the ‘integer part’ of q consistent with the built-in.

frac q overloads the built-in frac and returns the ‘fraction part’ of q consistent with the built-in.

Example 29 Standard integer and fraction parts:

```
> let nc = -22%7;
> int nc;
-3
> frac nc;
(-1L)%7L
```

9 Rounding to Multiples

round_to_multiple roundfun multOf q returns q rounded to an integer multiple of a non-zero value multOf, using **roundfun** as the rounding mode (see [Rounding to Integer](#)). Note that it is the multiple that is rounded in the prescribed way, and not the final result, which may make a difference in the case that multOf is negative. If that is

not the desired behaviour, pass this function the absolute value of `multOf` rather than `multOf`. Similar comments apply to the following functions.

`floor_multiple multOf q` returns `q` rounded to a downwards integer multiple of `multOf`.

`ceil_multiple multOf q` returns `q` rounded to an upwards integer multiple of `multOf`.

`trunc_multiple multOf q` returns `q` rounded towards zero to an integer multiple of `multOf`.

`round_multiple multOf q` returns `q` rounded towards the nearest integer multiple of `multOf`, with half-integer multiples rounded away from 0.

`round_multiple_zero_bias multOf q` returns `q` rounded towards the nearest integer multiple of `multOf`, with half-integer multiples rounded towards 0.

`round_multiple_unbiased multOf q` returns `q` rounded towards the nearest integer multiple of `multOf`, with half-integer multiples rounded to an even multiple.

Example 30 Round to multiple:

```
> let l = [34.9, 35, 35%1, 35.0, 35.1];
> map double l; // (just to show the values in a familiar format)
[34.9,35.0,35.0,35.0,35.1]
> map (floor_multiple 10) l;
[30.0,30L,30L,30.0,30.0]
> map (ceil_multiple 10) l;
[40.0,40L,40L,40.0,40.0]
> map (trunc_multiple 10) l;
[30.0,30L,30L,30.0,30.0]
> map (round_multiple 10) l;
[30.0,40L,40L,40.0,40.0]
> map (round_multiple_zero_bias 10) l;
[30.0,30L,30L,30.0,40.0]
> map (round_multiple_unbiased 10) l;
[30.0,40L,40L,40.0,40.0]
```

(See the function [double](#).)

The round multiple functions may be used to find a fixed denominator approximation of a number. (The simplest denominator may actually be a proper factor of the chosen value.) To approximate for a bounded (rather than particular) denominator, use rational approx max den instead (see [Best Approximation with Bounded Denominator](#)).

Example 31 Finding the nearest $q = n/d$ value to $1/e \approx 0.368$ where $d = 1000$ (actually, where $d \mid 1000$):

```
> let co_E = exp (-1);
co_E;
0.367879441171442
> round_multiple (1%1000) (rational co_E);
46L%125L
> 1000 * ans;
368L%1L
```

Example 32 Finding the nearest $q = n/d$ value to $1/\phi \approx 0.618$ where $d = 3^5 = 243$ (actually, where $d \mid 243$):

```
> let co_Phi = (sqrt 5 - 1) / 2;  
> round_multiple (1%243) (rational co_Phi);  
50L%81L
```

Other methods for obtaining a rational approximation of a number are described in [R -> Q — Approximation](#).

10 Q -> R — Conversion / Casting

double q (overload of built-in) returns a double having a value as close as possible to q. (Overflow, underflow and loss of accuracy are potential problems. rationals that are too absolutely large or too absolutely small may overflow or underflow; some rationals can not be represented exactly as a double.)

11 R -> Q — Approximation

This section describes functions that approximate a number (usually a double) by a rational. See [Rounding to Multiples](#) for approximation of a number by a rational with a fixed denominator. See [Numeral String -> Q — Approximation](#) for approximation by a rational of a string representation of a real number.

11.1 Intervals

Some of the approximation functions return an **interval**. The file `rat_interval.pure` is a basic implementation of interval arithmetic, and is not included in the default prelude. It is not intended to provide a complete implementation of interval arithmetic. The notions of ‘open’ and ‘closed’ intervals are not distinguished. Infinite and half-infinite intervals are not specifically provided. Some operations and functions may be missing. The most likely functions to be used are simply the ‘deconstructors’; see [Interval Constructors and ‘Deconstructors’](#).

Interval Constructors and ‘Deconstructors’

Intervals are constructed with the function `interval`.

interval (n1, n2) given a pair of numbers ($z1 \leq z2$), this returns the interval `z1..z2`. This is the inverse of `lo_up`.

Example 33 Constructing an interval:

```
> let v = interval (3, 8);  
> v;  
interval::Ivl 3 8
```

lower *v* returns the infimum (roughly, minimum) of *v*.

upper *v* returns the supremum (roughly, maximum) of *v*.

lo_up *v* returns a pair (*l*, *u*) containing the lower *l* and upper *u* extrema of the interval *v*.
This is the inverse of interval as defined on number pairs.

Example 34 Deconstructing an interval:

```
> lower v;  
3  
> upper v;  
8  
> lo_up v;  
3,8
```

Interval Type Tests

exactp *v* returns whether an interval *v* has exact extrema.

inexactp *v* returns whether an interval *v* has an inexact extremum.

intervalp *x* returns whether *x* is of type interval.

interval_valp *x* returns whether *x* has an interval value.

ratinterval_valp *x* returns whether *x* has an interval value with rational extrema.

intinterval_valp *x* returns whether *x* has an interval value with integral extrema.

Example 35 Interval value tests:

```
> let l = [interval(0,1), interval(0,1%1), interval(0,3%2), interval(0,1.5)];  
> map exactp l;  
[1,1,1,0]  
> map inexactp l;  
[0,0,0,1]  
> map intervalp l;  
[1,1,1,1]  
> map interval_valp l;  
[1,1,1,1]  
> map ratinterval_valp l;  
[1,1,1,1]  
> map intinterval_valp l;  
[1,1,0,0]
```

Interval Arithmetic Operators and Relations

The standard arithmetic operators (+), (−), (*), (/) and (%) are overloaded for intervals. The divide operators (/) and (%) do not produce a result if the right operand is an interval containing 0. **Example 36** Some intervals:

```
> let a = interval (11, 19);
> let b = interval (16, 24);
> let c = interval (21, 29);
> let d = interval (23, 27);
```

Example 37 Interval arithmetic:

```
> let p = interval (0, 1);
> let s = interval (-1, 1);
> a + b;
interval::Ivl 27 43
> a - b;
interval::Ivl (-13) 3
> a * b;
interval::Ivl 176 456
> p * 2;
interval::Ivl 0 2
> (-2) * p;
interval::Ivl (-2) 0
> -c;
interval::Ivl (-29) (-21)
> s * a;
interval::Ivl (-19) 19
> a % 2;
interval::Ivl (11L%2L) (19L%2L)
> a / 2;
interval::Ivl 5.5 9.5
> reciprocal a;
interval::Ivl (1L%19L) (1L%11L)
> 2 % a;
interval::Ivl (2L%19L) (2L%11L)
> a % b;
interval::Ivl (11L%24L) (19L%16L)
> a % a; // notice that the intervals are mutually independent here
interval::Ivl (11L%19L) (19L%11L)
```

There are also some relations defined for intervals. The standard relations (`==`) and (`~=`) are overloaded.

However, rather than overloading (`<`), (`<=`), (`>`), (`>=`), which could be used for either ordering or containment with some ambiguity, the module defines (`before`), (`within`), and so on. ‘Strictness’ refers to the properties at the end-points.

v1 before v2 returns whether v1 is entirely before v2.

v1 strictly_before v2 returns whether v1 is strictly entirely before v2.

v1 after v2 returns whether v1 is entirely after v2.

v1 strictly_after v2 returns whether v1 is strictly entirely after v2.

v1 within v2 returns whether v1 is entirely within v2; i.e. whether v1 is subinterval of v2.

v1 strictly_within v2 returns whether v1 is strictly entirely within v2; i.e. whether v1 is proper subinterval of v2.

v1 without v2 returns whether v1 entirely contains v2; i.e. whether v1 is superinterval of v2. 'Without' is used in the sense of outside or around.

v1 strictly_without v2 returns whether v1 strictly entirely contains v2; i.e. whether v1 is proper superinterval of v2.

v1 disjoint v2 returns whether v1 and v2 are entirely disjoint.

v strictly_disjoint v2 returns whether v1 and v2 are entirely strictly disjoint.

Example 38 Interval relations:

```
> a == b;
0
> a == a;
1
> a before b;
0
> a before c;
1
> c before a;
0
> a disjoint b;
0
> c disjoint a;
1
> a within b;
0
> a within c;
0
> d within c;
1
> c within d;
0
> a strictly_within a;
0
> a within a;
1
```

(The symbols a through d were defined in [Example 36](#).)

These may also be used with a simple (real) value, and in particular to test membership.

Example 39 Membership:

```
> 10 within a;
0
> 11 within a;
1
> 11.0 within a;
1
> 12 within a;
```

```
1
> 12.0 within a;
1
> 10 strictly_within a;
0
> 11 strictly_within a;
0
> (11%1) strictly_within a;
0
> 12 strictly_within a;
1
> (12%1) strictly_within a;
1
```

(The symbol `a` was defined in [Example 36](#).)

Interval Maths

Some standard functions are overloaded for intervals; some new functions are provided.

abs v returns the interval representing the range of (x) as x varies over v .

Example 40 Absolute interval:

```
> abs (interval (1, 5));
interval::Ivl 1 5
> abs (interval (-1, 5));
interval::Ivl 0 5
> abs (interval (-5, -1));
interval::Ivl 1 5
```

sgn v returns the interval representing the range of $\text{sgn}(x)$ as x varies over v .

#v returns the length of an interval.

Example 41 Absolute interval:

```
> #d;
4
```

(The symbol `d` was defined in [Example 36](#).)

11.2 Least Complex Approximation within Epsilon

rational_approx_epsilon ϵ r Find the least complex (see [Complexity Extrema](#)) rational approximation to r (usually a double) that is ϵ -close. That is find the q with the smallest possible denominator such that $|q - r| \leq \epsilon$. ($\epsilon > 0$.)

Example 42 Rational approximation to $\pi \approx 3.142 \approx 22/7$:

```
> rational_approx_epsilon .01 pi;
22L%7L
> abs (ans - pi);
0.00126448926734968
```

Example 43 The golden ratio $\phi = (1 + \sqrt{5}) / 2 \approx 1.618$:

```
> let phi = (1 + sqrt 5) / 2;
> rational_approx_epsilon .001 phi;
55L%34L
> abs (ans - phi);
0.000386929926365465
```

rational_approx_epsilon ϵ **r** Produce a list of ever better rational approximations to **r** (usually a double) that is eventually ϵ -close. ($\epsilon > 0$.)

Example 44 Rational approximations to π :

```
> rational_approx_epsilon .0001 pi;
[3L%1L,25L%8L,47L%15L,69L%22L,91L%29L,113L%36L,135L%43L,157L%50L,179L%57L,
201L%64L,223L%71L,245L%78L,267L%85L,289L%92L,311L%99L,333L%106L]
```

Example 45 Rational approximations to the golden ratio ϕ ; these approximations are always reverse consecutive Fibonacci numbers (from f1: 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, 89, ...):

```
> rational_approx_epsilon .0001 phi;
[1L%1L,3L%2L,8L%5L,21L%13L,55L%34L,144L%89L]
```

(The symbol `phi` was defined in [Example 43](#).)

rational_interval_epsilon ϵ **r** Find the least complex (see [Complexity Extrema](#)) rational interval containing **r** (usually a double) that is ϵ -small. That is find the least complex (see [Complexity Extrema](#)) $q1 \leq q2$ such that $r \in [q1, q2]$ and $q2 - q1 \leq \epsilon$. ($\epsilon > 0$.)

Example 46 Rational interval surrounding π :

```
> let i_Pi = rational_interval_epsilon .01 pi;
> i_Pi;
interval::Ivl (47L%15L) (22L%7L)
> double (lower i_Pi); pi; double (upper i_Pi);
3.133333333333333
3.14159265358979
3.14285714285714
```

(The functions `lower` and `upper` are described in [Interval Constructors and ‘Deconstructors’](#).)

Example 47 Rational interval surrounding the golden ratio ϕ :

```
> rational_interval_epsilon .001 phi;
interval::Ivl (55L%34L) (89L%55L)
> #ans;
1L%1870L
```

(The symbol `phi` was defined in [Example 43](#). The function `# (interval)` is described in [Interval Maths](#).)

11.3 Best Approximation with Bounded Denominator

rational_approx_max_den *maxDen* *r* Find the closest rational approximation to *r* (usually a double) that has a denominator no greater than *maxDen*. (*maxDen* > 0).

Example 48 Rational approximation to π :

```
> rational_approx_max_den 10 pi;  
22L%7L
```

Example 49 Rational approximation to the golden ratio ϕ :

```
> rational_approx_max_den 1000 phi;  
1597L%987L
```

(The symbol *phi* was defined in [Example 43](#).)

rational_approxs_max_den *maxDen* *r* Produce a list of ever better rational approximations to *r* (usually a double) while the denominator is bounded by *maxDen* (*maxDen* > 0).

Example 50 Rational approximations to π :

```
> rational_approxs_max_den 100 pi;  
[3L%1L,25L%8L,47L%15L,69L%22L,91L%29L,113L%36L,135L%43L,157L%50L,179L%57L,  
201L%64L,223L%71L,245L%78L,267L%85L,289L%92L,311L%99L]
```

Example 51 Rational approximations to the golden ratio ϕ :

```
> rational_approxs_max_den 100 phi;  
[1L%1L,3L%2L,8L%5L,21L%13L,55L%34L,144L%89L]
```

(The symbol *phi* was defined in [Example 43](#).)

rational_interval_max_den *maxDen* *r* Find the smallest rational interval containing *r* (usually a double) that has endpoints with denominators no greater than *maxDen* (*maxDen* > 0).

Example 52 Rational interval surrounding π :

```
> let i_Pi = rational_interval_max_den 100 pi ; i_Pi;  
interval::Ivl (311L%99L) (22L%7L)  
> double (lower i_Pi); pi; double (upper i_Pi);  
3.1414141414141414  
3.14159265358979  
3.14285714285714
```

Example 53 Rational interval surrounding the golden ratio ϕ :

```
> rational_interval_max_den 1000 phi;  
interval::Ivl (987L%610L) (1597L%987L)
```

(The symbol *phi* was defined in [Example 43](#).)

To approximate for a particular (rather than bounded) denominator, use `round` to multiple instead (see [Rounding to Multiples](#)).

12 Decomposition

There is more than one way to ‘decompose’ a rational number into its ‘components’. It might be split into an integer and a fraction part — see [Integer and Fraction Parts](#); or sign and absolute value — see [Absolute Value and Sign](#); or numerator and denominator — see [Deconstructors](#).

13 Continued Fractions

13.1 Introduction

In “pure-rational”, a continued fraction $a_0 + (1 / (a_1 + (1 / (a_2 + \dots + 1 / a_n))))$ where $\forall i > 0 \bullet a_i \neq 0$, is represented by $[a_0, a_1, a_2, \dots, a_n]$.

A ‘simple’ continued fraction is one in which $\forall i \bullet a_i \in \mathbf{Z}$ and $\forall i > 0 \bullet a_i > 0$.

Simple continued fractions for rationals are not quite unique since $[a_0, a_1, \dots, a_n, 1] = [a_0, a_1, \dots, a_{n+1}]$. We will refer to these as the ‘non-standard’ and ‘standard’ forms, respectively. The following functions return the standard form.

13.2 Generating Continued Fractions

Exact

continued_fraction q Find ‘the’ (exact) continued fraction of a rational (including, trivially, integer) value q.

Example 54 The rational 1234/1001:

```
> continued_fraction (1234%1001);  
[1L,4L,3L,2L,1L,1L,1L,8L]  
> evaluate_continued_fraction ans;  
1234L%1001L
```

Inexact

continued_fraction_max_terms n r Find up to n initial terms of continued fraction of the value r with the ‘remainder’, if any, in the final element. (If continued_fraction_max_terms n r returns a list of length n or less, then the result is exact.)

Example 55 First 5 terms of the continued fraction for the golden ratio ϕ :

```
> continued_fraction_max_terms 5 phi;  
[1.0,1.0,1.0,1.0,1.0,1.61803398874989]  
> evaluate_continued_fraction ans;  
1.61803398874989
```

(The symbol ϕ was defined in [Example 43](#).)

continued_fraction_epsilon ϵ **r** Find enough of the initial terms of a continued fraction to within ϵ of the value r with the ‘remainder’, if any, in the final element.

Example 56 First few terms of the value $\sqrt{2}$:

```
> continued_fraction_epsilon .001 (sqrt 2);
[1.0,2.0,2.0,2.0,2.0,2.41421356237241]
> map double (convergents ans);
[1.0,1.5,1.4,1.41666666666667,1.41379310344828,1.41421356237309]
```

13.3 Evaluating Continued Fractions

evaluate_continued_fraction **aa** Fold a continued fraction **aa** into the value it represents. This function is not limited to simple continued fractions. (Exact simple continued fractions are folded into a rational.)

Example 57 The continued fraction $[1, 2, 3, 4]$ and the non-standard form $[4, 3, 2, 1]$:

```
> evaluate_continued_fraction [1,2,3,4];
43L%30L
> continued_fraction ans;
[1L,2L,3L,4L]
> evaluate_continued_fraction [4,3,2,1];
43L%10L
> continued_fraction ans;
[4L,3L,3L]
```

Convergents

convergents **aa** Calculate the convergents of the continued fraction **aa**. This function is not limited to simple continued fractions.

Example 58 Convergents of a continued fraction approximation of the value $\sqrt{2}$:

```
> continued_fraction_max_terms 5 (sqrt 2);
[1.0,2.0,2.0,2.0,2.0,2.41421356237241]
> convergents ans;
[1L%1L,3L%2L,7L%5L,17L%12L,41L%29L,261169948480459693L%184675041612674289L]
```

14 Rational Complex Numbers

Pure together with `rational.pure` provide various types of number, including integers (**Z**), doubles (**R**, roughly), complex numbers (**C**) and Gaussian integers (**Z[i]**), rationals (**Q**) and rational complex numbers (**Q[i]**).

Functions concerning ‘complex numbers’ are named with ‘comp’, whereas functions concerning ‘complexity’ (see [Complexity](#)) are named with ‘cplx’.

14.1 Rational Complex Constructors and ‘Deconstructors’

Complex numbers can have rational parts.

Example 59 Forming a rational complex:

```
> 1 +: 1 * (1%2);
1+:1L%2L
> ans * ans;
3L%4L+:1L%1L
```

And rational numbers can be given complex parts.

Example 60 Complex rationals and complicated rationals:

```
> (1 +: 2) % (3 +: 4);
1L%25L+:2L%25L
> ans * (3 +: 4);
1L%1L+:2L%1L
> ((4%1) * (0 +: 1)) % 2;
0L%1L+:2L%1L
> ((4%1) * (0 +: 1)) % (1%2);
0L%1L+:8L%1L
> ((4%1) * (0 +: 1)) % (1 + (1%2) * (0 +: 1));
8L%5L+:16L%5L
> ans * (1+(1%2) * (0 +: 1));
0L%1L+:4L%1L
> ((4%1) * (0 +: 1)) / (1 + (1%2) * (0 +: 1));
1.6+:3.2
```

The various parts of a complex rational may be deconstructed using combinations of `num` and `den` and the standard functions `re` and `im`.

Thus, taking real and imaginary parts first, a rational complex number may be considered to be $(x_n / x_d) + (y_n / y_d) * i$ with $x_n, x_d, y_n, y_d \in \mathbb{Z}$.

A rational complex number may also be decomposed into its ‘numerator’ and ‘denominator’, where these are both integral complex numbers, or ‘Gaussian integers’, and the denominator is a minimal choice in some sense.

One way to do this is so that the denominator is the minimum positive integer. The denominator is a complex number with zero imaginary part.

Thus, taking numerator and denominator parts first, a rational complex number may be considered to be $(n_x + n_y * i) / (d + 0 * i)$ with $n_x, n_y, d \in \mathbb{Z}$.

Another way to do this is so that the denominator is a Gaussian integer with minimal absolute value. Thus, taking numerator and denominator parts first, a rational complex number may be considered to be $(n_x + n_y * i) / (d_x + d_y * i)$ with $n_x, n_y, d_x, d_y \in \mathbb{Z}$.

The d_x, d_y are not unique, but can be chosen such that $d_x > 0$ and either $|d_y| < d_x$ or $d_y = d_x > 0$.

num_den_nat *c* given a complex rational or integer *c*, returns a pair (*n*, *d*) containing an integral complex (Gaussian integral) numerator *n*, and the smallest natural (i.e. positive integral real) complex denominator *d*, i.e. a complex number where $\Re(d) \in \mathbf{Z}$, $\Re(d) > 0$, $\Im(d) = 0$; i.e. the numerator and denominator of one ‘normalised’ form of *c*.

This is an inverse (up to equivalence) of rational as defined on integral complex pairs (see [Constructors](#)).

num_den_gauss *c* given a complex rational or integer *c*, returns a pair (*n*, *d*) containing an integral complex (Gaussian integral) numerator *n*, and an absolutely smallest integral complex denominator *d* chosen s.t. $\Re(d), \Im(d) \in \mathbf{Z}$, $\Re(d) > 0$, and either $|\Re(d)| < \Im(d)$ or $\Re(d) = \Im(d) > 0$; i.e. the numerator and denominator of another ‘normalised’ form of *c*.

This is an inverse (up to equivalence) of rational as defined on integral complex pairs (see [Constructors](#)).

num_den *c* synonymous with `num_den_gauss`.

This is an inverse (up to equivalence) of rational as defined on integer pairs (see [Constructors](#)).

num *c* given a complex rational or integer *c*, returns just the numerator of the normalised form of *c* given by `num_den c`.

den *c* given a complex rational or integer *c*, returns just the denominator of the normalised form of *c* given by `num_den c`.

Example 61 Rational complex number deconstruction:

```
> let cq = (1+2*i)%(3+3*i); cq;
1L%2L+:1L%6L
> (re cq, im cq);
1L%2L,1L%6L
> (num . re) cq;
1L
> (den . re) cq;
2L
> (num . im) cq;
1L
> (den . im) cq;
6L
> let (n_nat,d_nat) = num_den_nat cq;
> (n_nat, d_nat);
3+:1,6+:0
> n_nat % d_nat;
1L%2L+:1L%6L
> abs d_nat;
6.0
> let (n, d) = num_den_gauss cq; (n, d);
1L+:2L,3L+:3L
```

```
> let (n,d) = num_den cq; (n, d);
1L+:2L,3L+:3L
> n % d;
1L%2L+:1L%6L
> abs d;
4.24264068711928
> (re . num) cq;
1L
> (im . num) cq;
2L
> (re . den) cq; //always > 0
3L
> (im . den) cq; //always <= (re.den)
3L
```

14.2 Rational Complex Type and Value Tests

Beware that `intcompvalp` and `ratcompvalp` may return 1 even if the value is of complex type with double parts. However, these functions may be combined with `exactp`.

complexp *x* standard function; returns whether *x* is of complex type.

compvalp *x* standard function; returns whether *x* has a complex value ($\in \mathbf{C} = \mathbf{R}[i]$).

ratcompvalp *x* returns whether *x* has a rational complex value ($\in \mathbf{Q}[i]$).

intcompvalp *x* returns whether *x* has an integral complex value ($\in \mathbf{Z}[i]$), i.e. a Gaussian integer value.

Example 62 Rational complex number value tests:

```
> let l = [9, 9%1, 9%2, 4.5, sqrt 2, 1+:1, 1%2+:1, 0.5+:1, inf, nan];
> map exactp l;
[1,1,1,0,0,1,1,0,0,0]
> map inexactp l;
[0,0,0,1,1,0,0,1,1,1]
> map complexp l;
[0,0,0,0,0,1,1,1,0,0]
> map compvalp l;
[1,1,1,1,1,1,1,1,1,1]
> map (\x -> (exactp x and compvalp x)) l; // "has exact complex value"
[1,1,1,0,0,1,1,0,0,0]
> map ratcompvalp l;
[1,1,1,1,1,1,1,1,0,0]
> map (\x -> (exactp x and ratcompvalp x)) l;
[1,1,1,0,0,1,1,0,0,0]
> map intcompvalp l;
[1,1,0,0,0,1,0,0,0,0]
> map (\x -> (exactp x and intcompvalp x)) l;
[1,1,0,0,0,1,0,0,0,0]
> map ratvalp l;
[1,1,1,1,1,0,0,0,0,0]
```

```
> map (\x -> (exactp x and ratvalp x)) l;  
[1,1,1,0,0,0,0,0,0]  
> map intvalp l; // for comparison  
[1,1,0,0,0,0,0,0,0]  
> map (\x -> (exactp x and intvalp x)) l;  
[1,1,0,0,0,0,0,0,0]
```

See ‘Type and Value Tests’_ for some details of rational type and value tests.

14.3 Rational Complex Arithmetic Operators and Relations

The standard arithmetic operators (+), (−), (*), (/), (%), (), (==) and (~=) are overloaded to have at least one complex and/or rational operand, but (<), (<=), (>), (>=) are not, as complex numbers are unordered.

Example 63 Rational complex arithmetic:

```
> let w = 1%2 +: 3%4;  
> let z = 5%6 +: 7%8;  
> w + z;  
4L%3L+:13L%8L  
> w % z;  
618L%841L+:108L%841L  
> w / z;  
0.734839476813318+:0.128418549346017  
> w ^ 2;  
-0.3125+:0.75  
> w == z;  
0  
> w == w;  
1
```

14.4 Rational Complex Maths

The standard functions `re` and `im` work with rational complex numbers (see [Rational Complex Constructors and ‘Deconstructors’](#)). The standard functions `polar`, `abs` and `arg` work with rational complex numbers, but the results are inexact.

Example 64 Rational complex maths:

```
> polar (1%2+:1%2);  
0.707106781186548<:0.785398163397448  
> abs (4%2+:3%2);  
2.5  
> arg (-1%1);  
3.14159265358979
```

There are some additional useful functions for calculating with rational complex numbers and more general mathematical values.

norm_gauss *c* returns the Gaussian norm $||c||$ of any complex (or real) number *c*; this is the square of the absolute value, and is returned as an (exact) integer.

div_mod_gauss *n d* performs Gaussian integer division, returning (*q*, *r*) where *q* is a (not always unique) quotient, and *r* is a (not always unique) remainder. *q* and *r* are such that $n = q * d + r$ and $||r|| < ||d||$ (equivalently, $|r| < |d|$).

n_div_gauss *d* returns just a quotient from Gaussian integer division as produced by `div_mod_gauss n d`.

n_mod_gauss *d* returns just a remainder from Gaussian integer division as produced by `div_mod_gauss n d`.

gcd_gauss *c1 c2* returns a GCD *G* of the Gaussian integers *c1*, *c2*. This is chosen so that s.t. $\Re(G) > 0$, and either $|\Im(G)| < \Re(G)$ or $\Im(G) = \Re(G) > 0$;

euclid_gcd *zerofun modfun x y* returns a (non-unique) GCD calculated by performing the Euclidean algorithm on the values *x* and *y* (of any type) where *zerofun* is a predicate for equality to 0, and *modfun* is a binary modulus (remainder) function.

euclid_alg *zerofun divfun x y* returns (*g*, *a*, *b*) where the *g* is a (non-unique) GCD and *a*, *b* are (arbitrary, non-unique) values such that $a * x + b * y = g$ calculated by performing the generalised Euclidean algorithm on the values *x* and *y* (of any type) where *zerofun* is a predicate for equality to 0, and *div* is a binary quotient function.

Example 65 More rational complex and other maths:

```
> norm_gauss (1 +: 3);
10
> abs (1 +: 3);
3.16227766016838
> norm_gauss (-5);
25
> let (q, r) = div_mod_gauss 100 (12 +: 5);
> (q, r);
7L+:-3L,1L+:1L
> q * (12 +: 5) + r;
100L+:0L
> 100 div_gauss (12 +: 5);
7L+:-3L
> 100 mod_gauss (12 +: 5);
1L+:1L
> div_mod_gauss 23 5;
5L+:0L,-2L+:0L
> gcd_gauss (1 +: 2) (3 +: 4);
1L+:0L
> gcd_gauss 25 15;
5L+:0L
> euclid_gcd (==0) (mod_gauss) (1+: 2) (3 +: 4);
1L+:0L
> euclid_gcd (==0) (mod) 25 15;
5
> let (g, a, b) = euclid_alg (==0) (div_gauss) (1 +: 2) (3 +: 4); g;
```



```
1L+:0L
> (a, b);
-2L+:0L, 1L+:0L
> a * (1 +: 2) + b * (3 +: 4);
1L+:0L
> let (g, a, b) = euclid_alg (==0) (div) 25 15; g;
5
> (a, b);
-1, 2
> a * 25 + b * 15;
5
```

14.5 Rational Complex Type Simplification

comp_simplify c returns *q* with complex numbers simplified to reals, if possible.

ratcomp_simplify c returns *q* with rationals simplified to integers, and complex numbers simplified to reals, if possible.

Example 66 Rational complex number type simplification:

```
> let l = [9+:1, 9%1+:1, 9%2+:1, 4.5+:1, 9%1+:0, 9%2+:0, 4.5+:0.0];
> l;
[9+:1, 9L%1L+:1, 9L%2L+:1, 4.5+:1, 9L%1L+:0, 9L%2L+:0, 4.5+:0.0]
> map comp_simplify l;
[9+:1, 9L%1L+:1, 9L%2L+:1, 4.5+:1, 9L%1L, 9L%2L, 4.5+:0.0]
> map ratcomp_simplify l;
[9+:1, 9+:1, 9L%2L+:1, 4.5+:1, 9, 9L%2L, 4.5+:0.0]
```

See ‘Rational Type Simplification’_ for some details of rational type simplification.

15 String Formatting and Evaluation

15.1 The Naming of the String Conversion Functions

There are several families of functions for converting between strings and rationals.

The functions that convert from rationals to strings have names based on that of the standard function `str`. The `str_*` functions convert to a formatted string, and depend on a ‘format structure’ parameter (see [Internationalisation and Format Structures](#)). The `strs_*` functions convert to a tuple of string fragments.

The functions that convert from strings to rationals have names based on that of the standard function `eval` (`val` in *Q*). The `val_*` functions convert from a formatted string, and depend on a format structure parameter. The `sval_*` functions convert from a tuple of string fragments.

There are also `join_*` and `split_*` functions to join string fragments into formatted strings, and to split formatted strings into string fragments, respectively; these depend on a format structure parameter. These functions are not always invertible, because some of the functions reduce an error term to just a sign, e.g. `str_real_approx_dp` may round a value. Thus sometimes the `join_*` and `split_*` pairs, and the `str_*` and `val_*` pairs are not quite mutual inverses.

15.2 Internationalisation and Format Structures

Many of the string formatting functions in the following sections are parameterised by a 'format structure'. Throughout this document, the formal parameter for the format structure will be `fmt`. This is simply a record mapping some string 'codes' to functions as follows. The functions are mostly from strings to a string, or from a string to a tuple of strings.

"sm" a function mapping a sign and an unsigned mantissa (or integer) strings to a signed mantissa (or integer) string.

"se" a function mapping a sign and an unsigned exponent string to a signed exponent string.

"-s" a function mapping a signed number string to a pair containing a sign and the unsigned number string.

"gi" a function mapping an integer representing the group size and an integer string to a grouped integer string.

"gf" a function mapping an integer representing the group size and a fraction-part string to a grouped fraction-part string.

"-g" a function mapping a grouped number string to an ungrouped number string.

"zi" a function mapping an integer number string to a number string. The input string representing zero integer part is `""`, which should be mapped to the desired representation of zero. All other number strings should be returned unaltered.

"zf" a function mapping a fraction-part number string to a number string. The input string representing zero fraction part is `""`, which should be mapped to the desired representation of zero. All other number strings should be returned unaltered.

"ir" a function mapping initial and recurring parts of a fraction part to the desired format.

"-ir" a function mapping a formatted fraction part to the component initial and recurring parts.

"if" a function mapping an integer string and fraction part string to the radix-point formatted string.

"-if" a function mapping a radix-point formatted string to the component integer fraction part strings

"me" a function mapping a mantissa string and exponent string to the formatted exponential string.

"-me" a function mapping a formatted exponential string to the component mantissa and exponent strings.

"e" a function mapping an 'error' number (not string) and a number string to a formatted number string indicating the sign of the error.

"-e" a function mapping a formatted number string indicating the sign of the error to the component 'error' string (not number) and number strings.

Depending upon the format structure, some parameters of some of the functions taking a format structure may have no effect. For example, an `intGroup` parameter specifying the size of the integer digit groups will have no effect if the integer group separator is the empty string.

`create_format options` is a function that provides an easy way to prepare a 'format structure' from the simpler 'options structure'. The options structure is another record, but from more descriptive strings to a string or tuple of strings.

For example, `format_uk` is generated from `options_uk` as follows:

```
public options_uk;
const options_uk =
{
  "sign" => ("-", "", "+"),           // alternative: ("-", " ", "+")
  "exponent sign" => ("-", "", "+"),  // alternative: ("-", " ", "+")
  "group separator" => "",           // might be " " or "." or "" elsewhere
  "zero" => "0",
  "radix point" => ".",              // might be ", " elsewhere
  "fraction group separator" => "",
  "fraction zero" => "0",            // alternative: ""
  "recur brackets" => ("[" , "...]"),
  "exponent" => "*10^",              // (poor) alternative: "e"
  "error sign" => ("-", "", "+"),
  "error brackets" => ("(", ")")
};

public format_uk;
const format_uk = create_format options_uk;
```

The exponent string need not depend on the radix, as the numerals for the number radix in that radix are always "10".

Beware of using "e" or "E" as an exponent string as these have the potential of being treated as digits in, e.g., hexadecimal.

Format structures do not have to be generated via `create format`; they may also be constructed directly.

15.3 Digit Grouping

Some functions take group parameters. A value of 0 means "don't group".

15.4 Radices

The functions that produce a decimal expansion take a Radix argument. The fraction parts are expanded in that radix (or ‘base’), in addition to the integer parts. The parameter Radix is not restricted to the usual {2, 8, 10, 16}, but may be any integer from 2 to 36; the numerals (‘digits’) are chosen from [“0”, ... , “9”, “A”, ... , “Z”]. The letter-digits are always upper case.

The functions do not attach a prefix (such as “0x” for hexadecimal) to the resulting string.

15.5 Error Terms

Some functions return a value including an ‘error’ term (in a tuple) or sign (at the end of a string). Such an error is represents what the next digit would be as a fraction of the radix.

Example 67 Error term in the tuple of string ‘fragments’:

```
> strs_real_approx_sf 10 floor 3 (234567%100000);  
"+", "2", "34", 567L%1000L  
> strs_real_approx_sf 10 ceil 3 (234567%100000);  
"+", "2", "35", (-433L)%1000L
```

(See the function [strs_real_approx_sf](#).)

In strings, only the sign of the error term is given. A “+” should be read as “and a bit more”; “-” as “but a bit less”.

Example 68 Error sign in the string:

```
> str_real_approx_sf format_uk 10 0 0 floor 3 (234567%100000);  
"2.34(+)"  
> str_real_approx_sf format_uk 10 0 0 ceil 3 (234567%100000);  
"2.35(-)"
```

(See the function [str_real_approx_sf](#).)

16 Q <-> Fraction String (“i + n/d”)

16.1 Formatting to Fraction Strings

str_vulgar q returns a String representing the rational (or integer) q in the form

- “[−]n/d”

str_vulgar_or_int q returns a String representing the rational (or integer) q in one of the forms

- “[−]n/d”
- “[−]i”

str_mixed q returns a String representing the rational (or integer) q in one of the forms

- “ $i + n/d$ ”
- “ $-(i + n/d)$ ”
- “ $[-]n/d$ ”
- “ $[-]i$ ”

Example 69 The fraction string representations:

```
> let l = iterwhile (<= 3%2) (+(1%2)) (-3%2);
> l;
[(-3L)%2L, (-1L)%1L, (-1L)%2L, 0L%1L, 1L%2L, 1L%1L, 3L%2L]
> map str_vulgar l;
["-3L/2L", "-1L/1L", "-1L/2L", "0L/1L", "1L/2L", "1L/1L", "3L/2L"]
> map str_vulgar_or_int l;
["-3L/2L", "-1L", "-1L/2L", "0L", "1L/2L", "1L", "3L/2L"]
> map str_mixed l;
["-(1L+1L/2L)", "-1L", "-1L/2L", "0L", "1L/2L", "1L", "1L+1L/2L"]
```

These might be compared to the behaviour of the standard function `str`.

str x returns a string representing the value `x`.

Example 70 The standard function `str`:

```
> map str l;
["(-3L)%2L", "(-1L)%1L", "(-1L)%2L", "0L%1L", "1L%2L", "1L%1L", "3L%2L"]
```

16.2 Evaluation of Fraction Strings

val_vulgar strg returns a rational `q` represented by the string `strg` in the form

- “ $[-]n/d$ ”

Such strings can also be evaluated by the `val_mixed` function.

val_mixed strg returns a rational `q` represented by the string `strg`

- “ $i + n/d$ ”
- “ $-(i + n/d)$ ”
- “ $[-]n/d$ ” — thus `val_mixed` strictly extends `val_vulgar`
- “ $[-]i$ ”

Example 71 Evaluating fraction strings:

```
> val_vulgar "-22/7";
(-22L)%7L
> val_mixed "1L+5L/6L";
11L%6L
```

These might be compared to the behaviour of the standard function `eval`.

eval *s* evaluates the string *s*.

Example 72 The standard function `eval`:

```
> eval "1+5%6";
11L%6L
> eval "1L+5L/6L";
1.833333333333333
```

17 Q <-> Recurring Numeral Expansion String (“I.FR”)

See [Internationalisation and Format Structures](#) for information about the formatting structure to be supplied in the `fmt` parameter.

17.1 Formatting to Recurring Expansion Strings

str_real_recur `fmt radix intGroup q` returns a string (exactly) representing the rational (or integer) *q* as base-*Radix* expansion of one the forms

- “[−]int.frac”
- “[−]int.init frac part[smallest recurring frac part ...]”

Note that there is no `fracGroup` parameter.

Beware that the string returned by this function can be very long. The length of the recurring part of such a decimal expansion may be up to one less than the simplest denominator of *q*.

Example 73 The recurring radix expansion-type string representations:

```
> str_real_recur format_uk 10 3 (4000001%4); // grouped with commas
"1,000,000.25"
> str_real_recur format_uk 10 0 (4000001%4); // no grouping
"1000000.25"
> str_real_recur format_uk 10 3 (1000000%3);
"333,333.[3...]"
> str_real_recur format_uk 10 3 (1000000%7);
"142,857.[142857...]"
> str_real_recur format_uk 10 3 (-1%700);
"-0.00[142857...]"
> str_real_recur format_uk 10 3 (127%128);
"0.9921875"
> str_real_recur format_uk 2 4 (-127%128);
"-0.1111111"
> str_real_recur format_uk 16 4 (127%128);
"0.FE"
> str_real_recur format_uk 10 0 (70057%350); // 1%7 + 10001%50;
"200.16[285714...]"
```

The function allows expansion to different radices (bases).

Example 74 The recurring radix expansion in decimal and hexadecimal:

```
> str_real_recur format_uk 10 0 (1%100);  
"0.01"  
> str_real_recur format_uk 16 0 (1%100);  
"0.0[28F5C...]"
```

Example 75 The recurring radix expansion in duodecimal:

```
> str_real_recur format_uk 12 0 (1%100);  
"0.0[15343A0B62A68781B059...]"
```

Note that this bracket notation is not standard in the literature. Usually the recurring numerals are indicated by a single dot over the initial and final numerals of the recurring part, or an overline over the recurring part. For example $1/70 = 0.0\overline{142857} = 0.0142857$ and $1/3 = 0.\overline{3} = 0.3$.

strs_real_recur radix q returns a quadruple of the four strings:

- the sign,
- integer part (which is empty for 0),
- initial fraction part
- and recurring fraction part (either and both of which may be empty).

Example 76 The recurring radix expansion in decimal — the fragments:

```
> strs_real_recur 10 (100%7);  
"+", "14", "", "285714"  
> strs_real_recur 10 (-1%700);  
"- ", "", "00", "142857"  
> strs_real_recur 10 (70057%350);  
"+", "200", "16", "285714"
```

This function may be used to also, e.g. format the integer part with comma-separated groupings.

join_str_real_recur fmt intGroup sign i fracInit fracRecur formats the parts in the quadruple returned by `strs_real_recur` to the sort of string as returned by `str_real_recur`.

17.2 Evaluation of Recurring Expansion Strings

The `str_*` and `val_*` functions depend on a ‘format structure’ parameter (`fmt`) such as `format_uk`. Conversions may be performed between rationals and differently formatted strings if a suitable alternative format structure is supplied. See [Internationalisation and Format Structures](#) for information about formatting structures.

val_real_recur fmt radix strg returns the rational `q` represented by the base-radix expansion string `strg` of one the forms

- “[−]int.frac”
- “[−]int.init frac part[recurring frac part ...]”

Example 77 Conversion from the recurring radix expansion-type string representations:

```
> val_real_recur format_uk 10 "-12.345";  
(-2469L)%200L  
> val_real_recur format_uk 10 "0.3";  
3L%10L  
> val_real_recur format_uk 10 "0.[3...]";  
1L%3L  
> val_real_recur format_uk 10 ".333[33...]";  
1L%3L  
> val_real_recur format_uk 10 ".[9...]";  
1L%1L
```

sval_real_recur radix sign iStr fracStr recurPartStr returns the rational q represented by the parts

- sign
- integer part
- initial fraction part
- recurring fraction part

split_str_real_recur Fmt strg returns a tuple containing the parts

- sign
- integer part
- initial fraction part
- recurring fraction part of one the forms - “[−]int.frac” - “[−]int.init frac part[recurring frac part ...]”

18 Q <-> Numeral Expansion String (“I.F × 10E”)

See [Internationalisation and Format Structures](#) for information about the formatting structure to be supplied in the `fmt` parameter.

The exponent string “*10^” need not depend on the radix, as the numerals for the number radix in that radix are always “10”.

18.1 Formatting to Expansion Strings

Functions for Fixed Decimal Places

str_real_approx_dp *fmt radix intGroup fracGroup roundfun dp q* returns a string representing a numeral expansion approximation of *q* to *dp* decimal places, using rounding mode *roundfun* (see [Rounding to Integer](#)) *roundfun* is usually [round](#) or [round_unbiased](#). (*dp* may be positive, zero or negative; non-positive *dps* may look misleading — use e.g. scientific notation instead.)

Example 78 Decimal places:

```
> str_real_approx_dp format_uk 10 3 3 round 2 (22%7);
"3.14(+)"
> str_real_approx_dp format_uk 10 3 3 ceil 2 (22%7);
"3.15(-)"
```

strs_real_approx_dp *radix roundfun do q* returns a tuple of strings

- sign
- integer part
- fraction part

representing an expansion to a number of decimal places, together with

- the rounding “error”: a fraction representing the next numerals.

Example 79 Decimal places — the fragments:

```
> strs_real_approx_dp 10 round 2 (22%7);
"+", "3", "14", 2L%7L
> strs_real_approx_dp 10 ceil 2 (22%7);
"+", "3", "15", (-5L)%7L
```

join_str_real_approx *fmt intGroup fracGroup sign i frac err* formats the parts in the quadruple returned by [strs_real_approx_dp](#) or [strs_real_approx_sf](#) to the sort of string as returned by [str_real_approx_dp](#) or [str_real_approx_sf](#).

Functions for Significant Figures

str_real_approx_sf *fmt radix intGroup fracGroup roundfun sf q* returns a string representing a numeral expansion approximation of *q* to *sf* significant figures, using rounding mode *roundfun* (see [Rounding to Integer](#)).

Round is usually [round](#) or [round_unbiased](#). (*sf* must be positive.)

Example 80 Significant figures:

```
> str_real_approx_sf format_uk 10 3 3 floor 2 (22%7);
"3.1(+)"
```

```
> str_real_approx_sf format_uk 10 3 3 floor 2 ((-22)%7);  
"-3.2(+)"
```

strs_real_approx_sf radix roundfun sf q returns a tuple of strings

- sign,
- integer part,
- fraction part, representing an expansion to a number of significant figures, together with
- the rounding “error”: a fraction representing the next numerals

join_str_real_approx see [join_str_real_approx](#).

Functions for Scientific Notation and Engineering Notation

str_real_approx_sci fmt radix intGroup fracGroup roundfun sf q returns a string expansion with a number of significant figures in scientific notation, using rounding mode [roundfun](#) (see [Rounding to Integer](#)).

(sf must be positive; expStep is usually 3, radix is usually 10, [roundfun](#) is usually [round](#) or [round_unbiased](#); [str_real_approx_sci](#) is equivalent to [str_real_approx_eng](#) (below) with expStep = 1.)

strs_real_approx_sci radix roundfun sf q returns a tuple of strings:

- sign of mantissa,
- integer part of mantissa,
- fraction part of mantissa,
- sign of exponent,
- exponent magnitude

representing an expansion to a number of significant figures in scientific notation together with

- the rounding “error”: a fraction representing the next numerals.

str_real_approx_eng fmt expStep radix intGroup fracGroup round sf q returns a string expansion with a number of significant figures in engineering notation, using rounding mode [roundfun](#).

The ExpStep parameter specifies the granularity of the exponent; specifically, the exponent will always be divisible by expStep.

(sf must be positive; expStep is usually 3 and must be positive, radix is usually 10, [roundfun](#) is usually [round](#) or [round_unbiased](#).)

Example 81 Engineering notation:

```
> str_real_approx_eng format_uk 3 10 3 3 round 7 (rational 999950);
"999.950,0*10^3"
> str_real_approx_eng format_uk 3 10 3 3 round 4 999950;
"1.000*10^6(-)"
```

strs_real_approx_eng expStep radix roundfun sf q returns a tuple of strings:

- sign of mantissa,
- integer part of mantissa,
- fraction part of mantissa,
- sign of exponent,
- exponent magnitude

representing an expansion to a number of significant figures in engineering notation together with

- the rounding “error”: a fraction representing the next numerals.

Example 82 Engineering notation — the fragments:

```
> strs_real_approx_eng 3 10 round 7 (rational 999950);
"+", "999", "9500", "+", "3", 0L%1L
> strs_real_approx_eng 3 10 round 4 999950;
"+", "1", "000", "+", "6", (-1L)%20L
```

join_str_real_eng fmt intGroup fracGroup mantSign mantI mantF rac expSign expI err formats the parts in the quadruple returned by `strs_real_approx_eng` or `strs_real_approx_sci` to the sort of string as returned by `str_real_approx_eng` or `str_real_approx_sci`.

18.2 Evaluation of Expansion Strings

The `str_*` and `val_*` functions depend on a ‘format structure’ parameter (`fmt`) such as `format_uk`. Conversions may be performed between rationals and differently formatted strings if a suitable alternative format structure is supplied. See [Internationalisation and Format Structures](#) for information about formatting structures.

val_real_eng fmt radix strg returns the rational `q` represented by the base-radix expansion string `strg` of one the forms

- “[−]int.frac”
- “[−]int.frace[−]exponent”

Example 83 Conversion from the recurring radix expansion-type string representations:

```
> val_real_eng format_uk 10 "-12.345";
(-2469L)%200L
> val_real_eng format_uk 10 "-12.345*10^2";
(-2469L)%2L
```

sval_real_eng radix signStr mantIStr mantF racStr expSignStr expStr returns the rational q represented by the parts

- sign
- integer part of mantissa
- fraction part of mantissa
- sign of exponent
- exponent

split_str_real_eng fmt strg — returns a tuple containing the string parts

- sign
- integer part of mantissa
- fraction part of mantissa
- sign of exponent
- exponent
- the “error” sign

of one the forms

- “[−]int.frac”
- “[−]int.frac $\times 10^{[−]exponent}$ ”

These functions can deal with the fixed decimal places, the significant figures and the scientific notation in addition to the engineering notation.

19 Numeral String -> Q — Approximation

This section describes functions to approximate by a rational a real number represented by a string. See [R -> Q — Approximation](#) for approximation by a rational of a double.

The `str_*` and `val_*` functions depend on a ‘format structure’ parameter (`fmt`) such as `format uk`. Conversions may be performed between rationals and differently formatted strings if a format structure is supplied. See [Internationalisation and Format Structures](#) for information about formatting structures.

val_eng_approx_epsilon fmt radix epsilon strg Find the least complex rational approximation q to the number represented by the base-radix expansion string `str` in one of the forms

- “[−]int.frac”
- “[−]int.frac $\times 10^{[−]exponent}$ ”

that is ϵ -close. That is find a q such that $|q - \text{eval str}| \leq \epsilon$.

Example 84 Rational from a long string:

```
> let strg = "123.456,789,876,543,212,345,678,987,654,321*10^27";
> let x = val_real_eng format_uk 10 strg;
> x;
123456789876543212345678987654321L%1000L
> let q = val_eng_approx_epsilon format_uk 10 (1%100) strg;
> q;
1975308638024691397530863802469L%16L
> double (x - q);
0.0085
> str_real_approx_eng format_uk 3 10 3 3 round 30 q;
"123.456,789,876,543,212,345,678,987,654*10^27(+)"
> str_real_approx_eng format_uk 3 10 3 3 round 42 q;
"123.456,789,876,543,212,345,678,987,654,312,500,000,000*10^27"
> double q;
1.23456789876543e+029
```

val_eng_interval_epsilon fmt radix epsilon strg Find the least complex rational interval containing the number represented by the base-radix expansion string strg in one of the forms

- "[−]int.frac"
- "[−]int.frac × 10^[−]exponent"

that is "-small.

val_eng_approx_max_den fmt radix maxDen strg Find the closest rational approximation to the number represented by the base-rRadix expansion string strg in one of the forms

- "[−]int.frac"
- "[−]int.frac × 10^[−]exponent"

that has a denominator no greater than maxDen. (maxDen > 0)

val_eng_interval_max_den fmt radix maxDen strg Find the smallest rational interval containing the number represented by the base-radix expansion string strg in one of the forms

- "[−]int.frac"
- "[−]int.frac × 10^[−]exponent"

that has endpoints with denominators no greater than maxDen. (maxDen > 0)

Example 85 Other rationals from a long string:

```
> val_eng_approx_epsilon format_uk 10 (1%100) strg;
1975308638024691397530863802469L%16L
> val_eng_interval_epsilon format_uk 10 (1%100) strg;
interval::Ivl (3086419746913580308641974691358L%25L)
(3456790116543209945679011654321L%28L)
> val_eng_approx_max_den format_uk 10 100 strg;
9999999980000000199999998000000L%81L
```

```
> val_eng_interval_max_den format_uk 10 100 strg;  
interval::Ivl 9999999980000000199999998000000L%81L  
3456790116543209945679011654321L%28L
```

Pure-CSV - Comma Separated Value Interface for the Pure Programming Language

Version 1.4, November 14, 2010

Eddie Rucker <erucker@bmc.edu>

Copyright (c) 2008, 2009, 2010 by Eddie Rucker

License: 3-clause BSD license, see the accompanying COPYING file for details

The CSV library provides an interface for reading and writing comma separated value files. The module is very loosely based on Python's CSV module (<http://docs.python.org/lib/module-csv.html>).

1 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-csv-1.4.tar.gz>.

Run `make` to compile the module and `make install` (as root) to install it in the Pure library directory. This requires GNU make. The `make install` step is only necessary for system-wide installation.

The `make` utility tries to guess your Pure installation directory and platform-specific setup. If it gets this wrong, you can set some variables manually. In particular, `make install prefix=/usr` sets the installation prefix, and `make PIC=-fPIC` or some similar flag might be needed for compilation on 64 bit systems. Please see the Makefile for details.

2 Usage

Data records are represented as vectors or lists of any Pure values. Values are converted as necessary and written as a group of strings, integers, or doubles separated by a delimiter. Three predefined dialects are provided; `default` (record terminator= `\n`), `RFC4180` (record

terminator= \r\n), and Excel. Procedures are provided to create other CSV dialects. See (<http://www.ietf.org/rfc/rfc4180.txt>) for more details about the RFC4180 standard.

2.1 Handling Errors

error msg is an error handling term. Operations resulting in parse errors, memory errors, or read/write errors produce a special `csv::error msg` term, where `msg` is a string describing the particular error. Your application should either check for these or have `csv::error` defined to directly handle errors in some way (e.g., provide a default value, or raise an exception).

2.2 Creating Dialects

dialect record creates a dialect from a record of dialect option pairs. The dialect object is freed automatically when exiting the pure script. The list of possible options and option values are presented below.

- **delimiter** - Character used to separate fields.
 - Value - any string.
 - Default - ", ".
- **escape** - Embedded escape character used to embed a delimiter, escape, or terminator into unquoted fields. If the escape character is not null, then the quote character is ignored.
 - Value - any string.
 - Default - "".
 - Reading - The escape character is dropped and the next char is inserted into the field.
 - Writing - The escape character is written into the output stream before the delimiter, escape, or return character.
- **quote** - Quotes are used to embed delimiters, quotes, or terminators into a field.
 - Value - any string.
 - Default - "\".
 - Notes - Embedded quotes are doubled. The escape option must be the null string.
- **terminator** - Record termination string.
 - Value - One of {LF, CRLF} where LF = "\n" and CRLF = "\r\n".
 - Default - LF.

- Reading - Carriage returns, "\r", are always ignored except when embedded in quotes or escaped.
- Writing - Every written record is terminated with the the defined terminator string.
- `quote_flag` - Sets the quoting style of strings and/or numbers.
 - Value - One of {ALL, STRINGS, MINIMAL}.
 - Default - ALL.
 - Reading -
 1. ALL - Every field is read as a string.
 2. STRING, MINIMAL - Fields within quotes and fields that cannot be converted to integers or doubles are read as strings.
 - Writing -
 1. ALL - Every field is written within quotes.
 2. STRING - Only fields of type string are quoted.
 3. MINIMAL - Only fields containing embedded quotes, terminators, or delimiters are written within quotes.
- `space_around_quoted_field` - Determines how white space between quotes and delimiters should be treated.
 - Value - One of {N0, LEFT, RIGHT, BOTH}.
 - Default - N0.
 - Reading -
 1. N0 - Follows RFC4180 rules.
 2. LEFT - Allows space before a quoted field.
 3. RIGHT - Allows space between a quoted field and a delimiter.
 4. BOTH - Allows space before and after a quoted field.
 - Writing - fields are never written with space before a quoted field or between a quoted field and a delimiter.
 - Notes this option does not affect space within quotes or fields written using the escape string option.
- `trim_space` - trim white space before or after field contents.
 - Value - One of {N0, LEFT, RIGHT, BOTH}.
 - Default - N0.
 - Reading -
 1. N0 - Reading follows RFC4180 rules.

2. LEFT, RIGHT, or BOTH - The field is trimmed accordingly. Use *caution* because trimming may allow automatic conversion of numbers if the `quote_flag` is set to MINIMAL.

– Writing –

1. NO - Reading follows RFC4180 rules
2. LEFT, RIGHT, or BOTH - Trimming space is probably a bad idea since leading or trailing space may be significant for other applications.

The following example illustrates the construction of a dialect for reading tab delimited files without quoted strings.

Example

```
> using csv;
> using namespace csv;
> let d = dialect {delimiter=>"\t", quote_flag=>STRING};
> d;
#<pointer 0x9d75280>
>
```

2.3 Opening CSV Files

open name::string opens a CSV file for reading using the default dialect. If the file does not exist, the error `msg` rule is invoked.

open (name::string, rw_flag::string) opens a CSV file for reading, writing, or appending using the default dialect. Valid `rw_flag` values are "r" for reading, "w" for writing, and "a" for appending. If the file does not exist when opened for reading, the error `msg` rule is invoked. When a file is opened for writing and the file exists, the old file is overwritten. If the file does not exist, a new empty file is created. When a file is opened for appending and the file exists, new records are appended to the end of the file, otherwise a new empty file is created.

open (name::string, rw_flag::string, d::pointer) exactly as above except reading/writing is done according to a user defined dialect `d`.

open (name::string, rw_flag::string, d::pointer, opts@(_:_)) exactly as above except allows for list output or header options when reading.

1. If `opts` contains LIST, the output of `getr`, `fgetr`, and `fgetr_lazy` is a list instead of a vector.
2. If `opts` contains HEADER, the first line of the file is automatically read and parsed as a record where entries are `key=>position` pairs where `key` is a string and `position` is an integer denoting the location of a field within the record. The header record may be accessed by `header`.

Examples

```
> using csv;
> using namespace csv;
> let d = dialect {delimiter=>"\t"};
> let f = open ("junk.csv", "w", d);
> putr f {"hello",123,"",3+:4,world};
()
> close f;
()
> let f = open ("junk.csv", "r", d);
> getr f;
{"hello","123","","3+:4","world"}
>
```

Suppose our file “test.csv” is as presented below.

```
ir$ more test.csv
NAME,TEST1,TEST2
"HOPE, BOB",90,95
"JONES, SALLY",88,72
"RED, FEEFEE",45,52
```

Notice how the LIST option affects the return of `getr` and how the HEADER option may be used to index records.

```
> using csv;
> using namespace csv;
> let d = dialect {quote_flag=>MINIMAL};
> let f = open ("test.csv", "r", d, [LIST,HEADER]);
> let r = getr f;
> r!0;
"HOPE, BOB"
> let k = header f;
> k;
["NAME"=>0,"TEST1"=>1,"TEST2"=>2]
> r!(k!"NAME");
"HOPE, BOB"
> r!!(k!!["NAME","TEST1"]);
["HOPE, BOB",95]
>
```

2.4 File Reading Functions

header csv_file::pointer returns the record of `key=>position` pairs when opened by `csv::open` using the header option. If the file was opened without the HEADER option, `{}` is returned.

getr csv_file::pointer reads from a `csv_file` opened by `csv::open` and returns a record represented as a row matrix. Reading from a file opened for writing or appending invokes the error `msg` rule.

fgetr csv_file::pointer reads a whole file and returns a list of records. This procedure

should only be used on data files that are small enough to fit in the computer's primary memory. Reading from a file opened for writing or appending invokes the error `msg` rule.

fgetr_lazy csv_file::pointer Lazy version of `fgetr`.

2.5 File Writing Functions

When modifying CSV files that will be imported into Microsoft Excel, fields with significant leading 0s should be written using a `"="0..."` formatting scheme. This same technique will work for preserving leading space too. Again, this quirk should only be necessary for files to be imported into MS Excel.

putr csv_file::pointer rec::matrix writes a record in row matrix format to `csv_file`. Writing to a file opened for reading invokes the error `msg` rule.

fputr csv_file::pointer l@(_:_) writes a list of records where each record is a row matrix to `csv_file`. Writing to a file opened for reading invokes the error `msg` rule.

2.6 Examples

The first example shows how to write and read a default CSV file.

```
> using csv;
> using namespace csv;
> let f = open ("testing.csv", "w");
> fputr f [{"bob",3.9,"",-2},{fred",-11.8,"",0},{mary",2.3,"$",11}];
()
> close f;
()
> let f = open "testing.csv";
> fgetr f;
[{"bob","3.9","",-2},{fred",-11.8,"",0},{mary","2.3","$","11"}]
> close f;
>
```

The second example illustrates how to write and read a CSV file using automatic conversions.

```
> using csv;
> using namespace csv;
> let d = dialect {quote_flag=>MINIMAL};
> let f = open ("test.csv", "w", d);
> putr f {"I","",-4,1.2,2%4,like};
()
> putr f {"playing","the",0,-0.2,1+:4,drums};
()
> close f;
()
> let f = open ("test.csv", "r", d);
```

```
> fgetr f;
[{"I","", -4, 1.2, "2%4", "like"}, {"playing", "the", 0, -0.2, "1+:4", "drums"}]
> close f;
()
>
```

Records containing quotes, delimiters, and line breaks are also properly handled.

```
> using csv;
> using namespace csv;
> let d = dialect {quote_flag=>STRING};
> let f = open ("test.csv", "w", d);
> fputr f [{"this\nis\n", 1}, {"a \"test\"", 2}];
()
> close f;
()
> let f = open ("test.csv", "r", d);
> fgetr f;
[{"this\nis\n", 1}, {"a \"test\"", 2}]
> close f;
()
>
```

Consider the following hand written CSV file. According to RFC4180, this is not a valid CSV file. However, by using the `space_around_quoted_field`, the file can still be read.

```
erucker:$ more test.csv
  "this",   "is",   "not", "valid"
```

```
> using csv;
> using namespace csv;
> let f = open "test.csv";
> getr f;
csv::error "parse error at line 1"
> let d = dialect {space_around_quoted_field=>BOTH};
> let f = open ("test.csv", "r", d);
> getr f;
{"this","is","not","valid"}
>
```

The `trim_space` flag should be used with caution. A field with space in front of a number should be interpreted as a string, but consider the following file.

```
erucker:$ more test.csv
"  this  ", 45 ,23, hello
```

Now observe the differences for the two dialects below.

```
> using csv;
> using namespace csv;
> let d = dialect {trim_space=>BOTH};
> let f = open ("test.csv", "r", d);
> getr f;
```

```
{"this","45","23","hello"}
> let d = dialect {trim_space=>BOTH, quote_flag=>MINIMAL};
> let f = open ("test.csv", "r", d);
> getr f;
{"this",45,23,"hello"}
>
```

The `trim_space` flag also affects writing.

```
> using csv;
> using namespace csv;
> let d = dialect {trim_space=>BOTH};
> let f = open ("test.csv", "w", d);
> putr f {"  this  ","  45  "};
()
> close f;
()
> quit
```

```
erucker:$ more test.csv
"this","45"
```

For the last example a tab delimiter is used, automatic conversions is on, and records are represented as lists. Files are automatically closed when the script is finished.

```
> using csv;
> using namespace csv;
> let d = dialect {quote_flag=>MINIMAL, delimiter=>"\t"};
> let f = open ("test.csv", "w", d, [LIST]);
> fputr f [{"a","b",-4.5,""}, {"c","d",2.3,"-"}];
()
> close f;
()
> let f = open ("test.csv", "r", d, [LIST]);
> fgetr f;
[{"a","b",-4.5,""}, {"c","d",2.3,"-"}]
> quit
```

pure-fastcgi: FastCGI module for Pure

Version 0.2, November 14, 2010

Albert Gräf <Dr.Graef@t-online.de>

This module lets you write **FastCGI** scripts with Pure, to be run by web servers like **Apache**. Compared to normal CGI scripts, this has the advantage that the script keeps running and may process as many requests from the web server as you like, in order to reduce startup times and enable caching techniques. Most ordinary CGI scripts can be converted to use FastCGI with minimal changes.

1 Copying

Copyright (c) 2009 by Albert Graef. pure-fastcgi is distributed under a 3-clause BSD-style license, please see the included COPYING file for details.

2 Installation

Besides Pure, you'll need to have the **FastCGI** library installed to compile this module. Also, to run FastCGI scripts, your web server must be configured accordingly; see the documentation of FastCGI and your web server for details.

Get the latest source from <http://pure-lang.googlecode.com/files/pure-fastcgi-0.2.tar.gz>.

Running `make` compiles the module, `make install` installs it in your Pure library directory. You might have to adjust the path to the `fcgi_stdio.h` header file in `fastcgi.c` and/or the option to link in the FastCGI library in the Makefile.

The Makefile tries to guess the host system type and Pure version, and set up some platform-specific things accordingly. If this doesn't work for your system then you'll have to edit the Makefile accordingly.

3 Usage

pure-fastcgi provides the accept function with which you tell the FastCGI server that your script is ready to accept another request. It also overrides a number of standard I/O functions so that they talk to the server instead. These routines are all in the fastcgi namespace. In your Pure script, you can set up a simple loop to process requests as follows:

```
#!/usr/local/bin/pure -x

using fastcgi;
using namespace fastcgi;

extern char *getenv(char*);

main count = main count when
    count = count+1;
    printf "Content-type: text/html\n\n\
<title>FastCGI Hello! (Pure, fcgi_stdio library)</title>\
<h1>FastCGI Hello! (Pure, fcgi_stdio library)</h1>\
Request number %d running on host <i>%s</i>\n"
    (count,getenv "SERVER_NAME");
end if accept >= 0;

main 0;
```

(You might have to adjust the “shebang” in the first line above, so that the shell finds the Pure interpreter. Also, remember to make the script executable. If you’re worried about startup times, or if your operating system doesn’t support shebangs, then you can also use the Pure interpreter to compile the script to a native executable instead.)

This script keeps running until accept returns -1 to indicate that the script should exit. Each call to accept blocks until either a request is available or the FastCGI server detects an error or other kind of termination condition. As with ordinary CGI, additional information about the request is available through various environment variables. A list of commonly supported environment variables and their meaning can be found in [The Common Gateway Interface](#) specification.

A number of other routines are provided to deal with data filters, finish a request and set an exit status for a request. These correspond to operations provided by the FastCGI library, see the FastCGI documentation and the FCGI_Accept(3), FCGI_StartFilterData(3), FCGI_Finish(3) and FCGI_SetExitStatus(3) manpages for details. An interface to the FCGI_ToFILE macro is also available. Note that in Pure these functions are called accept, start_filter_data, finish, set_exit_status and to_file, respectively, and are all declared in the fastcgi namespace. A detailed listing of all routines can be found in the fastcgi.pure module.

Please see the examples subdirectory in the pure-fastcgi sources for some more elaborate examples.

Note that to run your FastCGI scripts in a browser, your web server must have the FastCGI module loaded and must also be set up to execute the scripts. E.g., when using Apache, the

following configuration file entry will set up a directory for FastCGI scripts:

```
ScriptAlias /fastcgi-bin/ "/srv/www/fastcgi-bin/"
<Location /fastcgi-bin/>
    Options ExecCGI
    SetHandler fastcgi-script
    Order allow,deny
    Allow from all
</Location>
```

(Replace `fastcgi-script` with `fcgid-script` if you're running `mod_fcgid` rather than `mod_fastcgi`.)

Put this entry into `http.conf` or a similar file provided by your Apache installation (usually under `/etc/apache2`), and restart Apache. After that you can just throw your scripts into the `fastcgi-bin` directory to have them executed via an URL like `http://localhost/fastcgi-bin/myscript`.

You can also set up a handler for the `.pure` filename extension as follows:

```
<IfModule mod_fastcgi.c>
<FilesMatch "\.pure$">
    AddHandler fastcgi-script .pure
    Options +ExecCGI
</FilesMatch>
</IfModule>
```

(Again, you'll have to adjust the `IfModule` statement and replace `fastcgi-script` with `fcgid-script` if you're running `mod_fcgid`.) After that you should be able to execute scripts with the proper extension anywhere under your server's document root.

Pure-ODBC - ODBC interface for the Pure programming language

Version 0.5, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

Jiri Spitz <jiri.spitz@bluetone.cz>

This module provides a simple ODBC interface for the Pure programming language, which lets you access a large variety of open source and commercial database systems from Pure. ODBC a.k.a. “Open Database Connectivity” was originally developed by Microsoft for Windows, but is now available on many different platforms, and two open source implementations exist for Unix-like systems: iODBC (<http://www.iodbc.org>) and unixODBC (<http://www.unixodbc.org>).

ODBC has become the industry standard for portable and vendor independent database access. Most modern relational databases provide an ODBC interface so that they can be used with this module. This includes the popular open source DBMSs MySQL (<http://www.mysql.com>) and PostgreSQL (<http://www.postgresql.org>). The module provides the necessary operations to connect to an ODBC data source and retrieve or modify data using SQL statements.

To make this module work, you must have an ODBC installation on your system, as well as the driver backend for the DBMS you want to use (and, of course, the DBMS itself). You also have to configure the DBMS as a data source for the ODBC system. On Windows this is done with the ODBC applet in the system control panel. For iODBC and unixODBC you can either edit the corresponding configuration files (/etc/odbc.ini and/or ~/.odbc.ini) by hand, or use one of the available graphical setup tools. More information about the setup process can be found on the iODBC and unixODBC websites.

1 Copying

Copyright (c) 2009 by Albert Graef <Dr.Graef@t-online.de>.

Copyright (c) 2009 by Jiri Spitz <jiri.spitz@bluetone.cz>.

pure-odbc is free software: you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

pure-odbc is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

2 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-odbc-0.5.tar.gz>.

Run `make` to compile the module and `make install` (as root) to install it in the Pure library directory. This requires GNU make, and of course you need to have Pure installed. The only other dependency is the GNU Multiprecision Library (GMP).

`make` tries to guess your Pure installation directory and platform-specific setup. If it gets this wrong, you can set some variables manually. In particular, `make install prefix=/usr` sets the installation prefix, and `make PIC=-fPIC` or some similar flag might be needed for compilation on 64 bit systems. The variable `ODBCLIB` specifies the ODBC library to be linked with. The default value is `ODBCLIB=-lodbc`. Please see the Makefile for details.

3 Opening and Closing a Data Source

To open an ODBC connection, you have to specify a “connect string” which names the data source to be used with the `odbc::connect` function. A list of available data sources can be obtained with the `odbc::sources` function. For instance, on my Linux system running MySQL and PostgreSQL it shows the following:

```
> odbc::sources;
[("myodbc", "MySQL ODBC 2.50"), ("psqlodbc", "PostgreSQL ODBC")]
```

The first component in each entry of the list is the name of the data source, which can be used as the value of the `DSN` option in the connect string, the second component provides a short description of the data source.

Likewise, the list of ODBC drivers available on your system can be obtained with the `odbc::drivers` function which returns a list of pairs of driver names and attributes. (Older ODBC implementations on Unix lacked this feature, but it seems to be properly supported in recent unixODBC implementations at least.) On Windows this function can be used to determine a legal value for the `DRIVER` attribute in the connect string, see below.

The `odbc::connect` function is invoked with a single parameter, the connect string, which is used to describe the data source and various other parameters such as user id and password.

For instance, on my system I can connect to the local myodbc data source from above as follows:

```
> let db = odbc::connect "DSN=myodbc";
```

The `odbc::connect` function returns a pointer to an `ODBCHandle` object which is used to refer to the database connection in the other routines provided by this module. An `ODBCHandle` object is closed automatically when it is no longer accessible. You can also close it explicitly with a call to the `odbc::disconnect` function:

```
> odbc::disconnect db;
```

After `odbc::disconnect` has been invoked on a handle, any further operations on it will fail.

`odbc::connect` allows a number of attributes to be passed to the ODBC driver when opening the database connection. E.g., here's how to specify a username and password; note that the different attributes are separated with a semicolon:

```
> let db = odbc::connect "DSN=myodbc;UID=root;PWD=guess";
```

The precise set of attributes in the connect string depends on your ODBC driver, but at least the following options should be available on most systems. (Case is insignificant in the attribute names, so e.g. the `DATABASE` attribute may be specified as either `DATABASE`, `Database` or `database`.)

- `DSN=<data source name>`
- `HOST=<server host name>`
- `DATABASE=<database path>`
- `UID=<user name>`
- `PWD=<password>`

The following attributes appear to be Windows-specific:

- `FILEDSN=<DSN file name>`
- `DRIVER=<driver name>`
- `DBQ=<database file name>`

Using the `FILEDSN` option you can establish a connection to a data source described in a `.dsn` file on Windows, as follows:

```
> odbc::connect "FILEDSN=test.dsn";
```

On Windows it is also possible to directly connect to a driver and name a database file as the data source. For instance, using the MS Access ODBC driver you can connect to a database file `test.mdb` as follows:

```
> odbc::connect "DRIVER=Microsoft Access Driver (*.mdb);DBQ=test.mdb";
```

SQLite (<http://www.sqlite.org>) provides another way to get a database up and running quickly. For that you need the SQLite library and the SQLite ODBC driver available at <http://www.ch-werner.de/sqliteodbc>. Then you can open an SQLite database as follows (the database file is named with the DATABASE attribute and is created automatically if it doesn't exist):

```
> odbc::connect "DSN=SQLite3 Datasource;Database=test.db";
```

SQLite generally performs very well if you avoid some pitfalls (in particular, big batches of updates/inserts should be done within a transaction, otherwise they will take forever). It is certainly good enough for smaller databases and very easy to set up. Basically, after installing SQLite and its ODBC driver you're ready to go immediately. This makes it a very convenient alternative if you don't want to go through the tedium of setting up one of the big hulking DBMS.

4 Getting Information about a Data Source

You can get general information about an open database connection with the `odbc::info` function. This function returns a tuple of strings with the following items (see the description of the `SQLGetInfo()` function in the ODBC API reference for more information):

- `DATA_SOURCE_NAME`: the data source name
- `DATABASE_NAME`: the default database
- `DBMS_NAME`: the host DBMS name
- `DBMS_VER`: the host DBMS version
- `DRIVER_NAME`: the name of the ODBC driver
- `DRIVER_VER`: the version of the ODBC driver
- `DRIVER_ODBC_VER`: the ODBC version supported by the driver
- `ODBC_VER`: the ODBC version of the driver manager

E.g., here is what the connection to MySQL shows on my Linux system:

```
> odbc::info db;
"myodbc", "test", "MySQL", "5.0.18", "myodbc3.dll", "03.51.12", "03.51", "03.52"
```

The `odbc` module also provides a number of operations to retrieve a bunch of additional meta information about the given database connection. In particular, the `odbc::getinfo` function provides a direct interface to the `SQLGetInfo()` routine. The result of `odbc::getinfo` is a pointer which can be converted to an integer or string value, depending on the type of information requested. For instance:

```
> get_short $ odbc::getinfo db odbc::SQL_MAX_TABLES_IN_SELECT;
31
```

```
> cstring_dup $ odbcc::getinfo db odbcc::SQL_IDENTIFIER_QUOTE_CHAR;
""
```

Information about supported SQL data types is available with the `odbcc::typeinfo` routine (this returns a lot of data, see `odbcc.pure` for an explanation):

```
> odbcc::typeinfo db odbcc::SQL_ALL_TYPES;
```

Moreover, information about the tables in the current database, as well as the structure of the tables and their primary and foreign keys can be retrieved with the `odbcc::tables`, `odbcc::columns`, `odbcc::primary_keys` and `odbcc::foreign_keys` functions:

```
> odbcc::tables db;
[("event", "TABLE"), ("pet", "TABLE")]

> odbcc::columns db "pet";
[("name", "varchar", "NO", "''), ("owner", "varchar", "YES", odbcc::SQLNULL),
 ("species", "varchar", "YES", odbcc::SQLNULL), ("sex", "char", "YES", odbcc::SQLNULL),
 ("birth", "date", "YES", odbcc::SQLNULL), ("death", "date", "YES", odbcc::SQLNULL)]

> odbcc::primary_keys db "pet";
["name"]

> odbcc::foreign_keys db "event";
[("name", "pet", "name")]
```

This often provides a convenient and portable means to retrieve basic information about table structures, at least on RDBMS which properly implement the corresponding ODBC calls. Also note that while this information is also available through special system catalogs in most databases, the details of accessing these vary a lot among implementations.

5 Executing SQL Queries

As soon as a database connection has been opened, you can execute SQL queries on it using the `sql` function which executes a query and collects the results in a list. Note that SQL queries generally come in two different flavours: queries returning data (so-called *result sets*), and statements modifying the data (which have as their result the number of affected rows). The `sql` function returns a nonempty list of lists (where the first list denotes the column titles, and each subsequent list corresponds to a single row of the result set) in the former, and the row count in the latter case.

For instance, here is how you can select some entries from a table. (The following examples assume the sample “menagerie” tables from the MySQL documentation. The `initdb` function in the `examples/menagerie.pure` script can be used to create these tables in your default database.)

```
> odbcc::sql db "select name,species from pet where owner='Harold'" [];
[["name", "species"], ["Fluffy", "cat"], ["Buffy", "dog"]]
```

Often the third parameter of `sql`, as above, is just the empty list, indicating a parameter-less query. Queries involving marked input parameters can be executed by specifying the parameter values in the third argument of the `sql` call. For instance:

```
> odbcc::sql db "select name,species from pet where owner=?" ["Harold"];  
[["name","species"],["Fluffy","cat"],["Buffy","dog"]]
```

Multiple parameters are specified as a list:

```
> odbcc::sql db "select name,species from pet where owner=? and species=?"  
> ["Harold","cat"];  
[["name","species"],["Fluffy","cat"]]
```

Parameterized queries are particularly useful for the purpose of inserting data into a table:

```
> odbcc::sql db "insert into pet values (?, ?, ?, ?, ?, ?)"  
> ["Puffball","Diane","hamster","f","1999-03-30",odbcc::SQLNULL];  
1
```

In this case we could also have hard-coded the data to be inserted right into the SQL statement, but a parameterized query like the one above can easily be applied to a whole collection of data rows, e.g., as follows:

```
> do (odbcc::sql db "insert into pet values (?, ?, ?, ?, ?, ?)") data;
```

Parameterized queries also let you insert data which cannot be specified easily inside an SQL query, such as long strings or binary data.

The following SQL types of result and parameter values are recognized and converted to/from the corresponding Pure types:

SQL value/type	Pure value/type
SQL NULL (no value)	odbcc::SQLNULL
integer types (INTEGER and friends)	int
64-bit integers	bigint
floating point types (REAL, FLOAT and friends)	double
binary data (BINARY, BLOB, etc.)	(size, data)
character strings (CHAR, VARCHAR, TEXT, etc.)	string

Note the special constant (nullary symbol) `odbcc::SQLNULL` which is used to represent SQL NULL values.

Also note that binary data is specified as a pair `(size, data)` consisting of an `int` or `bigint` size which denotes the size of the data in bytes, and a pointer data (which must not be a null pointer unless size is 0 as well) pointing to the binary data itself.

All other SQL data (including, e.g., `TIME`, `DATE` and `TIMESTAMP`) is represented in Pure using its character representation, encoded as a Pure string.

Some databases also allow special types of queries (e.g., “batch” queries consisting of multiple SQL statements) which may return multiple result sets and/or row counts. The `sql` function only returns the first result set, which is appropriate in most cases. If you need to determine all result sets returned by a query, the `msql` function must be used. This function

is invoked in exactly the same way as the `sql` function, but returns a list with all the result sets and/or row counts of the query.

Example:

```
> odbc::msql db "select * from pet; select * from event" [];
```

This will return a list with two result sets, one for each query.

6 Low-Level Operations

The `sql` and `msql` operations are in fact just ordinary Pure functions which are implemented in terms of the low-level operations `sql_exec`, `sql_fetch`, `sql_more` and `sql_close`. You can also invoke these functions directly if necessary. The `sql_exec` function starts executing a query and returns either a row count or the column names of the first result set as a tuple of strings. After that you can use `sql_fetch` to obtain the results in the set one by one. When all rows have been delivered, `sql_fetch` fails. The `sql_more` function can then be used to check for additional result sets. If there are further results, `sql_more` returns either the next row count, or a tuple of column names, after which you can invoke `sql_fetch` again to obtain the data rows in the second set, etc. When the last result set has been processed, `sql_more` fails.

Example:

```
> odbc::sql_exec db "select name,species from pet where owner='Harold'" [];  
["name","species"]  
> odbc::sql_fetch db; // get the 1st row  
["Fluffy","cat"]  
> odbc::sql_fetch db; // get the 2nd row  
["Buffy","dog"]  
> odbc::sql_fetch db; // no more results  
odbc::sql_fetch #<pointer 0x24753e0>  
> odbc::sql_more db; // no more result sets  
odbc::sql_more #<pointer 0x24753e0>
```

Moreover, the `sql_close` function can be called at any time to terminate an SQL query, after which subsequent calls to `sql_fetch` and `sql_more` will fail:

```
> odbc::sql_close db; // terminate query  
( )
```

This is not strictly necessary (it will be done automatically as soon as the next SQL query is invoked), but it is useful in order to release all resources associated with the query, such as parameter values which have to be cached so that they remain accessible to the SQL server. Since these parameters in some cases may use a lot of memory it is better to call `sql_close` as soon as you are finished with a query. This is also done automatically by the `sql` and `msql` functions.

Also note that only a single query can be in progress per database connection at any one time. That is, if you invoke `sql_exec` to initiate a new query, a previous query will be terminated

automatically. (However, it is possible to execute multiple queries on the same database simultaneously, if you process them through different connections to that database.)

The low-level operations are useful when you have to deal with large result sets where you want to avoid to build the complete list of results in main memory. Instead, these functions allow you to process the individual elements immediately as they are delivered by the `sql_fetch` function. (An alternative method which combines the space efficiency of immediate processing with the convenience of the list representation is discussed in the following section.) Using the low-level operations you can also build your own specialized query engines; take the definitions of `sql` or `msql` as a start and change them according to your needs.

7 Lazy Processing

As an experimental feature, the `odbc` module also provides two operations `odbc::lsql` and `odbc::lmsql` which work like `odbc::sql` and `odbc::msql` (see [Executing SQL Queries](#) above), but return lazy lists (streams) instead. This offers the convenience of a list-based representation without the overhead of keeping entire result sets in memory, which can be prohibitive when working with large amounts of data.

These functions are invoked just like `odbc::sql` and `odbc::msql`, but they return a lazy list of rows (or a lazy list of lazy lists of rows in the case of `lmsql`). For instance:

```
> odbc::lsql db "select * from pet" [];
["name", "owner", "species", "sex", "birth", "death"]:#<thunk 0x7ffbb9aa2eb8>
```

Note that the tail of the result list is “thunked” and will only be produced on demand, as you traverse the list. As a simple example, suppose that we just want to print the name field of each data row:

```
> using system;
> do (\(name:_)>puts name) $ tail $ odbc::lsql db "select * from pet" [];
Fluffy
Claws
Buffy
Fang
Bowser
Chirpy
Whistler
Slim
()
```

Here only one row is in memory at any time while the `do` function is in progress. This keeps memory requirements much lower than when using the `odbc::sql` function which first loads the entire result set into memory. Another advantage is that only those data rows are fetched from the database which are actually needed in the course of the computation. This can speed up the processing significantly if only a part of the result set is needed. For instance, in the following example we only look at the first two data rows until the desired row is found, so the remaining rows are never fetched from the database:

```
> head [row | row@(name:_) = tail $ odbc::lsql db "select * from pet" [];  
>           name == "Claws"];  
["Claws", "Gwen", "cat", "m", "1994-03-17", odbc::SQLNULL]
```

On the other hand, `lsql/lmsql` will usually be somewhat slower than `sql/msql` if the entire result set is being processed. So you should always consider the time/space tradeoffs when deciding which functions to use in a given situation.

Also note that when using `lsql/lmsql`, the query remains in progress as long as the result list is still being processed. (This is different from `sql/msql` which load the complete result set(s) at once after which the query is terminated immediately.) Since only one query can be executed per database connection, this means that only one lazy result set can be processed per database connection at any time. However, as with the lowlevel operations it is possible to do several lazy queries simultaneously if you assign them to different database connections.

8 Error Handling

When one of the above operations fails because the SQL server reports an error, an error term of the form `odbc::error msg state` will be returned, which specifies an error message and the corresponding SQL state (i.e., error code). A detailed explanation of the state codes can be found in the ODBC documentation. For instance, a reference to a non-existent table will cause a report like the following:

```
> odbc::sql db "select * from pets" [];  
odbc::error "[TCX][MyODBC]Table 'test.pets' doesn't exist" "S1000"
```

You can check for such return values and take some appropriate action. By redefining `odbc::error` accordingly, you can also have it generate exceptions or print an error message. For instance:

```
odbc::error msg state = fprintf stderr "%s (%s)\n" (msg,state) $$ ();
```

NOTE: When redefining `odbc::error` in this manner, you should be aware that the return value of `odbc::error` is what will be returned by the other operations of this module in case of an error condition. These return values are checked by other functions such as `sql`. Thus the return value should still indicate that an error has happened, and not be something that might be interpreted as a legal return value, such as an integer or a nonempty tuple. It is usually safe to have `odbc::error` return an empty tuple or throw an exception, but other types of return values should be avoided.

9 Caveats and Bugs

Be warned that multiple result sets are not supported by all databases. I also found that some ODBC drivers do not properly implement this feature, even though the database supports

it. So you better stay away from this if you want your application to be portable. You can easily implement batched queries using a sequence of single queries instead.

Note that since the exact numeric SQL data types (NUMERIC, DECIMAL) are mapped to Pure double values (which are double precision floating point numbers), there might be a loss of precision in extreme cases. If this is a problem you should explicitly convert these values to strings in your query, which can be done using the SQL CAST function, as in `select cast(1234.56 as char)`.

10 Further Information and Examples

For further details about the operations provided by this module please see the `odbc.pure` file. A sample script illustrating the usage of the module can be found in the `examples` directory.

Pure-Sql3

Version 0.1, November 14, 2010

Peter Summerland

Albert Graef <Dr.Graef@t-online.de>

This document describes **Sql3**, a [SQLite](#) module for the [Pure](#) programming language.

1 Introduction

SQLite is a software library that implements an easy to use, self-contained, serverless, zero-configuration, transactional SQL database engine. SQLite is not intended to be an enterprise database engine like Oracle or PostgreSQL. Instead, SQLite strives to be small, fast, reliable, and above all simple. See [Appropriate Uses For SQLite](#).

Sql3 is a wrapper around SQLite's C interface that provides Pure programmers access to almost all of SQLite's features, including many that are not available through Pure's generic ODBC interface.

1.1 Simple Example

Here is a simple example that opens a database file "readme.db" (creating it if it does not exist), adds a table "RM", populates "RM" and executes a query.

```
pure-sql3$> pure -q
>

> using sql3; using namespace sql3;

> let dbp = open "readme.db";

> exec dbp "create table if not exists RM (name text, age integer)";

> exec dbp "delete from RM";
```

```
> let sp1 = prep dbp "ci" "insert into RM values (?,?)";  
  
> exec sp1 ("Sam",20);  
  
> exec sp1 ("Fred",22);  
  
> let sp2 = prep dbp "ci:i" "select * from RM where age > ?";  
  
> exec sp2 18;  
[["Sam",20],["Fred",22]]
```

The Sql3 functions, `open`, `prep` and `exec` encapsulate the core functionality of SQLite, and in many cases are all you need to use SQLite effectively.

1.2 More Examples

The examples subdirectory of pure-Sql3 contains several files that further illustrate basic usage as well as some of Sql3's more sophisticated features. These include `readme.pure`, a short file that contains the examples included herein.

1.3 SQLite Documentation

SQLite's home page provides excellent documentation regarding its SQL dialect as well as its C interface. Comments in this document regarding SQLite are not meant to be a substitute for the actual documentation and should not be relied upon, other than as general observations which may or may not be accurate. The best way to use Sql3 is to get familiar with SQLite and its C interface and go directly to the [SQLite Site Map](#) for authoritative answers to any specific questions that you might have.

In the rest of this document, it is assumed the reader has some familiarity with SQLite and has read [An Introduction To The SQLite C/C++ Interface](#).

1.4 Sqlite3 - The SQLite Command-Line Utility

The SQLite library includes a really nice command-line utility named `sqlite3` (or `sqlite3.exe` on windows) that allows the user to manually enter and execute SQL statements against a SQLite database (and much more).

This tool is an invaluable aid when working with SQLite in general and with Sql3 in the Pure interpreter in particular. For example, after entering the Pure statements from the Simple Example above, you could start a new terminal, `cd` to `pure-sql3`, type `"sqlite3 readme.db"` at the prompt, and see the effect the Pure statements had on the database:

```
pure-sql3$> sqlite3 readme.db  
SQLite version 3.6.16  
Enter ".help" for instructions
```

Enter SQL statements terminated with a ";"

```
sqlite> select * from RM;  
Sam|20  
Fred|22
```

For bottom up REPL development, sqlite3 and Pure are an excellent combination.

2 Copying

Copyright (c) 2010 by Peter Summerland <p.summerland@gmail.com>.

Copyright (c) 2010 by Albert Graef <Dr.Graef@t-online.de>.

All rights reserved.

Sql3 is free software: you can redistribute it and/or modify it under the terms of the New BSD License, often referred to as the 3 clause BSD license. Sql3 is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

Please see the COPYING file for the actual license applicable to Sql3.

3 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-sql3-0.1.tar.gz>.

Unless you already have them on your machine, download SQLite and sqlite3 from the SQLite website and install as indicated. To install Sql3, cd to the pure-sql3 directory, run make, and then run `sudo make install` (on Linux).

4 Data Structure

As far as the SQLite C interface is concerned, the most important “data structures” are `sqlite3*` and `sqlite3_stmt*`, which are pointers to the following two SQLite objects.

- Database connection object - `sqlite3`
- Prepared statement object - `sqlite3_stmt`

The Sql3 functions `open` and `prep` return “sentry-guarded” pointers to database connection objects and prepared statement objects, respectively. The pointers returned by these two functions are referred to herein as `db_ptr`s and `stmt_ptr`s.

A `db_ptr` or a `stmt_ptr` is, as far as SQLite is concerned, the same as a `sqlite3*` returned by `sqlite3_open_v2` or a `sqlite3_stmt*` returned by `sqlite3_prepare_v2`, the SQLite C interface

functions corresponding to open and prep. Thus, `db_ptr`s and `stmt_ptr`s can be used when calling functions provided by the SQLite C interface (See [Accessing the Rest of SQLite's C Interface](#)).

5 Core Database Operations

The core database operations are (a) opening and closing database connections and (b) preparing, executing and closing prepared statements.

5.1 Database Connections

Generally speaking, the first step in accessing a database is to obtain a `db_ptr` that references a database connection object. Once the `db_ptr` is obtained, it can be used to construct prepared statements for updating and querying the underlying database. The last step is usually to close the database connection (although this is often done automatically by `Sql3` when the `db_ptr` goes out of scope).

Opening a Database Connection

In `Sql3` `open` constructs a database connection and returns a `db_ptr` that refers to the connection.

`open (file_path::string [,access_mode::int,custom_bindings])` opens a SQLite database file whose name is given by the `file_path` argument and returns a `db_ptr` for the associated database connection object created by SQLite.

Example:

```
> let dbp2 = open "abc.db"; dbp2;  
#<pointer 0x992dff8>
```

If the filename is `":memory:"` a private, temporary in-memory database is created for the connection.

The basic access modes are:

- `SQLITE_OPEN_READONLY` - the database is opened in read-only mode. If the database does not already exist, an error is returned.
- `SQLITE_OPEN_READWRITE` - the database is opened for reading and writing if possible, or reading only if the file is write protected by the operating system. In either case the database must already exist, otherwise an error is returned.
- `SQLITE_OPEN_READWRITE | SQLITE_OPEN_CREATE` - the database is opened for reading and writing, and is created if it does not already exist. This is the default value that is used if the flags argument is omitted.

- `SQLITE_OPEN` - an alias for `SQLITE_OPEN_READWRITE` | `SQLITE_OPEN_CREATE` provided by `Sql3`.

These flags can be combined with `SQLITE_OPEN_NOMUTEX`, `SQLITE_OPEN_FULLMUTEX`, `SQLITE_OPEN_SHARED_CACHE`, `SQLITE_OPEN_PRIVATE_CACHE` to control SQLite's threading and shared cache features. All of these flags are exported by `Sql3`.

The optional `custom_bindings` argument allows the user to set up customized binding and fetching behavior for prepared statements associated with the returned `db_ptr`. (See [Custom Binding Types for Prepared Statements](#))

Failure to Open a Database Connection

If SQLite cannot open the connection, it still returns a pointer to a database connection object that must be closed. In this case, `open` automatically closes the connection object and then throws an exception. E.g.:

```
> catch error (open ("RM_zyx.db",SQLITE_OPEN_READONLY));
error (sql3::db_error 14 "unable to open database file [open RM_zyx.db]")
```

Apparently, SQLite does not verify that a file is a valid SQLite database when it opens a connection. However, if the file is corrupted SQLite will return an error when the connection is used.

Testing a `db_ptr`

For argument checking and other purposes `Sql3` provides the following functions:

`is_db_ptr ptr::pointer` returns 1 if `ptr` is a `db_ptr` returned by `open`, and 0 if it is not.

`is_open ptr::pointer` returns 1 if the database connection referenced by `ptr` is open.

Closing a Database Connection

When a database connection object is no longer needed, it should be closed so that SQLite can free the associated resources.

`close dbp::pointer` if the database connection referenced by the `db_ptr` `dbp` is open, close it using `sqlite3_close`; otherwise do nothing.

Before calling `sqlite3_close`, `close` finalizes all prepared statements associated with the connection being closed. `Sql3` will detect and throw a `db_error` if an attempt is subsequently made to execute a statement associated with the closed database connection.

```
> let dbp2_sp = prep dbp2 "ci:" "select * from RM";

> exec dbp2_sp ();
[["Sam",20],["Fred",22]]
```

```
> close dbp2;

> catch error (exec dbp2_sp);
error (sql3::db_error 0 "Attempt to exec on a closed db_ptr.")
```

If a `db_ptr` goes out of scope, `Sql3` will automatically call `sqlite3_close` to close the referenced database connection, but only if the connection has not already been closed by `close`.

When debugging, this activity can be observed by editing `sql3.pure`, changing “`const SHOW_OPEN_CLOSE = 0;`” to “`const SHOW_OPEN_CLOSE = 1;`” and running `sudo make install` in the `pure-sql3` directory. This will cause a message to be printed whenever a `db_ptr` or `stmt_ptr` is created or finalized.

N.B. You should never call `sqlite3_close` with a `db_ptr`. If the referenced database connection is closed by such a call, a subsequent call to `close` on this `db_ptr` (including the call that will automatically occur when the `db_ptr` goes out of scope) will cause a seg fault.

5.2 Prepared Statements

In SQLite, prepared statement objects are used to execute SQL statements using the following SQLite C interface functions.

- `sqlite3_prepare_v2`
- `sqlite3_bind`
- `sqlite3_step`
- `sqlite3_column`
- `sqlite3_finalize`

Using the C interface, the basic procedure is to prepare a statement using `sqlite3_prepare_v2`, bind its parameters using `sqlite3_bind`, step it using `sqlite3_step` one or more times until it is done and then finalize it using `sqlite3_finalize`. Each time `sqlite3_step` returns `SQLITE_ROW`, use `sqlite3_column` to fetch the row’s values. Here `sqlite3_bind` and `sqlite3_column` represent families of bind and column functions, rather than actual functions, with one member for each of the basic data types recognized by SQLite. Thus, for example, `sqlite_bind_double` is the function one would use to bind a prepared statement with an argument of type `double`.

`Sql3` encapsulates these procedures in four functions: `prep`, `exec`, `lexec` and `finalize`.

Constructing Prepared Statements

In `Sql3` you can use `prep` to construct a prepared statement and obtain a `stmt_ptr` that refers to it.

`prep dbp::pointer binding_string::string sql_statement::string` constructs a prepared statement object and returns a `stmt_ptr` that references it. `dbp` must be a `db_ptr`

or the rule will not match. `sql_statement` is the SQL statement that will be executed when the prepared statement is passed to `exec`.

Basically, `prep` just passes `dbp` and `sql_statement` on to `sqlite3_prepare_v2` and returns a sentry guarded version of the `sqlite3_stmt*` it receives back from `sqlite3_prepare_v2`. SQL statements passed to `prep` (and `sqlite3_prepare_v2`) can have argument placeholders, indicated by `"?"`, `"?nnn"`, `":AAA"`, etc, in which case the argument placeholders must be bound to values using `sqlite_bind` before the prepared statement is passed to `sqlite3_step`. Hence the `binding_string`, which is used by `Sql3` to determine how to bind the prepared statement's argument placeholders, if any. The binding string also tells `Sql3` how to fetch values in the `sqlite3_column` phase of the basic prepare, bind, step, fetch, finalize cycle dictated by the SQLite C interface.

In the following two examples, the `"c"` and `"i"` in the binding strings indicate that (a) a string and an int will be used to bind `sp1`, (b) an int will be used to bind `sp2` and (c) `sp2`, when executed, will return a result set in the form of a list of sublists each of which contains a string and an int.

```
> let sp1 = prep dbp "ci" "insert into RM values (?,?)";

> let sp2 = prep dbp "ci:i" "select * from RM where age > ?";
```

In general, the characters in the type string before the `":"`, if any, indicate the types in the result set. Those that occur after the `":"`, if any, indicate the types of the arguments used to bind the prepared statement object. If the type string does not contain a `":"`, the characters in the type string, if any, are the types of binding arguments.

`Sql3` provides the following set of “core” binding types:

Type	Pure Argument	SQLite Type
b	(int, pointer)	blob
c	string	text (utf8)
d	double	float
i	int	int
k	int or bigint	int64
l	bigint	blob
n	Sql3::SQLNULL	NULL
x	expression	blob
v	variant	variant

The `"b"` or blob type is different from the rest in that the Pure argument is specified as a pair. The first element of the pair indicates the length in bytes of the object to be stored and the second element indicates its location in memory. The `"c"` type stands for string (as in `"char*"`), `"d"` stands for double and `"i"` stands for int. The `"k"` type stands for “key” and maps Pure ints and bigints (within the range of int64) to int64 values in the database. This type is useful when dealing with SQLite’s “integer primary keys” and “rowids” both of which are int64. The `"l"` type, in contrast applies to all bigints (and not to ints) and it maps bigints onto blobs, which are generally meaningless in SQL math expressions. The `"n"` type can only appear on the binding side of a type string. The `"v"` type stands for any of `"b"`, `"c"`, `"d"`, `"i"` or `"n"`, based on the type of the binding argument. A `"v"` type will be fetched from

SQLite according to the native SQLite column type of the corresponding column. The “x” type is used to store and reconstruct Pure expressions as binary objects, using the `val` and `blob` functions provided by the Pure system module.

Users can define custom binding types and pass them as a third parameter to `open`. The resulting `db_ptr` can be used with the custom binding types to construct prepared statements using `prep`.

Testing a `stmt_ptr`

For argument checking and other purposes it is sometimes useful to determine if a given expression is a valid `stmt_ptr`.

`is_stmt_ptr ptr::pointer` returns 1 if `ptr` is a `stmt_ptr`, otherwise returns 0.

Executing Prepared Statements

In `Sql3`, the `bind`, `step`, `column`, `step`, `column` ... cycle is encapsulated in the `exec` and `lexec` functions.

`exec sp::pointer args` use `args` to bind the prepared statement referenced by `sp`, execute it and return the result set as a list. The first parameter, `sp` must be a valid `stmt_ptr` or the rule will fail.

The second parameter, `args`, is a tuple or list of arguments whose number and type correspond to the `bind` parameter types specified in the call to `prep` that produced the first parameter `sp`.

Example:

```
> exec sp1 ("Tom",30); //insert Tom
[]

> exec sp2 19;          //select age > 19
[["Sam",20],["Fred",22],["Tom",30]]

> exec dbp "delete from RM where name='Tom';
```

An error is thrown if the `args` do not correspond to the specified types.

```
> catch error (exec sp2 "a");
error (sql3::db_error 0 "\"a\" does not have type int")
```

If a prepared statement does not have any binding parameters, the call to `exec` should use `()` as the binding argument.

```
> let sp3 = prep dbp "c:" "select name from RM";

> exec sp3 ();
[["Sam"],["Fred"]]
```

Extra care is required when executing prepared statements that take a blob argument because it must be a pair. In order to preserve the tuple as a pair, binding arguments that include a blob should be passed to `exec` as a list. If passed as a member of a larger tuple, it will be treated as two arguments due to the nature of tuples.

```
> let blb = (100,ptr);
```

```
> (a,blb,c);  
a,100,ptr,c
```

```
> [a,blb,c];  
[a,(100,ptr),c]
```

Thus something like `"exec stpx [a,blb,c]"` would work fine, while `"exec stpx (a,blb,c)"` would produce a `Sql3` binding exception.

Executing Lazily

The `exec` function returns result sets as eager lists which can sometimes be inefficient or simply not feasible for large result sets. In such cases it is preferable to use `lexec` instead of `exec`.

`lexec stmp::pointer args` same as `exec` except that it returns a lazy list.

Example:

```
> lexec sp2 19;  
["Sam",20]:#<thunk 0xb6475ab0>
```

Note that no changes to `sp2` were required. In addition, for most purposes the lazy list returned by `lexec` can be processed by the same code that processed the eager list returned by `exec`.

Executing Directly on a `db_ptr`

For statements that have no parameters and which do not return results, `exec` can be applied to a `db_ptr`.

`exec dbp::pointer sql_statement::string` constructs a temporary prepared statement using `sql_statement`. The SQL statement cannot contain argument placeholders and cannot be a select statement.

Example:

```
> exec dbp "create table if not exists RM (name varchar, age integer)";
```

Executing Against a Busy Database

SQLite allows multiple processes to concurrently read a single database, but when any process wants to write, it locks the entire database file for the duration of its update.

When `sqlite3_step` tries to access a file that is locked by another process, it treats the database as “busy” and returns the `SQLITE_BUSY` error code. If this happens in a call to `exec` or `lexec`, a `db_busy` exception will be thrown.

You can adjust SQLite’s behavior using `sqlite3_busy_handler` or `sqlite3_busy_timeout`.

If the statement is a `COMMIT` or occurs outside of an explicit transaction, then you can retry the statement. If the statement is not a `COMMIT` and occurs within a explicit transaction then you should rollback the transaction before continuing.

Grouping Execution with Transactions

No changes can be made to a SQLite database file except within a transaction. Transactions can be started manually by executing a `BEGIN` statement. Manually started transactions persist until the next `COMMIT` or `ROLLBACK` statement is executed. Transactions are also ended if an error occurs before the transaction is manually ended using a `COMMIT` or `ROLLBACK` statement. This behavior provides the means make a series of changes “atomically.”

By default, SQLite operates in autocommit mode. In autocommit mode, any SQL statement that changes the database (basically, anything other than `SELECT`) will automatically start a transaction if one is not already in effect. As opposed to manually started transactions, automatically started transactions are committed as soon as the execution of the related statement completes.

The upshot of this, in `Sql3` terms, is that unless a transaction is started manually, the database will be updated each time `exec` is called. For a long series of updates or inserts this can be very slow. The way to avoid this problem is to manually begin and end transactions manually.

Transactions created using `BEGIN...COMMIT` do not nest. For nested transactions, use the `SAVEPOINT` and `RELEASE` statements.

`Sql3` provides the following convenience functions all of which simply call `exec` with the appropriate statement. For example `begin dbp` is exactly the same as `exec dbp "BEGIN"`.

- `begin dbp::pointer`
- `begin_exclusive dbp::pointer`
- `begin_immediate dbp::pointer`
- `commit dbp::pointer`
- `rollback dbp::pointer`
- `savepoint dbp::pointer save_point::string`
- `release dbp::pointer save_point::string`

- `rollback_to dbp::pointer save_point::string`

Finalizing Prepared Statements

When a prepared statement is no longer needed it should be finalized so that SQLite can free the associated resources.

finalize sp::pointer finalize the prepared statement referenced by `sp`, which must be a `stmt_ptr` previously returned by `prep`.

Often there is no need to call `finalize` for a given `stmt_ptr` because it will be automatically called when the `stmt_ptr` goes out of scope.

If the `stmt_ptr` is associated with a database connection that has been closed (which would have caused an exception to be thrown), an attempt to finalize it, including the automatic finalization can occur when `stmt_ptr` goes out of scope, will cause an exception to be thrown.

```
> catch error (finalize dbp2_sp);  
error (sql3::db_error 0 "finalize: STMT attached to a closed db_ptr.")
```

Multiple calls to `finalize` are fine. In contrast, the corresponding native C interface function, `sqlite3_finalize` will cause a seg fault if called with a pointer to a finalized prepared statement object. This is the main reason why you should never call `sqlite3_finalize` with a `stmt_ptr`. If the prepared statement referenced by the `stmt_ptr` is finalized by such a call, a subsequent call to `finalize` with the `stmt_ptr` (including the call that will automatically occur when the `stmt_ptr` goes out of scope) will cause a seg fault.

5.3 Exceptions

Sql3 throws two types of exceptions, one for outright errors and one for database “busy” conditions.

db_error

When a Sql3 function detects an error it throws an exception of the form “`db_error ec msg`” where `ec` is an error code and `msg` is the corresponding error message. If `ec>0`, the error was detected by SQLite itself, and `ec` and `msg` are those returned by SQLite. If `ec==0`, the error was detected by Sql3 and `msg` is a Sql3 specific description of the error. E.g.,

```
> db_error_handler (db_error ec msg) = ()  
> when  
>   source = if ec > 0 then "SQLite" else "Sql3";  
>   printf "%s db_error: ec %d, %s\n" (source,ec,msg);  
> end;  
> db_error_handler x = throw x;  
  
> catch db_error_handler (exec dbp "select * from NO_TABLE");  
SQLite db_error: ec 1, no such table: NO_TABLE
```

db_busy

Sql3 functions `exec` and `lexec` throw exceptions of the form “`db_busy dbp`”, where `dbp` is a `db_ptr`, if they are prevented from executing successfully because the database referenced by `dbp` is locked (See [Executing Against a Busy Database](#)).

SQLite Error Codes

Here is a list, as of April 2, 2010, SQLite’s error codes.

SQLite_ERROR	1	<i>/* SQL error or missing database */</i>
SQLite_INTERNAL	2	<i>/* Internal logic error in SQLite */</i>
SQLite_PERM	3	<i>/* Access permission denied */</i>
SQLite_ABORT	4	<i>/* Callback routine requested an abort */</i>
SQLite_BUSY	5	<i>/* The database file is locked */</i>
SQLite_LOCKED	6	<i>/* A table in the database is locked */</i>
SQLite_NOMEM	7	<i>/* A malloc() failed */</i>
SQLite_READONLY	8	<i>/* Attempt to write a readonly database */</i>
SQLite_INTERRUPT	9	<i>/* Operation terminated by sqlite3_interrupt()*/</i>
SQLite_IOERR	10	<i>/* Some kind of disk I/O error occurred */</i>
SQLite_CORRUPT	11	<i>/* The database disk image is malformed */</i>
SQLite_NOTFOUND	12	<i>/* NOT USED. Table or record not found */</i>
SQLite_FULL	13	<i>/* Insertion failed because database is full */</i>
SQLite_CANTOPEN	14	<i>/* Unable to open the database file */</i>
SQLite_PROTOCOL	15	<i>/* NOT USED. Database lock protocol error */</i>
SQLite_EMPTY	16	<i>/* Database is empty */</i>
SQLite_SCHEMA	17	<i>/* The database schema changed */</i>
SQLite_TOOBIG	18	<i>/* String or BLOB exceeds size limit */</i>
SQLite_CONSTRAINT	19	<i>/* Abort due to constraint violation */</i>
SQLite_MISMATCH	20	<i>/* Data type mismatch */</i>
SQLite_MISUSE	21	<i>/* Library used incorrectly */</i>
SQLite_NOLFS	22	<i>/* Uses OS features not supported on host */</i>
SQLite_AUTH	23	<i>/* Authorization denied */</i>
SQLite_FORMAT	24	<i>/* Auxiliary database format error */</i>
SQLite_RANGE	25	<i>/* 2nd parameter to sqlite3_bind out of range */</i>
SQLite_NOTADB	26	<i>/* File opened that is not a database file */</i>

New error codes may be added in future versions of SQLite. Note that the SQLite names of the error codes are not exported by the Sql3 module.

6 Advanced Features

Sql3’s advanced features include the ability to implement SQL functions in Pure, convenient access to the SQLite C interface and custom binding types.

6.1 Custom SQL Functions

An extremely powerful (albeit complex) feature of the SQLite C interface is the ability to add new SQL scalar or aggregate functions. The new functions can be used in SQL statements the same way as SQLite's prepackaged functions. Sql3 hides the complexity and seamlessly integrates all of this functionality, :, into Pure via `create_function`. This function is used to register both scalar SQL functions and aggregate SQL functions with SQLite.

Scalar SQL Functions

You can add a custom SQL scalar function to SQLite by passing a single Pure function to `create_function`.

`create_function dbp::pointer name::string nargs::int pure_fun` registers a new SQL scalar function of `nargs` arguments that can be called, as `name`, in SQL statements prepared with respect to `dbp`, a `db_ptr`. When the SQL function is called in a SQL statement, control is passed to `pure_fun`, a curried function written in Pure. If `nargs` is `(-1)`, the SQL function name is variadic, and the arguments will be passed to `pure_fun` as a single list.

Note that `create_function` can also register aggregate functions (see [Aggregate SQL Functions](#)).

Here is an example of a scalar function that takes two parameters. Note that any kind of Pure "function" can be passed here; local functions, global functions, lambdas or partial applications all work.

```
> create_function dbp "p_fn" 2 plus with plus x y = x + y; end;

> let sp4 = prep dbp "cii:"
>      "select p_fn('Hi ',name), age, p_fn(age,10) from RM";

> exec sp4 ();
[["Hi Sam",20,30],["Hi Fred",22,32]]
```

Here is an example of a variadic function:

```
> create_function dbp "p_qm" (-1) quasimodo with
>   quasimodo xs = "quasimodo: "+join ":" [str x | x<xs];
> end;
```

If the SQL function takes no arguments, the corresponding Pure function must, for technical reasons in Pure, take a single dummy argument. E.g.,

```
> create_function dbp "p_count" 0 counter with
>   counter () = put r (get r+1);
> end when r = ref 0 end;
```

Here is how `count` and `quasimodo` might be used:

```
> let sp5 = prep dbp "ic:" "select p_count(), p_qm(name,age) from RM";

> exec sp5 ();
[[1,"quasimodo: \"Sam\":20\"],[2,"quasimodo: \"Fred\":22"]]

> exec sp5 ();
[[3,"quasimodo: \"Sam\":20\"],[4,"quasimodo: \"Fred\":22"]]
```

Multiple SQL functions can be registered with the same name if they have differing numbers of arguments. Built-in SQL functions may be overloaded or replaced by new application-defined functions.

Generally, a custom function is permitted to call other Sql3 and native SQLite C interface functions. However, such calls must not close the database connection nor finalize or reset the prepared statement in which the function is running.

Aggregate SQL Functions

You can use `create_function` to register an aggregate SQL function with SQLite by passing a triple consisting of two Pure functions and a start value, in lieu of a single Pure function.

create_function dbp::pointer name::string nargs::int (step,final,start)
registers a new SQL aggregate function of `nargs` arguments that can be called, as `name` in SQL statements prepared with respect to `dbp`, a `db_ptr`. `step` and `final` are curried Pure functions and `start` is the initial value for the aggregation. The `step` function is called repeatedly to accumulate values from the database, starting from the given `start` value, and finally the `final` function is applied to the accumulated result.

Note that for a single-argument `step` function, this works exactly as if the functions were invoked as “`final (foldl step start values)`”, where `values` is the list of aggregated values from the database.

```
> create_function dbp "p_avg" 1 (step,final,(0,0.0)) with
>   step (n,a) x = n+1, a+x;
>   final (n,a) = a/n;
> end;

> let sp6 = prep dbp "id:" "select count(name), p_avg(age) from RM";

> exec sp6 ();
[[2,21.0]]
```

6.2 Accessing the Rest of SQLite's C Interface

The `db_ptrs` returned by `open` and `stmt_ptrs` returned by `prep` are sentry guarded versions of the actual pointers to the data base connection objects and prepared statement objects returned by their corresponding native C interface functions `sqlite3_open_v2` and `sqlite3_prepare_v2`. This makes it easy to call almost any external function in SQLite's C

interface directly, passing it the same `db_ptr` or `stmt_ptr` that is passed to Sql3's functions, such as `prep` or `exec`.

For example, you can override SQLite's default behavior with respect to a busy database as follows:

```
> extern int sqlite3_busy_timeout(sqlite3*, int);

> sqlite3_busy_timeout dbp 10;
```

This sets a busy handler that will “sleep and retry” multiple times until at least 10 milliseconds of sleeping have accumulated. Calling this routine with an argument less than or equal to zero turns off all busy handlers.

Another example is to query the number of database rows that were changed, inserted or deleted by the most recently completed SQL statement executed on a given database connection:

```
> extern int sqlite3_changes(sqlite3*);

> exec sp1 ("Harvey",30);

> sqlite3_changes dbp;
1
```

As a final example, in this case using a `stmt_ptr`, you can determine name assigned to a column in a result using `sqlite3_column_name`:

```
> extern char *sqlite3_column_name(sqlite3_stmt*, int);

> exec sp2 29;
[["Harvey",30]]

> sqlite3_column_name sp2 1;
"age"
```

In order to call a native C function you must first make it accessible using an extern statement.

Please note also that directly calling a function provided by the SQLite C interface can be dangerous, as is the case with any call from Pure code to an external C function. Sql3 users should be especially careful in this regard because using a `db_ptr` or a `stmt_ptr` in calls to certain native C interface functions, including in particular `sqlite3_close` and `sqlite3_finalize`, will corrupt data held by the `db_ptr` or `stmt_ptr`, leading to undefined behavior. The reason for this restriction is that Sql3 uses sentries to insure that the resources associated with a `db_ptr` or a `stmt_ptr` are automatically finalized by SQLite when they go out of scope. In addition, the sentries carry internal information used by Sql3 for other purposes.

6.3 Custom Binding Types for Prepared Statements

You can add your own binding types for preparing and executing prepared statements by specifying a third argument to `open`. The third argument must be a list of “hash rocket pairs” in which the left side is a character for the custom binding type and the right side is a list with three members. The second and third members of the list are functions that map objects from the new type to one of the Sql3 core types and back, respectively. The first member of the list is the character for the Sql3 core types referenced by the mapping functions.

The file `sql3_user_bind_types.pure` in the `examples` subdirectory shows how this might be done for a couple of user defined types. The example script deals with dates and certain Pure expressions as bigints and native Pure expressions, while the database stores these as utf-8 text. The following snippets show parts of the script that are relevant to this discussion:

```
const custom_binds = [
  "t"=>["c",day_to_str,str_to_day],
  "s"=>["c",str,eval]
];

d1 = str_to_day "2010-03-22";

db = open ("abc.db", SQLITE_OPEN, custom_binds);
stm1 = prep db "cts" "insert into TC values(?,?,?)";
exec stm1 ["Manny", d1, s_expr];
stm3a = sql3::prep db "t:" "select t_date from TC";
stm3b = sql3::prep db "c:" "select t_date from TC";
```

Executing `stm3a` and `stm3b` from the interpreter shows that TC’s date field is stored as a string, but returned to the Pure script as a bigint.

```
> sql3::exec stm3a ();
[[146911]]

> sql3::exec stm3b ();
[["2010-03-22"]]
```

The character designating the custom type must not be one of the letters used to designate Sql3 core binding types.

7 Threading Modes

SQLite supports three different threading modes:

1. Single-thread. In this mode, all mutexes are disabled and SQLite is unsafe to use in more than a single thread at once.
2. Multi-thread. In this mode, SQLite can be safely used by multiple threads provided that no single database connection is used simultaneously in two or more threads.

3. Serialized. In serialized mode, SQLite can be safely used by multiple threads with no restriction.

SQLite can be compiled with or without support for multithreading and the default is to support it.

In many cases, single-thread mode might be appropriate if only because it is measurably faster. This might be the case, for example, if you are using SQLite as the on-disk file format for a desktop application.

If your version of SQLite was compiled with support for multithreading, you can switch to single-thread mode at runtime by calling `sqlite3_config()` with the verb `SQLITE_CONFIG_SINGLETHREAD`.

If you must use threads, it is anticipated that Sql3 probably will not impose an additional burden. Hopefully, you will be fine if you apply the same precautions to a `db_ptr` or `stmt_ptr` that you would apply to the underlying `sqlite*` and `sqlite_stmt*`s if you were not using Sql3. It is strongly advised however that you look at the underlying Sql3 code to verify that this will work. Since everything that is imposed between the raw pointers returned by the SQLite interface and the corresponding `db_ptr` and `stmt_ptr`s is written in Pure, it should be relatively easy to determine how Sql3 and your multithreading strategy will interact. See [Is SQLite threadsafe?](#) , [Opening A New Database Connection](#) and [Test To See If The Library Is Threadsafe](#).

Pure-XML - XML/XSLT interface

Version 0.4, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

XML, the Extensible Markup Language, facilitates the exchange of complex structured data between applications and systems. **XSLT** allows you to transform XML documents to other XML-based formats such as HTML. Together, XML and XSLT let you create dynamic web content with ease. Both XML and XSLT are open standards by the W3C consortium (<http://www.w3.org>).

Pure's XML interface is based on the libxml2 and libxslt libraries from the GNOME project. If you have a Linux system then you most likely have these libraries, otherwise you can get them from <http://xmlsoft.org>. For Windows users, the required dlls are available from the GnuWin32 project (<http://gnuwin32.sourceforge.net>) and are already included in the Pure MSI package.

1 Copying

Copyright (c) 2009 by Albert Graef <Dr.Graef@t-online.de>.

pure-xml is free software: you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

pure-xml is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

2 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-xml-0.4.tar.gz>.

Run `make` and then `sudo make install` to compile and install this module. This requires `libxml2`, `libxslt` and Pure.

3 Usage

Use the following declaration to make the operations of this module available in your programs:

```
using xml;
```

The module defines two namespaces `xml` and `xslt` for the XML and the XSLT operations, respectively. For convenience, you can open these in your program as follows:

```
using namespace xml, xslt;
```

A number of complete examples illustrating the use of this module can be found in the `examples` directory in the source distribution.

4 Data Structure

This module represents XML documents using pointers to the `xmlDoc` and `xmlNode` structures provided by the `libxml2` library. Similarly, stylesheets are simply pointers to the `xmlStylesheet` structure from `libxslt` (cf. [Transformations](#)). This makes it possible to use these objects directly with the operations of the `libxml2` and `libxslt` libraries (via Pure’s C interface) if necessary. Note, however, that these are all “cooked” pointers which take care of freeing themselves automatically when they are no longer needed, therefore you shouldn’t free them manually.

You can also check for these types of pointers using the following predicates:

- `xml::docp x` checks whether `x` is an XML document pointer.
- `xml::nodep x` checks whether `x` is a pointer to a node in an XML document.
- `xslt::stylesheetp x` checks whether `x` is an XSLT stylesheet pointer.

4.1 The Document Tree

An XML document is a rooted tree which can be created, traversed and manipulated using the operations of this module. There are different types of nodes in the tree, each carrying their own type of data. In Pure land, the node data is described using the following “node info” constructors.

- `xml::element tag ns attrs`

An XML element with given (possibly qualified) name `tag`, a (possibly empty) list of namespace declarations `ns` and a (possibly empty) list of attributes `attrs`. Namespace

declarations normally take the form of a pair of strings (`prefix`, `href`), where `prefix` is the prefix associated with the namespace and `href` the corresponding URI (the name of the namespace), but they can also be just a string `href` if the namespace prefix is missing. Attributes are encoded as `key=>value` pairs, where `key` is the attribute name and `value` the associated value; both `key` and `value` are strings.

- `xml::element_text tag ns attrs content`

A convenience function which denotes a combination of an element node with a text child. This is only used when creating a new node, and indicates that a text node child is to be added to the node automatically.

- `xml::attr key val`

An attribute node. These only occur as results of the `select` and `attrs` functions, and cannot be inserted directly into a document.

- `xml::text content`

A text node with the given content (a string).

- `xml::cdata content`

Like `xml::text`, but contains unparsed character data.

- `xml::comment content`

A comment.

- `xml::entity_ref name`

An entity reference (`&name;`).

- `xml::pi name content`

Processing instructions. `name` is the application name, `content` the text of the processing instructions.

4.2 Document Types

Besides the node types described above, there are some additional node types used in the `document type definition` (DTD), which can be extracted from a document using the `int_subset` and `ext_subset` functions. These are for inspection purposes only; it is not possible to change the DTD of a document in-place. (However, you can create a new document and attach a DTD to it, using the `new_doc` function.)

- `xml::doctype name extid`

DTDs are represented using this special type of node, where `name` is the name of the root element, and `extid` is a pair consisting of the external identifier and the URI of the DTD (or just the URI if there is no external identifier). The `xml::doctype` node has as its children zero or more of the following kinds of DTD declaration nodes (these are just straightforward abstract syntax for the `!ELEMENT`, `!ATTLIST` and `!ENTITY` declarations inside a DTD declaration; see the XML specification for details).

Element declarations: Here, `name` is the element tag and `content` the definition of the element structure, see [element content](#) below. XML supports various kinds of element types, please refer to [document type definition](#) in the XML specification for details.

- `xml::undefined_element name`
An undefined element. This is in libxml2, but not in the XML specification, you shouldn't see this in normal operation.
- `xml::empty_element name`
An element without any content.
- `xml::any_element name`
An element with unrestricted content.
- `xml::mixed_element name content`
A “mixed” element which can contain character data, optionally interspersed with child elements, as given in the content specification.
- `xml::std_element name content`
A standard element consisting *only* of child elements, as given in the content specification.

Attribute declarations: These are used to declare the attributes of an element. `elem_name` is the name of an element which describes the attribute type, `name` is the name of the attribute itself, and `default` specifies the default value of the attribute, see [attribute defaults](#) below. XML supports a bunch of different attribute types, please refer to [document type definition](#) in the XML specification for details.

- `xml::cdata_attr elem_name name default`
- `xml::id_attr elem_name name default`
- `xml::idref_attr elem_name name default`
- `xml::idrefs_attr elem_name name default`
- `xml::entity_attr elem_name name default`
- `xml::entities_attr elem_name name default`
- `xml::nmtoken_attr elem_name name default`
- `xml::nmtokens_attr elem_name name default`
- `xml::enum_attr elem_name name vals default`
- `xml::notation_attr elem_name name vals default`

Entity declarations: These are used for internal and external entity declarations. `name` is the entity name and `content` its definition. External entities also have an `extid` (external identifier/URI pair) identifying the entity.

- `xml::int_entity name content`

- `xml::int_param_entity` name content
- `xml::ext_entity` name extid content
- `xml::ext_param_entity` name extid content

The element content type (content argument of the [element declaration](#) nodes) is a kind of regular expression formed with tags (specified as strings) and the following constructors:

- `xml::pcdata`: text data (`#PCDATA`)
- `xml::sequence` `xs`: concatenation (`x,y,z`)
- `xml::union` `xs`: alternatives (`x|y|z`)
- `xml::opt` `x`: optional element (`x?`)
- `xml::mult` `x`: repeated element (`x*`)
- `xml::plus` `x`: non-optional repeated element (`x+`)

Attribute defaults (the default argument of [attribute declaration](#) nodes) are represented using the following constructor symbols:

- `xml::required`: a required attribute, i.e., the user must specify this
- `xml::implied`: an implied attribute, i.e., the user does not have to specify this
- `xml::default` `val`: an attribute with the given default value `val`
- `xml::fixed` `val`: an attribute with the given fixed value `val`

5 Operations

This module provides all operations necessary to create, inspect and manipulate XML documents residing either in memory or on disk. Operations for formatting XML documents using XSLT stylesheets are also available.

5.1 Document Operations

The following functions allow you to create new XML documents, load them from or save them to a file or a string, and provide general information about a document.

- `xml::new_doc` version dtd info

This function creates an XML document. It returns a pointer to the new document. `version` is a string denoting the XML version (or `"` to indicate the default). `info` is the [node info](#) of the root node (which should denote an element node). `dtd` denotes the document type which can be `()` to denote an empty DTD, a string (the URI/filename of the DTD), or a pair (`pubid`, `sysid`) where `pubid` denotes the public identifier of the DTD and `sysid` its system identifier (URI).

Note that only simple kinds of documents with an internal DTD can be created this way. Use the `load_file` or `load_string` function below to create a skeleton document if a more elaborate prolog is required.

- `xml::load_file name flags, xml::load_string s flags`

Load an XML document from a file name or a string `s`. `flags` denotes the parser flags, a bitwise disjunction of any of the following constants, or 0 for the default:

- `xml::DTDLOAD`: load DTD
- `xml::DTDVALID`: validate
- `xml::PEDANTIC`: pedantic parse
- `xml::SUBENT`: substitute entities
- `xml::NOBLANKS`: suppress blank nodes

The return value is the document pointer. These operations may also fail if there is a fatal error parsing the document.

- `xml::save_file name doc encoding compression, xml::save_string doc`

Save an XML document `doc` to a file or a string. When saving to a file, `encoding` denotes the desired encoding (or "" for the default), `compression` the desired level of zlib compression (0 means none, 9 is maximum, -1 indicates the default). Note that with `xml::save_string`, the result is always encoded as UTF-8.

- `xml::doc_info doc`

Retrieve general information about a document. Returns a tuple `(version, encoding, url, compression, standalone)`, where `version` is the XML version of the document, `encoding` the external encoding (if any), `url` the name/location of the document (if any), `compression` the level of zlib compression, and `standalone` is a flag indicating whether the document contains any external markup declarations “which affect the information passed from the XML processor to the application”, see the section on the [standalone document declaration](#) in the XML spec for details. (Apparently, in libxml2 `standalone` is either a truth value or one of the special values -1, indicating that there’s no XML declaration in the prolog, or -2, indicating that there’s an XML declaration but no `standalone` attribute.)

- `xml::int_subset doc, xml::ext_subset doc`

Retrieve the internal and external DTD subset of a document. Returns a doctype node (fails if there’s no corresponding DTD).

Example

Read the `sample.xml` document distributed with the sources (ignoring blank nodes) and retrieve the document info:

```
> using xml;
> let sample = xml::load_file "sample.xml" xml::NOBLANKS;
```

```
> xml::doc_info sample;  
"1.0", "", "sample.xml", 0, -2
```

5.2 Traversing Documents

These operations are used to traverse the document tree, i.e., the nodes of the document. They take either a document pointer `doc` or a node pointer `node` as argument, and yield a corresponding node pointer (or a document pointer, in the case of `xml::doc`). The node pointers can then be used with the [Node Information](#) and [Node Manipulation](#) operations described below.

- `xml::root doc`: the root node of `doc`
- `xml::doc node`: the document node belongs to
- `xml::parent node`: the parent of `node`
- `xml::first node`, `xml::last node`: first and last child node
- `xml::next node`, `xml::prev node`: next and previous sibling
- `xml::first_attr node`, `xml::last_attr node`: first and last attribute node

All these operations fail if the corresponding target node does not exist, or if the type of the given node is not supported by this implementation.

There are also two convenience functions to retrieve the children and attribute nodes of a node:

- `xml::children node`: returns the list of all child nodes of `node`
- `xml::attrs node`: returns the list of all attribute nodes of `node`

Moreover, given a node pointer `node`, `node!i` can be used to retrieve the *i*th child of `node`.

Example

Peek at the root node of the sample document and its children:

```
> let r = xml::root sample; r;  
#<pointer 0xe15e10>  
> xml::node_info r;  
xml::element "story" [] []  
> #xml::children r;  
5  
> xml::node_info (r!0);  
xml::cdata "<greeting>Hello, world!</greeting>"
```

5.3 Node Information

These operations retrieve information about the nodes of an XML document.

- `xml::select doc xpath,xml::select doc (xpath,ns)`

Retrieve nodes using an [XPath](#) specification. Given an XPath (a string) `xpath`, this operation returns the list of all matching nodes in the given document `doc`. You can also specify a node as the first argument, in which case the document of the given node is searched and paths are interpreted relative to the given node (rather than the root node of the document).

Moreover, instead of just an XPath you can also specify a pair `(xpath,ns)` consisting of an XPath `xpath` and a list `ns` of `prefix=>uri` string pairs which describe the namespaces to be recognized in the XPath expression. This is necessary to select nodes by qualified tag or attribute names. Note that only the namespace URIs must match up with those used in the queried document; the corresponding namespace prefixes can be chosen freely, so you can use whatever prefixes are convenient to formulate the XPath query. However, for each namespace prefix used in the XPath expression (not the document!), there *must* be a corresponding binding in the `ns` list. Otherwise the underlying libxml2 function will complain about an undefined namespace prefix and `xml::select` will fail.

- `xml::node_info node`

Retrieve the node data from `node`. Returns a [node info](#) value, as described in [Data Structure](#) above. Fails if the node does not belong to one of the supported node types.

- `xml::is_blank_node`

Checks whether a node is a blank node (empty or whitespace only) and thus possibly ignorable.

- `xml::node_base node`

Returns the base URI of the given node.

- `xml::node_path node`

Returns the path of a node in the document, in the format accepted by [select](#).

- `xml::node_content node`

Returns the text carried by the node, if any (after entity substitution).

In addition, you can retrieve and change attributes of element nodes with the following operations:

- `xml::node_attr node name`

Retrieves the value of the attribute with the given name (after entity substitution).

- `xml::set_node_attr node name value,xml::unset_node_attr node name`

Sets or unsets an attribute value.

Examples

Set and unset a node attribute:

```
> xml::set_node_attr r "foo" "4711";
()
> xml::node_info r;
xml::element "story" [] ["foo"=>"4711"]
> xml::node_attr r "foo";
"4711"
> xml::unset_node_attr r "foo";
()
> xml::node_info r;
xml::element "story" [] []
```

The `select` function is *very* powerful, and probably the single most important operation of this module if you want to extract information from an existing XML document without traversing the entire structure. Here is a very simple example of its use:

```
> [xml::node_content n, xml::node_path n | n = xml::select sample "//author"];
[("John Fleck", "/story/storyinfo/author")]
```

Note that if the XPath expression contains qualified names, the corresponding namespace prefixes and their URIs must be given in the second argument along with the XPath, as follows:

```
xml::select doc ("//foo:bar", ["foo"=>"http://www.foo.org"]);
```

5.4 Node Manipulation

These operations enable you to manipulate the document structure by adding a new node to the document tree (specified through its `node info`), and by removing (unlinking) existing nodes from the tree.

- `xml::replace node info`
Add the new node specified by `info` in place of the given node node.
- `xml::add_first node info, xml::add_last node info`
Add the new node as the first or last child of node, respectively.
- `xml::add_next node info, xml::add_prev node info`
Add the new node as the next or previous sibling of node, respectively.

The operations above all return a pointer to the new XML node object.

- `xml::unlink node`
Deletes an existing node from the document tree. Returns `()`.

Examples

Replace the first child of the root node in the sample document:

```
> xml::node_info (r!0);
xml::cdata "<greeting>Hello, world!</greeting>"
> xml::replace (r!0) (xml::text "bla bla");
#<pointer 0xd40d80>
> xml::node_info (r!0);
xml::text "bla bla"
```

Delete that node:

```
> xml::unlink (r!0);
()
> xml::node_info (r!0);
xml::comment "This is a sample document for testing the xml interface."
```

5.5 Transformations

The following operations provide basic XSLT support. As already mentioned, stylesheets are represented as pointers to the `xsltStylesheet` structure provided by `libxslt`. Note that, in difference to XML document pointers, this is an opaque type, i.e., there is no direct means to inspect and manipulate parsed stylesheets in memory using the operations of this module. However, a stylesheet is just a special kind of XML document and thus can be manipulated after reading the stylesheet as an ordinary XML document. The [load_stylesheet](#) function then allows you to convert the document pointer to an XSLT Stylesheet object.

Applying a stylesheet to an XML document generally involves the following steps:

1. Load and parse the stylesheet using [load_stylesheet](#). The parameter to [load_stylesheet](#) can be either the name of a stylesheet file or a corresponding document pointer. The function returns a pointer to the stylesheet object which is used in the subsequent processing.
2. Invoke [apply_stylesheet](#) on the stylesheet and the target document. This returns a new document containing the transformed XML document.
3. Run [save_result_file](#) or [save_result_string](#) on the result and the stylesheet to save the transformed document in a file or a string.

Here is a brief summary of the XSLT operations. Please check the [XSLT](#) documentation for details of the transformation process.

- `xslt::load_stylesheet x`

Load a stylesheet. `x` can be either an XML document pointer, or a string denoting the desired `.xsl` file.

- `xslt::apply_stylesheet style doc params`

Apply the stylesheet `style` to the given document `doc` with the given parameters `params`. The third argument is a (possibly empty) list of `key=>value` string pairs which allows you to pass additional parameters to the stylesheet.

- `xslt::save_result_file` name doc style compression,
`xslt::save_result_string` doc style

Save the transformation result doc obtained with the stylesheet style to a file or a string. This works pretty much like [save_file](#) or [save_string](#), but also keeps track of some output-related information contained in the stylesheet.

Example

Load the `recipes.xml` document and the `recipes.xsl` stylesheet distributed with the sources:

```
> let recipes = xml::load_file "recipes.xml" xml::DTDVALID;  
> let style = xslt::load_stylesheet "recipes.xsl";
```

Apply the stylesheet to the document and save the result in a html file:

```
> let res = xslt::apply_stylesheet style recipes [];  
> xslt::save_result_file "recipes.html" res style 0;  
( )
```

That's all. You can now have a look at `recipes.html` in your favourite web browser.

pure-g2

Version 0.1, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

This is a straight wrapper of the g2 graphics library, see <http://g2.sf.net/>.

License: BSD-style, see the COPYING file for details.

Get the latest source from <http://pure-lang.googlecode.com/files/pure-g2-0.1.tar.gz>.

g2 is a simple, no-frills 2D graphics library, distributed under the LGPL. It's easy to use, portable and supports PostScript, X11, PNG and Win32. Just the kind of thing that you need if you want to quickly knock out some basic graphics, and whipping out the almighty OpenGL or GTK/Cairo seems overkill.

To use this module, you need to have libg2 installed as a shared library (libg2.so, .dll etc.) in a place where the Pure interpreter can find it.

Documentation still needs to be written, so for the time being please see g2.pure and have a look at the examples provided in the distribution.

Run `make install` to copy g2.pure to the Pure library directory. This tries to guess the prefix under which Pure is installed; if this doesn't work, you'll have to set the prefix variable in the Makefile accordingly.

The Makefile also provides the following targets:

- `make examples` compiles the examples to native executables.
- `make clean` deletes the native executables for the examples, as well as some graphics files which are produced by running g2_test.pure.
- `make generate` regenerates the g2.pure module. This requires that you have pure-gen installed, as well as the g2 header files (you can point pure-gen to the prefix under which g2 is installed with the g2prefix variable in the Makefile). This step shouldn't normally be necessary, unless you find that the provided wrapper doesn't work with your g2 version. The g2.pure in this release has been generated from g2 0.72.

Pure OpenGL Bindings

Version 0.8, November 14, 2010

Scott Dillard

Albert Graef <Dr.Graef@t-online.de>

These are fairly complete Pure bindings for the OpenGL graphics library, which allow you to do 2D and 3D graphics programming with Pure. The bindings should work out of the box on most contemporary systems which have OpenGL drivers installed, thanks to Scott's on-demand loading code for the GL functions, which accounts for the fact that different GL implementations will export different functions. (Mostly to account for Microsoft's Museum of Ancient OpenGL History, otherwise known as `opengl32.dll`.)

Information about OpenGL can be found at: <http://www.opengl.org/>

As of pure-gl 0.5, the bindings are now generated using pure-gen instead of Scott's original OpenGL-specific generator. The stuff needed to do this is included (except pure-gen, which is a separate package available from the Pure website), so that you can regenerate the bindings if necessary.

1 Copying

Copyright (c) 2009, Scott E Dillard

Copyright (c) 2009, Albert Graef

Copyright (c) 2002-2005, Sven Panne

pure-gl is distributed under a 3-clause BSD-style license, please see the accompanying COPYING file for details.

2 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-gl-0.8.tar.gz>.

Normally you just run `make` && `sudo make install`, as with the other Pure modules. (See The Makefile for further options.) This doesn't regenerate the bindings and so can be done on any system which has Pure, OpenGL and a C compiler installed.

If you miss some vendor-specific OpenGL functionality which is in your system's header files but not in the distributed bindings, with some effort you can fix that yourself by regenerating the bindings, see below.

3 Using the GL Bindings

The bindings mainly consist of 3 Pure files: `GL.pure`, `GLU.pure` and `GLUT.pure`.

In your Pure program, write something like:

```
using GL, GLU, GLUT;
```

`GL.pure` covers OpenGL up through version 2.1. To get access to extensions, you include `GL_XXX.pure` where `XXX` is the extensions suffix. Currently, there are `GL_ARB.pure`, `GL_EXT.pure`, `GL_NV.pure` and `GL_ATI.pure`, which should cover about 99% of the useful extensions out there. If you need more than that, it is straightforward to tweak the Makefile to scrape some of the more esoteric extensions from your headers. All OpenGL functions are loaded on first use. If your OpenGL implementation does not define a given function, a `gl_unsupported` exception is thrown with the name of the function as its only argument.

The functions are in namespaces `GL`, `GLU` and `GLUT` respectively. Functions are in curried form, i.e.:

```
GL::Vertex3d 1.0 2.0 3.0;
```

GL enumerants are in uppercase, as in C:

```
GL::Begin GL::LINE_STRIP;
```

Currently, if the `GLU` or `GLUT` bindings reference a function that your DLL does not contain, it echoes this to stdout. I'm working on a way to suppress this.

Some examples can be found in the `examples` subdirectory. This also includes a wrapper of Rasterman's `imlib2` library (also generated with `pure-gen`), and an example which uses this to render an image as a texture.

The `examples/flexi-line` directory contains Eduardo Cavazos' port of the flexi-line demo. Run `pure flexi-line-auto.pure`, sit back and enjoy. There's also an interactive version of the demo available in `flexi-line.pure`.

4 Regenerating the Bindings

You need to have pure-gen installed to do this.

Also make sure that you have the OpenGL headers installed. By default, the Makefile assumes that they are in the GL subdirectory of /usr/include, you can set the `glpath` variable in the Makefile accordingly to change this. (Set `glpath` to the path under which the GL subdirectory resides, not to the GL subdirectory itself. See below for an example.) Note that on Linux systems, /usr/include/GL usually contains the MESA headers. If available, you may want to use your GPU vendor's headers instead, to get all the extensions available on your system.

Alternatively, you can also just put the headers (`gl.h`, `glx.h`, `glu.h`, `glut.h`, and any other OpenGL headers that get `#included` in those) into the GL subdirectory of the pure-gl sources, by copying them over or creating symbolic links to them. This is particularly useful for maintainers, who may want to use a "staged" header set which is different from the installed OpenGL headers. The `."` directory will always be searched first, so you can also just put the vendor-specific headers there. For instance, if you're like Scott and you use Ubuntu with an Nvidia GPU, then you can do this:

```
cd pure-gl/GL
ln -s /usr/share/doc/nvidia-glx-new-dev/include/GL/gl.h
ln -s /usr/share/doc/nvidia-glx-new-dev/include/GL/glx.h
```

Finally, the Makefile also assumes that you have freeglut (an improved GLUT replacement) installed and want all the extensions offered by freeglut. To use the vanilla GLUT without the extensions instead, you only have to change the value of the source variable in the Makefile from `GL/all_gl_freeglut.h` to `GL/all_gl.h`. If you use `openglut` instead of `freeglut` you will have to change the `GL/all_gl_freeglut.h` file accordingly.

Once you have set up things to your liking, you can regenerate the bindings by running `make` as follows:

```
make generate
```

If you need a custom path to the OpenGL headers as described above (say, `/usr/local/include`) then do this instead:

```
make generate glpath=/usr/local/include
```

If you're lucky, this will regenerate all the `GL*.pure` and `GL*.c` files, and recompile the shared module from the `GL*.c` files after that. This shared module, instead of the OpenGL libraries themselves, is what gets loaded by the Pure modules.

If you're not so lucky, save a complete build log with all the error messages and ask on the pure-lang mailing list for help.

See the "Generator stuff" section in the Makefile for further options. Adding a rule for other extensions should be easy, just have a look at an existing one (e.g., `GL_EXT.c`) and modify it accordingly.

Pure GTK+ Bindings

Version 0.8, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

pure-gtk is a collection of bindings to use the GTK+ GUI toolkit version 2.x with Pure, see <http://www.gtk.org>. The bindings include the gtk (+gdk), glib, atk, cairo and pango libraries, each in their own Pure module.

At present these are just straight 1-1 wrappers of the C libraries, created with pure-gen. So they still lack some convenience, but they are perfectly usable already, and a higher-level API for accessing all the functionality will hopefully become available in time. In fact *you* can help make that happen. :) So please let me know if you'd like to give a helping hand in improving pure-gtk.

1 Copying

Copyright (c) 2008, 2009 by Albert Graef.

pure-gtk is free software: you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

pure-gtk is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

2 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-gtk-0.8.tar.gz>.

To install, do the usual `make && sudo make install` (see the Makefile for further options). This needs Pure and the GTK header files and libraries.

This release was prepared with GTK+ 2.14.4 on SUSE 11.1. If you're seeing a lot of warnings or even compilation errors when running `make`, your GTK headers are probably much different from these. In that case you should run `make generate` to regenerate the bindings; for this you also need to have `pure-gen` installed. (If you already have `pure-gen` then it's a good idea to do this anyway.)

3 Usage

See `examples/hello.pure` for a basic example. The files `uiexample.pure` and `uiexample.glade` show how to run a GUI created with the Glade-3 interface builder. This needs a recent version of the GtkBuilder API to work. (If you're still running Glade-2 and an older GTK+ version, you might want to use the older `libglade` interface instead. Support for that is in the Makefile, but it's not enabled by default.) NOTE: The examples start up much faster when they are compiled to native executables. To do this, just run `make examples` after `make`. (Be patient, this takes a while.)

`pure-gtk` can be discussed on the Pure mailing list at:
<http://groups.google.com/group/pure-lang>

pure-tk

Version 0.1, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

Pure's [Tcl/Tk](#) interface.

1 Introduction

This module provides a basic interface between Pure and Tcl/Tk. The operations of this module allow you to execute arbitrary commands in the Tcl interpreter, set and retrieve variable values in the interpreter, and invoke Pure callbacks from Tcl/Tk.

A recent version of Tcl/Tk is required (8.0 or later should do). You can get this from <http://www.tcl.tk>. Both releases in source form and binary releases for Windows and various Unix systems are provided there.

Some information on how to use this module can be found below. But you'll find that pure-tk is very easy to use, so you might just want to look at the programs in the examples folder to pick it up at a glance. A very basic example can be found in `tk_hello.pure`; a slightly more advanced example of a tiny but complete Tk application is in `tk_examp.pure`.

pure-tk also offers special support for Peter G. Baum's [Gnocl](#) extension which turns Tcl into a frontend for [GTK+](#) and [Gnome](#). If you have Gnocl installed then you can easily create GTK+/Gnome applications, either from Tcl sources or from [Glade](#) UI files, using the provided `gnocl.pure` module. See the included `uiexample.pure` and the accompanying Glade UI file for a simple example. Also, some basic information on using Gnocl with pure-tk can be found in the [Tips and Tricks](#) section below.

One nice thing about Tcl/Tk is that it provides a bridge to a lot of other useful libraries. A prominent example is [VTK](#), a powerful open-source 3D visualization toolkit which comes with full Tcl/Tk bindings. The examples directory contains a simple example (`earth.pure` and `earth.tcl`) which shows how you can employ these bindings to write cool animated 3D applications using either Tk or Gnocl as the GUI toolkit.

2 Copying

Copyright (c) 2010 by Albert Gräf, all rights reserved. pure-tk is distributed under a BSD-style license, see the COPYING file for details.

3 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-tk-0.1.tar.gz>.

As with the other addon modules for Pure, running `make` && `sudo make install` should usually do the trick. This requires that you have Pure and Tcl/Tk installed. `make` tries to guess your Pure installation directory and platform-specific setup. If it gets this wrong, you can set some variables manually. In particular, `make install prefix=/usr` sets the installation prefix. Please see the Makefile for details.

IMPORTANT: When starting a new interpreter, the Tcl/Tk initialization code looks for some initialization files which it executes before anything else happens. Usually these files will be found without any further ado, but if that does not happen automatically, you must set the `TCL_LIBRARY` and `TK_LIBRARY` environment variables to point to the Tcl and Tk library directories on your system.

All programs in the examples subdirectory have been set up so that they can be compiled to native executables, and a Makefile is provided in that directory to handle this. So after installing pure-tk you just need to type `make` there to compile the examples. (This step isn't necessary, though, you can also just run the examples with the Pure interpreter as usual.)

4 Basic Usage

You can submit a command to the Tcl interpreter with `tk cmd` where `cmd` is a string containing the command to be executed. If the Tcl command returns a value (i.e., a nonempty string) then `tk` returns that string, otherwise it returns `()`.

`tk` also starts a new instance of the Tcl interpreter if it is not already running. To stop the Tcl interpreter, you can use the `tk_quit` function. Note that, as far as pure-tk is concerned, there's only one Tcl interpreter per process, but of course you can create secondary interpreter instances in the Tcl interpreter using the appropriate Tcl commands.

Simple dialogs can be created directly using Tk's `tk_messageBox` and `tk_dialog` functions. For instance:

```
tk "tk_dialog .warning \"Warning\" \"Are you sure?\" warning 0 Yes No Cancel";
```

Other kinds of common dialogs are available; see the Tcl/Tk manual for information.

For more elaborate applications you probably have to explicitly create some widgets, add the appropriate callbacks and provide a main loop which takes care of processing events in the Tcl/Tk GUI. We discuss this in the following.

5 Callbacks

pure-tk installs a special Tcl command named `pure` in the interpreter which can be used to implement callbacks in Pure. This command is invoked from Tcl as follows:

```
pure function args ...
```

It calls the Pure function named by the first argument, passing any remaining (string) arguments to the callback. If the Pure callback returns a (nonempty) string, that value becomes the return value of the `pure` command, otherwise the result returned to the Tcl interpreter is empty.

Pure callbacks are installed on Tk widgets just like any other, just using the `pure` command as the actual callback command. For instance, you can define a callback which gets invoked when a button is pushed as follows:

```
using tk, system;
tk "button .b -text {Hello, world!} -command {pure hello}; pack .b";
hello = puts "Hello, world!";
```

6 The Main Loop

The easiest way to provide a main loop for your application is to just call `tk_main` which keeps processing events in the Tcl interpreter until the interpreter is exited. You can terminate the interpreter in a Pure callback by calling `tk_quit`. Thus a minimalistic Tcl/Tk application coded in Pure may look as follows:

```
using tk;
tk "button .b -text {Hello, world!} -command {pure tk_quit}; pack .b";
tk_main;
```

The main loop terminates as soon as the Tcl interpreter is exited, which can happen, e.g., in response to a callback which invokes the `tk_quit` function (as shown above) or Tcl code which destroys the main window (`destroy .`). The user can also close the main window from the window manager in order to exit the main loop.

7 Accessing Tcl Variables

pure-tk allows your script to set and retrieve variable values in the Tcl interpreter with the `tk_set`, `tk_unset` and `tk_get` functions. This is useful, e.g., to change the variables associated with entry and button widgets, and to retrieve the current values from the application. For instance:

```
> tk_set "entry_val" "some string";
"some string"
> tk_get "entry_val";
```

```
"some string"
> tk_unset "entry_val";
()
> tk_get "entry_val";
tk_get "entry_val"
```

Note that `tk_set` returns the assigned value, so it is possible to chain such calls if several variables have to be set to the same value:

```
> tk_set "foo" $ tk_set "bar" "yes";
"yes"
> map tk_get ["foo","bar"];
["yes","yes"]
```

8 Conversions Between Pure and Tcl Values

As far as pure-tk is concerned, all Tcl values are strings (in fact, that's just what they are at the Tcl language level, although the Tcl interpreter uses more elaborate representations of objects such as lists internally). There are no automatic conversions of any kind. Thus, the arguments passed to a Pure callback and the result returned by `tk` are simply strings in Pure land. The same holds for the `tk_set` and `tk_get` functions.

However, there are a few helper functions which can be used to convert between Tcl and Pure data. First, the following operations convert Pure lists to corresponding Tcl lists and vice versa:

```
> tk_join ["0","1.0","Hello, world!"];
"0 1.0 {Hello, world!}"
> tk_split ans;
["0","1.0","Hello, world!"]
```

The `tk_str` and `tk_val` operations work in a similar fashion, but they also do automatic conversions for numeric values (ints, bigints and doubles):

```
> tk_str [0,1.0,"Hello, world!"];
"0 1.0 {Hello, world!}"
> tk_val ans;
[0,1.0,"Hello, world!"]
```

In addition, these operations also convert single atomic values:

```
> tk_str 1.0;
"1.0"
> tk_val ans;
1.0
```

9 Tips and Tricks

Here are a few other things that are worth keeping in mind when working with pure-tk.

- Errors in Tcl/Tk commands can be handled by giving an appropriate definition of the `tk_error` function, which is invoked with an error message as its single argument. For instance, the following implementation of `tk_error` throws an exception:

```
tk_error msg = throw msg;
```

If no definition for this function is provided, then errors cause a literal `tk_error msg` expression to be returned as the result of the `tk` function. You can then check for such results and take an appropriate action.

- The Tcl interpreter, when started, displays a default main window, which is required by most Tk applications. If this is not desired (e.g., if only the basic Tcl commands are needed), you can hide this window using a `tk "wm withdraw ."` command. To redisplay the window when it is needed, use the `tk "wm deiconify ."` command. It is also common practice to use `wm withdraw` and `wm deiconify` while creating the widgets of an application, in order to reduce “flickering”.
- Instead of calling `tk_main`, you can also code your own main loop in Pure as follows:

```
main = do_something $$ main if tk_ready;  
      = () otherwise;
```

Note that the `tk_ready` function checks whether the Tcl interpreter is still up and running, after processing any pending events in the interpreter. This setup allows you to do your own custom idle processing in Pure while the application is running. However, you have to be careful that your `do_something` routine runs neither too short nor too long (a few milliseconds should usually be ok). Otherwise your main loop may turn into a busy loop and/or the GUI may become very sluggish and unresponsive. Thus it's usually better to just call `tk_main` and do any necessary background processing using the Tcl interpreter's own facilities (e.g., by setting up a Pure callback with the `Tcl after command`).

- The `tk` function can become rather tedious when coding larger Tk applications. Usually, you will prefer to put the commands making up your application into a separate Tcl script. One way to incorporate the Tcl script into your Pure program is to use the `Tcl source command`, e.g.:

```
tk "source myapp.tcl";
```

However, this always requires the script to be available at runtime. Another method is to read the script into a string which is assigned to a Pure constant, and then invoke the `tk` command on this string value:

```
using system;  
const ui = fget $ fopen "myapp.tcl" "r";  
tk ui;
```

This still reads the script at runtime if the Pure program is executed using the Pure interpreter. However, you can now compile the Pure program to a native executable (see the Pure manual for details on this), in which case the text of the Tcl script is included verbatim in the executable. The compiled program can then be run without having the original Tcl script file available:

```
$ pure -c myapp.pure -o myapp
$ ./myapp
```

This is also the method to use for running existing Tk applications, e.g., if you create the interface using some interface builder like [vtcl](#).

- The Tcl package command allows you to load additional extensions into the Tcl interpreter at runtime. For instance:

```
tk "package require Gnocl";
```

This loads Peter G. Baum's [Gnocl](#) extension which turns Tcl into a frontend for [GTK+](#) and [Gnome](#). In fact, pure-tk includes a special module to handle the nitty-gritty details of creating a GTK+/Gnome application from a [Glade](#) UI file and set up Pure callbacks as specified in the UI file. To use this, just import the `gnocl.pure` module into your Pure scripts:

```
using gnocl;
```

Note that the Glade interface requires that you have a fairly recent version of Gnocl installed (Gnocl 0.9.94g has been tested). The other facilities provided by the `gnocl.pure` module should also work with older Gnocl versions such as Gnocl 0.9.91. Please see the `gnocl.pure` module and the corresponding examples included in the sources for more information.

- The Tcl `exit` procedure, just as in `tcsh` or `wish`, causes exit from the current process. Since the Tcl interpreter hosted by the `pure-tk` module runs as part of a Pure program and not as a separate child process, this might not be what you want. If you'd like exit to only exit the Tcl interpreter, without exiting the Pure program, you can redefine the `exit` procedure, e.g., as follows:

```
tk "proc exit { {returnCode 0} } { pure tk_quit }";
```

If you want to do something with the exit code provided by `exit`, you will have to provide an appropriate callback function, e.g.:

```
tk "proc exit { {returnCode 0} } { pure quit_cb $returnCode }";
```

A suitable implementation of `quit_cb` might look as follows:

```
quit_cb 0 = puts "Application exited normally." $$ tk_quit;
quit_cb n = printf "Application exited with exit code %d.\n" n $$
               tk_quit otherwise;
```

- If you need dialogs beyond the standard kinds of message boxes and common dialogs, you will have to do these yourself using a secondary toplevel. The dialog toplevel is

just like the main window but will only be shown when the application needs it. You can construct both non-modal and modal dialogs this way, the latter can be implemented using Tk's `grab` command.

faust2pd: Pd Patch Generator for Faust

Version 2.2, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

This package contains software which makes it easier to use Faust DSPs with Pd and the Pure programming language. The main component is the faust2pd script which creates GUI wrappers for Faust DSPs. The package also includes a bunch of examples. The software is distributed under the GPL; see the COPYING file for details.

NOTE: This faust2pd version is written in Pure and was ported from an earlier version written in Pure's predecessor Q. The version of the script provided here should be 100% backward-compatible to those previous versions, except for the following changes:

- The (rarely used) -f (a.k.a. -fake-buttons) option was renamed to -b.
- A new -f (a.k.a. -font-size) option was added to change the font size of the GUI elements.
- Most command line options can now also be specified using special [pd: . . .] tags in the Faust source.

Also note that you can now run the script on 64 bit systems (Q never worked there).

As of version 2.1, the faust2pd script is now compiled to a native executable before installation. This makes the program start up much faster, which is a big advantage because most xml files don't take long to be processed. You'll need Pure 0.21 or later to make this work.

For Windows users, there's an MSI package of Faust available at the [Pure](#) website, which also includes faust2pd.

1 Copying

Copyright (c) 2009 by Albert Graef.

faust2pd is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

faust2pd is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

2 Requirements

faust2pd is known to work on Linux and Windows, and there shouldn't be any major roadblocks preventing it to work on other systems supported by Pure.

The faust2pd script is written in the [Pure](#) programming language and requires Pure's XML module, so you need to install these to make it work. Install the latest pure*.tar.gz and pure-xml*.tar.gz packages and you should be set. (Pure 0.21 or later is required.) Also make sure that the LLVM base package is installed, as described in the Pure INSTALL file, some LLVM utilities are needed to make Pure's batch compiler work.

To run the seqdemo example, you'll also need the Pd Pure external (pd-pure*.tar.gz), also available at the [Pure](#) website.

To compile the examples, you'll need GNU C++ and make, [Pd](#) and, of course, [Faust](#). Make sure you get a recent version of Faust; Faust releases >0.9.8 include the puredata architecture necessary to create Pd externals from Faust DSPs.

3 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/faust2pd-2.2.tar.gz>.

Run `make` and `make install` to compile and install the faust2pd program on your system. You can set the installation prefix by running `make install prefix=/some/path`. Default installation prefix is `/usr/local`, faust2pd is installed in the bin directory below that.

Optionally, you can also run `make install-pd` to copy the supporting Pd abstractions (faust*.pd) to your `lib/pd/extra` directory, so that you can use the patches generated by faust2pd without copying these abstractions to your working directory. The Makefile tries to guess the prefix of your Pd installation, if it guesses wrong, you can specify the prefix explicitly by running `make install-pd pdprefix=/some/path`.

The included faustxml.pure script provides access to Faust-generated dsp descriptions in xml files to Pure scripts. This module is described in its own [appendix](#) below. It may have uses beyond faust2pd, but isn't normally installed. If you want to use this module, you can just copy it to your Pure library directory.

4 Quickstart

Run `make examples` to compile the Faust examples included in this package to corresponding Pd plugins. After that you should be able to run the patches in the various subdirectories of the examples directory. Everything is set up so that you can try the examples “in-place”, without installing anything except the required software as noted in [Requirements](#) above. You can also run `make realclean` before `make` to regenerate everything from scratch (this requires `faust2pd`, so this will only work if you already installed the Pure interpreter).

Faust Pd plugins work in much the same way as the well-known `plugin~` object (which interfaces to LADSPA plugins), except that each Faust DSP is compiled to its own Pd external. Under Linux, the basic compilation process is as follows (taking the `freeverb` module from the Faust distribution as an example):

```
# compile the Faust source to a C++ module using the "puredata" architecture
faust -a puredata.cpp freeverb.dsp -o freeverb.cpp
# compile the C++ module to a Pd plugin
g++ -shared -Dmydsp=freeverb freeverb.cpp -o freeverb~.pd_linux
```

By these means, a Faust DSP named `xyz` with `n` audio inputs and `m` audio outputs becomes a Pd object `xyz~` with `n+1` inlets and `m+1` outlets. The leftmost inlet/outlet pair is for control messages only. This allows you to inspect and change the controls the unit provides, as detailed below. The remaining inlets and outlets are the audio inputs and outputs of the unit, respectively. For instance, `freeverb.dsp` becomes the Pd object `freeverb~` which, in addition to the control inlet/outlet pair, has 2 audio inputs and outputs.

When creating a Faust object it is also possible to specify, as optional creation parameters, an extra unit name (this is explained in the following section) and a sample rate. If no sample rate is specified explicitly, it defaults to the sample rate at which Pd is executing. (Usually it is not necessary or even desirable to override the default choice, but this might occasionally be useful for debugging purposes.)

As already mentioned, the main ingredient of this package is a Pure script named “`faust2pd`” which allows you to create Pd abstractions as “wrappers” around Faust units. The wrappers generated by `faust2pd` can be used in Pd patches just like any other Pd objects. They are much easier to operate than the “naked” Faust plugins themselves, as they also provide “graph-on-parent” GUI elements to inspect and change the control values.

The process to compile a plugin and build a wrapper patch is very similar to what we’ve seen above. You only have to add the `-xml` option when invoking the Faust compiler and run `faust2pd` on the resulting XML file:

```
# compile the Faust source and generate the xml file
faust -a puredata.cpp -xml freeverb.dsp -o freeverb.cpp
# compile the C++ module to a Pd plugin
g++ -shared -Dmydsp=freeverb freeverb.cpp -o freeverb~.pd_linux
# generate the Pd patch from the xml file
faust2pd freeverb.dsp.xml
```

Please see [Wrapping DSPs with faust2pd](#) below for further details.

Note that, just as with other Pd externals and abstractions, the compiled `.pd_linux` modules and wrapper patches must be put somewhere where Pd can find them. To these ends you can either move the files into the directory with the patches that use the plugin, or you can put them into the `lib/pd/extra` directory or some other directory on Pd's library path for system-wide use.

Also, `faust2pd`-generated wrappers use a number of supporting abstractions (the `faust-*.pd` files in the `faust2pd` sources), so you have to put these into the directory of the main patch or install them under `lib/pd/extra` as well. (The `make pd-install` step does the latter, see [Installation](#) above.)

5 Control Interface

The control inlet of a Faust plugin understands messages in one of the following forms:

- `bang`, which reports all available controls of the unit on the control outlet, in the form: `type name val init min max step`, where `type` is the type of the control as specified in the Faust source (`checkbox`, `nentry`, etc.), `name` its (fully qualified) name, `val` the current value, and `init`, `min`, `max`, `step` the initial value, minimum, maximum and stepsize of the control, respectively.
- `foo 0.99`, which sets the control `foo` to the value 0.99, and outputs nothing.
- Just `foo`, which outputs the (fully qualified) name and current value of the `foo` control on the control outlet.

Control names can be specified in their fully qualified form, like e.g. `/gnu/bar/foo` which indicates the control `foo` in the subgroup `bar` of the topmost group `gnu`, following the hierarchical group layout defined in the Faust source. This lets you distinguish between different controls with the same name which are located in different groups. To find out about all the controls of a unit and their fully qualified names, you can `bang` the control inlet of the unit as described above, and connect its control outlet to a `print` object, which will cause the descriptions of all controls to be printed in Pd's main window. (The same information can also be used, e.g., to initialize GUI elements with the proper values. Patches generated with `faust2pd` rely on this.)

You can also specify just a part of the control path (like `bar/foo` or just `foo` in the example above) which means that the message applies to *all* controls which have the given pathname as the final portion of their fully qualified name. Thus, if there is more than one `foo` control in different groups of the Faust unit then sending the message `foo` to the control inlet will report the fully qualified name and value for each of them. Likewise, sending `foo 0.99` will set the value of all controls named `foo` at once.

Concerning the naming of Faust controls in Pd you should also note the following:

- A unit name can be specified at object creation time, in which case the given symbol is used as a prefix for all control names of the unit. E.g., the control `/gnu/bar/foo` of an object `baz~` created with `baz~ baz1` has the fully qualified name `/baz1/gnu/bar/foo`.

This lets you distinguish different instances of an object such as, e.g., different voices of a polyphonic synth unit.

- Pd's input syntax for symbols is rather restrictive. Therefore group and control names in the Faust source are mangled into a form which only contains alphanumeric characters and hyphens, so that the control names are always legal Pd symbols. For instance, a Faust control name like `meter #1 (dB)` will become `meter -1 -dB` which can be input directly as a symbol in Pd without any problems.
- "Anonymous" groups and controls (groups and controls which have empty labels in the Faust source) are omitted from the path specification. E.g., if `foo` is a control located in a main group with an empty name then the fully qualified name of the control is just `/foo` rather than `//foo`. Likewise, an anonymous control in the group `/foo/bar` is named just `/foo/bar` instead of `/foo/bar/`.

Last but not least, there is also a special control named `active` which is generated automatically for your convenience. The default behaviour of this control is as follows:

- When `active` is nonzero (the default), the unit works as usual.
- When `active` is zero, and the unit's number of audio inputs and outputs match, then the audio input is simply passed through.
- When `active` is zero, but the unit's number of audio inputs and outputs do *not* match, then the unit generates silence.

The `active` control frequently alleviates the need for special "bypass" or "mute" controls in the Faust source. However, if the default behaviour of the generated control is not appropriate you can also define your own custom version of `active` explicitly in the Faust program; in this case the custom version will override the default one.

6 Examples

In the examples subdirectory you'll find a bunch of sample Faust DSPs and Pd patches illustrating how Faust units are used in Pd.

- The `examples/basic/test.pd` patch demonstrates the basics of operating "bare" Faust plugins in Pd. You'll rarely have to do this when using the wrappers generated with the `faust2pd` program, but it is a useful starting point to take a look behind the scenes anyway.
- The `examples/faust` directory contains all the examples from (an earlier version of) the Faust distribution, along with corresponding Pd wrappers generated with `faust2pd`. Have a look at `examples/faust/faustdemo.pd` to see some of the DSPs in action. Note that not all examples from the Faust distribution are working out of the box because of name clashes with Pd builtins, so we renamed those. We also edited some of the `.dsp` sources (e.g., turning buttons into checkboxes or sliders into `nentries`) where this seemed necessary to make it easier to operate the Pd patches.

- The examples/synth directory contains various plugins and patches showing how to implement polyphonic synthesizers using Faust units. Take a look at examples/synth/synth.pd for an example. If you have properly configured your interfaces then you should be able to play the synthesizer via Pd's MIDI input.
- The examples/seqdemo/seqdemo.pd patch demonstrates how to operate a multitimbral synth, built with Faust units, in an automatic fashion using a pattern sequencer programmed in Pure. This example requires the Pure interpreter as well as the pd-pure plugin available from <http://pure-lang.googlecode.com>.

7 Wrapping DSPs with faust2pd

The faust2pd script generates Pd patches from the dsp.xml files created by Faust when run with the -xml option. Most of the sample patches were actually created that way. After installation you can run the script as follows:

```
faust2pd [-hVbs] [-f size] [-o output-file] [-n #voices]
         [-r max] [-X patterns] [-x width] [-y height] input-file
```

The default output filename is input-file with new extension .pd. Thus, faust2pd filename.dsp.xml creates a Pd patch named filename.pd from the Faust XML description in filename.dsp.xml.

The faust2pd program understands a number of options which affect the layout of the GUI elements and the contents of the generated patch. Here is a brief list of the available options:

- h, --help** display a short help message and exit
- V, --version** display the version number and exit
- b, --fake-buttons** replace buttons (bangs) with checkboxes (toggles)
- f, --font-size** font size for GUI elements (10 by default)
- n, --nvoices** create a synth patch with the given number of voices
- o, --output-file** output file name (.pd file)
- r, --radio-sliders** radio controls for sliders
- s, --slider-nums** sliders with additional number control
- X, --exclude** exclude controls matching the given glob patterns
- x, --width** maximum width of the GUI area
- y, --height** maximum height of the GUI area

Just like the Faust plugin itself, the generated patch has a control input/output as the left-most inlet/outlet pair, and the remaining plugs are signal inlets and outlets for each audio input/output of the Faust unit. However, the control inlet/outlet pair works slightly different from that of the Faust plugin. Instead of being used for control replies, the control outlet of the patch simply passes through its control input (after processing messages which are

understood by the wrapped plugin). By these means control messages can flow along with the audio signal through an entire chain of Faust units. (You can find an example of this in `examples/faust/faustdemo.pd`.) Moreover, when generating a polyphonic synth patch using the `-n` option then there will actually be two control inlets, one for note messages and one for ordinary control messages; this is illustrated in the `examples/synth/synth.pd` example.

The generated patch also includes the necessary GUI elements to see and change all (active and passive) controls of the Faust unit. Faust control elements are mapped to Pd GUI elements in an obvious fashion, following the horizontal and vertical layout specified in the Faust source. The script also adds special buttons for resetting all controls to their defaults and to operate the special active control.

This generally works very well, but you should be aware that the control GUIs generated by `faust2pd` are somewhat hampered by the limited range of GUI elements available in a vanilla Pd installation. As a remedy, `faust2pd` provides various options to change the content of the generated wrapper and work around these limitations.

- There are no real button widgets as required by the Faust specification, so bangs are used instead. There is a global delay time for switching the control from 1 back to 0, which can be changed by sending a value in milliseconds to the `faust-delay` receiver. If you need interactive control over the switching time then it is better to use checkboxes instead, or you can have `faust2pd` automatically substitute checkboxes for all buttons in a patch by invoking it with the `-f` a.k.a. `-fake-buttons` option.
- Sliders in Pd do not display their value in numeric form so it may be hard to figure out what the current value is. Therefore `faust2pd` has an option `-s` a.k.a. `-slider-nums` which causes it to add a number box to each slider control. (This flag also applies to Faust's passive bargraph controls, as these are implemented using sliders, see below.)
- Pd's sliders also have no provision for specifying a stepsize, so they are an awkward way to input integral values from a small range. OTOH, Faust doesn't support the "radio" control elements which Pd provides for that purpose. As a remedy, `faust2pd` allows you to specify the option `-r max` (a.k.a. `-radio-sliders=max`) to indicate that sliders with integral values from the range `0..max-1` are to be mapped to corresponding Pd radio controls.
- Faust's bargraphs are emulated using sliders. Note that these are passive controls which just display a value computed by the Faust unit. A different background color is used for these widgets so that you can distinguish them from the ordinary (active) slider controls. The values shown in passive controls are sampled every 40 ms by default. You can change this value by sending an appropriate message to the global `faust-timer` receiver.
- Since Pd has no "tabbed" (notebook-like) GUI element, Faust's tgroups are mapped to hgroups instead. It may be difficult to present large and complicated control interfaces without tabbed dialogs, though. As a remedy, you can control the amount of horizontal or vertical space available for the GUI area with the `-x` and `-y` (a.k.a. `-width` and `-height`) options and `faust2pd` will then try to break rows and columns in the layout to make everything fit within that area. (This feature has only been tested with simple layouts so far, so beware.)

- You can also exclude certain controls from appearing in the GUI using the `-X` option. This option takes a comma-separated list of shell glob patterns indicating either just the names or the fully qualified paths of Faust controls which are to be excluded from the GUI. For instance, the option `-X 'volume,meter*,faust/resonator?/*'` will exclude all volume controls, all controls whose names start with `meter`, and all controls in groups matching `faust/resonator?`. (Note that the argument to `-X` has to be quoted if it contains any wildcards such as `*` and `?`, so that the shell doesn't try to expand the patterns beforehand. Also note that only one `-X` option is recognized, so you have to specify all controls to be excluded as a single option.)
- Faust group labels are not shown at all, since I haven't found an easy way to draw some kind of labelled frame in Pd yet.

Despite these limitations, `faust2pd` appears to work rather well, at least for the kind of DSPs found in the Faust distribution. (Still, for more complicated control surfaces and interfaces to be used on stage you'll probably have to edit the generated GUI layouts by hand.)

For convenience, all the content-related command line options mentioned above can also be specified in the Faust source, as special tags in the label of the toplevel group of the dsp. These take the form `[pd:option]` or `[pd:option=value]` where `option` is any of the (long) options understood by `faust2pd`. For instance:

```
process = vgroup("mysynth [pd:nvoices=8] [pd:slider-nums]", ...);
```

Source options carrying arguments, like `nvoices` in the above example, can also be overridden with corresponding command line options.

8 Conclusion

Creating Faust plugins for use with Pd has never been easier before, so I hope that you'll soon have much joy trying your Faust programs in Pd. Add Pd-Pure to this, and you can program all your specialized audio and control objects using two modern-style functional languages which are much more fun than C/C++. Of course there's an initial learning curve to be mastered, but IMHO it is well worth the effort. The bottomline is that Pd+Faust+Pure really makes an excellent combo which provides you with a powerful, programmable interactive environment for creating advanced computer music and multimedia applications with ease.

8.1 Acknowledgements

Thanks are due to Yann Orlarey for his wonderful Faust, which makes developing DSP algorithms so easy and fun.

9 Appendix: faustxml

The `faustxml` module is provided along with `faust2pd` to retrieve the description of a Faust DSP from its XML file as a data structure which is ready to be processed by Pure programs. It may also be useful in other Pure applications which need to inspect description of Faust DSPs.

The main entry point is the `info` function which takes the name of a Faust-generated XML file as argument and returns a tuple (`name`, `descr`, `version`, `in`, `out`, `controls`) with the name, description, version, number of inputs and outputs and the toplevel group with the descriptions of the controls of the dsp. A couple of other convenience functions are provided to deal with the control descriptions.

9.1 Usage

Use the following declaration to import this module in your programs:

```
using faustxml;
```

For convenience, you can also use the following to get access to the module's namespace:

```
using namespace faustxml;
```

9.2 Data Structure

The following constructors are used to represent the UI controls of Faust DSPs:

- `button|checkbox label`: A button or checkbox with the given label.
- `nentry|vslider|hslider (label,init,min,max,step)`: A numeric input control with the given label, initial value, range and stepwidth.
- `vbargraph|hbargraph (label,min,max)`: A numeric output control with the given label and range.
- `vgroup|hgroup|tgroup (label,controls)`: A group with the given label and list of controls in the group.

9.3 Operations

controlp x Check for control description values.

control_type x, control_label x, control_args x Access functions for the various components of a control description.

controls x This function returns a flat representation of a control group `x` as a list of basic control descriptions, which provides a quick way to access all the control values of

a Faust DSP. The grouping controls themselves are omitted. You can pass the last component of the return value of the [info](#) function to this function.

pcontrols x Works like the [controls](#) function above, but also replaces the label of each basic control with a fully qualified path consisting of all control labels leading up to the given control. Thus, e.g., the label of a slider "gain" inside a group "voice#0" inside the main "faust" group will be denoted by the label "faust/voice#0/gain".

info fname Extract the description of a Faust DSP from its XML file. This is the main entry point. Returns a tuple with the name, description and version of the DSP, as well as the number of inputs and outputs and the toplevel group with all the control descriptions. Raises an exception if the XML file doesn't exist or contains invalid contents.

Example:

```
> using faustxml;
> let name,descr,version,in,out,group =
>   faustxml::info "examples/basic/freeverb.dsp.xml";
> name,descr,version,in,out;
"freeverb","freeverb -- a Schroeder reverb","1.0",2,2
> using system;
> do (puts.str) $ faustxml::pcontrols group;
faustxml::hslider ("freeverb/damp",0.5,0.0,1.0,0.025)
faustxml::hslider ("freeverb/roomsize",0.5,0.0,1.0,0.025)
faustxml::hslider ("freeverb/wet",0.3333,0.0,1.0,0.025)
```

pd-pure: Pd loader for Pure scripts

Version 0.10, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

This plugin lets you write external **Pd** control objects in Pure. **Pure** is a functional programming language based on the term rewriting calculus.

Please note that the present version of the module is still somewhat experimental, but it seems to work fairly well at least with Pure versions 0.34 and later. In particular, note that Pure is a *compiled* language and thus there are some inevitable latencies at startup, when the embedded Pure interpreter loads your Pure scripts and compiles them on the fly. However, once the scripts have been compiled, they are executed very efficiently.

1 Copying

Copyright (c) 2009 by Albert Graef. pd-pure is distributed under a 3-clause BSD-style license, please see the included COPYING file for details.

2 Installation

MS Windows users please see [pd-pure on Windows](#) below.

Get the latest source from <http://pure-lang.googlecode.com/files/pd-pure-0.10.tar.gz>.

Usually, `make` && `sudo make install` should do the trick. This will compile the external (you need to have GNU make, Pd and Pure installed to do that) and install it in the `lib/pd/extra/pure` directory.

The Makefile tries to guess the installation prefix under which Pd is installed. If it guesses wrong, you can tell it the right prefix with `make prefix=/some/path`. Or you can specify the exact path of the `lib/pd` directory with `make pdlibdir=/some/path`; by default the Makefile assumes `$(prefix)/lib/pd`.

The Makefile also tries to guess the host system type and Pure version, and set up some platform-specific things accordingly. If this doesn't work for your system then you'll have to edit the Makefile accordingly.

After installation, you still have to tell Pd to load the Pure external at startup, either with the `-lib` option (`pd -lib pure`), or by specifying `pure` in the File/Startup options. (This setting can be saved so that the Pure loader is always available when you run Pd.)

2.1 pd-pure on Windows

There's a binary package in MSI format available at the Pure website. In addition, you'll need the Pure interpreter, and a version of Pd which has been compiled with `mingw`; for your convenience, we provide a package for that as well. You can find all packages on the Download page at <http://pure-lang.googlecode.com>. For instance, at the time of this writing, the available packages are:

- <http://pure-lang.googlecode.com/files/pure-0.45.msi>
- <http://pure-lang.googlecode.com/files/pd-mingw-0.42-5.msi>
- <http://pure-lang.googlecode.com/files/pd-pure-0.10.msi>

Make sure that you get the latest versions of these packages. After installing the packages, you still have to configure Pd to load the pure external at startup (see above).

Please note that our binary `pd-pure` package does *not* work with the official Windows binaries of Pd available at <http://crca.ucsd.edu/~msp/software.html>. Apparently, some `mingw`-compiled plugins fail to load in a Pd version compiled with `MSVC` because of version mismatches in required libraries. Unfortunately, `pd-pure` is one of those (`pthreadGC2.dll` seems to be the main culprit here). So you'll have to use the `mingw`-compiled Pd version we provide, or roll your own.

3 Usage

See the `pure-help.pd` patch for a few examples. Basically, to implement a Pd object named `foo`, you have to supply a Pure script named `foo.pure` which defines a function `foo` and anything else that it might need. Put that script in the same directory as the Pd patch in which you want to use the `foo` object (or anywhere on Pd's search path). The `foo` function gets evaluated at object creation time, receiving any additional parameters the object is created with. The resulting Pure expression (which can also just be `foo` itself) becomes the object executed at runtime, by passing Pd messages from the inlets as parameters, and routing the function results to the outlets of the object.

Pd messages are translated to corresponding Pure expressions and vice versa in a straight-forward fashion. Special support is provided for converting between the natural Pd and Pure representations of floating point numbers, symbols and lists. The following table summarizes the available conversions.

Message Type	Pd	Pure
symbol	foo	foo
string	a&b	"a&b"
float	float 1.23	1.23
list	list 1 2 3	[1.0,2.0,3.0]
other	foo a 2 3	foo a 2.0 3.0

Note that Pd symbols which are no valid Pure symbols become strings in Pure. Conversely, both symbols and strings in Pure are mapped to corresponding Pd symbols. Pure (machine) integers and floating point values both become float messages in Pd. Pd list messages are translated to Pure list values, while other aggregate messages are mapped to Pure applications (and vice versa).

3.1 Simple Objects

By default, a Pure object has just one inlet and one outlet and thus acts like a simple function with no internal state. For instance, the following object accepts Pd float messages and adds 5 to each received value:

```
add5 x = x+5;
```

In the Pd patch each [add5] object then has a single inlet supplying parameters and a single outlet for results of the add5 function.

3.2 Creation Arguments

You can parameterize an object with creation arguments, which are passed to the Pure function at object creation time. For instance:

```
add x y = x+y;
```

This object can then be invoked, e.g., as [add 5] in the Pd patch to supply the needed creation argument x.

3.3 The [pure] Object

For simple kinds of objects like these, the Pure loader also provides the generic [pure] object as a quick means to create Pure objects without actually preparing a script file. The creation parameter of [pure] is the object function. This can be a predefined Pure function, or you can define it on the fly in a with clause.

For instance, [pure succ] uses the predefined Pure function succ which adds 1 to its input, while the object [pure add 5 with add x y = x+y end] produces the same results as the [add 5] object defined using a separate add.pure script in the previous section. You can also generate constant values that way. E.g., the object [pure cst 1.618] responds to any message (such as bang) by producing the constant value 1.618, while the object [pure cst [1..10]] yields the constant list containing the numbers 1..10.

3.4 Configuring Inlets and Outlets

To create an object with multiple inlets and outlets, the object creation function must return the desired numbers of inlets and outlets, along with a second function to be applied at runtime, as a tuple `n,m,foo`. The input arguments to the runtime function as well as the corresponding function results are then encoded as pairs `k,val` where `k` denotes the inlet or outlet index. (Note that the `k` index is provided only if there actually is more than one inlet. Also, the outlet index is assumed to be zero if none is specified, so that it can be omitted if there's only one outlet.)

For instance, the following object, invoked as `[cross]` in the Pd patch, has two inlets and two outlets and routes messages from the left inlet to the right outlet and vice versa:

```
cross = 2,2,cross with cross (k,x) = (1-k,x) end;
```

You can also emit multiple messages, possibly to different outlets, in one go. These must be encoded as lists of values or `index,value` pairs, which are emitted in the order in which they are written. E.g., the following `[fan]` object implements an “n-fan” which routes its input to `n` outlets simultaneously:

```
fan n = 1,n,fan n with fan n x = reverse [k,x | k = 0..n-1] end;
```

(Note that, because of the use of the `reverse`, the `n` outlets are served in right-to-left order here. This is not strictly necessary, but matches the Pd convention.)

Another example is the following `[dup]` object with a single inlet and outlet, which just sends out each received message twice:

```
dup x = [x,x];
```

Note that if you want to output a *real* list value to an outlet, you'll either have to specify the outlet index explicitly, or enclose the list in an extra pair of brackets, since otherwise the list elements will be sent as separate messages instead (like in the `dup` example). Thus, to output the list `[x,x]` literally, rather than sending `x` twice, use the following code:

```
dup2 x = [[x,x]]; // or: 0,[x,x]
```

An object can also just “swallow” messages and generate no output at all. To these ends, make the object return either an empty list `[]` or the empty tuple `()`. For instance, the following object `[echo]` implements a sink which just prints received messages on standard output, which is useful for debugging purposes:

```
using system;
echo x = () when puts (str x) end;
```

You could also implement this object as follows, by just removing the superfluous outlet (in this case all return values from the function will be ignored anyway):

```
using system;
echo = 1,0,puts.str;
```


3.5 Local State

Local state can be kept in Pure reference values. For instance, the following [counter] object produces the next counter value when receiving a bang message:

```
nonfix bang;  
counter = next (ref 0) with  
  next r bang = put r (get r+1);  
  next _ _    = () otherwise;  
end;
```

4 Advanced Features

4.1 Asynchronous Messages

pd-pure provides a simple asynchronous messaging facility which allows a Pure object to schedule a message to be delivered to itself later. This is useful for implementing all kinds of delays and, more generally, any kind of object which, once triggered, does its own sequencing of output messages.

To these ends, the object function may return a special message of the form `pd_delay t msg` (either by itself or as an element of a result list) to indicate that the message `msg` should be delivered to the object function after `t` milliseconds (where `t` is either a machine int or a double value). After the prescribed delay the object function will then be invoked on the given message, and the results of this call are processed as usual (routing messages to outlets and/or scheduling new timer events in response to further `pd_delay` messages). Note that if the delay is zero or negative, the message is scheduled to be delivered immediately.

For instance, a simple kind of delay object can be implemented in Pure as follows:

```
mydelay _ (alarm msg) = msg;  
mydelay t msg = pd_delay t (alarm msg) otherwise;
```

The desired delay time is specified as a creation argument. The first equation handles messages of the form `alarm msg`; the action is to just output the delayed message given by the `msg` argument. All other input messages are scheduled by the second equation, which wraps the message in an `alarm` term so that it gets processed by the first equation when it is delivered.

Note that pd-pure only allows you to schedule a single asynchronous event per call of the object function. Thus, if the `mydelay` object above receives another message while it is still waiting for the previous one to be delivered, the old timer is cancelled and the new one is scheduled instead; this works like Pd's builtin `delay` object.

Moreover, scheduling a new event at an infinite (or nan) time value cancels any existing timer. (Note that you still have to specify the `msg` parameter, but it will be ignored in this case.) We can use this to equip our `mydelay` object with a stop message as follows:

```
nonfix stop;
mydelay _ (alarm msg) = msg;
mydelay _ stop = pd_delay inf ();
mydelay t msg = pd_delay t (alarm msg) otherwise;
```

More elaborate functionality can be built on top of the basic timer facility. The following example shows how to maintain a timed message queue in a Pure list, in order to implement a simple delay line similar to Pd's builtin pipe object. Here we also employ the `pd_time` function, which is provided by the Pure loader so that Pure scripts can access the current logical Pd time in milliseconds (see [Programming Interface](#) below). This is convenient if we need to deal with absolute time values, which we use in this example to keep track of the times at which messages in the queue are to be delivered:

```
extern double pd_time();
mypipe t = process (ref []) with
  process q () = case dequeue q of
    x,(t,-):_ = [x,pd_delay (t-pd_time) ()];
    x,_ = x;
  end;
  process q x = enqueue q x $$ pd_delay t () if null (get q);
               = enqueue q x $$ () otherwise;
  enqueue q x = put q $ get q + [(pd_time+t,x)];
  dequeue q   = x,put q xs when (_,x):xs = get q end;
end;
```

4.2 Reading and Writing Audio Data

At present, `pd-pure` doesn't support audio objects. However, it is possible to transfer audio data between Pd and Pure by means of the `pd_getbuffer` and `pd_setbuffer` routines. In Pure land, the audio data is represented as a vector of floating point values. Please see [Programming Interface](#) below for a closer description of the provided routines.

For instance, here is a `randomwave` object which fills a Pd array (whose name is given as the creation argument) with random values in response to a bang message:

```
extern int pure_random() = random;
randf = random/0x7fffffff;

extern int pd_getbuffersize(char *name);
extern void pd_setbuffer(char *name, expr* x);

nonfix bang;

randomwave name = 1,0,process with
  process bang = pd_setbuffer name {randf | i = 1..nsamples};
  nsamples     = pd_getbuffersize name;
end;
```

5 Programming Interface

The Pure loader also provides a number of interface routines which can be called by Pure scripts running in the Pd environment.

```
extern char *pd_version_s();
```

Returns the Pd version number as a string. Note that this routine will only be available when a script is running inside Pd, so you can quickly check if that's the case as follows:

```
let ok = stringp $ eval "extern char *pd_version_s(); pd_version_s;";
```

The ok variable will then be true iff the script is running inside Pd.

```
extern char *pd_libdir_s();
```

Returns the Pd library dir (as determined at compile time). This is useful if your Pure scripts need to access files in that directory.

```
extern char *pd_path_sl();
```

Returns the Pd path (set in the File/Path dialog or via the -path command line option) as a list of directory names. This is useful if your Pure scripts need to locate files on the Pd search path.

```
extern void pd_post(char *s);
```

Posts a message in the Pd main window. A trailing newline is added automatically. This provides an alternative interface to Pd's `post()` function which cannot be used in Pure because it is a `printf`-style routine expecting a variable number of arguments.

```
extern double pd_time();
```

Retrieves the current Pd time as a double value in milliseconds, which is useful, in particular, when used in conjunction with the asynchronous message facility described under [Asynchronous Messages](#).

```
extern expr* pd_getbuffer(char *name);
```

```
extern void pd_setbuffer(char *name, expr* x);
```

```
extern int pd_getbuffersize(char *name);
```

```
extern void pd_setbuffersize(char *name, uint32_t sz);
```

Routines to access the Pd array (sample buffer) with the given name. These functions can be used to transfer audio data between Pd and Pure scripts; see [Reading and Writing Audio Data](#) above for an example.

`pd_getbuffersize` and `pd_setbuffersize` gets or sets the size of the given buffer, respectively.

`pd_getbuffer` reads the contents of the buffer and returns it as a Pure vector (or fails if the array with the given name doesn't exist).

`pd_setbuffer` sets the contents of the buffer from the given Pure vector `x`. If the size of the vector exceeds the size of the buffer, the former is truncated. Conversely, if the size of the buffer exceeds the size of the Pure vector, the trailing samples are unaffected. **NOTE:** The second argument of `pd_setbuffer` can also be a pair `(i, x)` denoting an offset `i` into the array at which the sample data is to be written, so that this routine allows you to overwrite any part of the array.

6 Interactive Facilities

6.1 Livecoding Support

Livecoding means that you can quickly reload your Pure scripts after changes while the Pd patch keeps running. The Pure loader provides some experimental support for this by means of the special `pure-runtime` object. Right now, this isn't really as "live" as we'd like it to be, because of the considerable delays caused by scripts being recompiled on the fly when they are loaded. However, there's hope that this facility becomes more useful as Pure's LLVM-based JIT compiler and computers become faster.

Sending a bang to the `pure-runtime` object tells the plugin to reload all loaded scripts and update the Pure objects in your patch accordingly. The object also provides two outlets to deal with the inevitable latencies caused by the compilation process. The right outlet is banged when the compilation starts and the left outlet gets a bang when the compilation is finished, so that a patch using this facility can respond to these events in the appropriate way (e.g., disabling output during compilation).

Note that there can be any number of `pure-runtime` objects in a patch, which will all refer to the same instance of the Pure interpreter.

Also please note that the number of inlets and outlets of Pure objects never changes after reloading scripts. (Pd does not support this through its API right now.) Thus by editing and reloading the Pure scripts you can change the functionality of existing Pure objects in a running patch, but not their interfaces.

6.2 Remote Control

The distribution also includes an abstraction `pure-remote.pd` which you can include in your patch to enable live coding, as well as remote control of the patch through the `pdsend` program. Sending a bang causes a reload of all scripts, which can also be triggered directly by clicking the bang control of the abstraction. The bang control also provides visual feedback indicating whether the compilation is still in progress. Messages are also routed through the embedded `pure-runtime` object using the single inlet and the two outlets of the abstraction, so that `pure-remote` can also be controlled from within the patch itself.

For added convenience, the `pure-runtime` and `pure-remote` objects also accept any other message of the form `receiver message` and will route the given message to the given receiver. This is intended to provide remote control of various parameters in patches. For instance, by having `pdsend` send a `play 0` or `play 1` message, one might implement a single playback control, provided that your patch includes an appropriate receiver (often a GUI object). See the `pure-help.pd` patch for an example.

To make these features available in Emacs, there's an accompanying elisp program (`pure-remote.el`) which contains some convenient keybindings for the necessary `pdsend` invocations, so that you can operate the `pure-remote` patch with simple keystrokes directly from the text editor. Please see the `pure-remote.el` file for more details. Currently `pure-remote.el` implements the following "standard" keybindings:

C-C C-X	Reload	Sends a bang message, causing scripts to be reloaded.
C-C C-M	Message	Prompts for a message and sends it to pure-remote.
C-C C-S	Play	Sends a <code>play 1</code> message.
C-C C-T	Stop	Sends a <code>play 0</code> message.
C-C C-G	Restart	Sends a <code>play 0</code> message followed by <code>play 1</code> .
C-C C-I	Dsp On	Sends a <code>pd dsp 1</code> , which enables audio processing.
C-C C-O	Dsp Off	Sends <code>pd dsp 0</code> , which disables audio processing.

Of course you can easily add more like these, just have a look at how the keybindings are implemented in `pure-remote.el` and create your own in an analogous fashion. Combining `pure-remote.el` with the Emacs Pure programming mode gives you a nice interactive environment for developing `pd-pure` applications.

pure-audio

Version 0.1, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

This is a digital audio interface for the Pure programming language. It currently includes the following modules:

- `audio.pure`: A PortAudio wrapper which gives you portable access to realtime audio input and output on a variety of different host APIs. This uses the well-known PortAudio library by Ross Bencina, Phil Burk et al, see <http://www.portaudio.com/>.
- `fftw.pure`: Compute real-valued FFTs of audio signals using Matteo Frigo's and Steven G. Johnson's portable and fast FFTW library ("Fastest Fourier Transform in the West").
- `sndfile.pure`: Reading and writing audio files in various formats. This is a fairly straightforward wrapper for Erik de Castro Lopo's libsndfile library, see <http://www.mega-nerd.com/libsndfile/>.
- `samplerate.pure`: Perform sample rate conversion on audio data. This uses another of Erik's excellent libraries, libsamplerate (a.k.a. SRC), see <http://www.mega-nerd.com/SRC/>.
- `realtime.pure`: A little utility module which provides access to realtime scheduling to Pure programs. You may need this for low-latency realtime audio applications.

Documentation still needs to be written, so for the time being please read the source modules listed above and have a look at the examples provided in the distribution.

1 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-audio-0.1.tar.gz>.

You need to have libportaudio (v19), libsndfile (1.x), libsamplerate (0.1.x) and libfftw3 (3.x) installed on your system. Any fairly recent version of these libraries should do. For the realtime module you also need a POSIX threads library (libpthread) with the POSIX realtime thread extension; Linux, OSX and other Un*x systems should offer this. Windows probably needs some work, so by default the realtime module isn't built on Windows right now.

The Pure wrappers contained in the distribution are for 64 bit POSIX systems. If you're running a 32 bit system, or Windows, then you should regenerate them using 'make generate'. This requires the header files `portaudio.h`, `samplerate.h` and `sndfile.h` (and `pure-gen`, of course). If you do this, check the `includedir` variables defined in the Makefiles, these need to point to the directories where the corresponding header files are to be found (the default is `/usr/include`).

Then just run 'make' to compile the package. If you're lucky and everything compiles smoothly, you can install with 'sudo make install'.

If you're not so lucky, you can get help on the Pure mailing list, see <http://groups.google.com/group/pure-lang>.

2 License

pure-audio is Copyright (c) 2010 by Albert Graef, licensed under the 3-clause BSD license, see the `COPYING` file for details.

Please note that if you're using these modules, you're also bound by the license terms of the PortAudio, libsamplerate and libsndfile libraries they are based on, see the corresponding sources and websites for details.

pure-faust: Run Faust programs in Pure scripts

Version 0.2, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

This module lets you load and run Faust-generated signal processing modules in [Pure](#). [Faust](#) (an acronym for Functional AUdio STreams) is a functional programming language for real-time sound processing and synthesis developed at [Grame](#) and distributed as GPL'ed software.

Note: This module is now largely obsolete as there's much better direct support for Faust interoperability in the Pure core (as of Pure 0.45), including the ability to inline Faust code in Pure programs. This requires `faust2` (the latest Faust incarnation from Faust's git repository), however, so this package is still provided for those who are stuck with older faust versions.

1 Copying

Unless explicitly stated otherwise, this software is Copyright (c) 2009 by Albert Graef. Please also see the source for the copyright and license notes pertaining to individual source files.

`pure-faust` is free software: you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

`pure-faust` is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

2 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-faust-0.2.tar.gz>.

Binary packages can be found at <http://pure-lang.googlecode.com/>. To install from source, run the usual `make` && `sudo make install`. This requires Pure, of course. The Makefile tries to guess the installation prefix under which Pure is installed. If it guesses wrong, you can tell it the right prefix with `make prefix=/some/path`. Or you can specify the exact path of the `lib/pure` directory with `make libdir=/some/path`; by default the Makefile assumes `$(prefix)/lib/pure`. The Makefile also tries to guess the host system type and set up some platform-specific things accordingly. If this doesn't work for your system then you'll have to edit the Makefile accordingly.

The Faust compiler is not required to compile this module, but of course you'll need it to build the examples in the `examples` subdirectory and to compile your own Faust sources. You'll need Faust 0.9.9.6 or later.

To compile Faust programs for use with this module, you'll also need the `pure.cpp` architecture file. This should be included in recent Faust releases. If your Faust version doesn't have it yet, you can find a suitable version of this file in the `examples` folder. Simply copy the file to your Faust library directory (usually `/usr/local/lib/faust` or similar) or the directory holding the Faust sources to be compiled, and you should be set. (Note that `pure.cpp` is distributed under a simple all-permissive license, so you can use it for pretty much any purpose, including commercial applications.)

3 Usage

Once Faust and this module have been installed as described above, you should be able to compile a Faust DSP to a shared module loadable by `pure-faust` as follows:

```
$ faust -a pure.cpp -o mydsp.cpp mydsp.dsp
$ g++ -shared -o mydsp.so mydsp.cpp
```

Note that, by default, Faust generates code which does all internal computations with single precision. You can add the `-double` flag to the Faust command in order to use double precision instead. (In either case, all data will be represented as doubles on the Pure side.)

Also note that the above compile command is for a Linux or BSD system using `gcc`. Add `-fPIC` for 64 bit compilation. For Windows compilation, the output filename should be `mydsp.dll` instead of `mydsp.so`; on Mac OSX, it should be `mydsp.dylib`. There's a Makefile in the `examples` folder which automates this process.

Once the module has been compiled, you can fire up the Pure interpreter and load the DSP as follows:

```
> using faust;
> let dsp = faust_init "mydsp" 48000;
> dsp;
#<pointer 0xf09220>
```

The `faust_init` function loads the `"mydsp.so"` module (the `.so` suffix is supplied automatically) and returns a pointer to the Faust DSP object which can then be used in subsequent operations. The second parameter of `faust_init`, 48000 in this example, denotes the sample

rate in Hz. This can be an arbitrary integer value which is available to the hosted DSP (it's up to the DSP whether it actually uses this value in some way). The sample rate can also be changed on the fly with the `faust_reinit` function:

```
> faust_reinit dsp 44100;
```

It is also possible to create copies of an existing DSP with the `faust_clone` function, which is quite handy if multiple copies of the same DSP are needed (a case which commonly arises when implementing polyphonic synthesizers):

```
> let dsp2 = faust_clone dsp;
```

When you're done with a DSP, you can invoke the `faust_exit` function to unload it (this also happens automatically when a DSP object is garbage-collected):

```
> faust_exit dsp2;
```

Note that after invoking this operation the DSP pointer becomes invalid and must not be used any more.

The `faust_info` function can be used to determine the number of input/output channels as well as the "UI" (a data structure describing the available control variables) of the loaded DSP:

```
> let n,m,ui = faust_info dsp;
```

To actually run the DSP, you'll need two buffers capable of holding the required number of audio samples for input and output. For convenience, the `faust_compute` routine lets you specify these as Pure double matrices. `faust_compute` is invoked as follows:

```
> faust_compute dsp in out;
```

Here, `in` and `out` must be double matrices of the same row size (which corresponds to the number of samples which will be processed) which have at least `n` or `m` rows, respectively (corresponding to the number of input and output channels of the Faust DSP). The `out` matrix will be modified in-place and also returned as the result of the call.

Some DSPs (e.g., synthesizers) only take control input without processing any audio input; others (e.g., pitch detectors) might produce just control output without any audio output. In such cases you can just specify an empty `in` or `out` matrix, respectively. For instance:

```
> faust_compute dsp {} out;
```

Most DSPs take additional control input. The control variables are listed in the "UI" component of the `faust_info` return value. For instance, suppose that there's a gain parameter listed there, it might look as follows:

```
> controls ui!0;  
hslider #<pointer 0x12780a4> ("gain",1.0,0.0,10.0,0.1)
```

The constructor itself denotes the type of control, which matches the name of the Faust builtin used to create the control (see the Faust documentation for more details on this).

The second parameter is a tuple which indicates the arguments the control was created with in the Faust program. The first parameter is a C `double*` which points to the current value of the control variable. You can inspect and change this value with the `get_double` and `put_double` routines available in the Pure prelude. (Note that changes of control variables only take effect between different invocations of `faust_compute`.) Example:

```
> let gain = control_ref (controls ui!0);
> get_double gain;
1.0
> put_double gain 2.0;
()
> faust_compute dsp in out;
```

Output controls such as `hbargraph` and `vbargraph` are handled in a similar fashion, only that the Faust DSP updates these values for each call to `faust_compute` and Pure scripts can then read the values with `get_double`.

Ok, I hope that this is enough to get you started. Further examples can be found in the `examples` subdirectory. Please send bug reports, patches, comments and suggestions to: Dr.Graef@t-online.de.

4 Acknowledgements

Many thanks to Yann Olarey at Grame, the principal author of Faust!

pure-liblo

Version 0.5, November 14, 2010

Albert Gräf <Dr.Graef@t-online.de>

1 Copying

Copyright (c) 2009 by Albert Graef.

pure-liblo is free software: you can redistribute it and/or modify it under the terms of the GNU Lesser General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

pure-liblo is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU Lesser General Public License for more details.

You should have received a copy of the GNU Lesser General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

2 Description

This is a quick and dirty Pure wrapper for the [liblo](#) library by Steve Harris and others, which implements Berkeley's [Open Sound Control](#) (OSC) protocol.

OSC is a protocol for exchanging data between multimedia devices and software across the network (TCP, UDP and UNIX domain sockets are supported as the transport layer). It is also useful as a general communication mechanism for both hard- and software. In difference to the plain socket interface (on which it builds), OSC provides you with an efficient means to send around binary data packets along with the corresponding type and timing information, which makes it well-suited for both realtime and non-realtime applications.

The OSC protocol is [standardized](#) and is supported by an abundance of different [implementations](#), which includes controller hardware of all sorts and computer music software like CSound, Pd and SuperCollider. Lots of implementations exist for different programming languages. liblo aims to provide a lightweight and ubiquitous OSC implementation for the C programming language.

The `lo.pure` module provides a fairly straight wrapper of the C library. A more high-level and Purified interface is available in `osc.pure`. Most of the time, you'll want to use the latter for convenience, but if you need utmost flexibility then it is worth having a look at `lo.pure`, too.

- Get the latest source from <http://pure-lang.googlecode.com/files/pure-liblo-0.5.tar.gz>.
- To install, run `make` and `sudo make install`. This will try to guess your Pure installation directory; if it guesses wrong, you can set the `prefix` variable accordingly, see the Makefile for details.
- You can also regenerate the wrapper by running `make generate`; this requires the `pure-gen` utility and the `liblo` headers. The present version was generated from `liblo` 0.26. If your `liblo` version differs from that then it's always a good idea to run `make generate`.
- Have a look at `lo.pure` and `osc.pure` for a description of the API provided to Pure programmers.
- The `examples` folder contains some Pure code which illustrates how to use these modules.

pure-midi

Version 0.1, November 14, 2010

Albert Graef <Dr.Graef@t-online.de>

This is a MIDI interface for the Pure programming language (Pure 0.45 or later is required). It includes the following modules:

- `midi.pure`: A PortMidi/PortTime wrapper which gives you portable access to realtime MIDI input and output. This uses PortMidi (by Roger B. Dannenberg et al) from the PortMedia project, see <http://portmedia.sourceforge.net/>.
- `midifile.pure`: Reading and writing standard MIDI files. This is based on David G. Slomin's light-weight midifile library, which comes bundled with the pure-midi sources.

Documentation still needs to be written, so for the time being please read the source modules listed above and have a look at the examples provided in the distribution.

1 Installation

Get the latest source from <http://pure-lang.googlecode.com/files/pure-midi-0.1.tar.gz>.

You need to have the PortMidi library installed on your system. This release was tested with PortMidi 2.00 (I recommend using the svn version of PortMidi, since it fixes some 64 bit compilation problems). If you have to use some earlier PortMidi version then you may have to fiddle with `portmidi.pure` and/or `midi.pure` to make it work. (You can also just regenerate the wrapper by copying `portmidi.h` and `porttime.h` from your PortMidi installation to the pure-midi source directory and running 'make generate'. This requires pure-gen. See the toplevel Makefile for details.)

Run 'make' to compile the package. If you're lucky and everything compiles smoothly, you can install with 'sudo make install'.

If you're not so lucky, you can get help on the Pure mailing list, see <http://groups.google.com/group/pure-lang>.

NOTE: You may also want to install the related pure-audio package. In particular, pure-audio also provides `realtime.pure`, a little utility module which gives Pure programs access to realtime scheduling.

2 License

pure-midi is Copyright (c) 2010 by Albert Graef, licensed under the 3-clause BSD license, see the COPYING file for details.

For convenience, I've bundled some (BSD-licensed or compatible) source files from other packages with this release. portmidi.h and porttime.h are from PortMidi 2.00 (<http://portmedia.sourceforge.net/>) which is

Copyright (c) 1999-2000 Ross Bencina and Phil Burk Copyright (c) 2001-2006 Roger B. Dannenberg

midifile.c and midifile.h in the midifile subdirectory are from "Div's midi utilities" (<http://public.sreal.com:8000/~div/midi-utilities/>) which is

Copyright (c) 2003-2006 David G. Slomin

Please see portmidi.h and midifile.h for the pertaining copyrights and license conditions.

Index

Symbols

`*`, 133
`*` (bag), 159
`*` (dictionary), 155
`*` (gsl::matrix), 345
`*` (interval), 388
`*` (rational complex), 399
`*` (rational), 376
`*` (set), 159
`+`, 133
`+` (bag), 159
`+` (dictionary), 155
`+` (gsl::matrix), 345
`+` (interval), 388
`+` (list), 117
`+` (rational complex), 399
`+` (rational), 376
`+` (set), 159
`+` (string), 121
`-`, 133
`-` (bag), 159
`-` (dictionary), 155
`-` (gsl::matrix), 345
`-` (interval), 388
`-` (rational complex), 399
`-` (rational), 376
`-` (set), 159
`.,` 114
`.*`, 345
`..`, 117
`./`, 345
`.^`, 345
`/`, 133
`/` (gsl::matrix), 345
`/` (interval), 388
`/` (rational complex), 399
`/` (rational), 376
`;`, 116
`==`, 133
`==` (interval), 389
`==` (rational complex), 399
`==` (rational), 377
`===`, 135
`#`, 116
`#` (array), 151
`#` (bag), 159
`#` (dictionary), 155
`#` (heap), 153
`#` (interval), 391
`#` (matrix), 125
`#` (record), 132
`#` (set), 159
`#` (string), 122
`$`, 114
`%`, 148, 377
`%` (interval), 388
`%` (rational complex), 399
`&&`, 133
`__locals__`, 136
`^`, 133
`^` (gsl::matrix), 345
`^` (rational), 376
`~`, 133
`~=`, 133
`~=` (interval), 389
`~=` (rational complex), 399
`~=` (rational), 377
`~==`, 135
`\\`, 345
`||`, 133
`>`, 133

> (rational), 377
 >=, 133
 >= (rational), 377
 >>, 133
 >> (gsl::matrix), 346
 <, 133
 < (rational), 377
 <=, 133
 <= (rational), 377
 <<, 133
 << (gsl::matrix), 346

A

abs, 134
 abs (complex), 147
 abs (interval), 391
 abs (rational complex), 399
 abs (rational), 378
 absdev, 349
 absdev_m, 349
 acos, 147
 acosh, 147
 add_first, 463
 add_last, 463
 add_next, 463
 add_prev, 463
 addr, 140
 after, 389
 airy_Ai, 329
 airy_Ai_deriv, 329
 airy_Ai_deriv_e, 329
 airy_Ai_deriv_scaled, 329
 airy_Ai_deriv_scaled_e, 329
 airy_Ai_e, 329
 airy_Ai_scaled, 329
 airy_Ai_scaled_e, 329
 airy_Bi, 329
 airy_Bi_deriv, 329
 airy_Bi_deriv_e, 329
 airy_Bi_deriv_scaled, 329
 airy_Bi_deriv_scaled_e, 330
 airy_Bi_e, 329
 airy_Bi_scaled, 329
 airy_Bi_scaled_e, 329
 airy_zero_Ai, 330
 airy_zero_Ai_deriv, 330
 airy_zero_Ai_deriv_e, 330

airy_zero_Ai_e, 330
 airy_zero_Bi, 330
 airy_zero_Bi_deriv, 330
 airy_zero_Bi_deriv_e, 330
 airy_zero_Bi_e, 330
 all, 119
 and, 133
 and (gsl::matrix), 346
 ans, 136
 any, 119
 append (array), 151
 applp, 135
 apply_stylesheet, 464
 arg, 147
 arg (rational complex), 399
 arithmetic sequences, 117
 arity, 136
 array, 151
 array2, 151
 arrayp, 151
 asctime, 164
 asin, 147
 asinh, 147
 atan, 147
 atan2, 147
 atanh, 147
 atomp, 135
 attribute declaration, 458
 attribute defaults, 459
 attrs, 461

B

bad_list_value, 113
 bad_matrix_value, 113
 bad_string_value, 113
 bad_tuple_value, 113
 bag, 159
 bagp, 159
 before, 389
 bernoulli_pdf, 354
 bessell_I0 (regular modified cylindrical), 332
 bessell_I0_e (regular modified cylindrical), 332
 bessell_I0_scaled (regular modified cylindrical), 332
 bessell_i0_scaled (regular modified spherical), 334

`bessel_I0_scaled_e` (regular modified cylindrical), 332
`bessel_i0_scaled_e` (regular modified spherical), 334
`bessel_I1` (regular modified cylindrical), 332
`bessel_I1_e` (regular modified cylindrical), 332
`bessel_I1_scaled` (regular modified cylindrical), 332
`bessel_i1_scaled` (regular modified spherical), 334
`bessel_i1_scaled_e` (regular modified spherical), 334
`bessel_i2_e` (regular modified spherical), 334
`bessel_i2_scaled` (regular modified spherical), 334
`bessel_il_scaled` (regular modified spherical), 334
`bessel_il_scaled_array` (regular modified spherical), 334
`bessel_il_scaled_e` (regular modified spherical), 334
`bessel_In` (regular modified cylindrical), 332
`bessel_In_array` (regular modified cylindrical), 332
`bessel_In_e` (regular modified cylindrical), 332
`bessel_In_scaled` (regular modified cylindrical), 332
`bessel_In_scaled_array` (regular modified cylindrical), 333
`bessel_In_scaled_e` (regular modified cylindrical), 333
`bessel_Inu` (regular modified), 334
`bessel_Inu_e` (regular modified), 334
`bessel_Inu_scaled` (regular modified), 334
`bessel_Inu_scaled_e` (regular modified), 334
`bessel_J0` (cylindrical), 332
`bessel_j0` (regular spherical), 333
`bessel_J0_e` (cylindrical), 332
`bessel_j0_e` (regular spherical), 333
`bessel_J1` (cylindrical), 332
`bessel_j1` (regular spherical), 333
`bessel_J1_e` (cylindrical), 332
`bessel_j1_e` (regular spherical), 333
`bessel_j2` (regular spherical), 333
`bessel_j2_e` (regular spherical), 333
`bessel_jl` (regular spherical), 333
`bessel_jl_array` (regular spherical), 333
`bessel_jl_e` (regular spherical), 333
`bessel_jl_stepped_array` (regular spherical), 333
`bessel_Jn` (cylindrical), 332
`bessel_Jn_array` (cylindrical), 332
`bessel_Jn_e` (cylindrical), 332
`bessel_Jnu` (regular), 334
`bessel_Jnu_e` (regular), 334
`bessel_K0` (irregular modified cylindrical), 333
`bessel_K0_e` (irregular modified cylindrical), 333
`bessel_K0_scaled` (irregular modified cylindrical), 333
`bessel_k0_scaled` (irregular modified spherical), 334
`bessel_K0_scaled_e` (irregular modified cylindrical), 333
`bessel_k0_scaled_e` (irregular modified spherical), 334
`bessel_K1` (irregular modified cylindrical), 333
`bessel_K1_e` (irregular modified cylindrical), 333
`bessel_K1_scaled` (irregular modified cylindrical), 333
`bessel_k1_scaled` (irregular modified spherical), 334
`bessel_K1_scaled_e` (irregular modified cylindrical), 333
`bessel_k1_scaled_e` (irregular modified spherical), 334
`bessel_k2_e` (irregular modified spherical), 334
`bessel_k2_scaled` (irregular modified spherical), 334
`bessel_kl_scaled` (irregular modified spherical), 334
`bessel_kl_scaled_array` (irregular modified spherical), 334
`bessel_kl_scaled_e` (irregular modified spherical), 334
`bessel_Kn` (irregular modified cylindrical), 333
`bessel_Kn_array` (irregular modified cylindrical), 333

`bessel_Kn_e` (irregular modified cylindrical), 333
`bessel_Kn_scaled` (irregular modified cylindrical), 333
`bessel_Kn_scaled_array` (irregular modified cylindrical), 333
`bessel_Kn_scaled_e` (irregular modified cylindrical), 333
`bessel_Knu` (irregular modified), 335
`bessel_Knu_e` (irregular modified), 335
`bessel_Knu_scaled` (irregular modified), 335
`bessel_Knu_scaled_e` (irregular modified), 335
`bessel_lnKnu` (irregular modified), 335
`bessel_lnKnu_e` (irregular modified), 335
`bessel_sequence_Jnu_e` (regular), 334
`bessel_Y0` (irregular cylindrical), 332
`bessel_y0` (irregular spherical), 333
`bessel_Y0_e` (irregular cylindrical), 332
`bessel_y0_e` (irregular spherical), 333
`bessel_Y1` (irregular cylindrical), 332
`bessel_y1` (irregular spherical), 333
`bessel_Y1_e` (irregular cylindrical), 332
`bessel_y1_e` (irregular spherical), 333
`bessel_y2` (irregular spherical), 333
`bessel_y2_e` (irregular spherical), 334
`bessel_y1` (irregular spherical), 334
`bessel_y1_array` (irregular spherical), 334
`bessel_y1_e` (irregular spherical), 334
`bessel_Yn` (irregular cylindrical), 332
`bessel_Yn_array` (irregular cylindrical), 332
`bessel_Yn_e` (irregular cylindrical), 332
`bessel_Ynu` (irregular), 334
`bessel_Ynu_e` (irregular), 334
`bessel_zero_J0` (regular), 335
`bessel_zero_J0_e` (regular), 335
`bessel_zero_J1` (regular), 335
`bessel_zero_J1_e` (regular), 335
`bessel_zero_Jnu` (regular), 335
`bessel_zero_Jnu_e` (regular), 335
`beta_P`, 356
`beta_pdf`, 354
`beta_Pinv`, 356
`beta_Q`, 356
`beta_Qinv`, 356
`bigint`, 134
`bigintp`, 135
`bigintvalp`, 149
`binomial_P`, 357
`binomial_pdf`, 354
`binomial_Q`, 357
`bivariate_gaussian_pdf`, 353
`blob`, 139
`blob_crc`, 140
`blob_size`, 140
`blobp`, 139
`byte_cstring`, 124
`byte_cstring_pointer`, 124
`byte_matrix`, 131
`byte_pointer`, 130
`byte_string`, 124
`byte_string_pointer`, 124
C
`calloc`, 141
`cat`, 119
`catmap`, 119
`cauchy_P`, 355
`cauchy_pdf`, 353
`cauchy_Pinv`, 355
`cauchy_Q`, 355
`cauchy_Qinv`, 355
`ceil`, 135
`ceil (rational)`, 384
`ceil_multiple`, 386
`ceye`, 344
`character arithmetic`, 122
`charp`, 135
`chars`, 123
`check_ptrtag`, 143
`children`, 461
`chisq_P`, 356
`chisq_pdf`, 354
`chisq_Pinv`, 356
`chisq_Q`, 356
`chisq_Qinv`, 356
`chr`, 122
`cis`, 147
`clausen`, 339
`clausen_e`, 339
`clear_sentry`, 141
`clearerr`, 167
`clock`, 163
`close`, 441

[closurep](#), 135
[cmatrix](#), 127
[cmatrixp](#), 128
[cmp](#), 378
[cmp_complexity](#), 381
[col](#), 128
[colcat](#), 129
[colcatmap](#), 129
[colrev](#), 130
[cols](#), 128
[colvector](#), 127
[colvectorp](#), 128
[colvectorseq](#), 127
[combinators](#), 114
[comp_simplify](#), 401
[complex](#), 147
[complex_colve_cubic](#), 327
[complex_float_matrix](#), 131
[complex_float_pointer](#), 130
[complex_matrix](#), 131
[complex_matrix_view](#), 131
[complex_pointer](#), 130
[complex_solve](#), 327
[complex_solve_quadratic](#), 327
[complexity_rel](#), 382
[complexp](#), 135
[compvalp](#), 149
[cones](#), 344
[conj](#), 148
[conj \(matrix\)](#), 129
[continued_fraction](#), 394
[continued_fraction_epsilon](#), 395
[continued_fraction_max_terms](#), 394
[convergents](#), 395
[cooked](#), 141
[cookedp](#), 141
[correlation](#), 350
[cos](#), 147
[cosh](#), 147
[coulomb_wave_CL_array](#), 340
[coulomb_wave_CL_e](#), 340
[coulomb_wave_F_array](#), 339
[coulomb_wave_FG_array](#), 340
[coulomb_wave_FG_e](#), 339
[coulomb_wave_FGp_array](#), 340
[coulomb_wave_sphF_array](#), 340
[coupling_3j](#), 340

[coupling_3j_e](#), 340
[coupling_6j](#), 340
[coupling_6j_e](#), 341
[coupling_9j](#), 341
[coupling_9j_e](#), 341
[covariance](#), 350
[covariance_m](#), 350
[create_format](#), 403
[create_function](#), 449
[cst](#), 114
[cstring](#), 124
[cstring_dup](#), 124
[cstring_list](#), 124
[cstring_vector](#), 124
[ctime](#), 164
[curry](#), 115
[curry3](#), 115
[cycle](#), 120
[cyclen](#), 120
[czeros](#), 344

D

[dawson](#), 341
[dawson_e](#), 341
[daylight](#), 164
[db_busy](#), 447
[db_error](#), 447
[dd_eval](#), 327
[dd_init](#), 327
[dd_talyor](#), 327
[debye_1](#), 342
[debye_1_e](#), 342
[debye_2](#), 342
[debye_2_e](#), 342
[debye_3](#), 342
[debye_3_e](#), 342
[debye_4](#), 342
[debye_4_e](#), 342
[debye_5](#), 342
[debye_5_e](#), 342
[debye_6](#), 342
[debye_6_e](#), 342
[delete \(bag\)](#), 160
[delete \(dictionary\)](#), 156
[delete \(record\)](#), 132
[delete \(set\)](#), 160
[delete_all \(bag\)](#), 160

[delete_all \(dictionary\)](#), 156
[delete_all \(set\)](#), 160
[delete_val \(dictionary\)](#), 156
[den](#), 149, 374
[den \(rational complex\)](#), 397
[diag](#), 128
[diagmat](#), 129
[dict](#), 155
[dictp](#), 155
[dilog](#), 343
[dilog_e](#), 343
[dim](#), 125
[dirichlet_lnpdf](#), 354
[dirichlet_pdf](#), 354
[discrete_free](#), 354
[discrete_pdf](#), 354
[discrete_preproc](#), 354
[disjoint](#), 390
[div](#), 133
[div \(gsl::matrix\)](#), 345
[div \(rational\)](#), 377
[div_mod_gauss](#), 400
[divide](#), 377
[dmatrix](#), 127
[dmatrixp](#), 128
[do](#), 119
[doc](#), 461
[doc_info](#), 460
[docp](#), 456
[double](#), 134
[double \(rational\)](#), 387
[double_matrix](#), 131
[double_matrix_view](#), 131
[double_pointer](#), 130
[doublep](#), 135
[dowith](#), 121
[dowith3](#), 121
[drop](#), 119
[dropwhile](#), 119
[dtd](#), 457

E

[element content](#), 459
[element declaration](#), 458
[emptyarray](#), 151
[emptybag](#), 159
[emptydict](#), 155

[emptyhbag](#), 159
[emptyhdict](#), 155
[emptyheap](#), 153
[emptyhmdict](#), 155
[emptyhset](#), 159
[emptymdict](#), 155
[emptyset](#), 159
[entity declaration](#), 458
[eq_cplx](#), 380
[errno](#), 162
[euclid_alg](#), 400
[euclid_gcd](#), 400
[eval](#), 137
[eval \(gsl::poly\)](#), 327
[evalcmd](#), 137
[evaluate_continued_fraction](#), 395
[exactp](#), 135, 375
[exactp \(interval\)](#), 388
[exec](#), 444
[exec dbp](#), 445
[execv](#), 166
[execve](#), 166
[execvp](#), 166
[exit](#), 140
[exp](#), 147
[expoential_pdf](#), 353
[exponential_P](#), 355
[exponential_Pinv](#), 355
[exponential_Q](#), 355
[exponential_Qinv](#), 355
[exppow_P](#), 355
[exppow_pdf](#), 353
[exppow_Q](#), 355
[ext_subset](#), 460
[eye](#), 344

F

[failed_cond](#), 113
[failed_match](#), 113
[false](#), 113
[farey](#), 383
[faustxml::control_args](#), 491
[faustxml::control_label](#), 491
[faustxml::control_type](#), 491
[faustxml::controlp](#), 491
[faustxml::controls](#), 491
[faustxml::info](#), 492

faustxml::pcontrols, 492
fclose, 167
fdist_P, 356
fdist_pdf, 354
fdist_Pinv, 356
fdist_Q, 356
fdist_Qinv, 356
fdopen, 167
feof, 167
ferror, 167
fflush, 167
fget, 167
fgets, 167
fileno, 167
filter, 119
finalize, 447
first (array), 151
first (bag), 160
first (dictionary), 155
first (heap), 153
first (set), 160
first (xml), 461
first_attr, 461
fix, 116
fixity, 136
flat_P, 355
flat_pdf, 353
flat_Pinv, 355
flat_Q, 355
flat_Qinv, 355
flip, 115
float_matrix, 131
float_pointer, 130
floor, 135
floor (rational), 384
floor_multiple, 386
fmt, 402
fnmatch, 170
foldl, 119
foldl1, 119
foldr, 119
foldr1, 119
fopen, 167
fork, 176
fprintf, 169
fputs, 167
frac, 135

frac (rational), 385
fraction, 385
fread, 167
free, 141
freopen, 167
fscanf, 169
fseek, 168
fstat, 170
ftell, 168
funp, 135
fwrite, 167

G

gamma_P, 355
gamma_pdf, 353
gamma_Pinv, 355
gamma_Q, 355
gamma_Qinv, 355
gaussian_P, 354
gaussian_pdf, 353
gaussian_Pinv, 355
gaussian_Q, 354
gaussian_Qinv, 355
gaussian_tail_pdf, 353
gcd, 134
gcd (rational), 378
gcd_gauss, 400
geometric_P, 357
geometric_pdf, 354
geometric_Q, 357
get, 146
get_byte, 141
get_double, 141
get_float, 141
get_int, 141
get_int64, 141
get_long, 141
get_pointer, 141
get_ptrtag, 143
get_sentry, 141
get_short, 141
get_string, 141
getopt, 176
getpid, 176
getppid, 176
gets, 167
gettimeofday, 163

glob, 170
gmtime, 164
gumbel1_P, 356
gumbel1_pdf, 354
gumbel1_Pinv, 357
gumbel1_Q, 356
gumbel1_Qinv, 357
gumbel2_P, 357
gumbel2_pdf, 354
gumbel2_Pinv, 357
gumbel2_Q, 357
gumbel2_Qinv, 357

H

hash, 134
hbag, 159
hbagp, 159
hdict, 155
hdictp, 155
head, 119
heap, 153
heapp, 153
hmdict, 155
hmdictp, 155
hset, 159
hsetp, 159
hydrogenicR, 339
hydrogenicR_1, 339
hydrogenicR_1_e, 339
hydrogenicR_e, 339
hypergeometric_P, 357
hypergeometric_pdf, 354
hypergeometric_Q, 357

I

id, 114
ieye, 344
im, 148
im (matrix), 129
im (rational), 396
imatrix, 127
imatrixp, 128
index, 119
index (string), 122
inexactp, 135, 375
inexactp (interval), 388
init, 120

insert (array), 151
insert (bag), 160
insert (dictionary), 156
insert (heap), 153
insert (record), 132
insert (set), 160
int, 134
int (rational), 385
int_matrix, 131
int_matrix_view, 131
int_pointer, 130
int_subset, 460
intcompvalp, 398
integerp, 135
interval, 387
interval_valp, 388
intervalp, 388
intinterval_valp, 388
intp, 135
intvalp, 149
iones, 344
is_blank_node, 462
is_db_ptr, 441
is_open, 441
is_stmt_ptr, 444
iterate, 120
iteraten, 120
iterwhile, 120
izeros, 344

J

join, 123
join_str_real_approx, 409
join_str_real_eng, 411
join_str_real_recur, 407

K

key, 116
keys (dictionary), 155
keys (record), 132
kill, 176
kurtosis, 349
kurtosis_m_sd, 349

L

lag1_autocorrelation, 349
lag1_autocorrelation_m, 349

- lambdap, 135
- landau_pdf, 353
- laplace_P, 355
- laplace_pdf, 353
- laplace_Pinv, 355
- laplace_Q, 355
- laplace_Qinv, 355
- last, 120
- last (array), 151
- last (bag), 160
- last (dictionary), 155
- last (set), 160
- last (xml), 461
- last_attr, 461
- lasterr, 138
- lcd, 134
- lcm (rational), 379
- least_cplx, 382
- less_cplx, 381
- less_eq_cplx, 381
- lexec, 445
- linear, 347
- linear_est, 347
- list, 118
- list (array), 151
- list (bag), 160
- list (dictionary), 155
- list (heap), 153
- list (matrix), 128
- list (set), 160
- list (string), 123
- list concatenation, 117
- list indexing, 118
- list size, 116
- list slicing, 118
- list2 (array), 151
- list2 (matrix), 128
- listmap, 120
- listnp, 135
- listp, 135
- lists, 116
- ln, 147
- lo_up, 388
- load_file, 460
- load_string, 460
- load_stylesheet, 464
- localtime, 164

- log, 147
- logarithmic_pdf, 354
- logistic_P, 356
- logistic_pdf, 354
- logistic_Pinv, 356
- logistic_Q, 356
- logistic_Qinv, 356
- lognormal_P, 355
- lognormal_pdf, 353
- lognormal_Pinv, 355
- lognormal_Q, 355
- lognormal_Qinv, 356
- lower, 387
- lstat, 169

M

- make_ptrtag, 144
- malloc, 141
- malloc_error, 113
- map, 120
- matcat, 129
- matrix, 126
- matrix (string), 123
- matrix dimensions, 125
- matrix size, 125
- matrixp, 135
- max, 134
- max (gsl::stats), 351
- max (rational), 380
- max_index, 351
- mdict, 155
- mdictp, 155
- mean, 349
- median_from_sorted_data, 351
- mediant, 382
- member (bag), 160
- member (dictionary), 155
- member (record), 132
- member (set), 160
- members (array), 151
- members (bag), 160
- members (dictionary), 155
- members (heap), 153
- members (set), 160
- members2 (array), 151
- min, 134
- min (gsl::stats), 351

- [min \(rational\)](#), 380
- [min_index](#), 351
- [minmax](#), 351
- [minmax_index](#), 351
- [mkarray](#), 151
- [mkarray2](#), 151
- [mkdict](#), 155
- [mkhdict](#), 155
- [mkhmdict](#), 155
- [mkmdict](#), 155
- [mktime](#), 164
- [mod](#), 133
- [mod \(gsl::matrix\)](#), 345
- [mod \(rational\)](#), 377
- [modulus](#), 377
- [more_cplx](#), 381
- [more_eq_cplx](#), 381
- [most_cplx](#), 382
- [mul](#), 347
- [mul_est](#), 347
- [multinomial_lnpdf](#), 354
- [multinomial_pdf](#), 354
- [multiply_e](#), 343
- [multiply_err_e](#), 343

N

- [n_div_gauss](#), 400
- [n_mod_gauss](#), 400
- [nanosleep](#), 163
- [nargs](#), 136
- [negative_binomial_P](#), 357
- [negative_binomial_pdf](#), 354
- [negative_binomial_Q](#), 357
- [new_doc](#), 459
- [next](#), 461
- [nmatrixp](#), 128
- [node info](#), 456
- [node_attr](#), 462
- [node_base](#), 462
- [node_content](#), 462
- [node_info](#), 462
- [node_path](#), 462
- [nodep](#), 456
- [norm_gauss](#), 399
- [not](#), 133
- [not \(gsl::matrix\)](#), 345
- [not_eq_cplx](#), 381

- [null](#), 116
- [null \(array\)](#), 151
- [null \(bag\)](#), 159
- [null \(dictionary\)](#), 155
- [null \(heap\)](#), 153
- [null \(matrix\)](#), 125
- [null \(set\)](#), 159
- [null \(string\)](#), 122
- [num](#), 149, 374
- [num \(rational complex\)](#), 397
- [num_den](#), 374
- [num_den \(rational complex\)](#), 397
- [num_den_gauss](#), 397
- [num_den_nat](#), 397
- [numberp](#), 135

O

- [ones](#), 344
- [open](#), 440
- [operators](#), 114
- [or](#), 133
- [or \(gsl::matrix\)](#), 346
- [ord](#), 122
- [out_of_bounds](#), 114

P

- [pack](#), 129
- [packed](#), 129
- [parent](#), 461
- [pareto_P](#), 356
- [pareto_pdf](#), 354
- [pareto_Pinv](#), 356
- [pareto_Q](#), 356
- [pareto_Qinv](#), 356
- [pascal_P](#), 357
- [pascal_pdf](#), 354
- [pause](#), 176
- [pclose](#), 167
- [perror](#), 162
- [pinv](#), 346
- [pointer](#), 134
- [pointer \(matrix\)](#), 130
- [pointer_cast](#), 144
- [pointer_tag](#), 144
- [pointer_type](#), 144
- [pointerp](#), 135
- [poisson_P](#), 357

poisson_pdf, 354
poisson_Q, 357
polar, 147
polar (rational complex), 399
popen, 167
pow, 134
pow (gsl::matrix), 346
pow (rational), 376
pred, 134
prefix_complexity_rel, 382
prep, 442
prev, 461
printf, 169
ptrtag, 143
put, 146
put_byte, 141
put_double, 141
put_float, 141
put_int, 141
put_int64, 141
put_long, 141
put_pointer, 141
put_short, 141
put_string, 141
puts, 167

Q

quantile_from_sorted_data, 351
quotient, 377

R

raise, 176
random, 146
random31, 146
random53, 146
rat_simplify, 383
ratcomp_simplify, 401
ratcompvalp, 398
ratinterval_valp, 388
rational, 148, 374
rational_approx_epsilon, 391
rational_approx_max_den, 393
rational_approx_epsilon, 392
rational_approx_max_den, 393
rational_interval_epsilon, 392
rational_interval_max_den, 393
rationalp, 135, 375

ratvalp, 149, 375
rayleigh_P, 355
rayleigh_pdf, 353
rayleigh_Pinv, 355
rayleigh_Q, 355
rayleigh_Qinv, 355
rayleigh_tail_pdf, 353
re, 148
re (matrix), 129
re (rational), 396
readdir, 170
realloc, 141
realp, 135
realvalp, 149
reciprocal, 377
record, 132
recordp, 132
rect, 147
redim, 129
reduce, 138
reduce_with, 138
ref, 145
refp, 146
reg, 172
reg_info, 172
reg_result, 172
regex, 171
regexg, 172
regexgg, 172
regexggs, 173
regexgs, 173
regs, 172
regsplit, 173
regplits, 173
regsub, 172
repeat, 120
repeatn, 120
replace, 463
reverse, 116
reverse (matrix), 130
rewind, 168
rmfirst (array), 151
rmfirst (bag), 160
rmfirst (dictionary), 156
rmfirst (heap), 153
rmfirst (set), 160
rmlast (array), 151

- [rmlast \(bag\)](#), 160
- [rmlast \(dictionary\)](#), 156
- [rmlast \(set\)](#), 160
- [root](#), 461
- [round](#), 135
- [round \(rational\)](#), 384
- [round_multiple](#), 386
- [round_multiple_unbiased](#), 386
- [round_multiple_zero_bias](#), 386
- [round_unbiased](#), 384
- [round_zero_bias](#), 384
- [roundfun](#), 384
- [row](#), 128
- [rowcat](#), 129
- [rowcatmap](#), 129
- [rowrev](#), 130
- [rows](#), 128
- [rowvector](#), 127
- [rowvectorp](#), 128
- [rowvectorseq](#), 127

S

- [same](#), 135
- [save_file](#), 460
- [save_result_file](#), 464
- [save_result_string](#), 464
- [save_string](#), 460
- [scanf](#), 169
- [scanl](#), 120
- [scanl1](#), 120
- [scanr](#), 120
- [scanr1](#), 120
- [sd](#), 349
- [sd_m](#), 349
- [sd_with_fixed_mean](#), 349
- [select](#), 461
- [sentry](#), 141
- [set](#), 159
- [set_errno](#), 162
- [set_node_attr](#), 462
- [setbuf](#), 168
- [setlocale](#), 162
- [setp](#), 159
- [setvbuf](#), 168
- [sgn](#), 134
- [sgn \(interval\)](#), 391
- [sgn \(rational\)](#), 378

- [short_matrix](#), 131
- [short_pointer](#), 130
- [sin](#), 147
- [sinh](#), 147
- [skew](#), 349
- [skew_m_sd](#), 349
- [sleep](#), 163
- [smatrix](#), 127
- [smatrixp](#), 128
- [solve_cubic](#), 327
- [solve_quadratic](#), 327
- [sort](#), 117
- [sort \(list\)](#), 120
- [sort \(matrix\)](#), 129
- [sort_vector_index](#), 363
- [spawnv](#), 166
- [spawnve](#), 166
- [spawnvp](#), 166
- [split](#), 123
- [split_str_real_eng](#), 412
- [split_str_real_recur](#), 408
- [sprintf](#), 169
- [sqrt](#), 147
- [srandom](#), 146
- [sscanf](#), 169
- [stack_fault](#), 113
- [stat](#), 169
- [str](#), 137
- [str_mixed](#), 404
- [str_real_approx_dp](#), 409
- [str_real_approx_eng](#), 410
- [str_real_approx_sci](#), 410
- [str_real_approx_sf](#), 409
- [str_real_recur](#), 406
- [str_vulgar](#), 404
- [str_vulgar_or_int](#), 404
- [strcat](#), 123
- [stream](#), 118
- [strerror](#), 162
- [strftime](#), 164
- [strictly_after](#), 389
- [strictly_before](#), 389
- [strictly_disjoint](#), 390
- [strictly_within](#), 389
- [strictly_without](#), 390
- [stride](#), 128
- [string](#), 123

- [string \(pointer\)](#), 124
- [string concatenation](#), 121
- [string indexing](#), 121
- [string size](#), 122
- [string slicing](#), 121
- [string_dup](#), 124
- [string_list](#), 124
- [string_vector](#), 124
- [stringp](#), 135
- [strings](#), 121
- [strptime](#), 164
- [strs_real_approx_dp](#), 409
- [strs_real_approx_eng](#), 411
- [strs_real_approx_sci](#), 410
- [strs_real_approx_sf](#), 410
- [strs_real_recur](#), 407
- [stylesheetp](#), 456
- [subdiag](#), 128
- [subdiagmat](#), 129
- [submat](#), 128
- [substr](#), 122
- [succ](#), 134
- [supdiag](#), 128
- [supdiagmat](#), 129
- [sval_real_eng](#), 411
- [sval_real_recur](#), 408
- [svd](#), 346
- [svd_jacobi](#), 346
- [svd_mod](#), 346
- [svd_solve](#), 346
- [system](#), 166

T

- [tail](#), 120
- [take](#), 120
- [takewhile](#), 120
- [tan](#), 147
- [tanh](#), 147
- [tdist_P](#), 356
- [tdist_pdf](#), 354
- [tdist_Pinv](#), 356
- [tdist_Q](#), 356
- [tdist_Qinv](#), 356
- [throw](#), 140
- [thunkp](#), 135
- [time](#), 163
- [timezone](#), 164

- [tmpfile](#), 167
- [transpose \(matrix\)](#), 130
- [trap](#), 162
- [true](#), 113
- [trunc](#), 135
- [trunc \(rational\)](#), 384
- [trunc_multiple](#), 386
- [tss](#), 349
- [tss_m](#), 349
- [tuple](#), 118
- [tuple \(matrix\)](#), 128
- [tuple \(string\)](#), 123
- [tuple indexing](#), 118
- [tuple size](#), 116
- [tuple slicing](#), 118
- [tuplep](#), 135
- [tuples](#), 116
- [tzname](#), 164
- [tzset](#), 164

U

- [ubyte](#), 134
- [ugaussian_P](#), 354
- [ugaussian_pdf](#), 353
- [ugaussian_Pinv](#), 354
- [ugaussian_Q](#), 354
- [ugaussian_Qinv](#), 354
- [ugaussian_tail_pdf](#), 353
- [uint](#), 134
- [uint64](#), 134
- [ulong](#), 134
- [uncurry](#), 115
- [uncurry3](#), 115
- [unlink](#), 463
- [unref](#), 146
- [unset_node_attr](#), 462
- [unzip](#), 121
- [unzip3](#), 121
- [update \(array\)](#), 151
- [update \(dictionary\)](#), 156
- [update \(record\)](#), 132
- [update2 \(array\)](#), 151
- [upper](#), 388
- [ushort](#), 134

V

- [val](#), 116

[val \(blob\)](#), 139
[val \(string\)](#), 137
[val_eng_approx_epsilon](#), 412
[val_eng_approx_max_den](#), 413
[val_eng_interval_epsilon](#), 413
[val_eng_interval_max_den](#), 413
[val_mixed](#), 405
[val_real_eng](#), 411
[val_real_recur](#), 407
[val_vulgar](#), 405
[vals \(dictionary\)](#), 155
[vals \(record\)](#), 132
[variance](#), 349
[variance_m](#), 349
[variance_with_fixed_mean](#), 349
[varp](#), 135
[vector](#), 127
[vectorp](#), 128
[vectorseq](#), 127
[void](#), 115

W

[wabsdev](#), 350
[wabsdev_m](#), 350
[wait](#), 176
[waitpid](#), 176
[weibull_P](#), 356
[weibull_pdf](#), 354
[weibull_Pinv](#), 356
[weibull_Q](#), 356
[weibull_Qinv](#), 356
[within](#), 389
[without](#), 390
[wkurtosis](#), 350
[wkurtosis_m_sd](#), 350
[wlinear](#), 347
[wmean](#), 350
[wmul](#), 347
[wsd](#), 350
[wsd_m](#), 350
[wsd_with_fixed_mean](#), 350
[wskew](#), 350
[wskew_m_sd](#), 350
[wtss](#), 350
[wtss_m](#), 350
[wvariance](#), 350
[wvariance_m](#), 350

[wvariance_with_fixed_mean](#), 350

Z

[zeros](#), 344
[zip](#), 121
[zip3](#), 121
[zipwith](#), 121
[zipwith3](#), 121